

SQL: Queries, Constraints, Triggers

Chapter 5

ECS 165A – Winter 2026



Mohammad Sadoghi
Exploratory Systems Lab
Department of Computer Science

UCDAVIS
UNIVERSITY OF CALIFORNIA



Apache
ResilientDB
Incubating



Example Instances

- ❖ We will use these instances of the Sailors and Reserves relations in our examples.
- ❖ If the key for the Reserves relation contained only the attributes *sid* and *bid*, how would the semantics differ?

<i>R1</i>	<u>sid</u>	<u>bid</u>	<u>day</u>
	22	101	10/10/96
	58	103	11/12/96

<i>S1</i>	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

<i>S2</i>	<u>sid</u>	sname	rating	age
	28	yuppy	9	35.0
	31	lubber	8	55.5
	44	guppy	5	35.0
	58	rusty	10	35.0

Basic SQL Query

```
SELECT [DISTINCT] target-list
       FROM relation-list
      WHERE qualification
```

- ❖ *relation-list* A list of relation names (possibly with a *range-variable* after each name).
- ❖ *target-list* A list of attributes of relations in *relation-list*
- ❖ *qualification* Comparisons (Attr *op* const OR Attr1 *op* Attr2, where *op* is one of <, >, =, ≤, ≥, ≠) combined using AND, OR and NOT.
- ❖ **DISTINCT** is an optional keyword indicating that the answer should not contain duplicates. Default is that duplicates are not eliminated!

Conceptual Evaluation Strategy

- ❖ Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they fail *qualifications*.
 - Delete attributes that are not in *target-list*.
 - If **DISTINCT** is specified, eliminate duplicate rows.
- ❖ This strategy is probably the least efficient way to compute a query! An optimizer will find more efficient strategies to compute *the same answers*.

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

sid	bid	day
22	101	10/10/96
58	103	11/12/96

R1

S1

Example of Conceptual Evaluation

```
SELECT S.sname
      FROM Sailors S, Reserves R
     WHERE S.sid=R.sid AND R.bid=103
```

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

A Note on Range Variables

- ❖ Really needed only if the same relation appears twice in the FROM clause. The previous query can also be written as:

SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND bid=103

OR

SELECT sname
FROM Sailors, Reserves
WHERE Sailors.sid=Reserves.sid
AND bid=103

*It is good style,
however, to use
range variables
always!*

Expressions and Strings

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
      FROM Sailors S  
     WHERE S.sname LIKE 'B_%B'
```

- ❖ Illustrates use of arithmetic expressions and string pattern matching: *Find triples (of ages of sailors and two fields defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
- ❖ **AS** and **=** are two ways to name fields in result.
- ❖ **LIKE** is used for string matching. **'_'** stands for any one character and **'%'** stands for 0 or more arbitrary characters.

Find sid's of sailors who've reserved a red or a green boat

- ❖ If we replace **OR** by **AND** in the first version, what do we get?
- ❖ **UNION**: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).
- ❖ Also available: **EXCEPT** (What do we get if we replace **UNION** by **EXCEPT**?)

```
SELECT S.sid  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid AND  
(B.color='red' OR B.color='green')
```

```
SELECT S.sid  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid AND  
B.color='red'  
UNION
```

```
SELECT S.sid  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid AND  
B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat

```
SELECT S.sid  
FROM Sailors S, Boats B1, Reserves R1,  
      Boats B2, Reserves R2  
WHERE S.sid=R1.sid AND R1.bid=B1.bid AND  
      S.sid=R2.sid AND R2.bid=B2.bid AND  
      (B1.color='red' AND B2.color='green')
```

SELECT S.sid Key field!

```
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid AND  
      B.color='red'
```

INTERSECT

```
SELECT S.sid  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid AND  
      B.color='green'
```

Nested Queries

Find names of sailors who've reserved boat #103:

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid IN (SELECT R.sid  
                 FROM Reserves R  
                 WHERE R.bid=103)
```

- ❖ A very powerful feature of SQL: a WHERE clause can itself contain an SQL query! (Actually, so can FROM and HAVING clauses.)
- ❖ To find sailors who've *not* reserved #103, use NOT IN.
- ❖ To understand semantics of nested queries, think of a *nested loops* evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

Nested Queries with Correlation

Find names of sailors who've reserved boat # 103:

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS (SELECT *  
              FROM Reserves R  
              WHERE R.bid=103 AND S.sid=R.sid)
```

- ❖ **EXISTS** is another set comparison operator, like **IN**.
- ❖ Illustrates why, in general, sub-query must be re-computed for each Sailors tuple.

More on Set-Comparison Operators

- ❖ We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- ❖ Also available: *op ANY, op ALL, IN >,<,=,≥,≤,≠*
- ❖ Find sailors whose rating is greater than that of some sailor called Horatio:

```
SELECT *
  FROM Sailors S
 WHERE S.rating > ANY (SELECT S2.rating
                           FROM Sailors S2
                          WHERE S2.sname='Horatio')
```

Division in SQL

Find sailors who've reserved all boats.

- ❖ Let's do it the hard way,
without EXCEPT:

```
SELECT S.sname  
      FROM Sailors S  
 WHERE NOT EXISTS (SELECT B.bid  
                      FROM Boats B
```

Sailors S such that ...

there is no boat B without ...

a Reserves tuple showing S reserved B

```
SELECT S.sname  
      FROM Sailors S  
 WHERE NOT EXISTS  
       ((SELECT B.bid  
          FROM Boats B)  
        EXCEPT  
       (SELECT R.bid  
          FROM Reserves R  
         WHERE R.sid=S.sid))
```

```
          WHERE NOT EXISTS (SELECT R.bid  
                           FROM Reserves R  
                           WHERE R.bid=B.bid  
                                 AND R.sid=S.sid))
```

Aggregate Operators

- ❖ Significant extension of relational algebra.

```
SELECT COUNT (*)  
      FROM Sailors S
```

```
SELECT AVG (S.age)  
      FROM Sailors S
```

```
WHERE S.rating=10
```

```
SELECT COUNT (DISTINCT S.rating)  
      FROM Sailors S
```

```
WHERE S.sname='Bob'
```

```
COUNT (*)  
COUNT ( [DISTINCT] A)  
SUM ( [DISTINCT] A)  
AVG ( [DISTINCT] A)  
MAX (A)  
MIN (A)
```

single column

```
SELECT *
```

```
      FROM S  
      WHERE column = (SELECT ... FROM R)
```

```
SELECT AVG ( DISTINCT S.age)  
      FROM Sailors S  
      WHERE S.rating=10
```

Find name and age of the oldest sailor(s)

- ❖ The first query is illegal!
(We'll look into the reason a bit later, when we discuss **GROUP BY**.)
- ❖ The third query is equivalent to the second query, and is allowed in the SQL/92 standard, but is not supported in some systems.

```
SELECT S.sname, MAX (S.age)  
FROM Sailors S
```

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
(SELECT MAX (S2.age)  
FROM Sailors S2)
```

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE (SELECT MAX (S2.age)  
FROM Sailors S2)  
= S.age
```

Motivation for Grouping

- ❖ So far, we've applied aggregate operators to all (qualifying) tuples. Sometimes, we want to apply them to each of several *groups* of tuples.
- ❖ Consider: *Find the age of the youngest sailor for each rating level.*
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

For $i = 1, 2, \dots, 10$:

```
SELECT MIN (S.age)
  FROM Sailors S
 WHERE S.rating = i
```

Queries With GROUP BY and HAVING

```
SELECT [DISTINCT] attribute-list, aggregate operations
        FROM relation-list
        WHERE qualification
GROUP BY grouping-list
HAVING group-qualification
```

- The attribute list (i) must be a subset of *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group. (A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)
- Terms with aggregate operations are of form $\text{MIN} (S.\text{age})$), for example

Conceptual Evaluation

- ❖ The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, `unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- ❖ The *group-qualification* is then applied to eliminate some groups. Expressions in *group-qualification* must have a *single value per group!*
 - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*. (SQL does not exploit primary key semantics here!)
- ❖ One answer tuple is generated per qualifying group.

*Find age of the youngest sailor with age ≥ 18 ,
for each rating with at least 2 such sailors*

```
SELECT S.rating, MIN(S.age)
      AS minage
  FROM Sailors S
 WHERE S.age >= 18
GROUP BY S.rating
 HAVING COUNT(*) > 1
```

Answer relation:

rating	minage
3	25.5
7	35.0
8	25.5

Sailors instance:

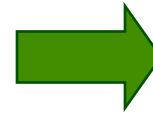
sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5

*Find age of the youngest sailor with age ≥ 18 ,
for each rating with at least 2 such sailors.*

rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
10	16.0
9	35.0
3	25.5
3	63.5
3	25.5



rating	age
1	33.0
3	25.5
3	63.5
3	25.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0



rating	minage
3	25.5
7	35.0
8	25.5

Null Values

- ❖ Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
 - SQL provides a special value *null* for such situations.
- ❖ The presence of *null* complicates many issues. E.g.:
 - Special operators needed to check if value is/is not *null*.
 - Is *rating>8* true or false when *rating* is equal to *null*? What about AND, OR and NOT connectives?
 - We need a 3-valued logic (true, false and *unknown*).
 - Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
 - New operators (in particular, *outer joins*) possible/needed.

Integrity Constraints (Review)

- ❖ An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts / deletes / updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- ❖ Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - *Domain constraints*: Field values must be of right type.
Always enforced.

General Constraints

- ❖ Useful when more general ICs than keys are involved.
- ❖ Can use queries to express constraint.
- ❖ Constraints can be named.

```
CREATE TABLE Sailors  
(sid INTEGER,  
 sname CHAR(10),  
 rating INTEGER,  
 age REAL,  
 PRIMARY KEY (sid),  
 CHECK ( rating >= 1 AND  
 rating <= 10 )  
  
CREATE TABLE Reserves  
( sname CHAR(10),  
 bid INTEGER,  
 day DATE,  
 PRIMARY KEY (bid,day),  
 CONSTRAINT noInterlakeRes  
 CHECK ('Interlake' <>  
 ( SELECT B.bname  
   FROM Boats B  
 WHERE B.bid=bid)))
```

Constraints Over Multiple Relations

- ❖ Awkward and wrong!
- ❖ If Sailors is empty, the number of Boats tuples can be anything!
- ❖ ASSERTION is the right solution; not associated with either table.

```
CREATE TABLE Sailors  
  ( sid INTEGER,  
    sname CHAR(10),  
    rating INTEGER,  
    age REAL,  
    PRIMARY KEY (sid),  
    CHECK  
      ( (SELECT COUNT (S.sid) FROM Sailors S)  
      + (SELECT COUNT (B.bid) FROM Boats B) < 100 )
```

Number of boats plus number of sailors is < 100

```
CREATE ASSERTION smallClub  
  CHECK  
    ( (SELECT COUNT (S.sid) FROM Sailors S)  
    + (SELECT COUNT (B.bid) FROM Boats B) < 100 )
```

Triggers

- ❖ Trigger: procedure that starts automatically if specified changes occur to the DBMS
- ❖ Three parts:
 - Event (activates the trigger)
 - Condition (tests whether the triggers should run)
 - Action (what happens if the trigger runs)

Triggers: Example (SQL:1999)

```
CREATE TRIGGER youngSailorUpdate
    AFTER INSERT ON SAILORS
    REFERENCING NEW TABLE NewSailors
    FOR EACH STATEMENT
        INSERT
            INTO YoungSailors(sid, name, age, rating)
            SELECT sid, name, age, rating
            FROM NewSailors N
            WHERE N.age <= 18
```

Summary

- ❖ SQL was an important factor in the early acceptance of the relational model; more natural than earlier, procedural query languages.
- ❖ Relationally complete; in fact, significantly more expressive power than relational algebra.
- ❖ Even queries that can be expressed in RA can often be expressed more naturally in SQL.
- ❖ Many alternative ways to write a query; optimizer should look for most efficient evaluation plan.
 - In practice, users need to be aware of how queries are optimized and evaluated for best results.

Summary (Contd.)

- ❖ NULL for unknown field values brings many complications
- ❖ SQL allows specification of rich integrity constraints
- ❖ Triggers respond to changes in the database