

Concurrency Control

Chapter 17

ECS 165A – Winter 2026



Mohammad Sadoghi
Exploratory Systems Lab
Department of Computer Science

UCDAVIS
UNIVERSITY OF CALIFORNIA



Apache
ResilientDB
Incubating



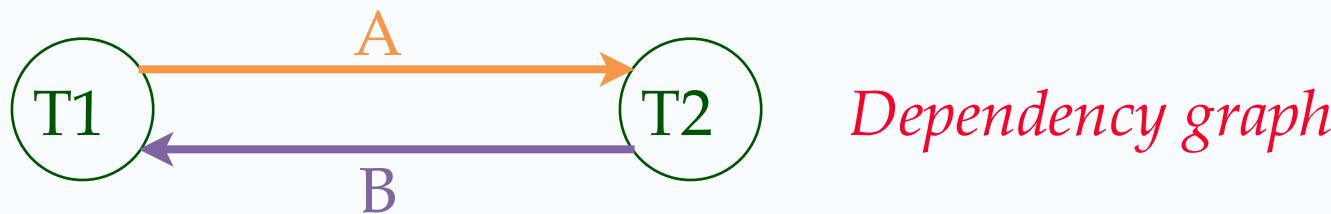
Conflict Serializable Schedules

- ❖ *Serial schedule:* Schedule that does not interleave the actions of different transactions.
- ❖ *Equivalent schedules:* For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- ❖ *Serializable schedule:* A schedule that is equivalent to some serial execution of the transactions.
- ❖ *Two schedules are conflict equivalent if:*
 - Involve the same actions of the same transactions
 - Every pair of conflicting actions is ordered the same way
- ❖ Schedule S is **conflict serializable** if S is conflict equivalent to some serial schedule

Example

- ❖ A schedule that is not conflict serializable:

T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	

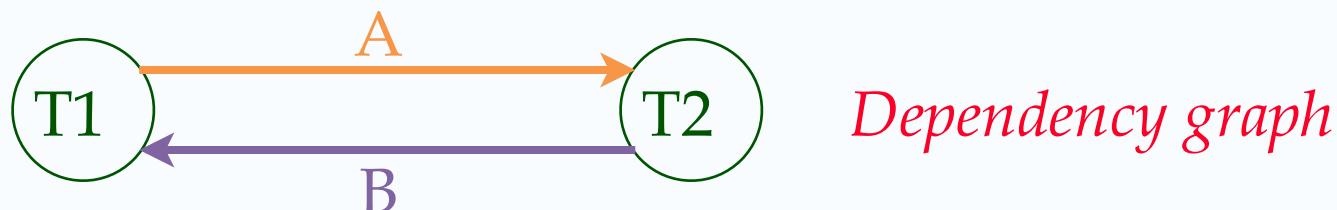


- ❖ The cycle in the graph reveals the problem. The output of T1 depends on T2, and vice-versa.

Dependency Graph

- ❖ Dependency graph: One node per Xact; edge from T_i to T_j if T_j reads/writes an object last written by T_i .
- ❖ Theorem: Schedule is conflict serializable if and only if its dependency graph is acyclic

T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	



Review: Strict 2PL

- ❖ *Strict Two-phase Locking (Strict 2PL) Protocol:*
 - Each Xact must obtain a **S (shared)** lock on object before reading, and an **X (exclusive)** lock on object before writing.
 - All locks held by a transaction are released when the transaction completes
 - If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object.
- ❖ Strict 2PL allows only schedules whose precedence graph is acyclic

Two-Phase Locking (2PL)

- ❖ Two-Phase Locking Protocol
 - Each Xact must obtain a S (*shared*) lock on object before reading, and an X (*exclusive*) lock on object before writing.
 - A transaction can not request additional locks once it releases any locks.
 - If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object.

View Serializability

- ❖ Schedules S1 and S2 are **view equivalent** if:
 - If Ti reads initial value of A in S1, then Ti also reads initial value of A in S2 (*initial values*)
 - If Ti reads value of A written by Tj in S1, then Ti also reads value of A written by Tj in S2 (*intermediate values*)
 - If Ti writes final value of A in S1, then Ti also writes final value of A in S2 (*final values*)

T1: R(A)	W(A)
----------	------

T2:	W(A)
-----	------

T3:	W(A)
-----	------

T1: R(A), W(A)

T2:	W(A)
-----	------

T3:	W(A)
-----	------

Lock Management

- ❖ Lock and unlock requests are handled by the lock manager
- ❖ Lock table entry:
 - Number of transactions currently holding a lock
 - Type of lock held (shared or exclusive)
 - Pointer to queue of lock requests
- ❖ Locking and unlocking have to be atomic operations
- ❖ Lock upgrade: transaction that holds a shared lock can be upgraded to hold an exclusive lock

Deadlocks

- ❖ Deadlock: Cycle of transactions waiting for locks to be released by each other.
- ❖ Two ways of dealing with deadlocks:
 - Deadlock prevention
 - Deadlock detection

Deadlock Prevention

- ❖ Assign priorities based on timestamps. Assume T_i wants a lock that T_j holds. Two policies are possible:
 - *Wait-Die*: If T_i has higher priority, T_i waits for T_j ; otherwise T_i aborts
 - *Wound-wait*: If T_i has higher priority, T_j aborts; otherwise T_i waits
- ❖ If a transaction re-starts, make sure it has its *original timestamp* (*why?*)

Deadlock Detection

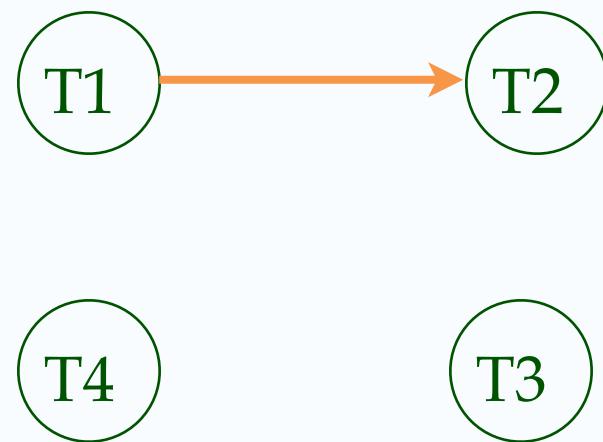
- ❖ Create a **waits-for graph**:
 - Nodes are transactions
 - There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock
- ❖ Periodically check for cycles in the waits-for graph

Deadlock Detection (Continued)

Example:

There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock

T1: S(A), R(A), S(B)
T2: X(B), W(B)
T3:
T4:



Deadlock Detection (Continued)

Example:

There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock

T1: S(A), R(A),

S(B)

T2:

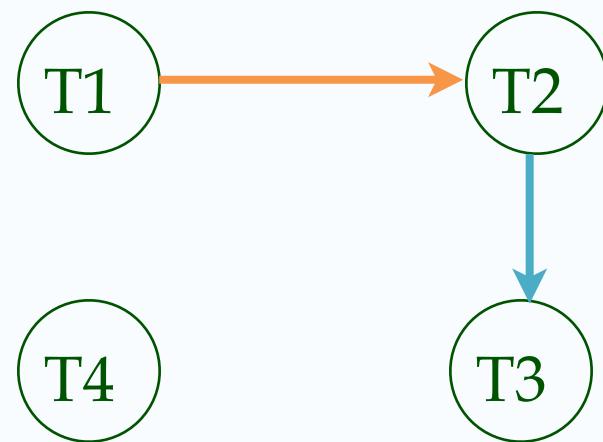
X(B),W(B)

X(C)

T3:

S(C), R(C)

T4:



Deadlock Detection (Continued)

Example:

There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock

T1: S(A), R(A),

S(B)

T2:

X(B), W(B)

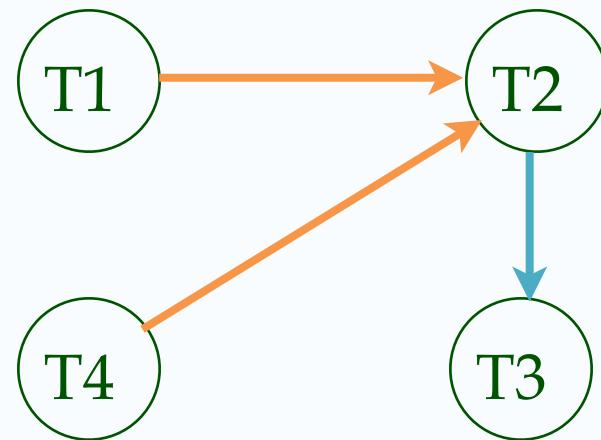
X(C)

T3:

S(C), R(C)

T4:

X(B)



Deadlock Detection (Continued)

Example:

There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock

T1: S(A), R(A),

S(B)

T2:

X(B), W(B)

X(C)

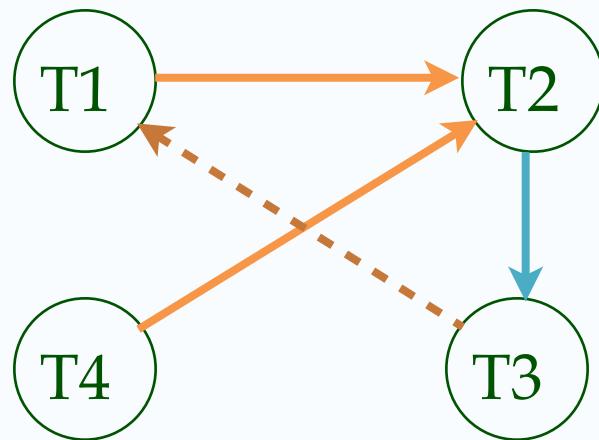
T3:

S(C), R(C)

X(A)

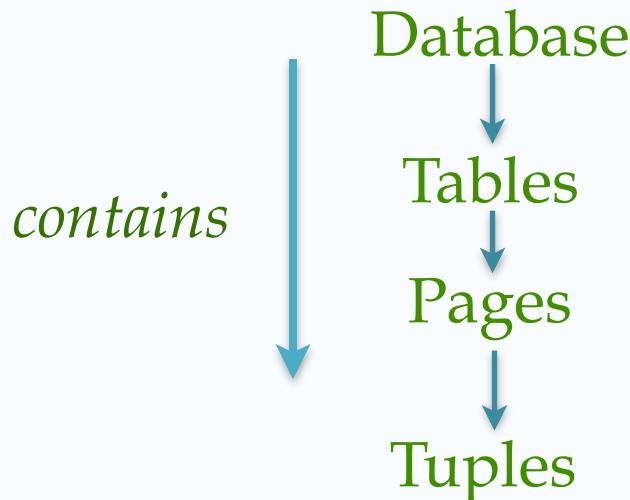
T4:

X(B)



Multiple-Granularity Locks

- ❖ Hard to decide what granularity to lock (*tuples vs. pages vs. tables*).
- ❖ Shouldn't have to decide!
- ❖ Data “containers” are nested:



Solution: New Lock Modes, Protocol

- ❖ Allow Xacts to lock at each level, but with a special protocol using new “intention” locks:
- Before locking an item, Xact must set “intention locks” on all its ancestors (i.e., *top-bottom*).
- For unlock, go from specific to general (i.e., *bottom-up*).
- **SIX mode:** Like S & IX at the same time.

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓			✓
X	✓				

Multiple Granularity Lock Protocol

- ❖ Each Xact starts from the *root* of the hierarchy.
- ❖ To get S or IS lock on a node, must hold IS or IX on parent node.
 - What if Xact holds SIX on parent? S on parent?
- ❖ To get X or IX or SIX on a node, must hold IX or SIX on parent node.
- ❖ Must release locks in *bottom-up* order.

Protocol is correct in that it is equivalent to directly setting locks at the leaf levels of the hierarchy.

Examples

- ❖ T1 scans R, and updates a few tuples:
 - T1 gets an SIX lock on R and occasionally upgrades to X on the tuples.
- ❖ T2 uses an index to read only part of R:
 - T2 gets an IS lock on R, and repeatedly gets an S lock on tuples of R.
- ❖ T3 reads all of R:
 - T3 gets an S lock on R.
 - OR, T3 could behave like T2; can use **lock escalation** to decide which.

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓			✓
X	✓				

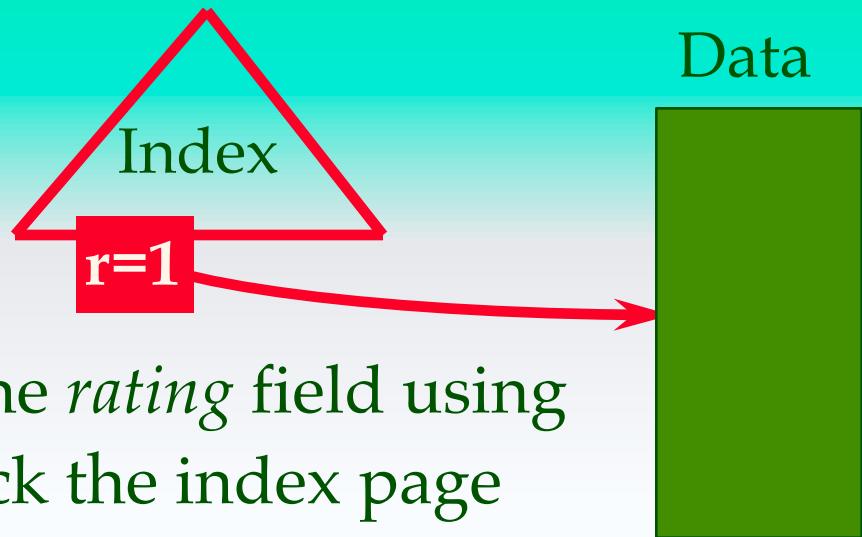
Dynamic Databases

- ❖ If we relax the assumption that the DB is a fixed collection of objects, even Strict 2PL will not assure serializability:
 - T1 locks all pages containing sailor records with $rating = 1$, and finds **oldest** sailor (say, $age = 71$).
 - Next, T2 inserts a new sailor; $rating = 1$, $age = 96$.
 - T2 also deletes oldest sailor with $rating = 2$ (and, say, $age = 80$), and commits.
 - T1 now locks all pages containing sailor records with $rating = 2$, and finds **oldest** (say, $age = 63$).
- ❖ There is no consistent DB state where T1 is “correct”!

The Problem

- ❖ T1 implicitly assumes that it has locked the set of all sailor records with *rating* = 1.
 - Assumption only holds if no sailor records are added while T1 is executing!
 - Need some mechanism to enforce this assumption.
(Index locking and predicate locking.)
- ❖ Example shows that conflict serializability guarantees serializability only if the set of objects is fixed!

Index Locking



- ❖ If there is a dense index on the *rating* field using Alternative (2), T1 should lock the index page containing the data entries with *rating* = 1.
 - If there are no records with *rating* = 1, T1 must lock the index page where such a data entry *would* be, if it existed!
- ❖ If there is no suitable index, T1 must lock all pages, and lock the file/table to prevent new pages from being added, to ensure that no new records with *rating* = 1 are added.

Predicate Locking

- ❖ Grant lock on all records that satisfy some logical predicate, e.g. $age > 2 * salary$.
- ❖ Index locking is a special case of predicate locking for which an index supports efficient implementation of the predicate lock.
 - What is the predicate in the sailor example?
- ❖ In general, predicate locking has a lot of locking overhead.

Transaction Support in SQL-92

- ❖ Each transaction has an access mode, a diagnostics size, and an isolation level.

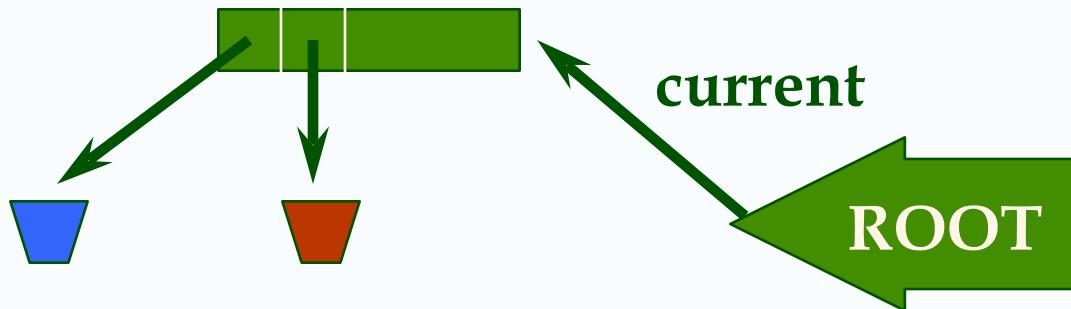
Isolation Level	Dirty Read	Unrepeatable Read	Phantom Problem
Read Uncommitted	Maybe	Maybe	Maybe
Read Committed	No	Maybe	Maybe
Repeatable Reads	No	No	Maybe
Serializable	No	No	No

Optimistic CC (Kung-Robinson)

- ❖ Locking is a conservative (pessimistic) approach in which conflicts are prevented. Disadvantages:
 - Lock management overhead.
 - Deadlock detection / resolution.
 - Lock contention for heavily used objects.
- ❖ If conflicts are rare, we might be able to gain concurrency by not locking, and instead checking for conflicts before Xacts commit.

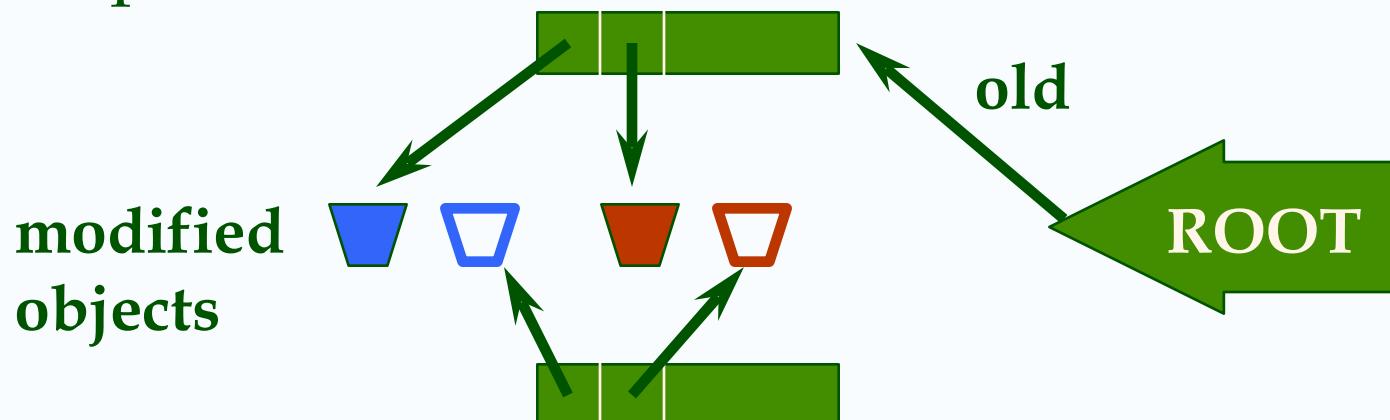
Kung-Robinson Model

- ❖ Xacts have three phases:
 - **READ:** Xacts read from the database, but make changes to private copies of objects.
 - **VALIDATE:** Check for conflicts.
 - **WRITE:** Make local copies of changes public.



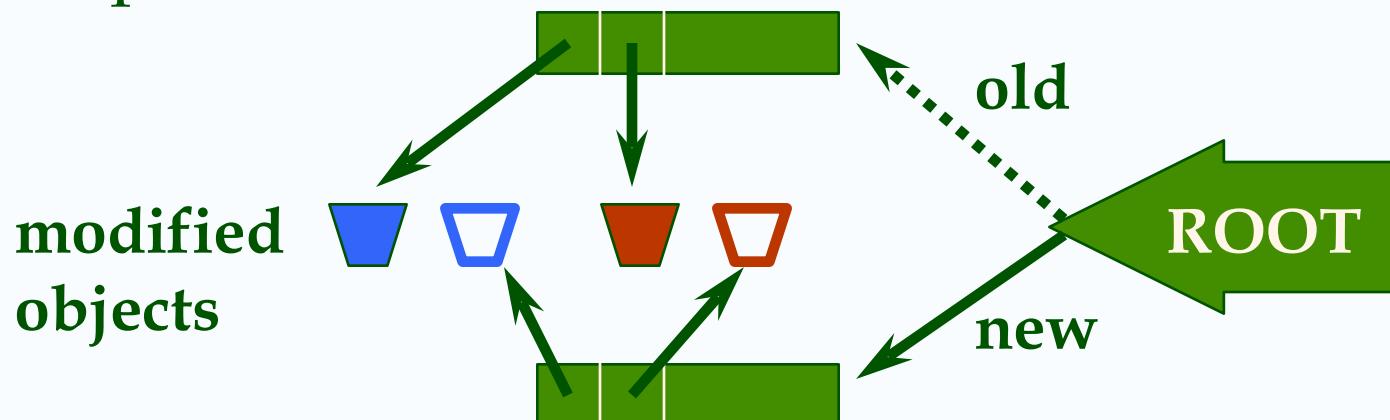
Kung-Robinson Model

- ❖ Xacts have three phases:
 - **READ**: Xacts read from the database, but make changes to private copies of objects.
 - **VALIDATE**: Check for conflicts.
 - **WRITE**: Make local copies of changes public.



Kung-Robinson Model

- ❖ Xacts have three phases:
 - **READ**: Xacts read from the database, but make changes to private copies of objects.
 - **VALIDATE**: Check for conflicts.
 - **WRITE**: Make local copies of changes public.



Validation

- ❖ Test conditions that are **sufficient** to ensure that no conflict occurred.
- ❖ Each Xact is assigned a numeric id.
 - Just use a **timestamp**.
- ❖ Xact ids assigned at end of READ phase, just before validation begins. (Why then?)
- ❖ **ReadSet(T_i):** Set of objects read by Xact T_i .
- ❖ **WriteSet(T_i):** Set of objects modified by T_i .

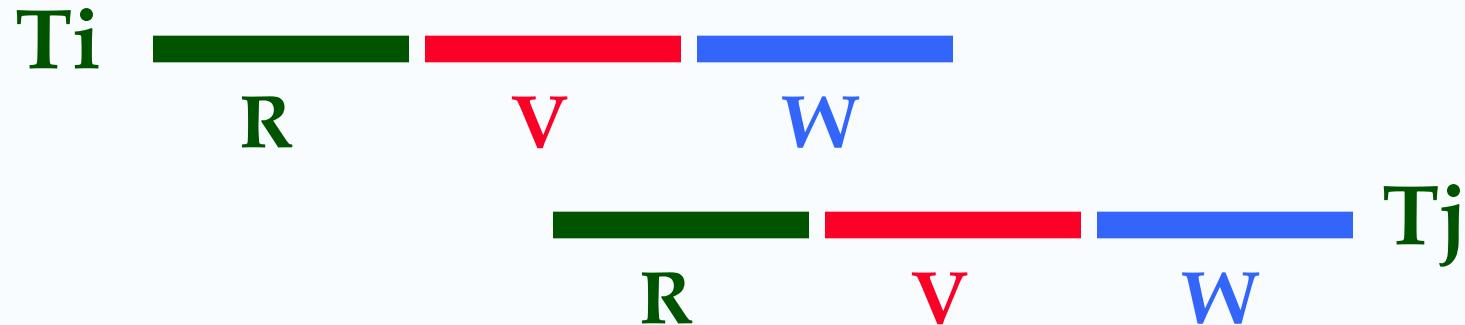
Test 1

- ❖ For all i and j such that $T_i < T_j$, check that T_i completes before T_j begins.



Test 2

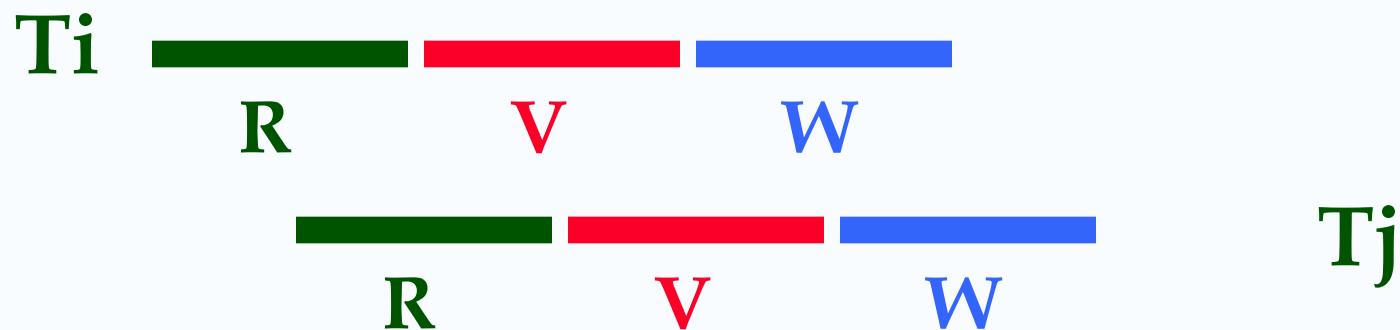
- ❖ For all i and j such that $T_i < T_j$, check that:
 - T_i completes before T_j begins its Write phase +
 - $\text{WriteSet}(T_i) \cap \text{ReadSet}(T_j)$ is empty.



Does T_j read dirty data? Does T_i overwrite T_j 's writes?

Test 3

- ❖ For all i and j such that $T_i < T_j$, check that:
 - T_i completes Read phase before T_j does +
 - $\text{WriteSet}(T_i) \cap \text{ReadSet}(T_j)$ is empty +
 - $\text{WriteSet}(T_i) \cap \text{WriteSet}(T_j)$ is empty.



Does T_j read dirty data? Does T_i overwrite T_j 's writes?

Comments on Serial Validation

- ❖ Assignment of Xact id, validation, and the Write phase are inside a **critical section!**
 - I.e., Nothing else goes on concurrently.
 - If Write phase is long, major drawback.
- ❖ Optimization for Read-only Xacts:
 - Don't need critical section (because there is no Write phase).

Overheads in Optimistic CC

- ❖ Must record read / write activity in ReadSet and WriteSet per Xact.
 - Must create and destroy these sets as needed.
- ❖ Must check for conflicts during validation, and must make validated writes ``global''.
 - Critical section can reduce concurrency.
 - Scheme for making writes global can reduce clustering of objects.
- ❖ Optimistic CC restarts Xacts that fail validation.
 - Work done so far is wasted; requires clean-up.

``Optimistic'' 2PL (analogous to 2VCC)

- ❖ If desired, we can do the following:
 - Set S locks as usual.
 - Make changes to private copies of objects.
 - Obtain all X locks at end of Xact, make writes global, then release all locks.
- ❖ In contrast to Optimistic CC as in Kung-Robinson, this scheme results in Xacts being blocked, waiting for locks.
 - However, no validation phase, no restarts (modulo deadlocks).

Timestamp CC

- ❖ **Idea:** Give each object a read-timestamp (RTS) and a write-timestamp (WTS), give each Xact a timestamp (TS) when it begins:
 - If action a_i of Xact T_i conflicts with action a_j of Xact T_j , and $TS(T_i) < TS(T_j)$, then a_i must occur before a_j . Otherwise, restart violating Xact.

When Xact T wants to Read Object O

- ❖ If $TS(T) < WTS(O)$, this violates timestamp order of T w.r.t. writer of O.
 - So, abort T and **restart it with a new larger TS**.
(If restarted with same TS, T will fail again! Contrast use of timestamps in 2PL for deadlock prevention.)
- ❖ If $TS(T) > WTS(O)$:
 - Allow T to read O.
 - Reset RTS(O) to $\max(RTS(O), TS(T))$
- ❖ Change to RTS(O) on reads must be written to disk! This and restarts represent overheads.

When Xact T wants to *Write* Object O

- ❖ If $TS(T) < RTS(O)$, this violates timestamp order of T w.r.t. writer of O; abort and restart T.
- ❖ If $TS(T) < WTS(O)$, violates timestamp order of T w.r.t. writer of O.
 - *Thomas Write Rule:* We can safely ignore such outdated writes; need not restart T! (T's write is effectively followed by another write, with no intervening reads.) Allows some serializable but non conflict serializable schedules:
- ❖ Else, allow T to write O.

T1	T2
R(A)	W(A) Commit
W(A) Commit	

Timestamp CC and Recoverability

- Unfortunately, unrecoverable schedules are allowed:

Read is aborted if $TS(T) < WTS(O)$

Write is aborted if $TS(T) < RTS(O)$ or $TS(T) < WTS(O)$

T1	T2
W(A)	
	R(A)
	W(B)
	Commit

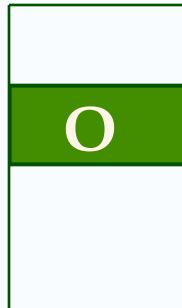
- ❖ Timestamp CC can be modified to allow only recoverable schedules (*any similarity to 2VCC?*):
 - **Buffer all writes** until writer commits (but update $WTS(O)$ when the write is **allowed**.)
 - **Block readers** T (where $TS(T) > WTS(O)$) until writer of O commits.
- ❖ Similar to writers holding X locks until commit, but still not quite 2PL.

Multiversion Timestamp CC

(Any Similarity to L-Store?)

Idea: Let writers make a “new” copy while readers use an appropriate “old” copy:

**MAIN
SEGMENT
(Current
versions of
DB objects)**

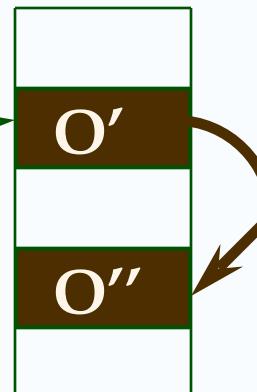


Multiversion Timestamp CC

(Any Similarity to L-Store?)

Idea: Let writers make a “new” copy while readers use an appropriate “old” copy:

**MAIN
SEGMENT
(Current
versions of
DB objects)**

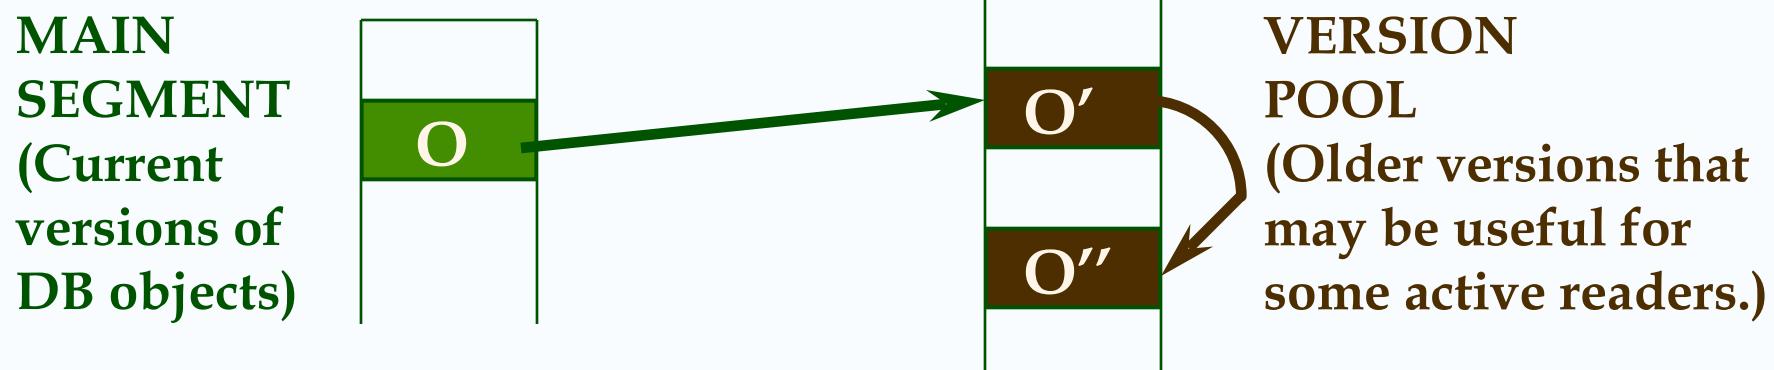


**VERSION
POOL
(Older versions that
may be useful for
some active readers.)**

Multiversion Timestamp CC

(Any Similarity to L-Store?)

Idea: Let writers make a “new” copy while readers use an appropriate “old” copy:



Readers are always allowed to proceed.
But may be blocked until writer commits.

Multiversion CC (Contd.)

- ❖ Each version of an object has its writer's TS as its **WTS**, and the TS of the Xact that most recently read this version as its **RTS**.
- ❖ Versions are chained backward; we can discard versions that are “too old to be of interest”.
- ❖ Each Xact is classified as **Reader** or **Writer**.
 - Writer *may* write some object; Reader never will.
 - Xact declares whether it is a Reader when it begins.

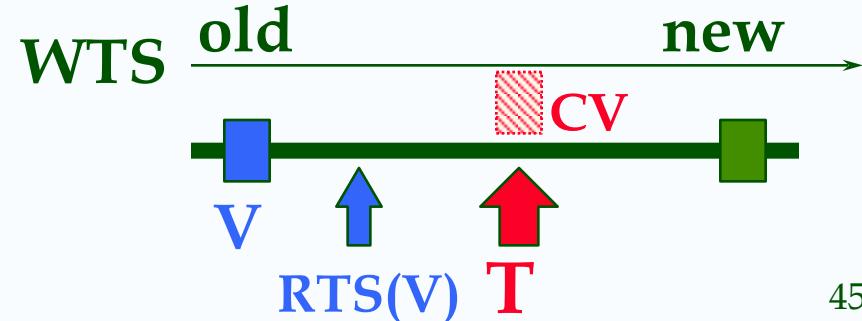
Reader Xact



- ❖ For each object to be read:
 - Finds **newest version** with $\text{WTS} < \text{TS}(T)$. (Starts with current version in the main segment and chains backward through earlier versions.)
- ❖ Assuming that some version of every object exists from the beginning of time, **Reader Xacts are never restarted**.
 - However, might block until writer of the appropriate version commits.

Writer Xact

- ❖ To read an object, follows reader protocol.
- ❖ To write an object:
 - Finds **newest version V** s.t. $WTS < TS(T)$.
 - If $RTS(V) < TS(T)$,
 - T makes a copy **CV** of V, with a pointer to V, with $WTS(CV) = TS(T)$, $RTS(CV) = TS(T)$.
 - Write is buffered until T commits; other Xacts can see TS values but can't read version **CV**.
 - **Else**, reject write.



Summary

- ❖ There are several lock-based concurrency control schemes (Strict 2PL, 2PL). Conflicts between transactions can be detected in the dependency graph
- ❖ The lock manager keeps track of the locks issued. Deadlocks can either be prevented or detected.
- ❖ Naïve locking strategies may have the phantom problem

Summary (Contd.)

- ❖ Index locking is common, and affects performance significantly.
 - Needed when accessing records via index.
 - Needed for **locking logical sets of records** (index locking/predicate locking).
- ❖ In practice, better techniques now known; do record-level, rather than page-level locking.

Summary (Contd.)

- ❖ Multiple granularity locking reduces the overhead involved in setting locks for nested collections of objects (e.g., a file of pages); should not be confused with tree index locking!
- ❖ Optimistic CC aims to minimize CC overheads in an ``optimistic'' environment where reads are common and writes are rare.
- ❖ Optimistic CC has its own overheads however; most real systems use locking.
- ❖ SQL-92 provides different isolation levels that control the degree of concurrency

Summary (Contd.)

- ❖ Timestamp CC is another alternative to 2PL; allows some serializable schedules that 2PL does not (although converse is also true).
- ❖ Ensuring recoverability with Timestamp CC requires ability to block Xacts, which is similar to locking.
- ❖ Multiversion Timestamp CC is a variant which ensures that read-only Xacts are never restarted; they can always read a suitable older version. Additional overhead of version maintenance.