

# *Storage & Indexing in Modern Databases*

**ECS 165A – Winter 2026**

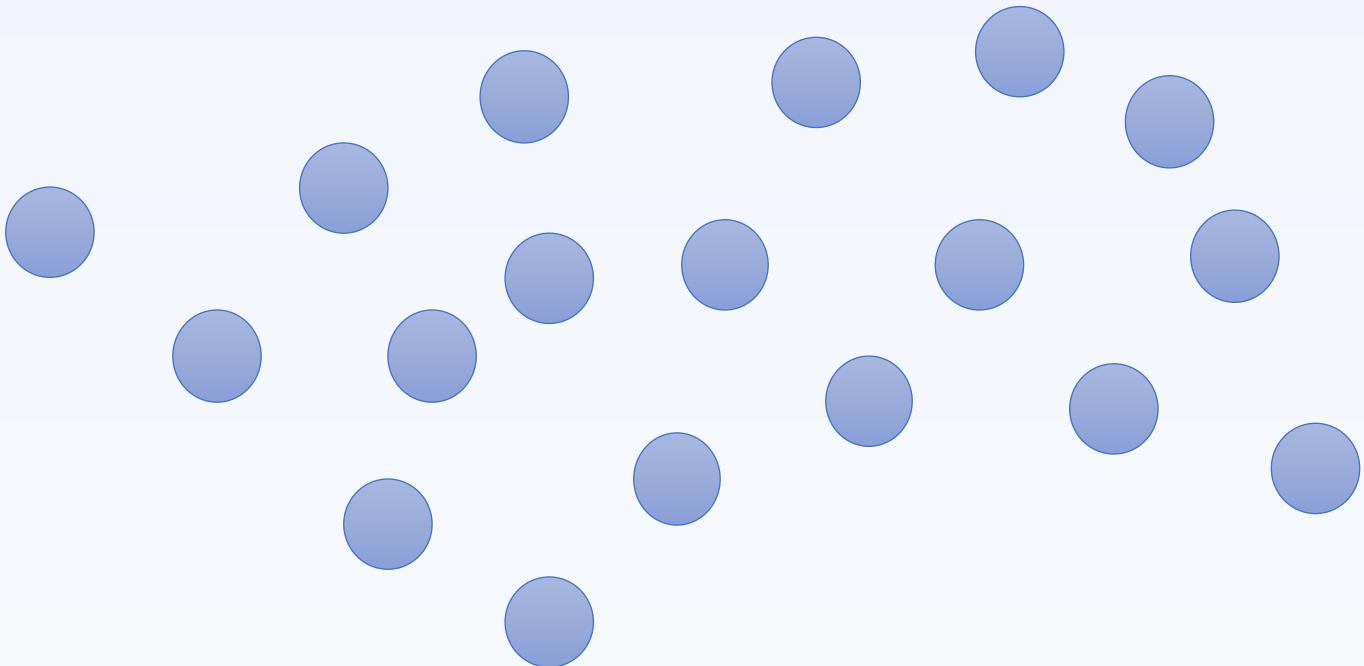


**Mohammad Sadoghi**  
Exploratory Systems Lab  
Department of Computer Science

**UCDAVIS**  
UNIVERSITY OF CALIFORNIA

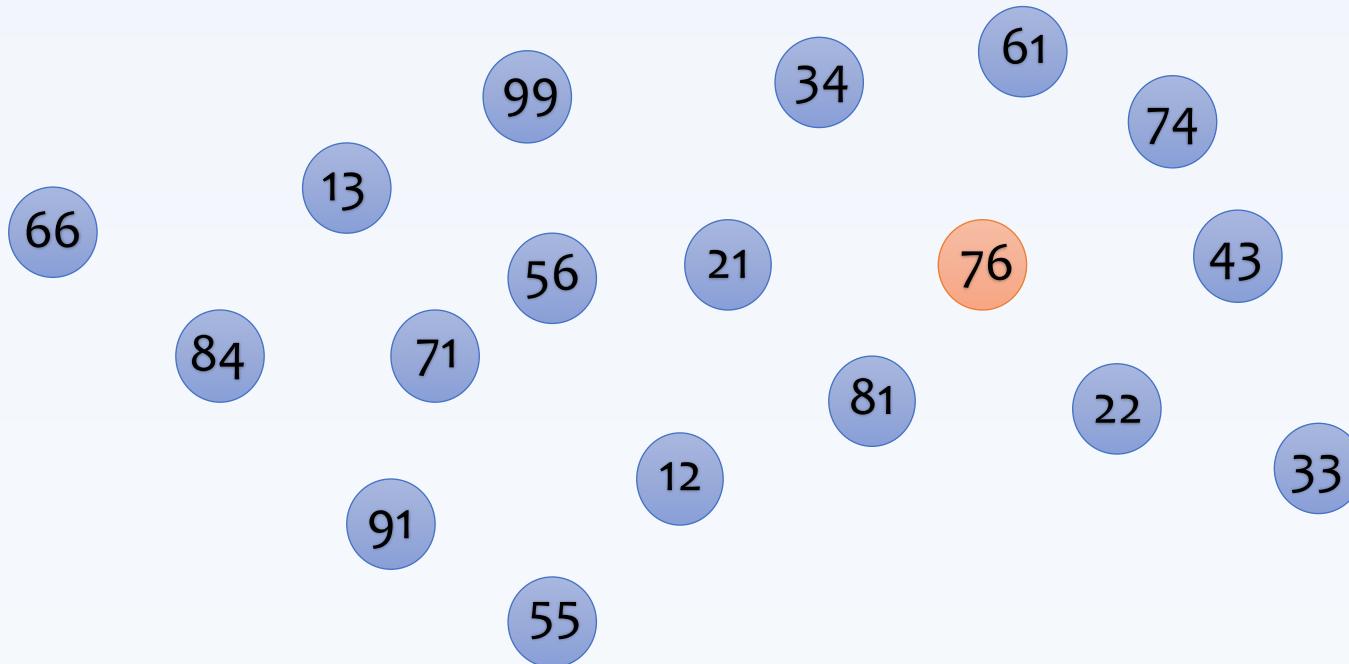


**How to quickly search for the desired information?**

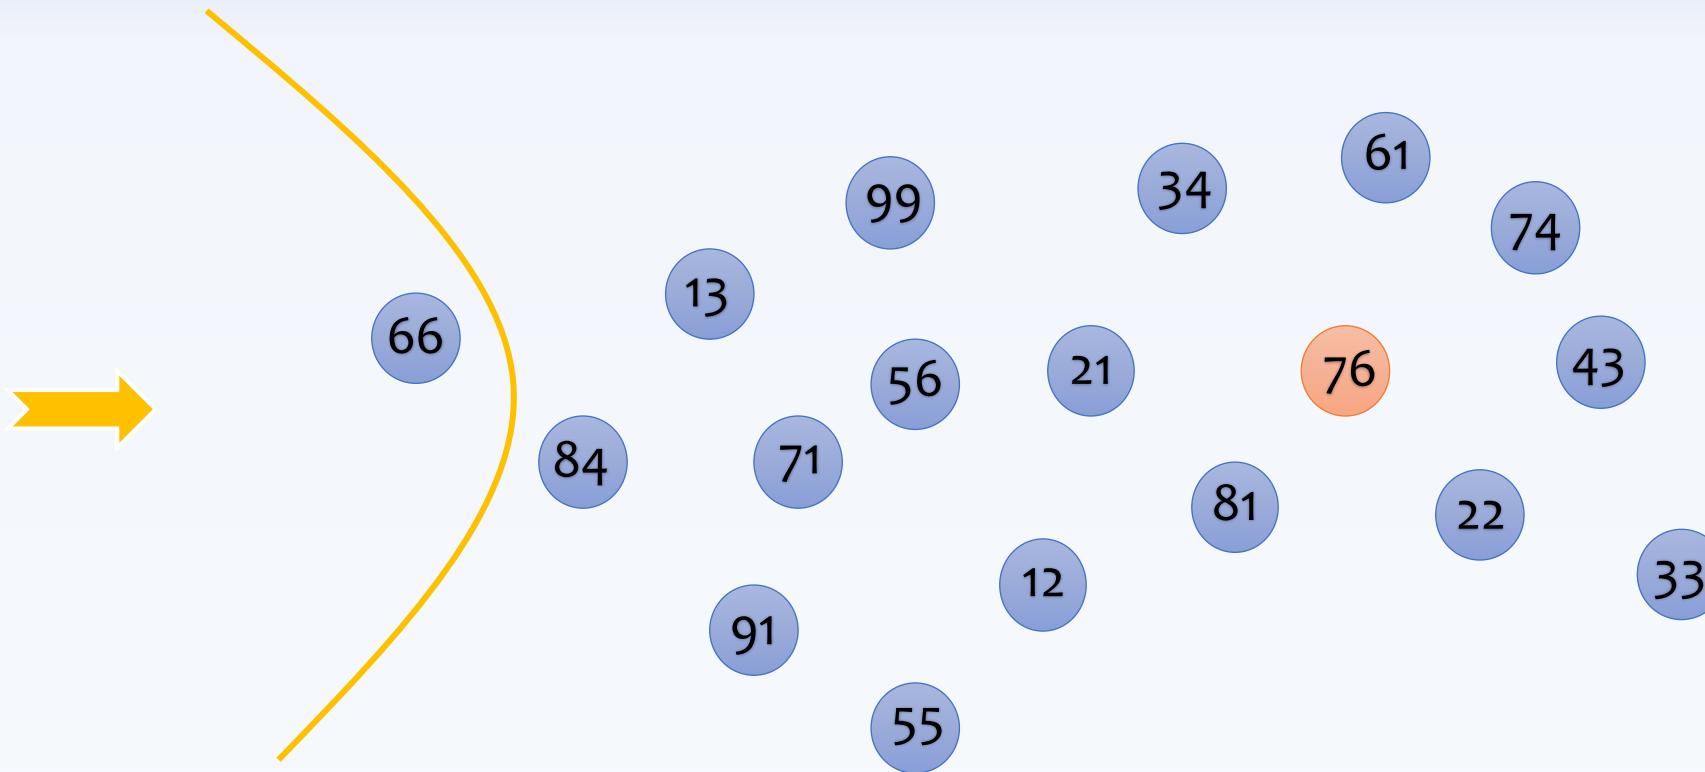




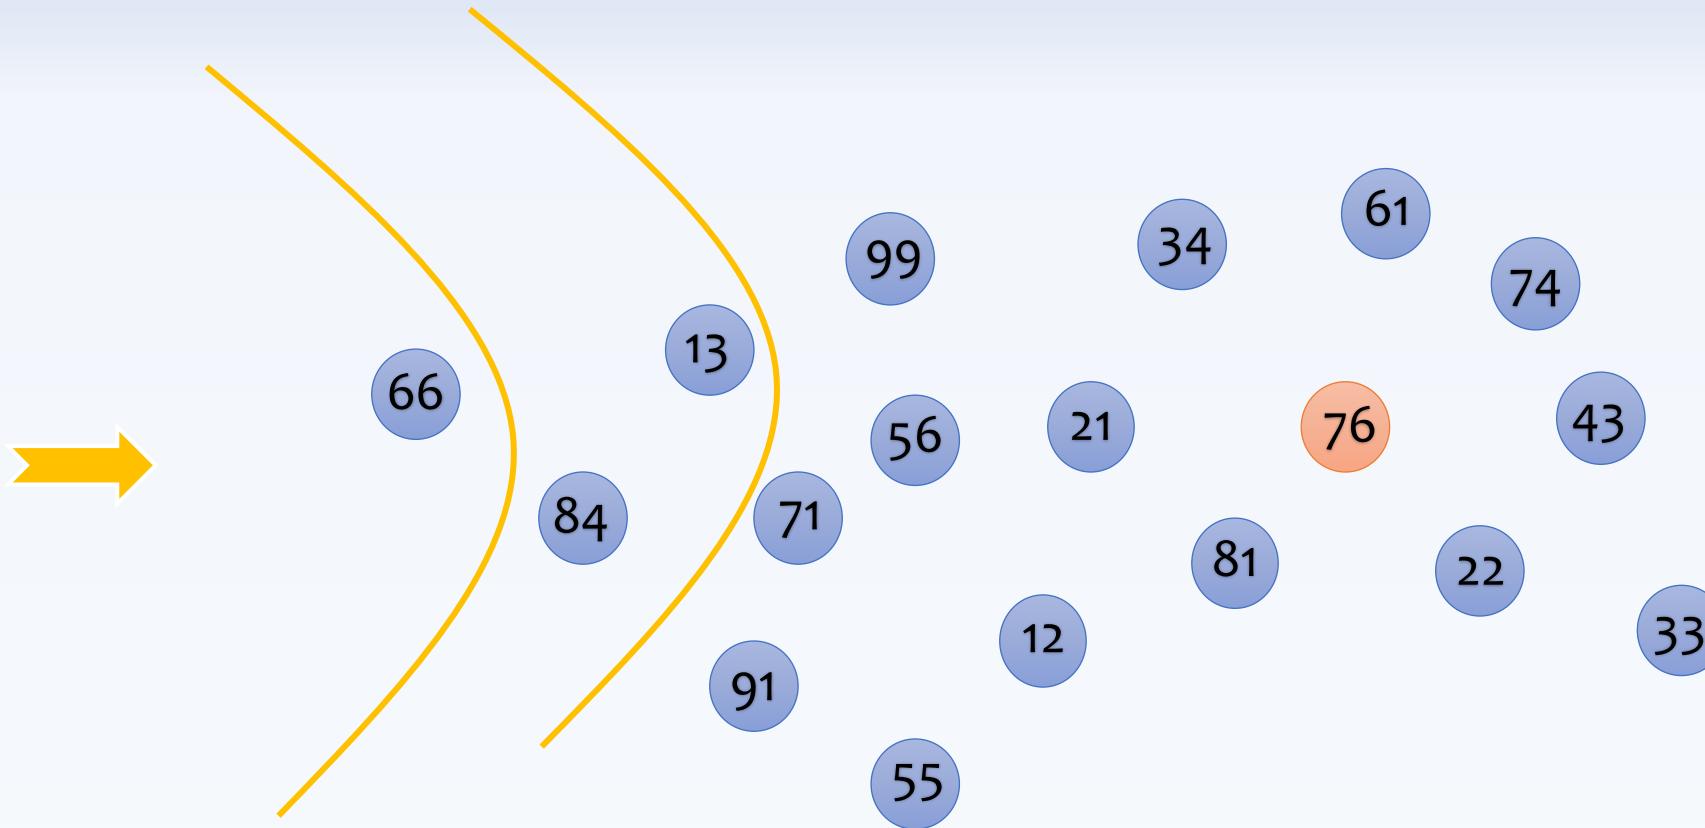
Searching for 76



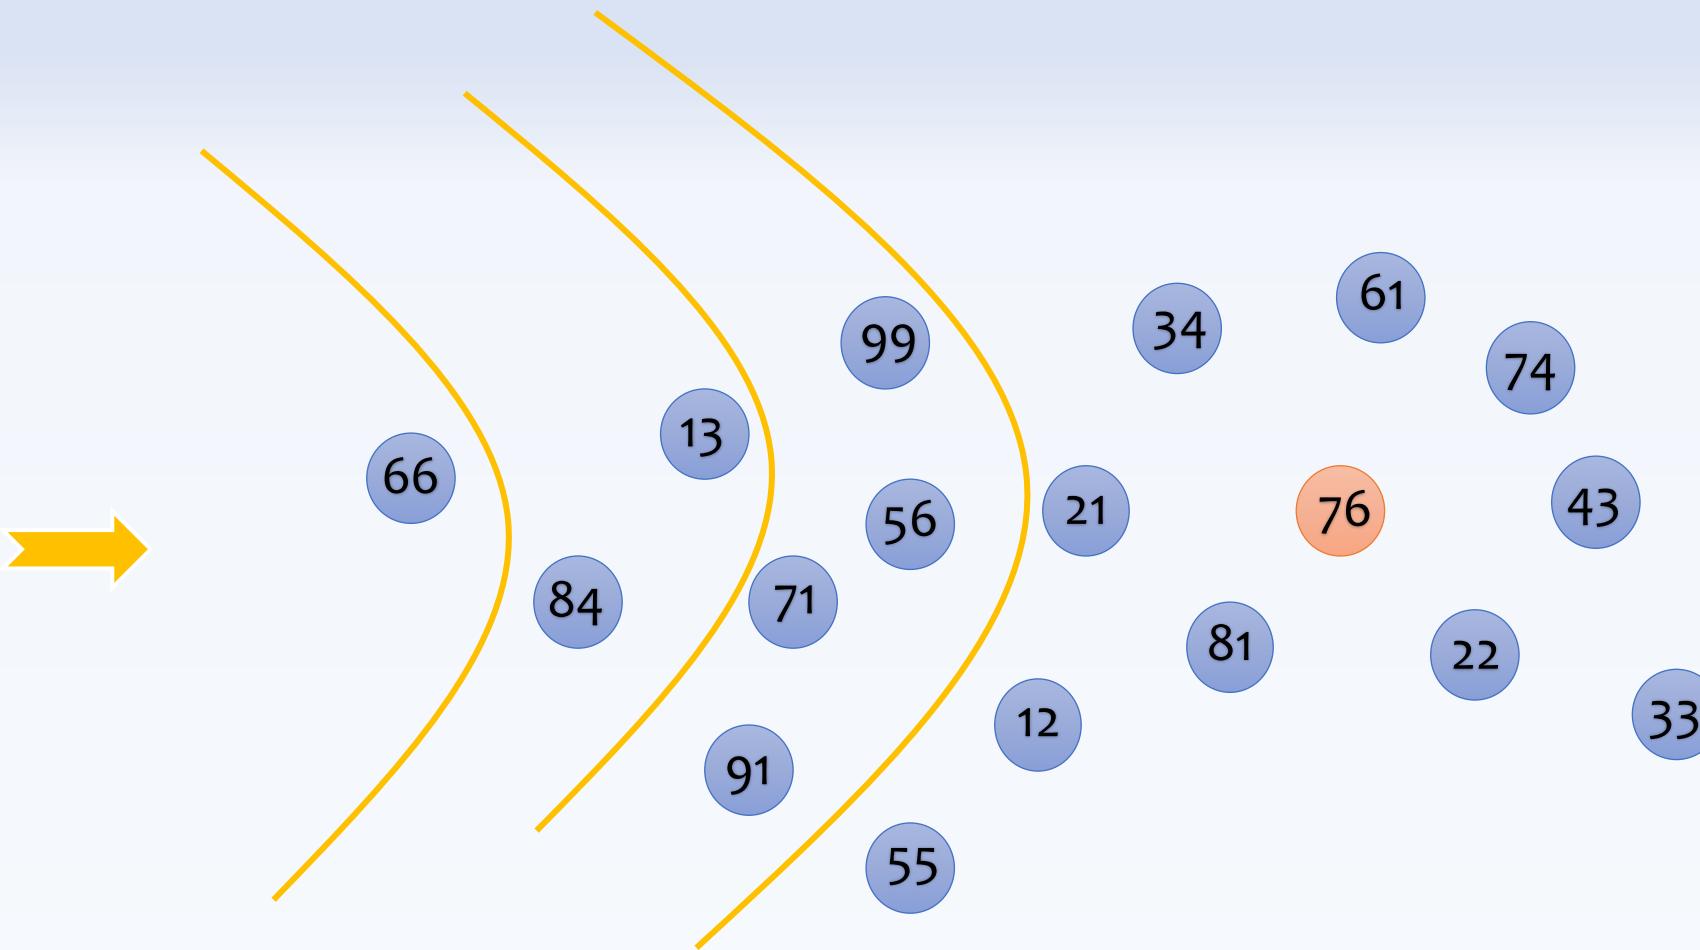
# Searching for 76



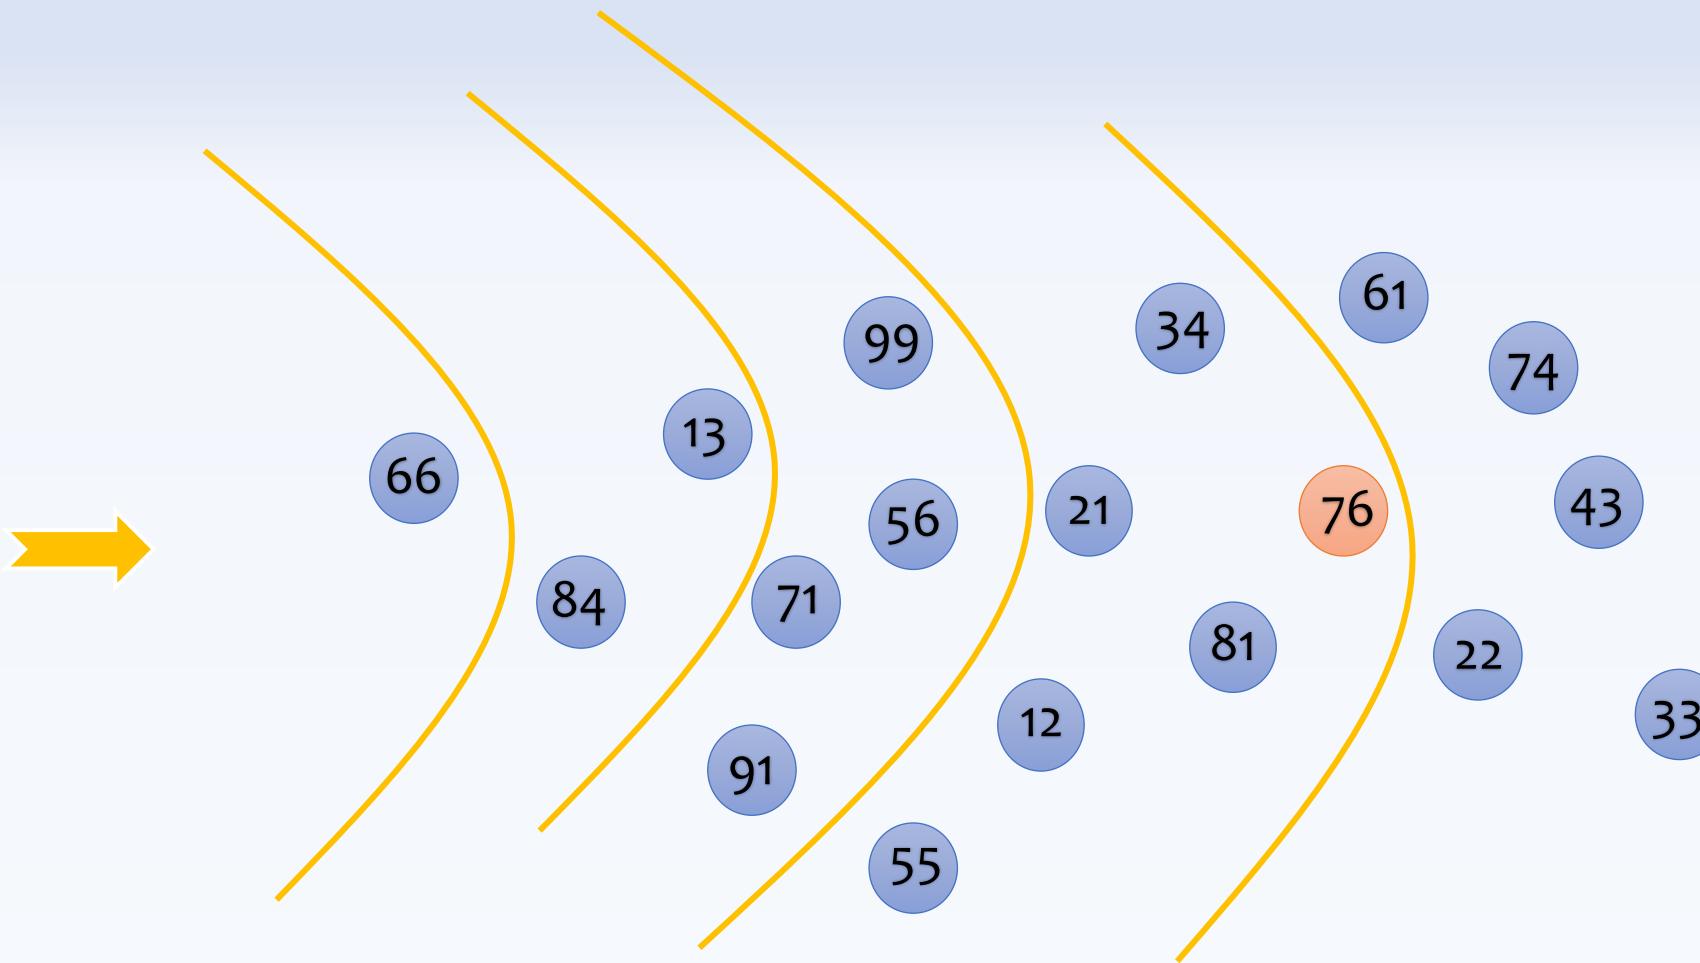
# Searching for 76



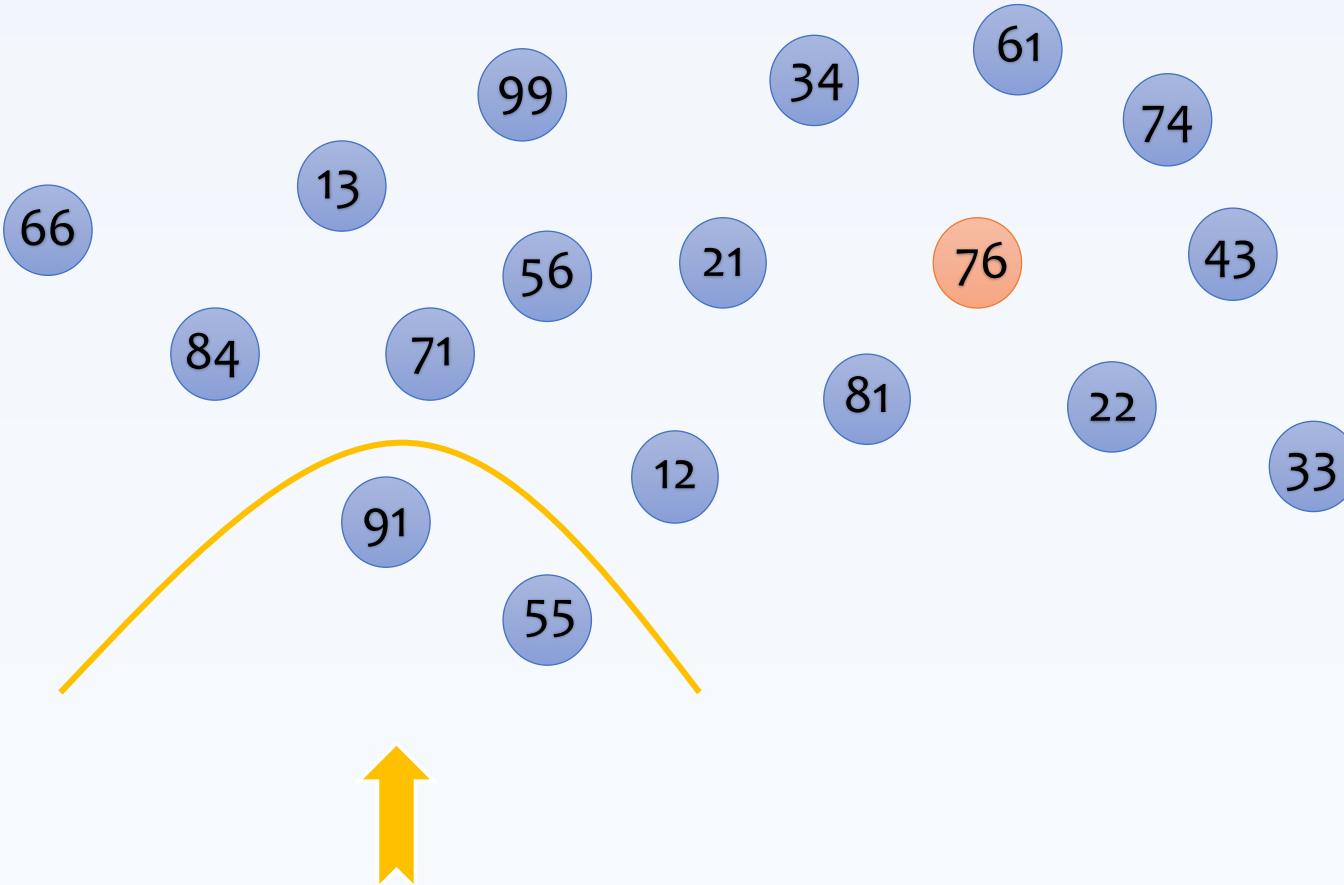
# Searching for 76



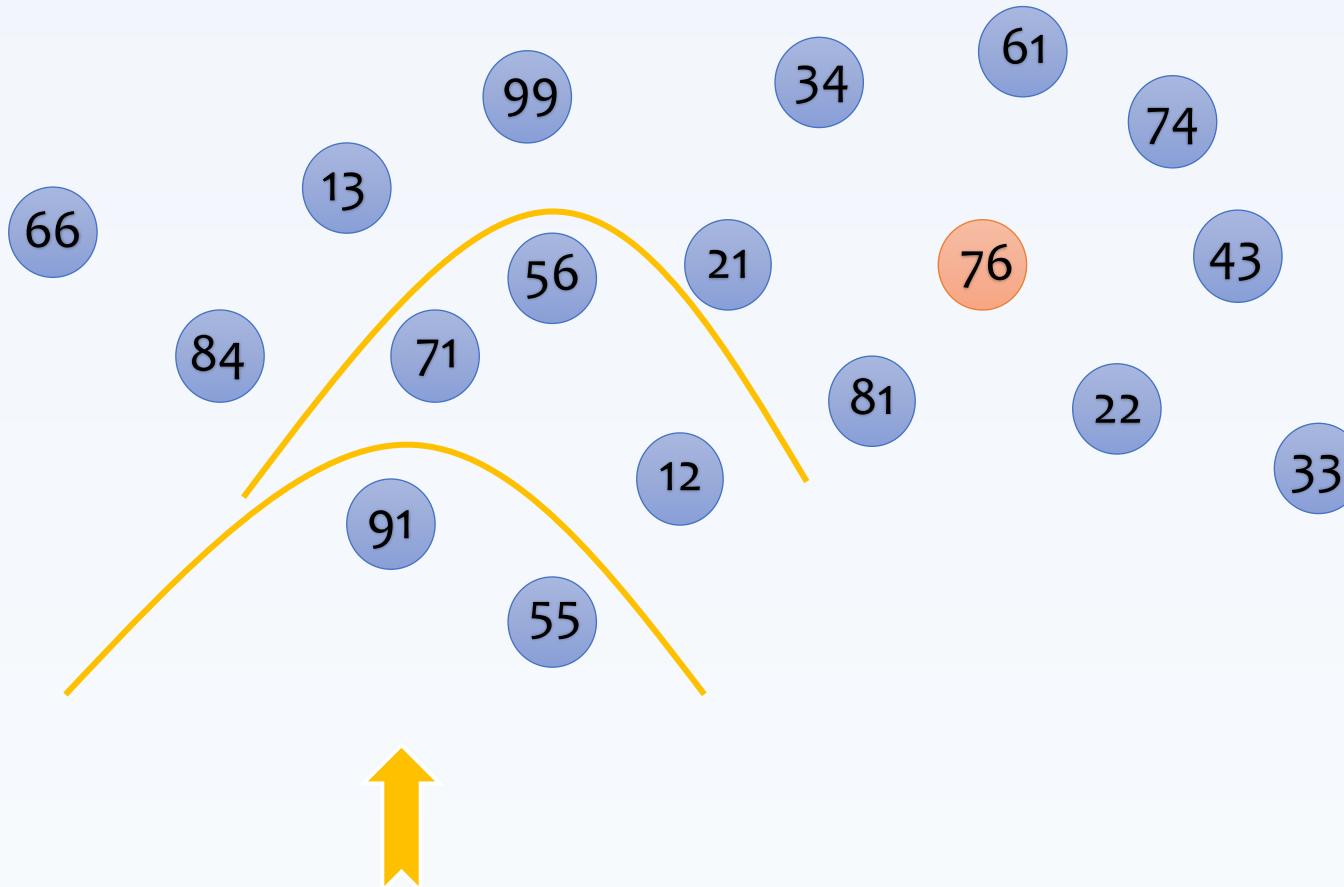
Searching for 76



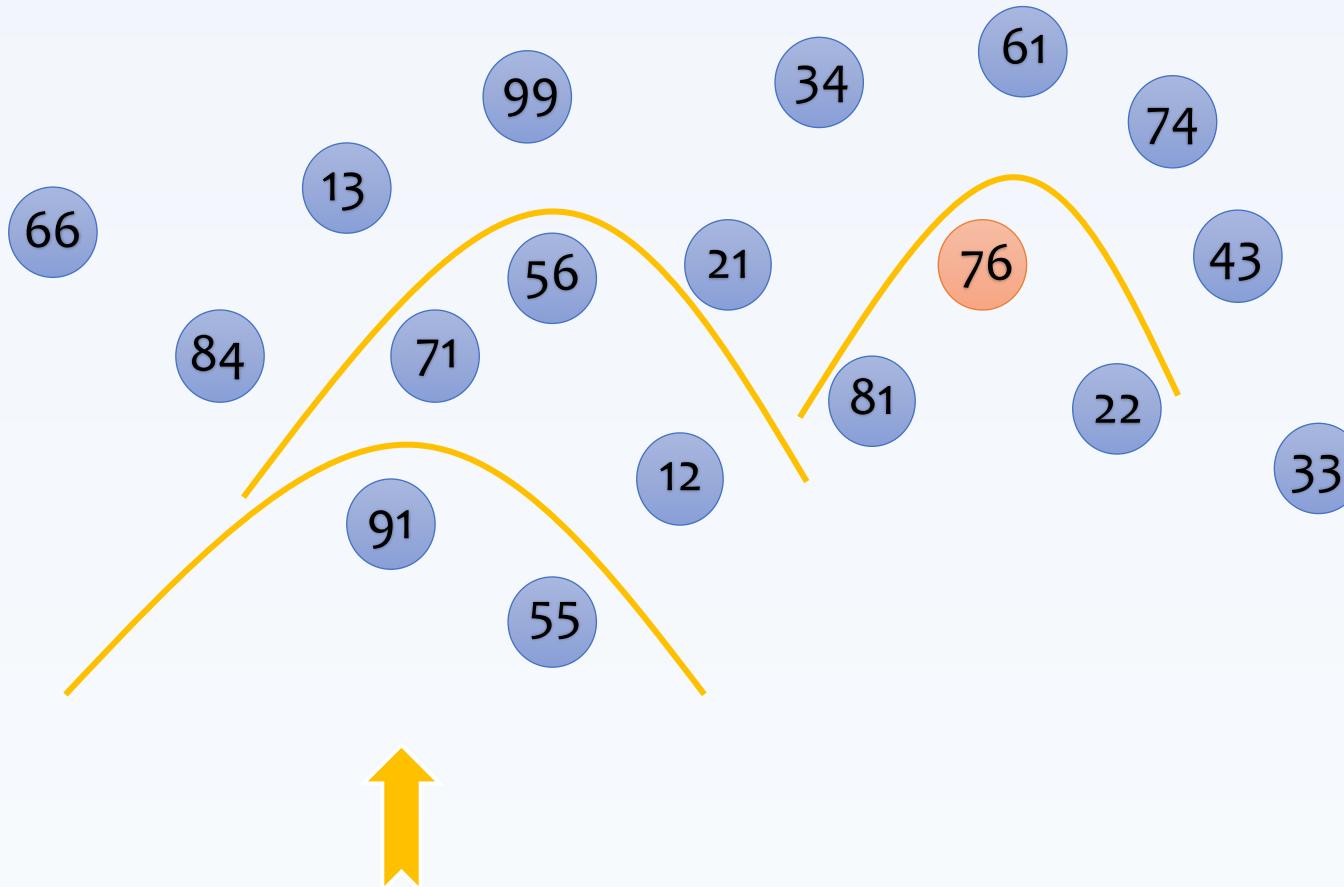
Searching for 76



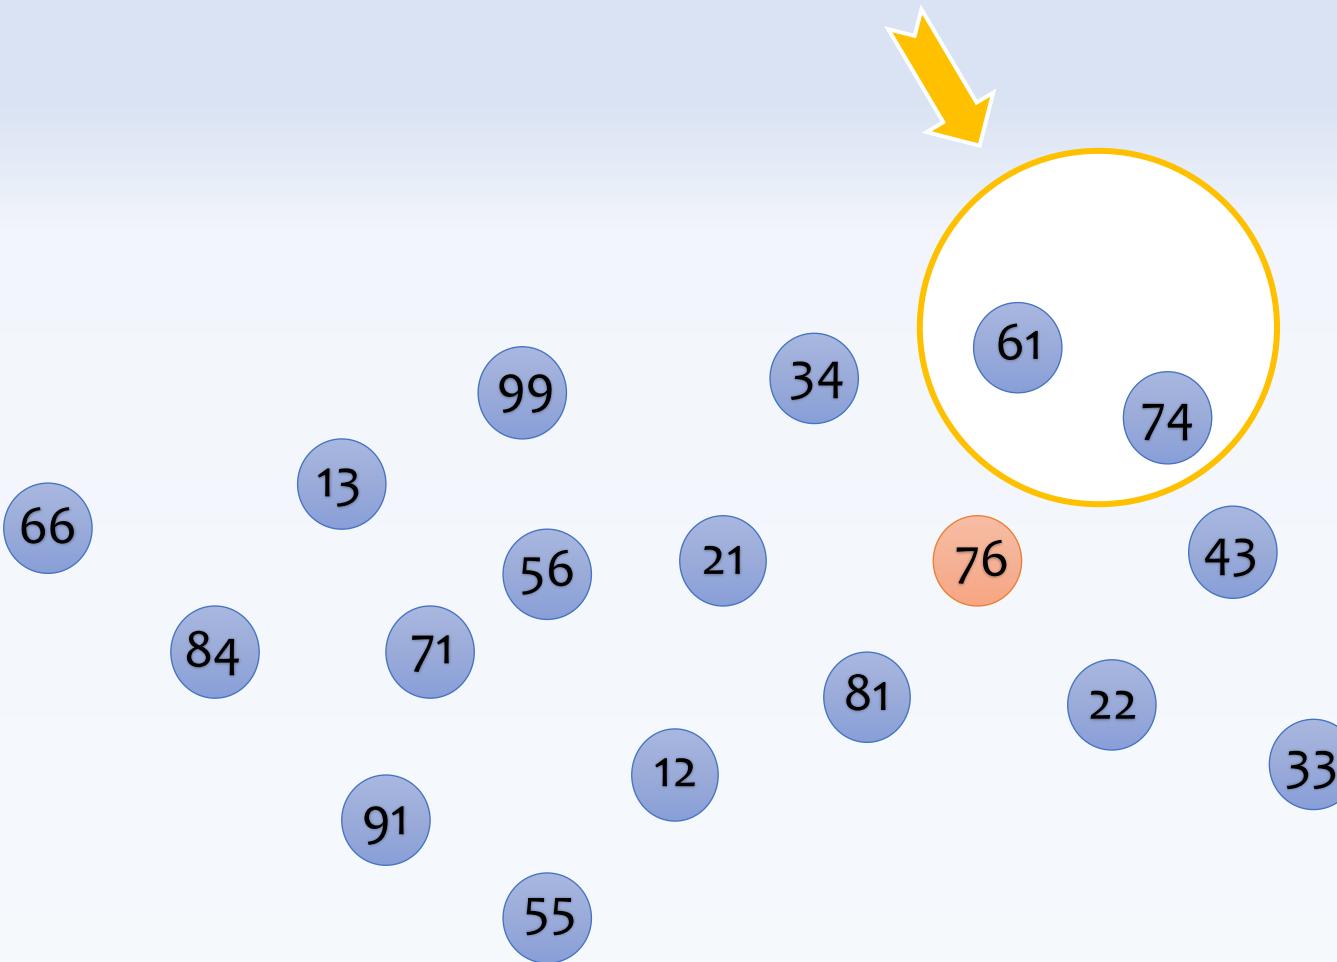
# Searching for 76



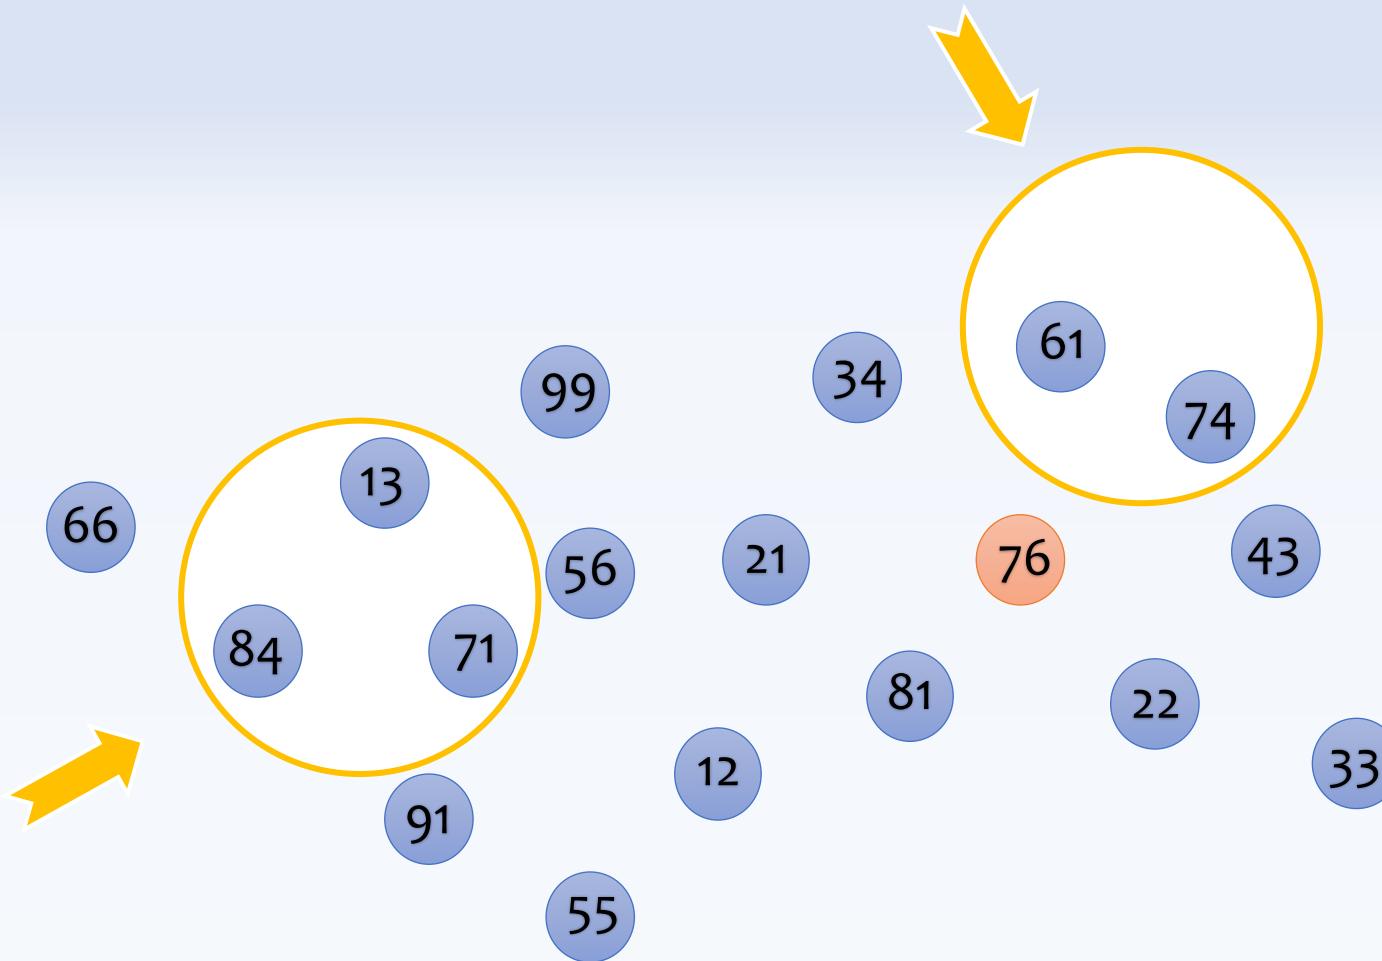
# Searching for 76



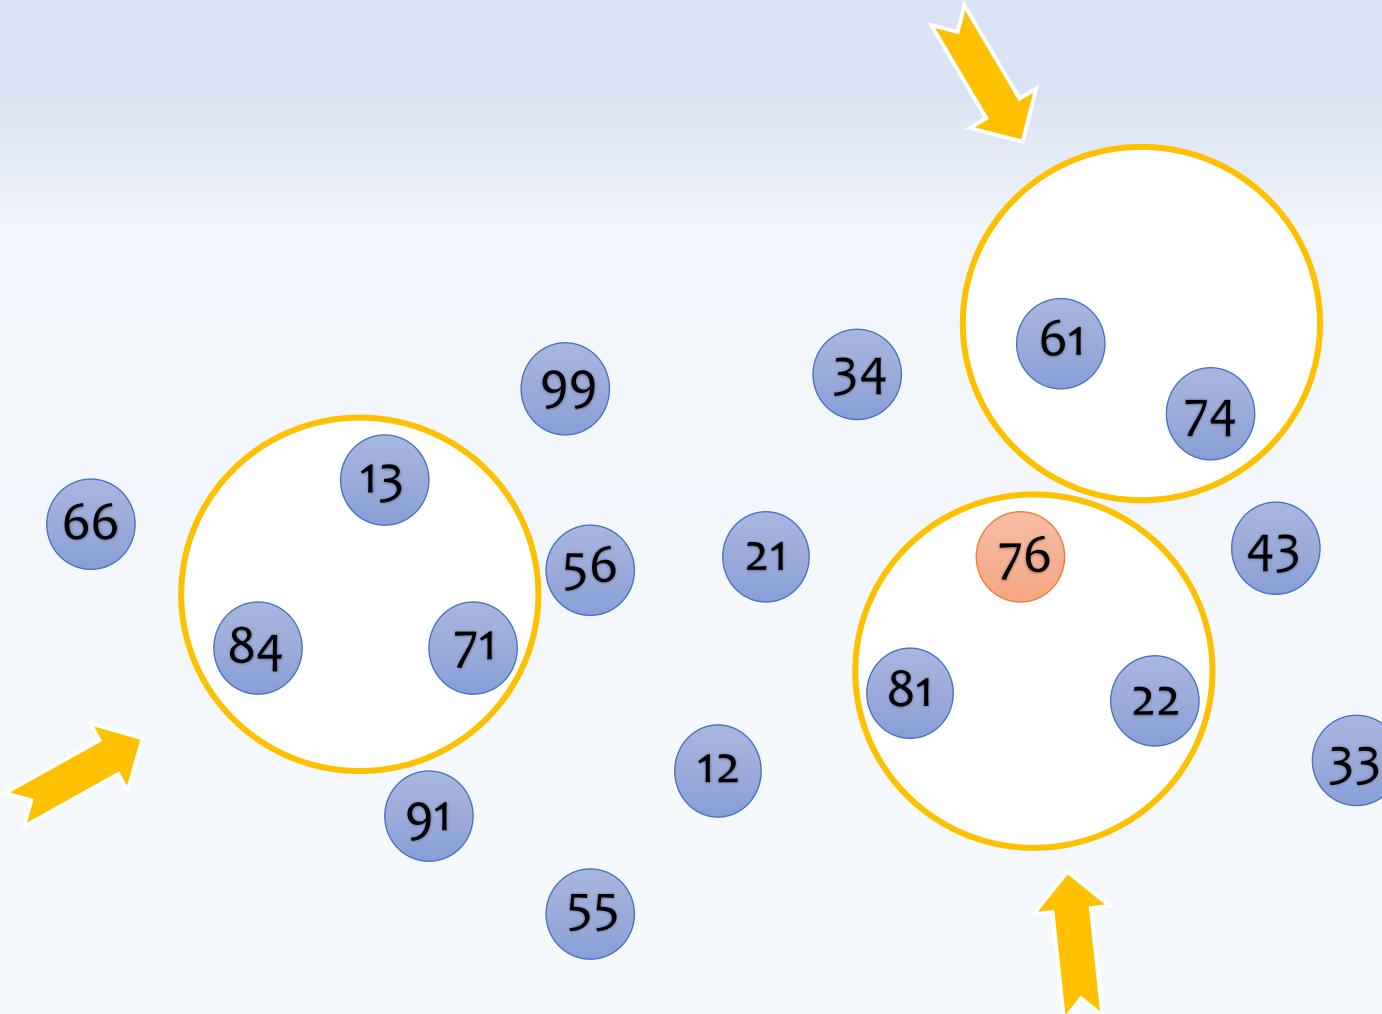
Searching for 76



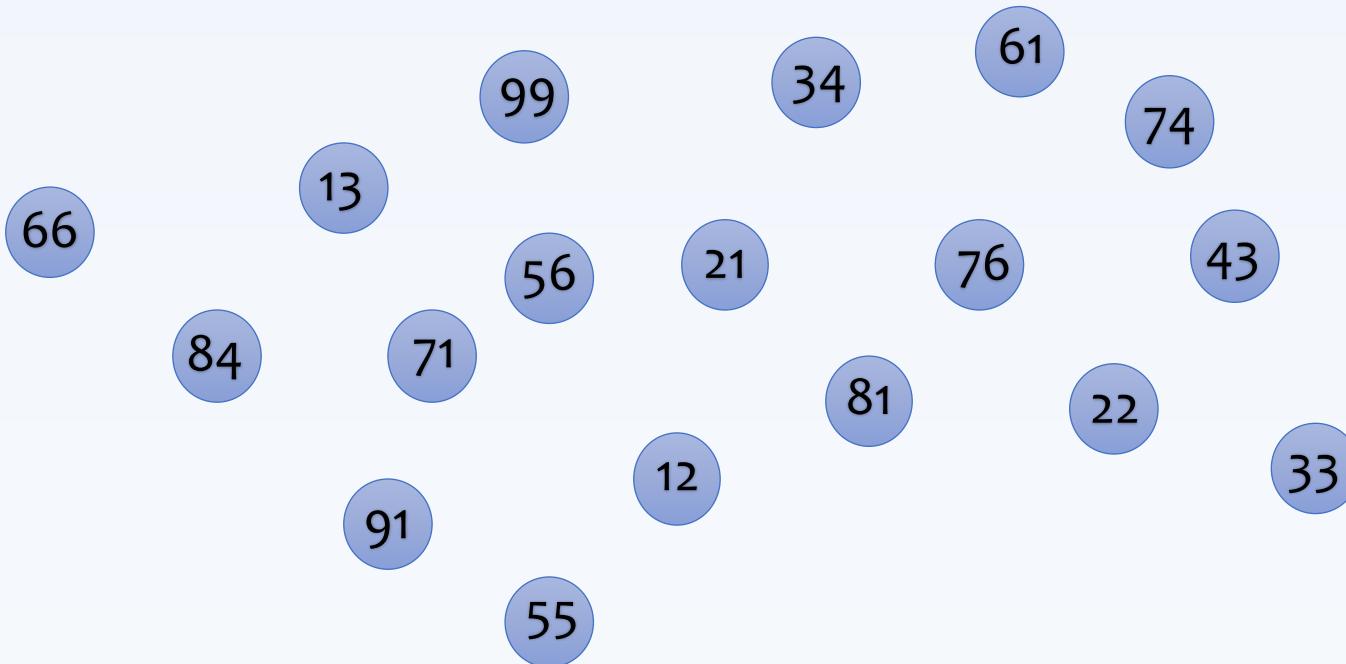
Searching for 76



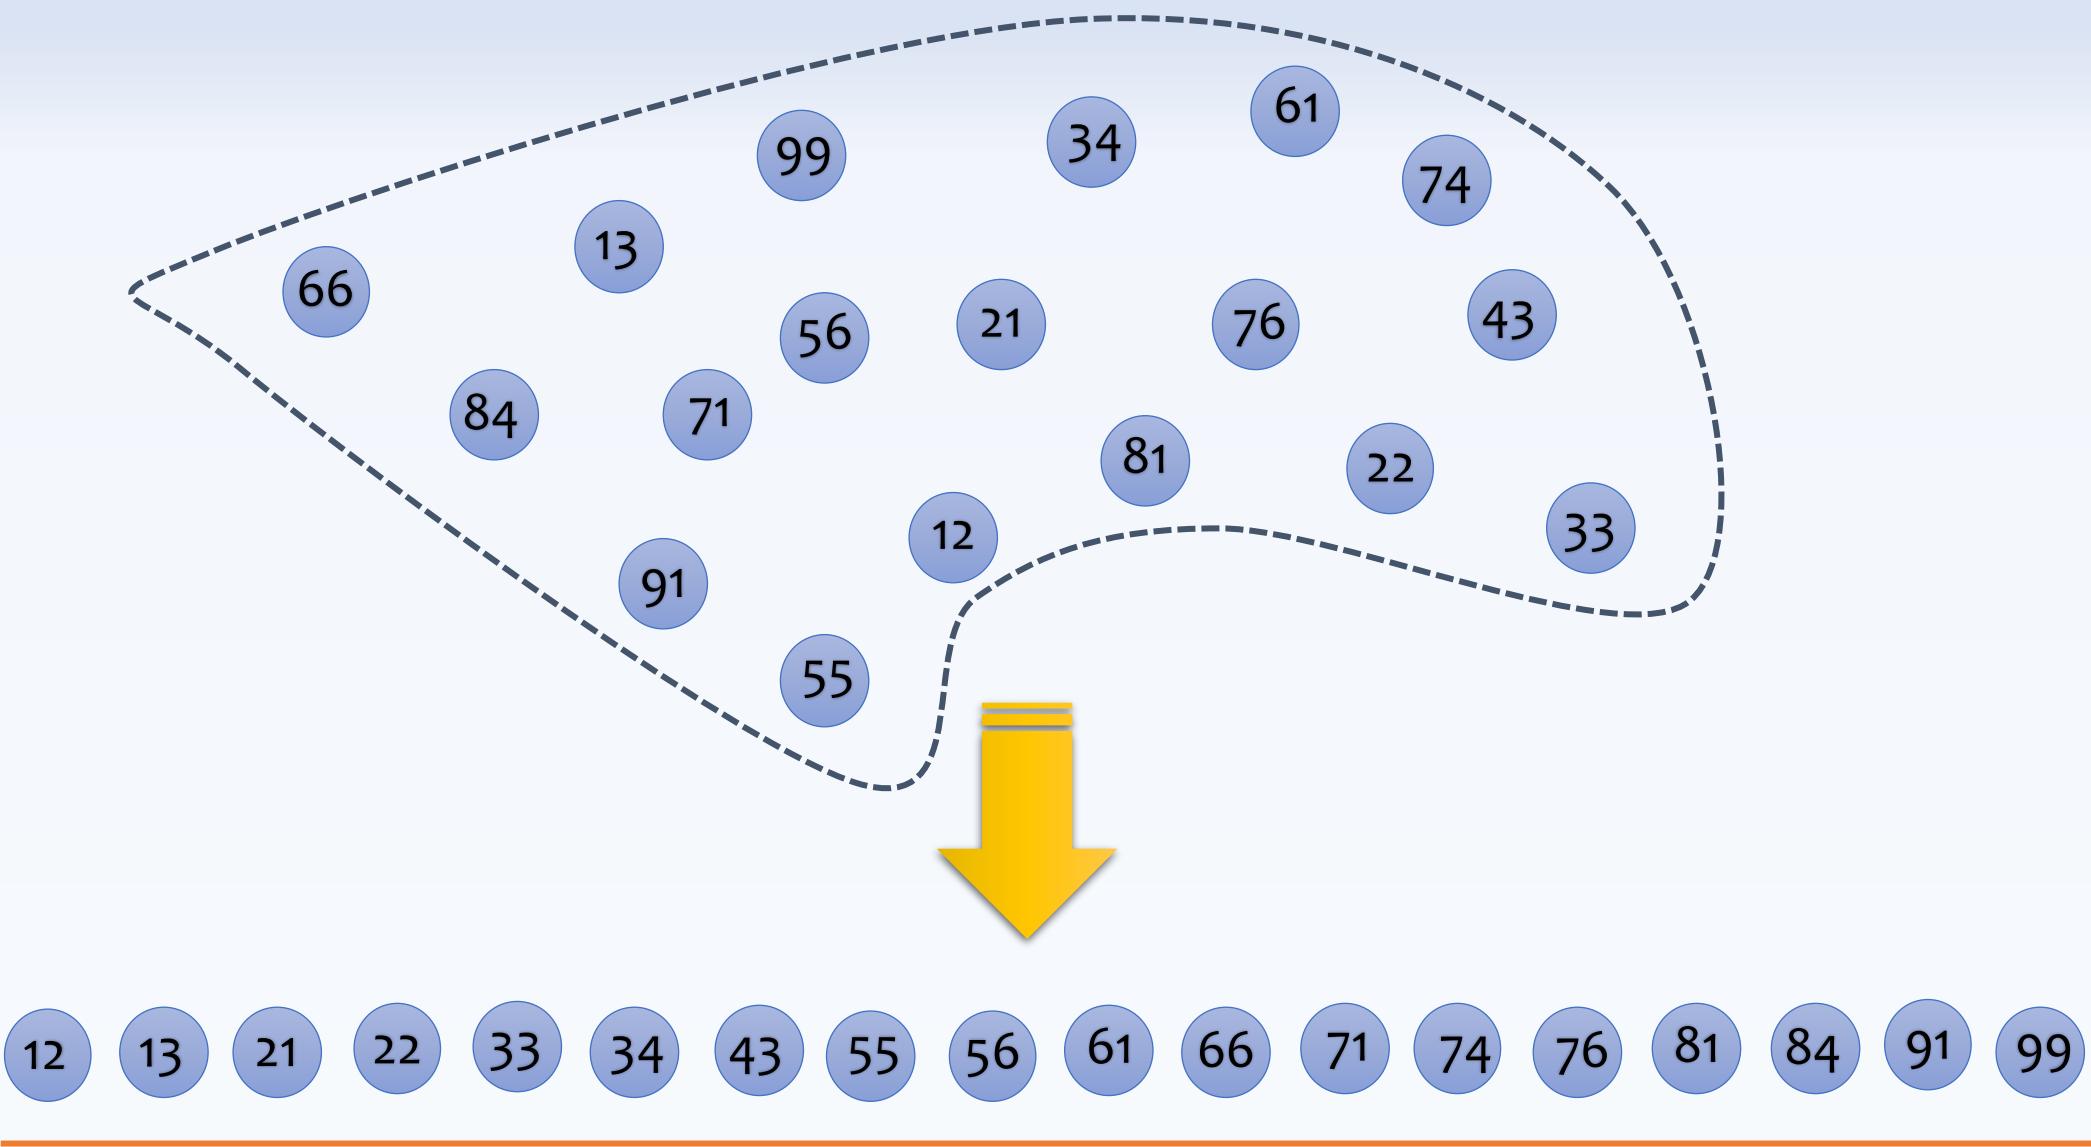
Searching for 76



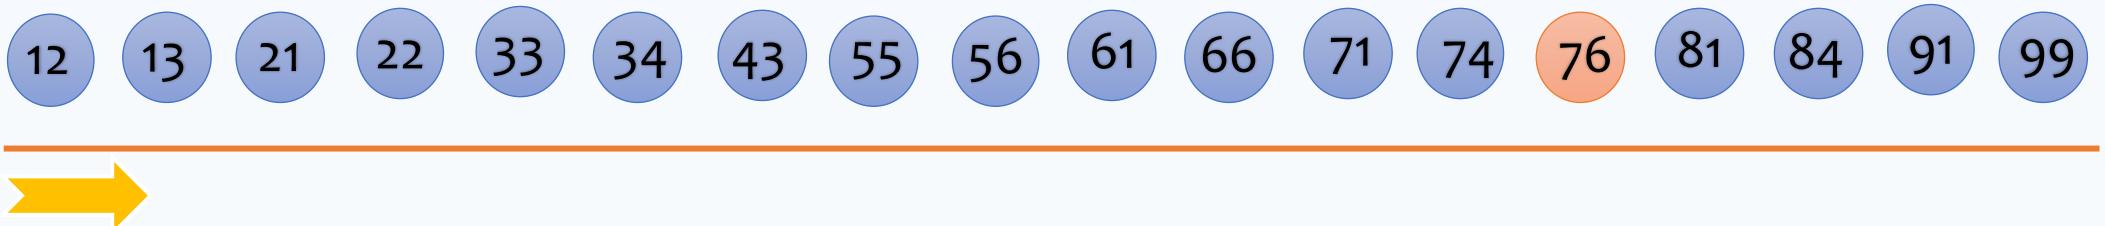
**Searching for 44?**  
(what-if the value does not exist)  
(could we have an early termination?)



**Could we impose an order to improve the search?**



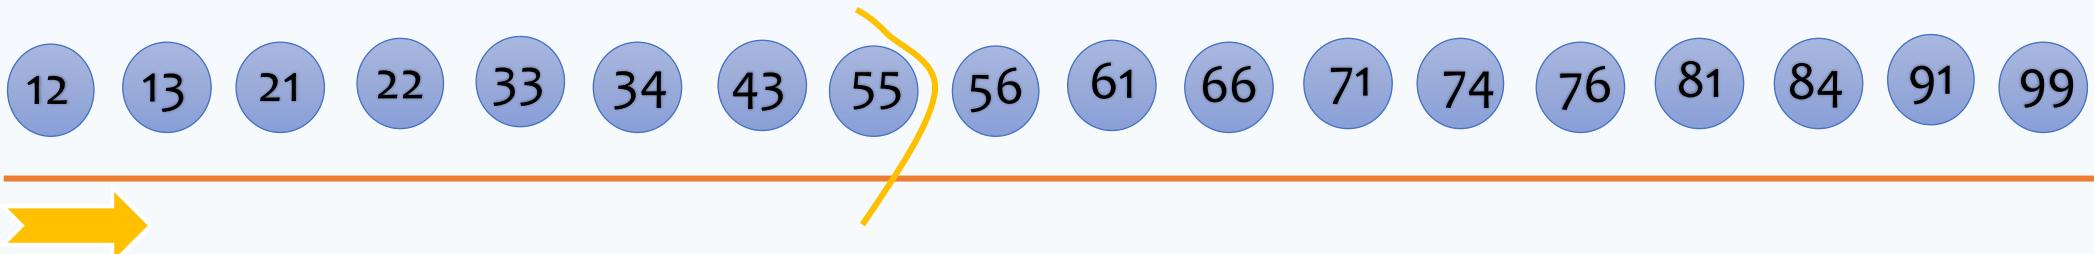
Searching for 76



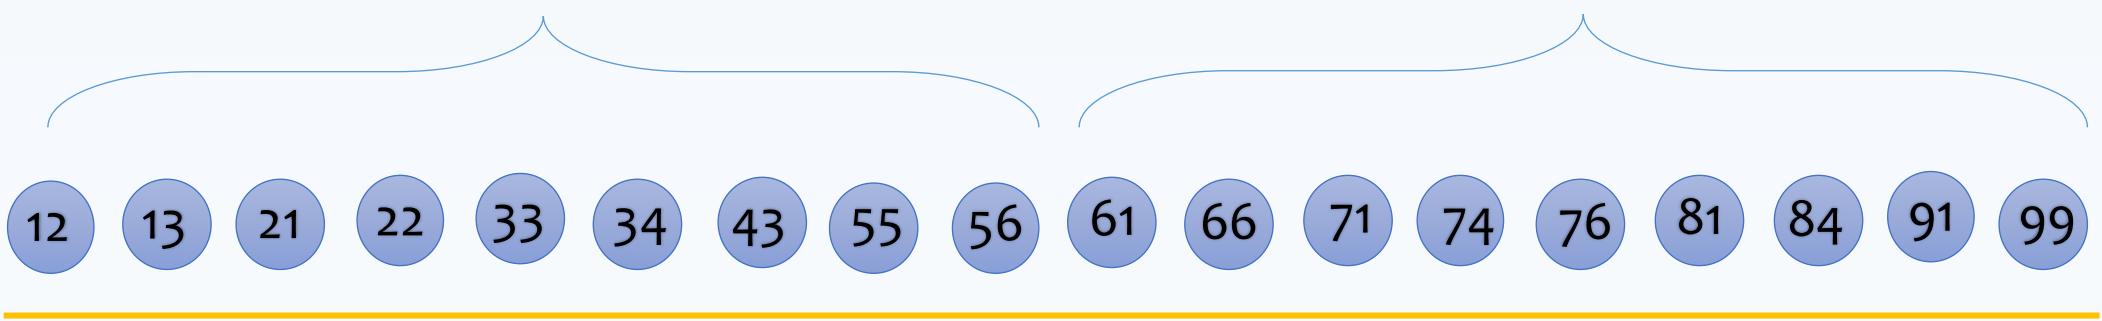
Searching for 76

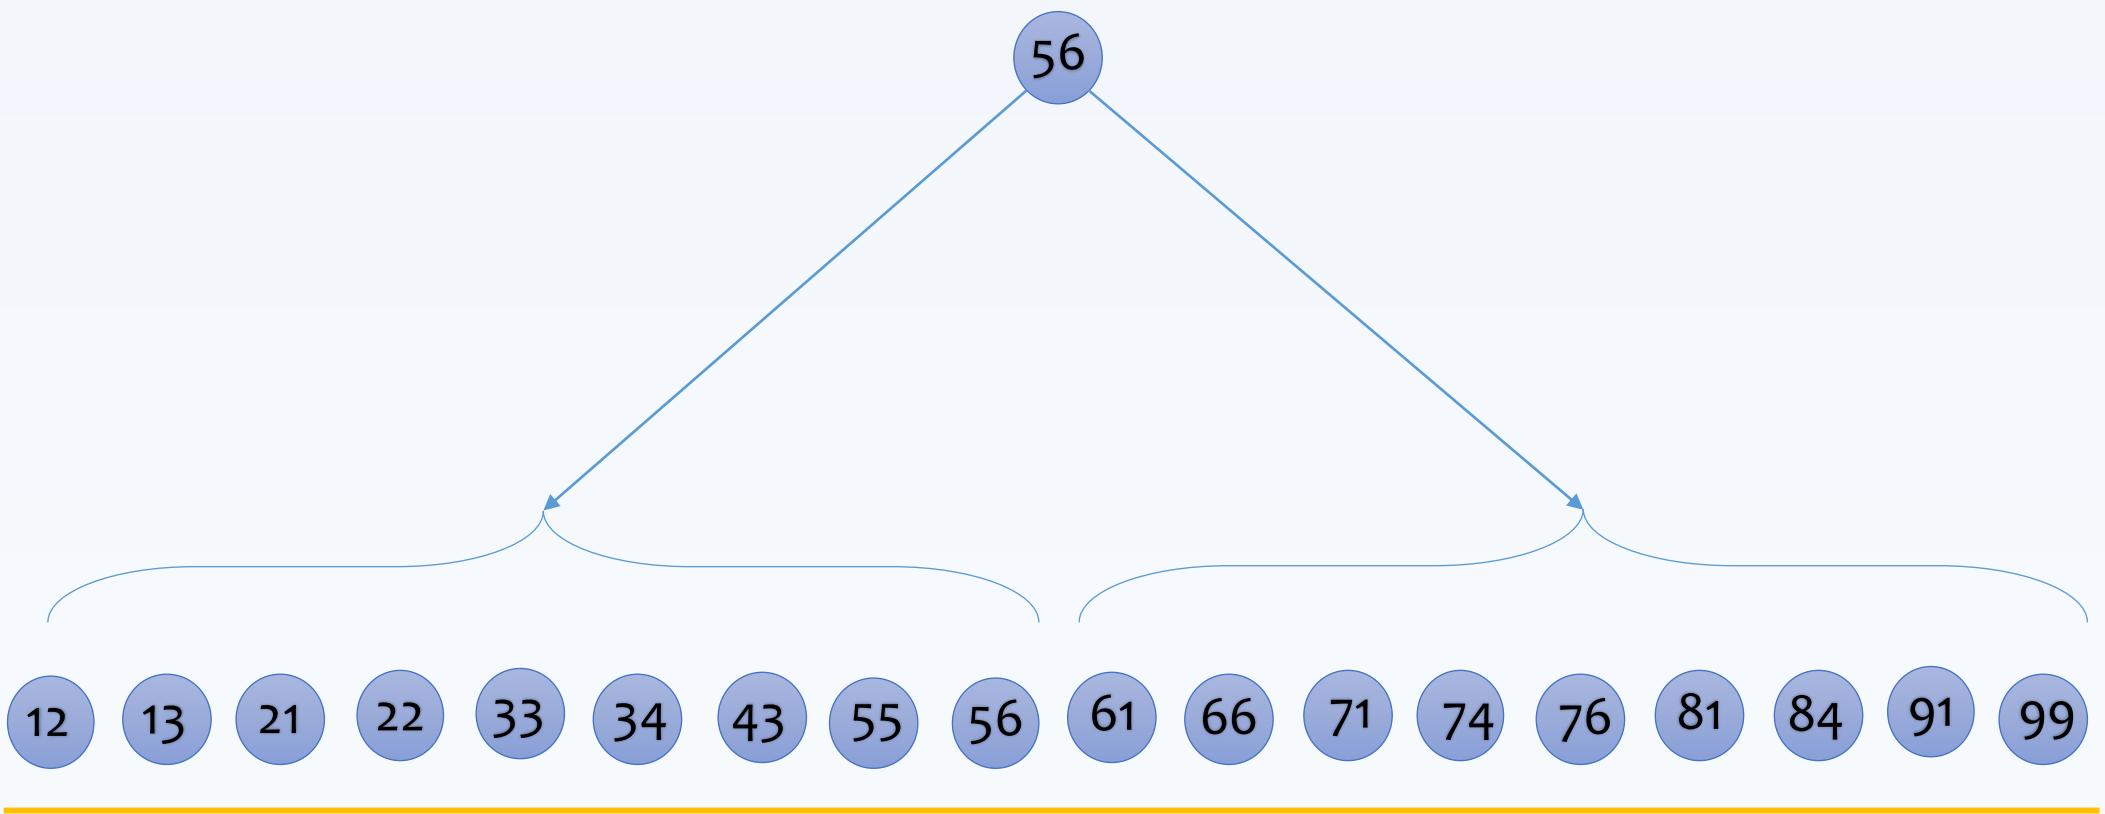


**Searching for 44?  
(could we have an early termination?)**

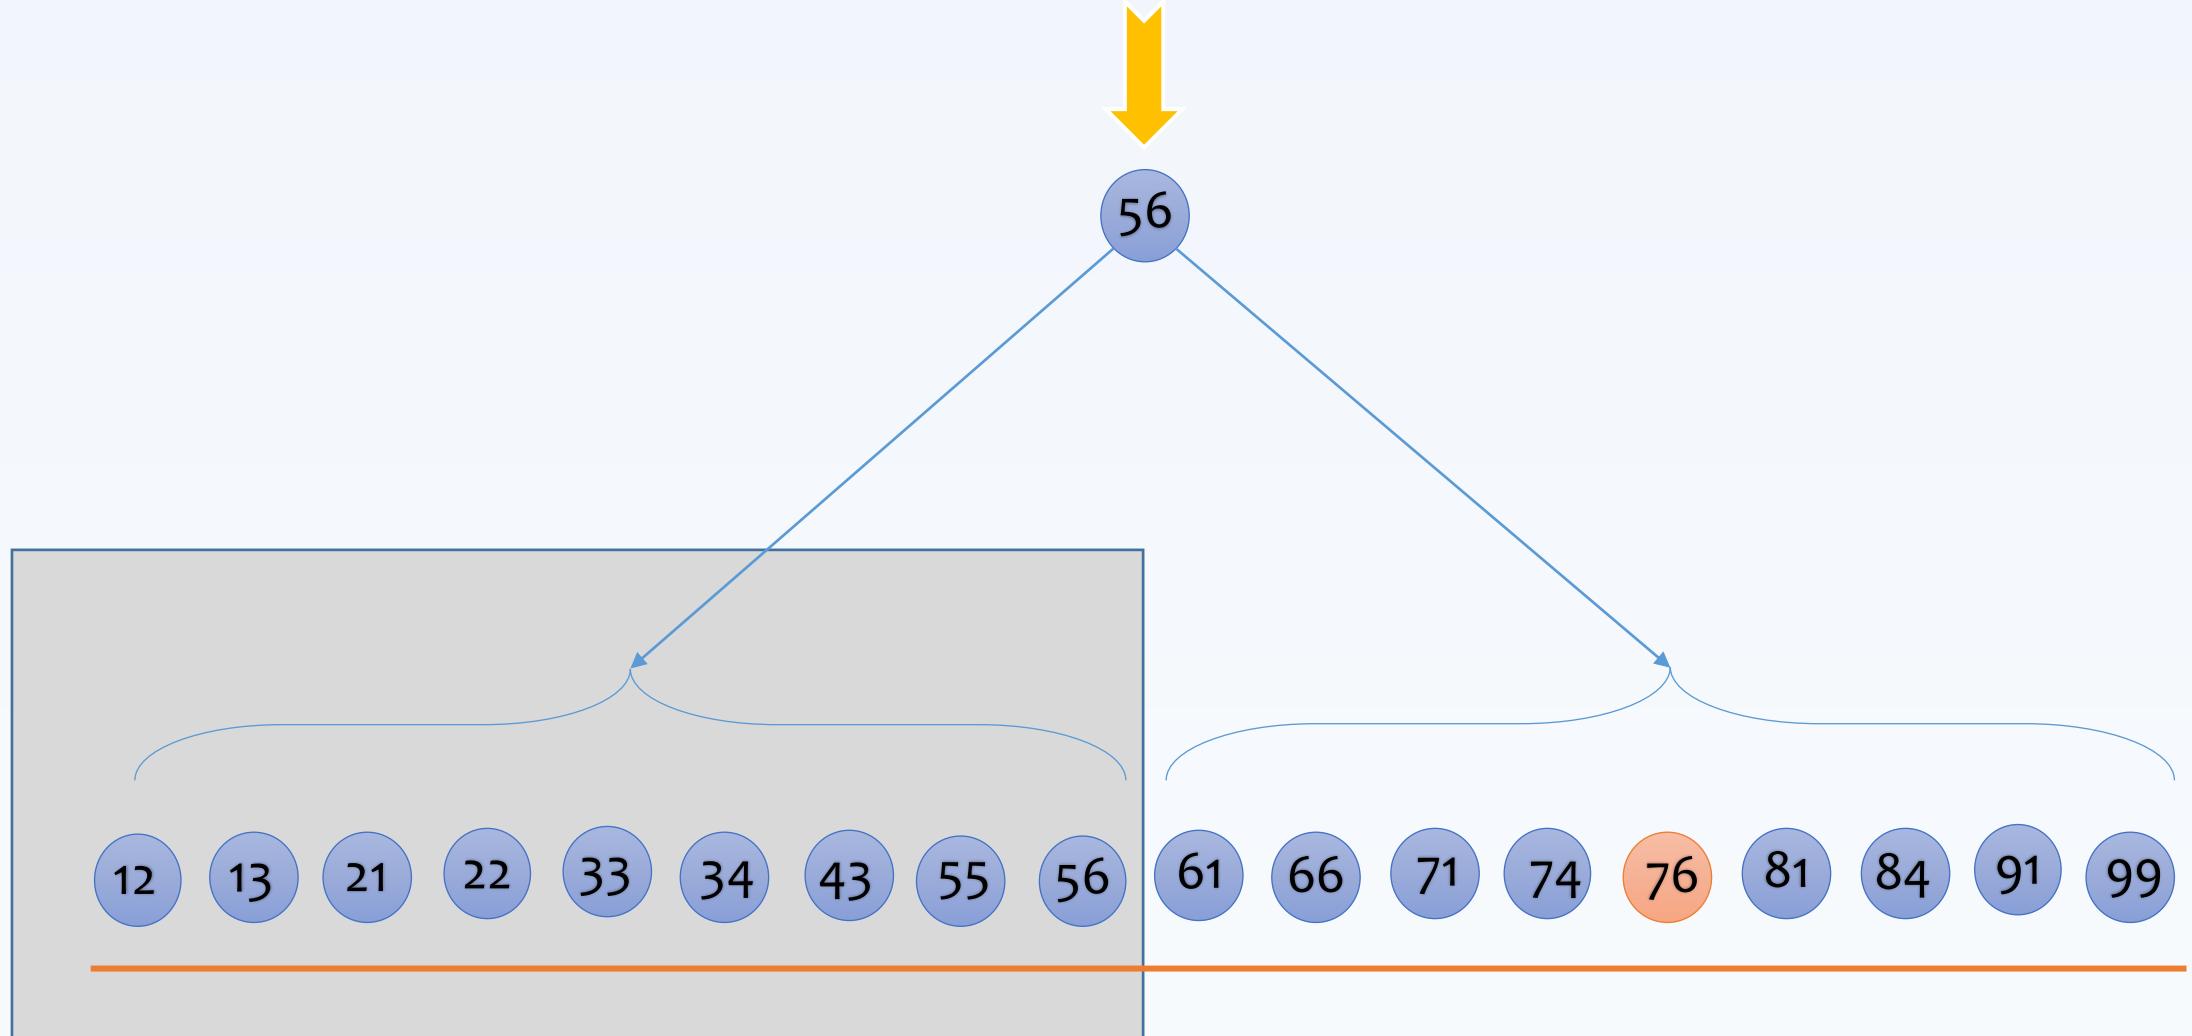


**Could we impose a structure to further improve the search?**

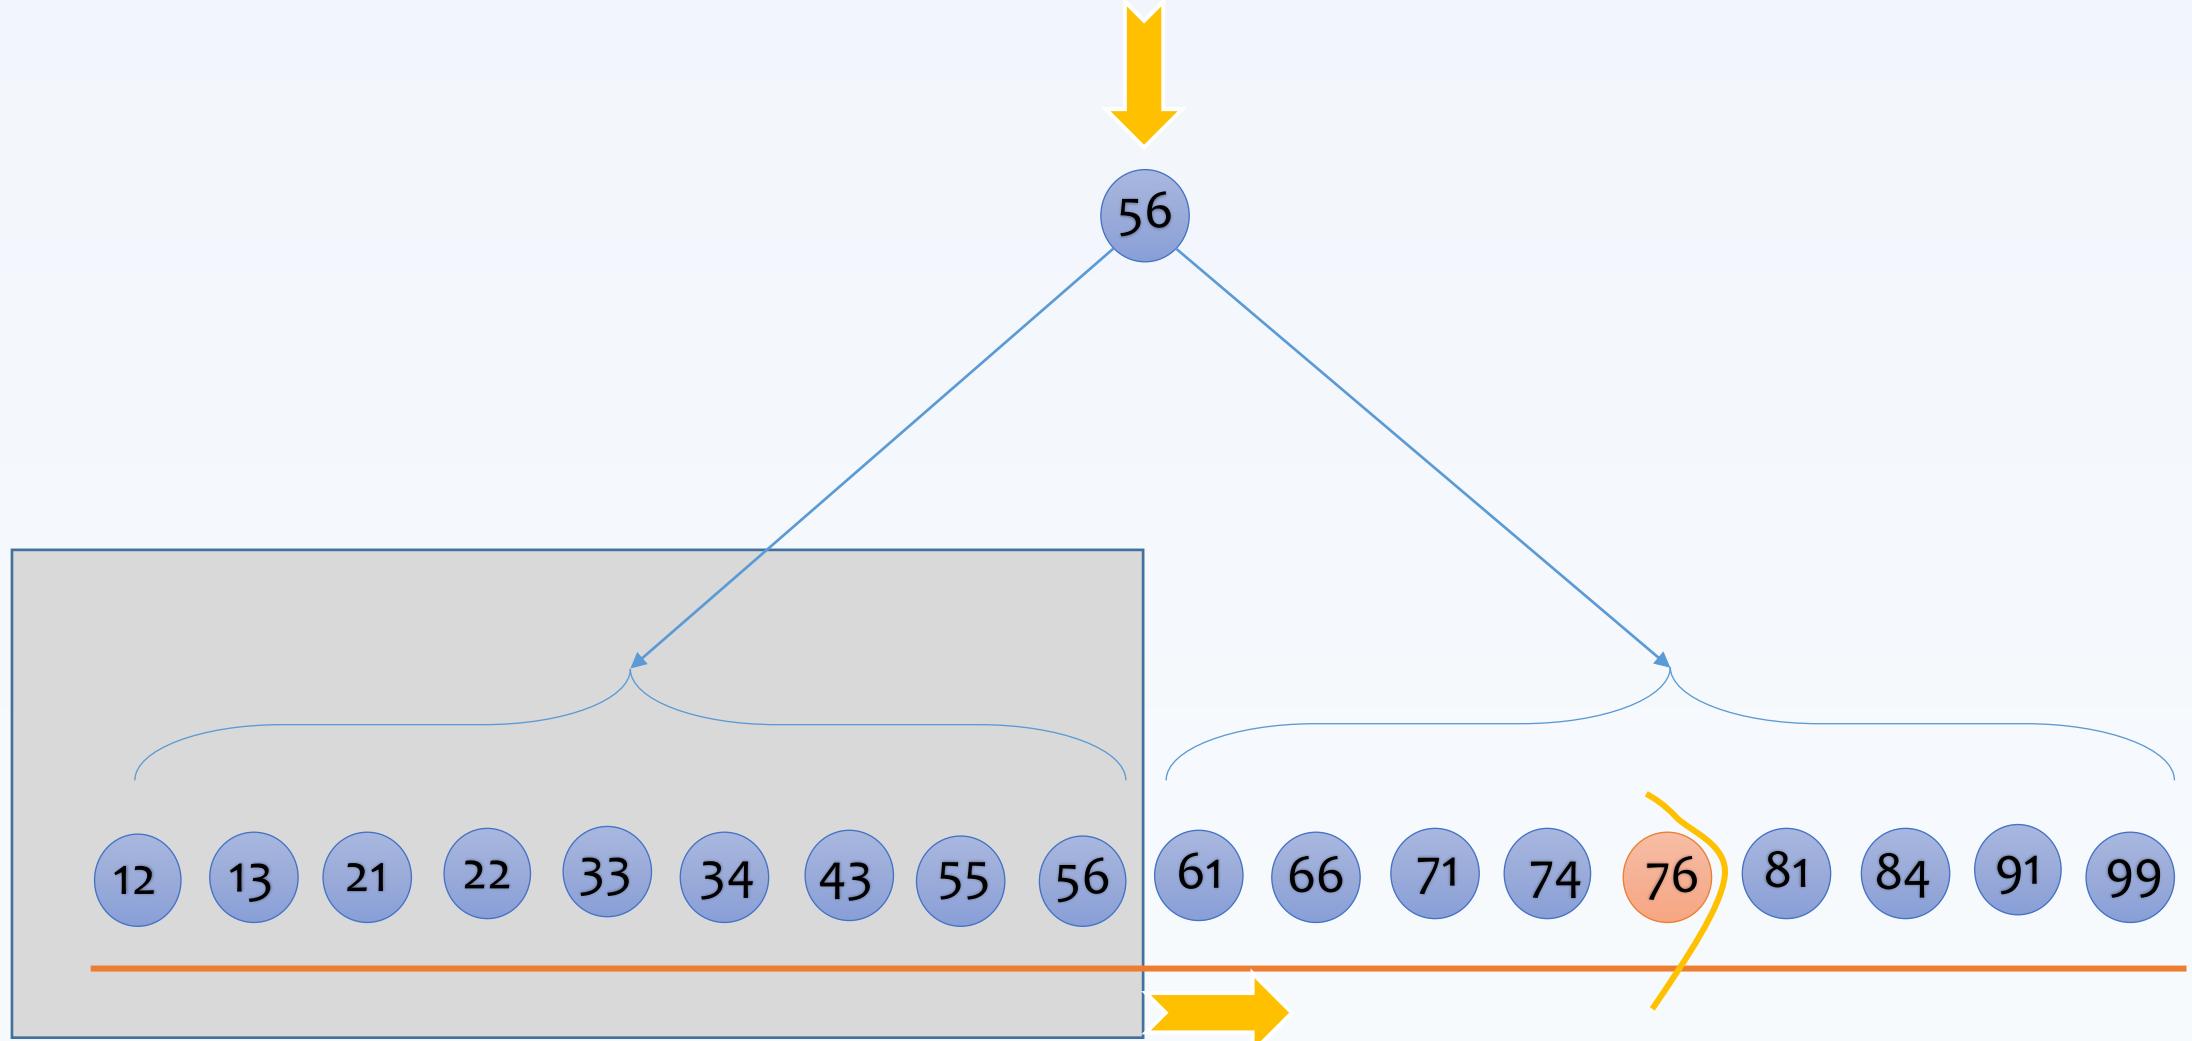


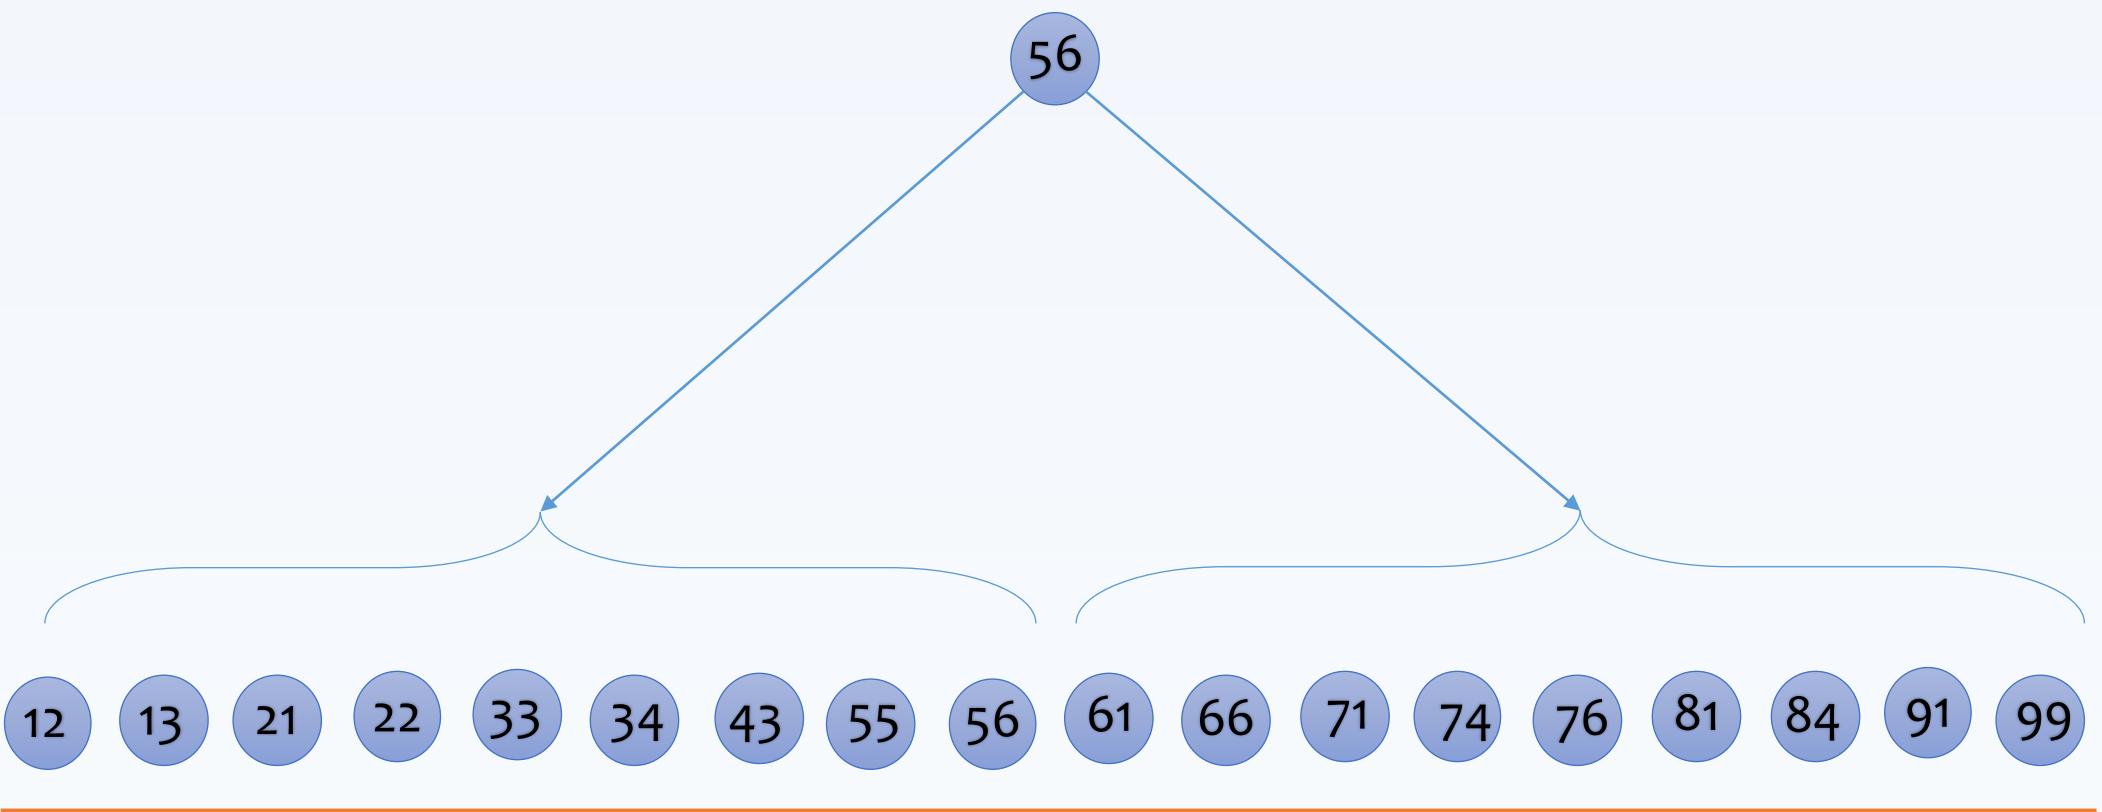


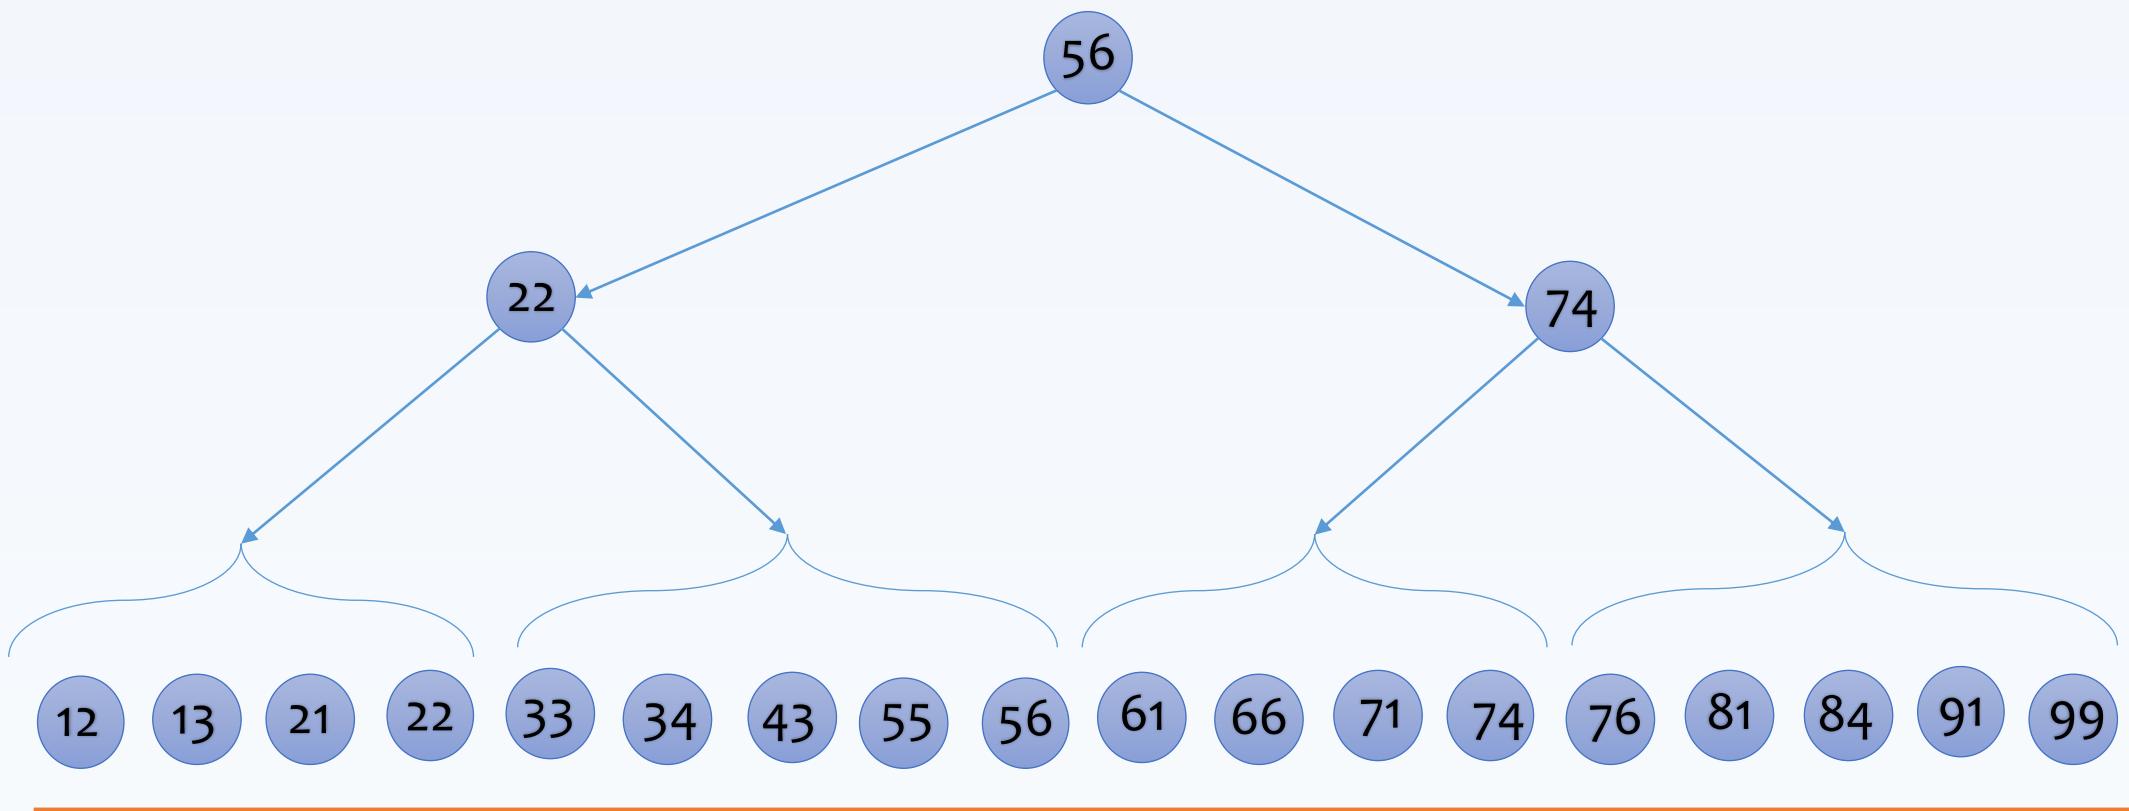
Searching for 76



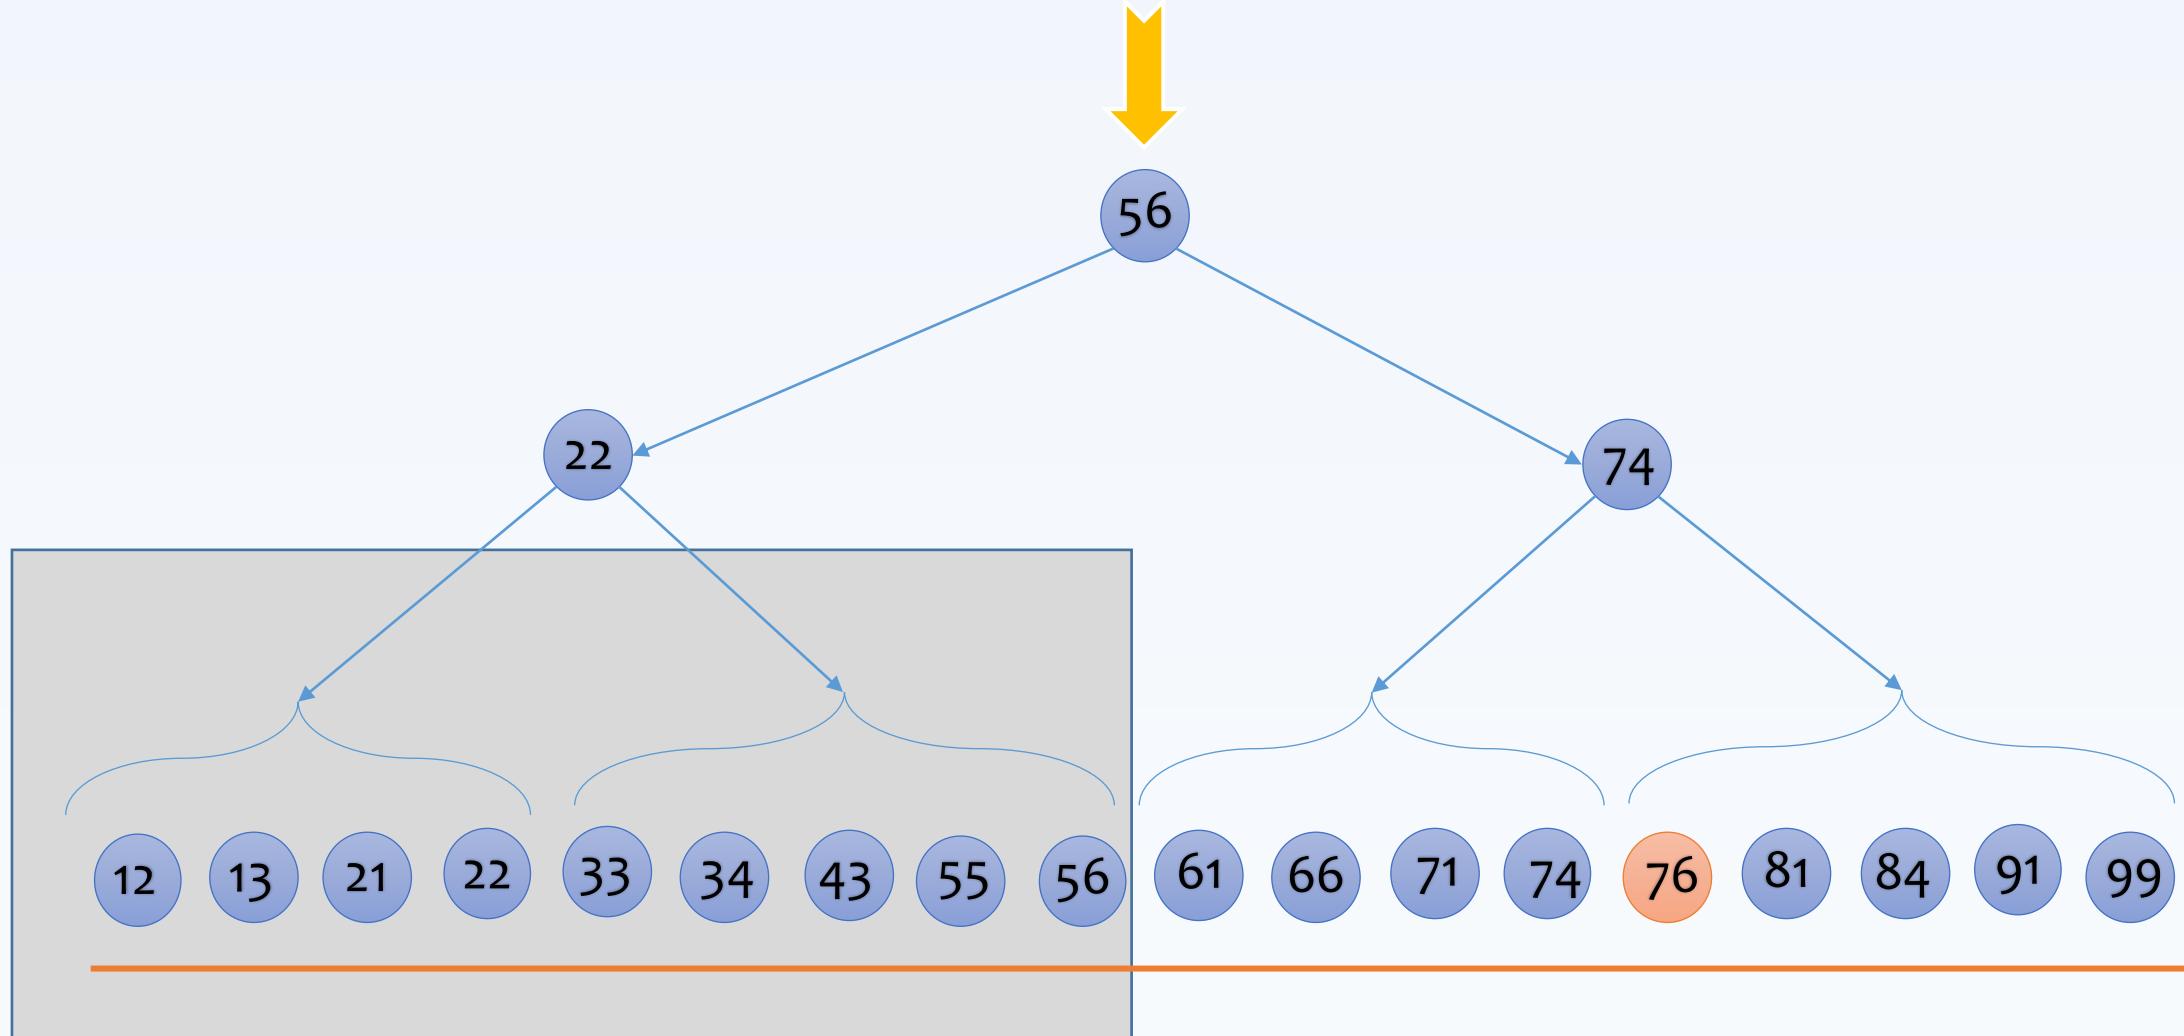
Searching for 76



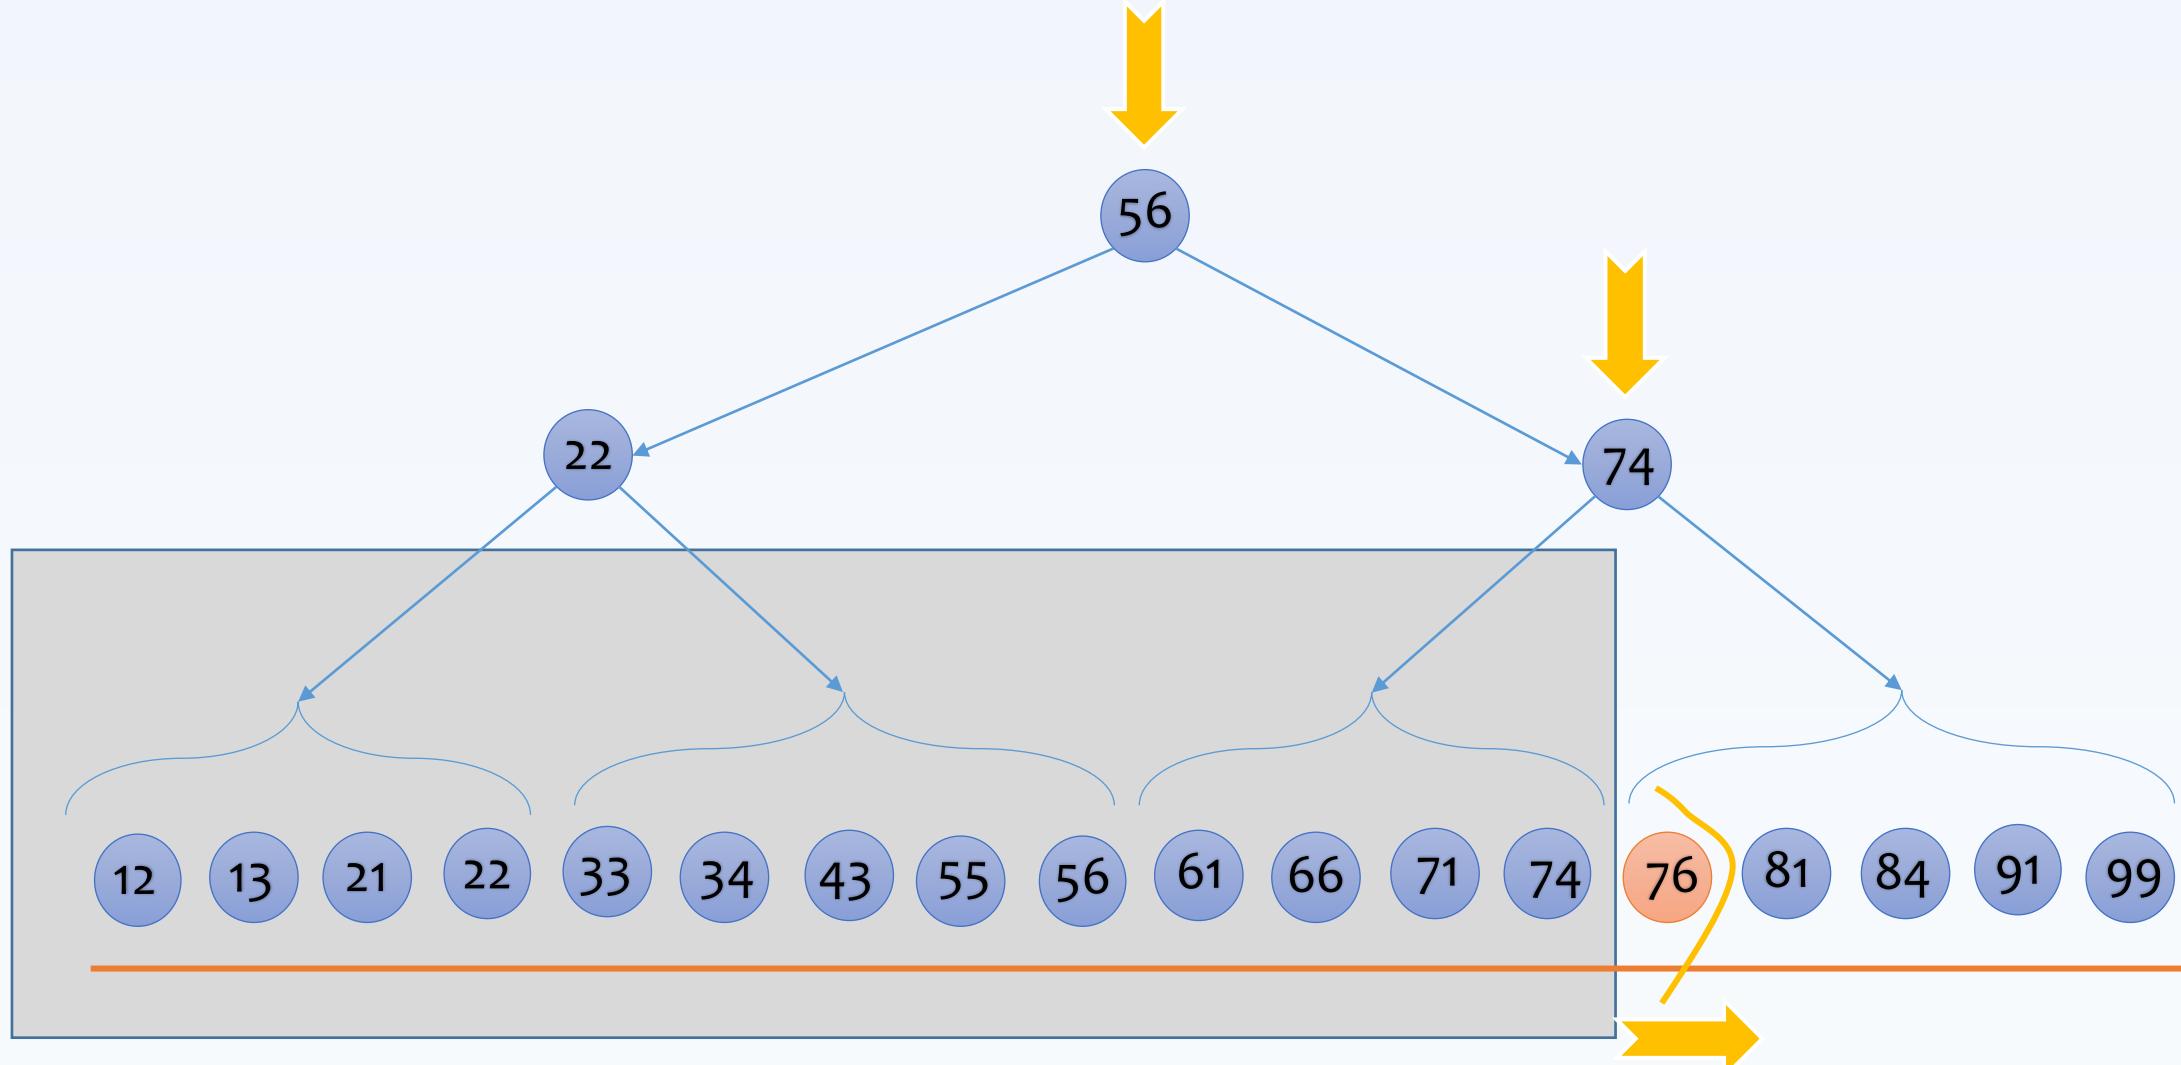




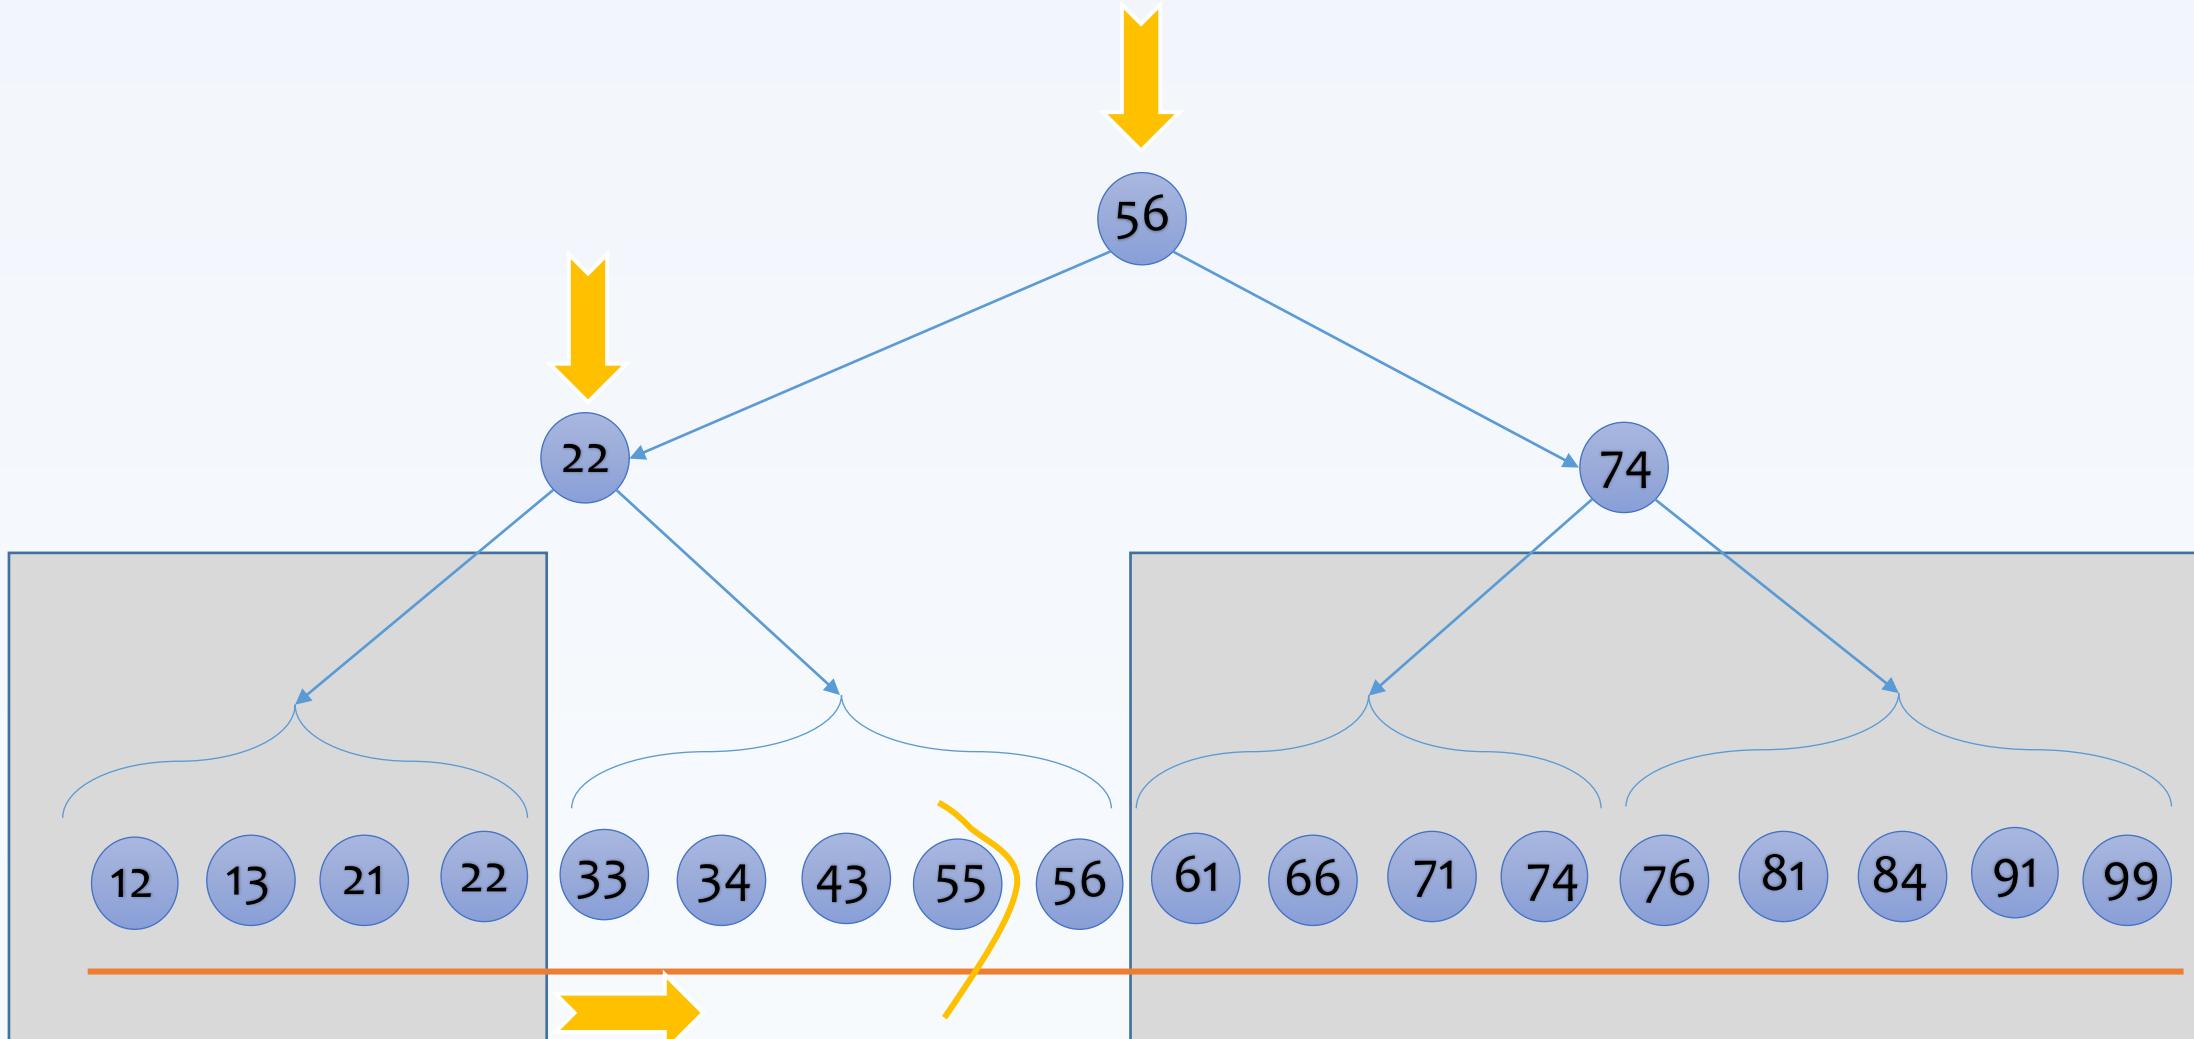
Searching for 76



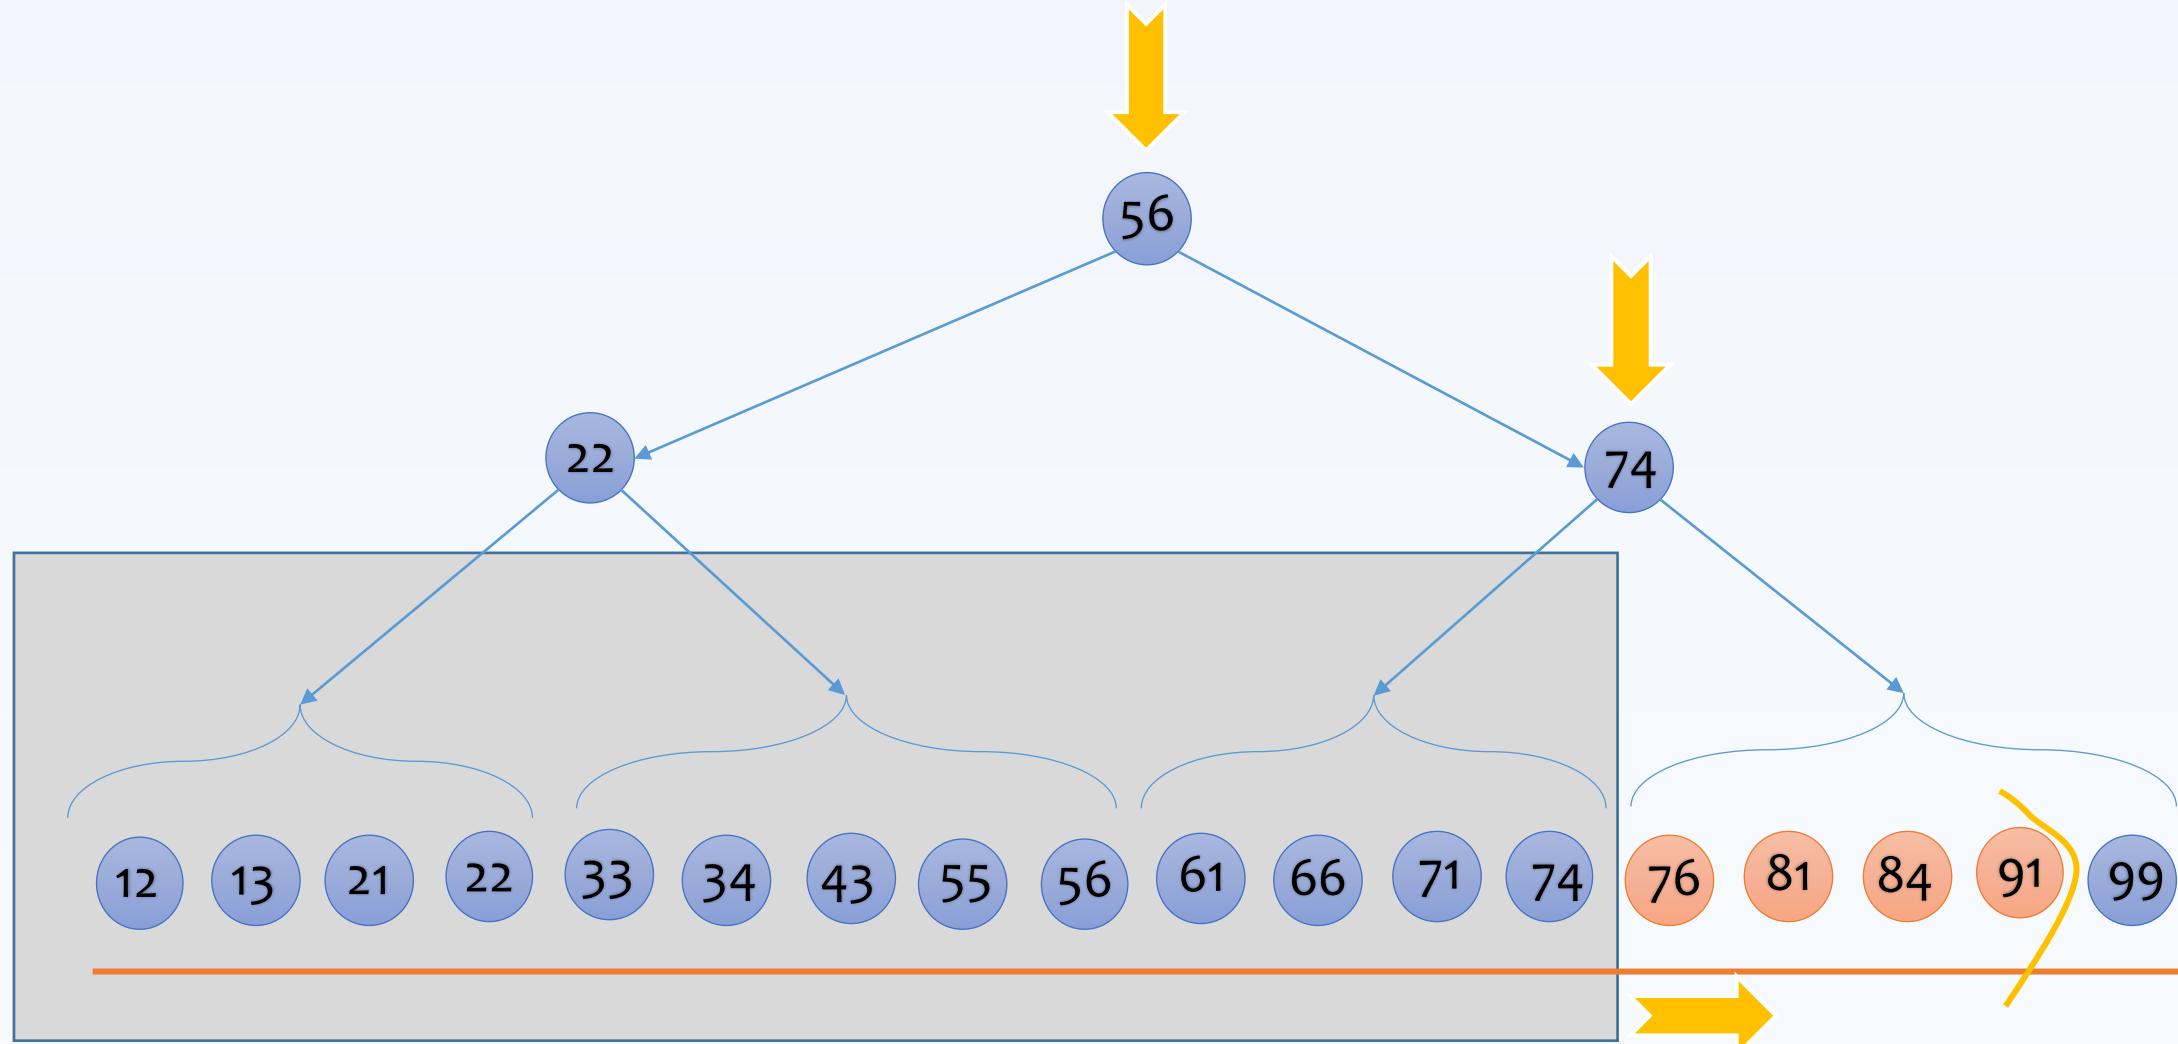
Searching for 76



# Searching for 44? (could we have an early termination?)



Searching for 76-91



**Could we spread the data cleverly to improve the search?**

## hashtable



bucket

Hashing ( ) = ?

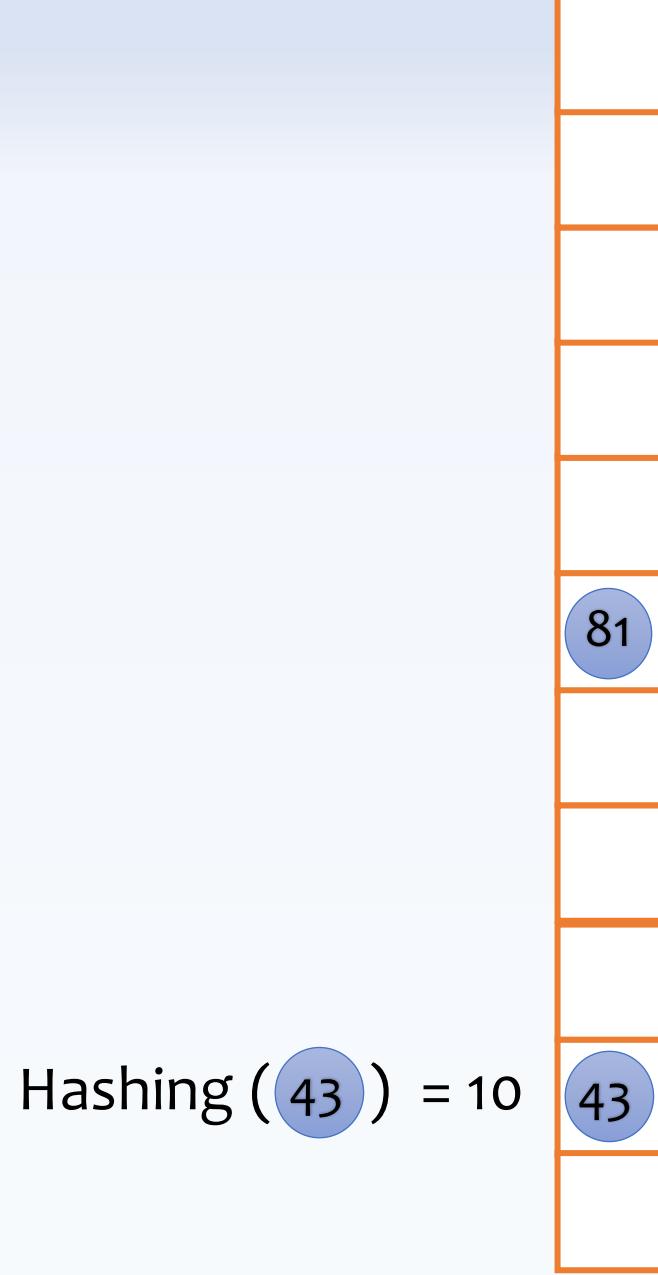
(returns a value  
between 1 to n,  
where n is the  
number of buckets)

## Inserting 81

Hashing ( 81 ) = 6



## Inserting 43

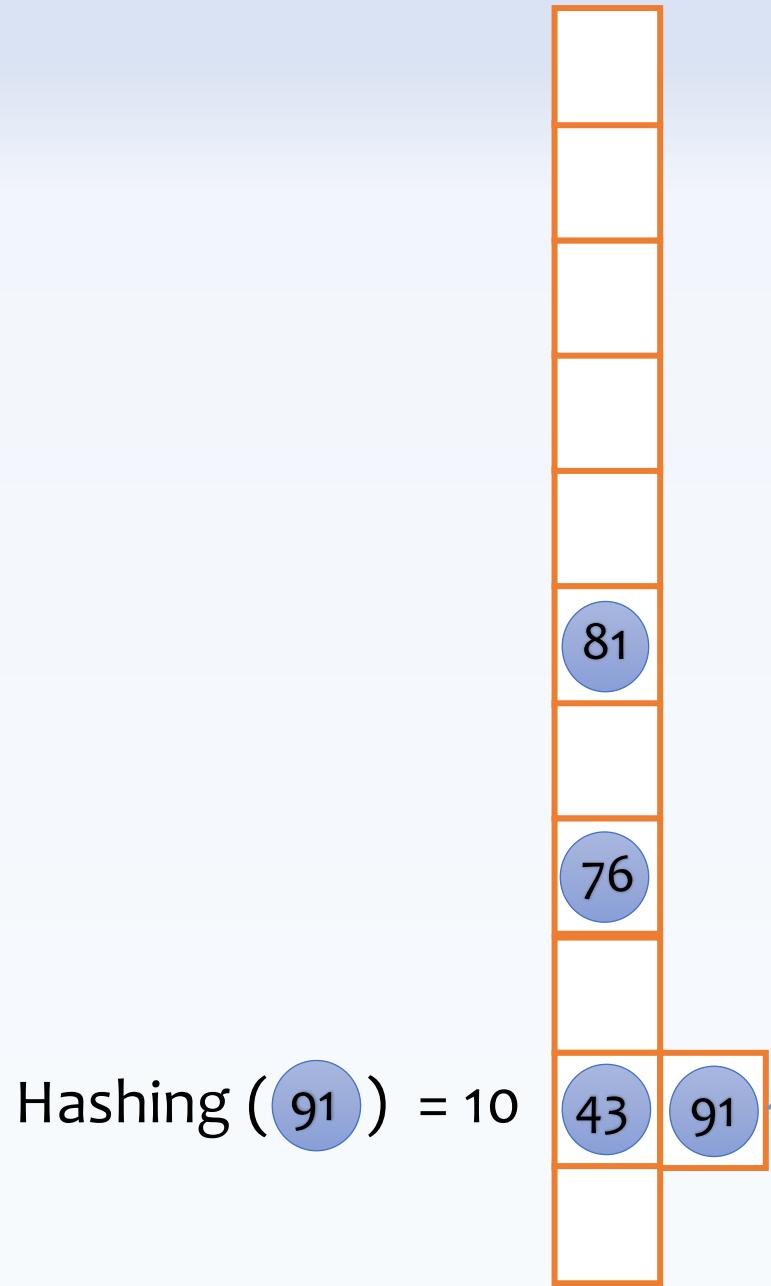


## Inserting 76

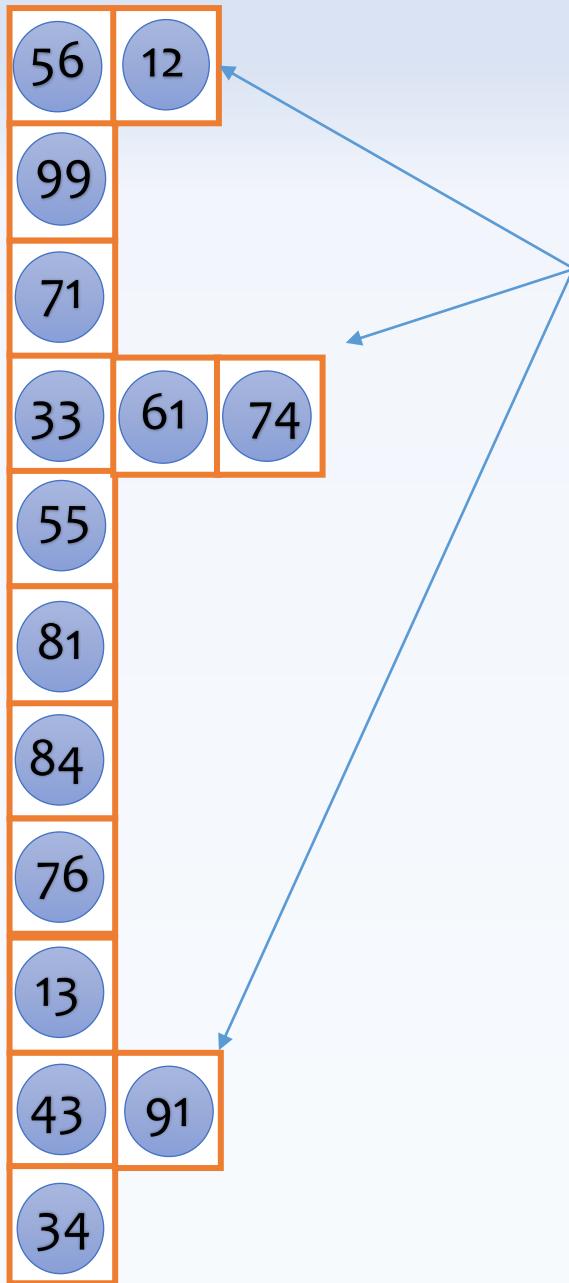
Hashing ( 76 ) = 8



## Inserting 91



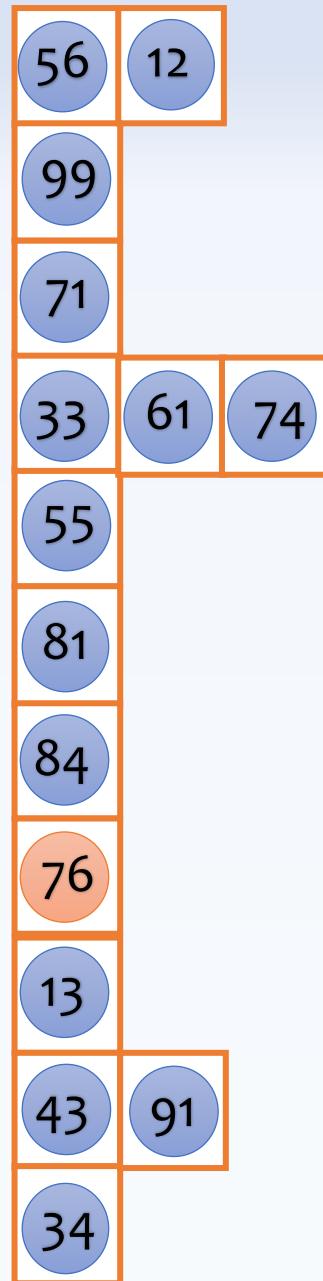
collisions  
(when multiple values  
hash to the same bucket)



collisions  
(when multiple values  
hash to the same bucket)

Searching for 76

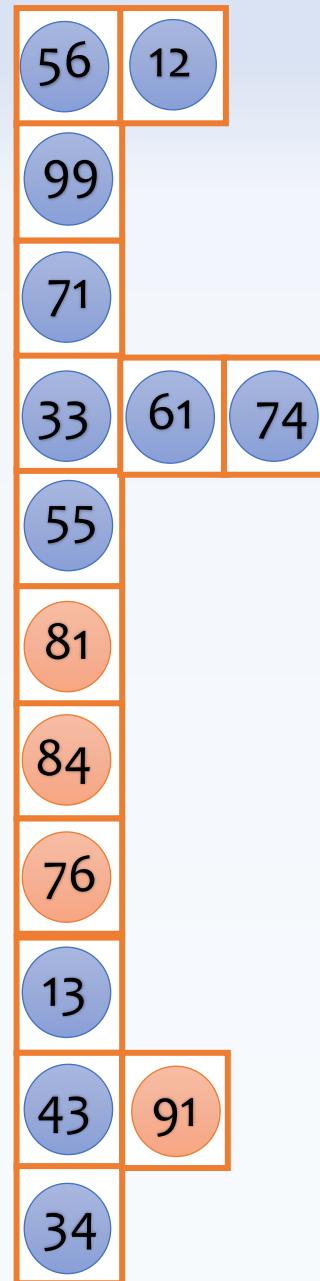
(now we can have a constant lookup cost)



Hashing ( 76 ) = 8

Searching for 76-91?

Could we instead search for  
76, 77, 78, ..., 90, 91?



Hashing ( 76 ) = 8

Hashing ( 77 ) = 1

Hashing ( 78 ) = 3

|

Hashing ( 81 ) = 6

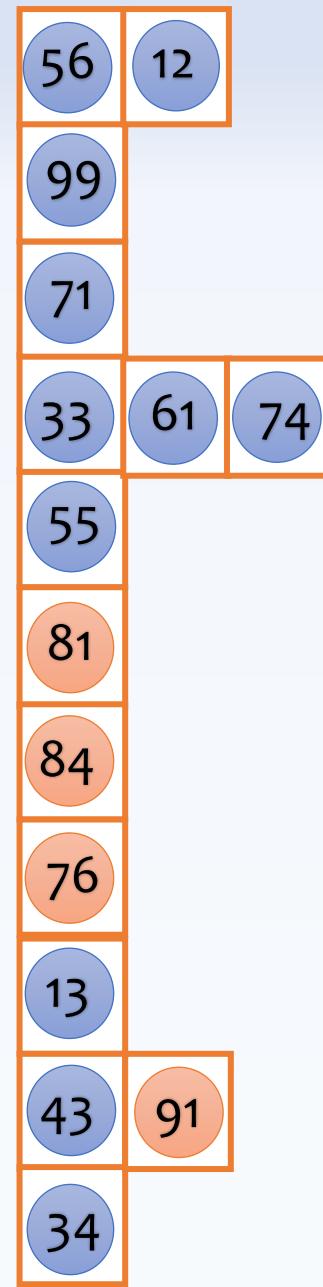
|

Hashing ( 84 ) = 7

|

Hashing ( 90 ) = 8

Hashing ( 91 ) = 10



Searching for 76-91  
Could we instead search for  
76, 77, 78, ..., 90, 91?

Hashing ( 76 ) = 8

Hashing ( 77 ) = 1

Hashing ( 78 ) = 3

|

Hashing ( 81 ) = 6

|

Hashing ( 84 ) = 7

|

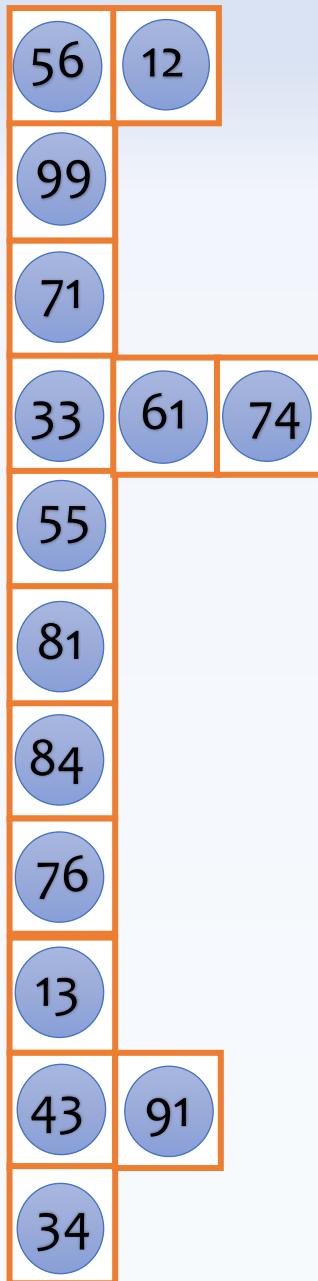
Hashing ( 90 ) = 8

Hashing ( 91 ) = 10

56	12	
99		
71		
33	61	74
55		
81		
84		
76		
13		
43	91	
34		

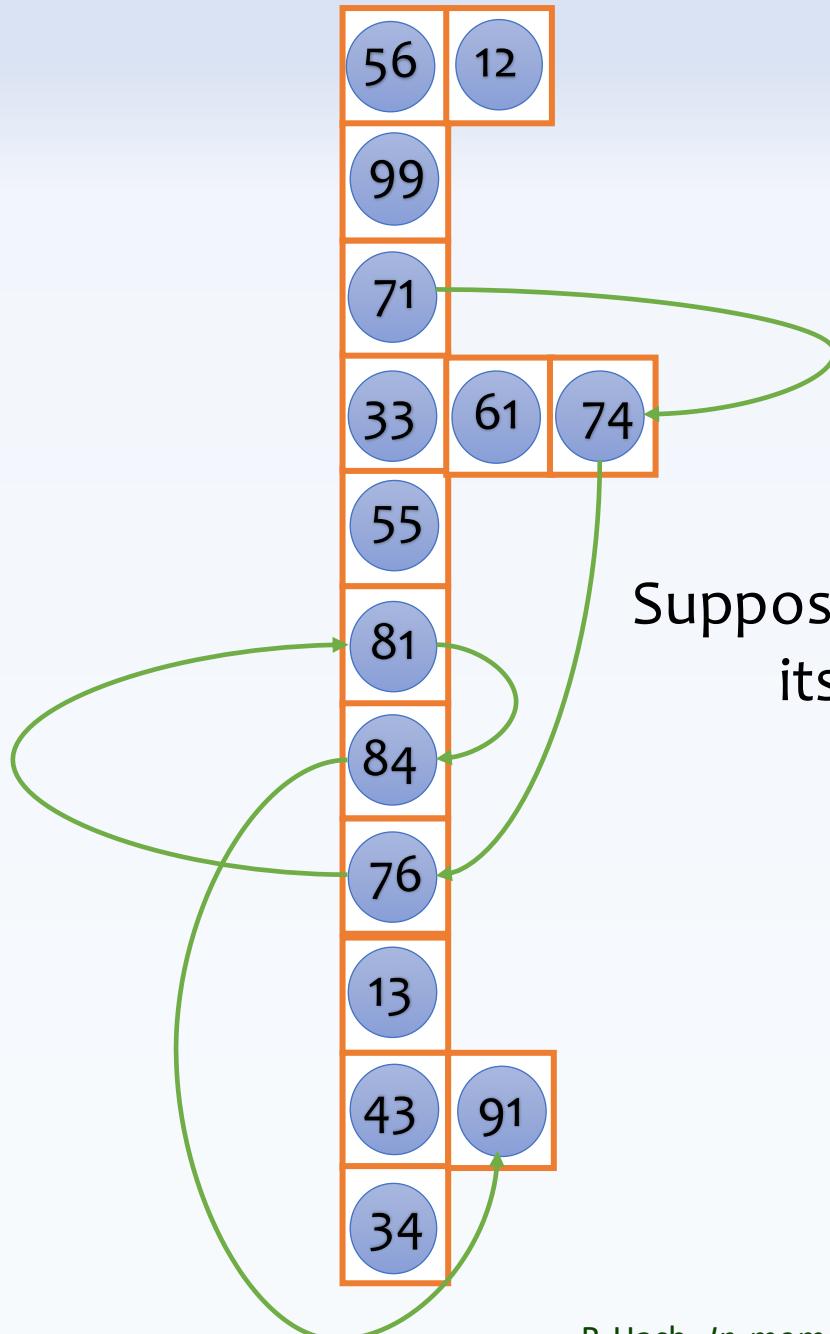
Searching for 76-91  
How about 76.01, 76.02, 76.03, ...?  
(simply not practical)

**Could we imagine a new design to support searching  
for a range of values efficiently?**



Let's promote a subset of values as seeds

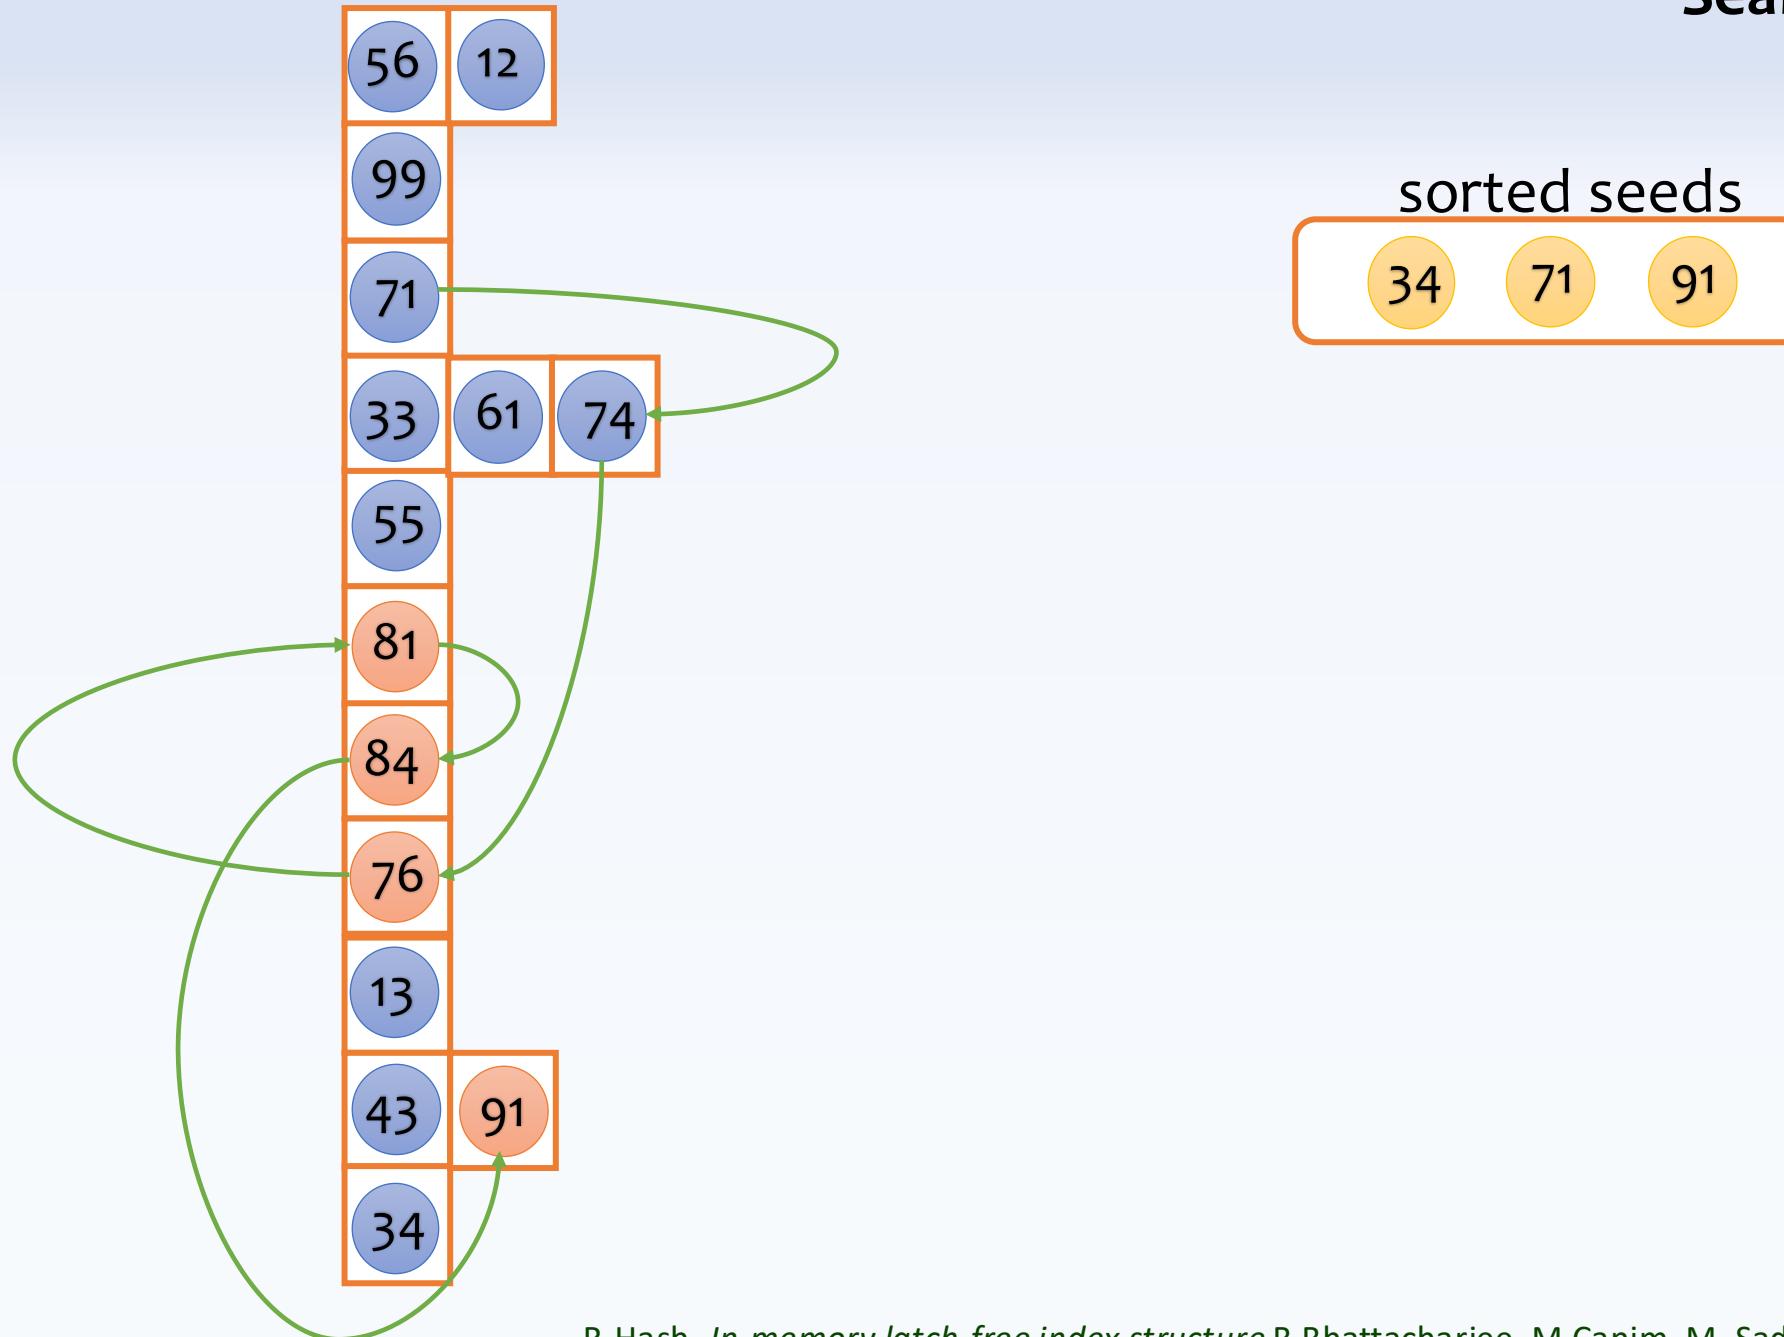




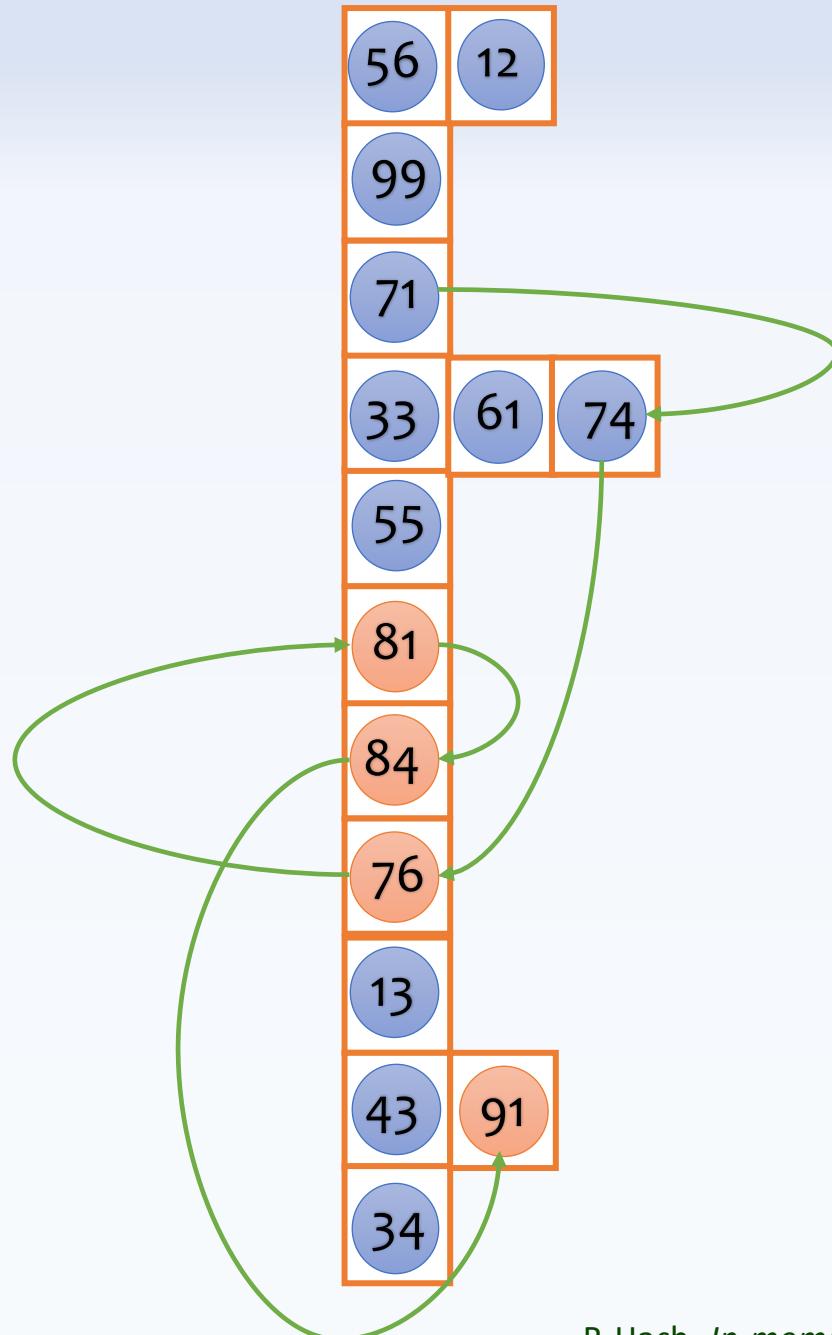
Let's promote a subset of values as seeds



Suppose every value points to  
its next larger value



# Searching for 76-91

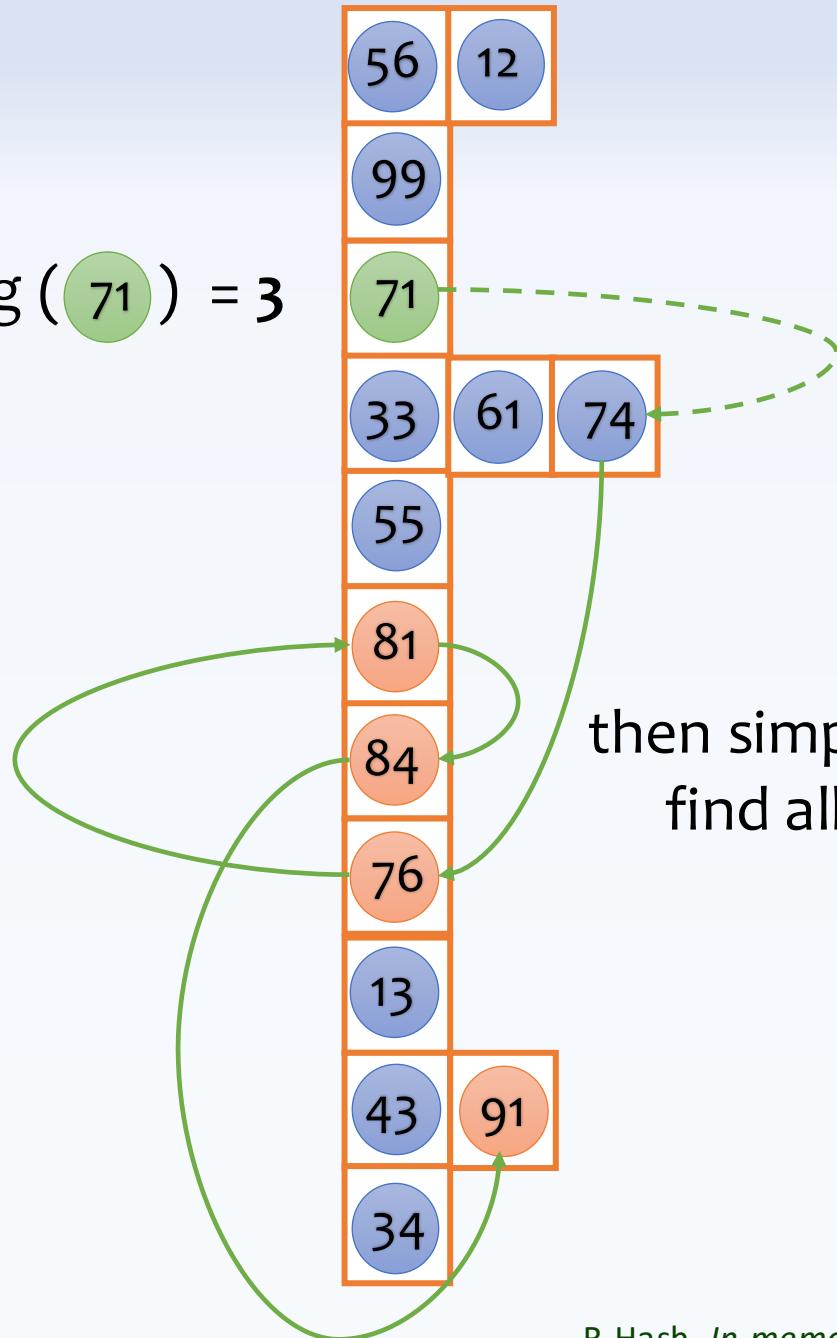


sorted seeds

34 71 91

Find the largest seed smaller than 76: 71

Hashing ( 71 ) = 3



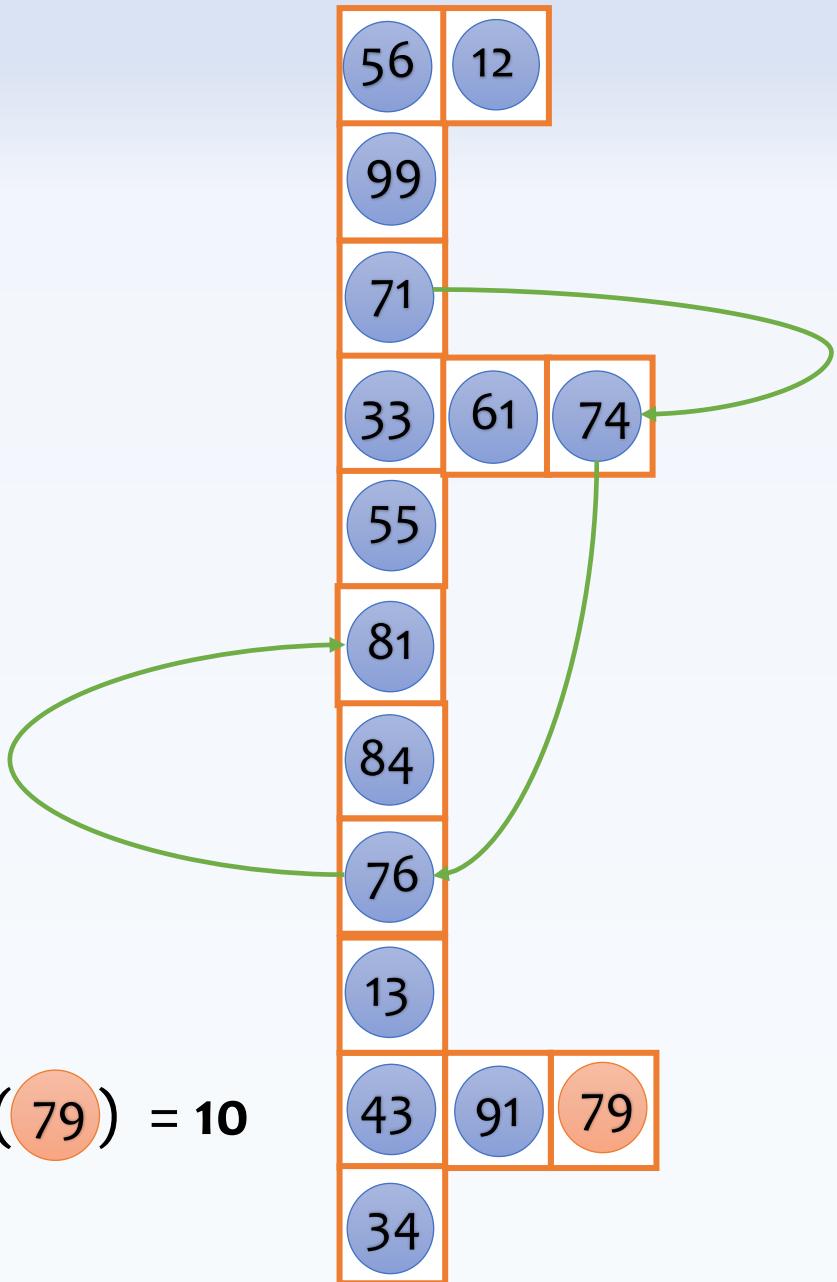
sorted seeds

34 71 91

Find the largest seed smaller than 76: 71

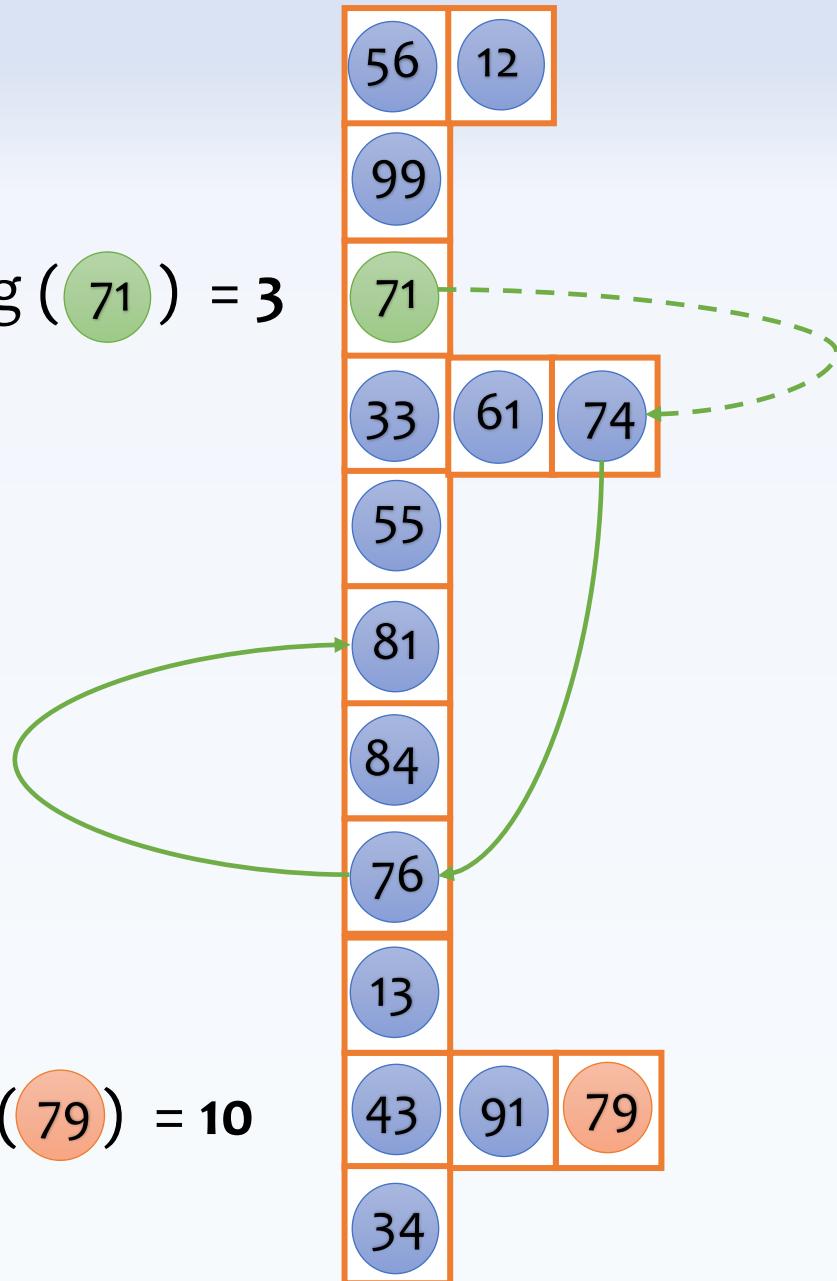
then simply follow the pointers to  
find all values between 76-91

# Inserting 79



## Inserting 79

Hashing ( 71 ) = 3



sorted seeds

34 71 91

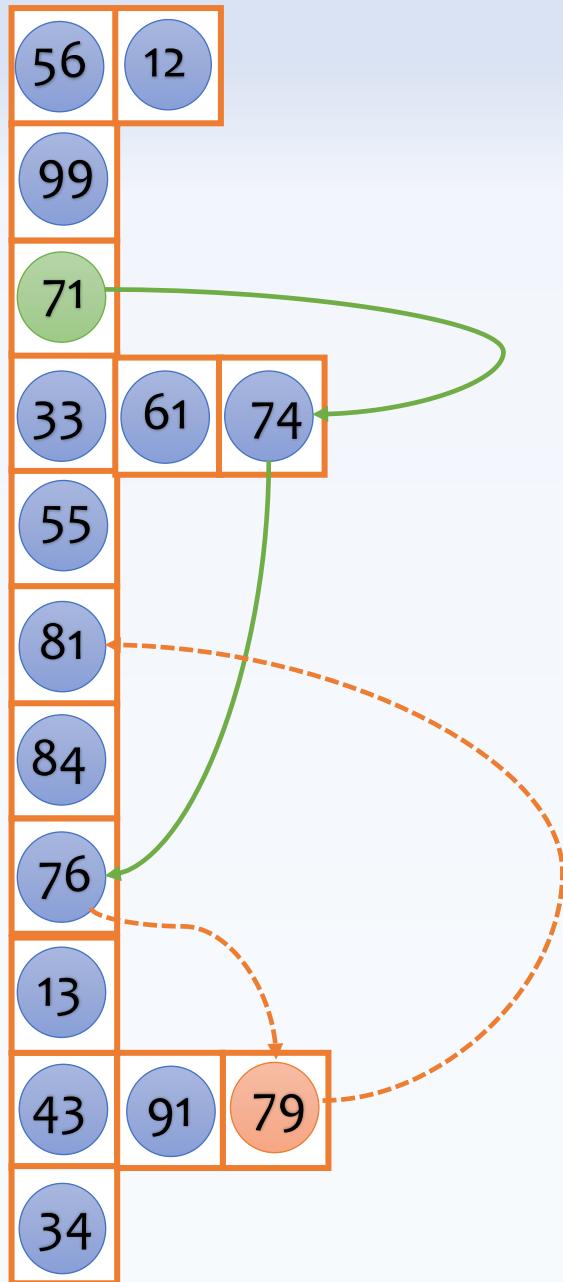
Find the largest seed smaller than 79: 71

Hashing ( 79 ) = 10

## Inserting 79

Hashing ( 71 ) = 3

Hashing ( 79 ) = 10



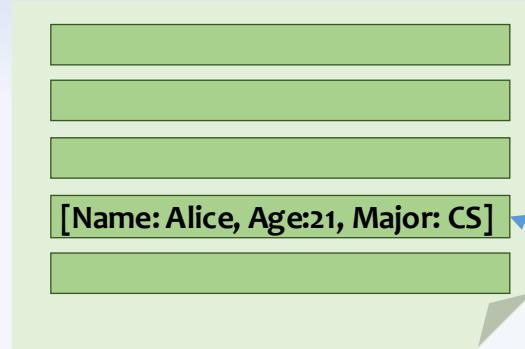
sorted seeds

Find the largest seed smaller than 79: 71

adjust the pointers accordingly

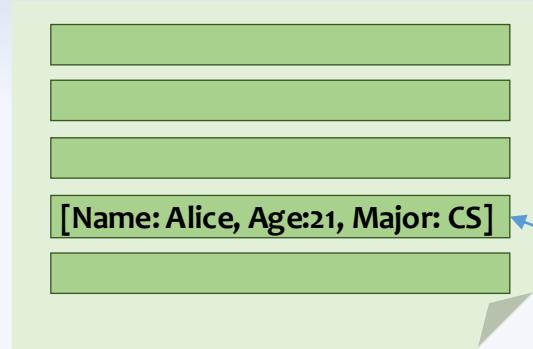
# **Database Storage Layouts**

## **(how likely that we need an index for range queries?)**



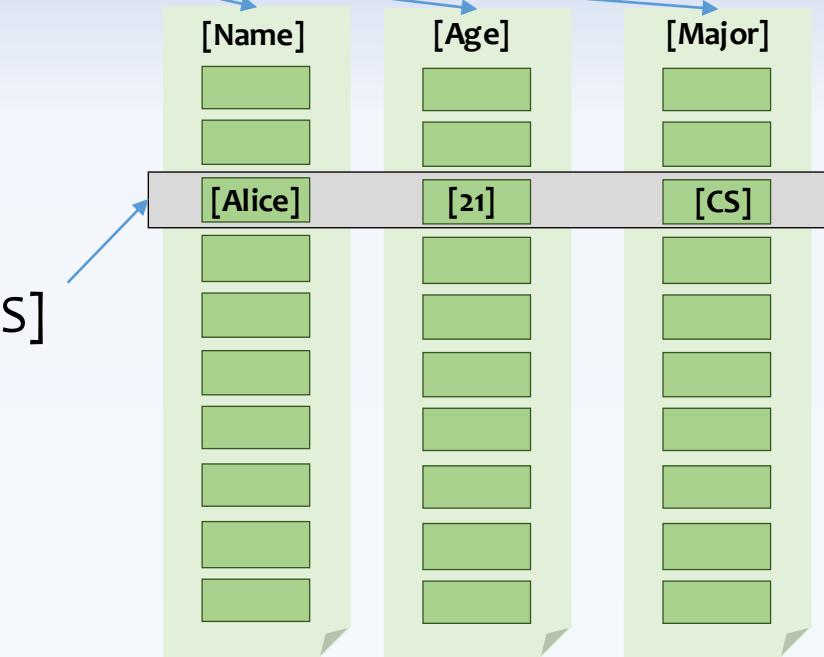
database pages  
(containing a set of records)

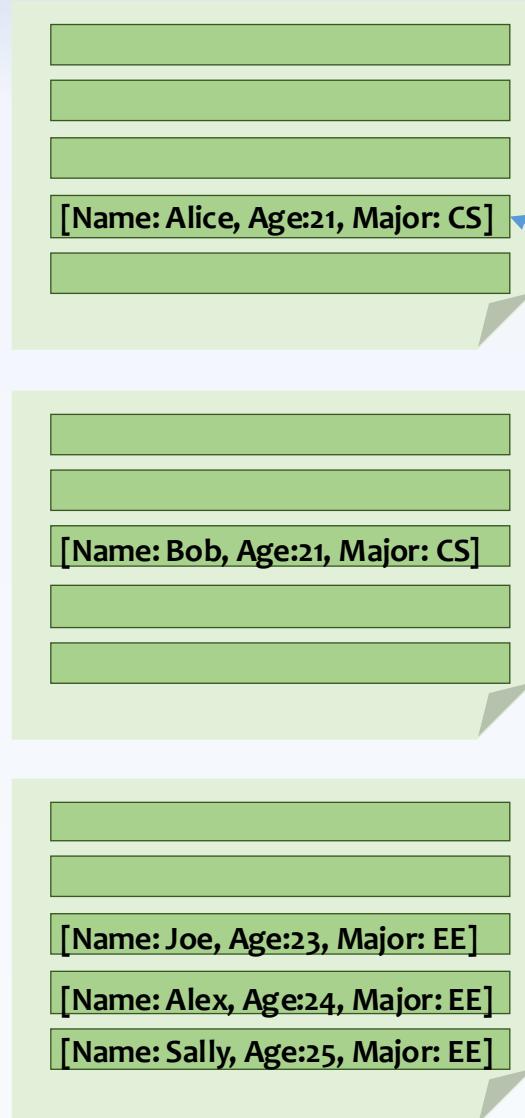
a database record, e.g.,  
[Name: Alice, Age:21, Major: CS]



database pages  
(containing a set of records)

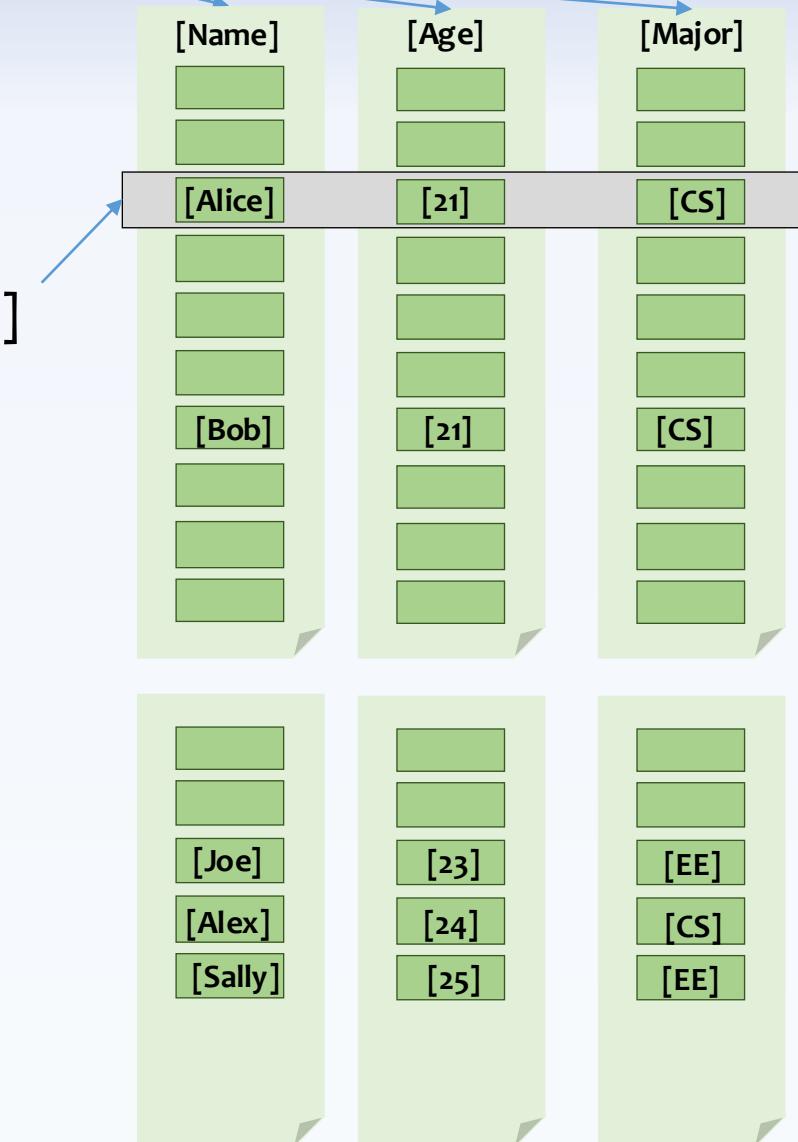
a database record, e.g.,  
[Name: Alice, Age:21, Major: CS]



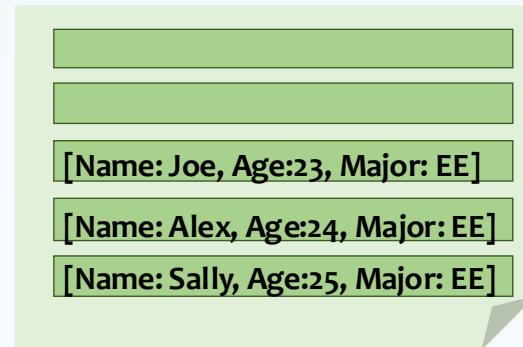
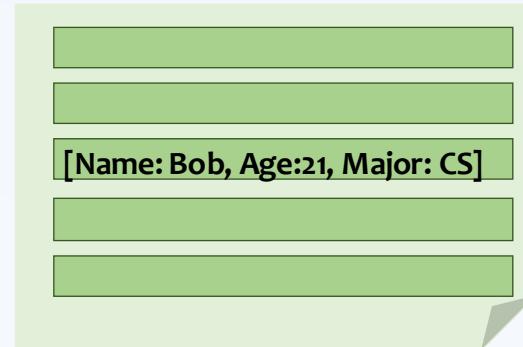
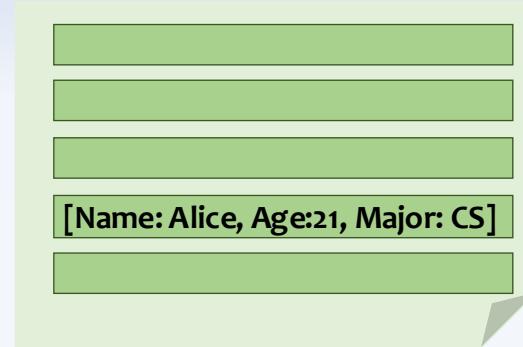


database pages  
(containing a set of records)

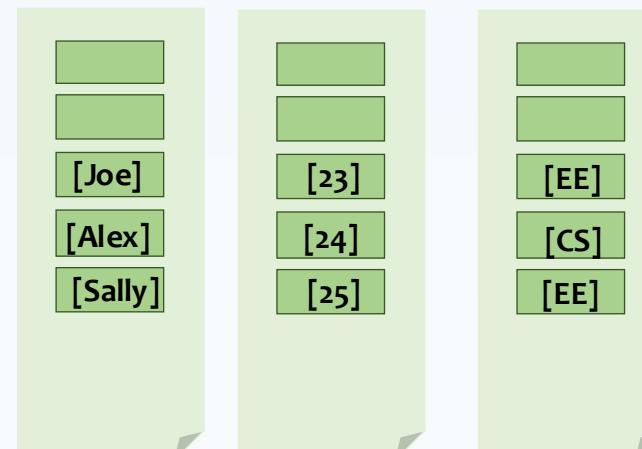
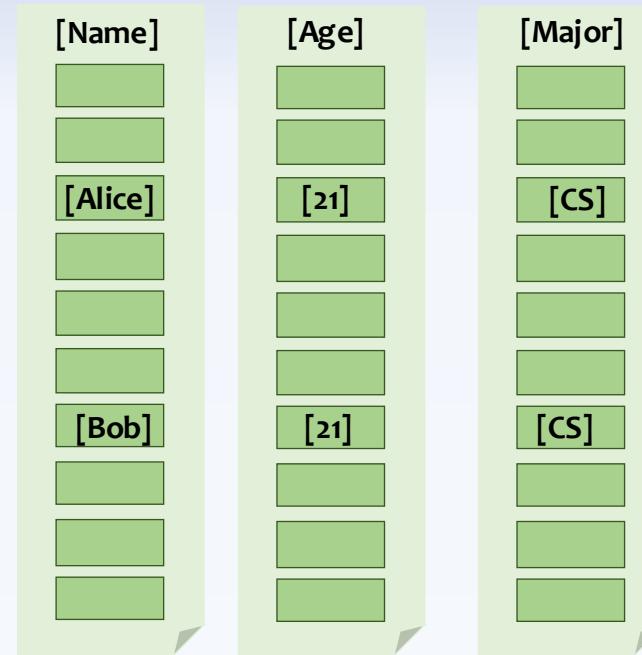
a database record, e.g.,  
[Name: Alice, Age:21, Major: CS]



# Searching for all students between the age of 21 to 24 (may return many students)

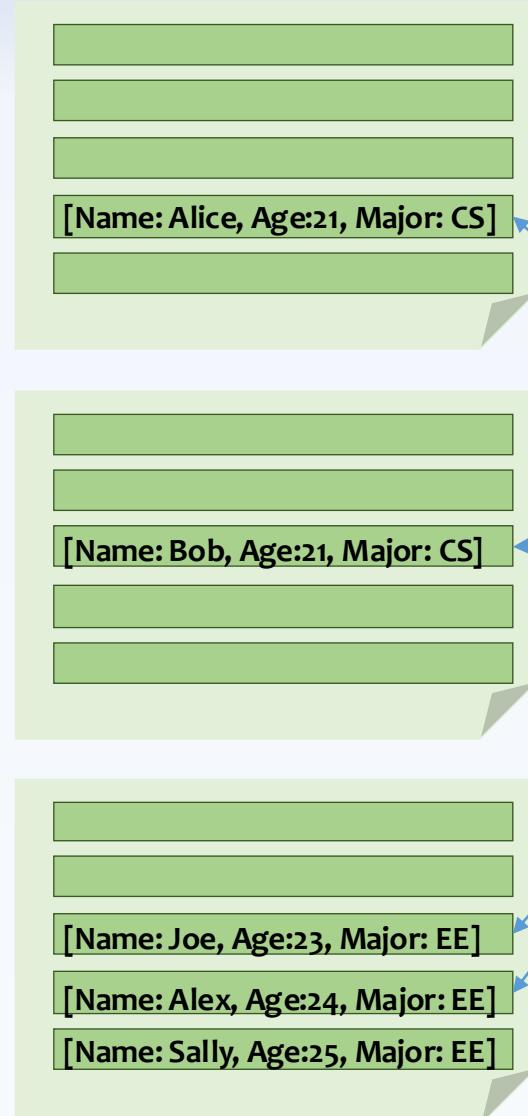


Row-based Layout



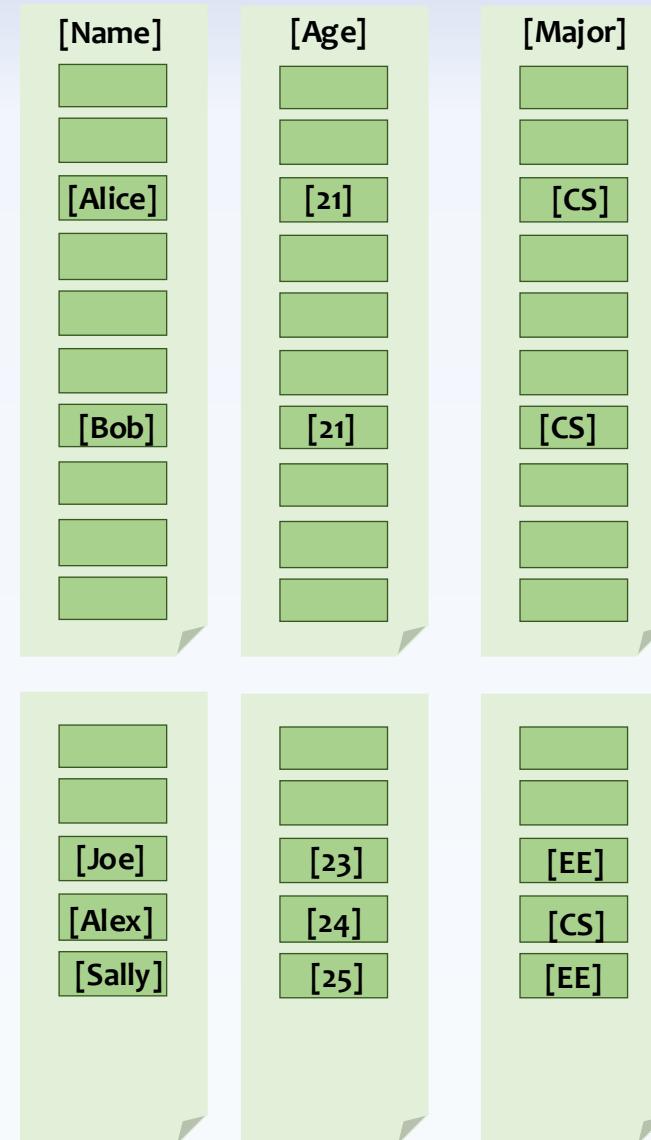
Column-based Layout

# Searching for all students between the age of 21 to 24 (may return many students)

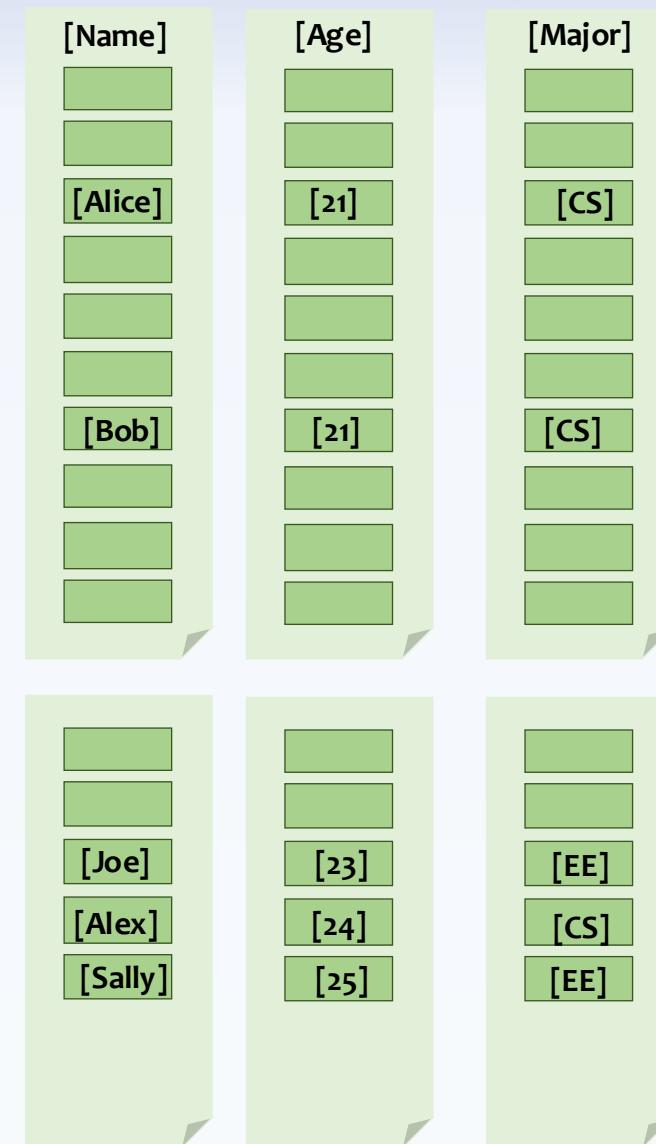
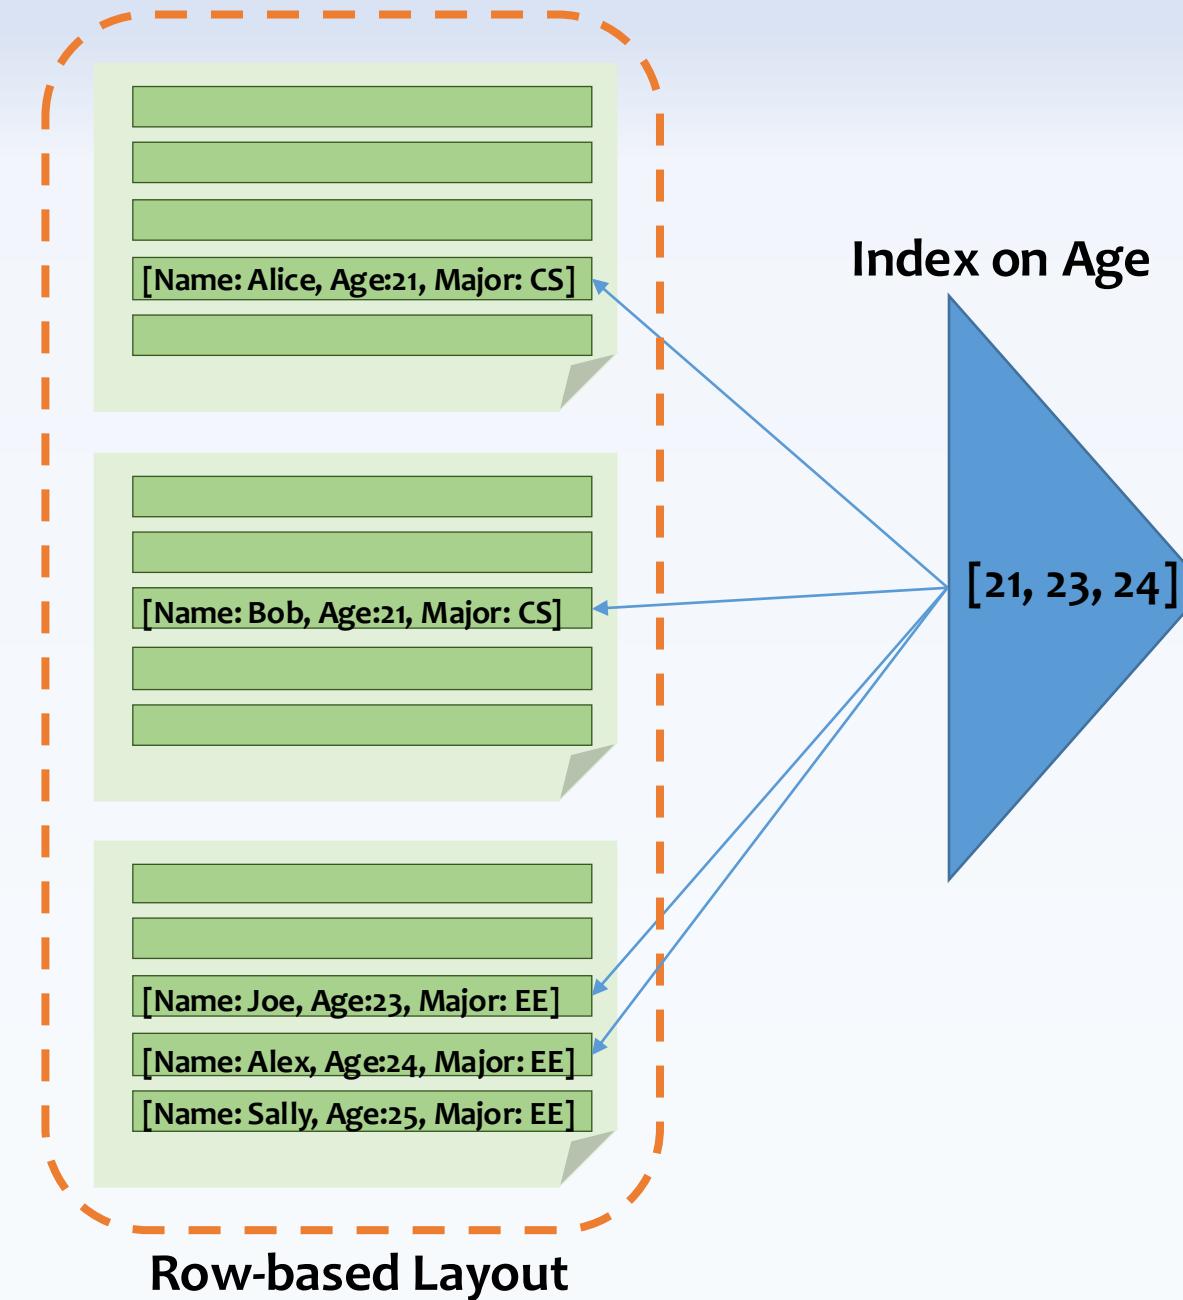


Index on Age

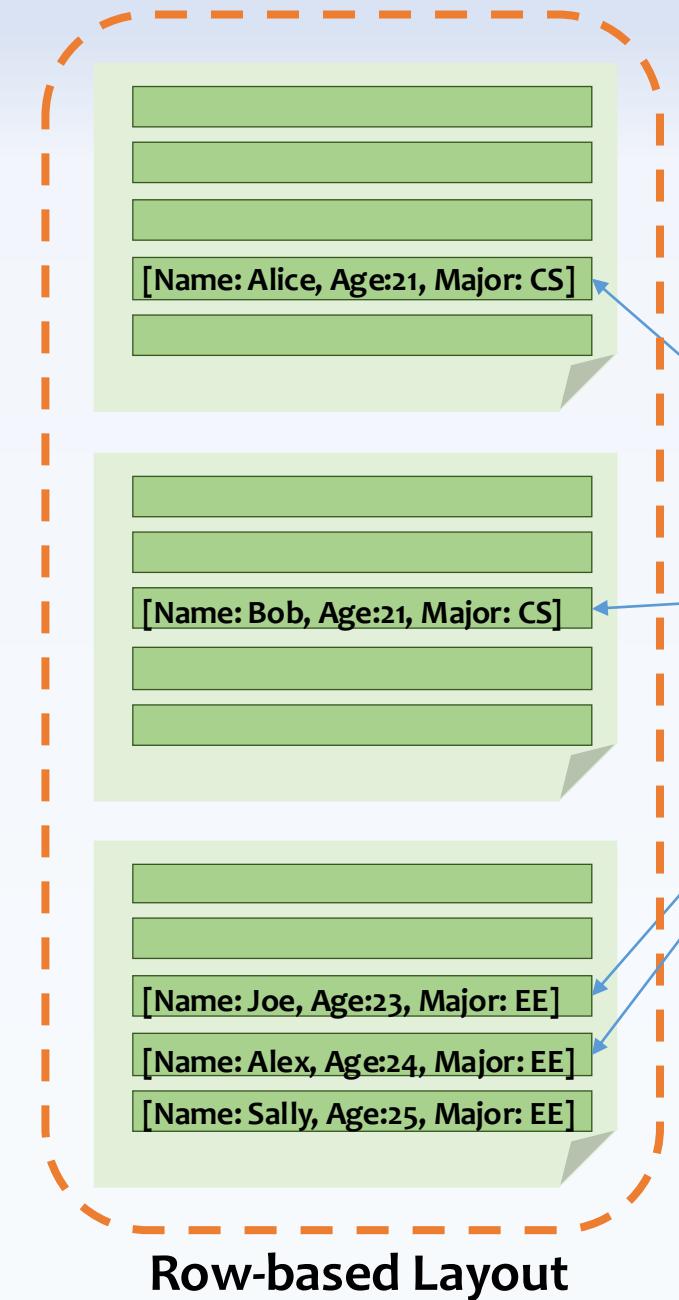
[21, 23, 24]



# Searching for all students between the age of 21 to 24 (may return many students)



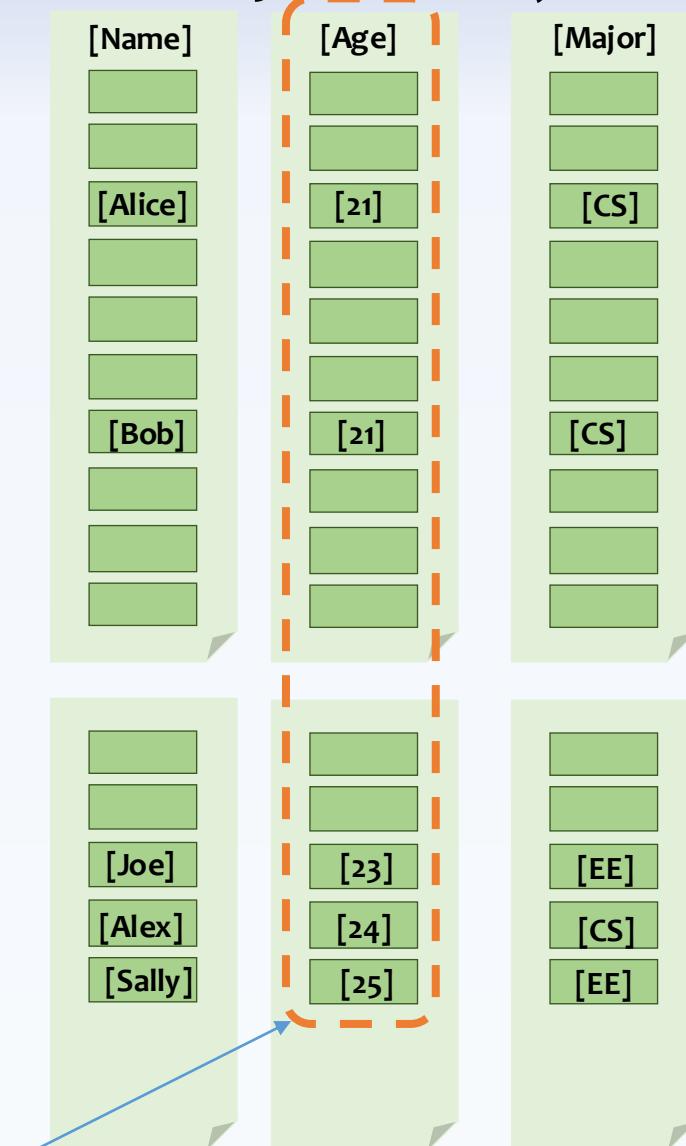
# Searching for all students between the age of 21 to 24 (may return many students)



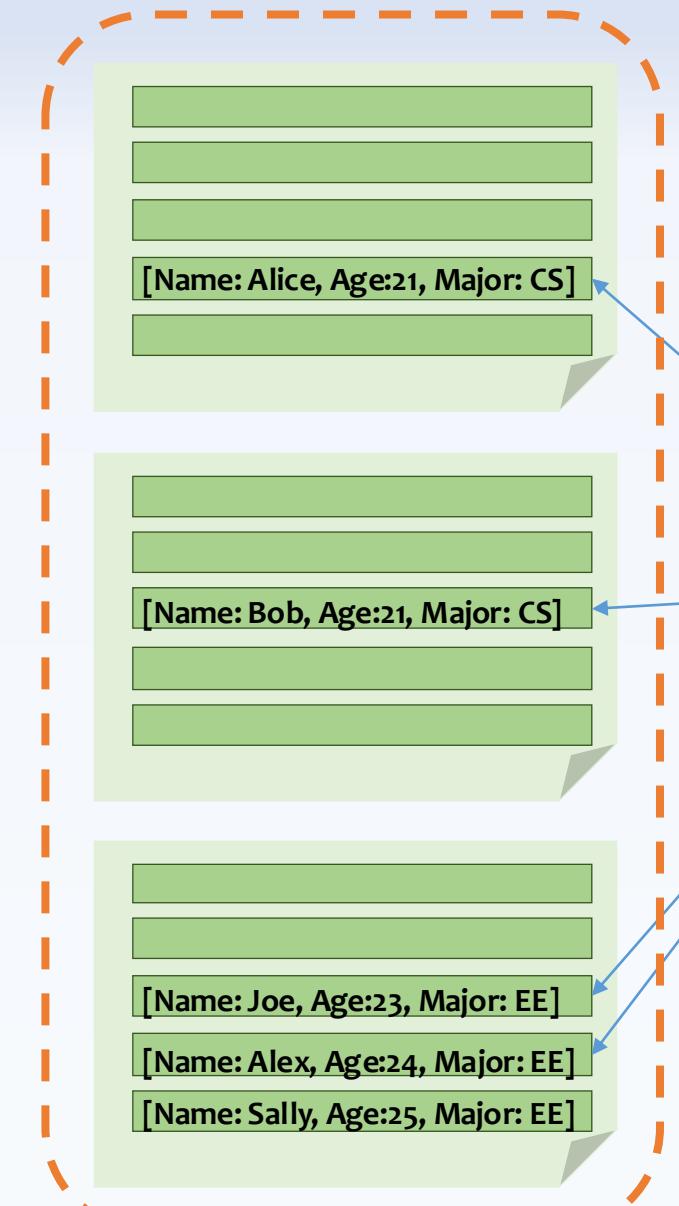
Index on Age

[21, 23, 24]

Alternatively read only the Age column to find the relevant values



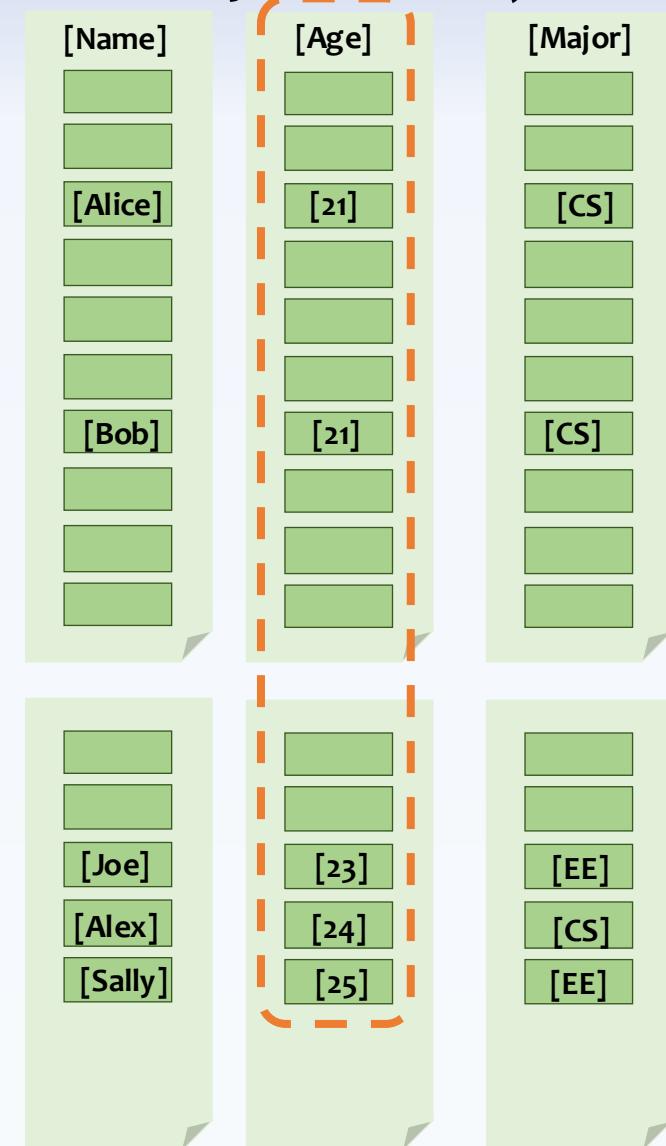
# Searching for all students between the age of 21 to 24 (may return many students)



Index on Age

Is an index really useful here?

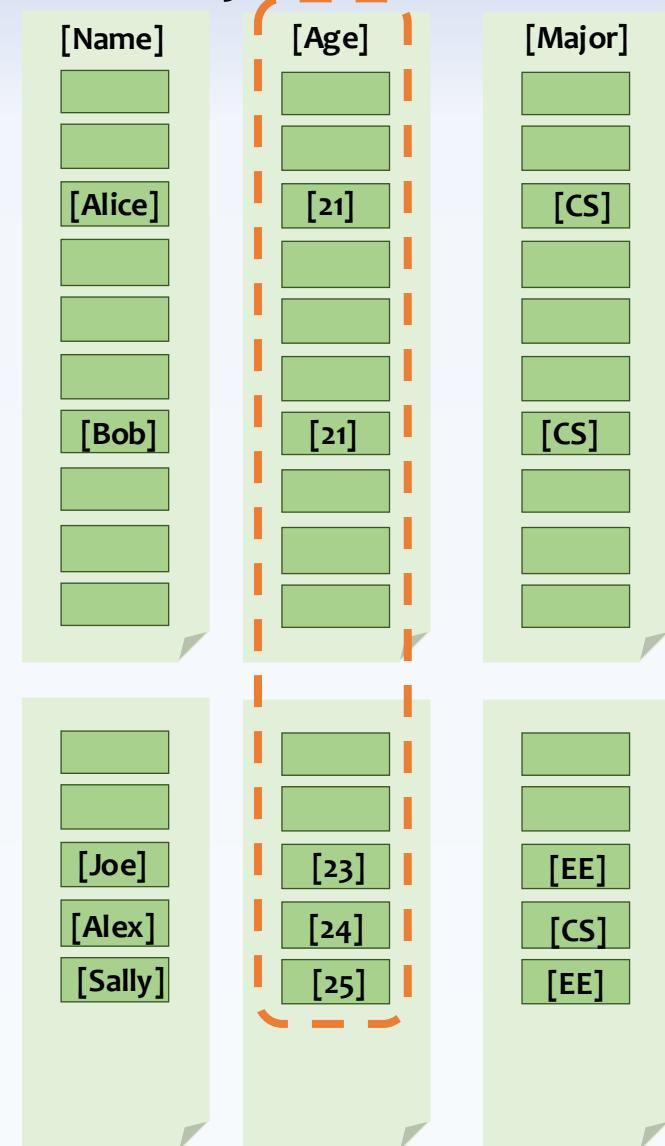
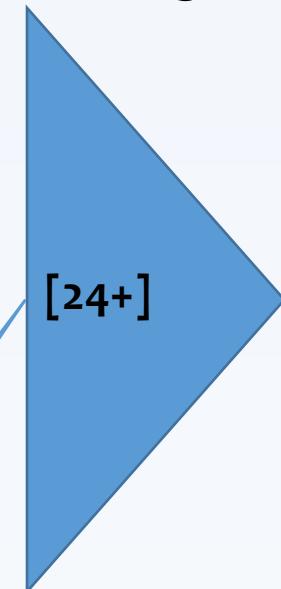
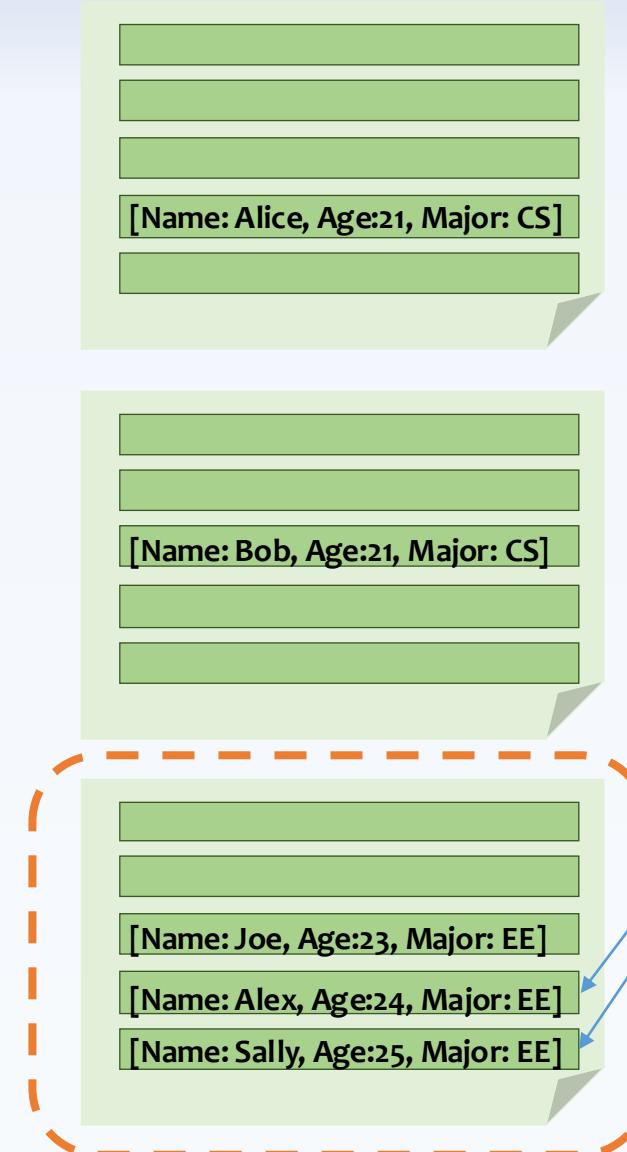
Row-based Layout



Column-based Layout

# Searching for all students over the age of 24 (may return only a few students)

Index on Age



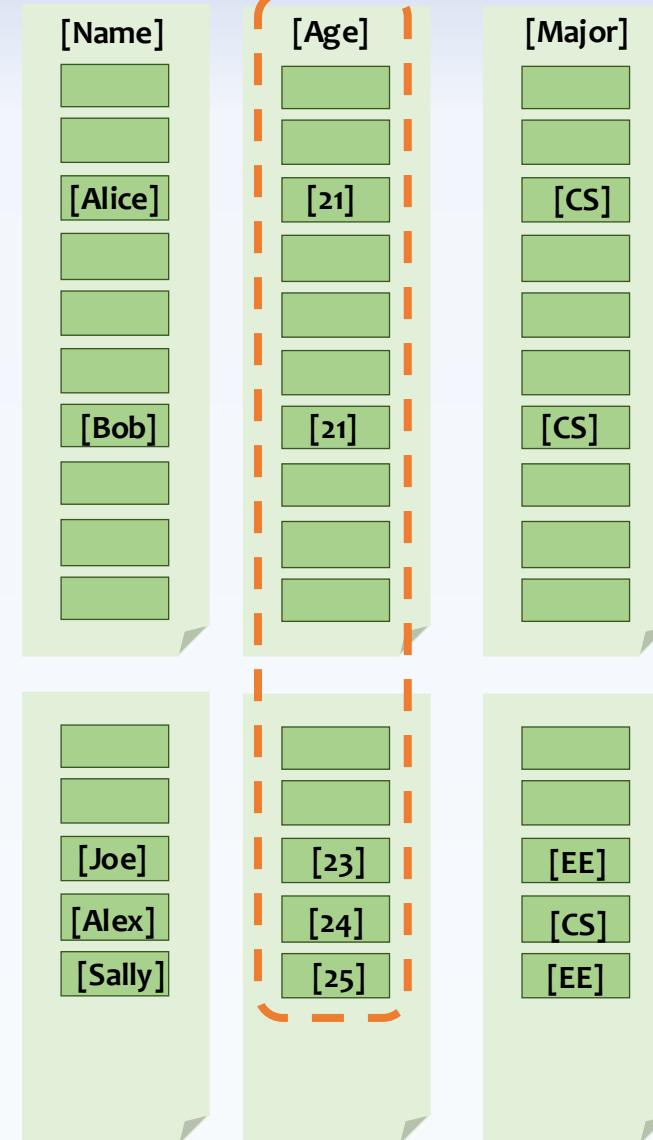
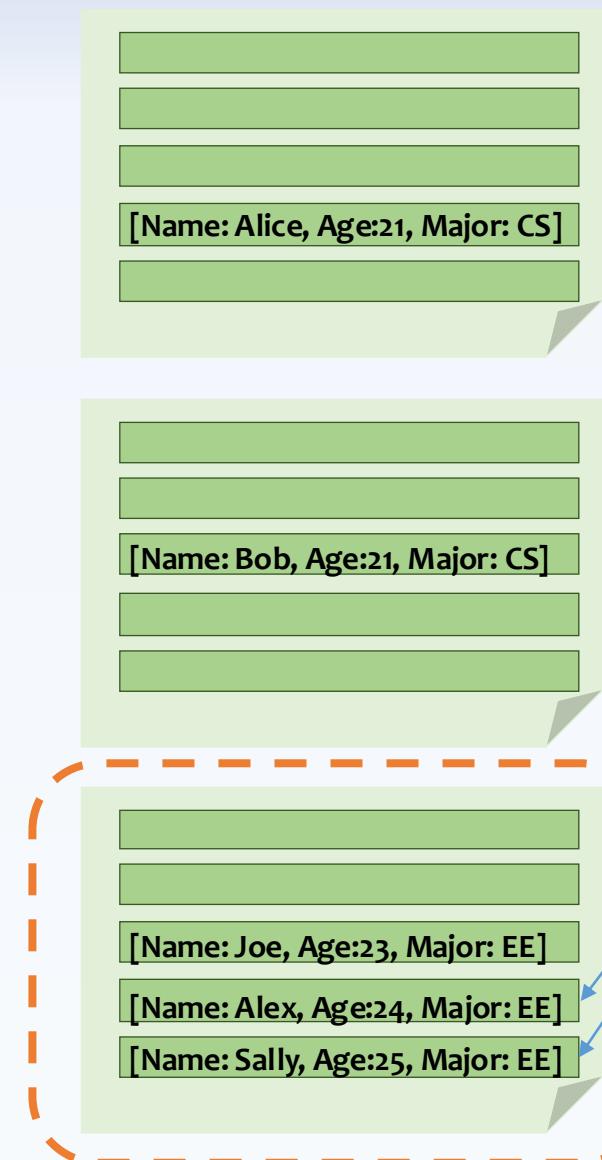
# Searching for all students over the age of 24 (may return only a few students)

Index on Age

[24+]

Could we instead employ  
hashing with the seeding idea?

Row-based Layout



Column-based Layout

**Thank You  
Questions?**