

Foundations of Bitcoin

Shreemoy Mishra, IOV Labs

About me

- Researcher at [IOVLabs.org](#)
 - RSK Blockchain: <https://www.rsk.co/>
 - EVM smart contracts 'with Bitcoin'
- Former academic economist and engineer
- IIITD, Vassar, Oberlin, UT Austin (PhD)
- Interface of econ + blockchain infrastructure

Great to be here!

Thanks to ExpoLab, UC Davis, BAF and Blockchain
at Davis

Pre-reqs for blockchain?

- None! 

Pre-reqs for blockchain?

- None! 
- But technical background helps
- Touches a lot of areas
- Willingness to learn is all that's needed

Pre-reqs for blockchain

- Databases, cryptography, comm protocols
- software engineering, security, UI/UX
- economics: incentives, game theory
- social/legal: trust, securities, money transfer
- Hence: no strict pre-reqs

Questions?

Questions?

- Remote presentations can be tricky
- If something goes off track ...
 - these slides (HackMd version)
 - <https://hackmd.io/@optimalbrew/rJnko444v>
- <https://github.com/optimalbrew/UCDavisBAF/>
 - these slides (markdown)
 - simple example (ipython notebook)

This Meeting ... I assume

- Very basic awareness of
- Cryptographic / Secure Hashing
 - **one way** functions (infeasible to invert)
 - “**unique**” (no collisions)

This Meeting ... I assume

- Very basic awareness of
- Cryptographic / Secure Hashing
 - **one way** functions (infeasible to invert)
 - "**unique**" (no collisions)
- Public Key Cryptography
 - Based on **private** and **public** keys

Secure Hashing Example

- Cryptographic (secure) Hash Functions

```
>>> import hashlib
>>> m0 = hashlib.sha256(b'this is a simple example')
>>> m0.hexdigest() #the sha256 hash in hexadecimal
'b973545e5472e4a5b7570d65467b5ec3fd5a82195d1593a9815bb18ca425
>>>
>>> m1 = hashlib.sha256(b'this is a simple example ')
>>> m1.hexdigest()
'59a0b81d2d91800f86645ba3d0ee1f7a40322733edebfc9483b24e274a82
>>>
>>> m2 = hashlib.sha256(b'this is a simple exAMPL')
>>> m2.hexdigest()
'4ca6aafe93b7ed4173421cb247119fbcfae66cca138d234f61ba3b861897
```

Secure Hashing Example

- Cryptographic (secure) Hash Functions

```
>>> m0 = hashlib.sha256(b'this is a simple example')
>>> m0.hexdigest() #the sha256 hash in hexadecimal
'b973545e5472e4a5b7570d65467b5ec3fd5a82195d1593a9815bb18ca425
>>>
```

- ? how long is a sha256 hash?

Public Key Crypto

- Based on **pair of keys**:
 - a *private key*, and a *public key*
- Uses:
 - authentication (e.g. SSH, GPG on Github)
 - **signing** (core)
 - encryption

Public Key Crypto

- Keys are very large integers
- large (pseudo) random number -> private key
- private key -> public key
- **crucial:** public key reveals nothing about private key

Public Key Crypto

- These large numbers are drawn from elaborate mathematical structures (*fields*, geometry)
- Special rules for (modular) arithmetic (+, \times), origin, *overflow* (e.g. 256 bit keys)

Digital Signatures Algo. (DSA)

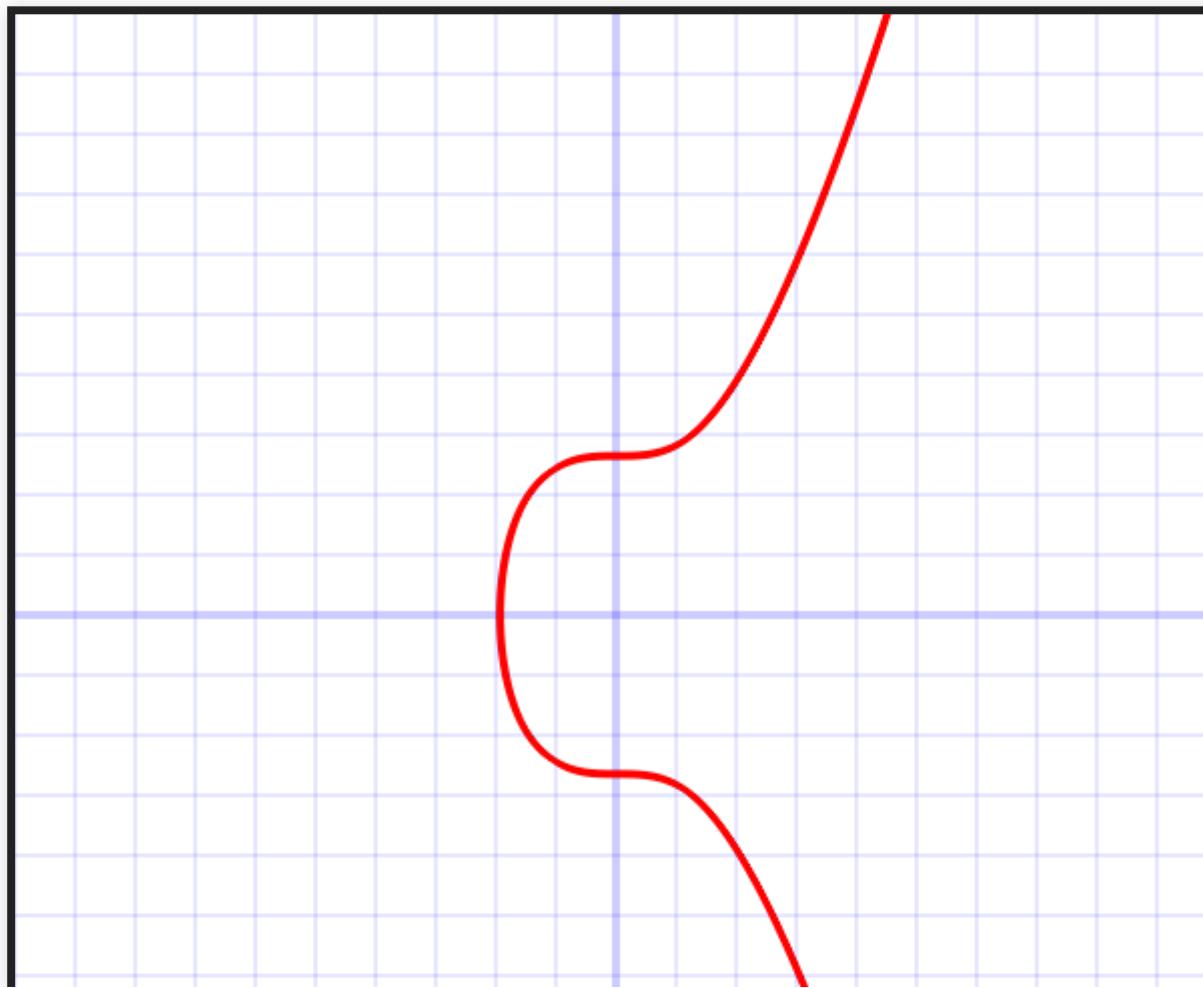
- Bitcoin uses Elliptic Curves (EC) cryptography
- ECDSA signatures consists of a **pair**:
- Typically denoted by r, s

Digital Signatures Algo. (DSA)

- Bitcoin uses Elliptic Curves (EC) cryptography
- ECDSA signatures consists of a **pair**:
- Typically denoted by r and s
- I was like what?? (the 1st time...)
- Also, in ETH and RSK: tuple $\langle v, r, s \rangle$

Public Key Crypto

- Using (simplistic) images of ECC can lead to flawed ideas





Public Key Crypto

- We won't discuss it
 - Leave it for cryptography
- Interesting aside:
 - in Bitcoin, each pvt_key -> 2 pub_keys!
 - hence 2 addresses
 - (what's an address? coming up)

Signature Example (Go module)

```
package main

import (
    "crypto/ecdsa"          // key gen, sign, verify
    "crypto/elliptic"        // specific curve families
    "crypto/rand"            // PRNG
    "crypto/sha256"
    "fmt"
)

func main() {
    //generate key from secp256k1 curve for ECDSA
    privateKey, err := ecdsa.GenerateKey(elliptic.P256(), rand.Reader)
    if err != nil { fmt.Println(err) }

    //sign message
    message := []byte("Hello World")
    signature, err := privateKey.Sign(rand.Reader, message)
    if err != nil { fmt.Println(err) }

    //verify message
    recoveredPubKey, err := ecdsa.RecoverFromSignature(privateKey, message, signature)
    if err != nil { fmt.Println(err) }

    if recoveredPubKey != privateKey.Public() {
        fmt.Println("Recovered public key does not match private key")
    }
}
```

Moving on to Bitcoin

Wiki <https://en.bitcoin.it>

- ...first successful implementation of a **distributed crypto-currency** ...
- ...described in part in 1998 by **Wei Dai**...

Wiki <https://en.bitcoin.it>

- ...first successful implementation of a **distributed crypto-currency** ...
 - ...described in part in 1998 by **Wei Dai**...
 - ...using cryptography to control creation and transfer of money
 - ...**rather than** relying on **central authorities**.
- Question:** what's a central authority?

Bitcoin origins

- Shrouded in mystery
- Started by unknown individual(s) using a pseudonym: Satoshi Nakamoto
- Continuing speculation about identity, motives
 - not our concern today 

A P2P Monetary System (how fun!)

- Physical currency  has *fixed denomination*
- E-money can have any value 10^{-8} BTC
- **Coins** are
 - chains of digitally signed messages
 - to transfer ownership
 - can have nearly arbitrary value
- Smallest unit in Bitcoin: 1 Satoshi

A P2P Monetary System

- You receive 2 coins: value 3.0, 4.0
 - you can create a coin worth 4.5 to pay someone
 - redirect 2.4 in a new coin to yourself
 - leftover ($7 - 6.9 = 0.1$) is transaction fee
- Will revisit this example in a bit

Fundamental security problem

- “**Double spending**”
- I pay you 100 coins (e.g. game/src/)
 - 1st signed message
- later, I transfer the *same* coins to myself
 - 2nd signed message
- I (try to) convince others that 1st msg **is invalid**

Fundamental security problem

- If not for this problem (double spending), e-currency is easy
- Goal:
 - Avoid multiple histories (conflict messages)
 - Agree upon a *single history* of payments

Basic jargon ... Bitcoin Address

- Addresses are generated from public key

```
Key hash = Version + RIPEMD-160(SHA-256(public key))
```

```
Checksum = 1st 4 bytes of SHA-256(SHA-256(Key hash))
```

```
Bitcoin Address = Base58Encode(Key hash concatenated with Checksum)
```

- **Not** double Sha256, checksum, and Base58

Basic jargon

- Transaction (shorthand TX)
 - a transfer of coins
 - from 1+ addresses (**TX inputs**)
 - to 1+ addresses (**TX outputs**)
- How much BTC do yo have?
- The net sum of all your *unspent* TX outputs

Basic jargon - Block

- A collection of transactions processed around the same time
- like a container class
- 1 block created approx every 10 mins
 - Who creates them?

Basic jargon - Block Chain

- **Chain** blocks together by including
 - a **pointer to previous (parent) block**
- similar to a *linked list* (single, not double)
- Done using a **secure hash** (of data in parent)

Basic jargon - Block Chain

- Chaining block hashes together...
- ... makes it practically impossible to change info without being detected.
- important, because the goal is to agree upon a single historical record of transactions

Basic jargon - Block Chain

- Chaining block hashes together...
- ... makes it practically impossible to change info without being detected.
- important, because the goal is to agree upon a single historical record of transactions

TX: recall example

- Receive 2 coins: value 3.0, 4.0
 - you can create a coin worth 4.5 to pay someone
 - redirect 2.4 in a new coin to yourself
- 2 TX inputs: (TXIs)
 - one for 3.0 and another for 4.0
- 2 TX outputs: (TXOs)
 - one for 4.25 and another 2.4

TX inputs and outputs

- $\sum TXOs \leq \sum TXIs$
- Any difference is collected by **miners** (block creator) as **fee**

Examples using block explorers

<https://www.blockchain.com/explorer>

- or any explorer
- search for block no 650750
- select BTC (bitcoin) chain (depending on explorer)
- <https://www.blockchain.com/btc/block/650750>

Checkout block info

on <https://www.blockchain.com/btc/block/650750>

- **Hash** = starts with lots of 0's
- Number of TXs = 2834, Size = 1,396,390
- **Timestamp** = 2020-09-30 20:11
- **Merkle Root** = d349...
- **Version** = 20, **Nonce** = 25,376,585
- Block reward = 6.25, Fee = 0.592 BTC
- Difficulty = 19,314,656,404,097.00, **Bits** = 386,831,018

Explorer

- ? What's average fee per TX? (in USD?)
- ? How would you find the hash of the **previous** block?
- ? what is the smallest block height/number?

Genesis: Block 0

- Find block 0 in explorer :
<https://www.blockchain.com/btc/block/0>
- Look at Inputs of Coinbase TC
- Sigsript (some of you know what's coming)
 - paste the numbers starting from 5468 ... into Hex to Ascii converter
- <https://www.rapidtables.com/convert/number/hex-to-ascii.html>

Time to meet Satoshi
(just the paper)

“The” Bitcoin Paper

A Peer-to-Peer Electronic Cash System

<https://bitcoin.org/en/bitcoin-paper>

Abstract

- **Motivation:** Purely p2p e-cash payments
- Without banks or financial intermediaries
- **No need to trust** anyone

Abstract

- Digital signatures are necessary...
- but not sufficient ...
- because of the possibility of *cheating* (double spending problem)

Abstract

- **Solution:** transaction timestamps ... so only the 1st one counts... rest are discarded
- **?** how to reach **consensus** on timestamp when (block) data can arrive in different order?

Agreement on TX orders

- Timestamping TXs to agree on a system for **single history**
- Create **cryptographic links between timestamps**
- Make them impractical to falsify

Agreement on TX orders

- Timestamping TXs to agree on a system for **single history**
- Create **cryptographic links between timestamps**
- Make them impractical to falsify
- (Note: impractical... not impossible)

Agreement on TX orders

- This record (linked hashes) cannot be changed (easily)
- The **difficulty of re-writing history** is what **blockchain security** is all about

Blockchain can be modified

- But to modify the record
- ... one **needs to overcome** all the associated computational work that went into constructing the chain ...
- ... since the timepoint from which they wish to rewrite it
- Such changes are called **chain reorgs**

Wanna re-write history?

- Suppose you want to change a **single block** from Oct 2019.
- You have to redo the **entire chain**... from that point on.

Wanna re-write history?

- Currently it costs \$10 million a day to create blocks (entire ecosystem) ...
- You need to burn through at least 1/2 of that... and then some
- The math is not quite right... (I am skipping stuff)

re-writing history

- But even if you do that, all you can do is double spend some of your own coins.
- Is the cost worth it?

re-writing history

- But even if you do that, all you can do is double spend some of your own coins.
- Is the cost worth it?
- Not for BTC ... but possible for other blockchains (ETClassic)

Majority is tied to longest chain length

- As long as **the majority** of computational power **is not attacking** the network...
- ... the honest nodes will (probabilistically) outpace any attackers...
- ... and generate the **longest chain** (or **“heaviest”**)
- Satoshi provides examples w/ C code

nodes come, nodes leave

- Nodes can go offline and rejoin at will...
- ... and accept the longest PoW chain
- ... as evidence of what happened while they were gone.

Transactions

- *Define a bitcoin* (or any coin) as - **a chain of digital signatures**
- Each owner transfers the coin by
 1. using their private key to sign (a hash) of **the previous transaction**
 - (where they got the coin from!)
 2. and **add public key** of next owner

Transactions

- To clarify... you don't have to do anything to receive coins
- Conditions are on spending

Transactions

- Anyone can verify the signatures and chain of ownership
- **Main problem: double spending**
 - recipient cannot verify that the signed coin has not been spent before...

Double spending

There should be no previous transaction

- with the same coin
- that someone is offering you now.

Double spending

- The only way to be sure ...
- is to keep a record of ALL PAST transactions.
- which is why running a fully synced bitcoin node offers max protection
- but not practical for most people

Timestamp server

- Satoshi's solution? **A p2p time server**
 1. Takes a block of TXs to be timestamped
 2. computes **a hash** of data from block
 3. widely distributes it across the network

Timestamp server

- This **block hash** helps proves that the transaction was known when timestamped
- Proves ? (coming up... Merkle Trees)

Timestamp server

- Each timestamp (hash) **also includes a hash of the previous timestamp (hash)**
- ... re-inforcing ones before it... forming a chain...

Timestamp server

- BTW: the hash *is* the timestamp
- It is okay to use block number or height (unique) to talk about a TX's timestamp...
- But no one uses the date-time format

Questions before we move to some implementation details?

Coming up... proof of work

Implementing the system: POW

- Creating a block requires **guessing** a value (**the nonce**)
- which when combined with other data
- and **then hashed...**
- gives a number that **begins with a certain number** of 0's

Block Hash (double hash)

- PoW is a double-hash sha256(sha256())
- But what data are we hashing?
- **6 fields** from the block:
 - 3 integers: **version**, **bits**, **nonce**
 - 2 hashes: **previous block**, TX Merkle root
(explain later)
 - The block's **time stamp** (UTC)
- Apart from nonce: limited degrees of freedom

Implementing the system: POW

- Nonce is 4 bytes: 2^{32} combinations, approx 4Billion
- Sha256? 2^{256} combinations
- The computational work is **exponential in the number of zeros...**

Implementing the system: POW

- finding a solution is very hard... lots of combinations
 - verifying a valid solution is trivial!
- It is like playing the lottery or mining for gold

Block Hash (double hash): format

- **Order:** version, prev_Hash, Merkle_Root, timestamp, bits, nonce
- Format
 - **hexstring + big endian**
 - concatenate (order)
- After double sha -> convert back to **little** endian
- Let's go through an example

Example computation

depends on time...

Implementing the system: PoW

- PoW is how “**someone**” gets to create the next block
- and collect rewards...

Implementing the system: PoW

- PoW is how “**someone**” gets to create the next block
- and collect rewards...
- but PoW is also about reaching consensus!

Implementing the system: PoW

- PoW is also how **everyone** can reach agreement (**consensus**) ...
- i.e. that a solution has been found...
- and move on to working on the next block
- ... not everyone *has to* agree... but if a majority do, then that chain gets longer

Implementing the system: POW

- This resembles a majority voting system
 - where **1 hash computation equals 1 vote**
 - (in Satoshi's words, one-CPU-one-vote)

Implementing the system: PoW

- All decentralized blockchains have some voting mechanism...
 - next to PoW, the most well known is Proof of Stake...
 - where the vote is based on staking coins...
 - each coin staked is one vote (See Eth2.0)

Implementing the system: POW

- HOW do miners “vote” for a solution?
- by starting to work on the next block
- AND pointing to the accepted solution (block hash) - as the parent

Implementing the system: POW

- How many starting 0's are enough?
- that depends on a parameter called **difficulty**
- difficulty is readjusted periodically
- using an algorithm that tracks a moving average of the last 2K blocks
- and keep that average around 10 minutes.

Implementing the system: PoW

- In contrast... RSK block time is about 30 seconds.
- RSK is tied to BTC (security via **merged mining**)
 - something you may learn later
- RSK difficulty lower... only 1 in 20 RSK PoW meets Bitcoin's difficulty

Money creation! Block rewards

- When a miner creates a new block, they earn a new coin
 - 6.25 BTC (block reward)
 - TX fees (variable)
- the reward coin is implemented as a special transaction
 - called the **coinbase** transaction
 - By convention, the 1st TX (position 0) in the block

Money Supply

- Rate of new coin creation is called *inflation*
 - different from economics
 - Econ: inflation (π) is about price level
- In BTC, the number started with 50 BTC per block
 - reduced by 1/2 every 4 years (210K blocks)
- Using geometric series formula
 - theoretical max of 21m BTC.

Block rewards ... miner incentives

- Recall the basic security problem...
- **if someone has majority compute power**
 - they can spend their own bitcoins more than once...
 - if this happens repeatedly... bitcoin loses value
 - miners will lose too

Block rewards ... miner incentives

- However, if they have that much power,
 - they are better off just mining new blocks.
 - reward is an important for incentives for honest behavior.

Block rewards ... miner incentives

- and it is worth repeating...
- even if someone has all the majority hash power...
 - they cannot steal other people's coins ... why ?

Stop for a sec for Q's

(next stop... verifying your payment (TX) was included in a block)

- Given:
 - the TX's hash
 - and the block number where it was included

Simplified Payment Verification

- How to verify that a TX was processed without running a full node.
- this is important, for example, for smart phone based apps and wallets

Simplified Payment Verification

- Two main ideas:
- store only the **block headers**
 - (80-100 bytes each)
- Use **Merkle** proofs

Merkle Trees

- Binary **hash** trees
- Leaves are TX hashes
 - TXs are identified by their hashes
- Each parent is the hash of its children
 $\text{sha256}(\text{sha256}(L + R))$
- Odd number?? (no right child?)
 - repeat the last one
- Repeat until we get to root
 - **root included in Block header**

Merkle Trees

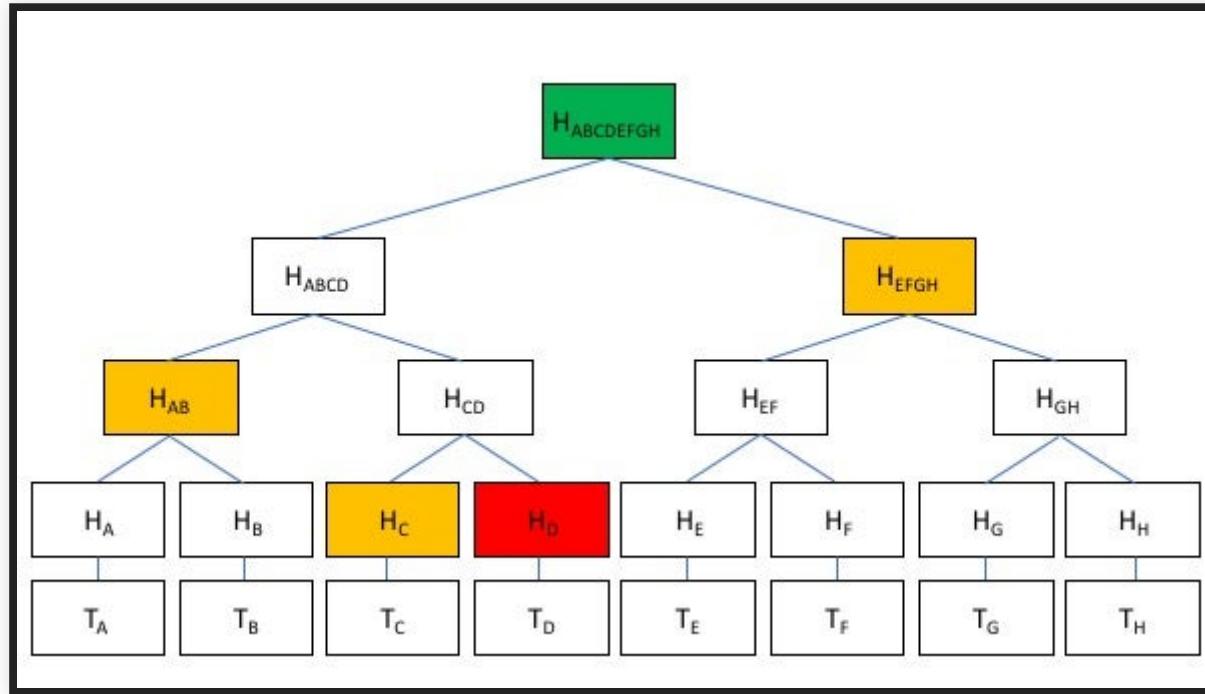
- Example: Coinbase TX + 2 user TXs:
 - Hashes: H0, H1, H2
 - Just 3 leaves in this binary tree (will have to repeat one)
 - Hashes: H0, H1, H2, H2 (again)

Merkle Trees

- parent of first two
 - $L = \text{sha256}(\text{sha256}(H_0+H_1))$
- parent of next 'two':
 - $R = \text{sha256}(\text{sha256}(H_2+H_2))$
- One level up? just the root in this example
 - `merkle_root =`
 $\text{sha256}(\text{sha256}(L+R))$

Merkle Tree and TX verification

-



- Image source: [investopedia.com](https://www.investopedia.com/terms/m/merkle-tree-1500000.html) (hence no T_0 , ..., T_i , T_{i+1} , ...)

SPV

- On top of the “Merkle proof” of TX inclusion,
 - we also wait a certain number of blocks to be added on top of the one where our TX.
 - often 6 blocks...
 - less for small amounts.

Merkle Trees

-  What if a block has no TX? e.g. block 0, 1?
- Someone mined a block with no TXs
- What's the Merkle root?

Merkle Trees

- Merkle trees are used for other things too...
 - distributed file storage (IPFS, Swarm, libtorrent)

Proof of Work illustration

(time check)

Network: Tie breaking

- A node in Japan may receive a block from China faster than a block from Ireland
- more than 1 valid block ... which one to vote for?
- **Question** what do you think happens?
- the program starts working on the next block, and points to whichever valid solution it received first.
- But it does NOT throw away the other one.

Network: Tie breaking

- It starts a fork...
- as more blocks come in... some will point to one and some to the other...
- in time... one of the forks will grow “longer” (based on majority)
- when a node switches to a new fork (because it has more cumulative PoW), that's called a **reorganization** or *reorg*

Network: Tie breaking

- such short-lived hard forks happen all the time
- Which is why, depending on the sums involved, the recipient of a transaction should wait a certain number of blocks before treating it as final.
- A common rule of thumb is 6 blocks... or about 55 minutes
- For small sums, you may choose to wait 10-15 minutes.
- If you are paying ... it does not matter :)

Bitcoin mining involves

- **Reapeatedly compute hashes** of ...
- a block header template + nonce
- until resulting hash is *smaller* than a **target**
- **before** someone else does (a race)

How is target selected?

- Probability a Hash satisfies the target

$$\frac{1}{D \times 2^{32}}$$

- D is difficulty (was 1 when it started)
- why 32? 32 comes from 32 bytes, i.e. 256 bits
- Readjusted periodically (2 weeks)

Payoff from mining

With simple assumptions

- constant hashrate h
- mining for time t
- constant difficulty D
- will find $ht/(D * 2^{32})$ blocks

Payoff from mining

- earn reward B per block
 $htB/(D * 2^{32})$
- neglecting fees

Example

Pooled mining

- Very high variance in solo mining
- Income smoothing via pooled mining
- Shares: used for accounting
- Shares are hashes that meet a lower target
(hence more frequent)

Professional Mining

- Bitcoin mining
 - started CPU, then GPU, then FPGA + ASICs
- <https://bitmain.com/product/antminer-s19-pro-110th-s/> 3250W

Professional Mining

- ASIC: <https://bitmain.com/product/antminer-s19-pro-110th-s/> 3250W
- like 2-3 space heaters on full power
- weighs about 35 lbs
- \$ 2500 purchase +
- electricity cost: about \$10/day in Seattle
- 110TH (Terahash)
- BTC n/w about 150EH (ExaHash) [btw, tera->peta->Exa]

Mining Costs (energy)

- So one S19 pro is less than $1/10^6$ network hashing power
- A probability of less than 1 in a million
- Alternate view: \$10M a day in elec costs
- $24 \text{ hours} * 6 \text{ blocks} = 144 \text{ blocks} * 6.25 \text{ BTC} * \$11K = \$9.9M$
- so our energy estimate is somewhat higher, but same order of magnitude.
- Also: neglected TX fees in payoff
- Also: Hashrate is not constant... it is a market ... miners come online and go offline based price of BTC and electricity.

Spending scripts

- In explorers, you will often see scripts for spending UTXOs
- OP_HASH160 OP_DUP OP_EQUALVERIFY
<PUBKEYHASH> OP_CHECKSIG

Spending coins P2PKH

- Pay to public key hash
- when redeeming / spending a UTXO...
 - provide PubKey
 - also provide signature (which will be matched with pubkey)
 - don't use same pubkey again (wallets do this automatically)
- PKH is shorter than PK (orig script was P2PK)

P2SH

- Pay to Script Hash
- The recipient provides the redeem script
- Address starts with 3 (instead of 1)
- As sender, you know nothing about how the address is controlled
- How to spend?
 - provide a script whose hash matches the **ScriptHash**
 - provide data (for computing) so the script evaluates to **True**
 - look for examples online

Information sources

