

Narwhal and Tusk: A DAG-based Mempool and Efficient BFT Consensus



Presented by:

Aditya Kumar Bej, Madhumitha Santhanakrishnan, Shreya Gundu and Shristi Suman

The slides are closely adapted from Alberto Sonnino's slides

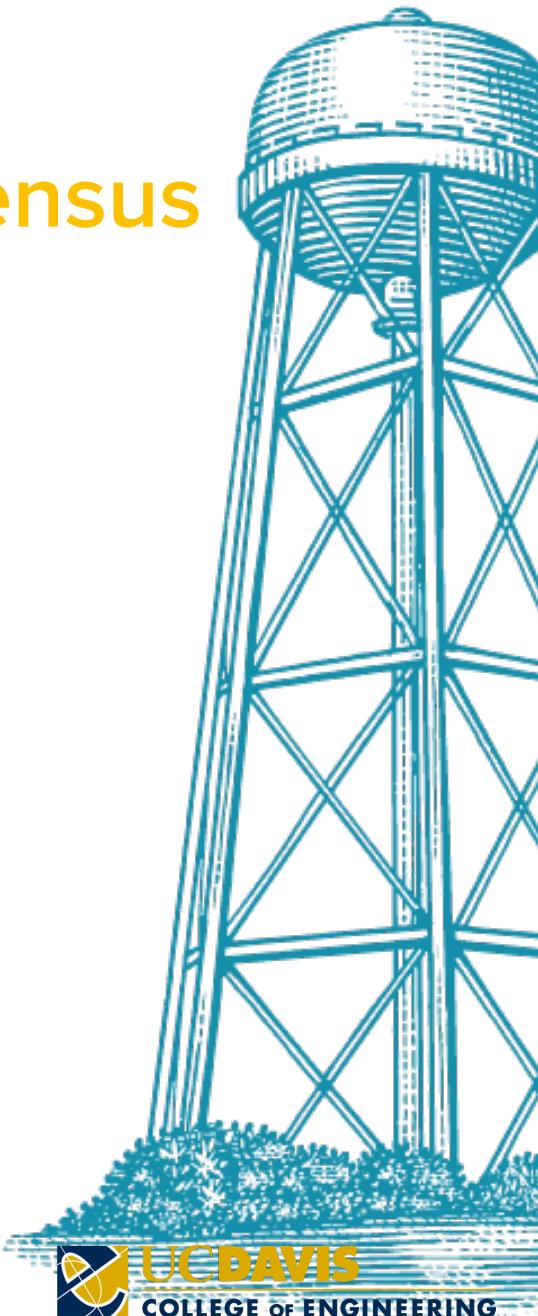


Table Of Contents

- Goal of the paper
- DAG and Mempool 101
- Risks in a Blockchain Transaction
- Narwhal
- Tusk
- Narwhal-Hotstuff
- Implementation and Metrics
- References

Goal of the Paper

To build (really) high performance blockchains

Goal of the Paper

To build (really) high performance blockchains

But how?

INTRODUCING NARWHAL AND TUSK

Narwhal

A Directed Acyclic Graph (DAG) based
Mempool protocol

Tusk

Zero-message overhead asynchronous consensus protocol

NARWHAL AND TUSK

Narwhal

A Directed Acyclic Graph (DAG) based
Mempool protocol

Questions

What is DAG?

Tusk

Zero-message overhead asynchronous consensus protocol

NARWHAL AND TUSK

Narwhal

A Directed Acyclic Graph (DAG) based
Mempool protocol

Questions

What is DAG?

What is Mempool?

Tusk

Zero-message overhead asynchronous consensus protocol

NARWHAL AND TUSK

Narwhal

A Directed Acyclic Graph (DAG) based
Mempool protocol

Questions

What is DAG?

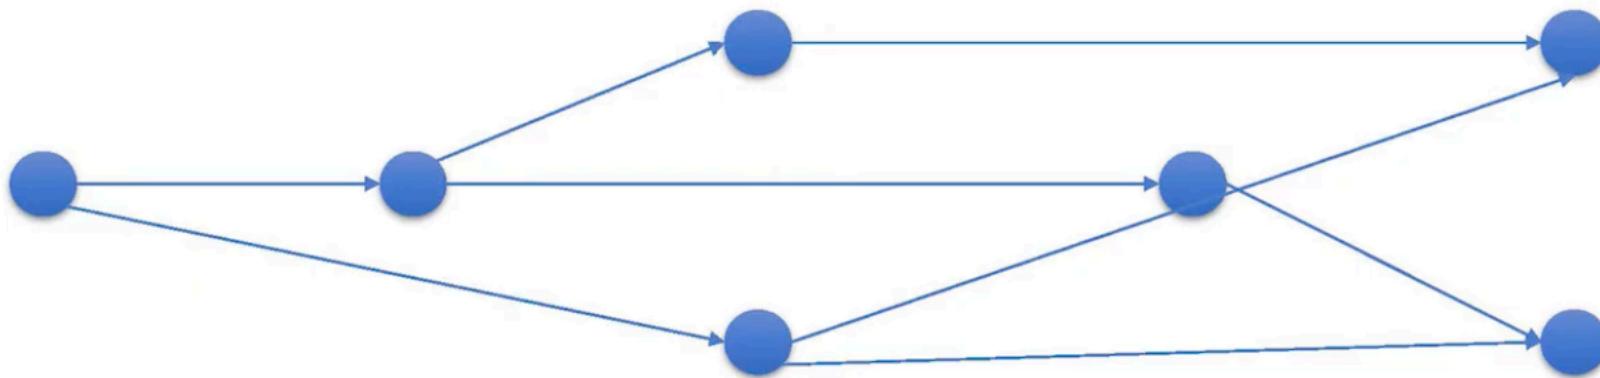
What is Mempool?

What is zero-message overhead?

Tusk

Zero-message overhead asynchronous consensus protocol

DAG - Directed Acyclic Graph



Graph - Node connected to other nodes

Directed - Connection between nodes have direction ($A \rightarrow B \neq B \rightarrow A$)

Acyclic - Non-circular. No closed loop in the connections.

Mempool 101

Fundamental Understanding

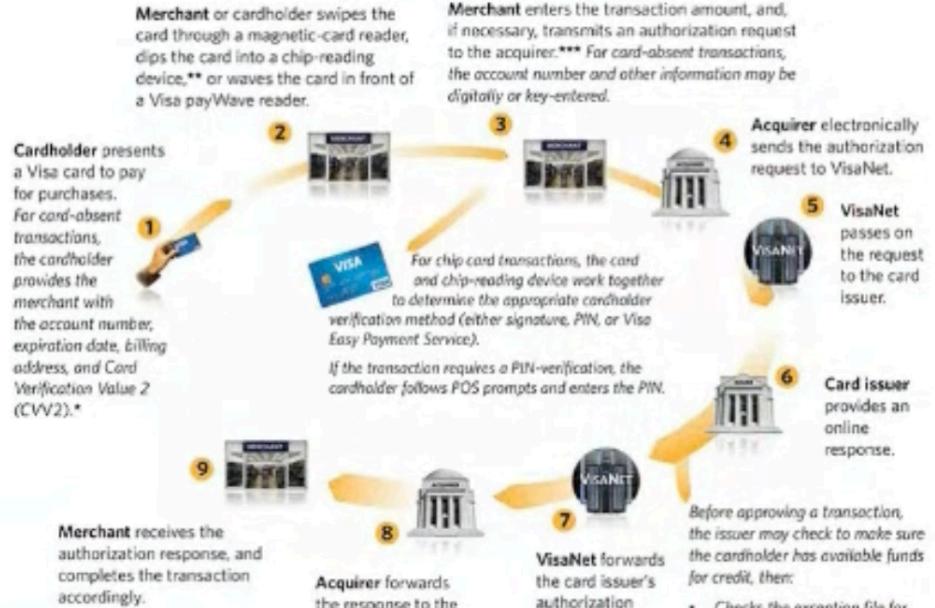
Transactions =
Complexity

What happens when you swipe and wait?

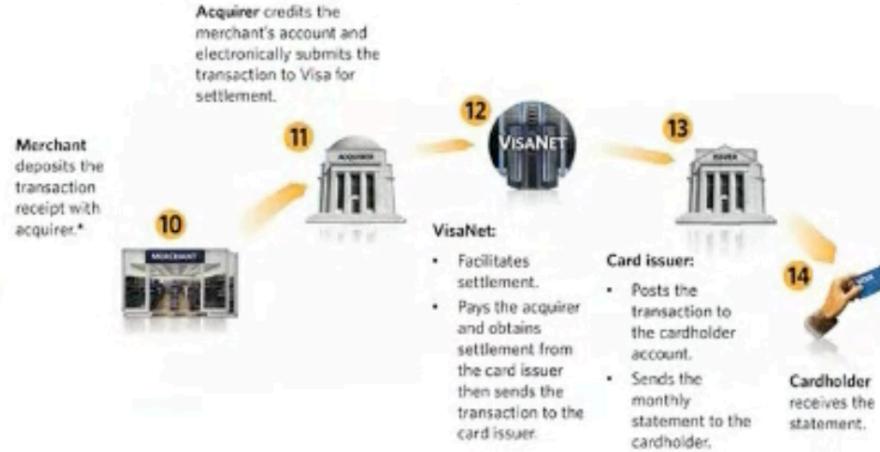


This happens

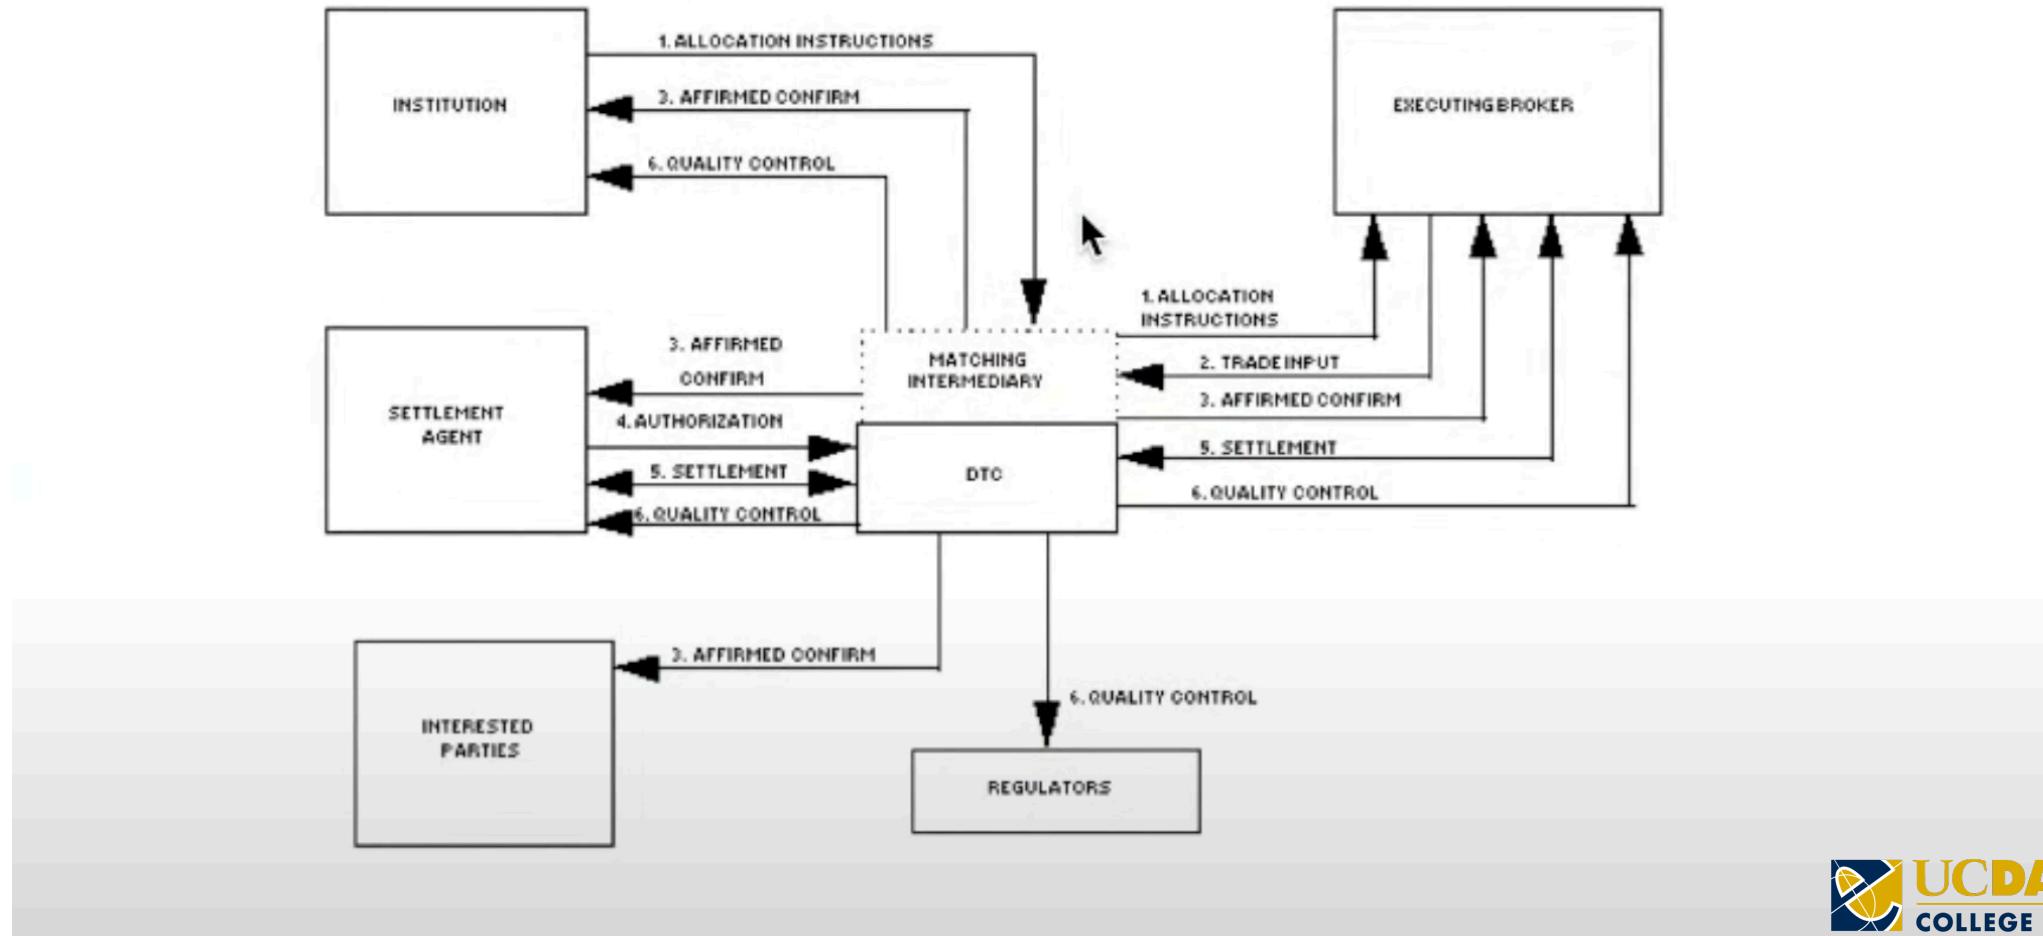
Magnetic-Stripe and Chip Card—Credit or Debit Authorization



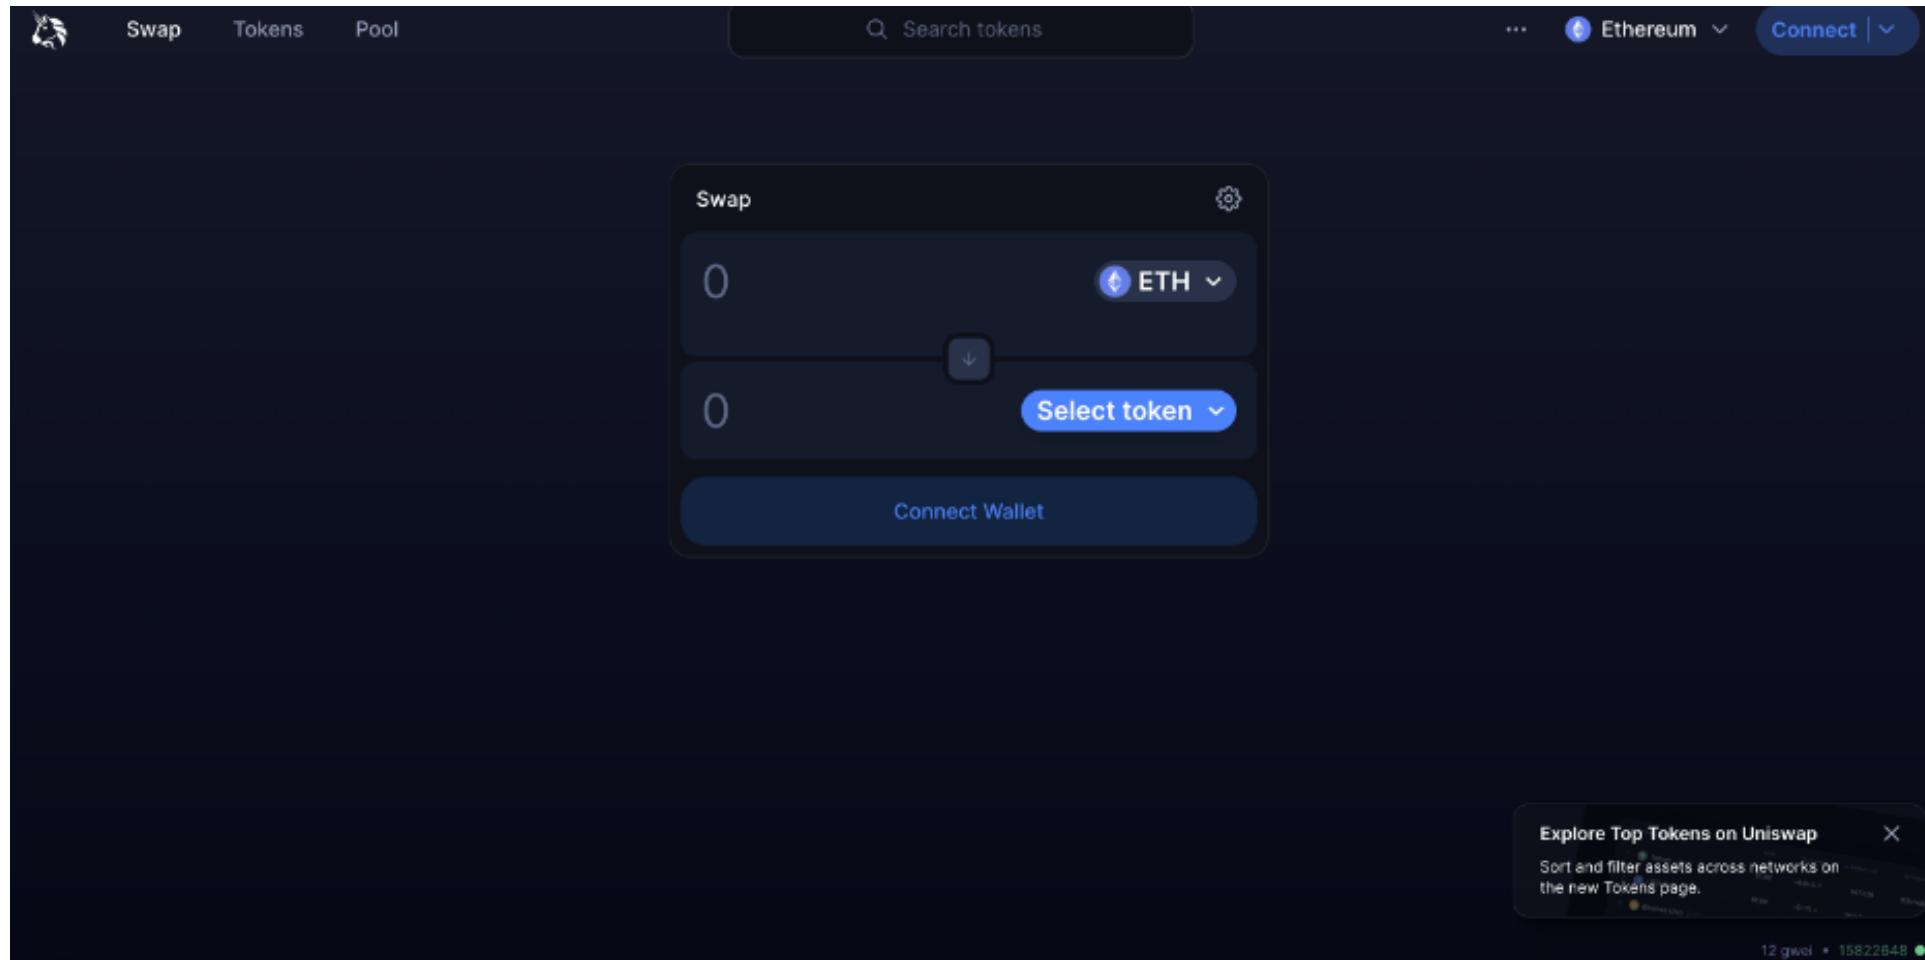
Magnetic-Stripe and Chip Card—Clearing and Settlement



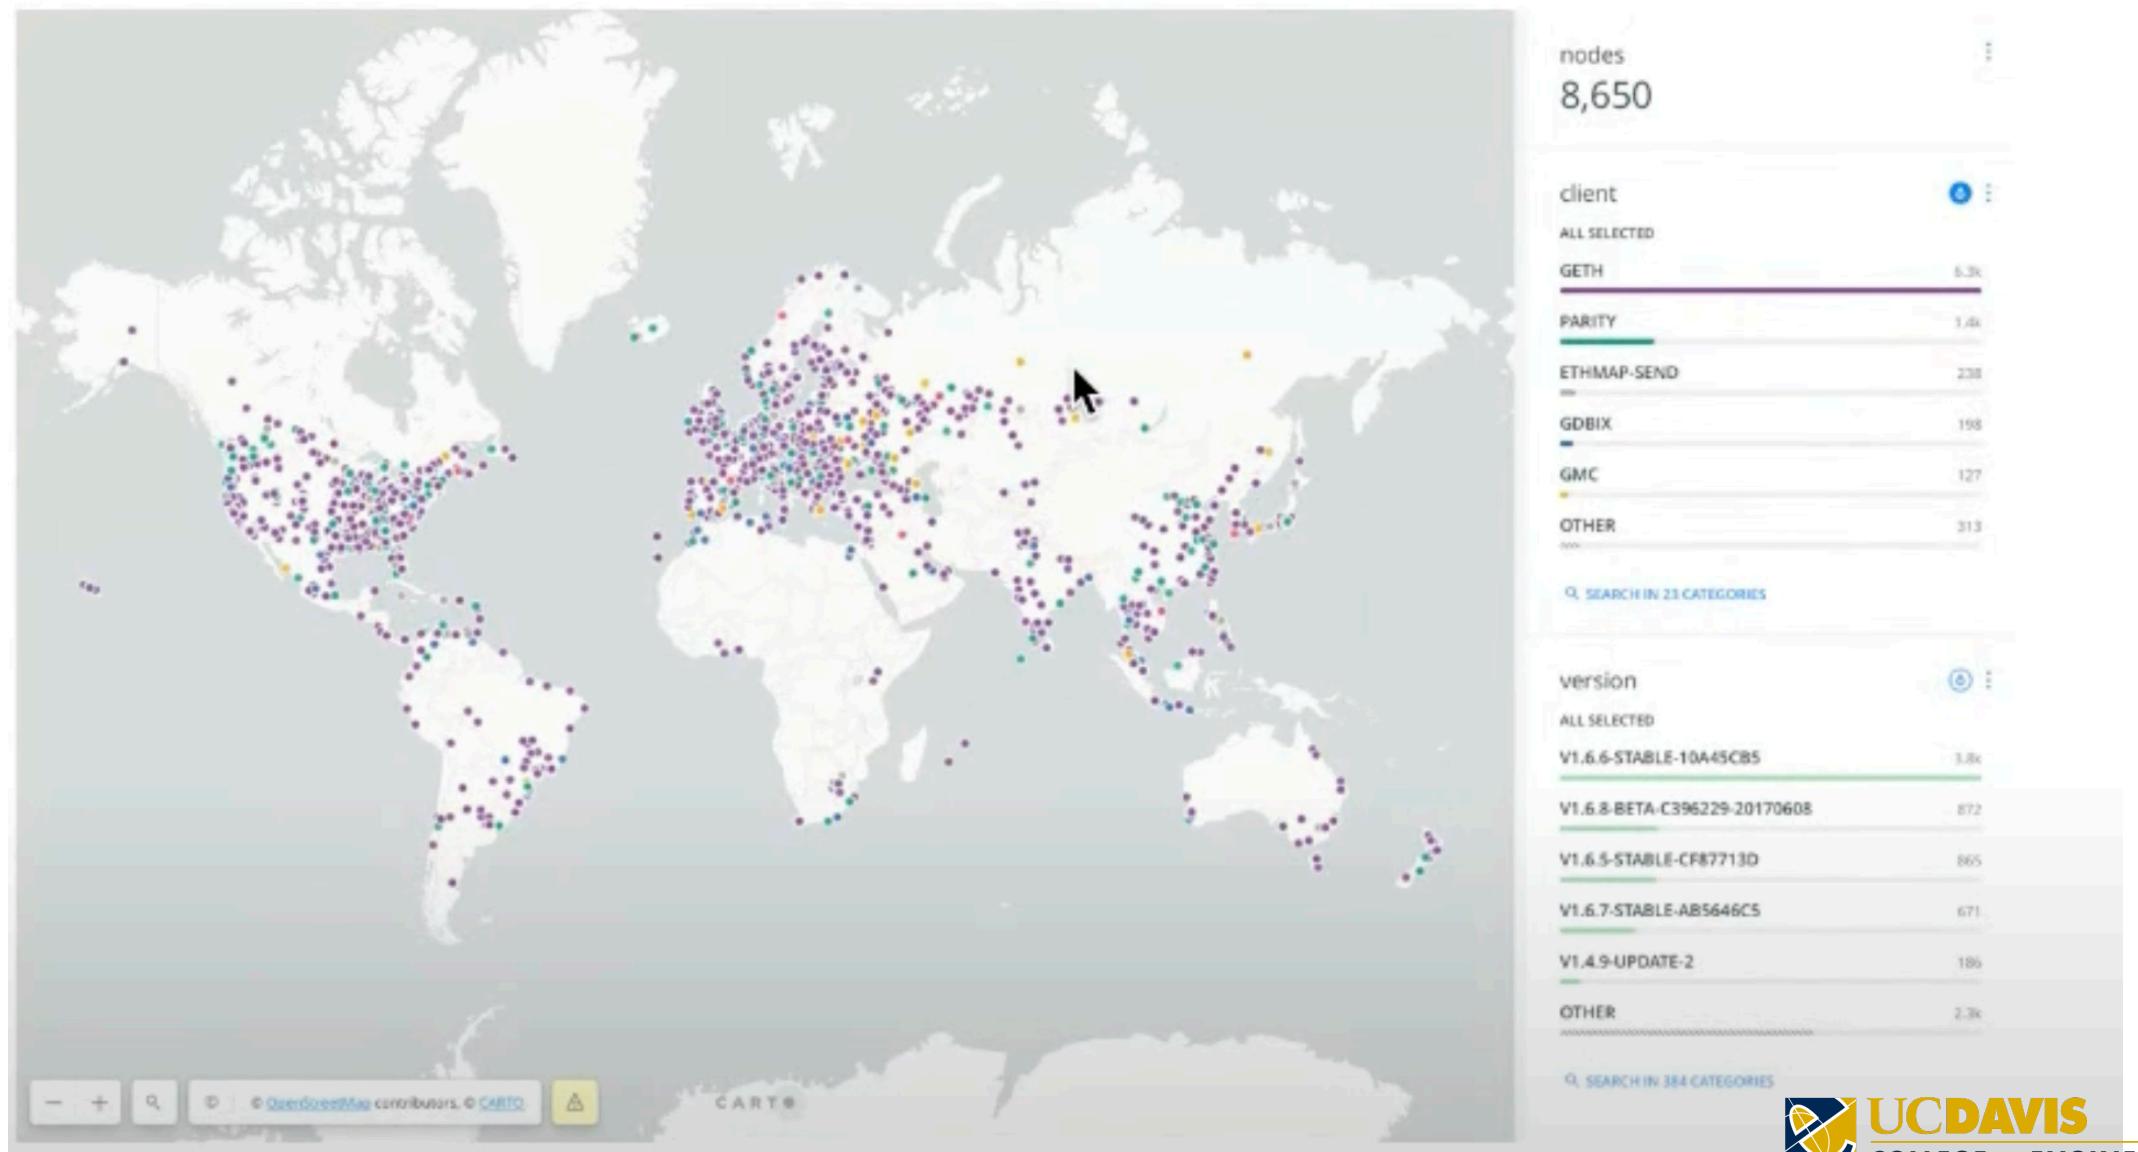
Same is true for a stock market application



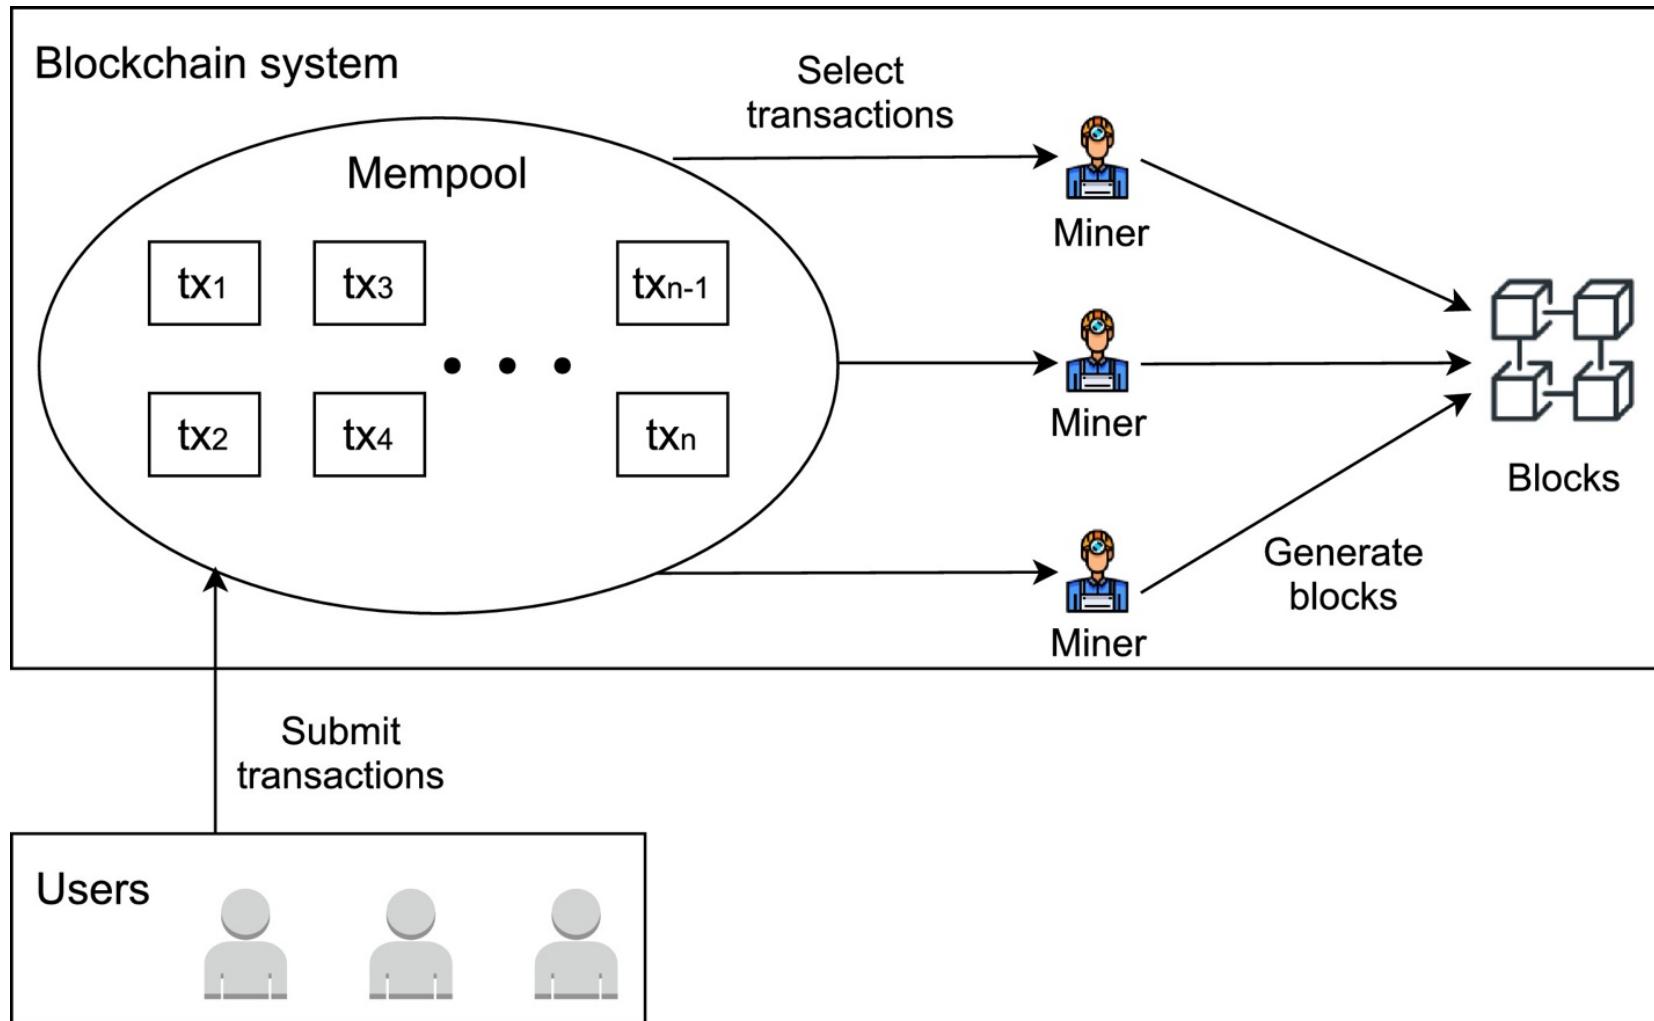
Onto a Crypto Trading application



Ethereum Network



Mempool



1

Mempool = PRE-CONSENSUS

2

There is no single mempool

3

Each mempool has its own
unique set of transactions

4

Each mempool has its **own unique characteristics** based on client type, configuration, peers

5

The mempool is mutable due to
Replacement Transactions,
including Speed ups & Cancels

6

The mempool is an **inconsistent ‘fabric’** that grows lumpy during congestion/pressure

7

Your transaction competes in
the mempool for block space

Traditional Transactions



Like a ballroom dance

- Very orderly
- Plays by the rules

Blockchain Transactions: Expectations



Like a flash mob

- More spontaneous
- Organic

Blockchain Transactions: Reality

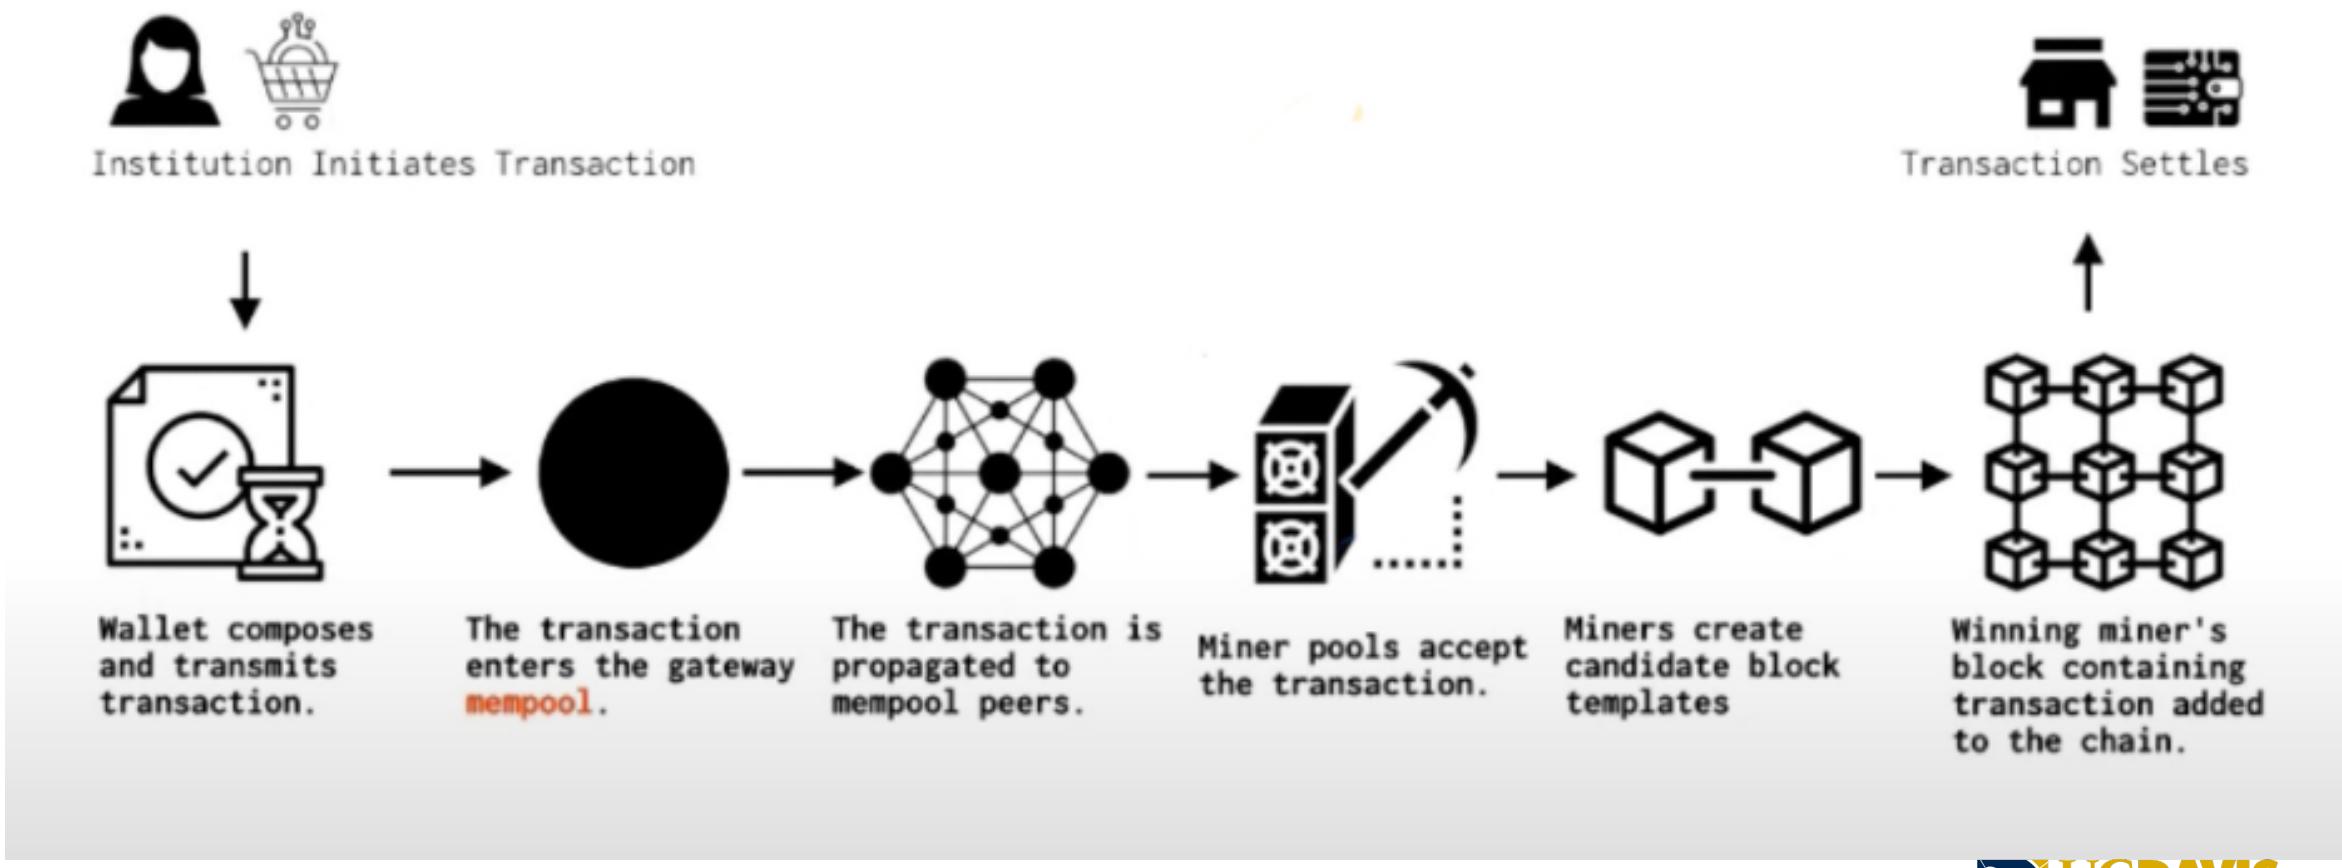


Like a mosh pit

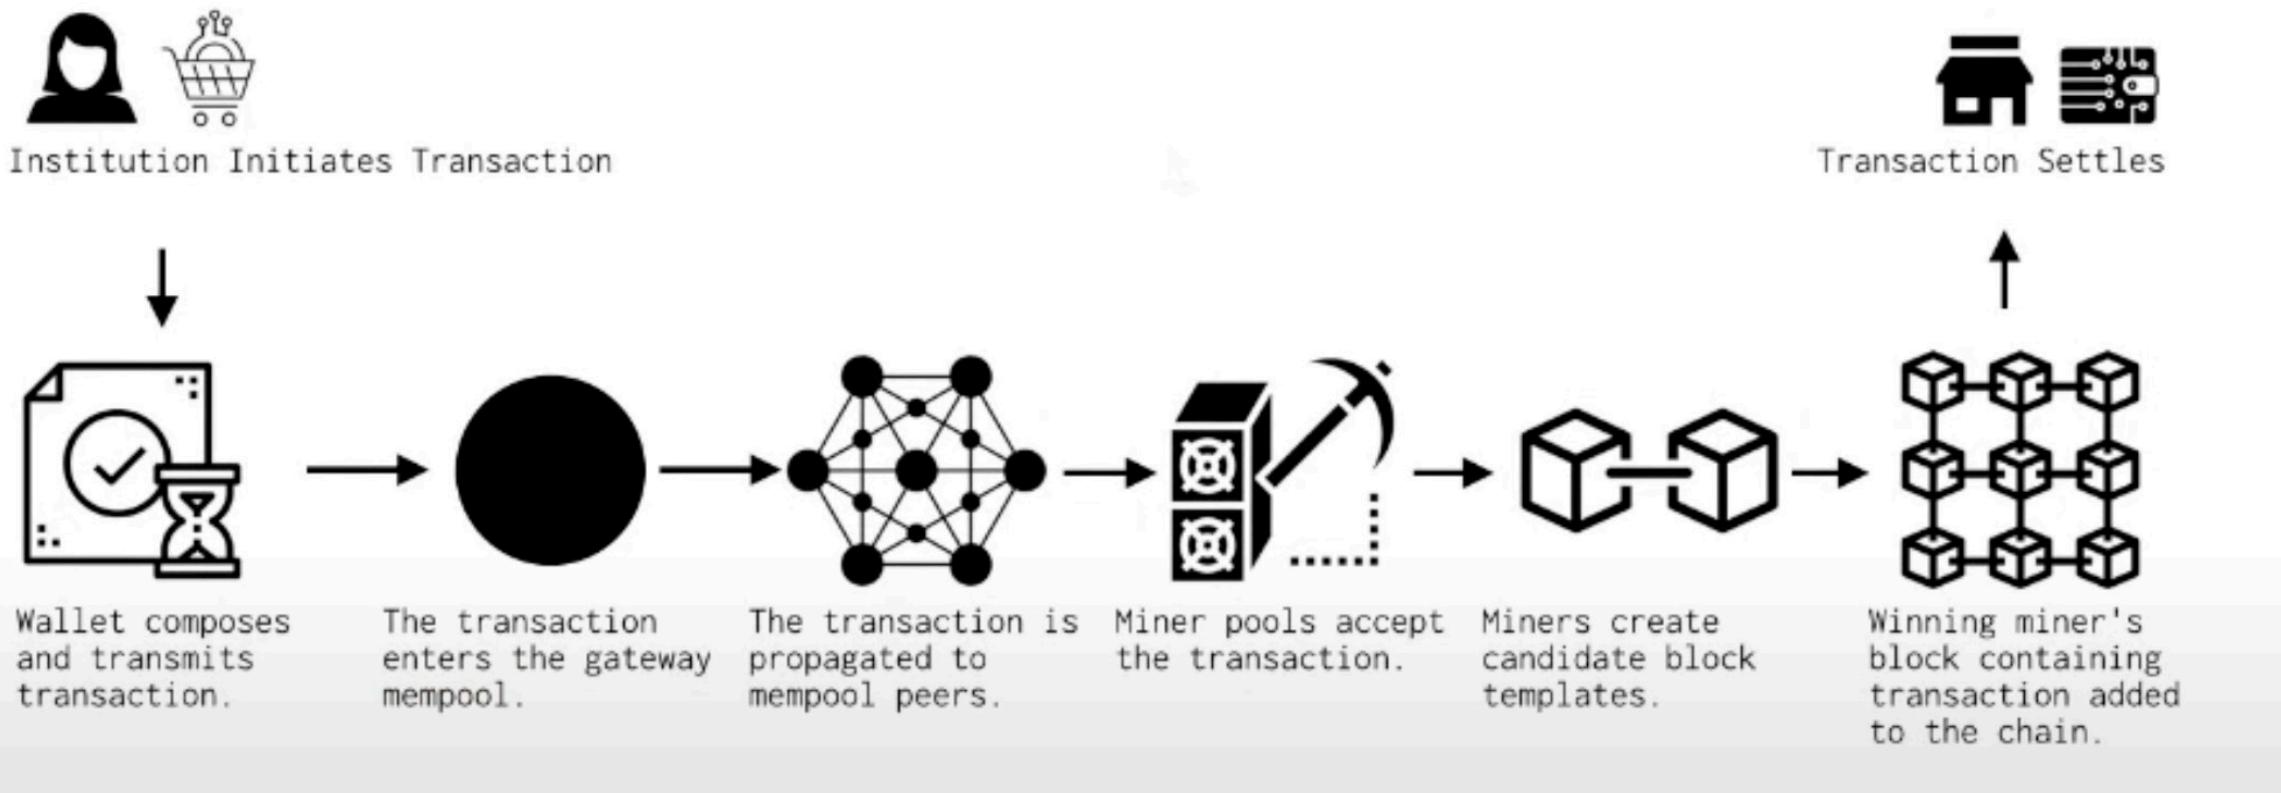
- Tx's competing for confirmation
- Un-orchestrated fashion

Behind the scenes of Blockchain Transaction Operations and our **Area of Focus**

Blockchain Transaction Operations: Behind The Scenes



Blockchain Transaction Operations: Risks



Blockchain Transaction Operations: Risks

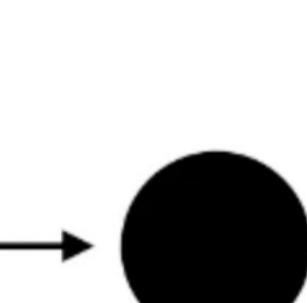


Institution Initiates Transaction

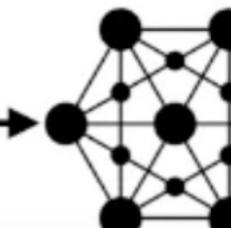
Trader cancels.



Wallet composes
and transmits
transaction.



The transaction
enters the gateway
mempool.



The transaction is
propagated to
mempool peers.



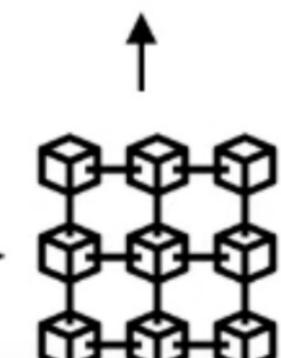
Miner pools accept
the transaction.



Miners create
candidate block
templates.

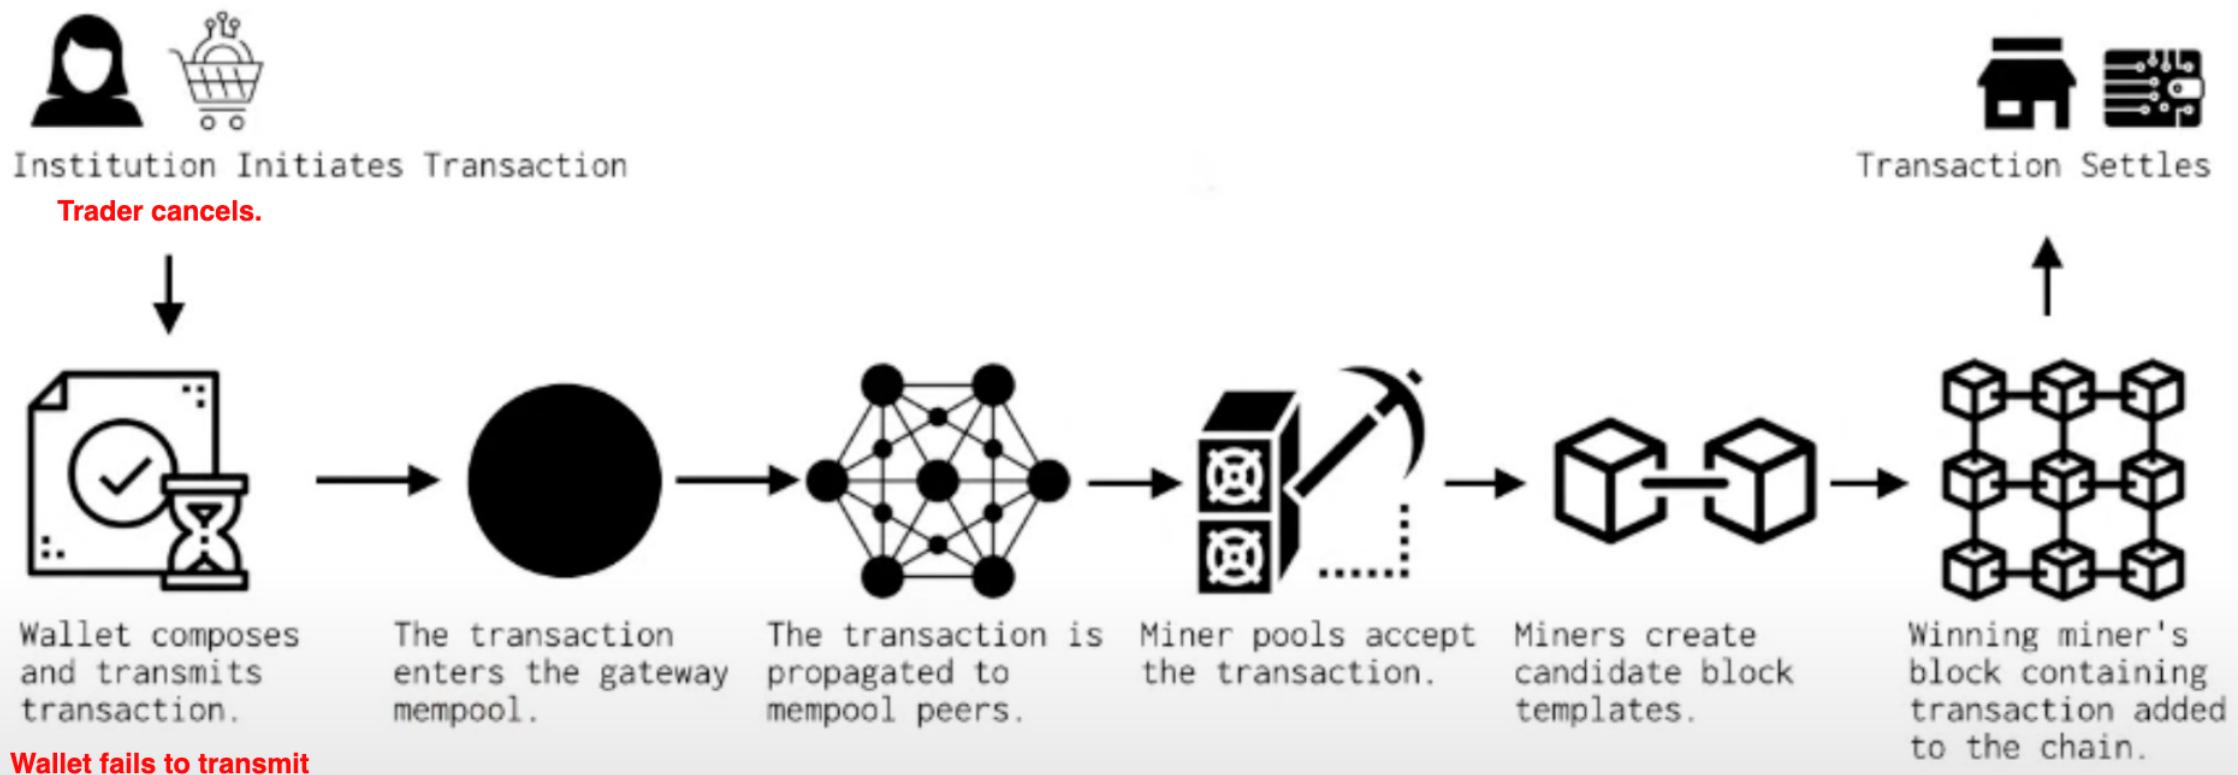


Transaction Settles



Winning miner's
block containing
transaction added
to the chain.

Blockchain Transaction Operations: Risks



Blockchain Transaction Operations: Risks



Institution Initiates Transaction

Trader cancels.



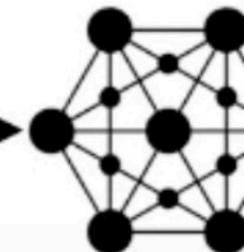
Wallet composes
and transmits
transaction.

Wallet fails to transmit



The transaction
enters the gateway
mempool.

Gateway fails to accept



The transaction is
propagated to
mempool peers.



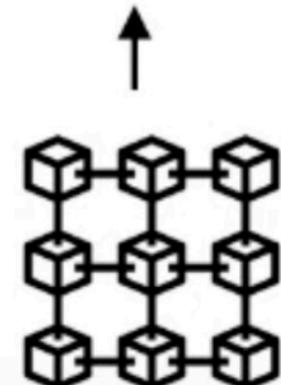
Miner pools accept
the transaction.



Miners create
candidate block
templates.

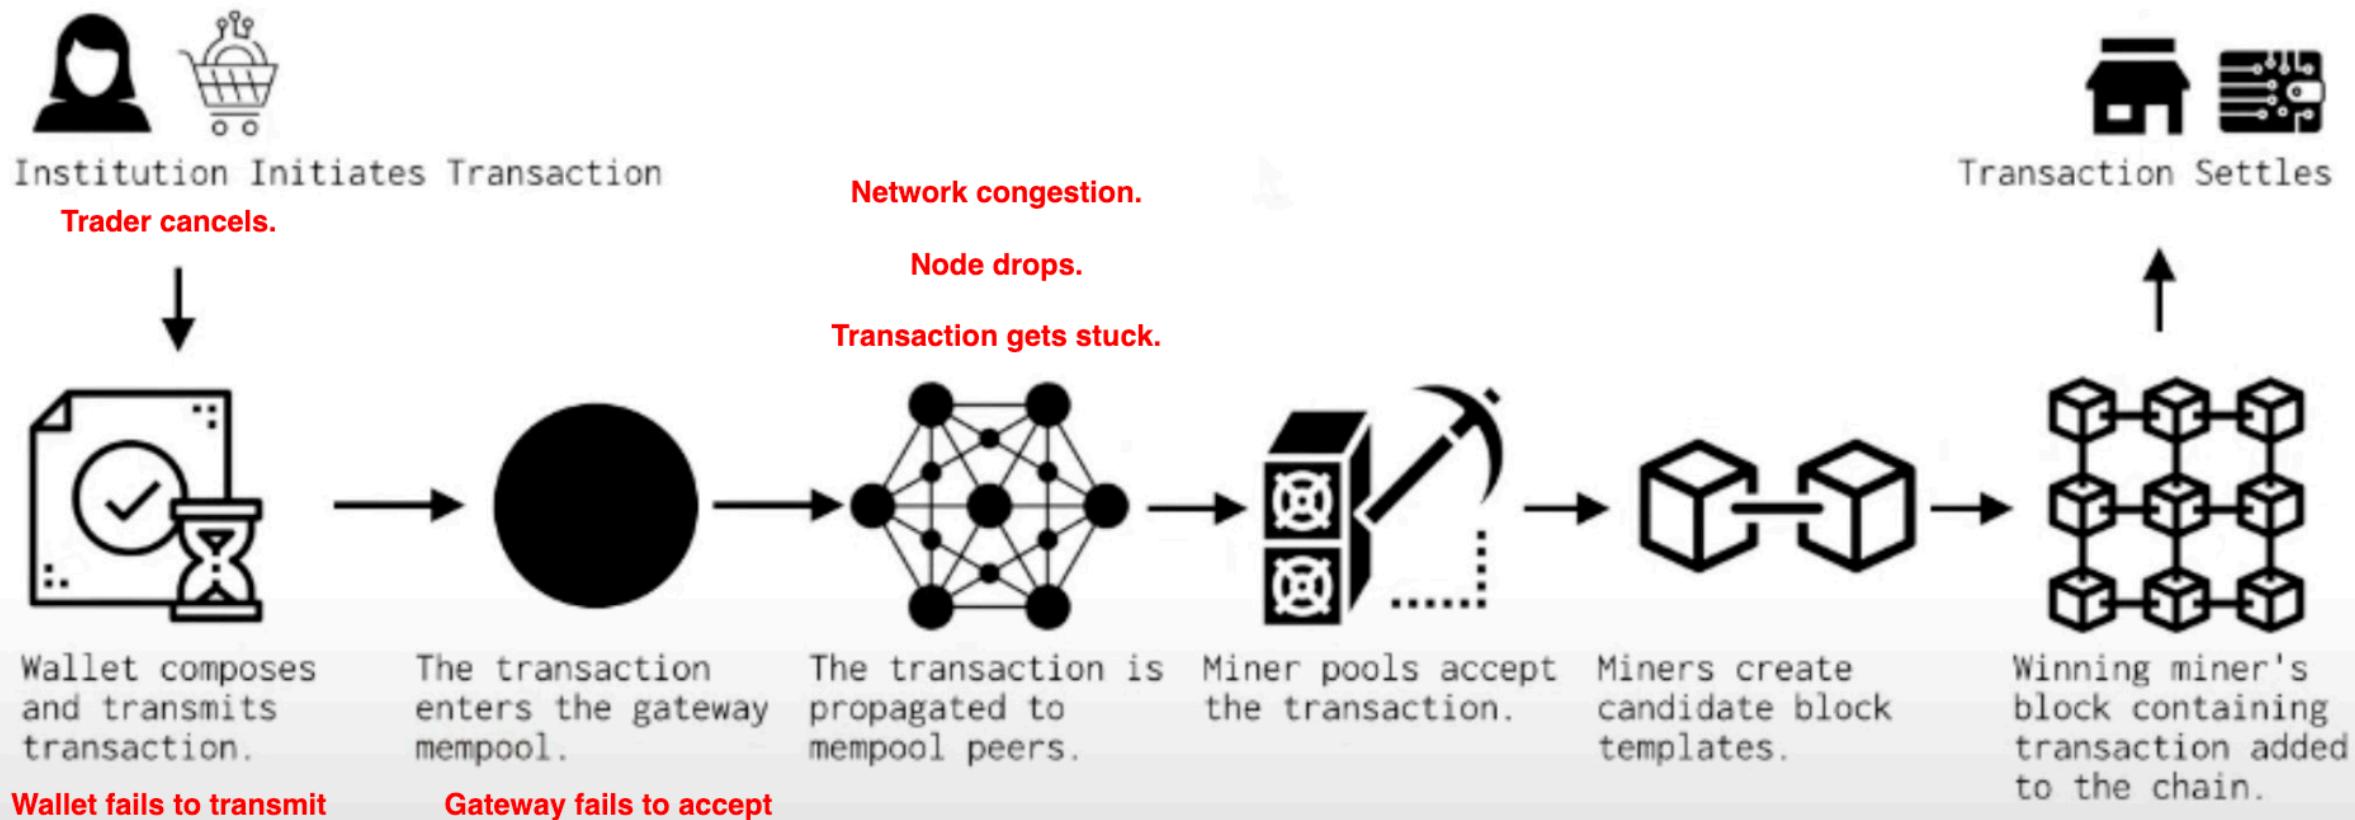


Transaction Settles



Winning miner's
block containing
transaction added
to the chain.

Blockchain Transaction Operations: Risks



Blockchain Transaction Operations: Risks



Institution Initiates Transaction
Trader cancels.

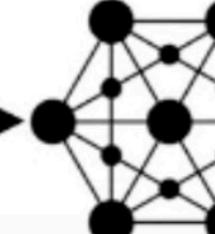


Wallet composes
and transmits
transaction.

Wallet fails to transmit

Network congestion.

Node drops.
Transaction gets stuck.

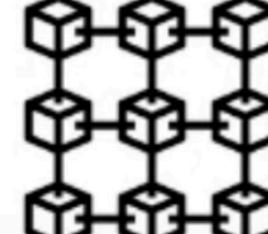
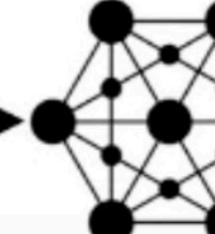


The transaction
enters the gateway
mempool.

Gateway fails to accept

Front running

**Outcompeted by higher
fee TX**

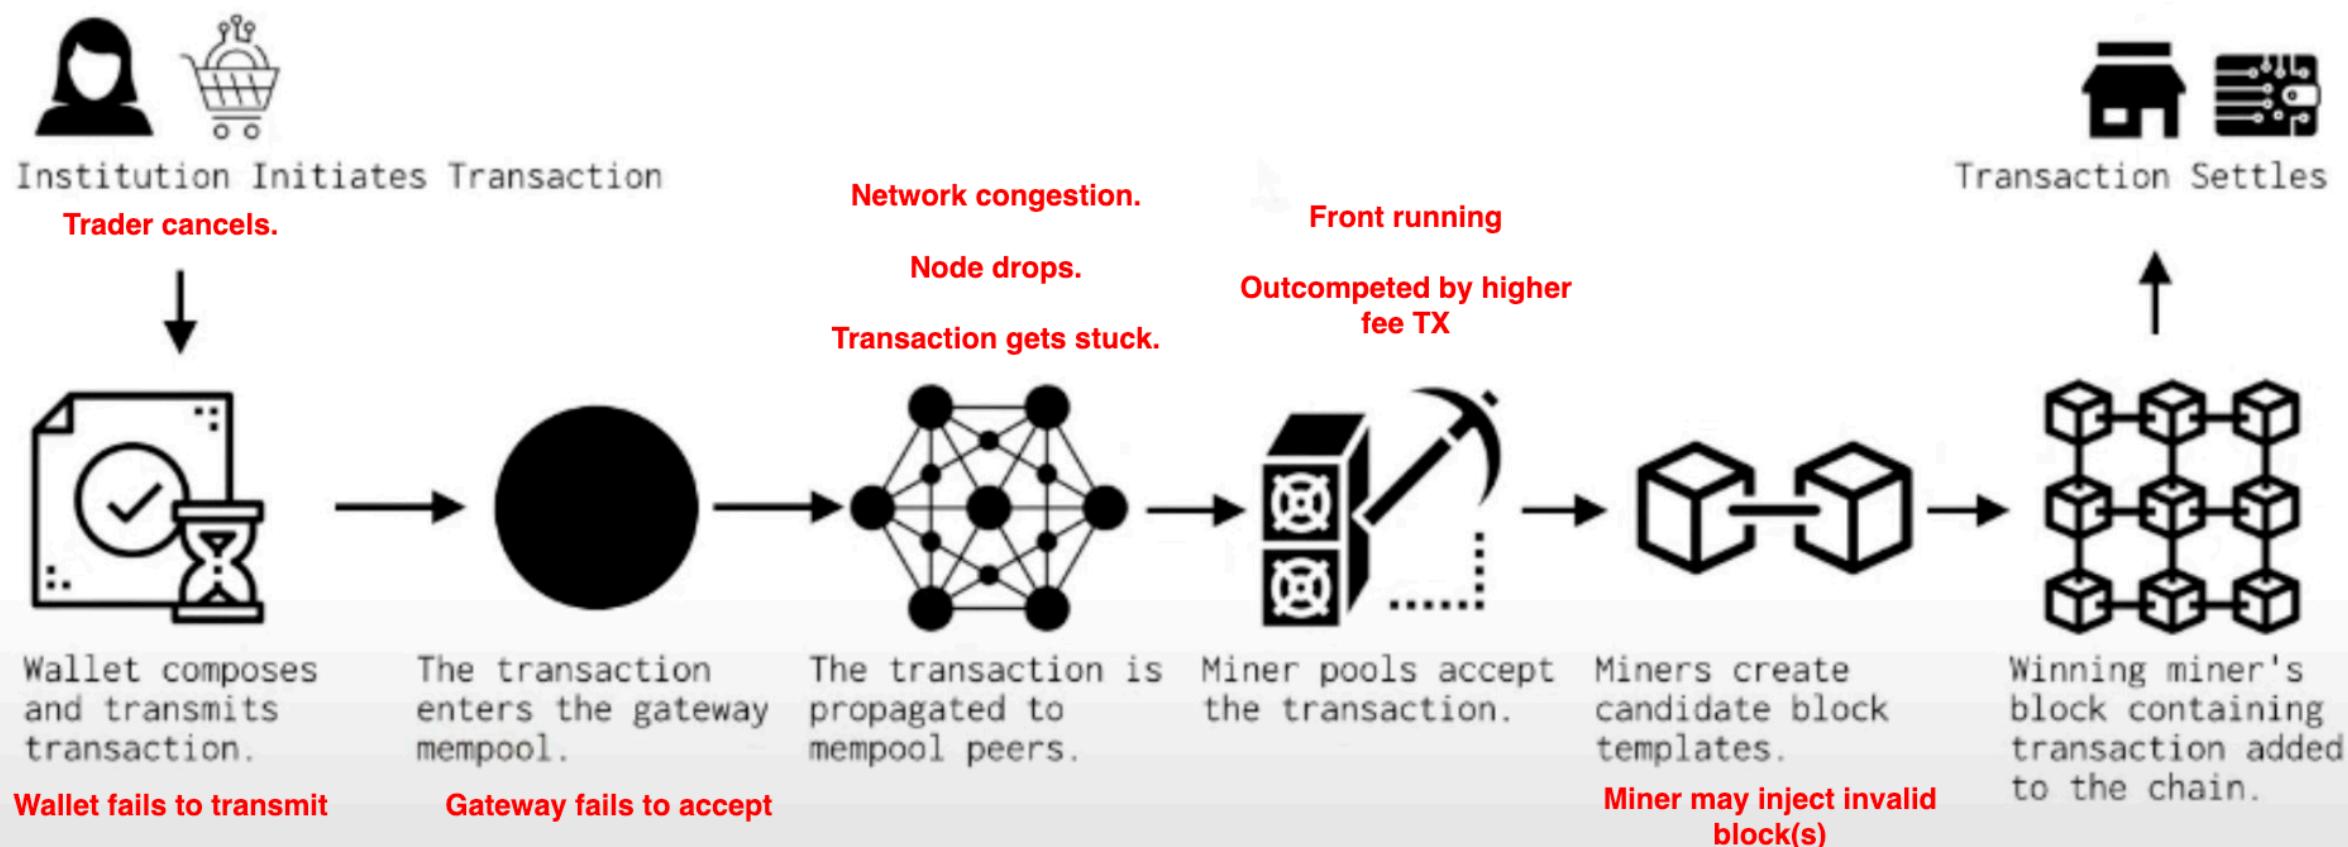


Transaction Settles

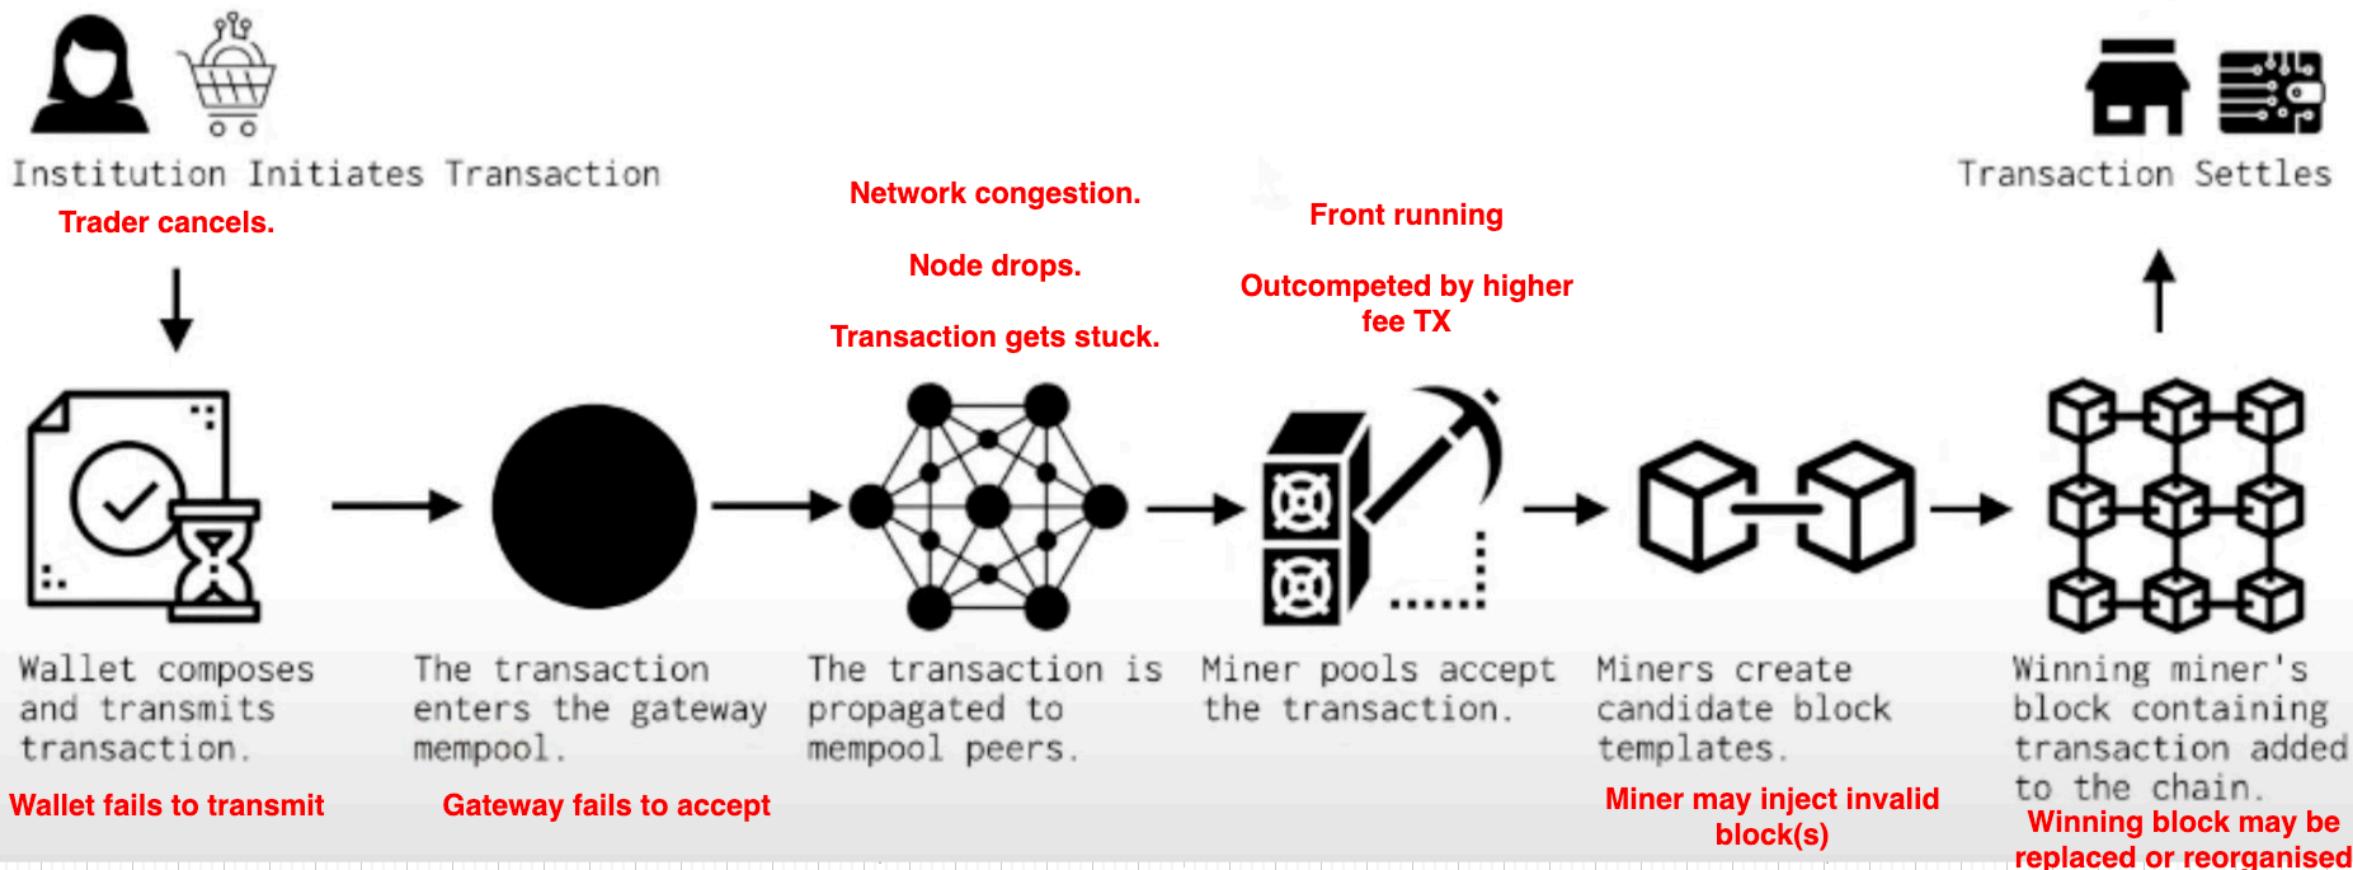


Upward arrow indicating the flow from initiation to settlement.

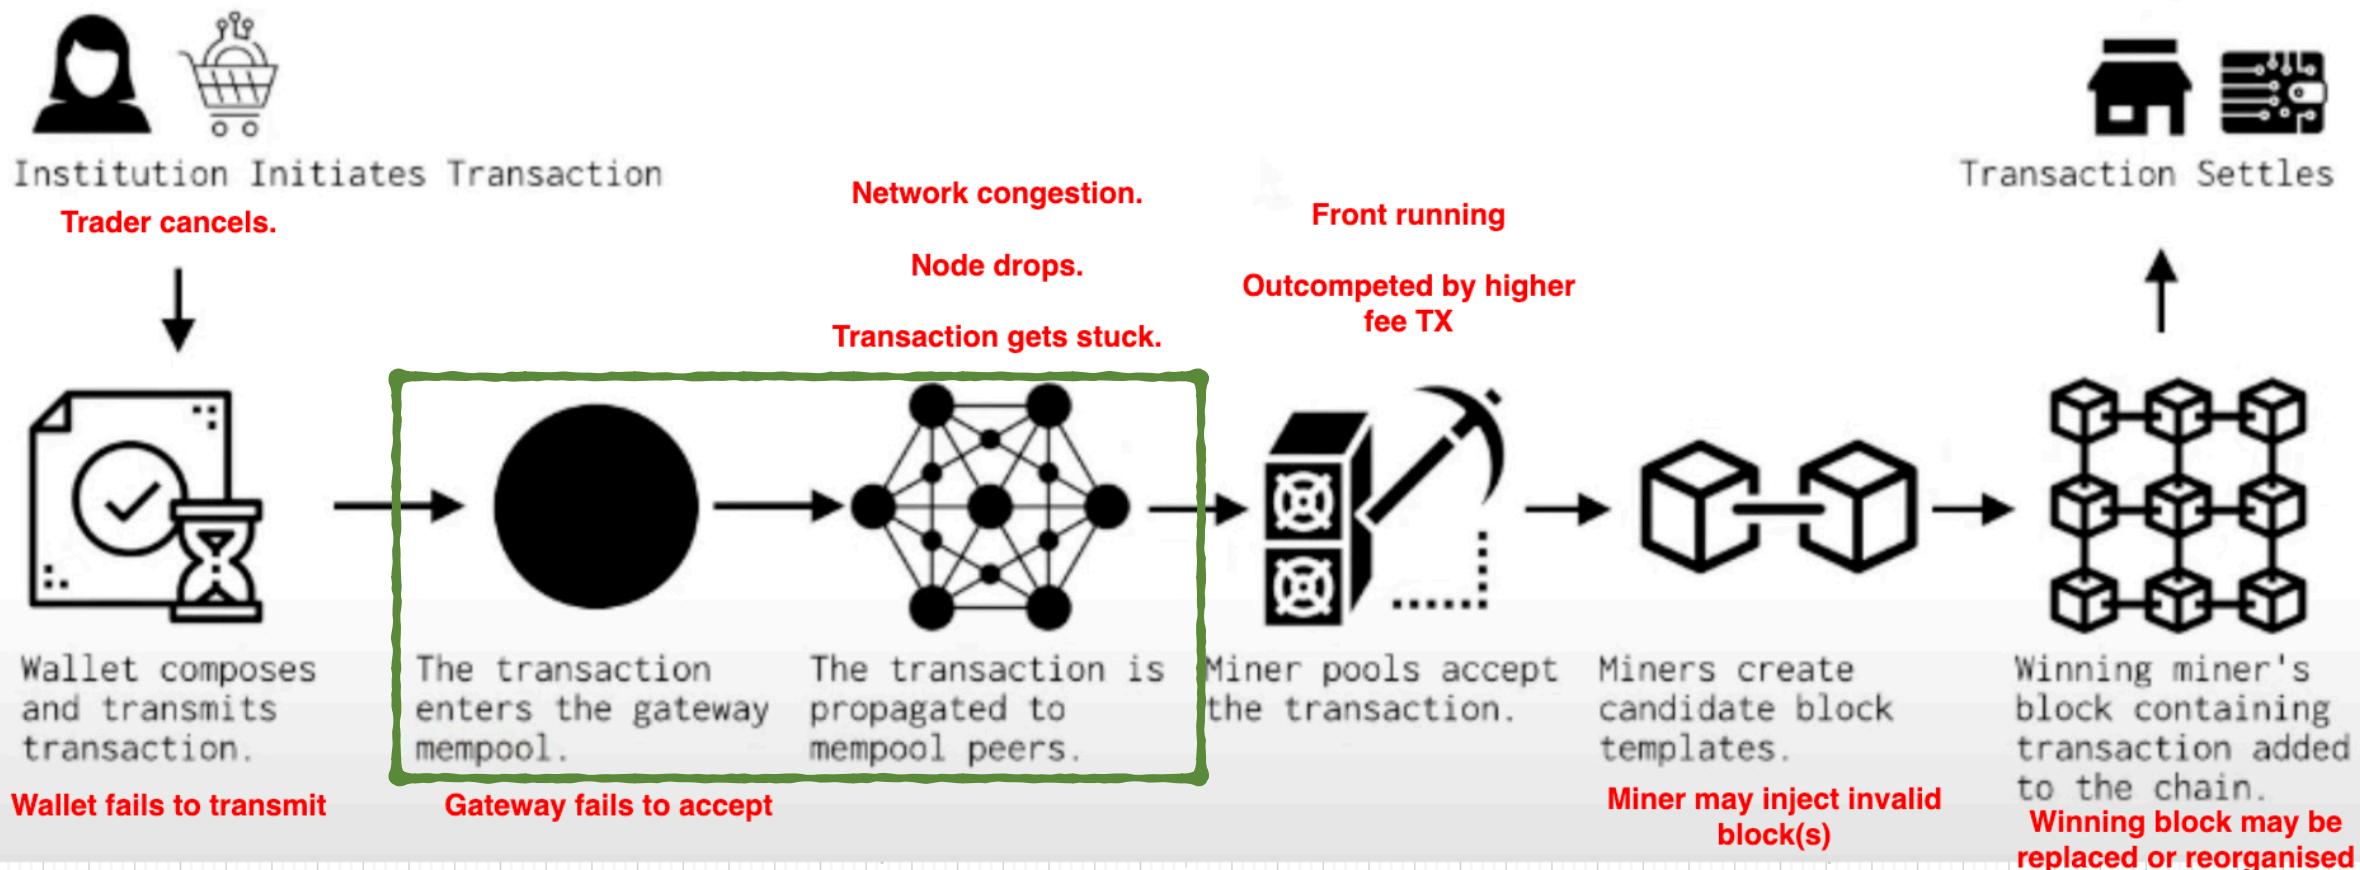
Blockchain Transaction Operations: Risks



Blockchain Transaction Operations: Risks



Blockchain Transaction Operations: Risks



NARWHAL AND TUSK

Narwhal

A Directed Acyclic Graph (DAG) based
Mempool protocol

Questions

What is DAG?

What is Mempool?

What is zero-message overhead?

Tusk

Zero-message overhead asynchronous consensus protocol

NARWHAL AND TUSK

Narwhal

A Directed Acyclic Graph (DAG) based
Mempool protocol

Questions

What is DAG?

What is Mempool?

What is zero-message overhead?

Tusk

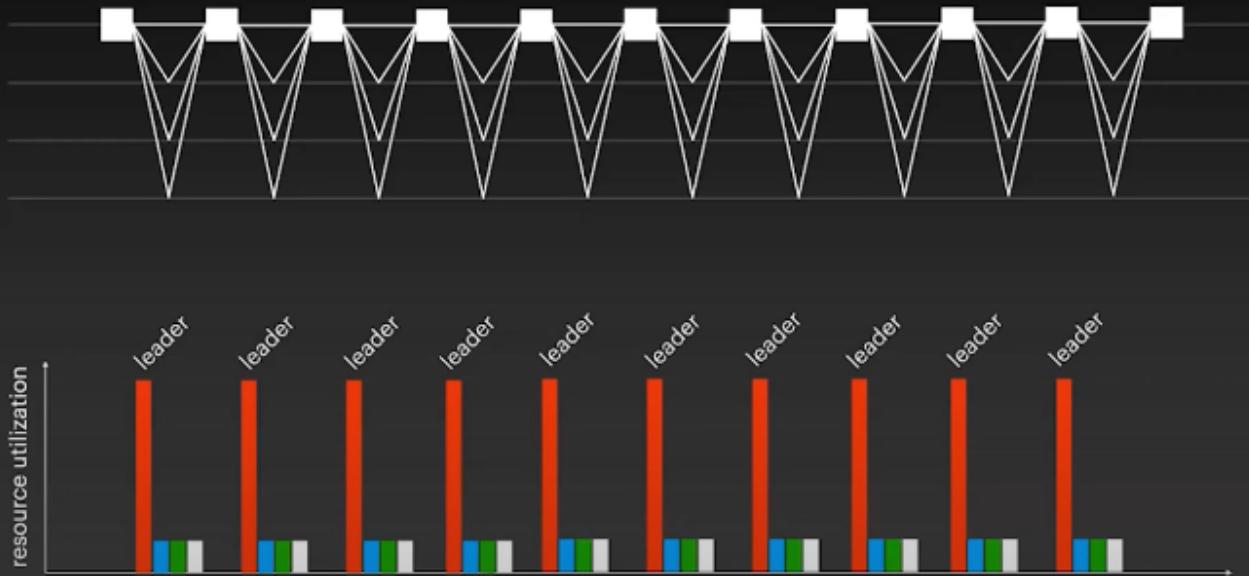
Zero-message overhead asynchronous consensus protocol

Consensus engine **does not send**
any message at all. Hence the term.

NON-EGALITARIAN RESOURCE UTILIZATIONS

Current Designs

Typical leader-based solutions

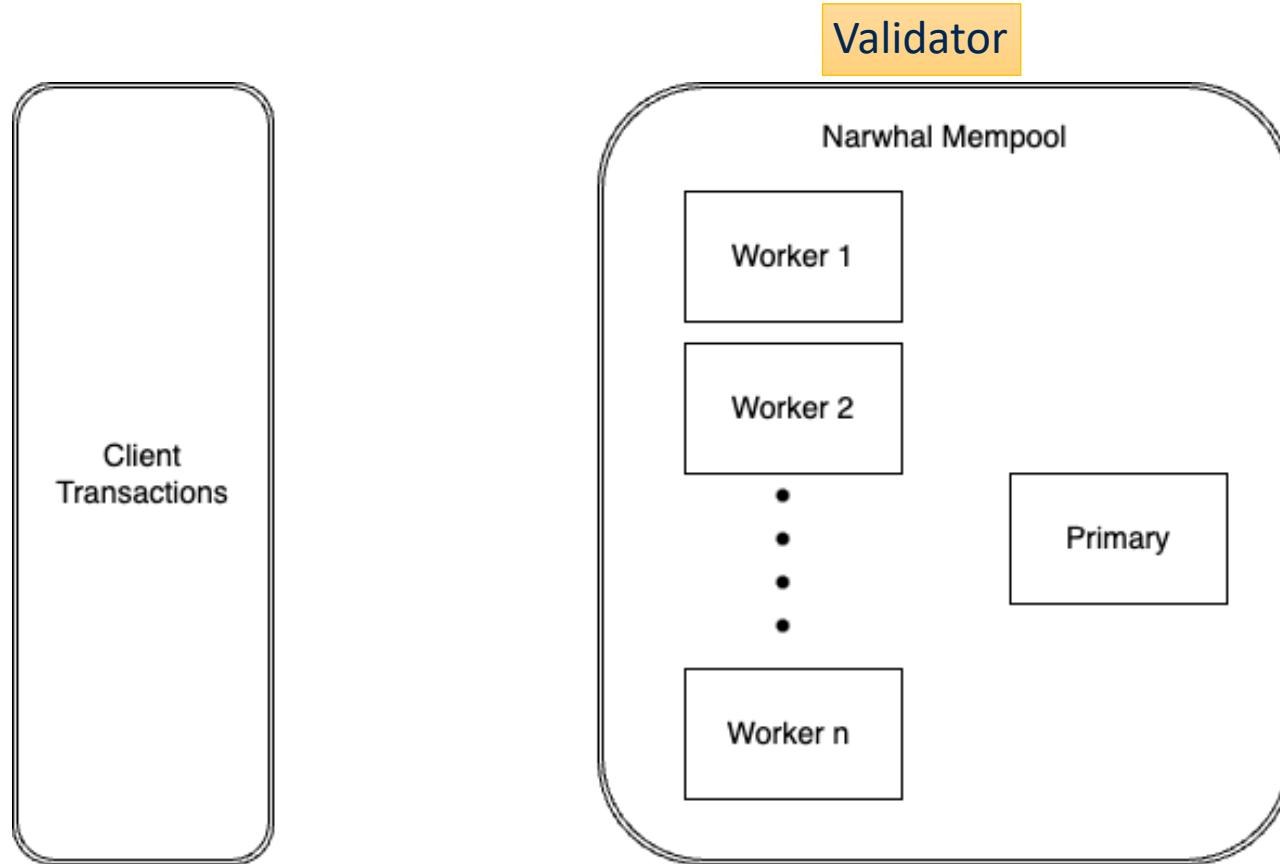


MEMPOOL IS THE KEY

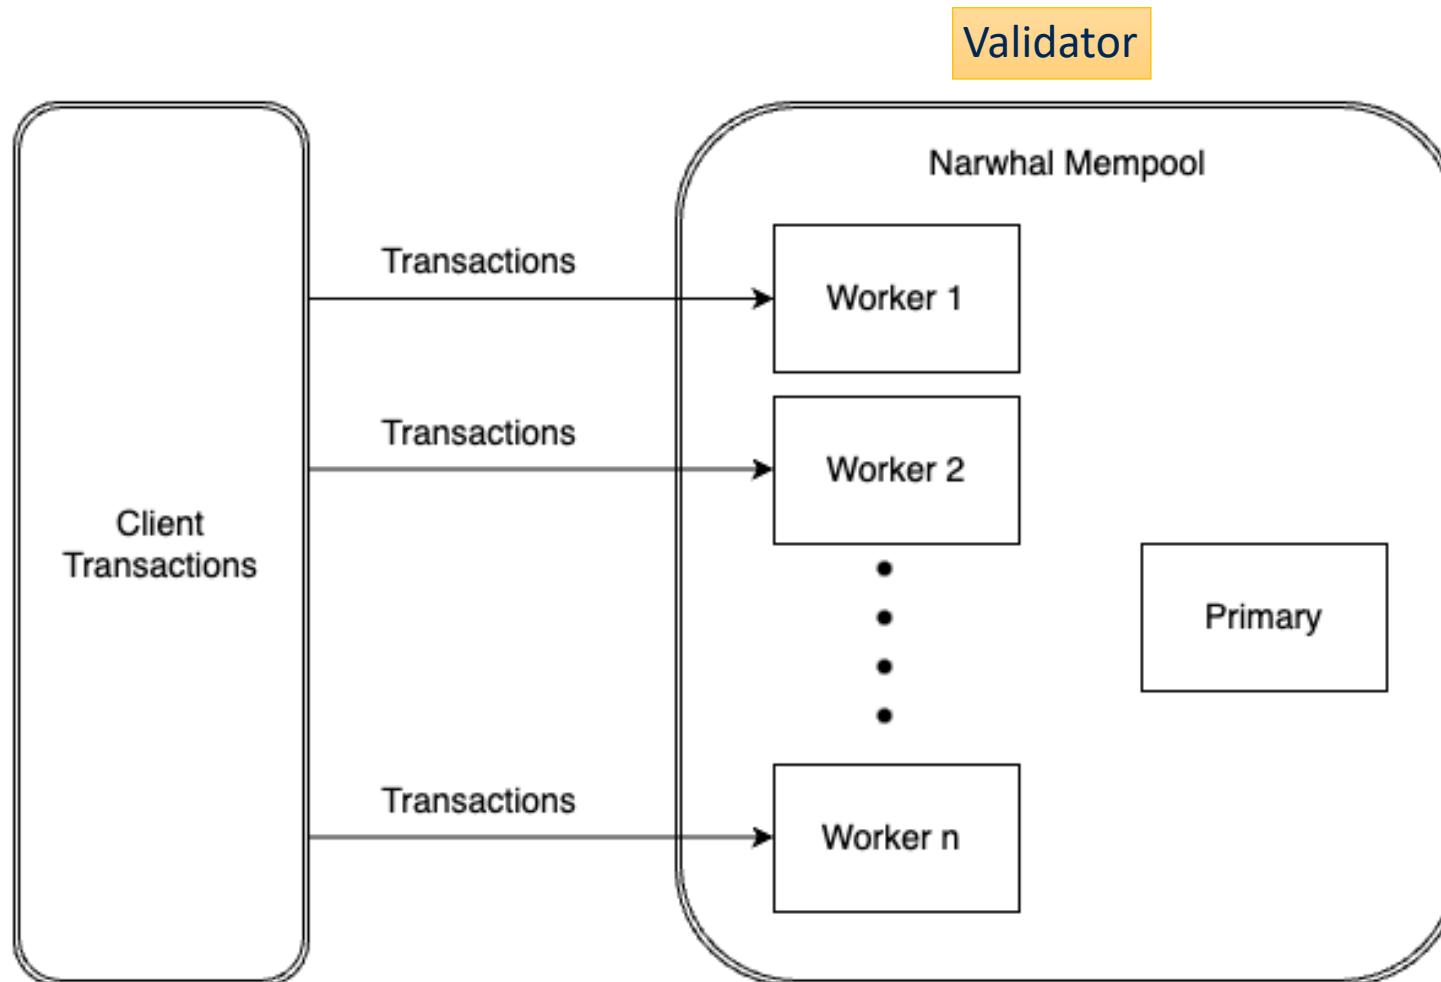
NARWHAL

Narwhal

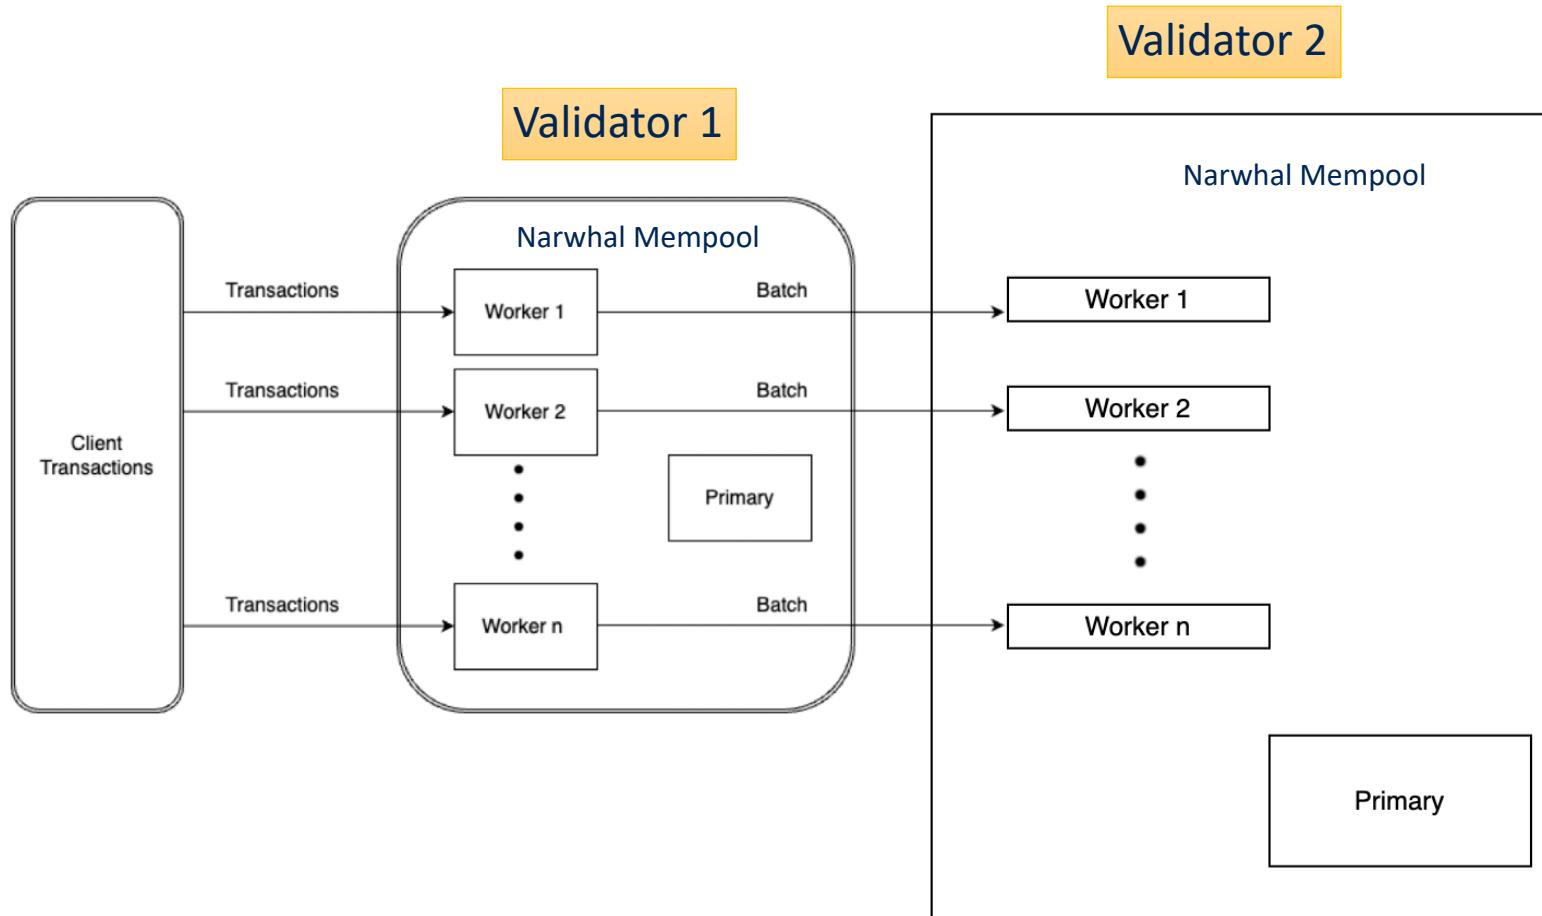
The workers and the primary in Mempool



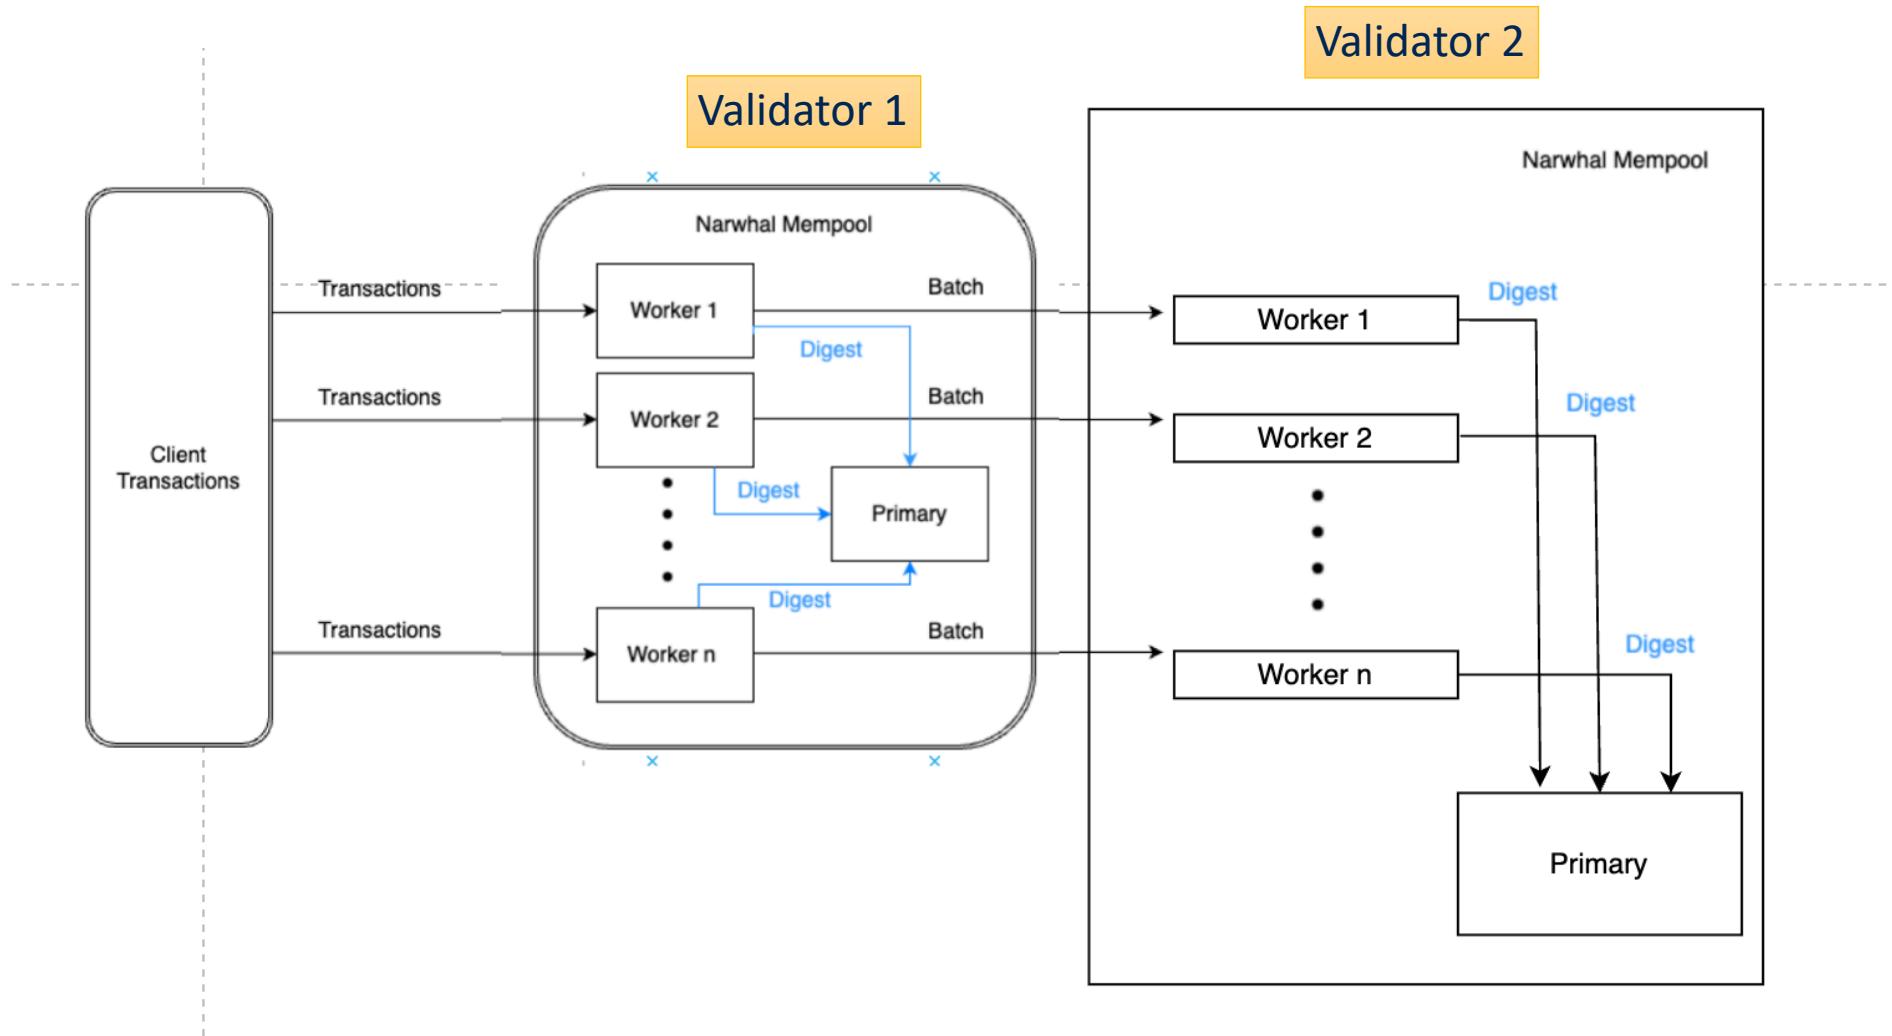
Client sends transactions



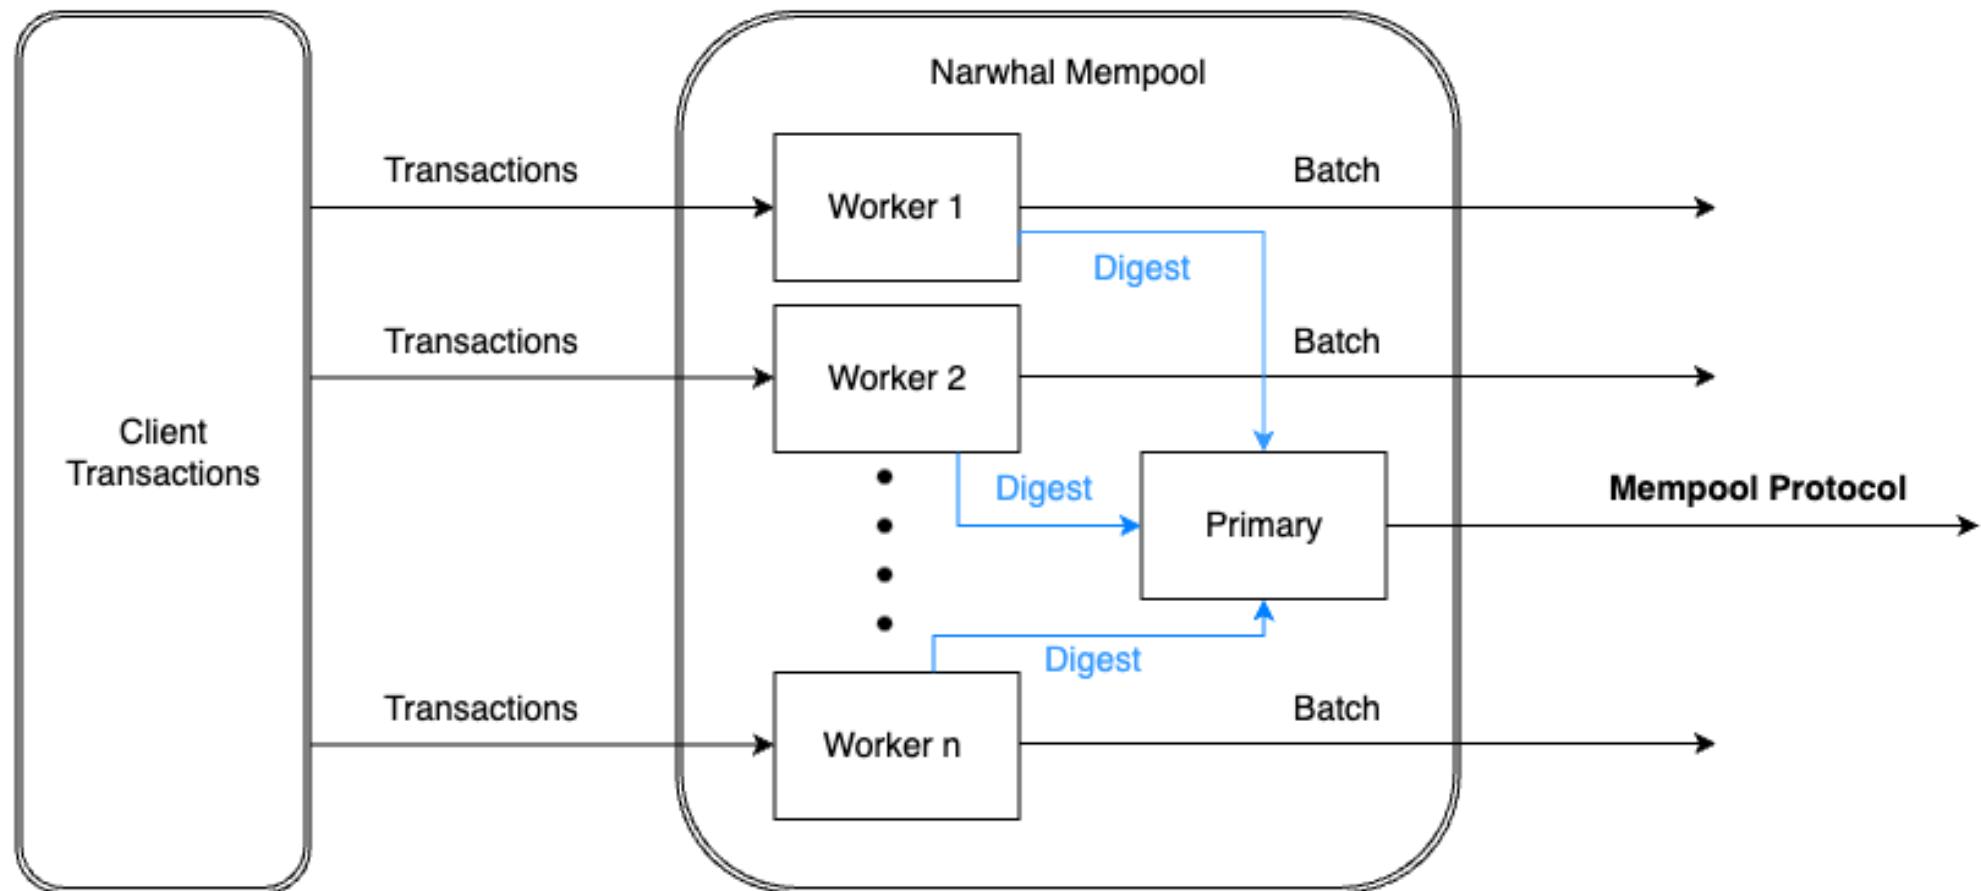
Batching Transactions



Digest sent to primary



Mempool Protocol Initialization

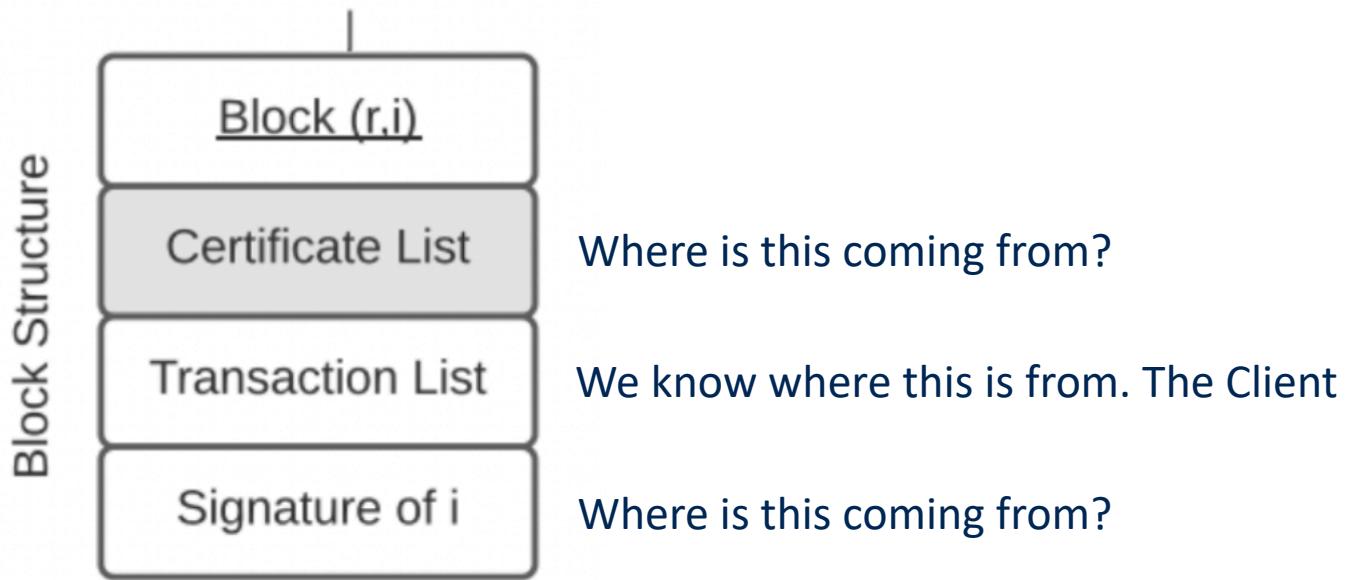


Every Primary creates a Block

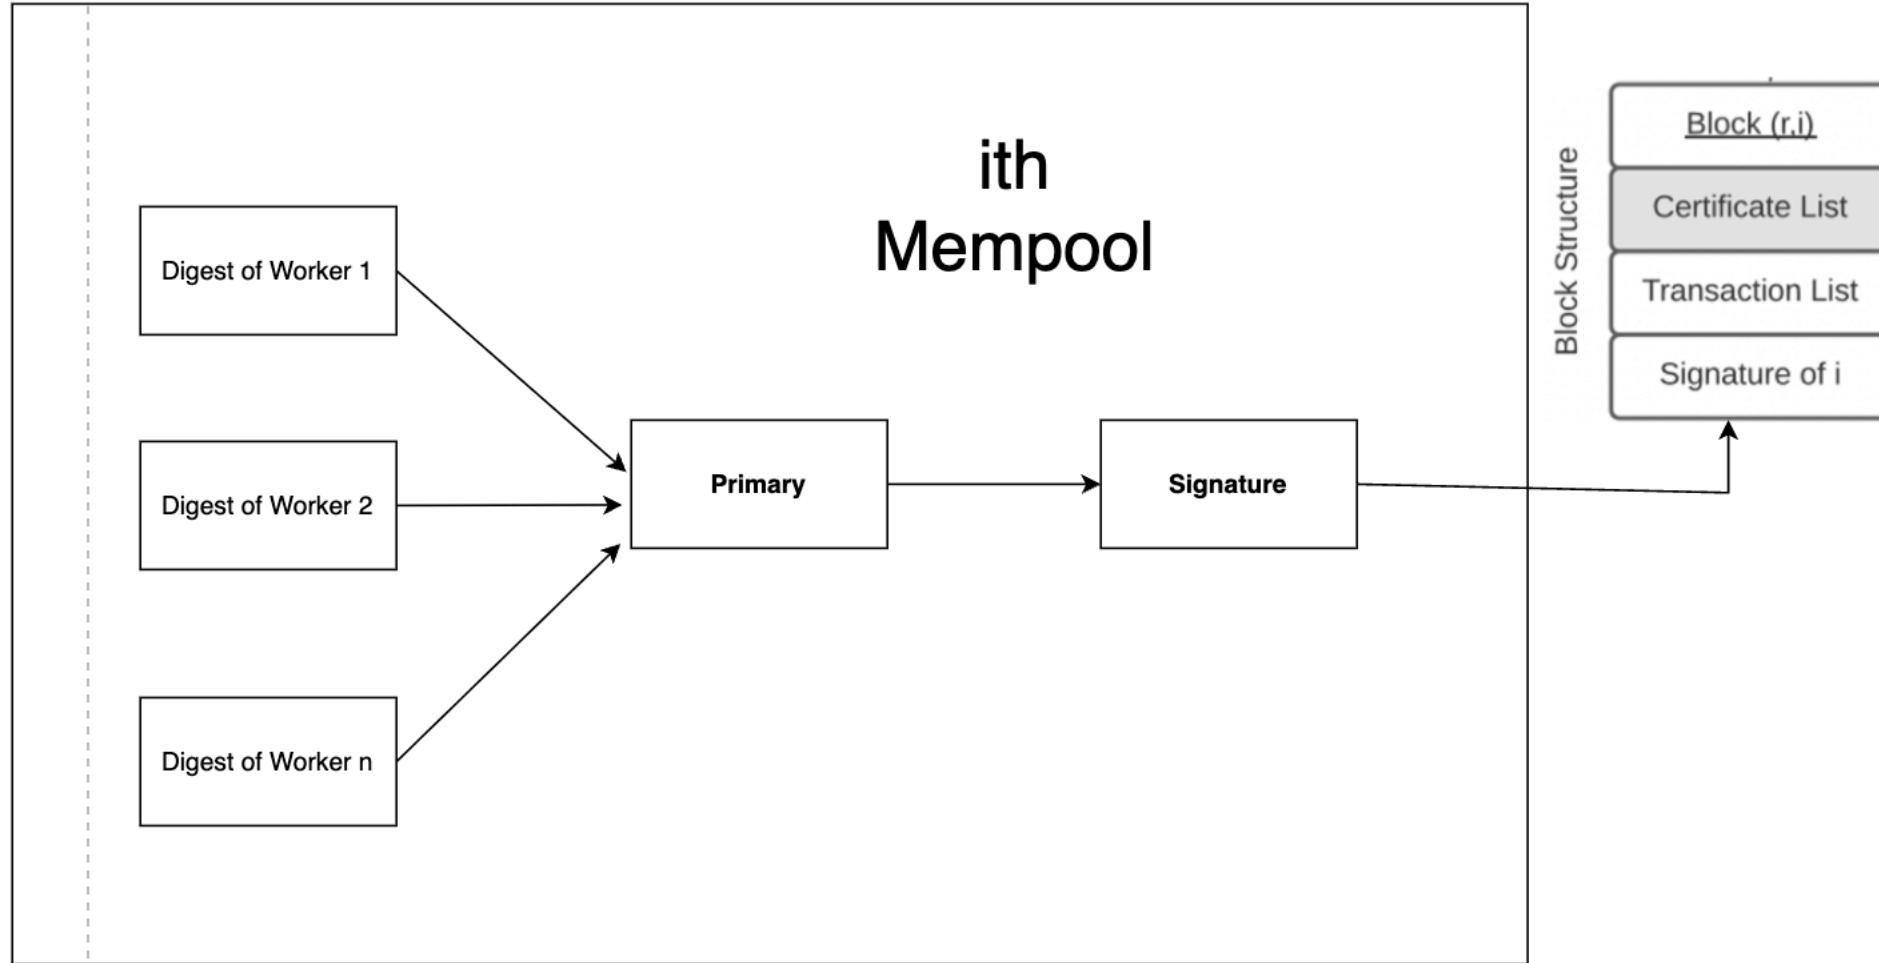
Every Primary creates a Block

What does it contain?

Block Structure

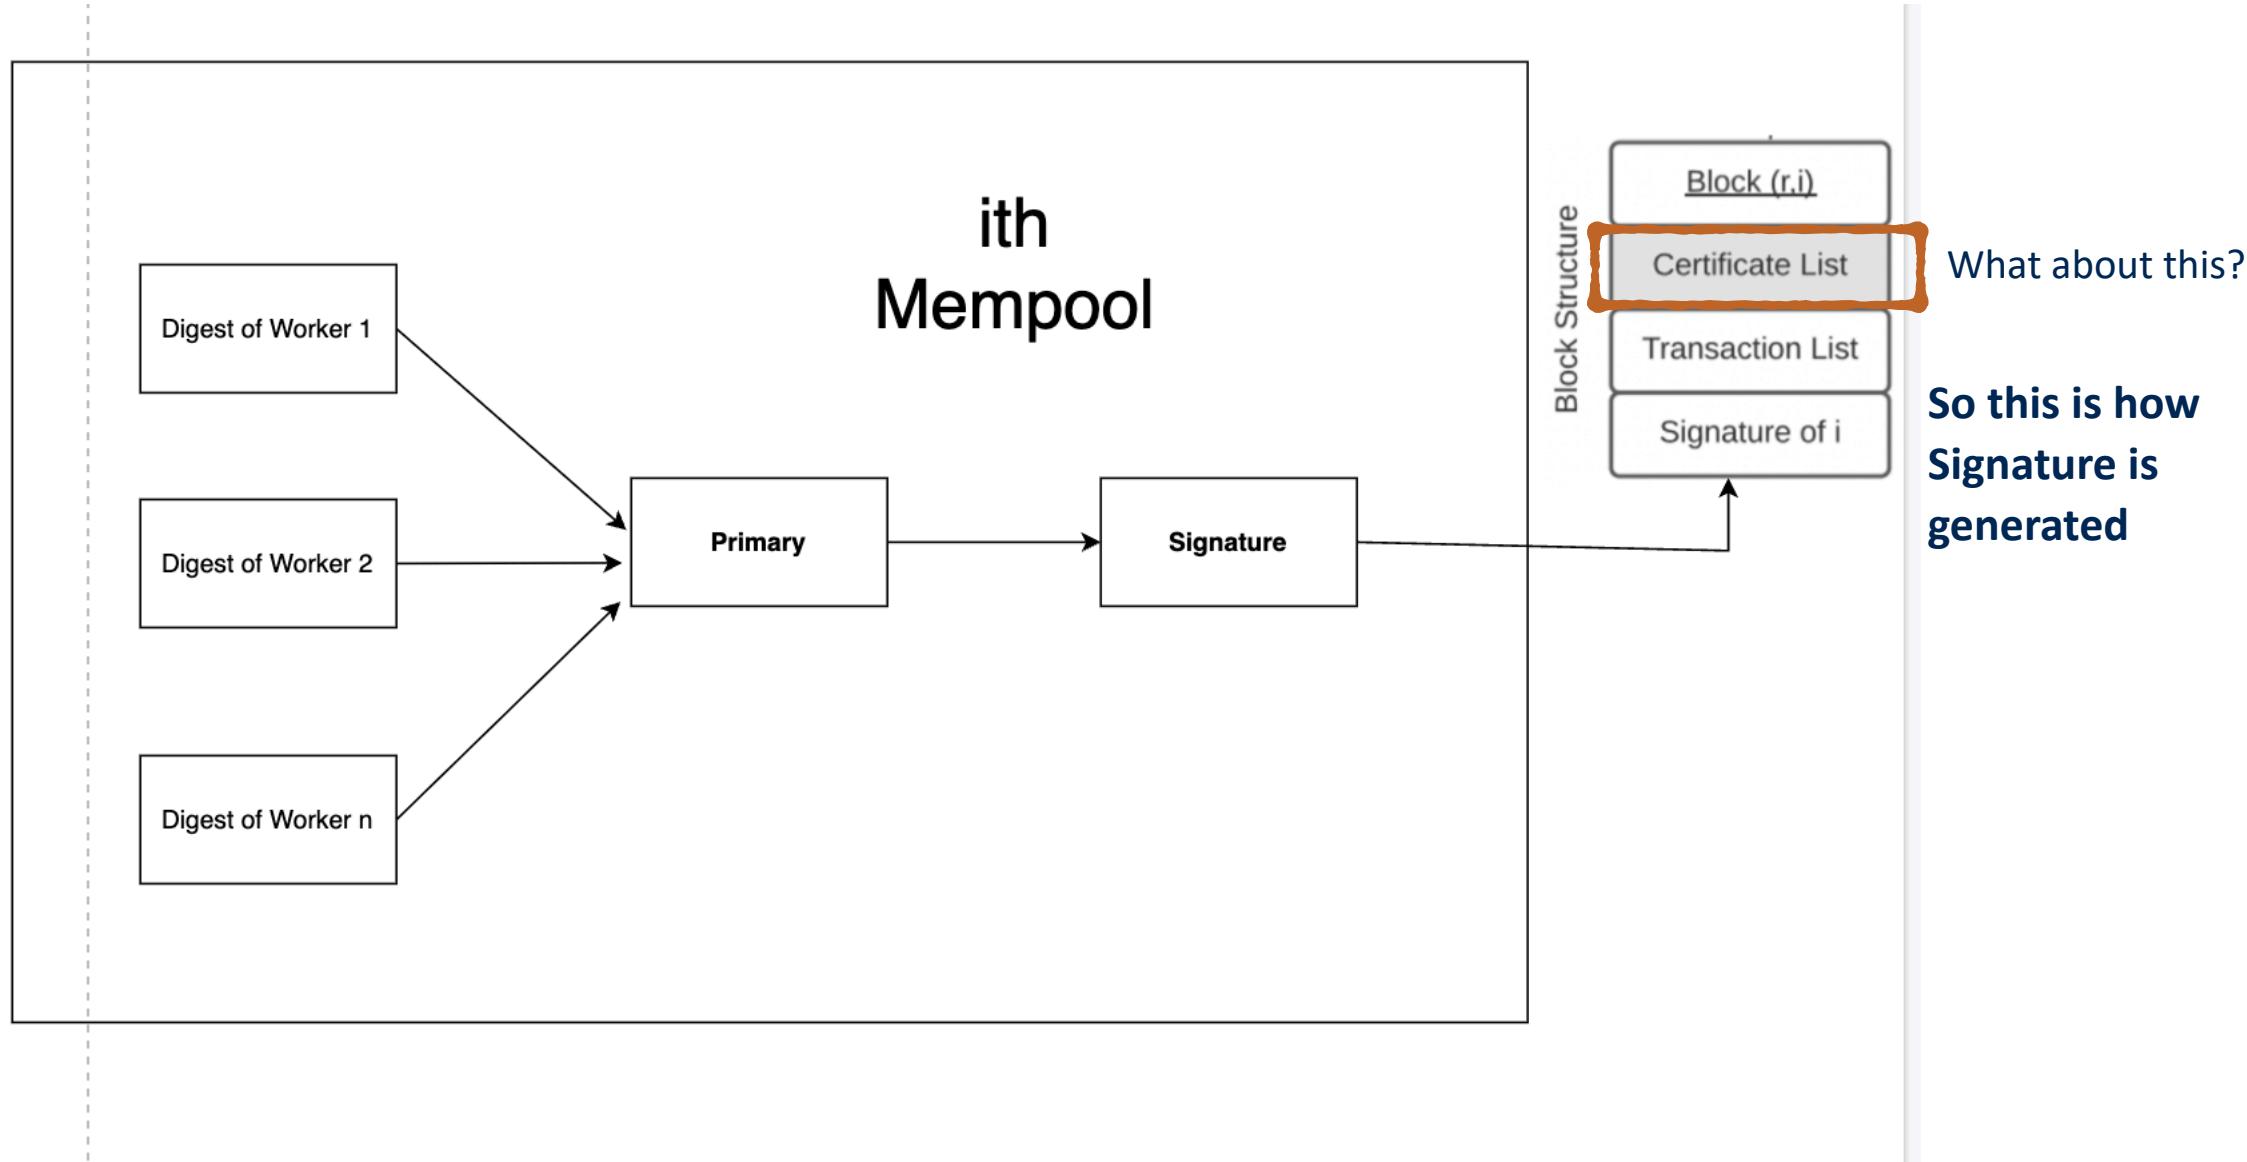


Signature of the Block

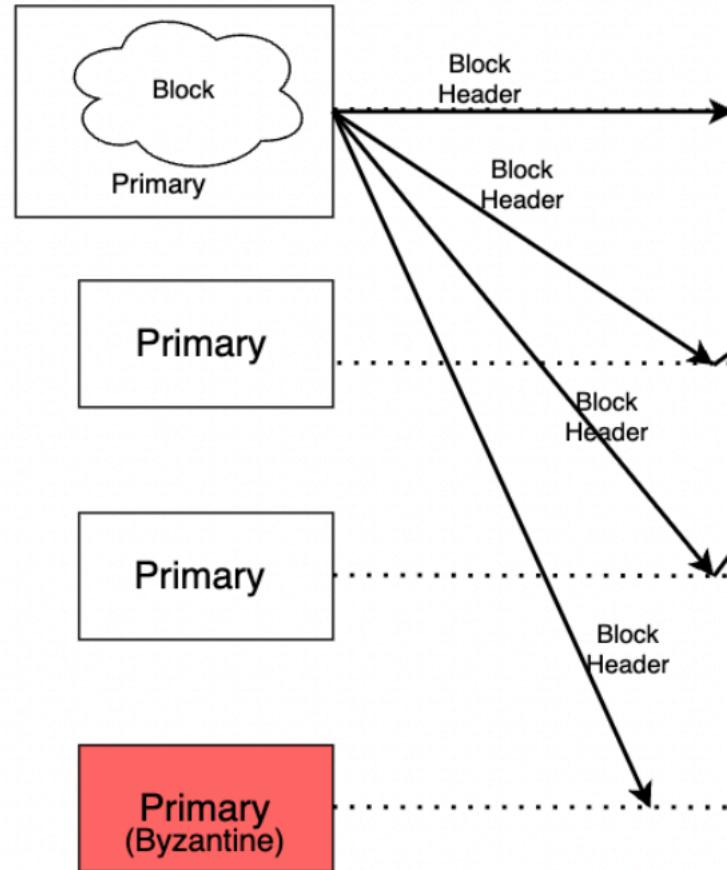


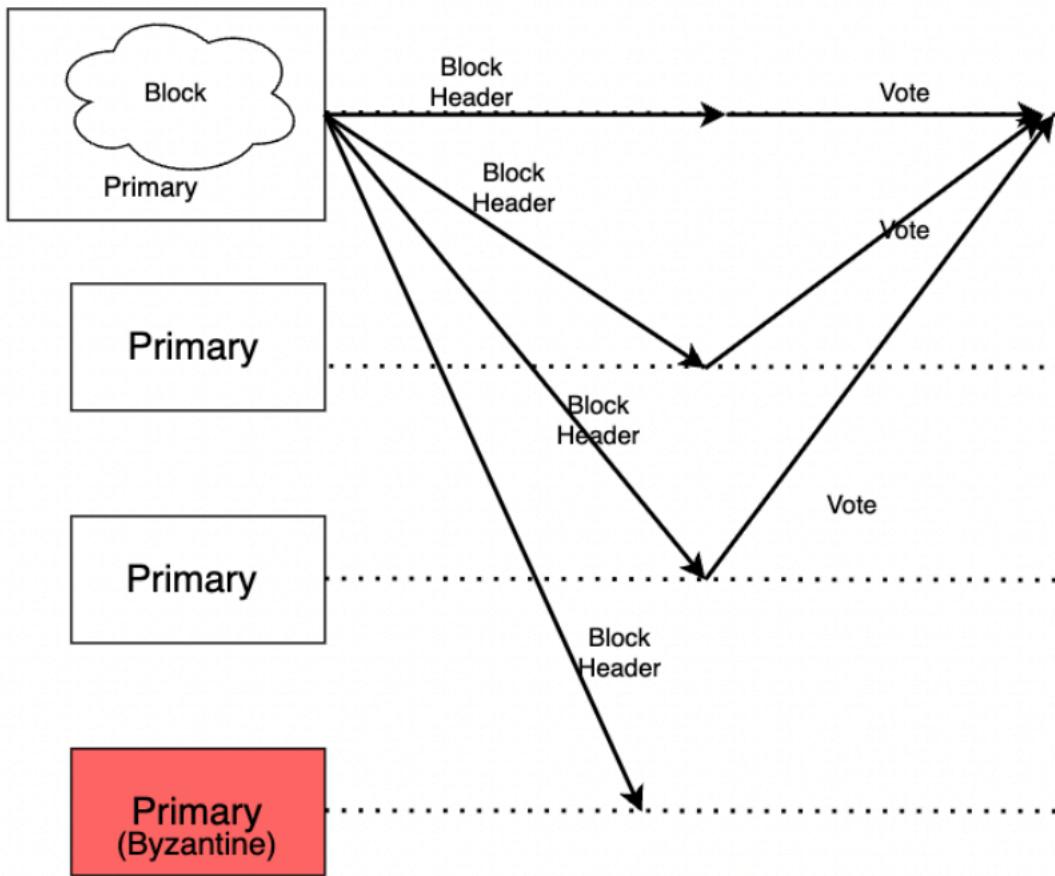
**So this is how
Signature is
generated**

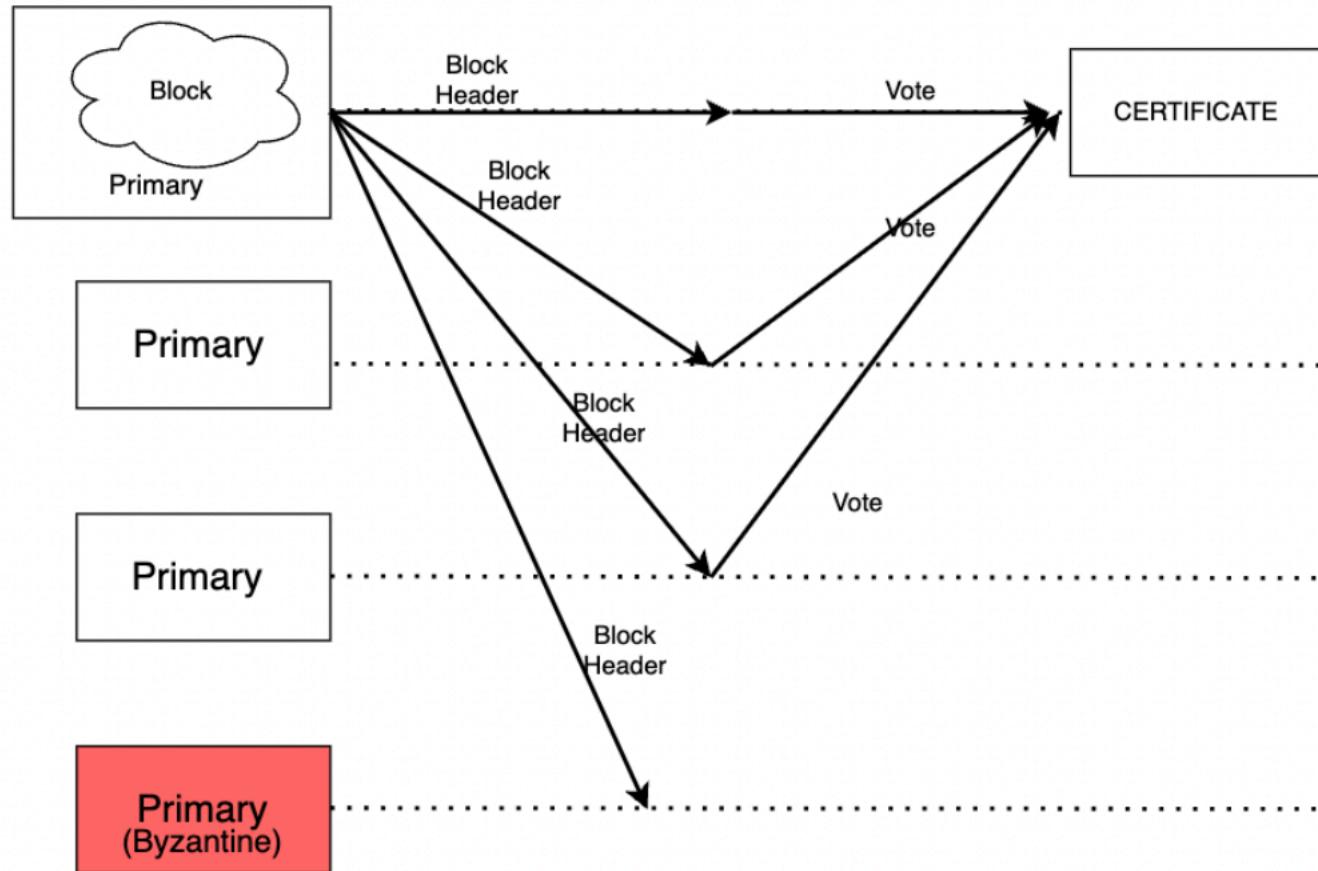
Signature of the Block

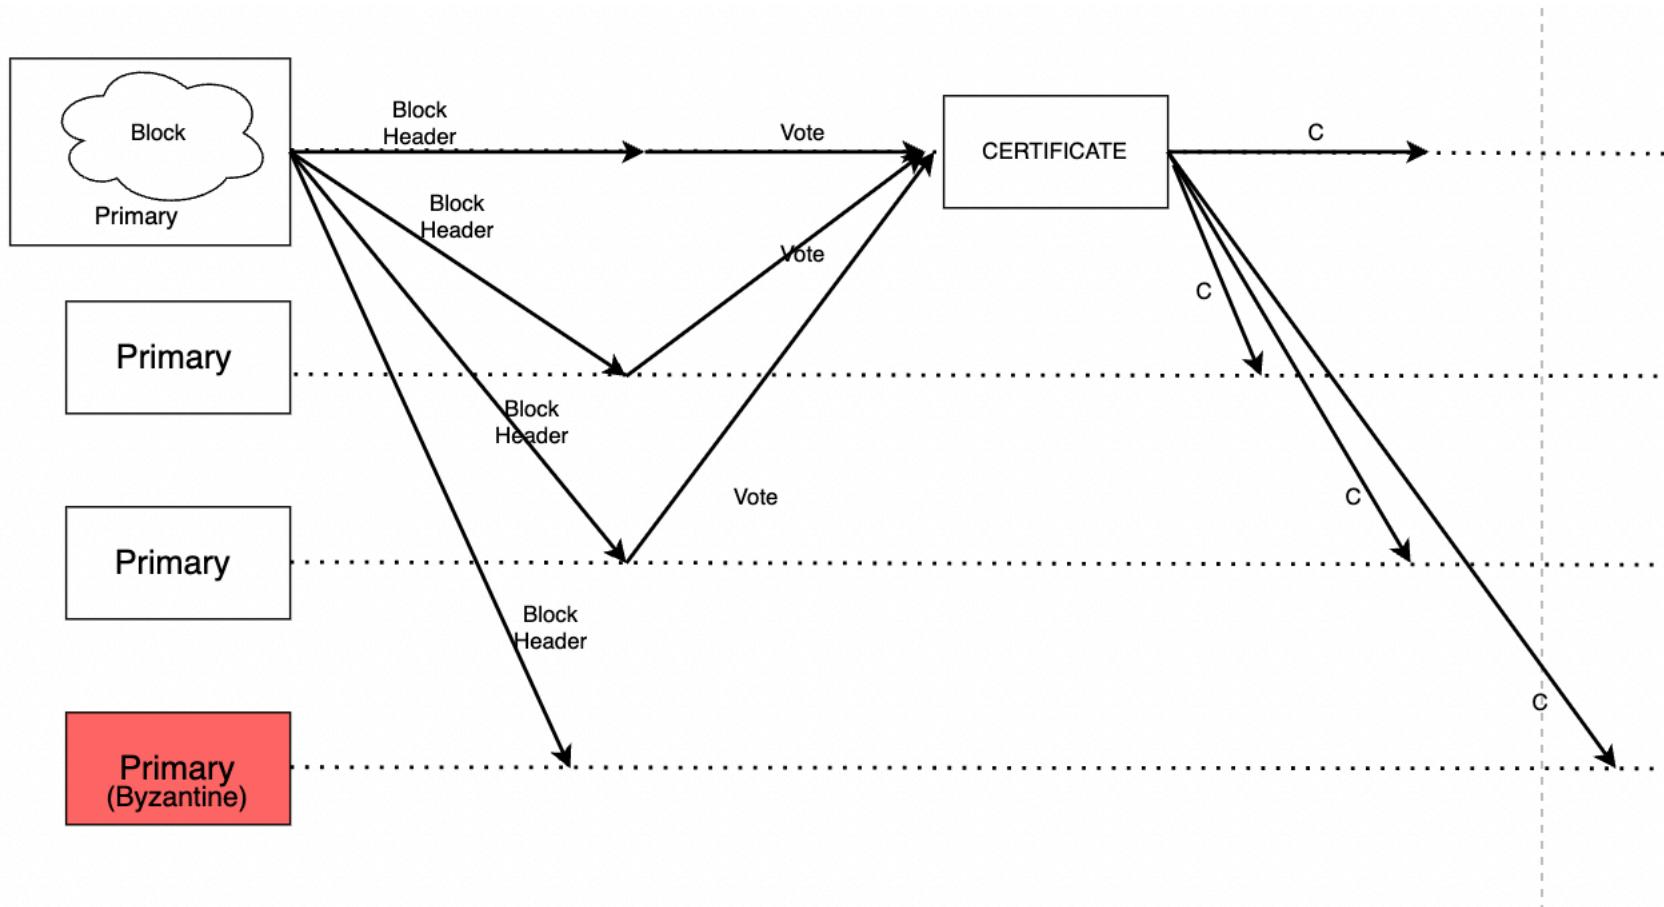


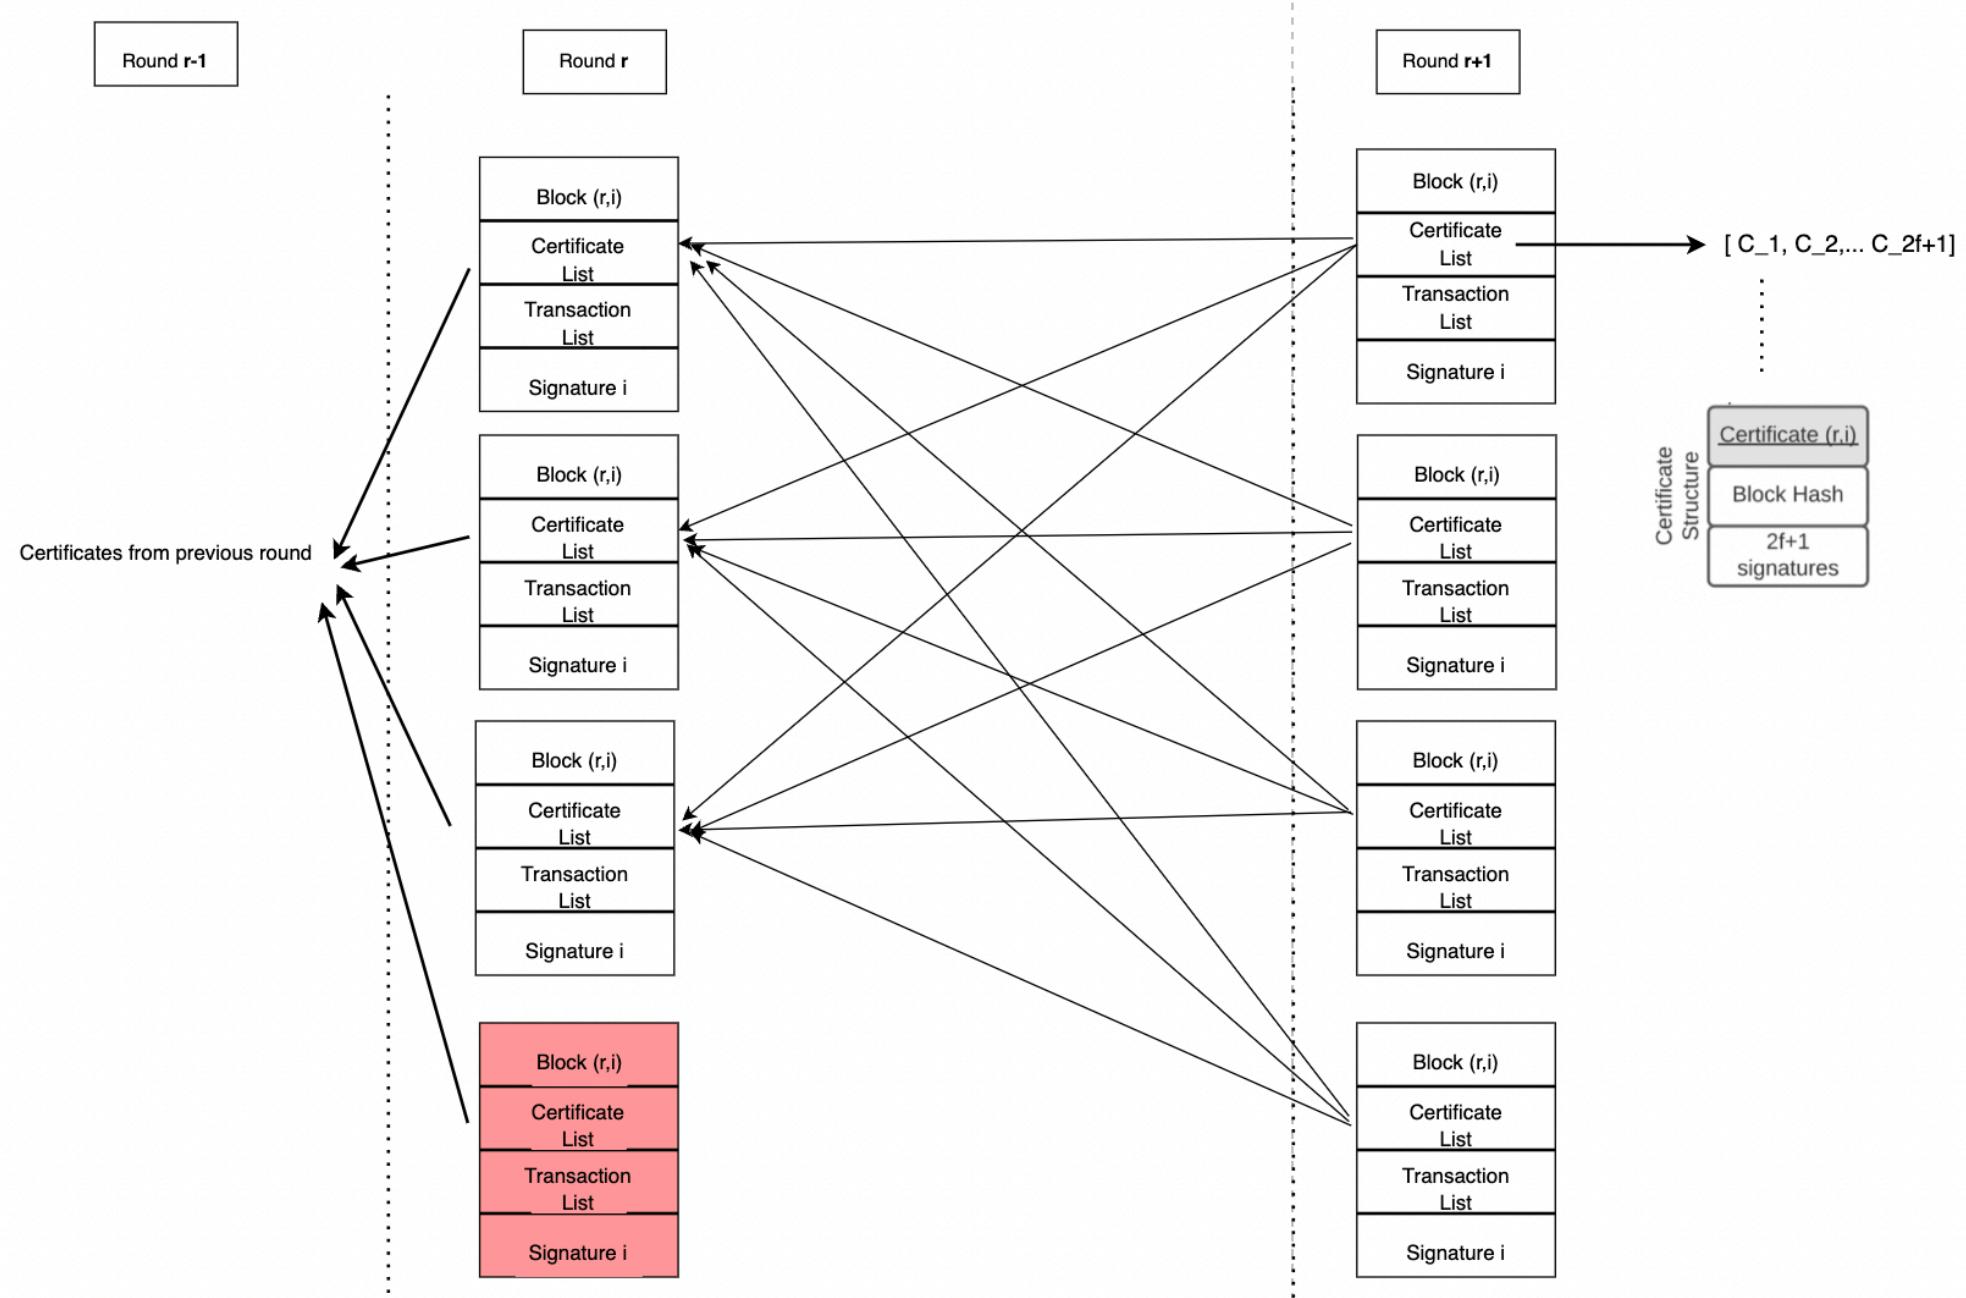
Where are these certificates coming from?



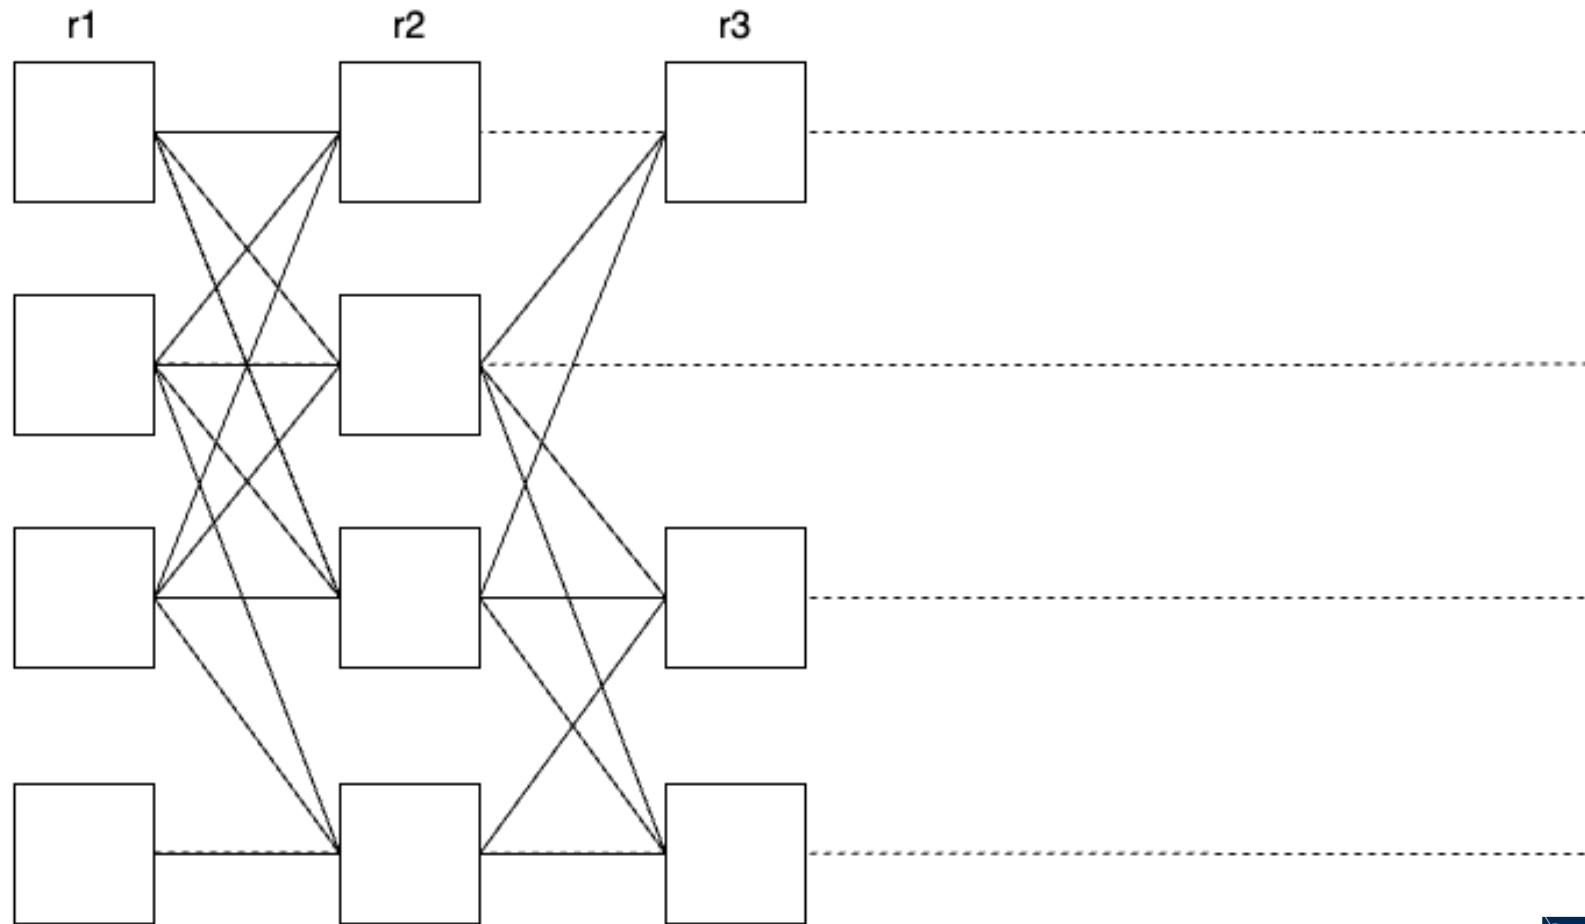








This is the DAG created in Narwhal



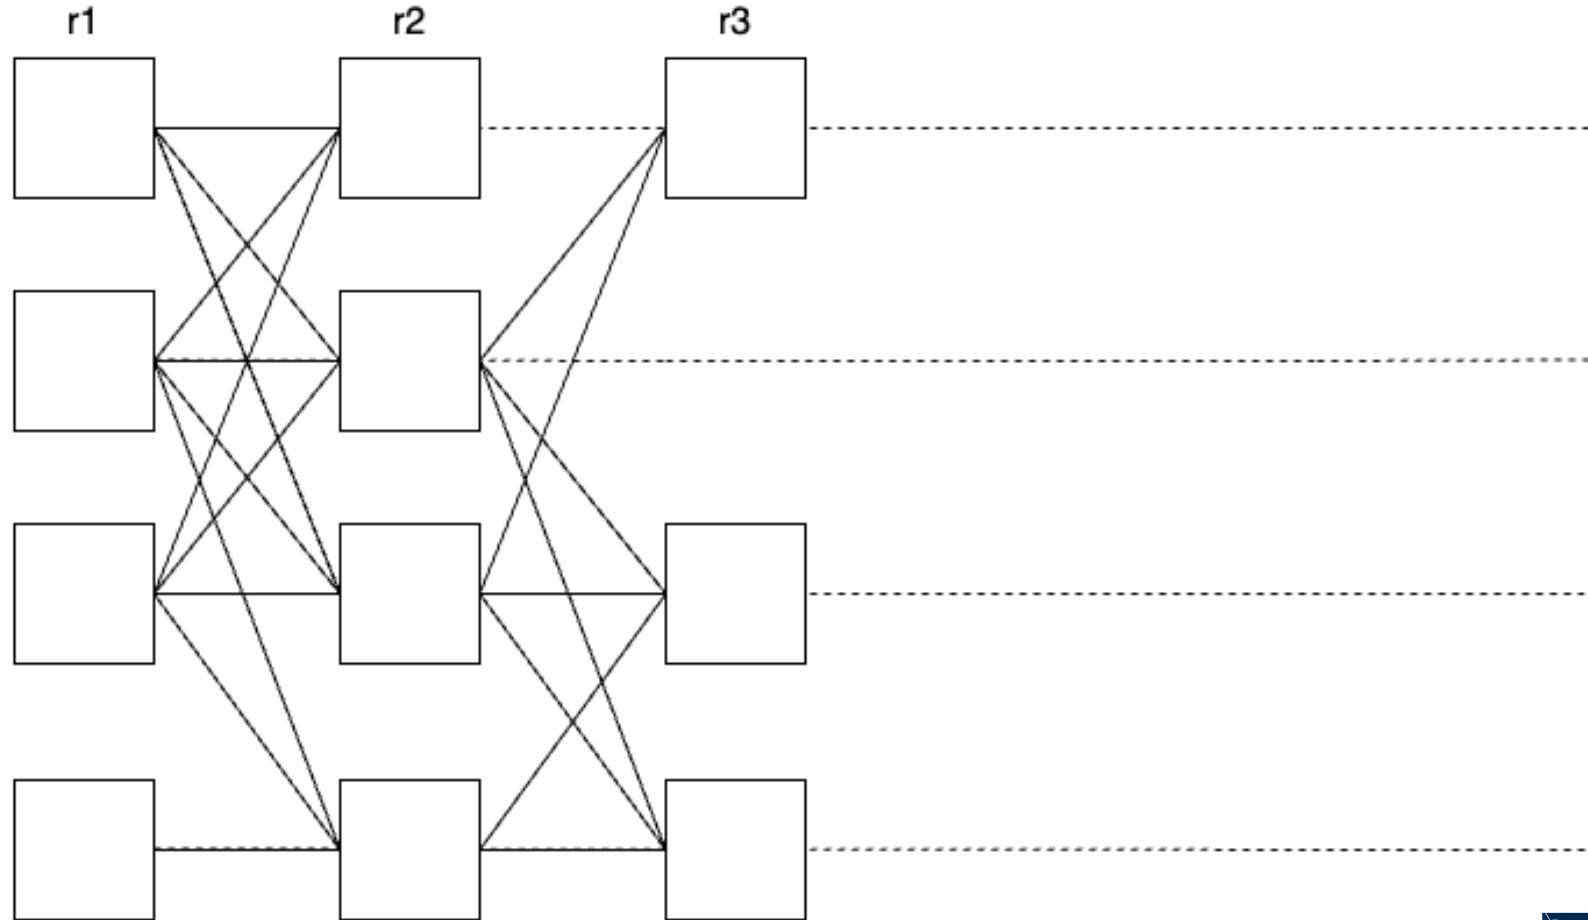
Tusk

What
is zero-message
overhead?

- A zero-message overhead asynchronous consensus protocol
- DAG → Starting point
- Three consecutive rounds
 - First round - Each validator proposes its block and **causal** history (**PROPOSE**)
 - Second round - Each validator votes on the proposal by including them in their block (**VOTE**)
 - Third round - Validators produce randomness to elect a random leader's block.

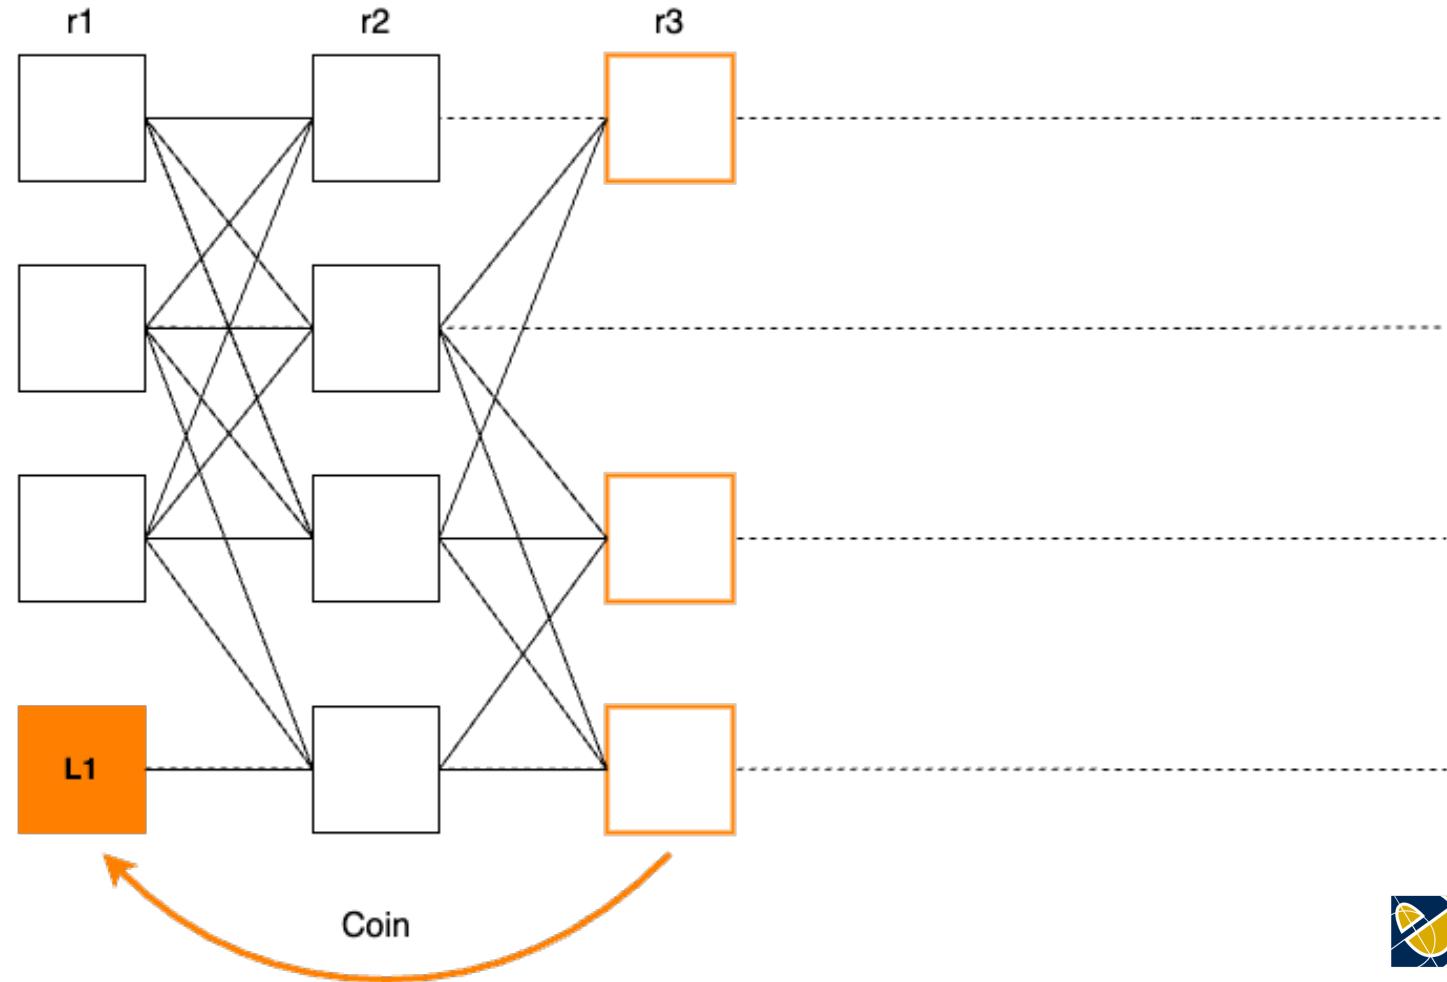
Tusk

Just interpret the DAG created in Narwhal



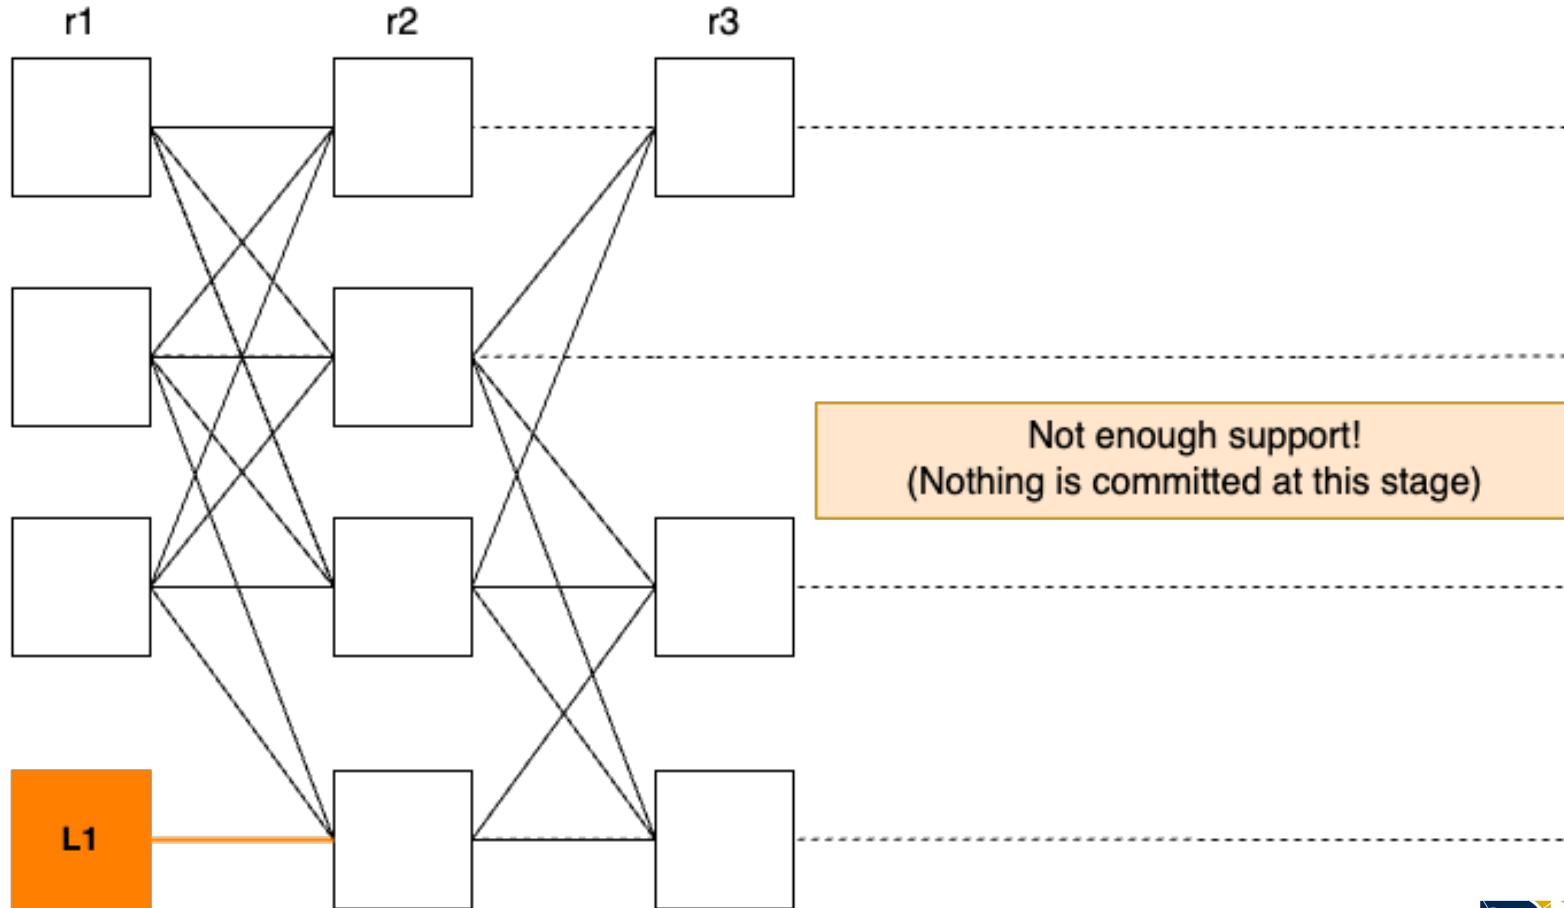
Tusk

The random coin elects the leader of r-2



Tusk

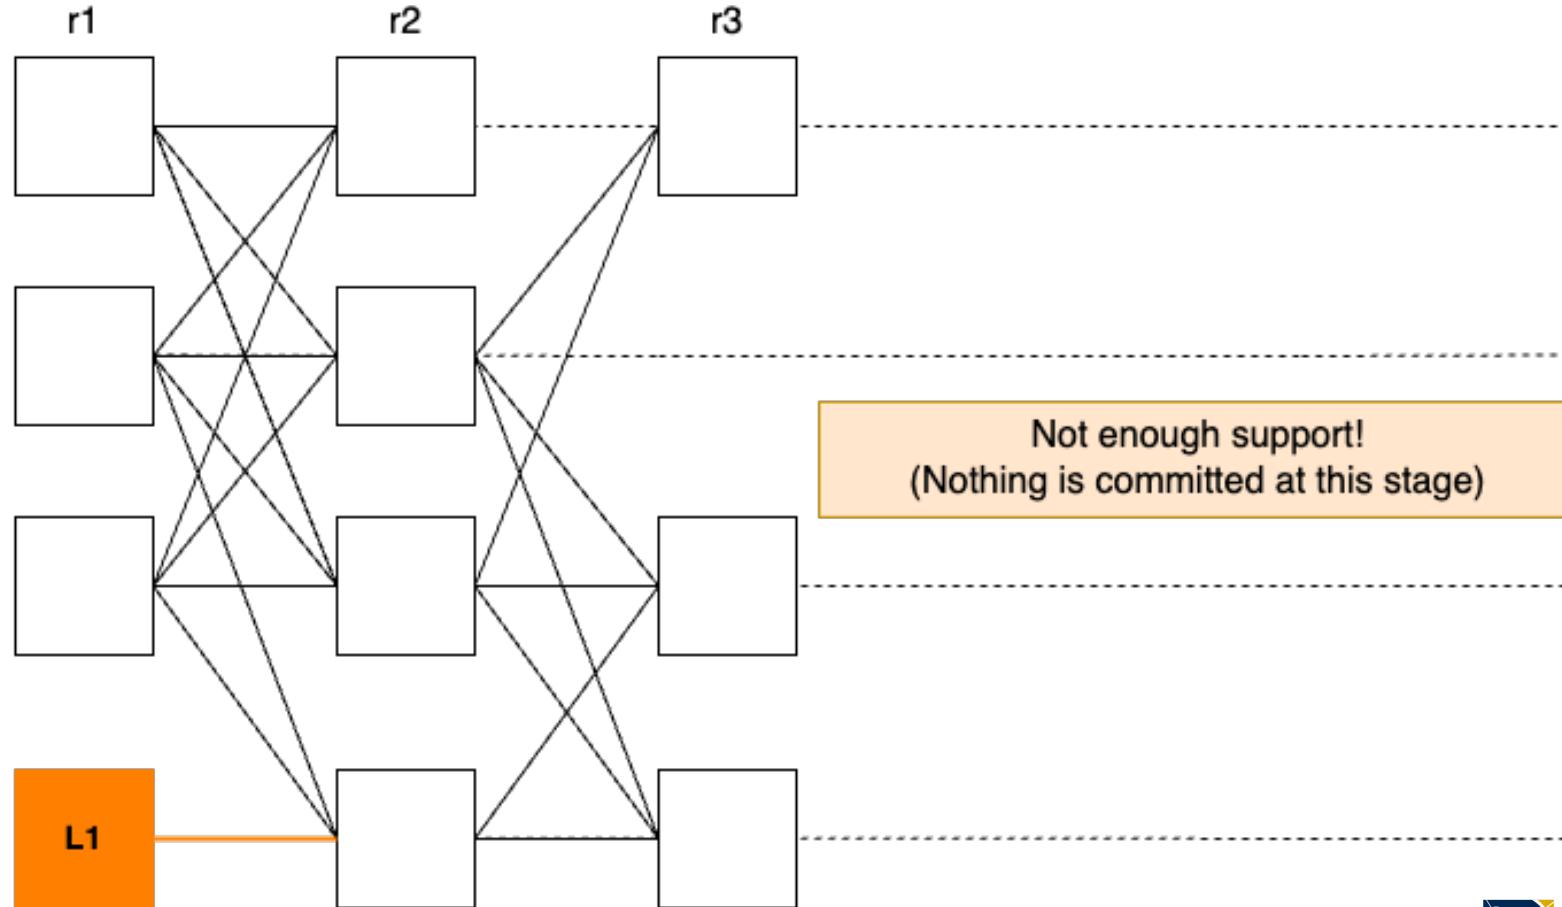
The chosen leader needs **f+1** links from round **r-1** to get committed



Tusk

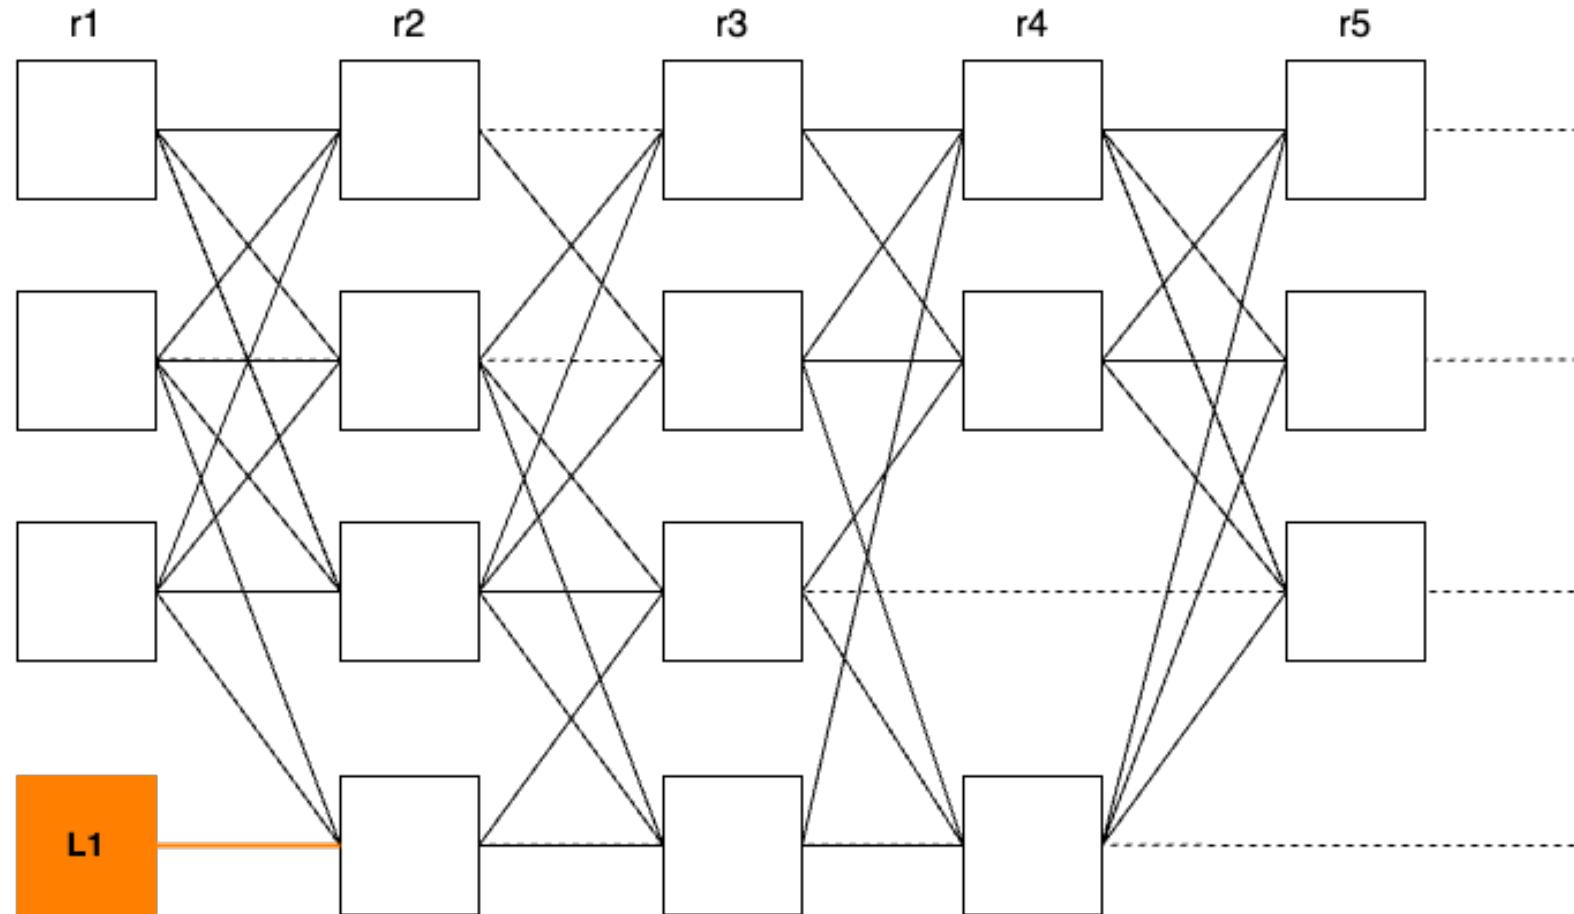
Let us discuss the intuition behind $f+1$ later.

The chosen leader needs $f+1$ links from round $r-1$ to get committed



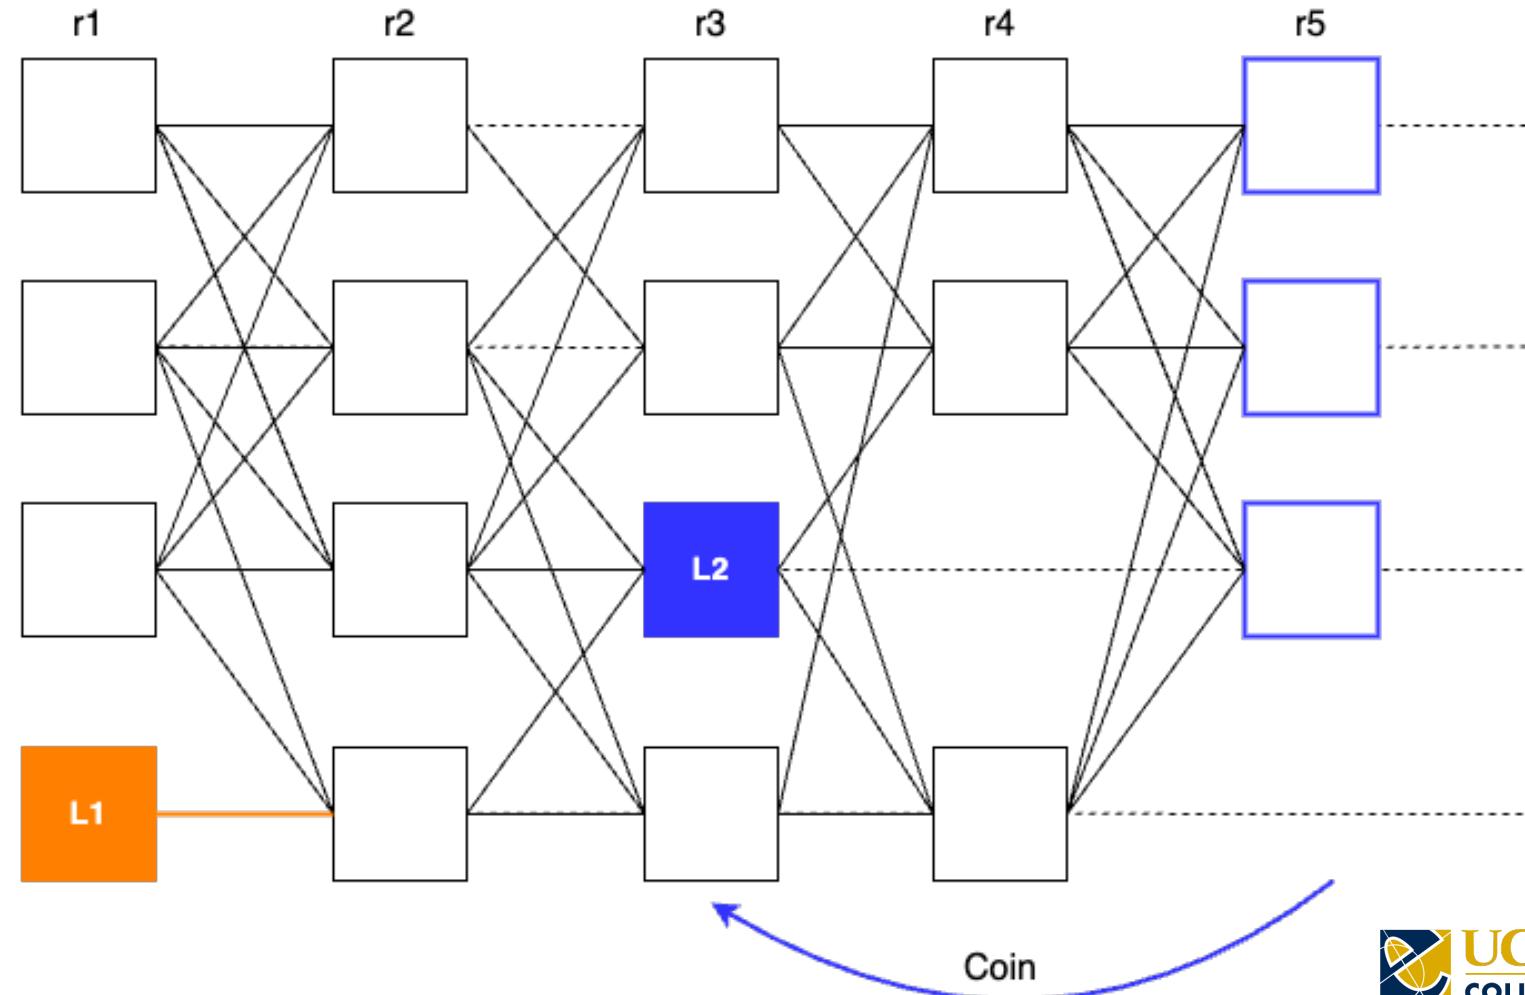
Tusk

Nothing is committed so we keep building the DAG



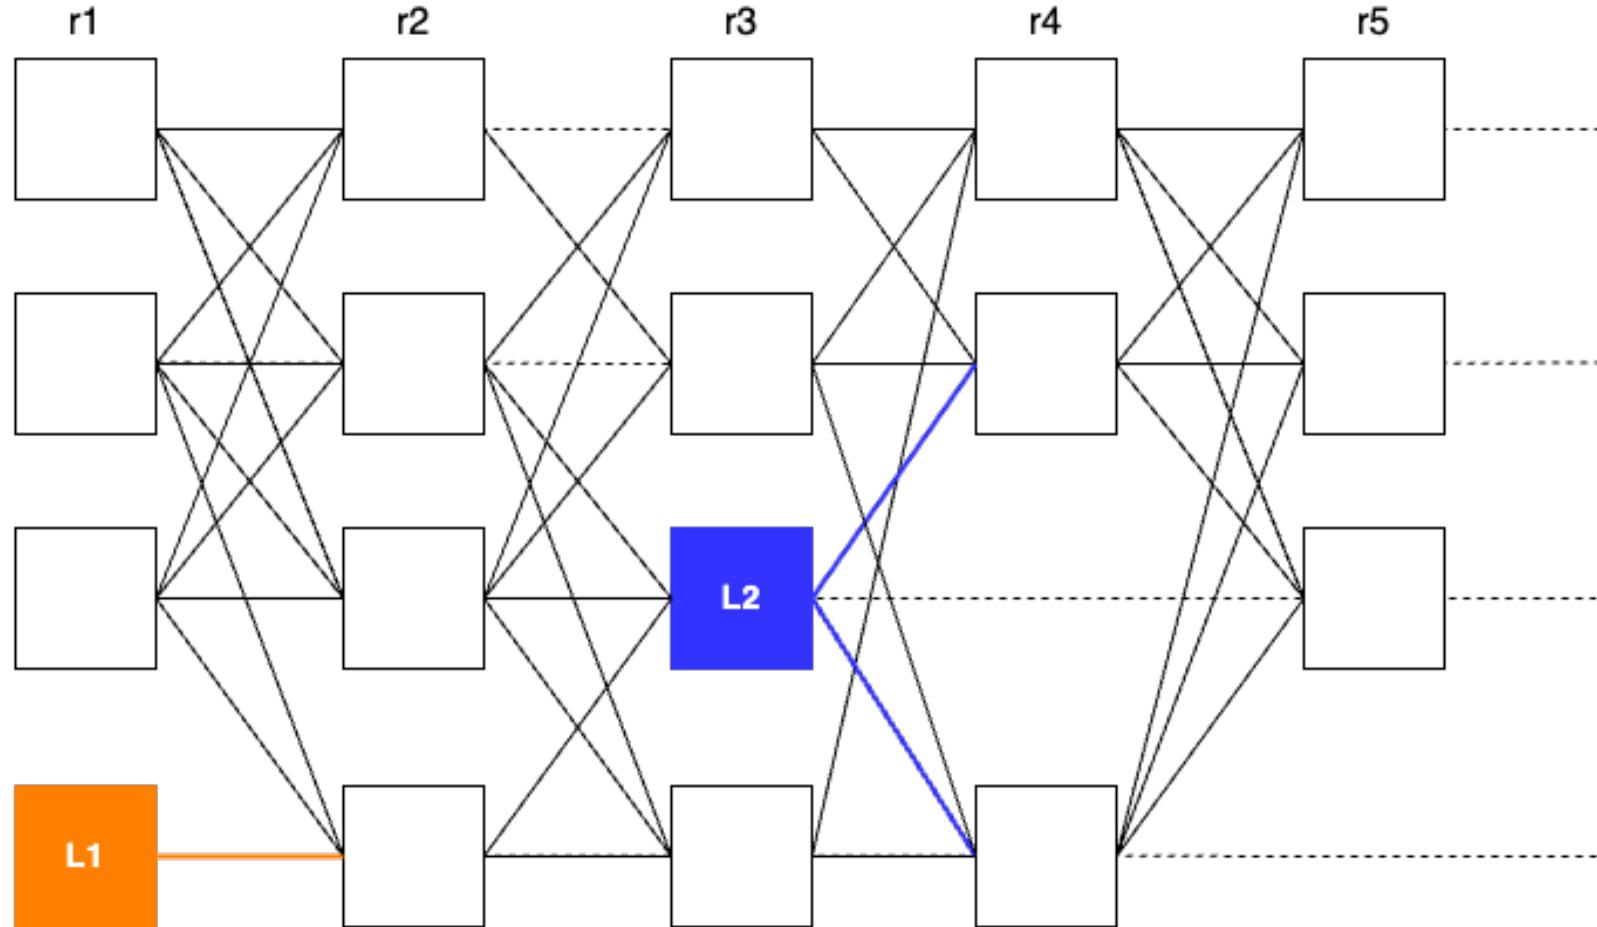
Tusk

Elect the leader of r3



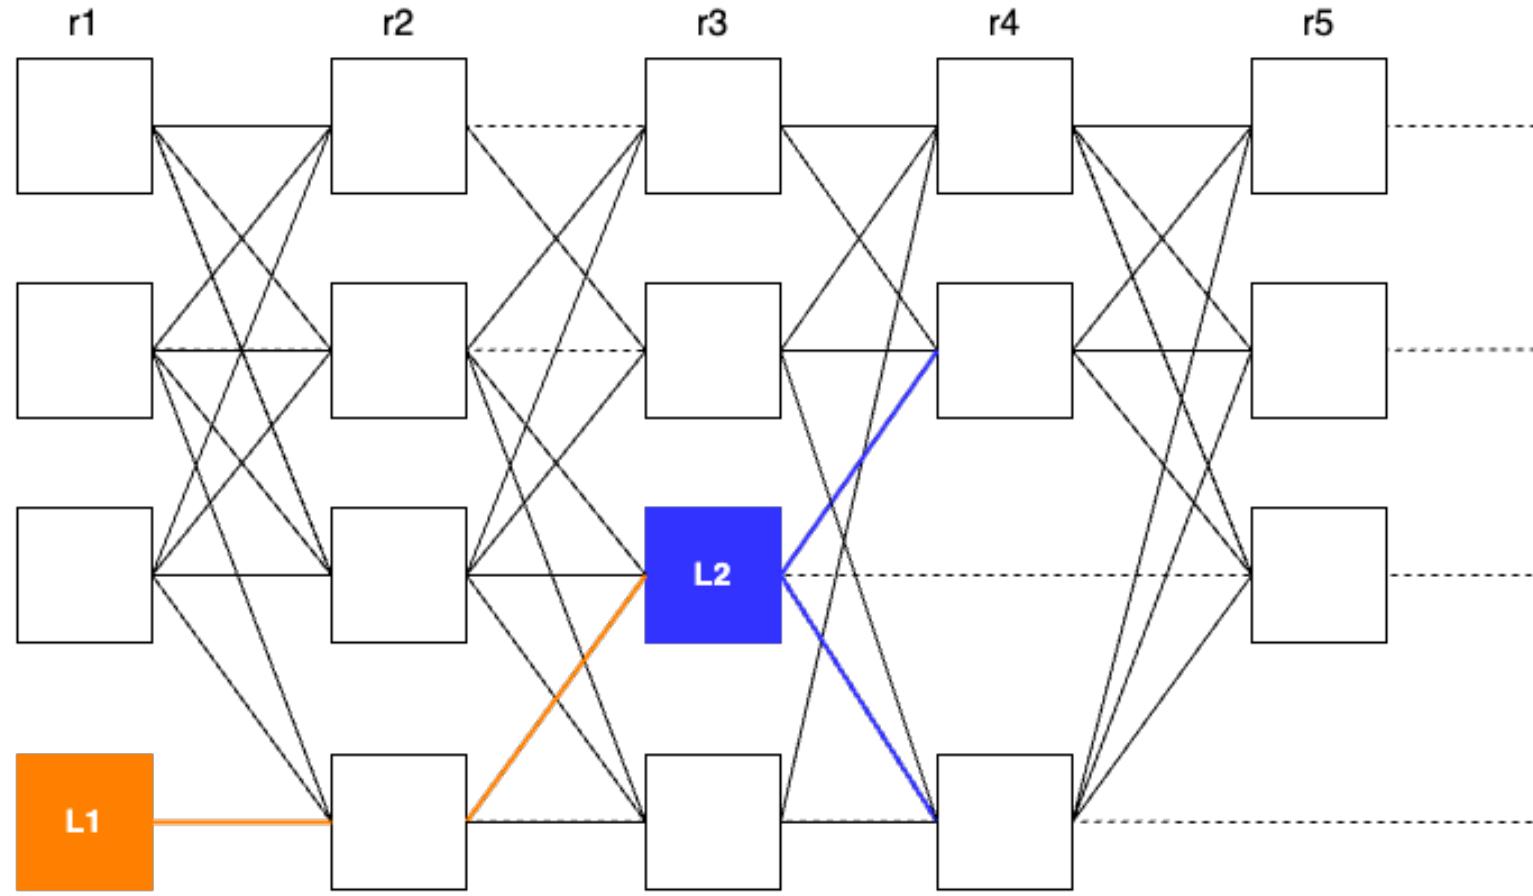
Tusk

Leader L2 has enough support



Tusk

Leader L2 has links to Leader L1

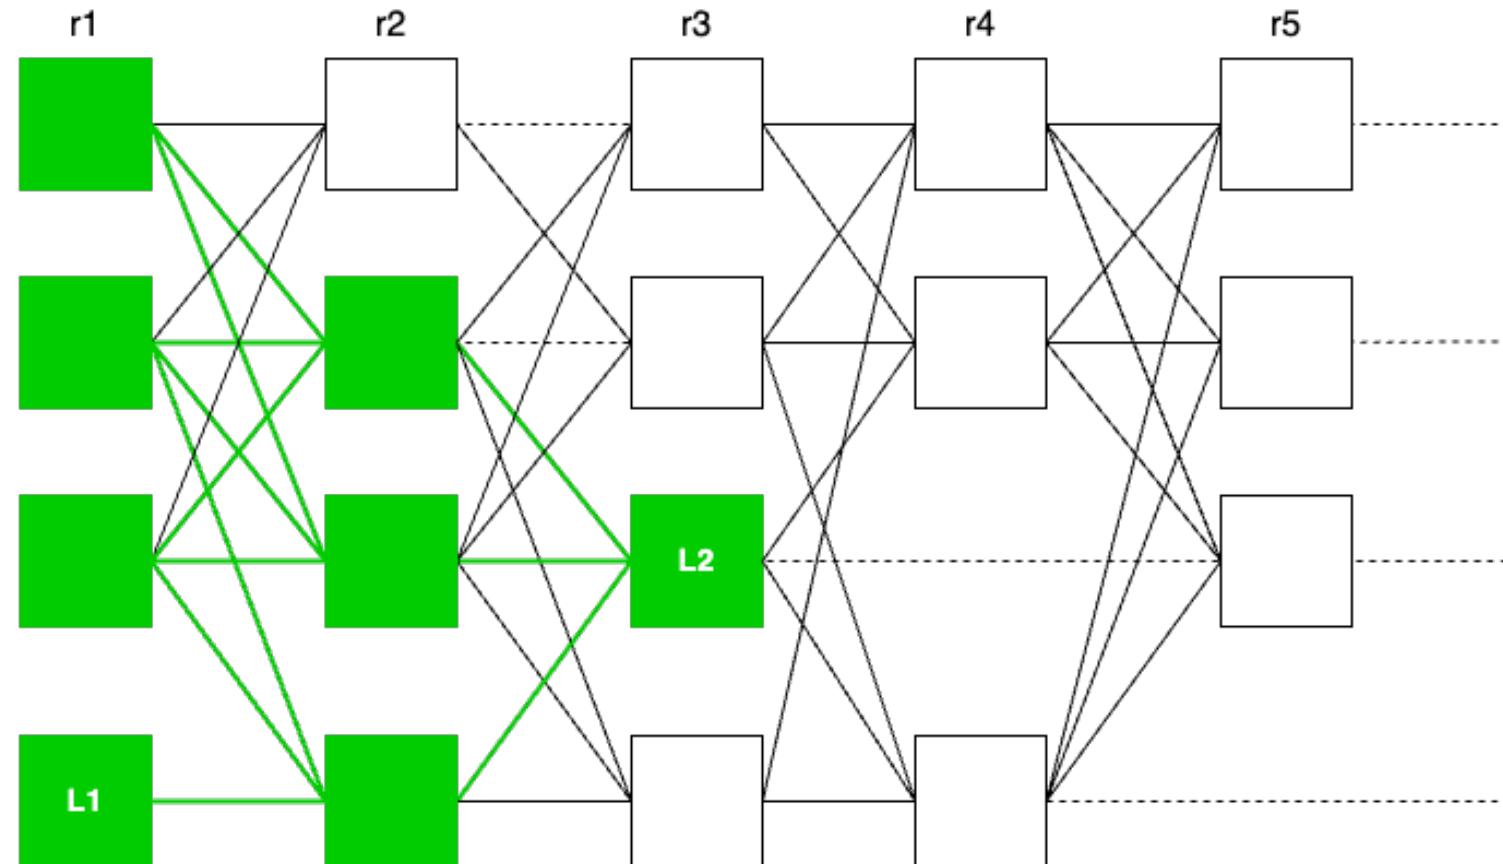


First commit L1

Then commit L2

Tusk

Commit all the sub-DAG of the leader



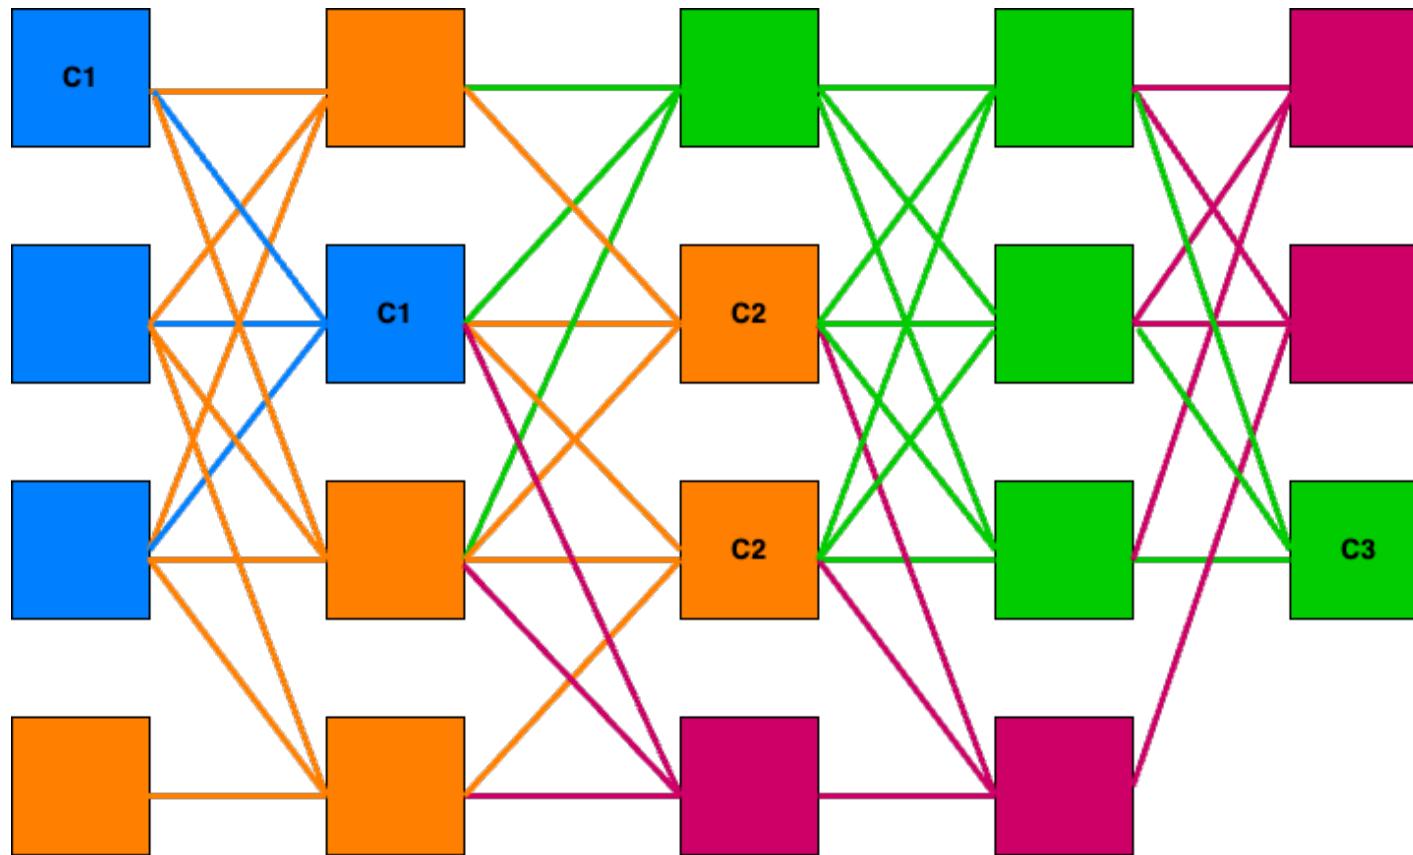
What is the intuition behind $f+1$?

What is the intuition behind $f+1$?

For safety

How to ensure liveness?

HotStuff on Narwhal



Implementation and Metrics

- Written in Rust
- Networking: Tokio (TCP)
- Storage: RocksDB
- Cryptography: ed25519-dalek

<https://github.com/facebookresearch/narwhal>

Evaluation - Experimental setup on AWS



Evaluation - Throughput latency graph

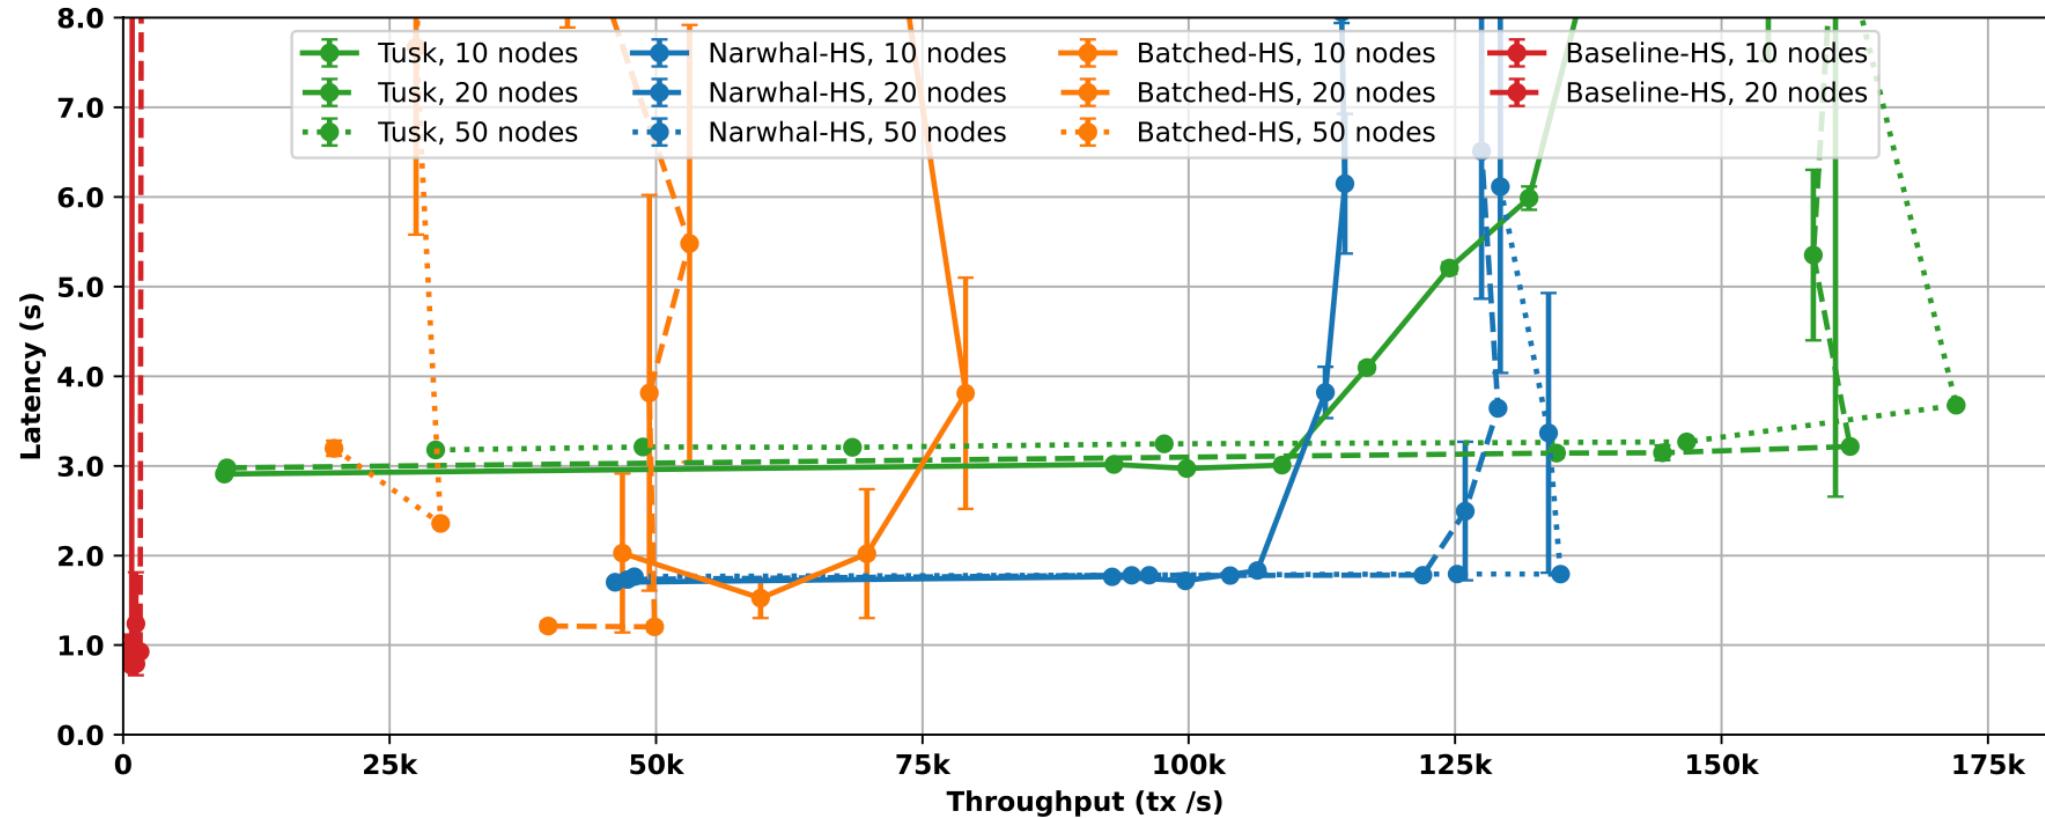
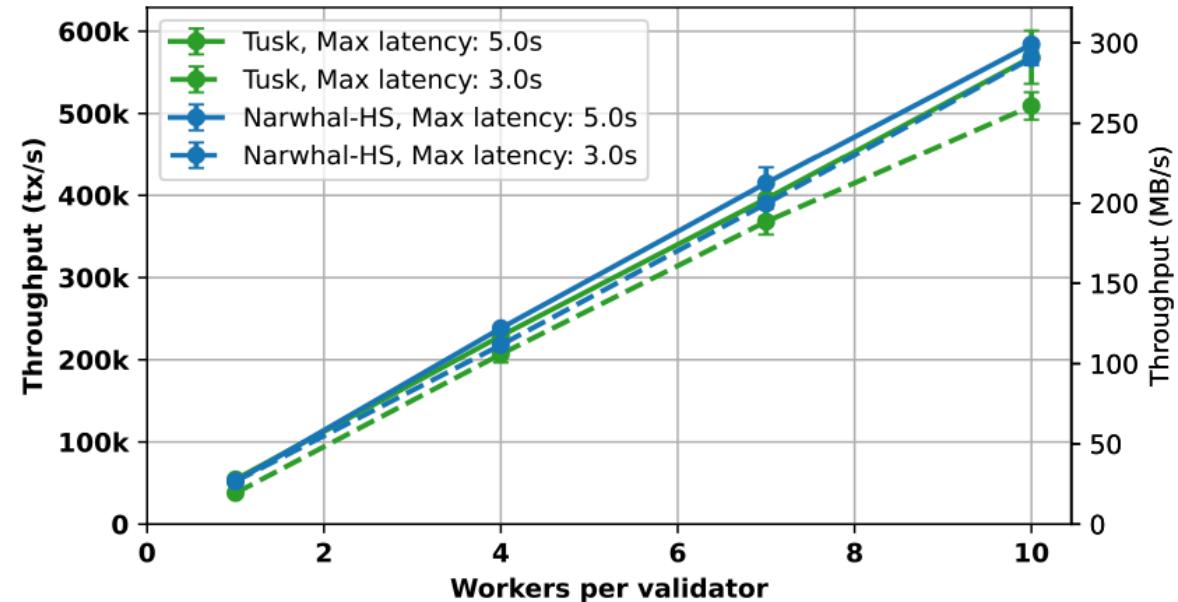
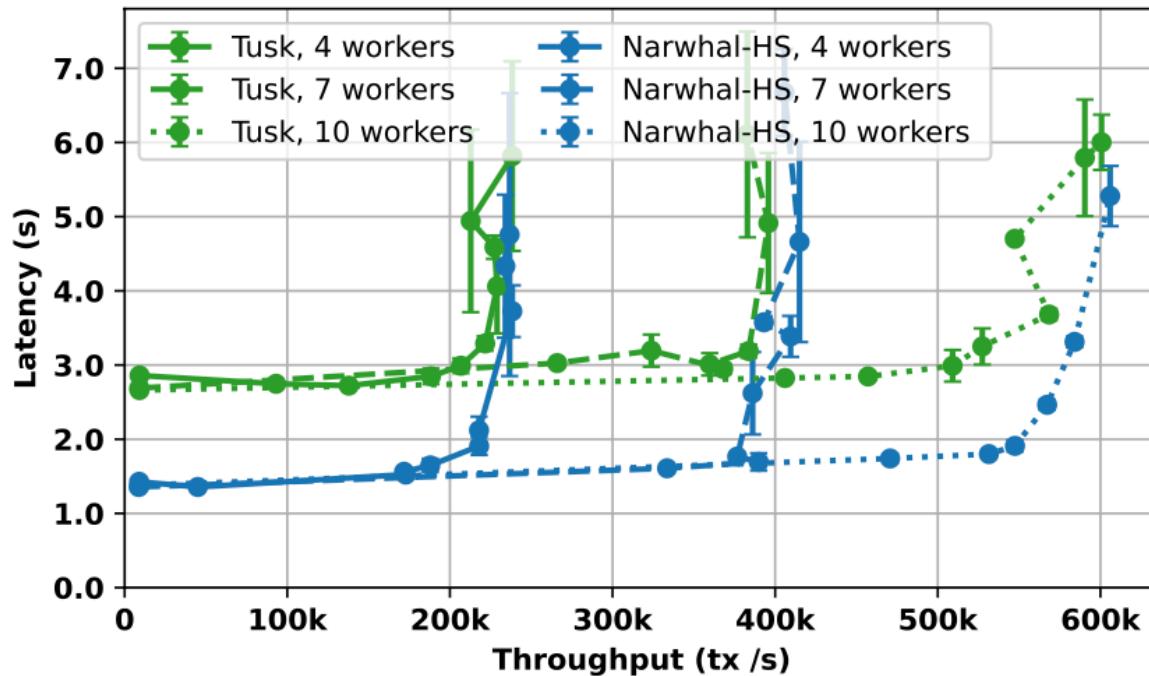


Figure 6. Comparative throughput-latency performance for the novel Narwhal-HotStuff, Tusk, batched-HotStuff and the baseline HotStuff. WAN measurements with 10, 20, and 50 validators, using 1 worker collocated with the primary. No validator faults, 500KB max. block size and 512B transaction size.

Evaluation - Scalability

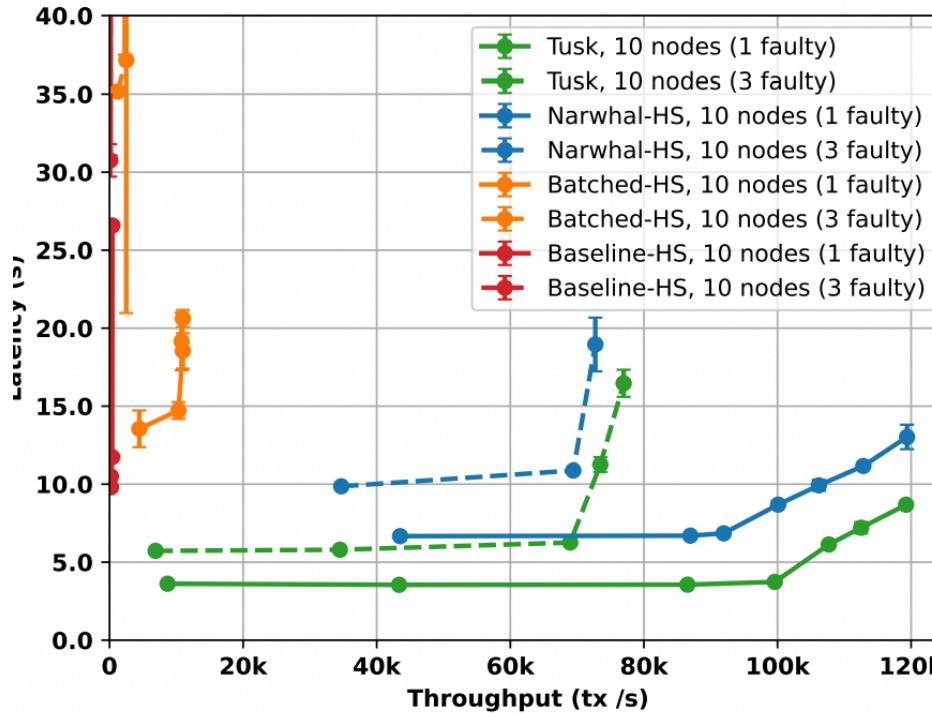


Conclusion

NARWHAL AND TUSK

- Separate consensus and data dissemination for high performance.
- Scalable design, egalitarian resource utilizations.

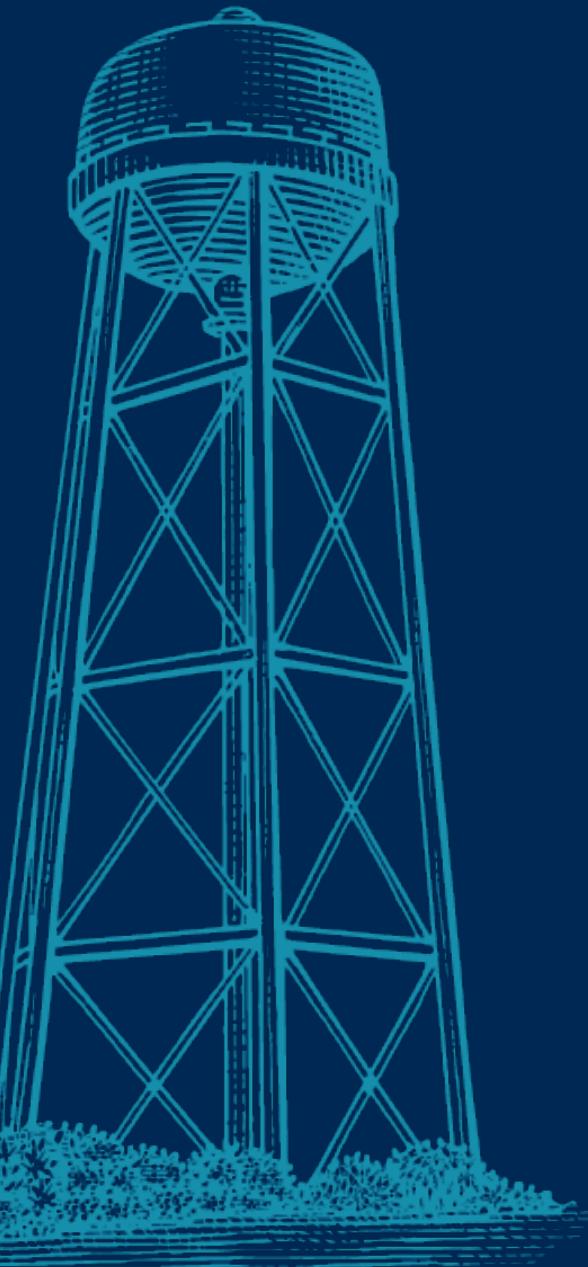
Performance under faults



Comparative throughput-latency under faults. WAN measurements with 10 validators using 1 worker co-located with the primary. 1 and 3 faults, 500KB max block size, and 512B transaction size.

References

- Danezis, G., Kogias, E. K., Sonnino, & Spiegelman, A. (2021). Narwhal and Tusk: A DAG-based Mempool and Efficient BFT Consensus. *arXiv*. <https://doi.org/10.48550/arXiv.2105.11827>
- <https://decentralizedthoughts.github.io/2022-06-28-DAG-meets-BFT/>
- <https://research.protocol.ai/sites/consensusday21/programme/>
- <https://docs.sui.io/learn/architecture/consensus>
- <https://eprint.iacr.org/2018/992.pdf>
- <https://salemal.medium.com/a-review-of-narwhal-and-tusk-a-dag-based-mempool-and-efficient-bft-consensus-bf2099908c63>
- <https://www.blocknative.com/blog/mempool-intro>
- <https://blog.unibulmerchantservices.com/how-visas-payment-system-works/>



THANK YOU!

ANY QUESTIONS? ASK AWAY!