

# **HOTSTUFF**

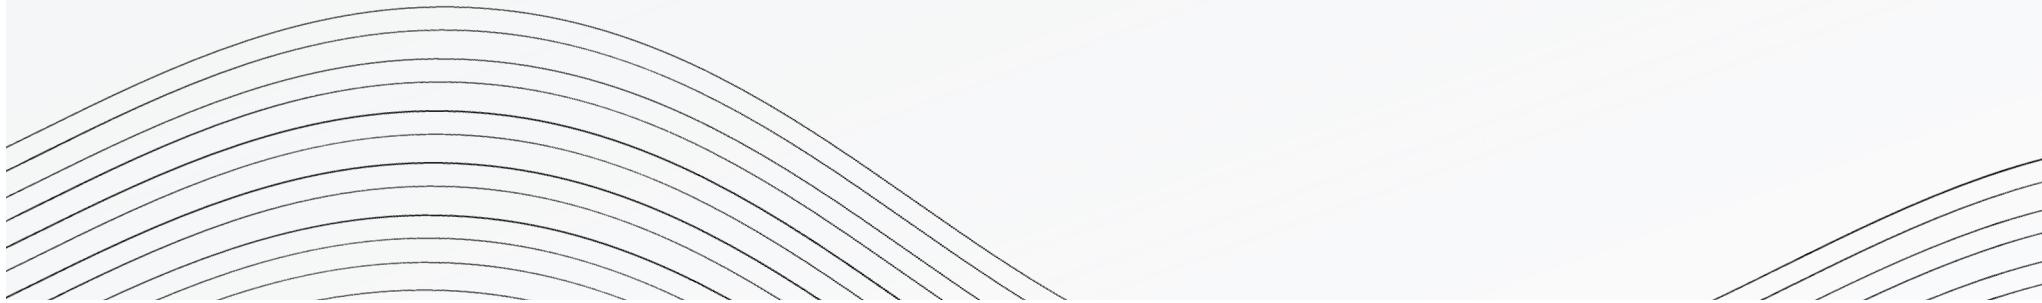
## **Consensus in the Lens of Blockchain.**

**PRESENTATION BY-**

**Amy Vu,  
Alyssa Yee,  
Alex Gao,  
Guransh Singh Anand**

# BACKGROUND

- PBFT Supports  $n \geq 3f+1$ , where  $n$  = total nodes,  $f$  = # of faulty nodes
- BFT(precursor to PBFT) was designed for system size  $n=4$  or  $n=7$ , not 1000's of nodes in modern systems



# PROBLEM?

- Not practical, makes many assumptions
- Provides inefficient way of solving problem (only theoretical solution)
- Does not assume delays in the network

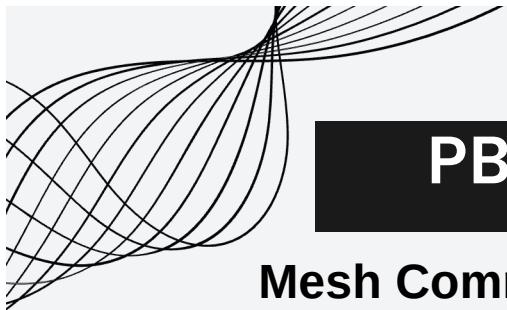
PBFT made to be used in practical environment,  
BFT made for theoretical concept

# WHY HOTSTUFF?

01

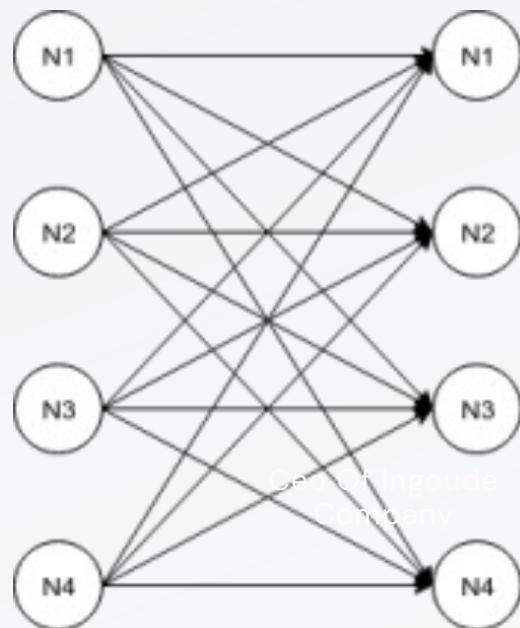
**Linear View Change:** Reduces communication complexity

Protocol	Normal Operation	View-Change (Failure)	Continuous Failure
PBFT	$O(n^2)$	$O(n^3)$	$O(fn^3)$
HotStuff	$O(n)$	$O(n)$	$O(fn)$



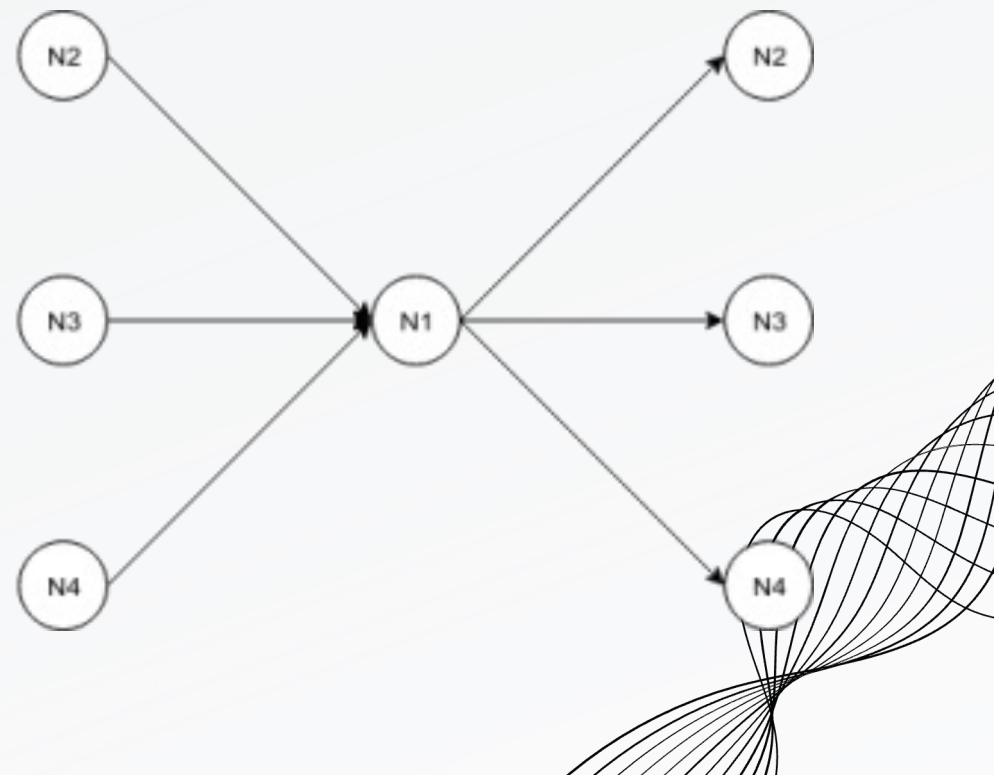
# PBFT

Mesh Communication Network



# Hotstuff

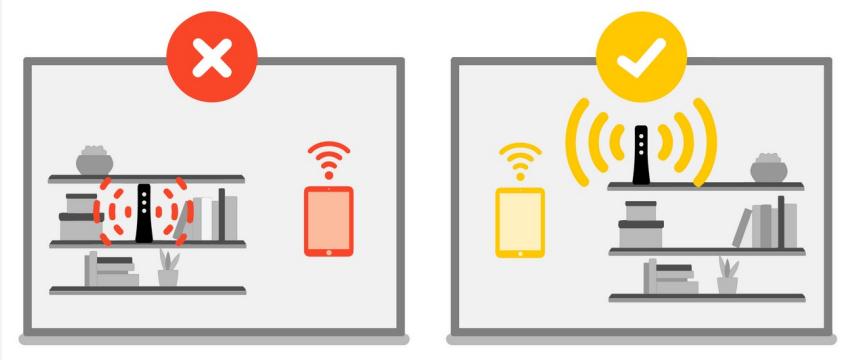
Star Communication Network

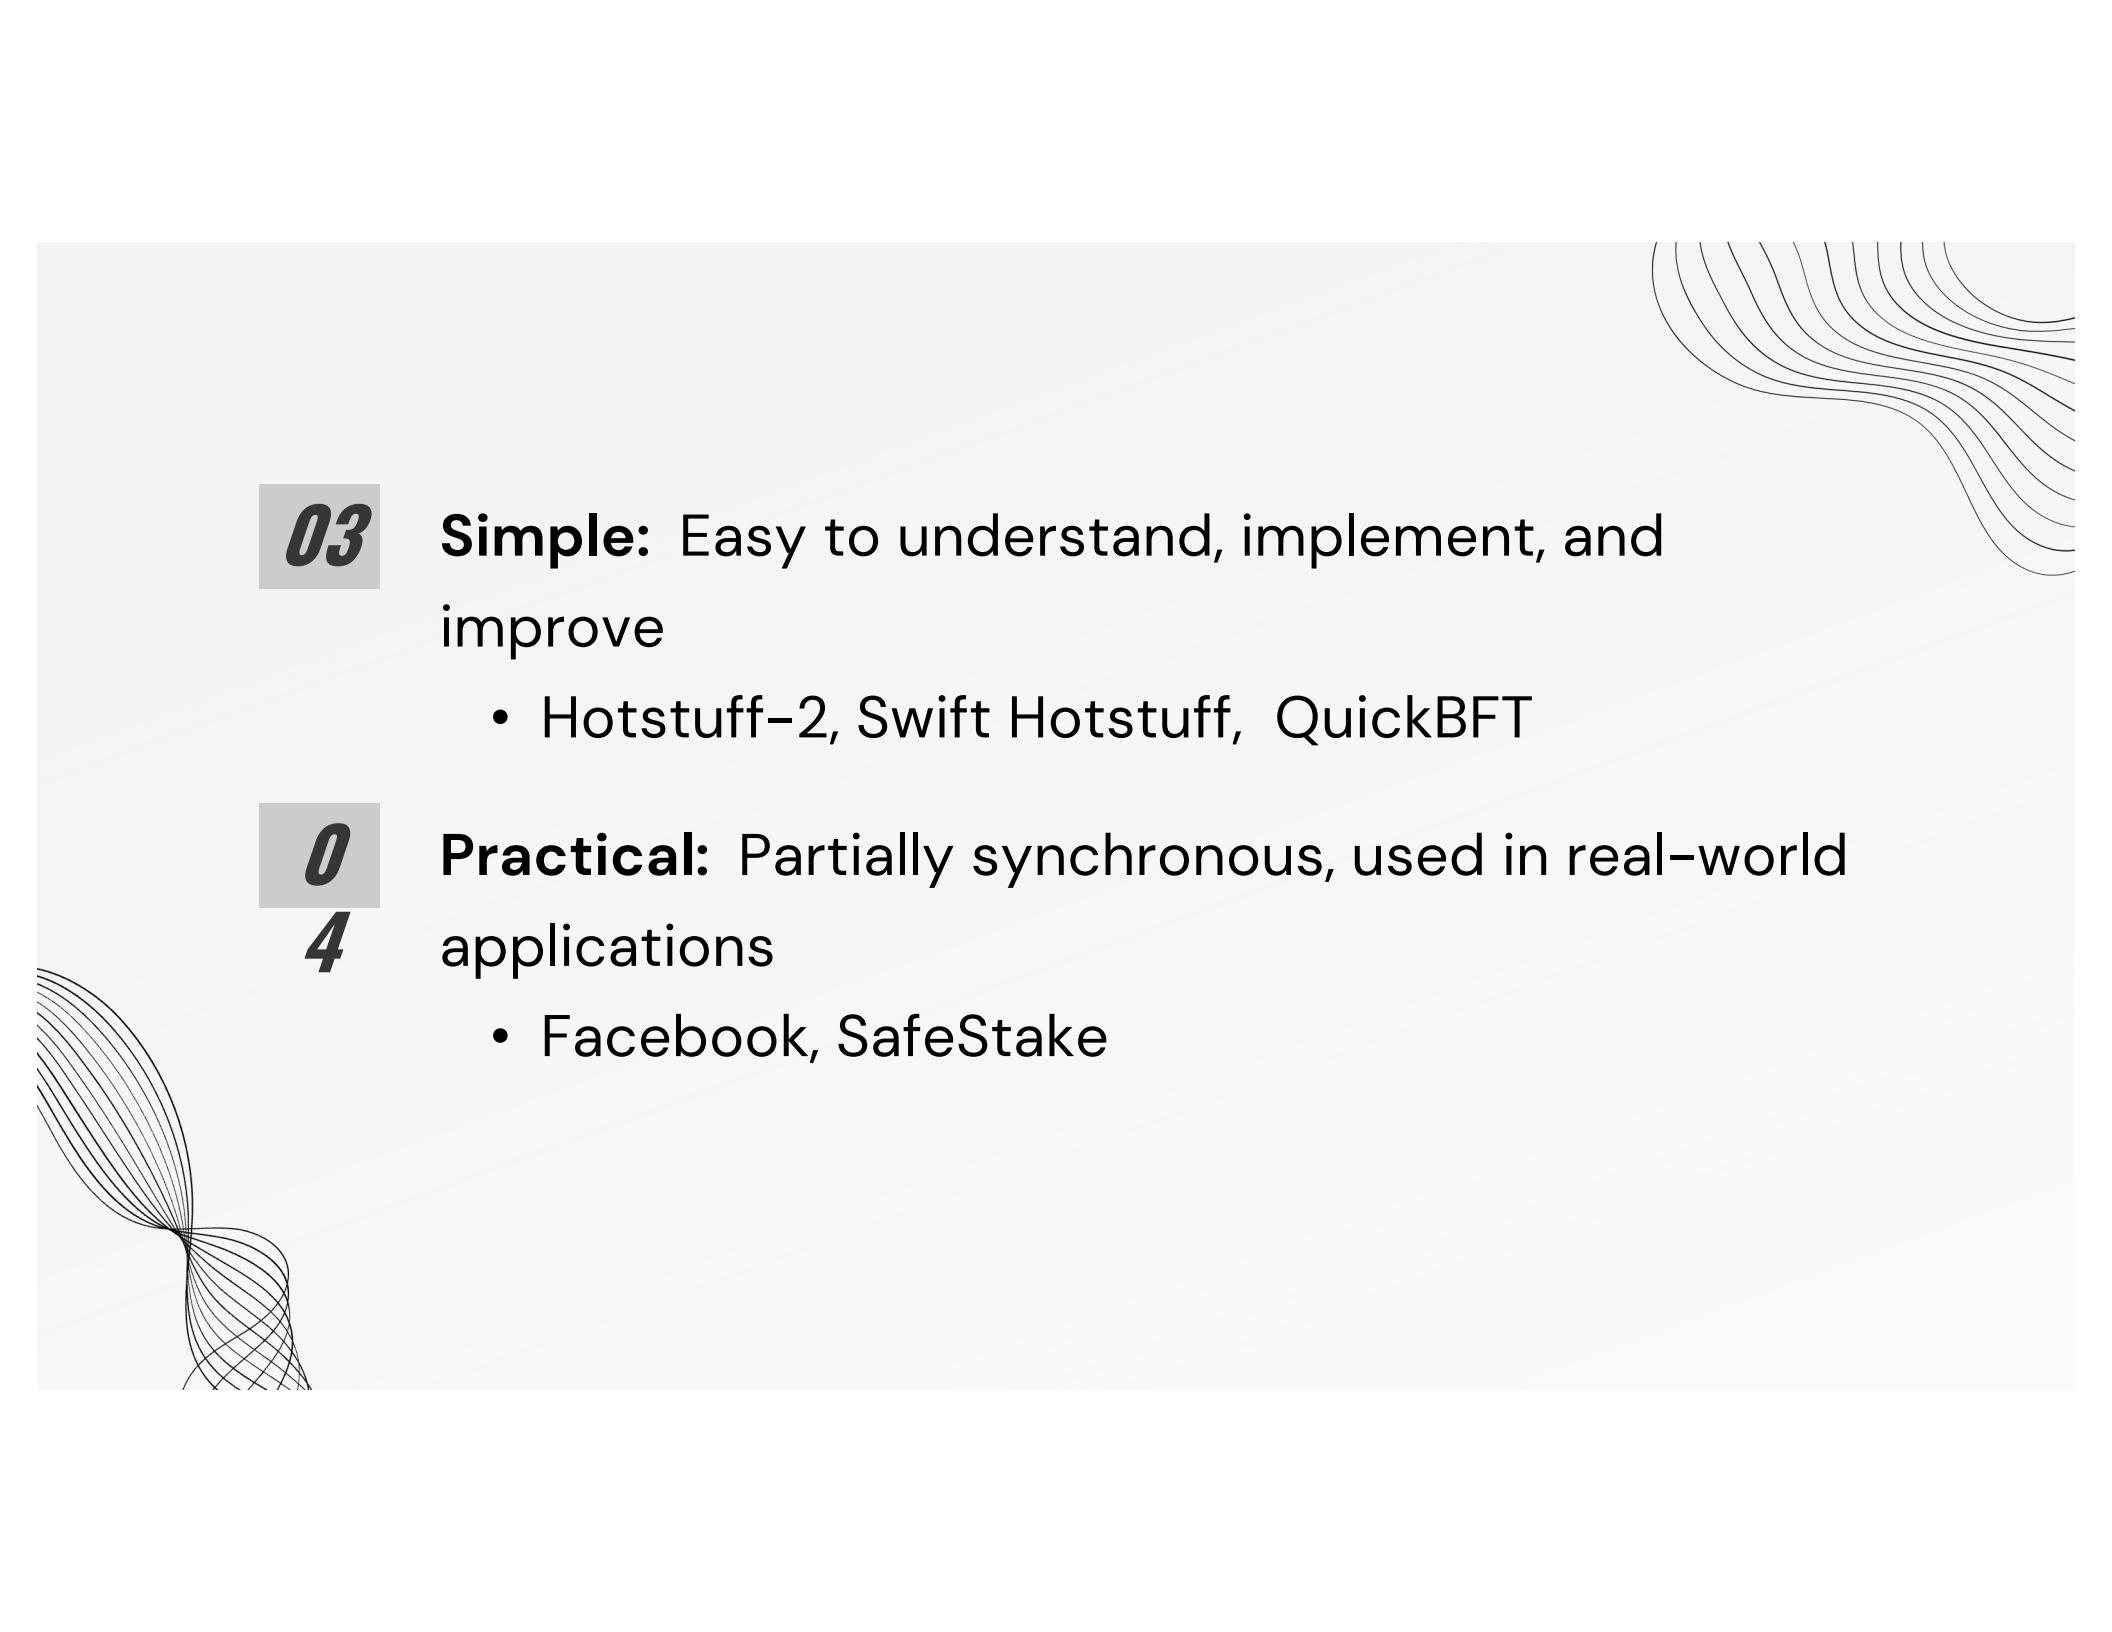


**02**

**Optimistic Responsiveness:** performs on actual network conditions (not worse-case)

- Receives majority votes, instantly commits





03

**Simple:** Easy to understand, implement, and improve

- Hotstuff-2, Swift Hotstuff, QuickBFT

0

4

**Practical:** Partially synchronous, used in real-world applications

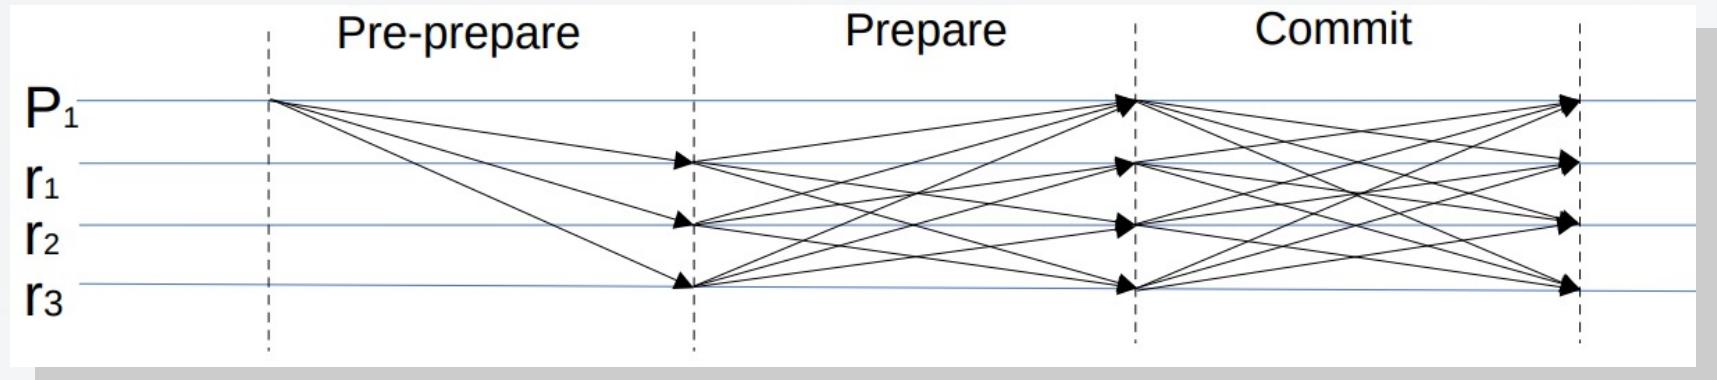
- Facebook, SafeStake



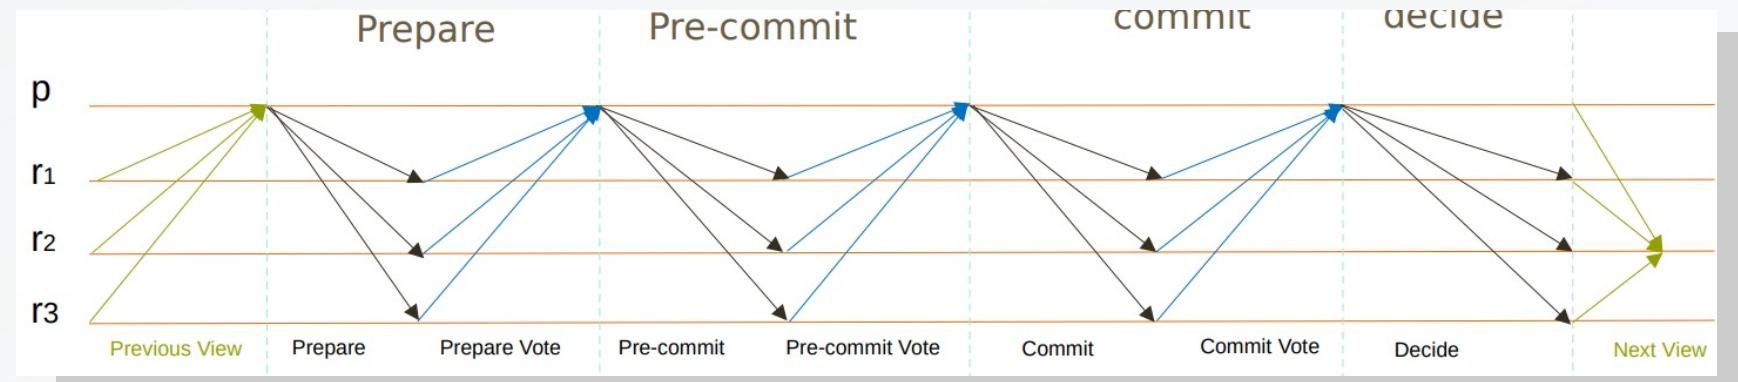
# PHASES

## HOTSTUFF PROTOCOL

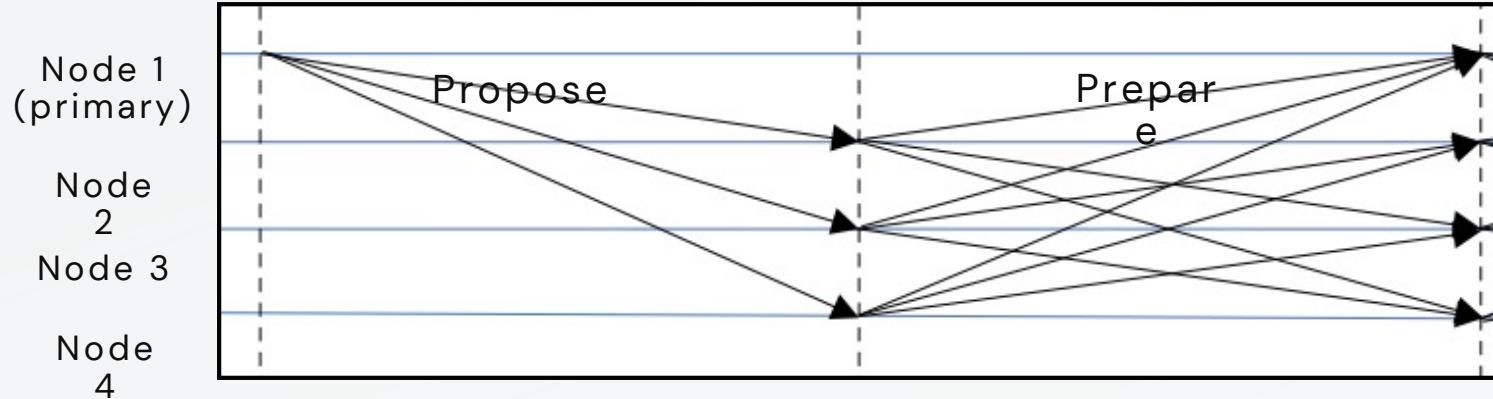
## PBFT



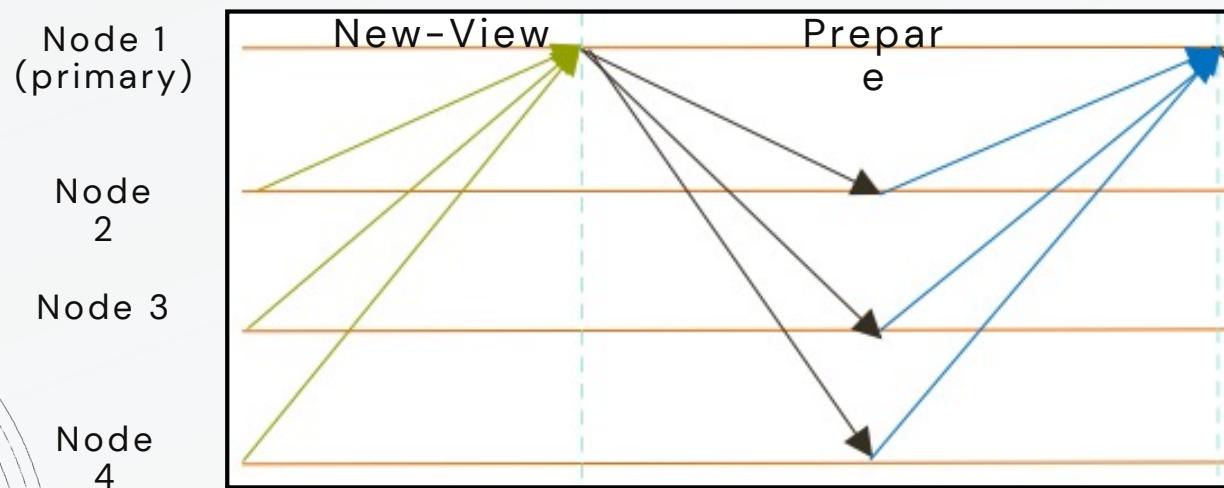
## Hotstuff



# PREPARE PHASE



PBF  
T



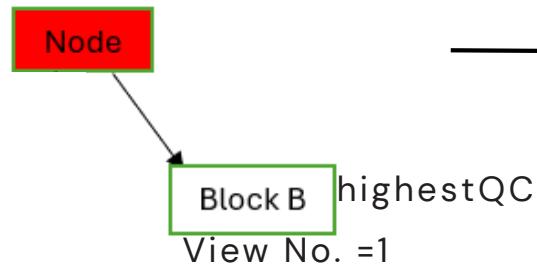
HOTSTUF  
F

## TREE and BRANCHES

- Commands in the protocol are wrapped in nodes that are part of a larger **tree** structure
- Each node contains parent link which is the hash digest of the parent node
- **Branches** extend from nodes and represent different possible paths

# Trees-the differentiating factor

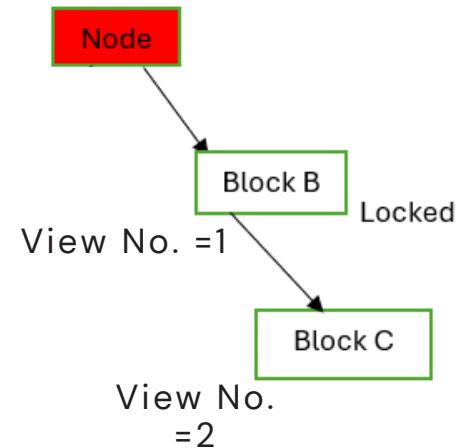
Local tree structure in all nodes after committing transaction B in last view.  
When B is committed it gets Block B is locked.

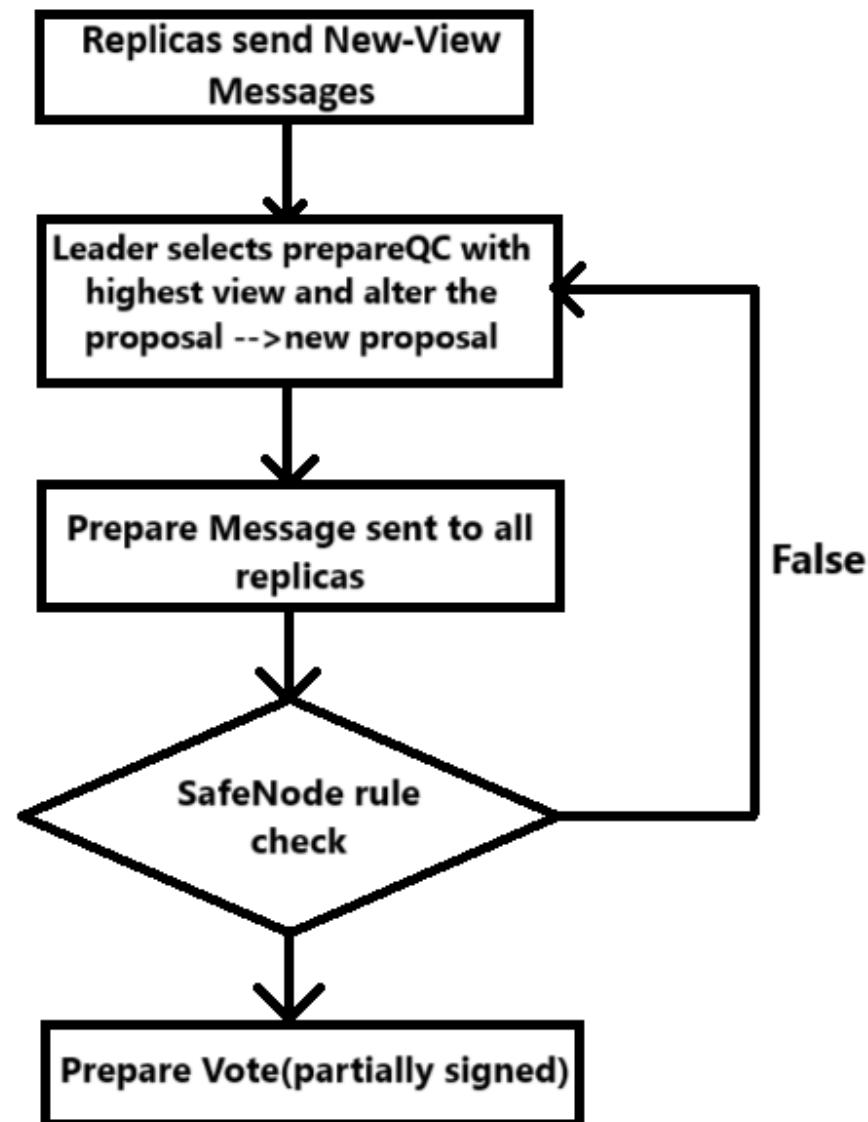


Now leader want to commit transaction C

For safety, it will build on block that had highest approval in the last view.  
Here its B.  
The leader will then add a new leaf node(block C) pointing to B.

Local tree in primary.



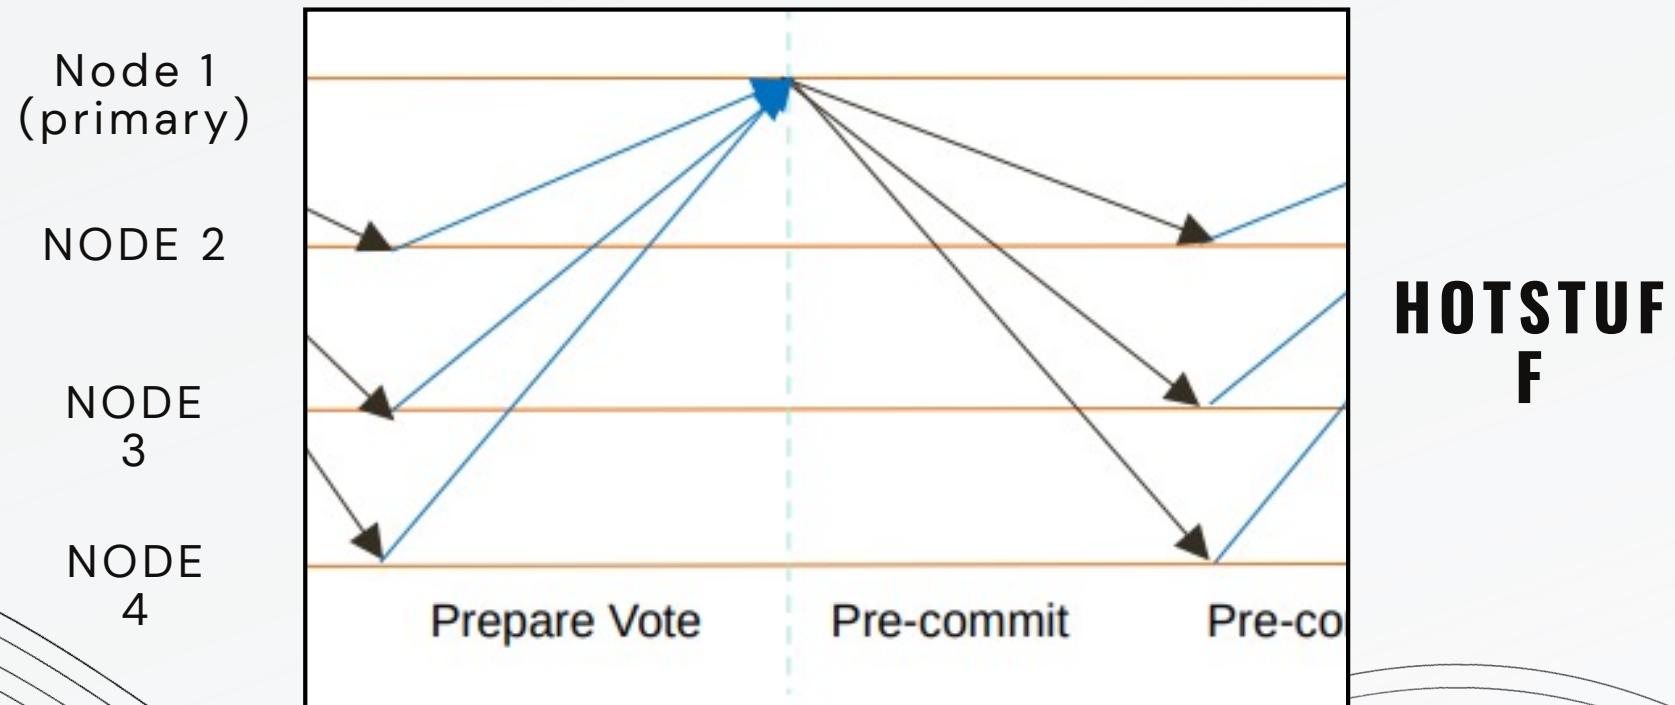


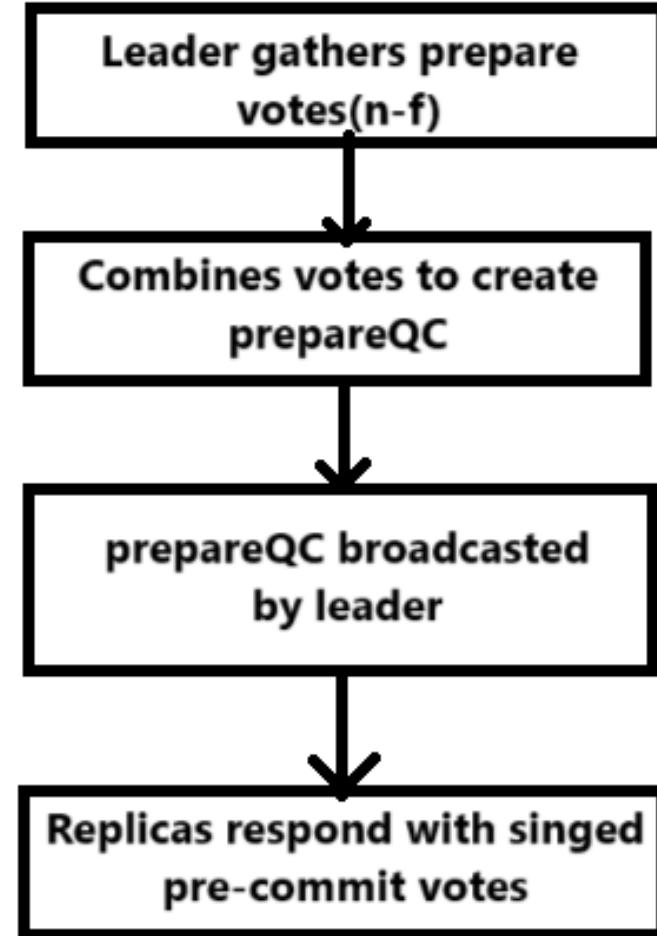
# SAFENODE

**Safety:** node extends from previous locked node

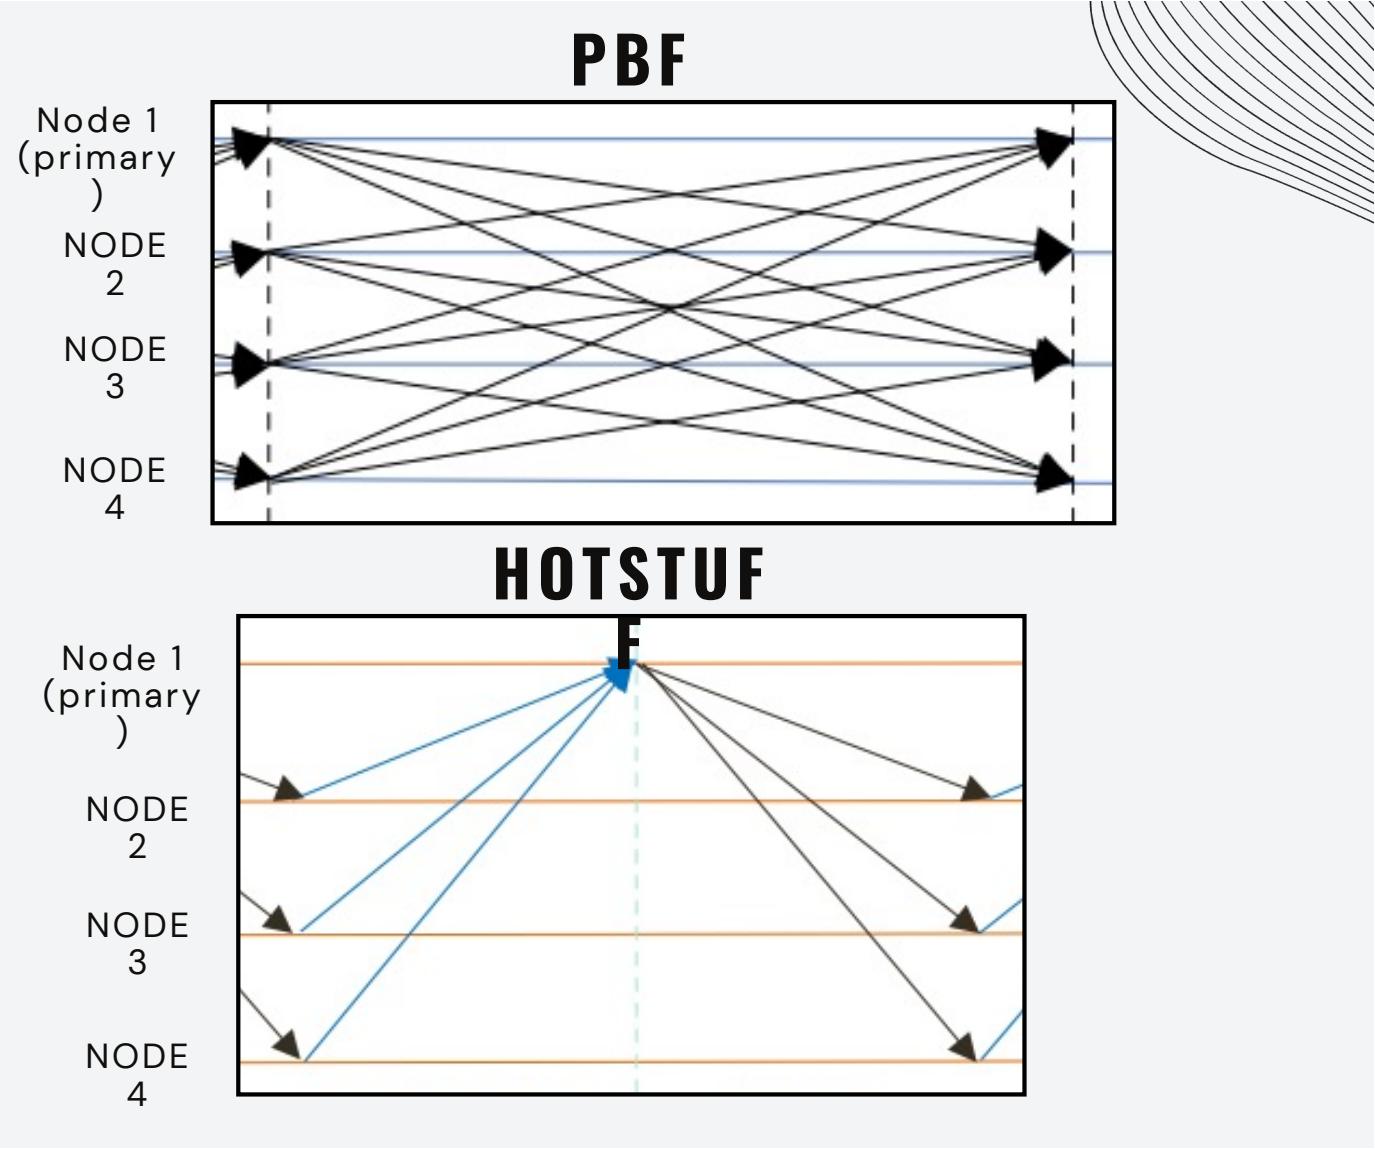
**Liveliness:** viewNumber higher than previous

# PRE-COMMIT PHASE

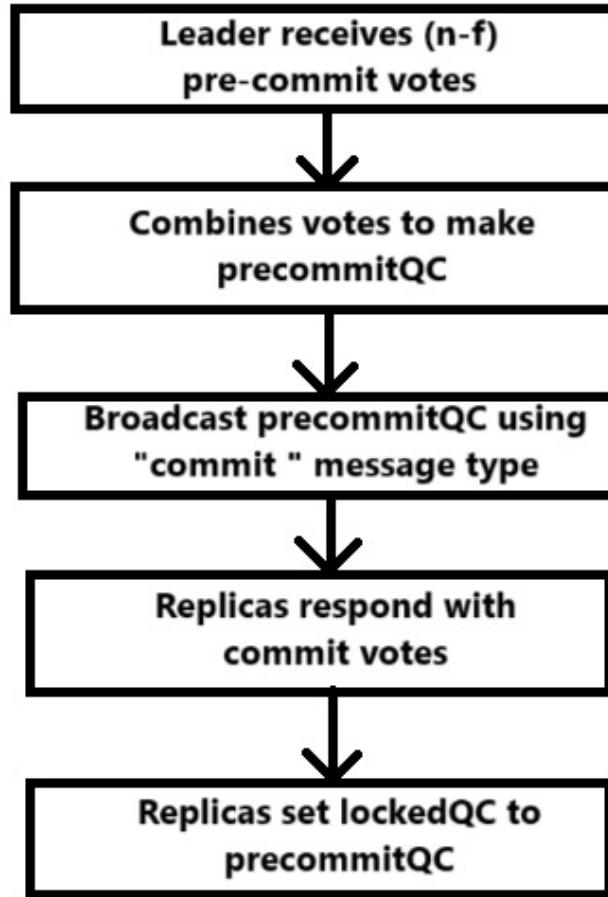




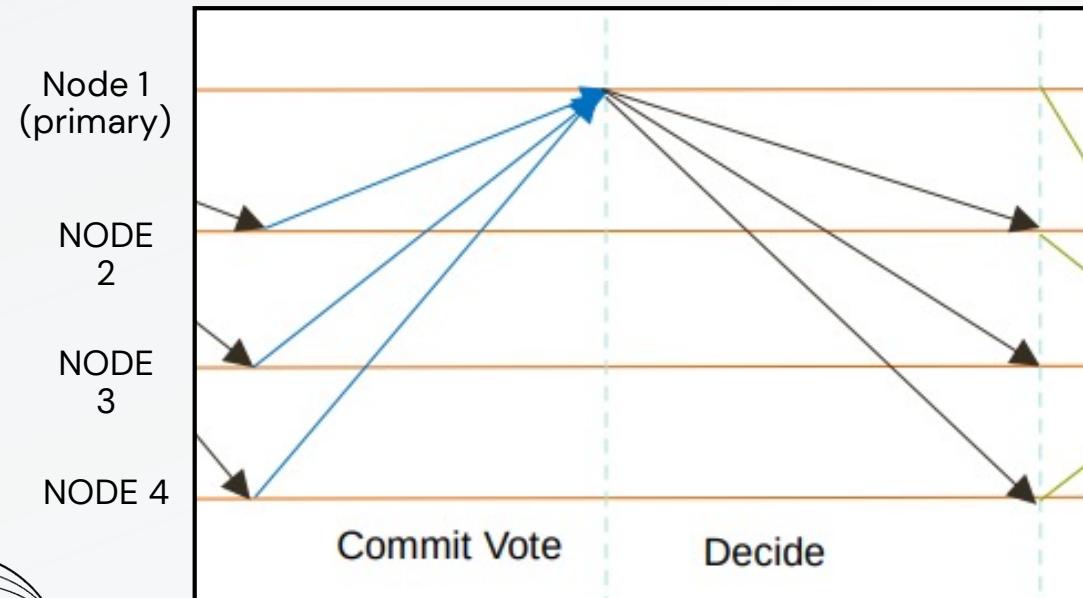
# COMMIT PHASE



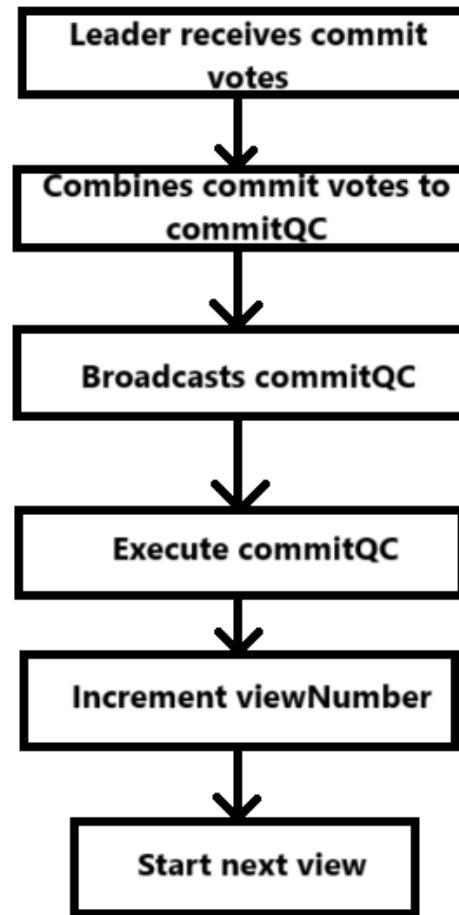
The last step provides safety incase the proposal become a consensus decision



# DECIDE PHASE



**HOTSTUF**  
F



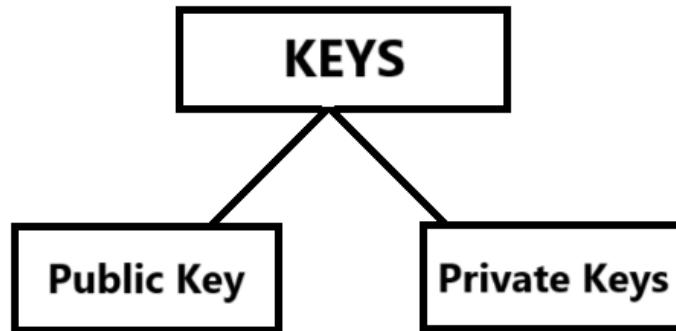
# NEXTVIEW INTERRUPT

- All phases, replica waits for a message at view  $viewNumber$  for a timeout period.
- If nextView interrupt, replica increments  $viewNumber$  and starts the next view.

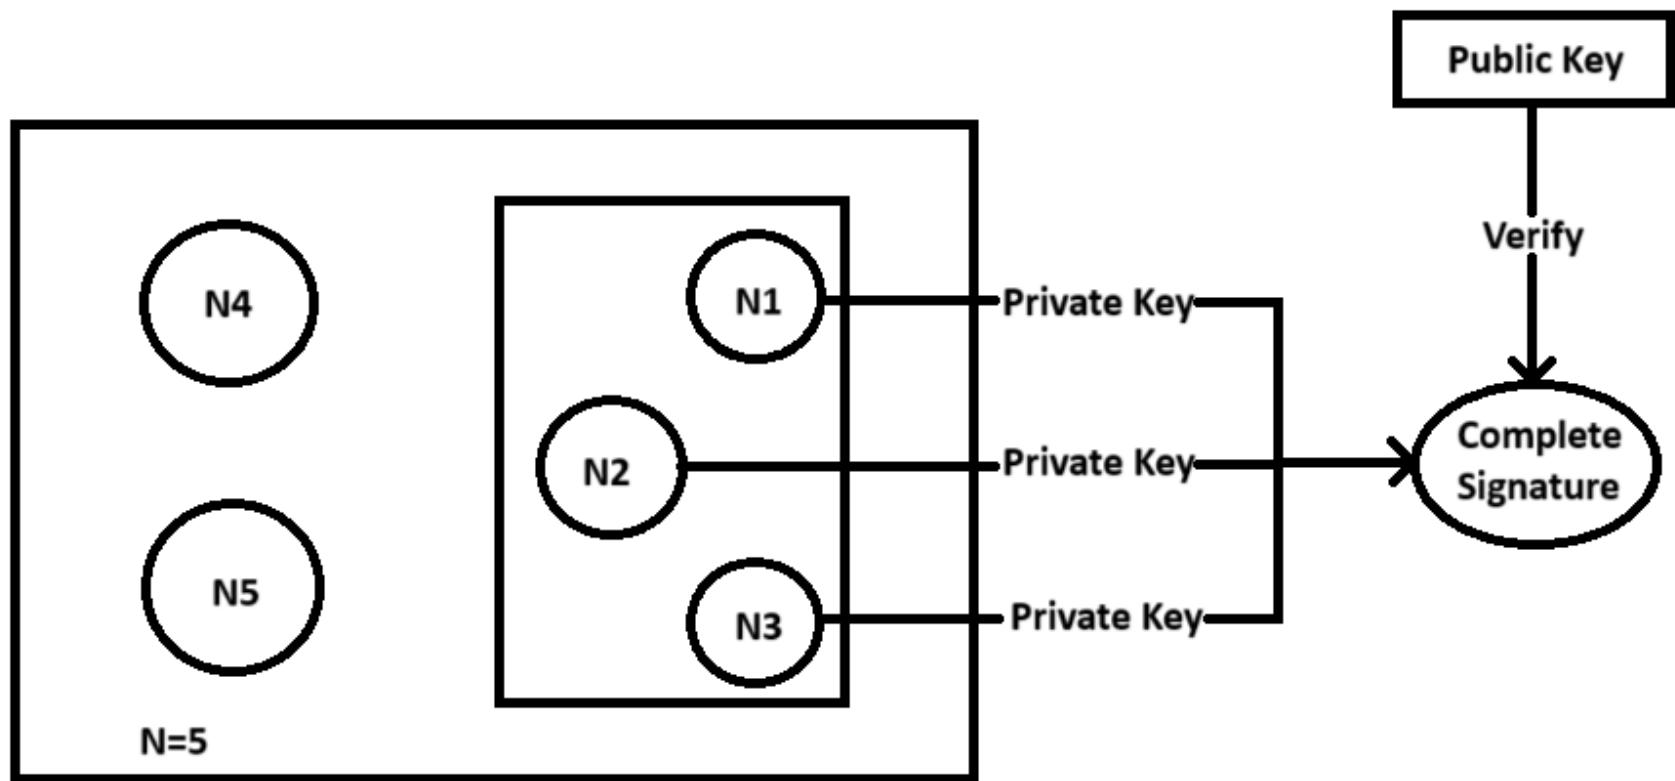
## CRYPTOGRAPHIC PRIMITIVES

Used for security against Byzantine faults caused by faulty nodes.

### (k,n)-Threshold Signature Scheme

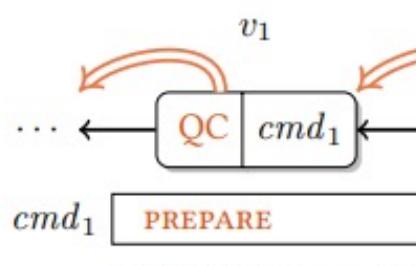


- **Public Key:** Single public key is shared by all replicas.
- **Private Keys:** Each replica has a private key which it will use to generate a partial signature.



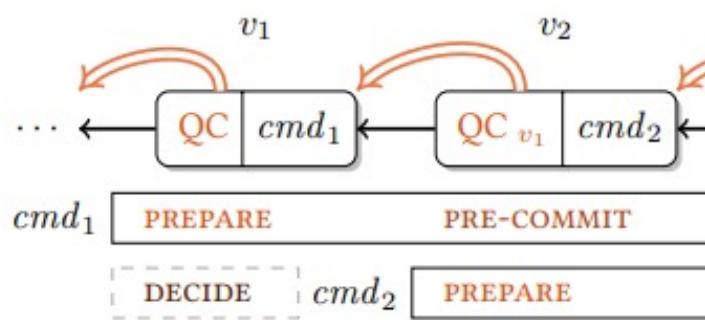
# CHAINED HOTSTUFF

- Pipelines process, reduces number of messages
- genericQC: Send to next leader to delegate responsibility



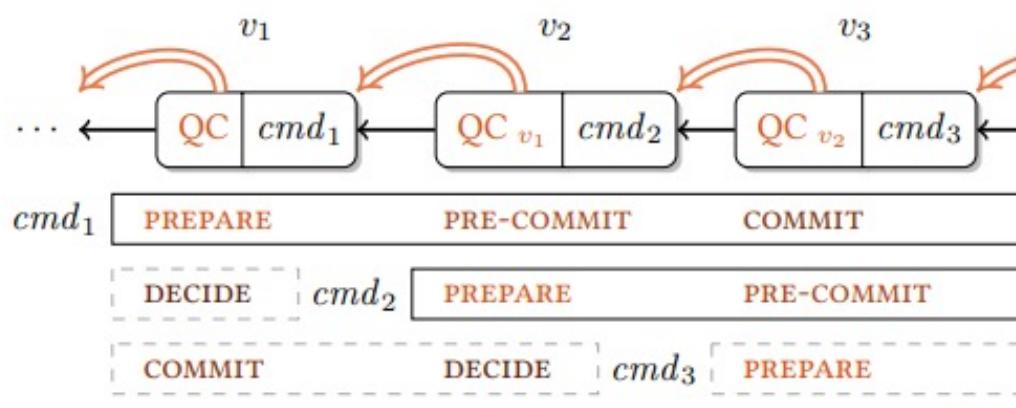
# CHAINED HOTSTUFF

- Pipelines process, reduces number of messages
- genericQC: Send to next leader to delegate responsibility



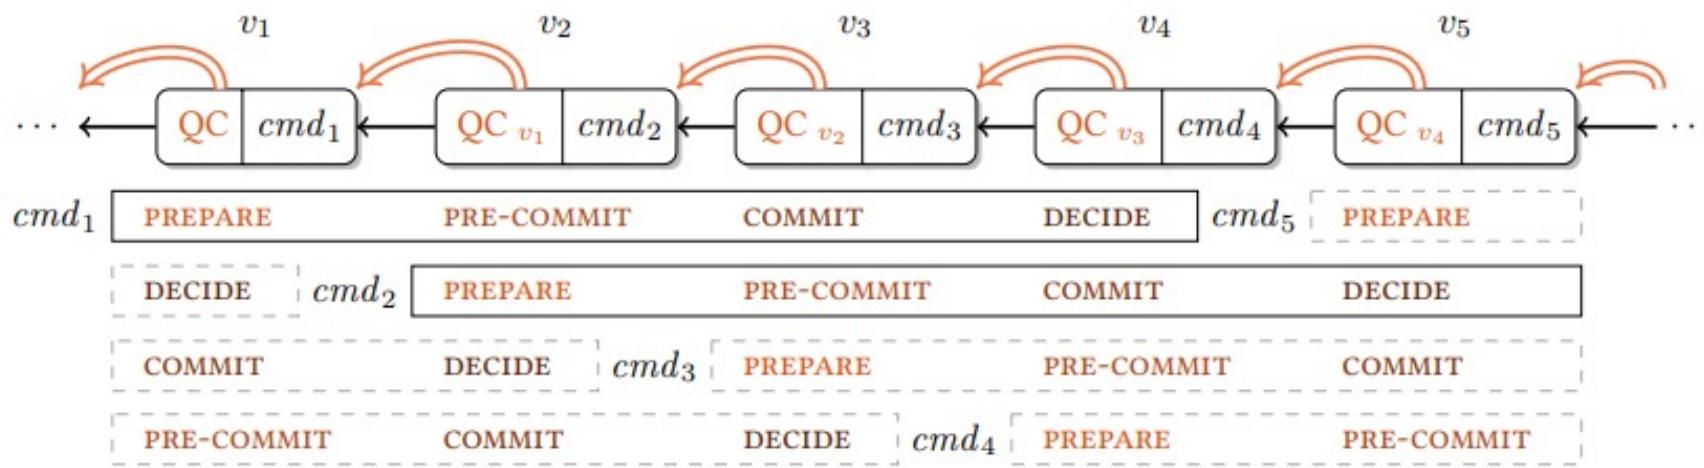
# CHAINED HOTSTUFF

- Pipelines process, reduces number of messages
- genericQC: Send to next leader to delegate responsibility



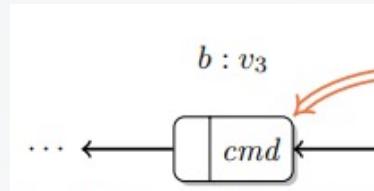
# CHAINED HOTSTUFF

- Pipelines process, reduces number of messages
- genericQC: Send to next leader to delegate responsibility



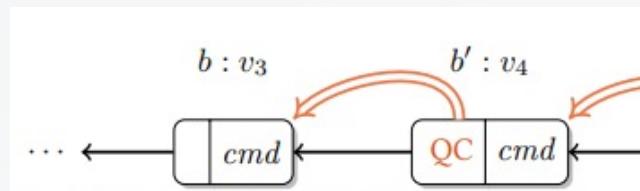
# NUMBERED CHAINS

- One-Chain: PREPARE b succeed



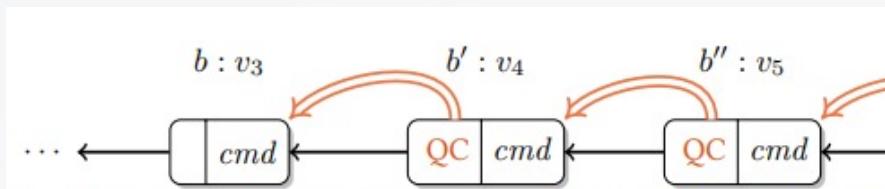
# NUMBERED CHAINS

- One-Chain: PREPARE b' succeed
- Two-Chain: PRE-COMMIT b succeed



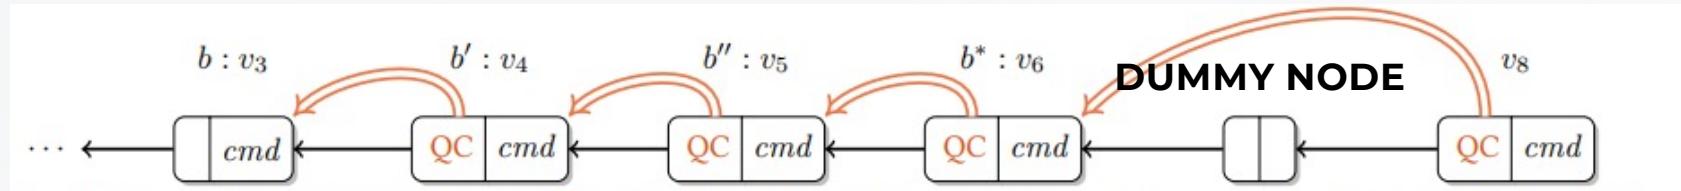
# NUMBERED CHAINS

- One-Chain: PREPARE b" succeed
- Two-Chain: PRE-COMMIT b' succeed
- Three-Chain: COMMIT b succeed



# NUMBERED CHAINS

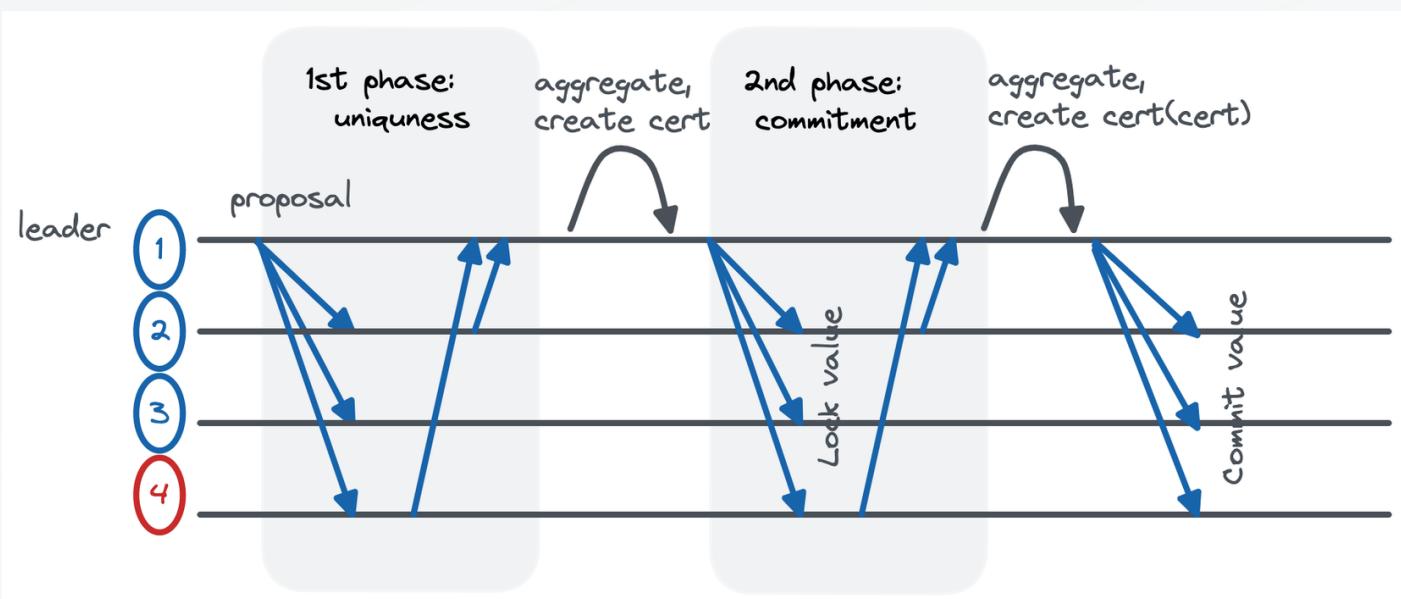
- One/Two/Three - Chain
- If at any point fail due to fault, chain breaks and calls nextView
  - Create dummy node in tree to maintain height



# HOTSTUFF-2

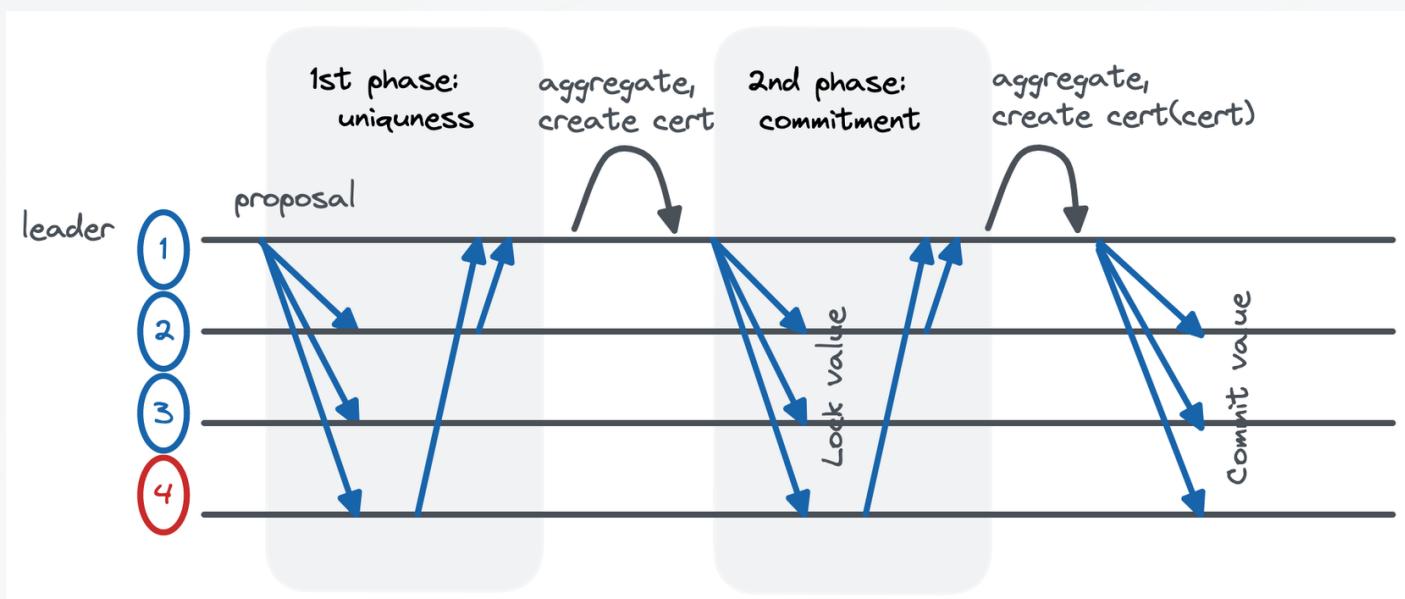
Two phases while maintaining all necessary properties

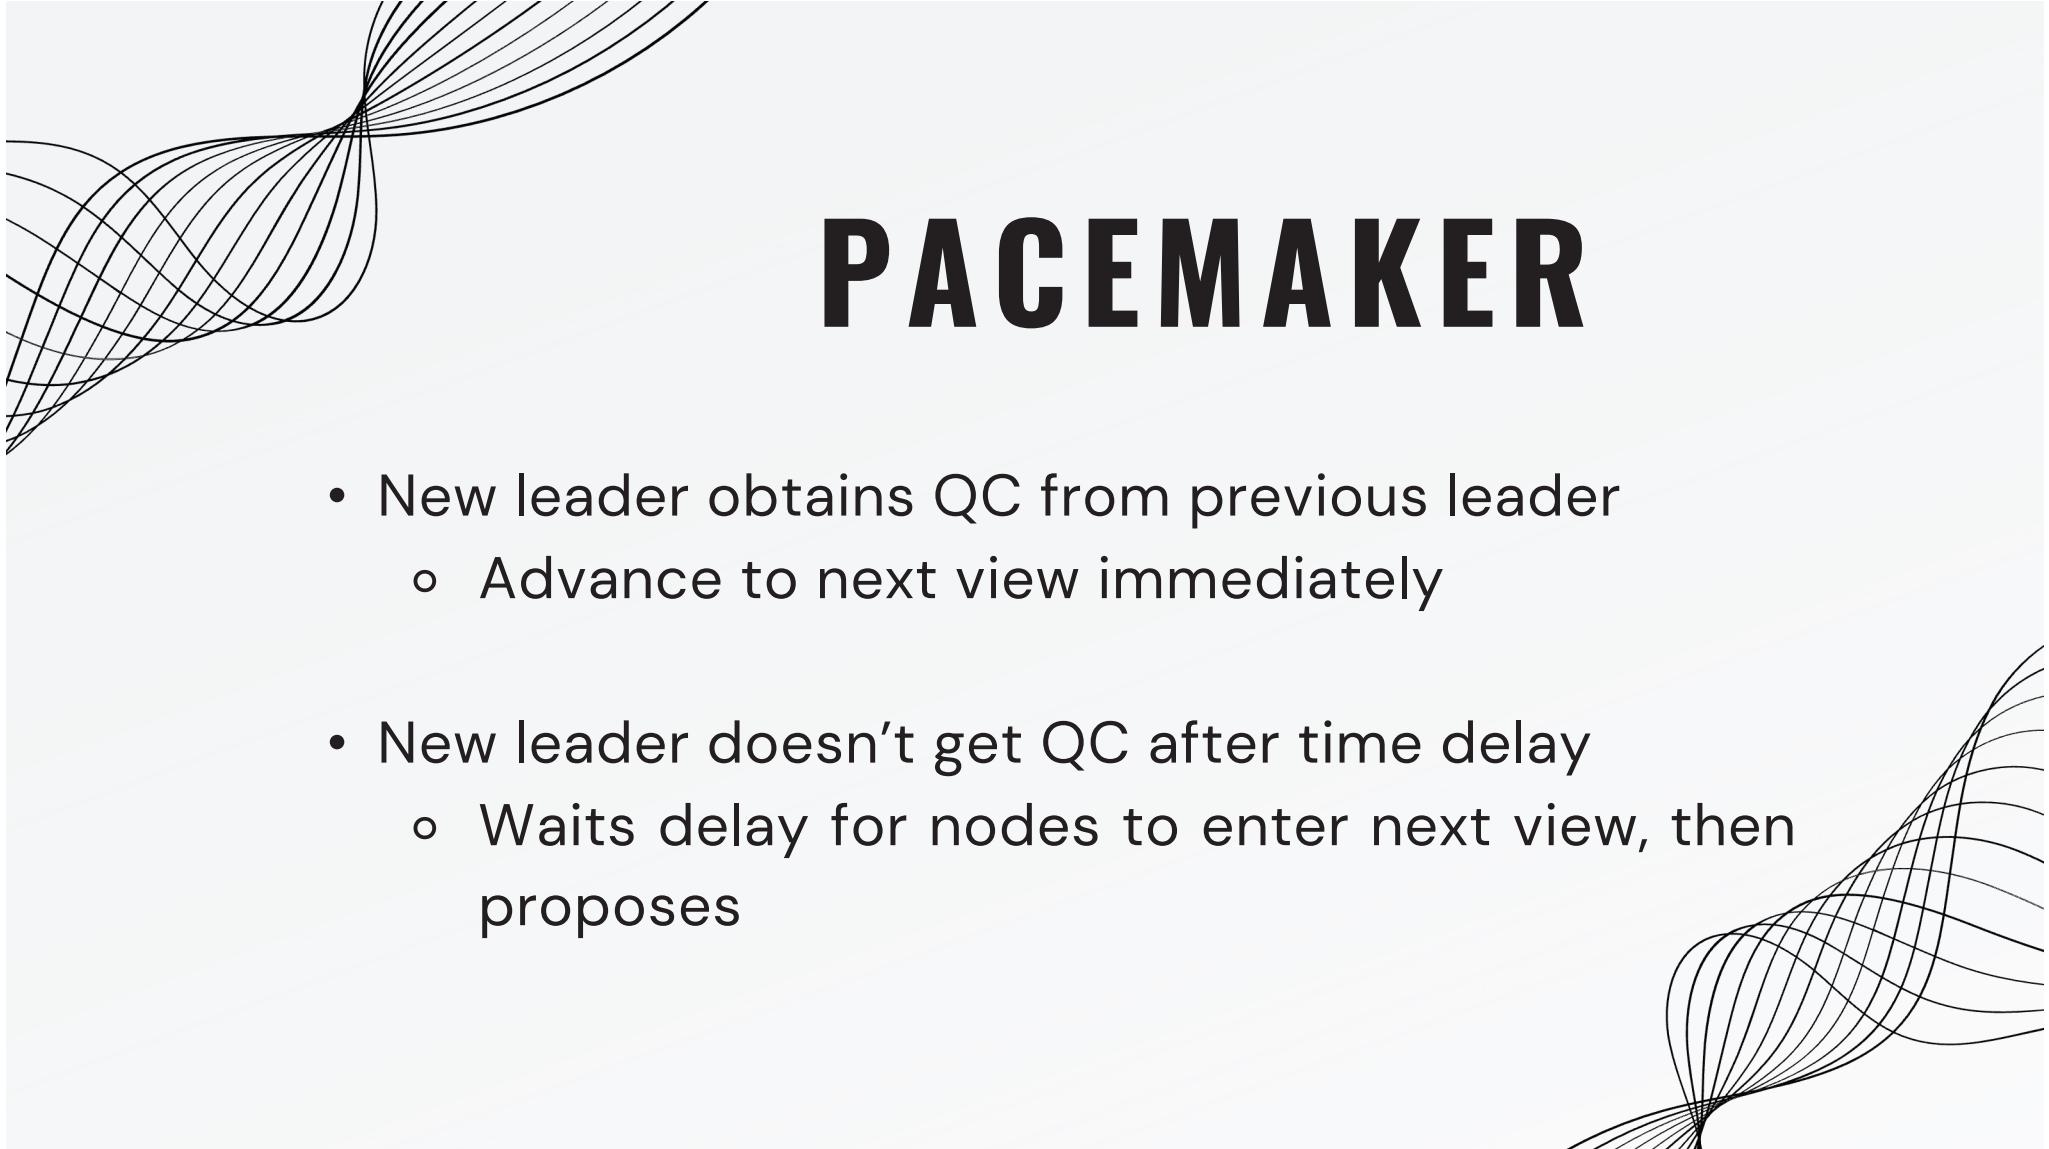
- No PRE-COMMIT phase, locked on prepareQC



# HOTSTUFF-2

Locked nodes will not vote on a conflicting value unless proof is shown that it is safe





# PACEMAKER

- New leader obtains QC from previous leader
  - Advance to next view immediately
- New leader doesn't get QC after time delay
  - Waits delay for nodes to enter next view, then proposes

# EVALUATION

HotStuff (2 and 3-phase)  
versus BFT-SMaRT

# LATENCY/THROUGHPUT

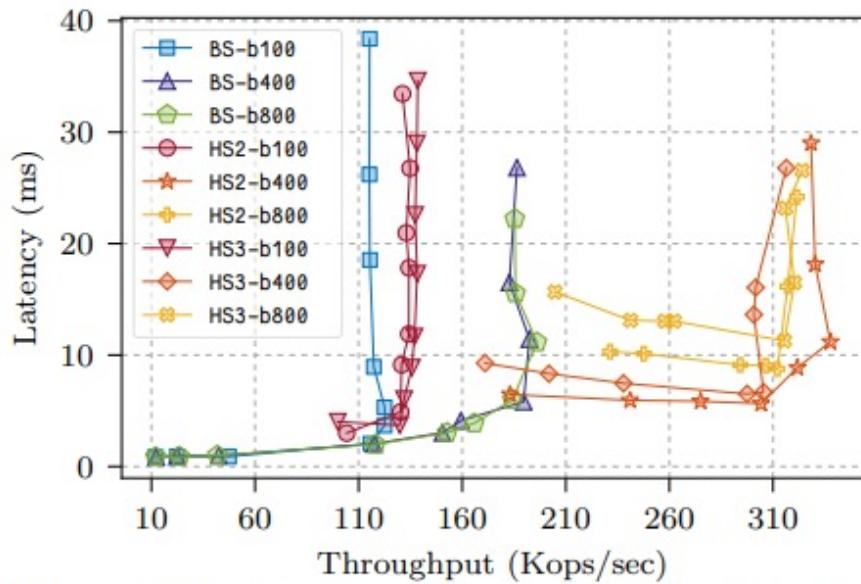


Figure 4: Throughput vs. latency with different choices of batch size, 4 replicas, 0/0 payload.

# LATENCY/THROUGHPUT

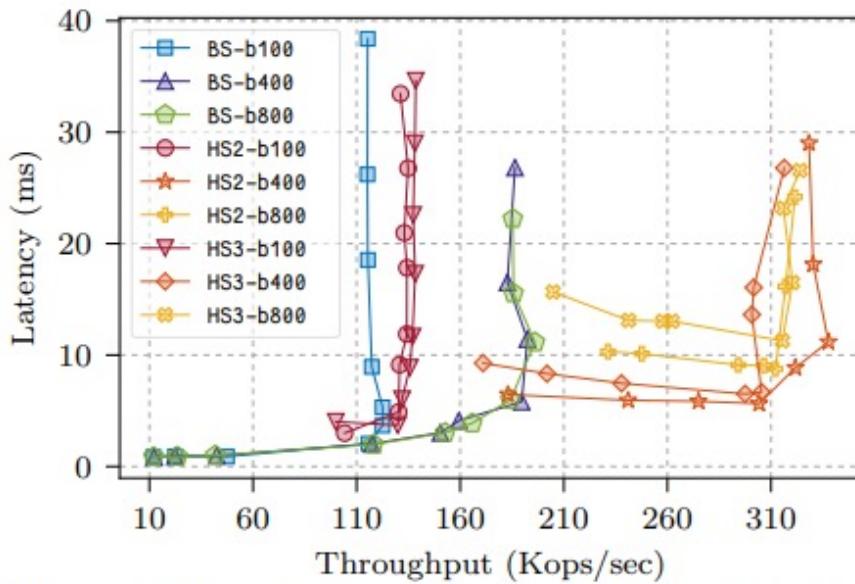


Figure 4: Throughput vs. latency with different choices of batch size, 4 replicas, 0/0 payload.

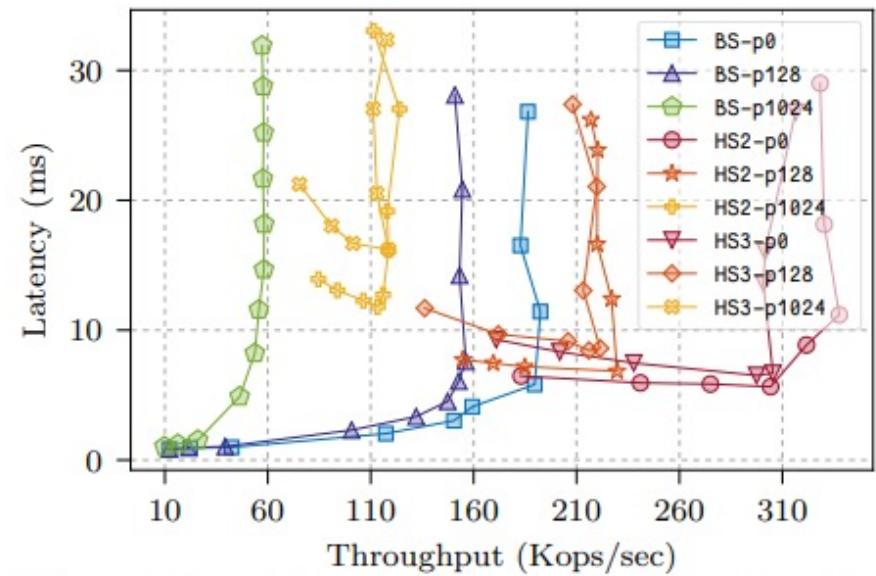
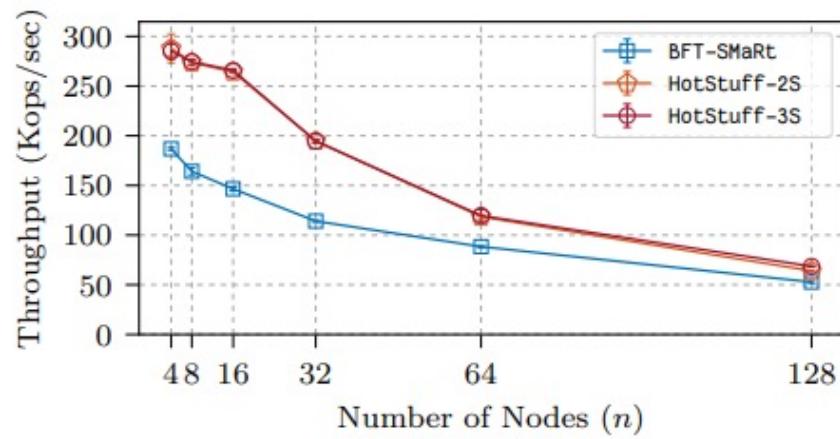
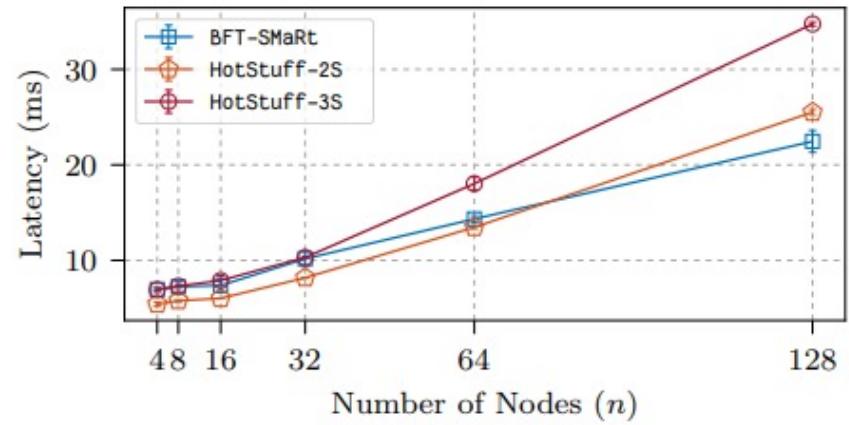


Figure 5: Throughput vs. latency with different choices of payload size, 4 replicas, batch size of 400.

# SCALABILITY



(a) Throughput



(b) Latency

Figure 6: Scalability with 0/0 payload, batch size of 400.

# SCALABILITY

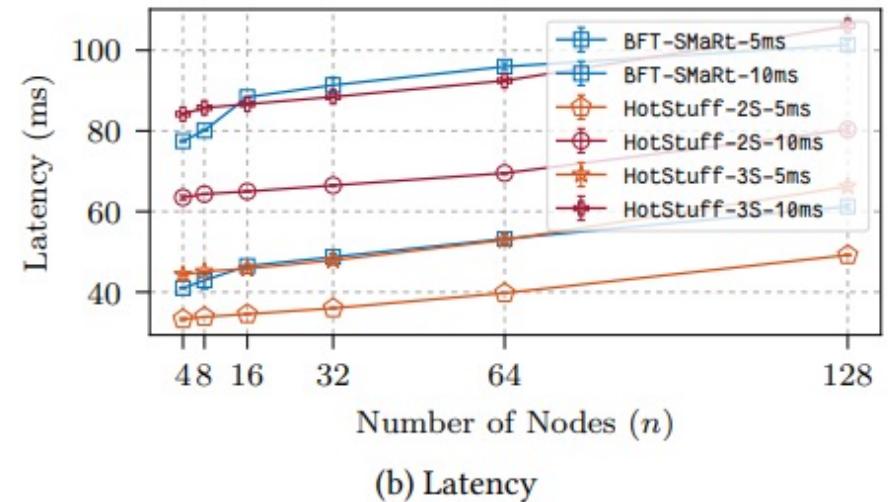
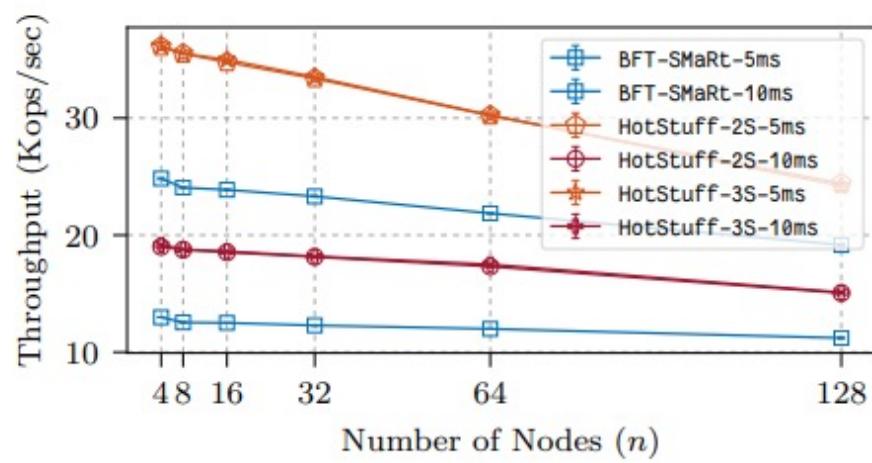
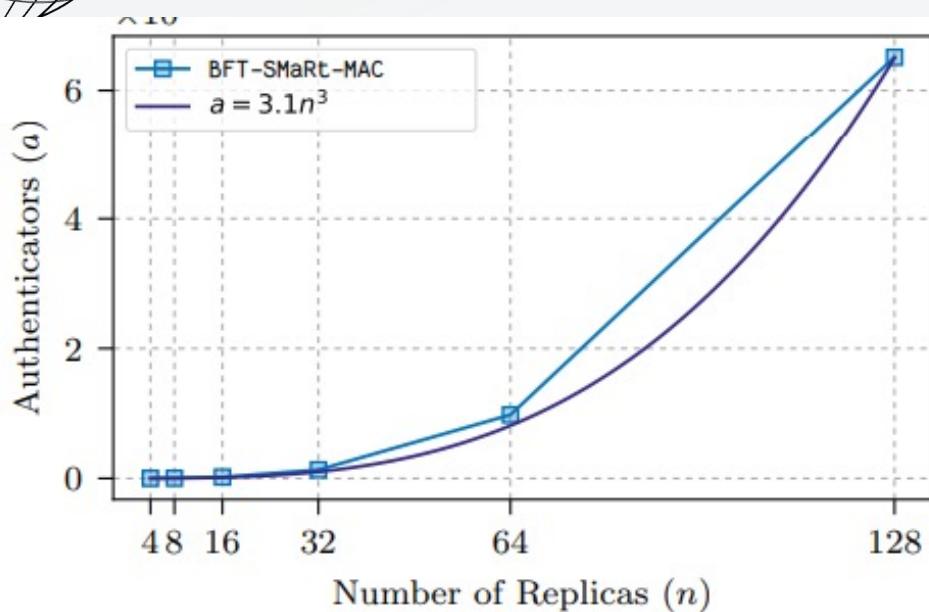
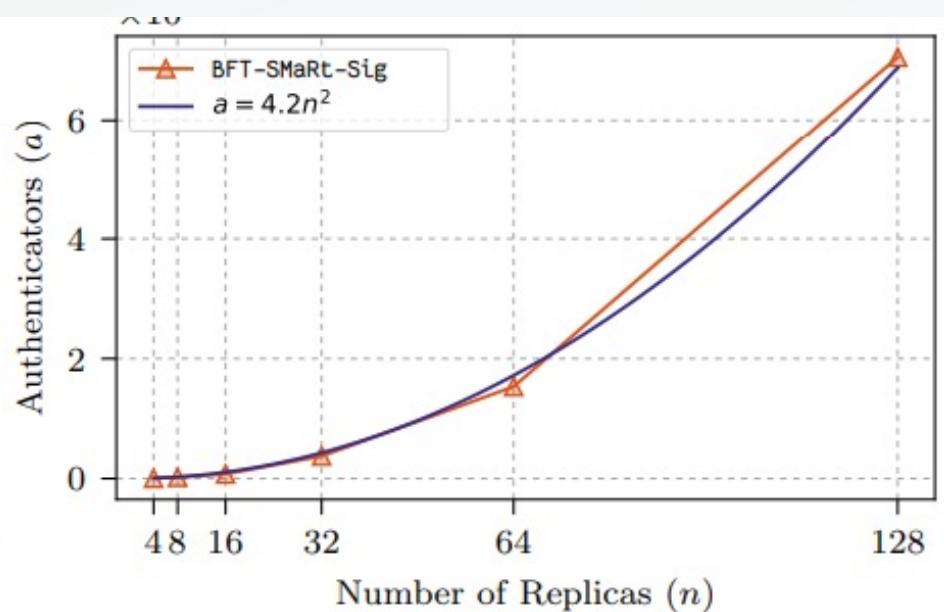


Figure 8: Scalability for inter-replica latency  $5\text{ms} \pm 0.5\text{ms}$  or  $10\text{ms} \pm 1.0\text{ms}$ , with 0/0 payload, batch size of 400.

# EXTRA AUTHENTICATORS



(a) MACs



(b) Signatures

Figure 10: Number of extra authenticators used for each BFT-SMaRt view change.

# RESOURCES

**Main Paper:** Yin, Maofan, et al. "HotStuff: BFT consensus in the lens of blockchain." arXiv preprint arXiv:1803.05069 (2018).

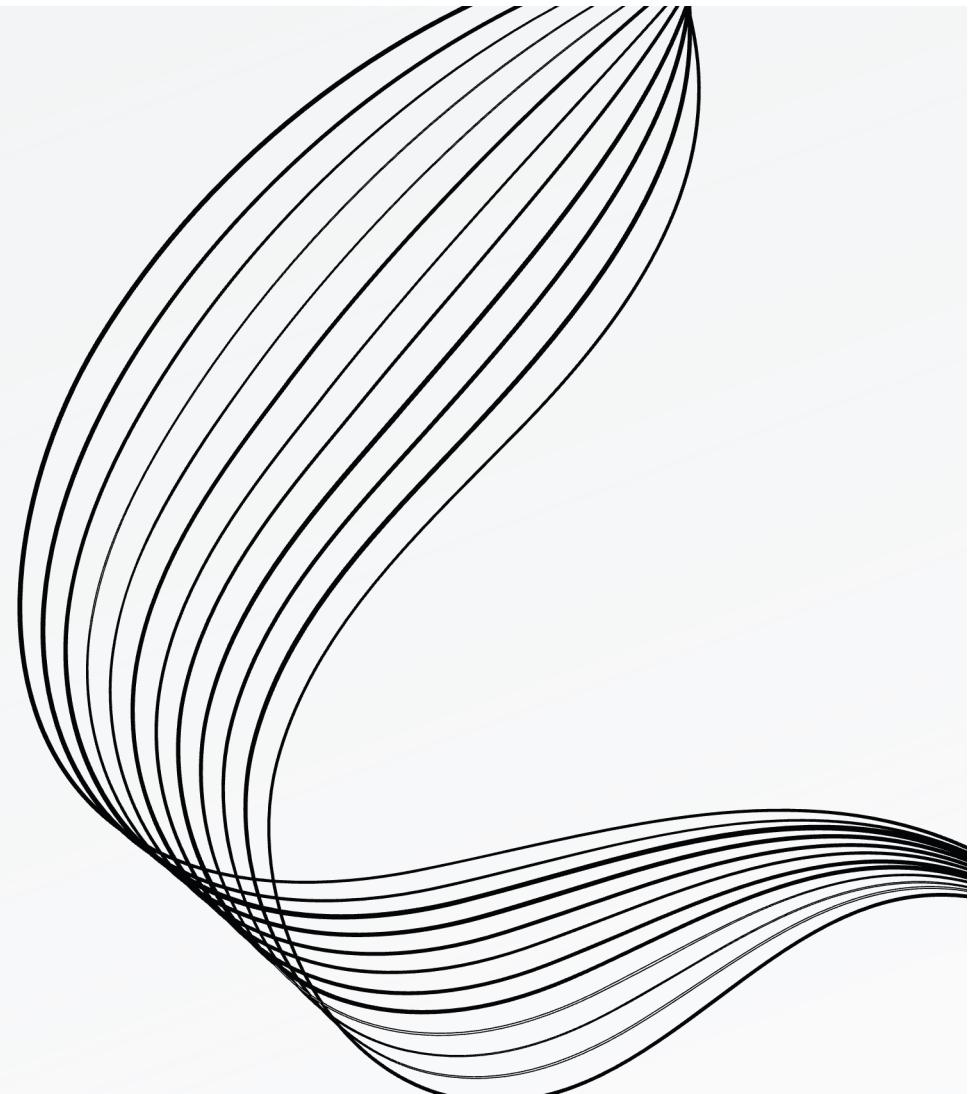
## **Additional Resources:**

Malkhi, Dahlia and Yin, Maofan. "Lessons from HotStuff." Proceedings of the 5th workshop on Advanced tools, programming languages, and PLatforms for Implementing and Evaluating algorithms for Distributed systems. 2023.

Malkhi, Dahlia and Nayak, Kartik. "HotStuff-2: Optimal Two-Phase Responsive BFT." Cryptology ePrint Archive, 2023.

Zhao, Siyuan, Wu, Yanqi, and Wang, Zheng. "HotStuff-2 vs. HotStuff: The Difference and Advantage." arXiv preprint arXiv:2403.18300 (2024).

# QUESTIONS?



# Synchronous

Known time bound on  
messages

- Less realistic
- Doesn't factor in delay and reordering
- Assume messages are all received within time bound

# Asynchronous

Unknown time bound on  
messages

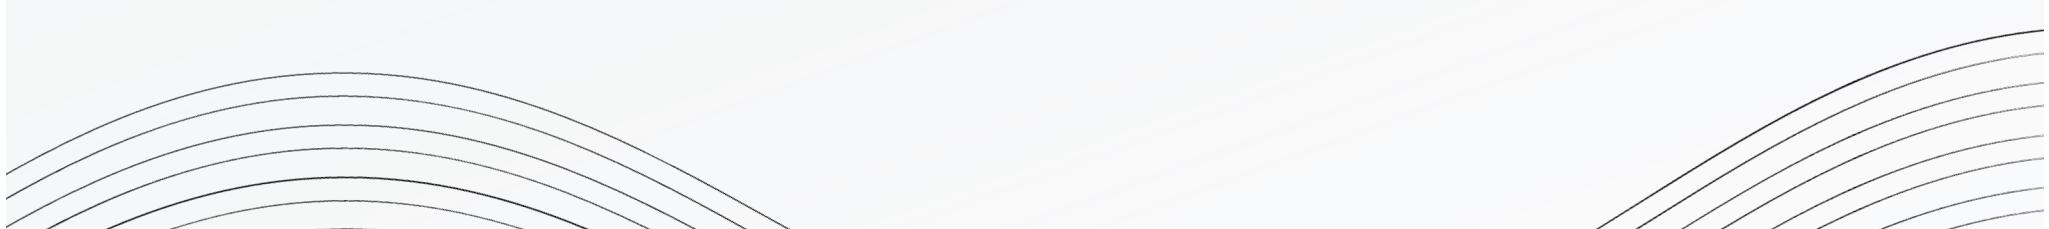
- More realistic
- Factors in delay and reordering
- Messages delayed by unknown time period
- Complicates correctness and liveness

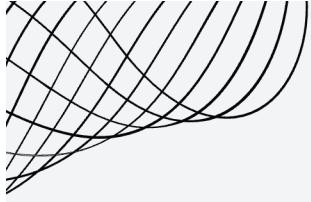
# Partial Synchrony

- Middle ground between asynchrony & synchrony
- Asynchronous before GST,  
synchronous after GST
- Unknown time bounds on  
communication + processing

# **GST (Global Stabilization Time)**

- Time bounds do not hold at all times
- Holds up to limited time after GST  
(assume infinite time for simplicity)
- behavior can only be stable and predictable after GST





# Precommit Phase

- Replica locks onto the block on separate phase before commit
  - Compare own QC with leader QC
  - Transition to new leader
    - Picks up from where old leader left off
    - Highest QC
- 