

# Lumiere: Making Optimal BFT for Partial Synchrony

## Practical

ANDREW LEWIS-PYE, DAHLIA MALKHI, ODED NAOR, KARTIK NAYAK

Henry Chou

Varun Ringnekar

Chris Ruan

Shaokang Xie

Yubo Bai

# Background

Introduction to core concepts



# Partially Synchronous BFT model

The Lumiere protocol operates on a Partially Synchronous Byzantine Fault Tolerant model

$3f + 1$  replicas

$f$  faults

$f_a$  Actual Faults

GST (Global Stabilisation Time)

$\Delta$  Known network delay

A message sent at time  $t$  must arrive by time  $\max\{\text{GST}, t\} + \Delta$ .

# Synchronized Clocks

All the nodes in the system have a local clock  $lc$


Lumiere tries to keep these clocks synchronized so that at least  $f+1$  clocks are synchronized closely together

All the nodes know which view will have which leader  $lead(v)$


All the nodes in the system know when a view should reach consensus, this time is also called the view timer  $C_v$

All the nodes also know when the next view is supposed to start  $\Gamma_v$

# Epochs



# Different Types of Views



# View Synchronisation

Making sure that all honest processors in a BFT system move between views in a coordinated and consistent manner, even in the presence of byzantine failures

**Epoch Certificate(EC):** Can be formed by any honest replica, when it receives  **$2f+1$**  **epoch view messages** from other replicas.  $O(n^2)$  communication. Formed to enter a new epoch, ensured heavy synchronization.

**Quorum Certificate(QC) :** Is formed by the leader when consensus is reached in a view i.e.  **$2f+1$**  replicas vote on a block. Created in initial views to enter non initial views

**View Certificate(VC) :** Is formed by the leader of the next view when enough replicas ( **$f+1$** ) send a **view message** to that leader to enter its view. Created to enter initial views

# Overview

Quick preview of the upcoming content





# Liveness

- **RECAP:** guarantee to make progress, even though faulty node exists
  - Timeout - change view with faulty leader/replica (Fever)
  - Epoch-based Design (LP22)

# Optimal Responsiveness

- **RECAP:** make progress at the speed of the network
  - Quorum Certificates (HotStuff) and Clock-bumping (Fever)
  - Epoch-based Design: Light & Heavy Synchronization (LP22)

# Fever protocol: Key characteristics

- Makes use of **optimistic responsiveness** - scales with actual network conditions rather than worst case assumptions, clocks of processors can be **bumped forward**
- View synchronization protocol that operates under a **partial synchrony model**
  - **Partial clock synchronization** after GST - clocks may drift before GST, but processor clocks are synchronized after network stabilization
- Designed to handle **frequent view changes** efficiently
- **Initial and non-initial views** - The leader for view  $v$  is  $(processor(v/2) \bmod n)$ , initial views have time based clock synchronization, optimistic responsiveness is achieved in non-initial views with QC
- **Fundamental idea of Fever - Honest Gap**

# Fever protocol: Honest Gap

- Measures the disparity in clock synchronization between the most advanced honest processors
- Definition: the  $i$ -th honest gap at time  $C$ , denoted as  $hg_{i,C}$ :  $hg_{i,C} := lc(p_1,C, C) - lc(p_i,C, C)$
- Gap needs to stay within bound to allow the processors to stay synchronized
- Despite processors bumping clocks forward, two properties are guaranteed:
  - $hg_{f+1,t} \leq \Gamma$  for all  $t \geq 0$ .
  - If  $hg_{f+1,t} \leq \Gamma$  at  $t \geq GST$  which is the first time an honest processor enters the view  $v$  with honest leader, then the leader will produce a QC.
- **Non-standard assumption**
  - Clocks assumed to be synchronized within  $\Gamma$  at the beginning for Fever protocol

# Drawback of PBFT

PBFT Properties that could be improved with Lumiere



# PBFT

- View-change - heavy communication exchange
  - All-to-All broadcast ( $O(n^2)$ ) for each view-change
  - When suffering from  $f$  consecutive faulty leader, at worst  $O((f+1)*n^2) \rightarrow O(n^3)$
- Not optimistically responsive
- Conclusion
  - practical for smaller, more controlled environments, but limited scalability

## Contribution of Lumiere:

1. Worst-case communication complexity  $O(n^2)$
2. Improve latency -  $O(n\Delta)$
3. The protocol is smoothly optimistically responsive
4. Eventual worst case communication complexity  $O(f_a n + n)$

## Overview of LP22

- Epoch-Based Structure
  - Time is divided into epochs, each consisting of  $f + 1$  views.
- Quadratic Worst-Case Communication Complexity:
  - Communication complexity of  $O(n^2)$  worst-case.

# Drawbacks of LP22

- Byzantine Leader Delays:
  - Even one Byzantine leader can cause a worst case delay, which is  $O(f + 1)\Gamma$
- High Communication Cost:
  - Heavy synchronization incurs quadratic cost ( $O(n^2)$ ), even without Byzantine faults.



# Lumiere Design

The Design of Lumiere



# Lumiere

## Key Contributions:

- **Latency Reduction:**


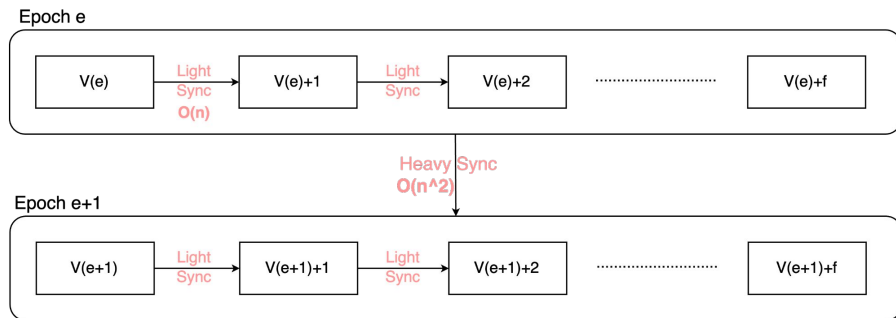
LP22's view sync process suffers from  $O((f + 1)\Gamma)$  latency due to Byzantine leaders. Lumiere reduces this by combining **LP22's epoch structure** with **Fever's clock bumping**.

- **Communication Optimization:**

Only a **constant-bounded number of heavy synchronizations ( $O(n^2)$ )** occur after GST, lowering communication overhead.


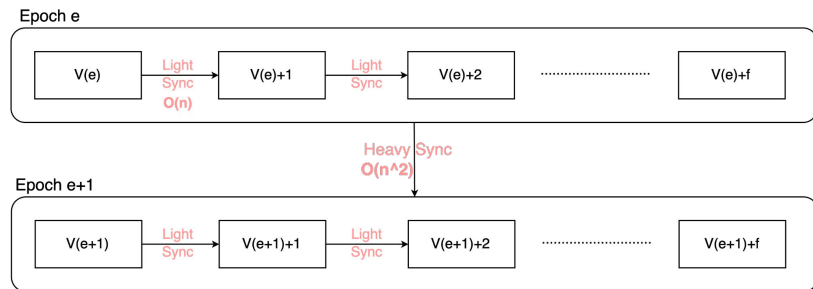
# Epoch Change in Lumiere

- **Heavy Sync** in Epoch change:  $(O(n^2))$ 
  - Phase 1: Send “**epoch e**” msg to all replicas
  - Phase 2.1: Any honest replica combine  $2f+1$  “**epoch e**” msg as **EC**(Epoch Certificate), then broadcast EC
  - Phase 2.2: When sb see a **EC**, enter **epoch e**
- **Clock Synchronization**: If a correct replica  $i$  enters epoch  $e$  at time  $t$ , then all correct replicas will have entered an epoch  $e' \geq e$  by time  $\max\{t, GST\} + \Delta$ .
  - Each replica will enter epoch  $e$  in a time interval  $\Delta$ .




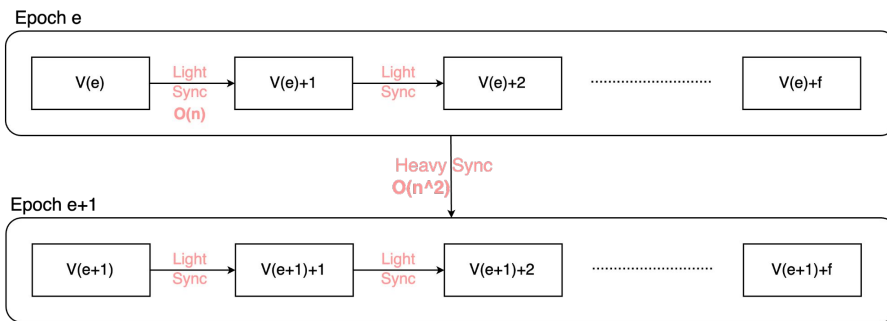
# Epoch Change in Lumiere

- **Worst-case Message Complexity:  $O(n^2)$** 
  - **Heavy sync** in each Epoch Change
- **Heavy Sync** is **not always necessary** after GST, once it satisfies “**success criterion**” (see QCs).
  - **At least  $f+1$  honest replicas have synchronized clock.**
- **After GST, the Message Complexity:  $O(f_a n + n)$** 
  - In each epoch,  $f_a$  faulty leader will cause  $f_a n$  message complexity, and  $n$  for a true leader



# View Change in Lumiere

- Whenever receiving QC, it would perform “Clock Bumping”.
- **Light Sync** in View change ( $O(n)$ )
  - Phase 1: Send “**view v**” msg to the certain leader
  - Phase 2.1: Leader combining **f+1** “**view v**” msg as VC(View Certificate)
  - Phase 2.2: When sb see a **VC**, enter view v




# Clock Bumping (Forwarding) overview

- Purpose of forwarding clocks
  - Ensures that nodes can adjust their clocks to stay in sync with the rest of the network despite potential network delays or Byzantine behavior
- Occurs when receiving a QC in a non initial view or VC for the initial view
- Importance of clock forwarding
  - Prevent nodes with delayed clocks from being stuck in old views


# Clock Bumping (Clock Forwarding) in Lumiere:

$\Gamma$ : Expectation  
time for a view

**LP22:**



**Lumiere:**

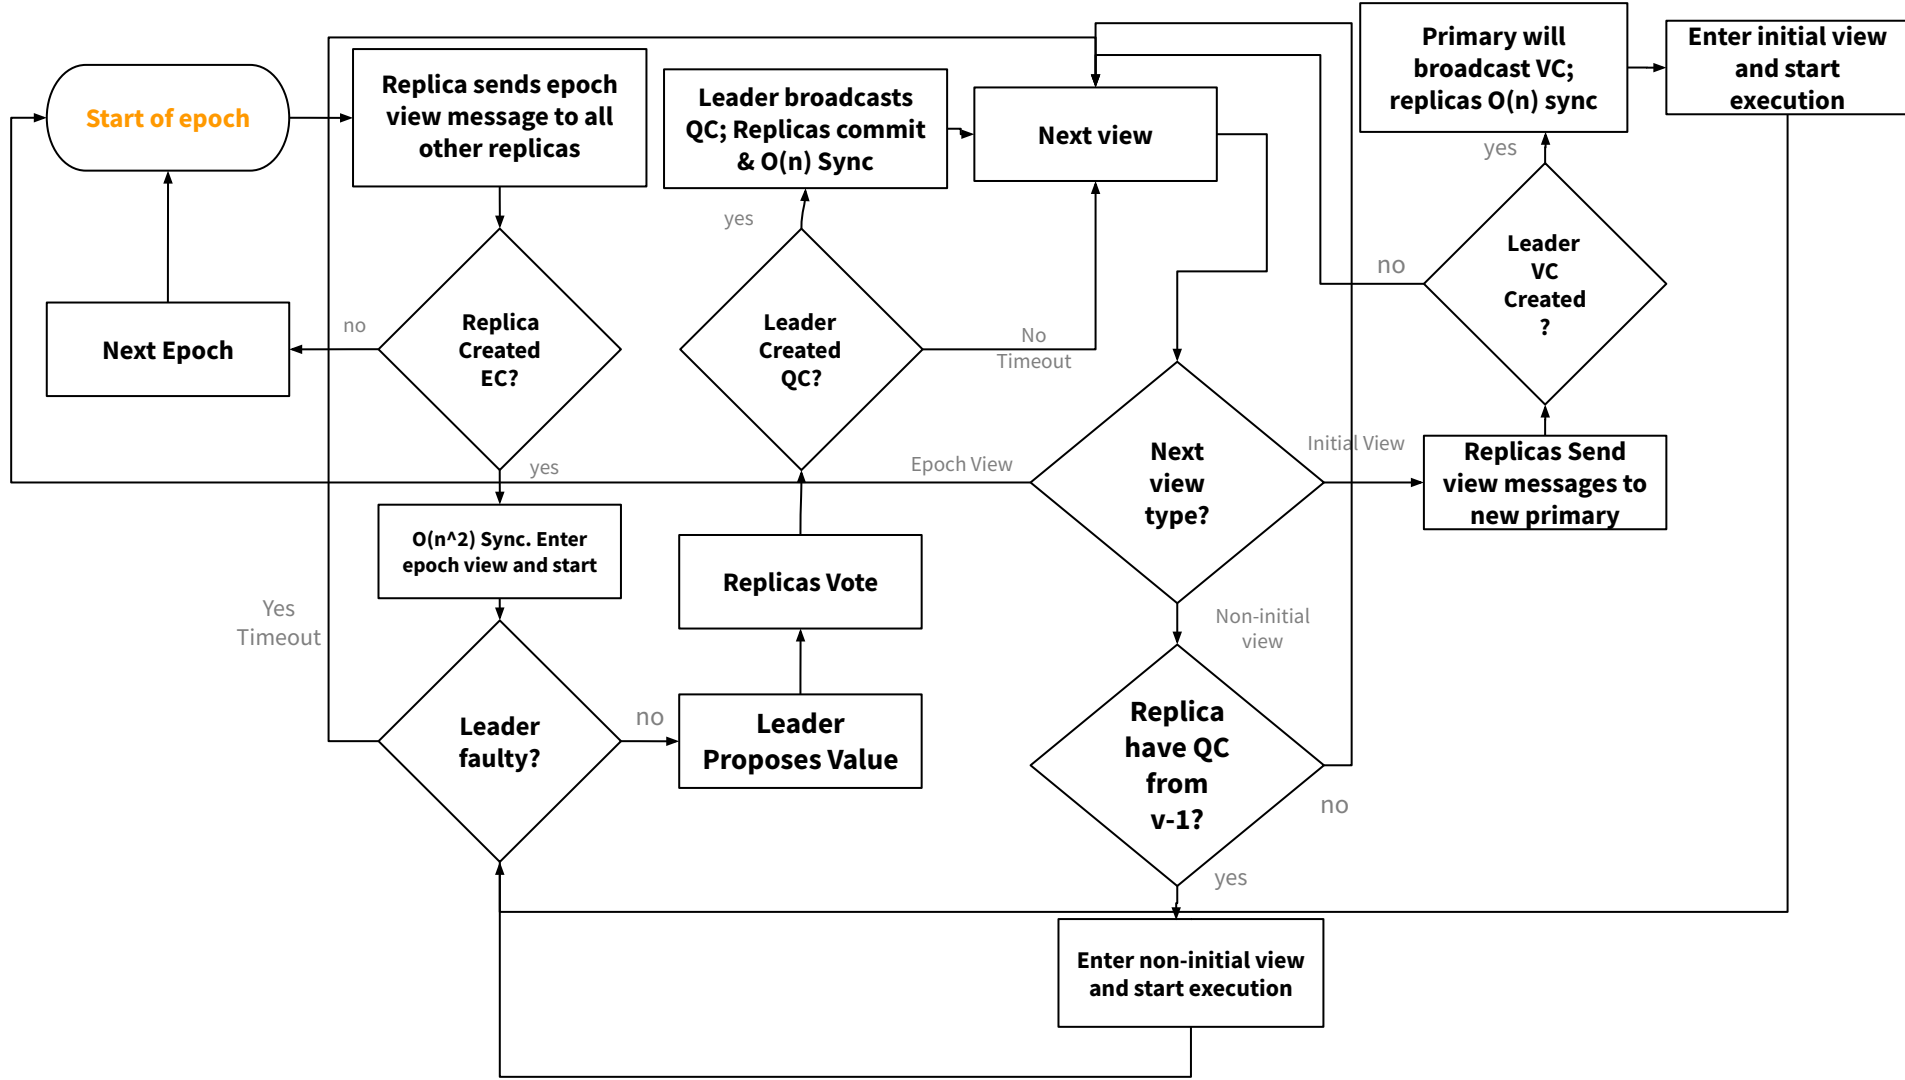


# The Flow

Flowchart of the protocol







**Thank You**