

# **REACHING AGREEMENT IN THE PRESENCE OF FAULTS**

M. PEASE, R. SHOSTAK, AND L. LAMPORT

**Ashay Katkar**

**Sameeksha Mahesh**

**Sarvesh Halbe**

**Ritesh Patro**

# ABSTRACT

- Each nonfaulty processor has a private value of the information that must be communicated to each other nonfaulty processor.
- Nonfaulty processors always communicate honestly, whereas faulty processors may lie.
- The problem is to devise an algorithm in which processors communicate their own values and relay values received from others that allows each nonfaulty processor to refer a value for each other processor

# INTERACTIVE CONSISTENCY

- Given  $m, n \geq 0$ , is it possible to devise an algorithm based on an exchange of messages that will allow each nonfaulty processor 'p' to compute a vector of values with an element for each of the  $n$  processors, such that:
  1. the non-faulty processors compute exactly the same vector;
  2. the element of this vector corresponding to a given non-faulty processor is the private value of the processor.
- The computed vector is called an interactive consistency (i.c) vector.

# SINGLE-FAULT CASE

- Round 1 (direct receives)
- A receives: A:24, B:24, C:24, D:30
- B receives: A:24, B:24, C:24, D:18
- C receives: A:24, B:24, C:24, D:100

# SINGLE-FAULT CASE

- Round 1 (direct receives)
  - A receives: A:24, B:24, C:24, D:30
  - B receives: A:24, B:24, C:24, D:18
  - C receives: A:24, B:24, C:24, D:100
- 
- Each processor forwards the values it heard in Round-1.
  - A forwards to everyone: “B told me 24”, “C told me 24”, “D told me 30”
  - B forwards to everyone: “A told me 24”, “C told me 24”, “D told me 18”
  - C forwards to everyone: “A told me 24”, “B told me 24”, “D told me 100”

# SINGLE-FAULT CASE

A's reports:

- For A: Private value (24)

# SINGLE-FAULT CASE

A's reports:

- For A: Private value (24)
- For B: [24 (direct from B), 24 (from C), 34 (from D)] → majority 24

# SINGLE-FAULT CASE

A's reports:

- For A: Private value (24)
- For B: [24 (direct from B), 24 (from C), 34 (from D)] → majority 24
- For C: [24 (direct from C), 24 (from B), 21 (from D)] → majority 24
- For D: [30 (direct from D), 18 (forward from B), 100 (forward from C)] → no majority → NIL

# SINGLE-FAULT CASE

A's reports:

- For A: Private value (24)
- For B: [24 (direct from B), 24 (from C), 34 (from D)] → majority 24
- For C: [24 (direct from C), 24 (from B), 21 (from D)] → majority 24
- For D: [30 (direct from D), 18 (forward from B), 100 (forward from C)] → no majority → NIL

Final IC vectors (all non-faulty record the same vector)

- A final: [A:24, B:24, C:24, D:NIL]
- B final: [A:24, B:24, C:24, D:NIL]
- C final: [A:24, B:24, C:24, D:NIL]
- IC vector of faulty processor not relevant.

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- $m = 1$ , round 1:  $A \rightarrow B, C, D$

$n-1$

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- $m = 1$ , round 1:  $A \rightarrow B, C, D$  (other 3 processors do the same)  
 $(n * n-1)$

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- $m = 1$ , round 1:  $n * n-1$

round 2: A → forwards to B what it heard from C, D

→ forwards to C what it heard from B, D

→ forwards to D what it heard from B, C

$n-1 * n-2$

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- $m = 1$ , round 1:  $n * n-1$

round 2: A → forwards to B what it heard from C, D

→ forwards to C what it heard from B, D

→ forwards to D what it heard from B, C

$n-1 * n-2$

(every processor does this)

$n * n-1 * n-2$

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- $m = 1$ , round 1:  $n * n-1$

round 2: A → forwards to B what it heard from C, D

→ forwards to C what it heard from B, D

→ forwards to D what it heard from B, C

(every processor does this)

$$(n * n-1) + (n * n-1 * n-2) \rightarrow \text{order of } n^3$$

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- $m = 1$ , round 1:  $n * n-1$

round 2:  $(n * n-1) + (n * n-1 * n-2) \rightarrow$  order of  $n^3$

- $m = 2$ , round 1:  $n * n-1$

- round 2:  $(n * n-1) + (n * n-1 * n-2)$

- round 3:  $((n * n-1) + (n * n-1 * n-2)) + (n * n-1 * n-2 * n-3) \rightarrow$  order of  $n^4$

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- $m = 1$ , round 1:  $n * n-1$

round 2:  $(n * n-1) + (n * n-1 * n-2) \rightarrow$  order of  $n^3$

- $m = 2$ , round 1:  $n * n-1$

- round 2:  $(n * n-1) + (n * n-1 * n-2)$

- round 3:  $((n * n-1) + (n * n-1 * n-2)) + (n * n-1 * n-2 * n-3) \rightarrow$  order of  $n^4$

$$M_{IC}(n, m) = \Theta(n^{m+2})$$

# **N >= 3M+1**

- Message rounds:

In the general case of  $m$  faults,  $m + 1$  rounds are required in order to describe the algorithm.

Each extra round gives honest processors a chance to verify and cross-check the information that might have been tampered with by faulty ones.

- Number of processors:

For consensus to hold, you need more than twice as many honest nodes as faulty ones.

$$\text{Total} = \text{honest} + \text{faulty} = (2m+1) + m = 3m + 1$$

(smallest number of honest processors that can still form a majority against any set of  $m$  liars)

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

- $p$  and  $p'$  are both nonfaulty processors

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

- p and p' are both nonfaulty processors

p found set with more than  $(n+m)/2$  members, all agreeing on some value for processor q

Vector p → { \_ \_ \_ q( v ) \_ \_ \_ }

Voted by greater than  $(n+m)/2$  other processors.

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

- $p'$  found set with more than  $(n+m)/2$  members, all agreeing on some value for processor  $q$

Vector  $p' \rightarrow \{ \_ \_ \_ q(v') \_ \_ \_ \}$

Voted by greater than  $(n+m)/2$  other processors.

Vector  $p \rightarrow \{ \_ \_ \_ q(v) \_ \_ \_ \}$

Vector  $p' \rightarrow \{ \_ \_ \_ q(v') \_ \_ \_ \}$

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

- $p'$  found set with more than  $(n+m)/2$  members, all agreeing on some value for processor  $q$

Vector  $p' \rightarrow \{ \_ \_ \_ q(v') \_ \_ \_ \}$

Voted by greater than  $(n+m)/2$  other processors.

Vector  $p \rightarrow \{ \_ \_ \_ q(v) \_ \_ \_ \}$

Vector  $p' \rightarrow \{ \_ \_ \_ q(v') \_ \_ \_ \}$

$$(n+m)/2 + (n+m)/2 = n+m$$

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

- $p'$  found set with more than  $(n+m)/2$  members, all agreeing on some value for processor  $q$

Vector  $p' \rightarrow \{ \_ \_ \_ q(v') \_ \_ \_ \}$

Voted by greater than  $(n+m)/2$  other processors.

Vector  $p \rightarrow \{ \_ \_ \_ q(v) \_ \_ \_ \}$

Vector  $p' \rightarrow \{ \_ \_ \_ q(v') \_ \_ \_ \}$

$$(n+m)/2 + (n+m)/2 = n+m \text{ (only } n \text{ processors!)}$$

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

$$\begin{aligned} q(v) + q(v') \\ &> (n+m)/2 + > (n+m)/2 \\ &> (n+m) \end{aligned}$$

Remove n processor to find overlapping processors.

> m

- Number of faulty procesors is ‘m’.

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

$$\begin{aligned} q(v) + q(v') \\ &> (n+m)/2 + > (n+m)/2 \\ &> (n+m) \end{aligned}$$

Remove n processor to find overlapping processors.

> m

- Number of faulty procesors is ‘m’.
- Therefore, at least one processor in the overlap must be nonfaulty.

# **N >= 3M+1**

Threshold:  $Q >(n+m)/2$

$$\begin{aligned} q(v) + q(v') \\ &> (n+m)/2 + > (n+m)/2 \\ &> (n+m) \end{aligned}$$

Remove n processor to find overlapping processors.

> m

- Number of faulty procesors is ‘m’.
- Therefore, at least one processor in the overlap must be nonfaulty.
- The two sets must give rise to the same value v.

# **N < 3M+1 (IMPOSSIBILITY)**

**A**  
**10**

**B**  
**10**

**C**  
**10**

- Group C sends private (two-party) messages to each group.  
To Group A, it claims its value is “v”.  
To Group B, it claims its value is “w”.

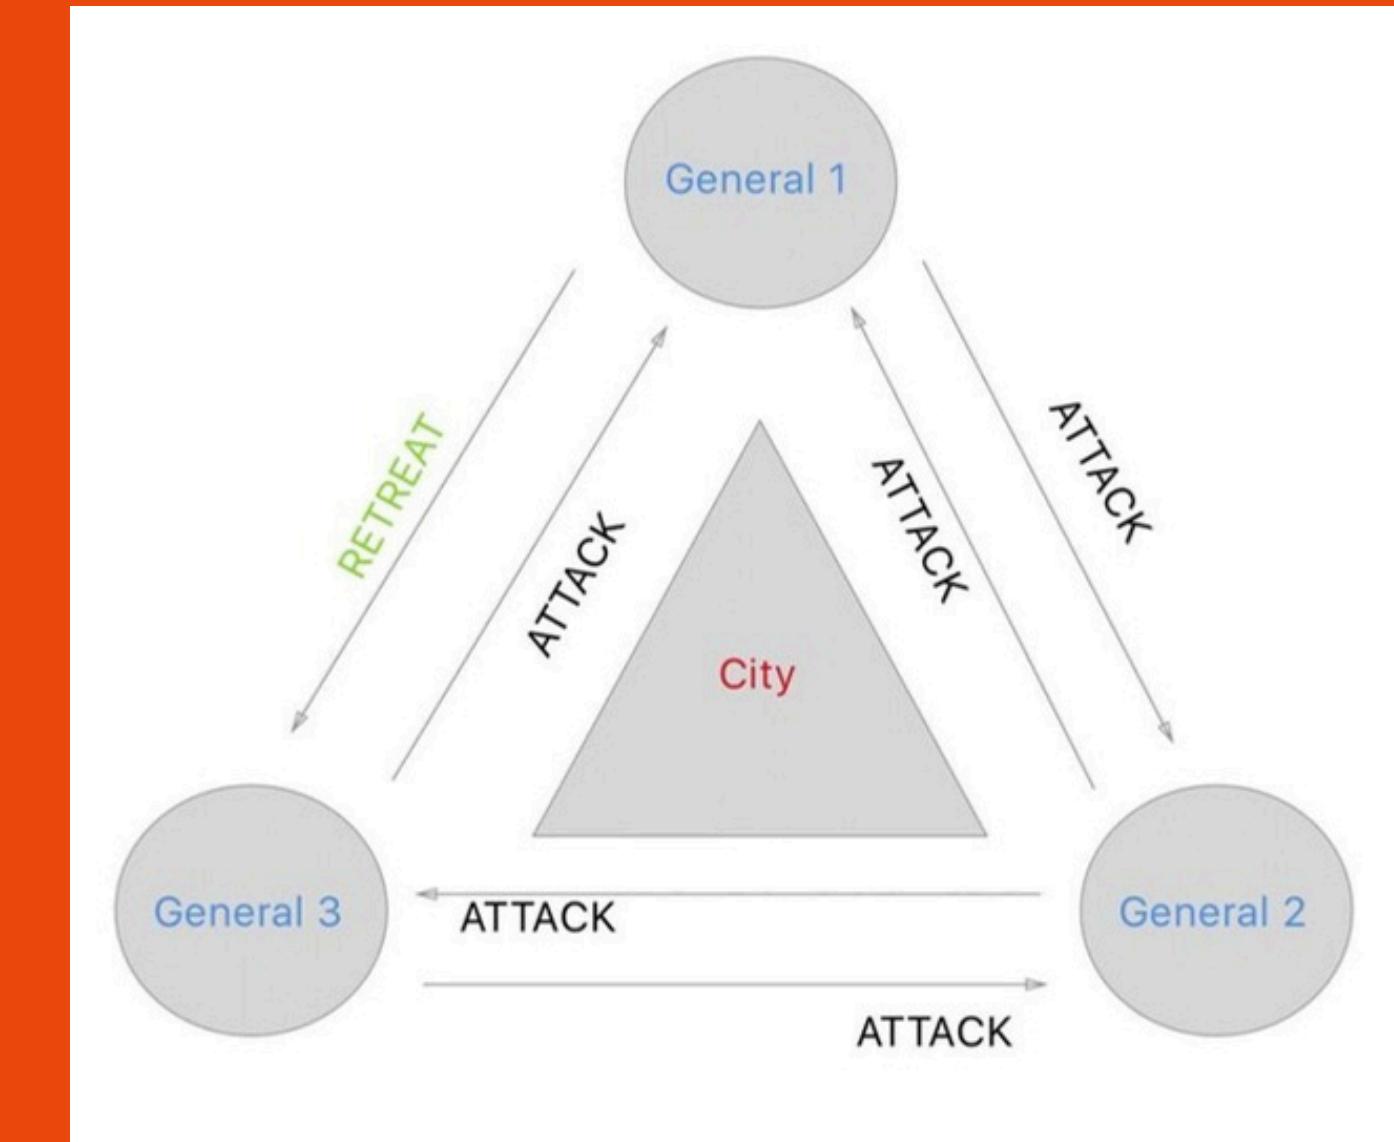
Since A and B have conflicting entries for C’s value (v vs. w),  
there is no common majority among the nonfaulty processors.

# **THE BYZANTINE GENERALS PROBLEM**

LESLIE LAMPORT, ROBERT SHOSTAK, AND  
MARSHALL PEASE

# INTRODUCTION

- All loyal generals decide upon the same plan of action.
- A small number of traitors cannot cause the loyal generals to adopt a bad plan.



# INTRODUCTION

- Each general sends their value  $v(i)$  using a Byzantine solution.
- Other generals act as lieutenants, receiving and forwarding messages.
- $v(i)$  be the information communicated by the  $i$ th general.
  - a. Every loyal general must obtain the same information  $v(1) \dots v(n)$ .
  - b. If the  $i$ th general is loyal, then the value that he sends must be used by every loyal general as the value of  $v(i)$ .

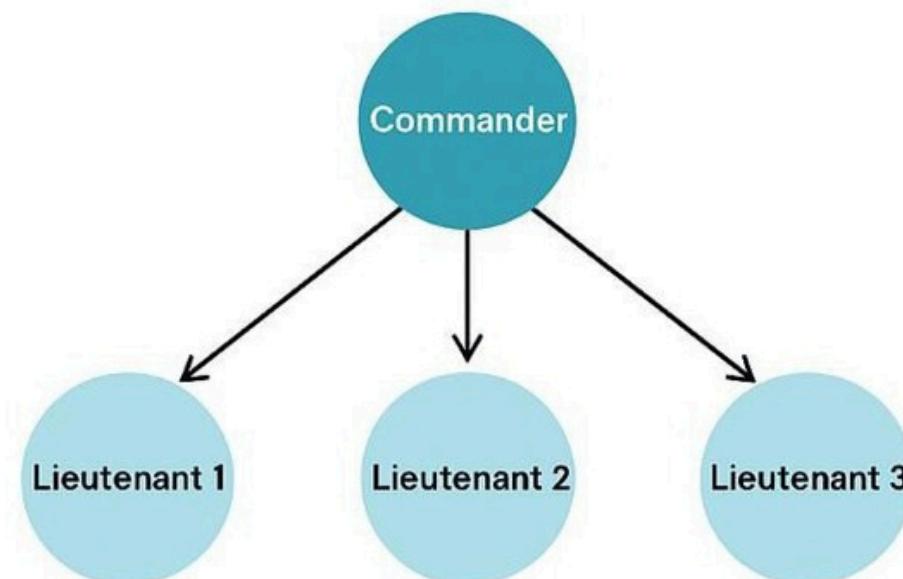
# INTRODUCTION

- A commanding general sends an order to  $n-1$  lieutenants.
- Interactive Consistency Conditions:
  1. IC1  $\rightarrow$  All loyal lieutenants obey the same order.
  2. IC2  $\rightarrow$  If the commander is loyal, all loyal lieutenants obey his order.

## RELIABLE SYSTEMS

All loyal generals decide upon the same plan of action

A small number of traitors cannot cause the loyal generals to adopt a bad plan



# IMPOSSIBILITY RESULTS

- Prove that any problem using oral messages with fewer than  $3m+1$  generals cannot handle  $m$  traitors
- First prove with only three generals that single traitor makes problem unsolvable, then extrapolate

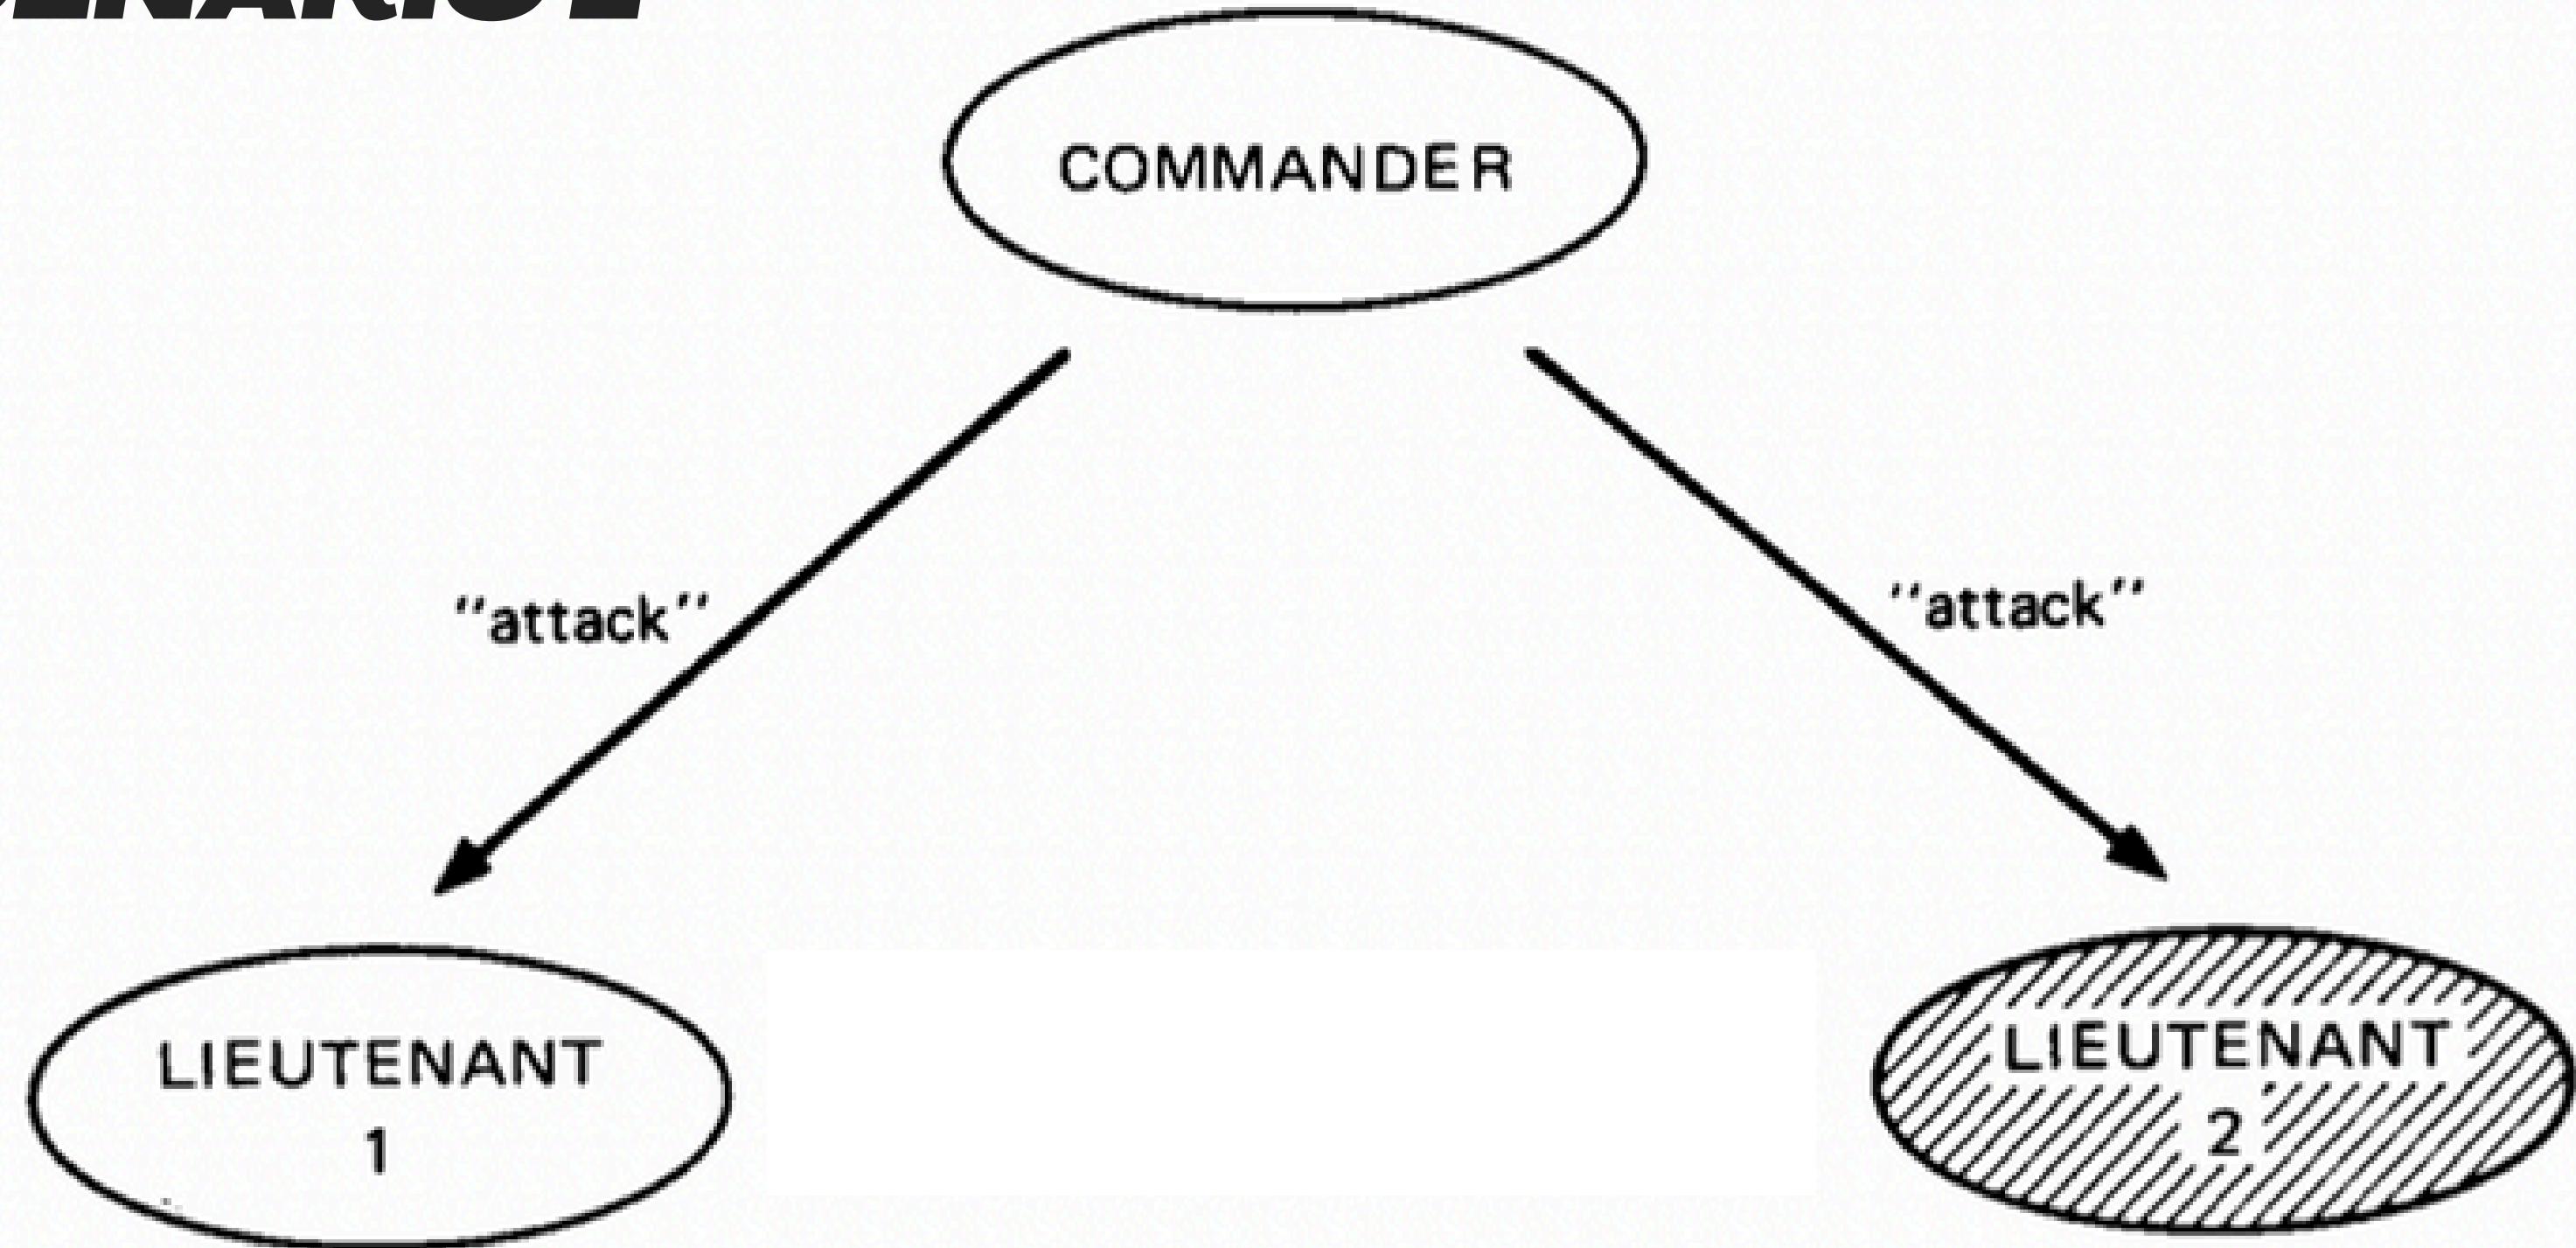
# IMPOSSIBILITY RESULTS

## SCENARIO 1



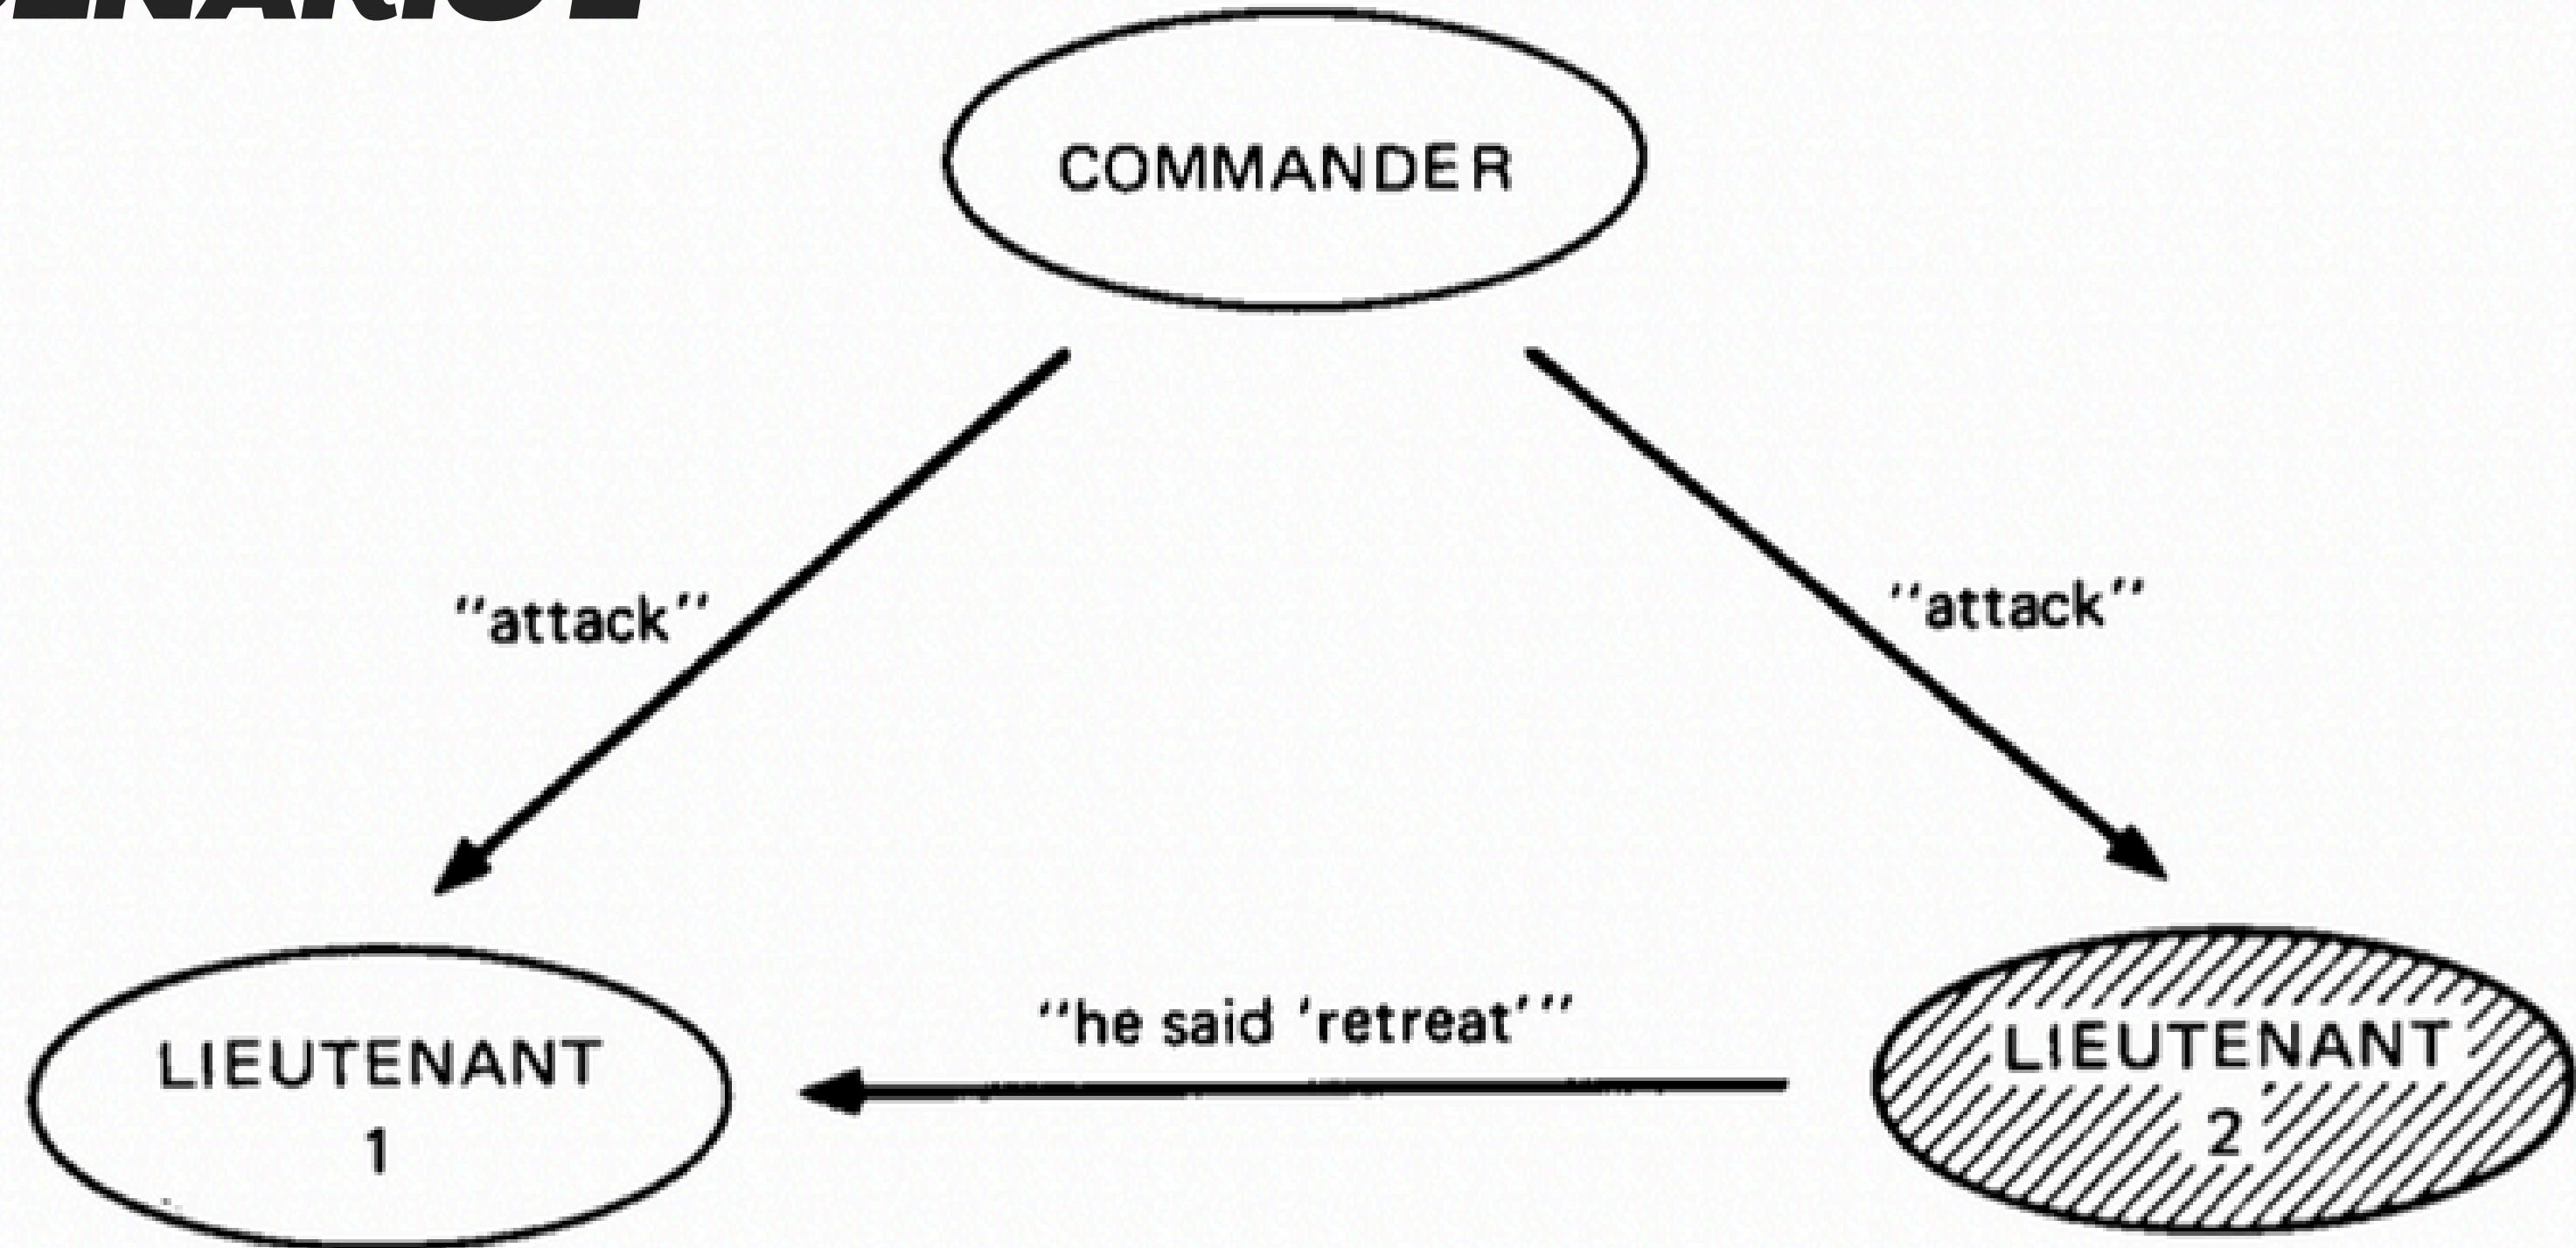
# IMPOSSIBILITY RESULTS

## SCENARIO 1



# IMPOSSIBILITY RESULTS

## SCENARIO 1



# IMPOSSIBILITY RESULTS

## SCENARIO 2

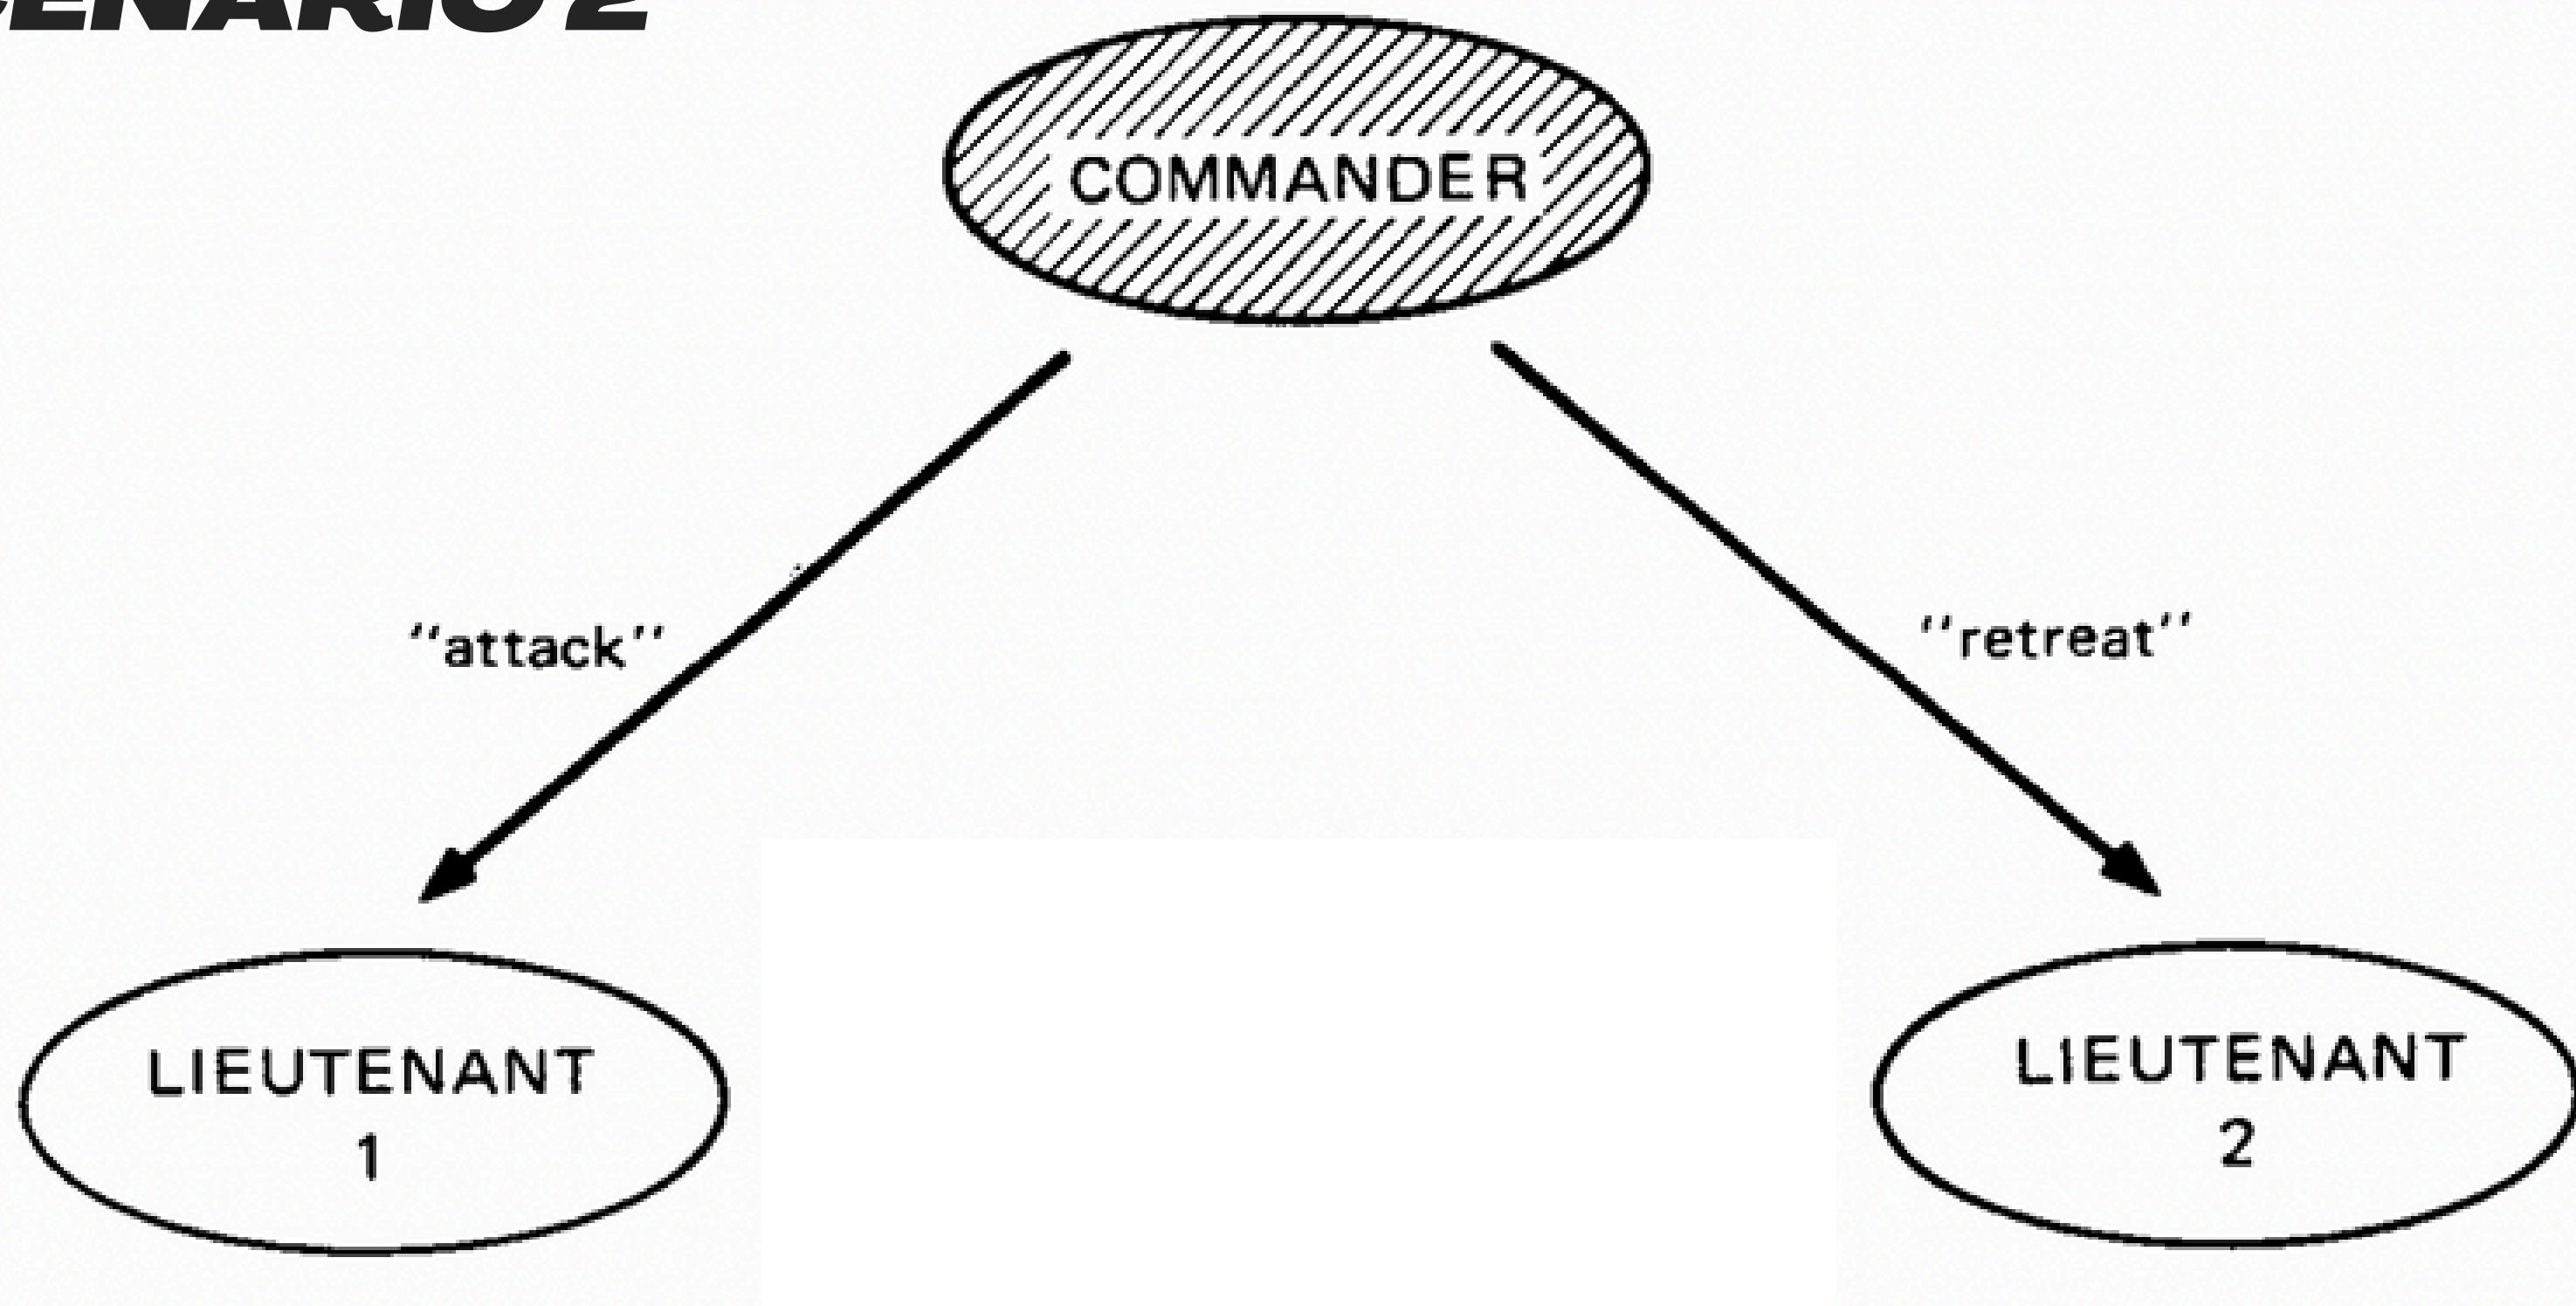


LIEUTENANT  
1

LIEUTENANT  
2

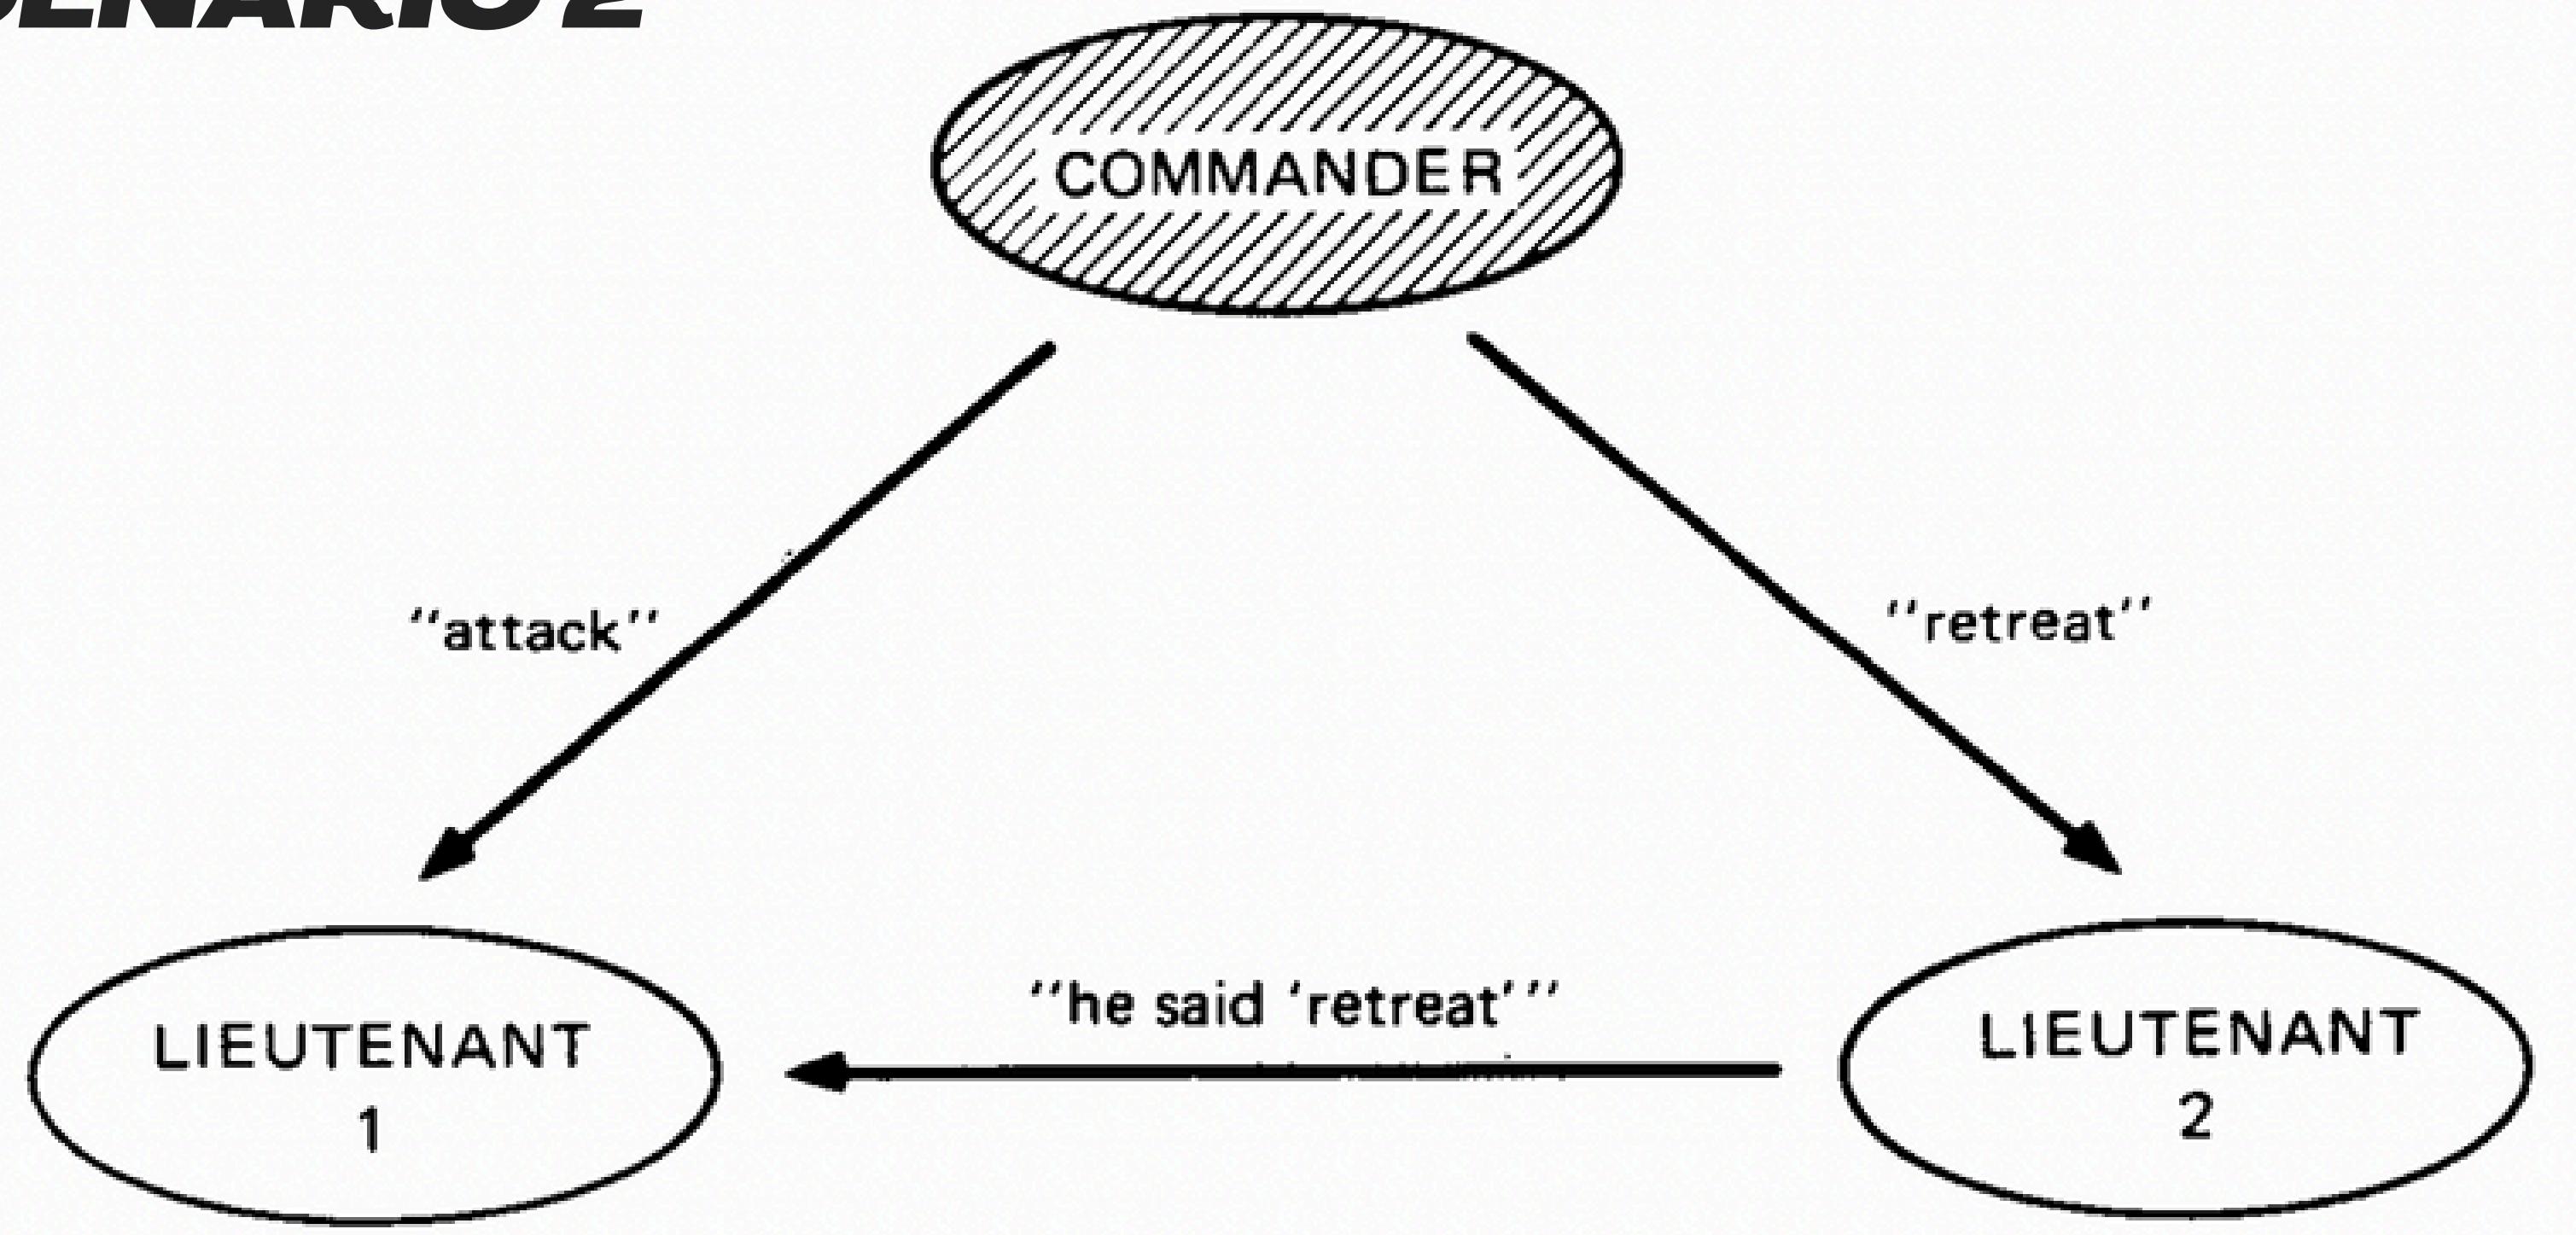
# IMPOSSIBILITY RESULTS

## SCENARIO 2



# IMPOSSIBILITY RESULTS

## SCENARIO 2



# **SOLUTION WITH ORAL MESSAGES**

## **MESSAGE SYSTEM ASSUMPTIONS**

A1: Every message that is sent is delivered correctly.

A2: The receiver of a message knows who sent it.

A3: The absence of a message can be detected.

### Majority function assumptions:

Majority value among inputs if it exists(takes median value if from ordered set), otherwise the value RETREAT

# SOLUTION WITH ORAL MESSAGES

## **ALGORITHM OM(*m*)**

*m* = 0:

1. The commander sends his value to every lieutenant.
2. Each lieutenant uses the value he receives from the commander, or uses the value RETREAT if he receives no value.

# SOLUTION WITH ORAL MESSAGES

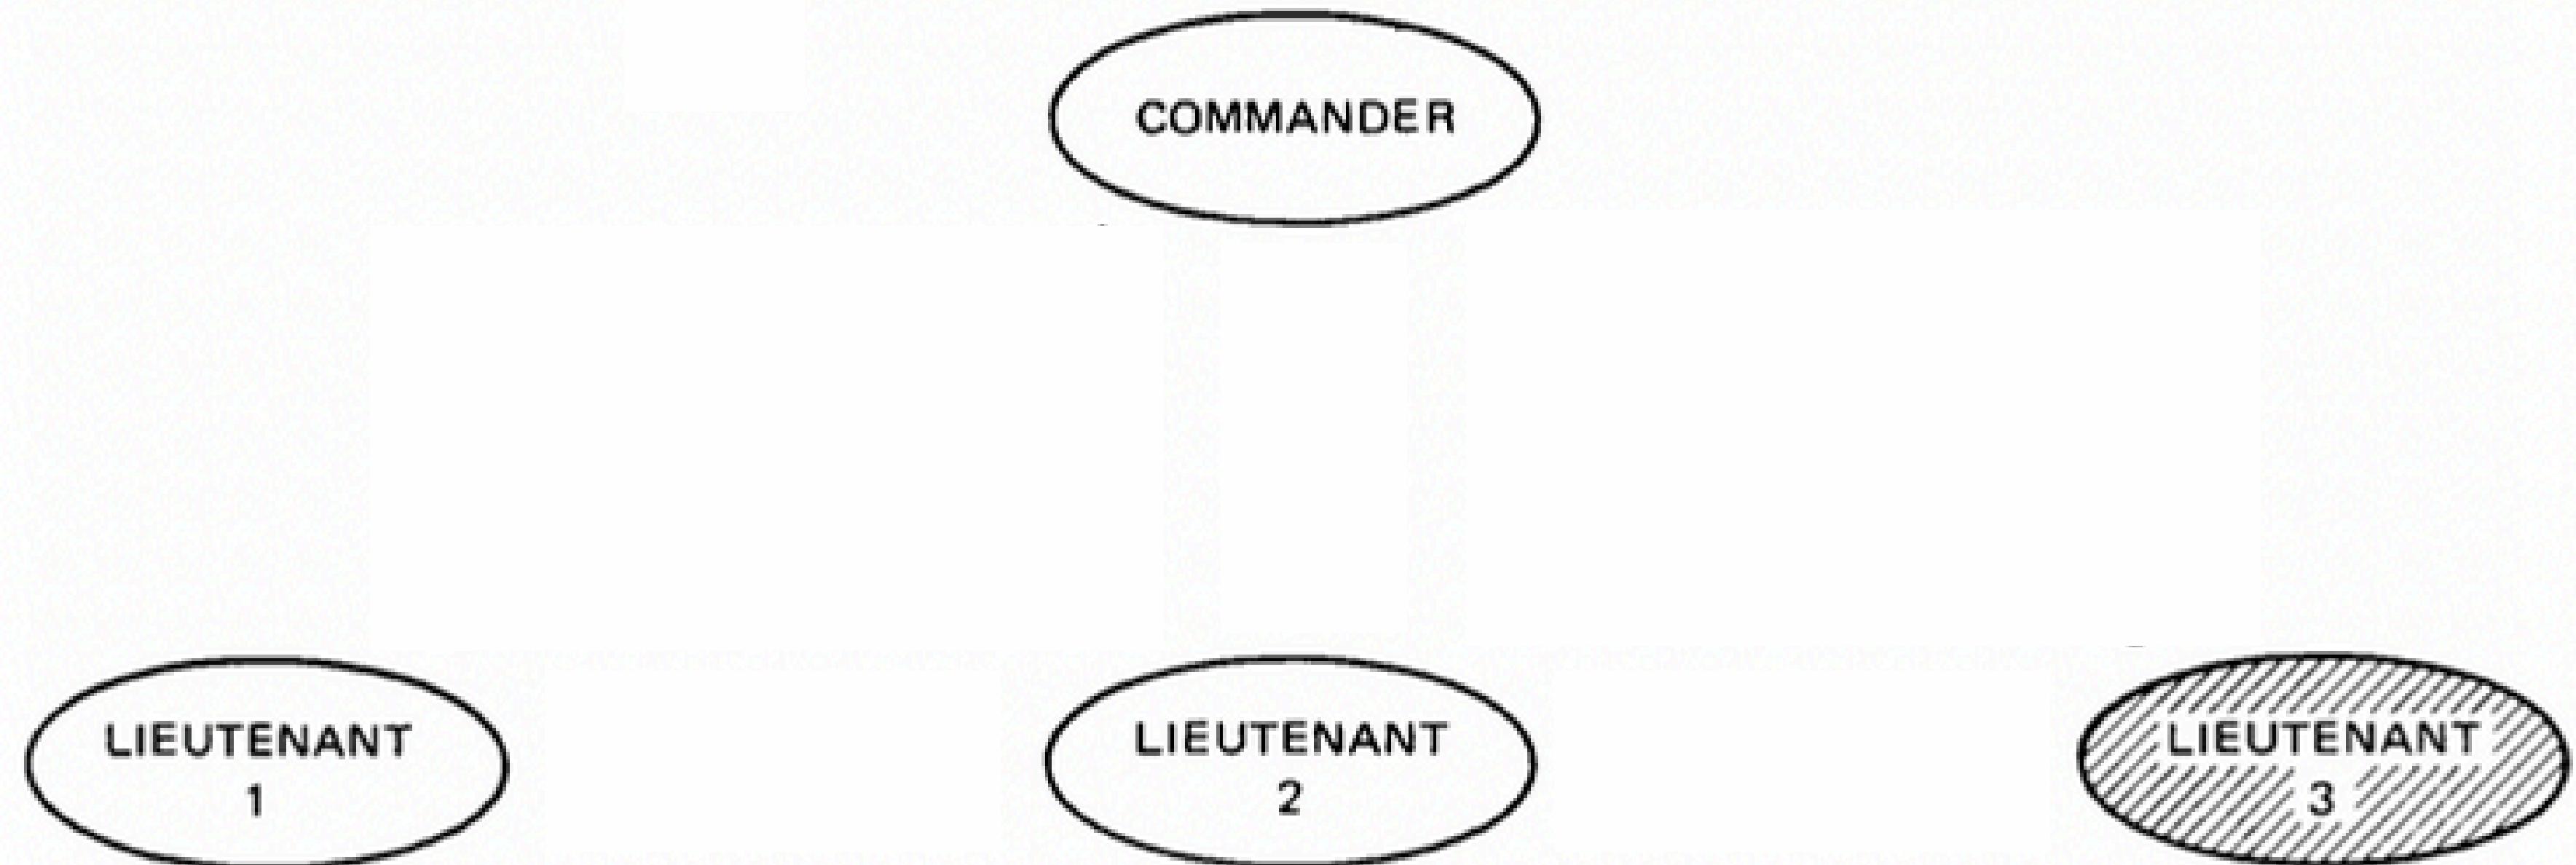
## **ALGORITHM OM( $m$ )**

$m > 0$ :

1. The commander sends his value to every lieutenant.
2. For each  $i$ , let  $v_i$  be the value Lieutenant  $i$  receives from the commander, or else be RETREAT if he receives no value.  
Lieutenant  $i$  acts as the commander in Algorithm OM( $m - 1$ ) to send the value  $v_i$  to each of the  $n - 2$  other lieutenants.
3. For each  $i$ , and each  $j \neq i$ , let  $v_j$  be the value Lieutenant  $i$  received from Lieutenant  $j$  in step (2) (using Algorithm OM( $m - 1$ )), or else RETREAT if he received no such value. Lieutenant  $i$  uses the value majority ( $v_1, \dots, v_{n-1}$ ).

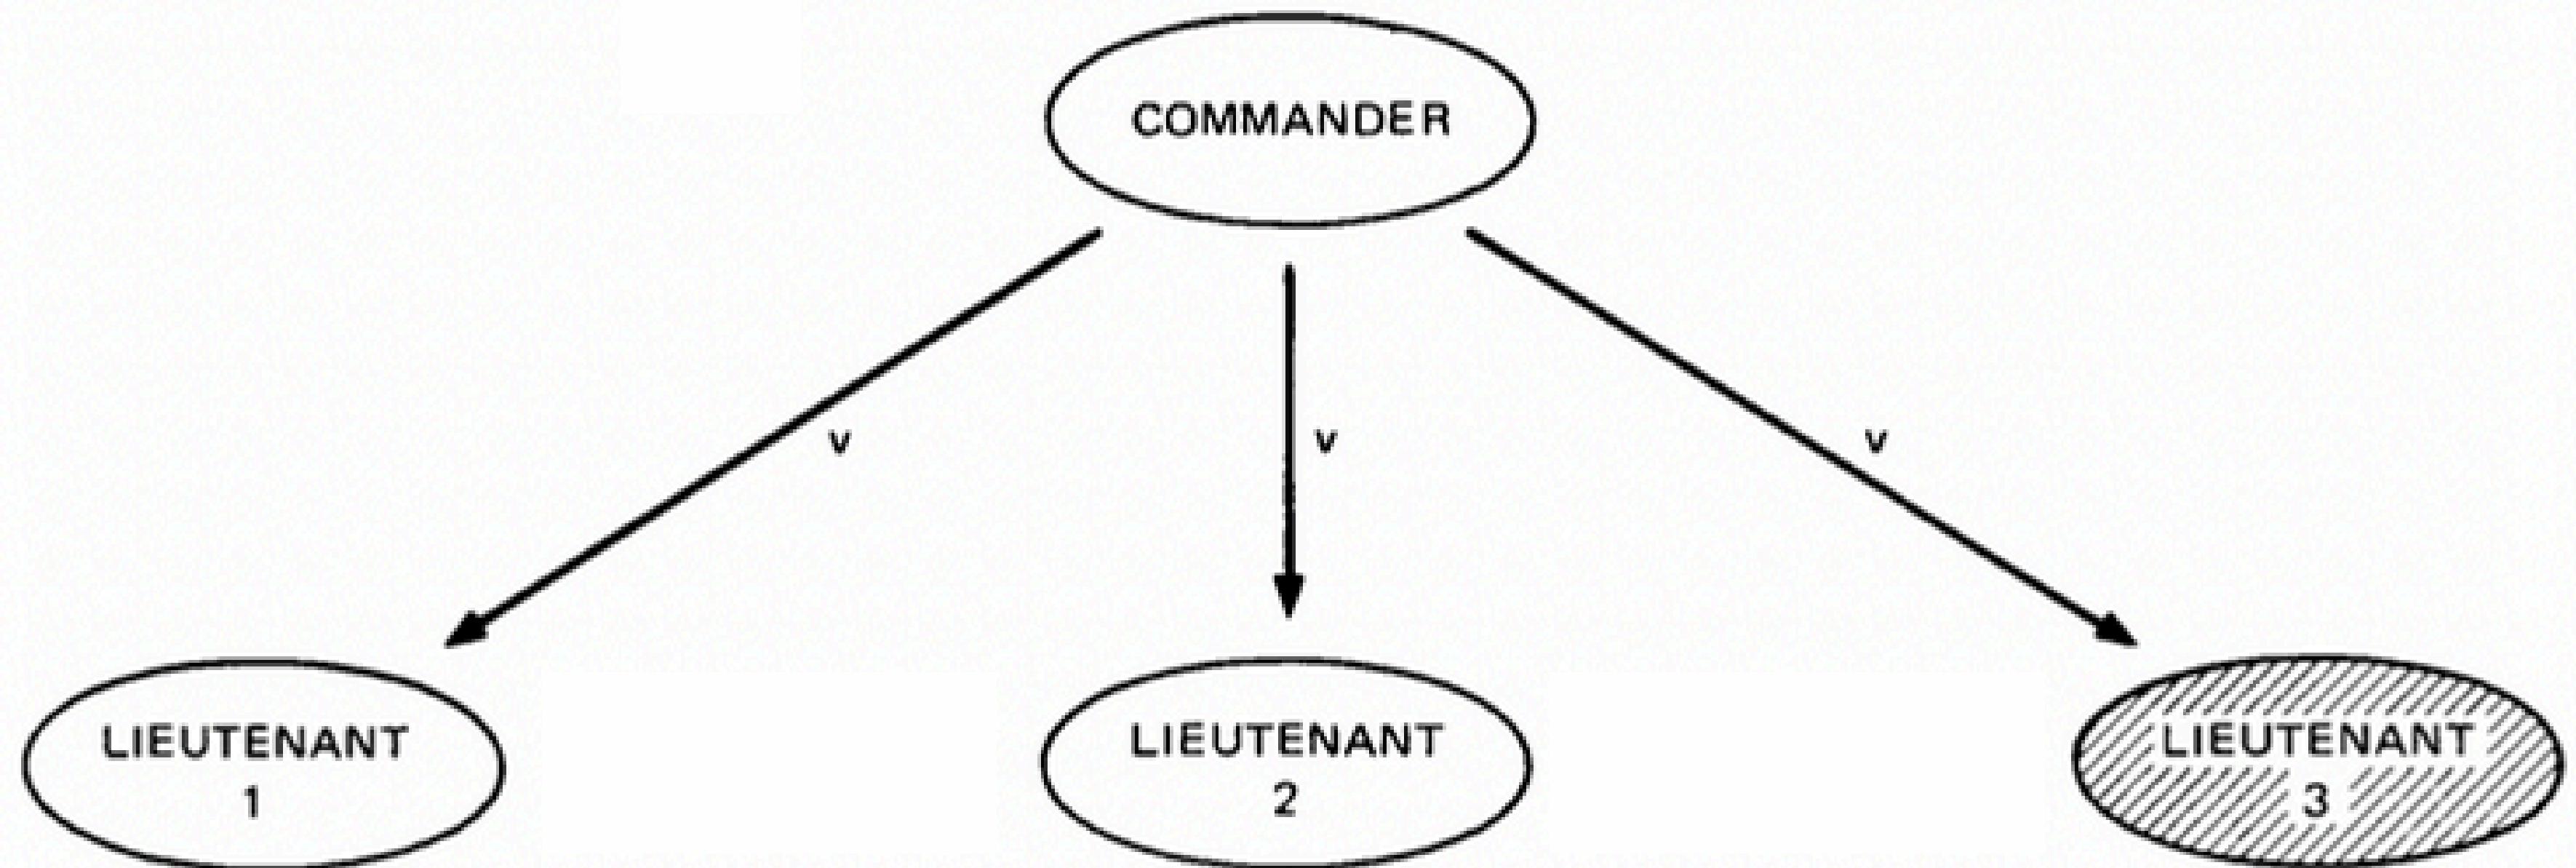
# **SOLUTION WITH ORAL MESSAGES**

## **SCENARIO 1**



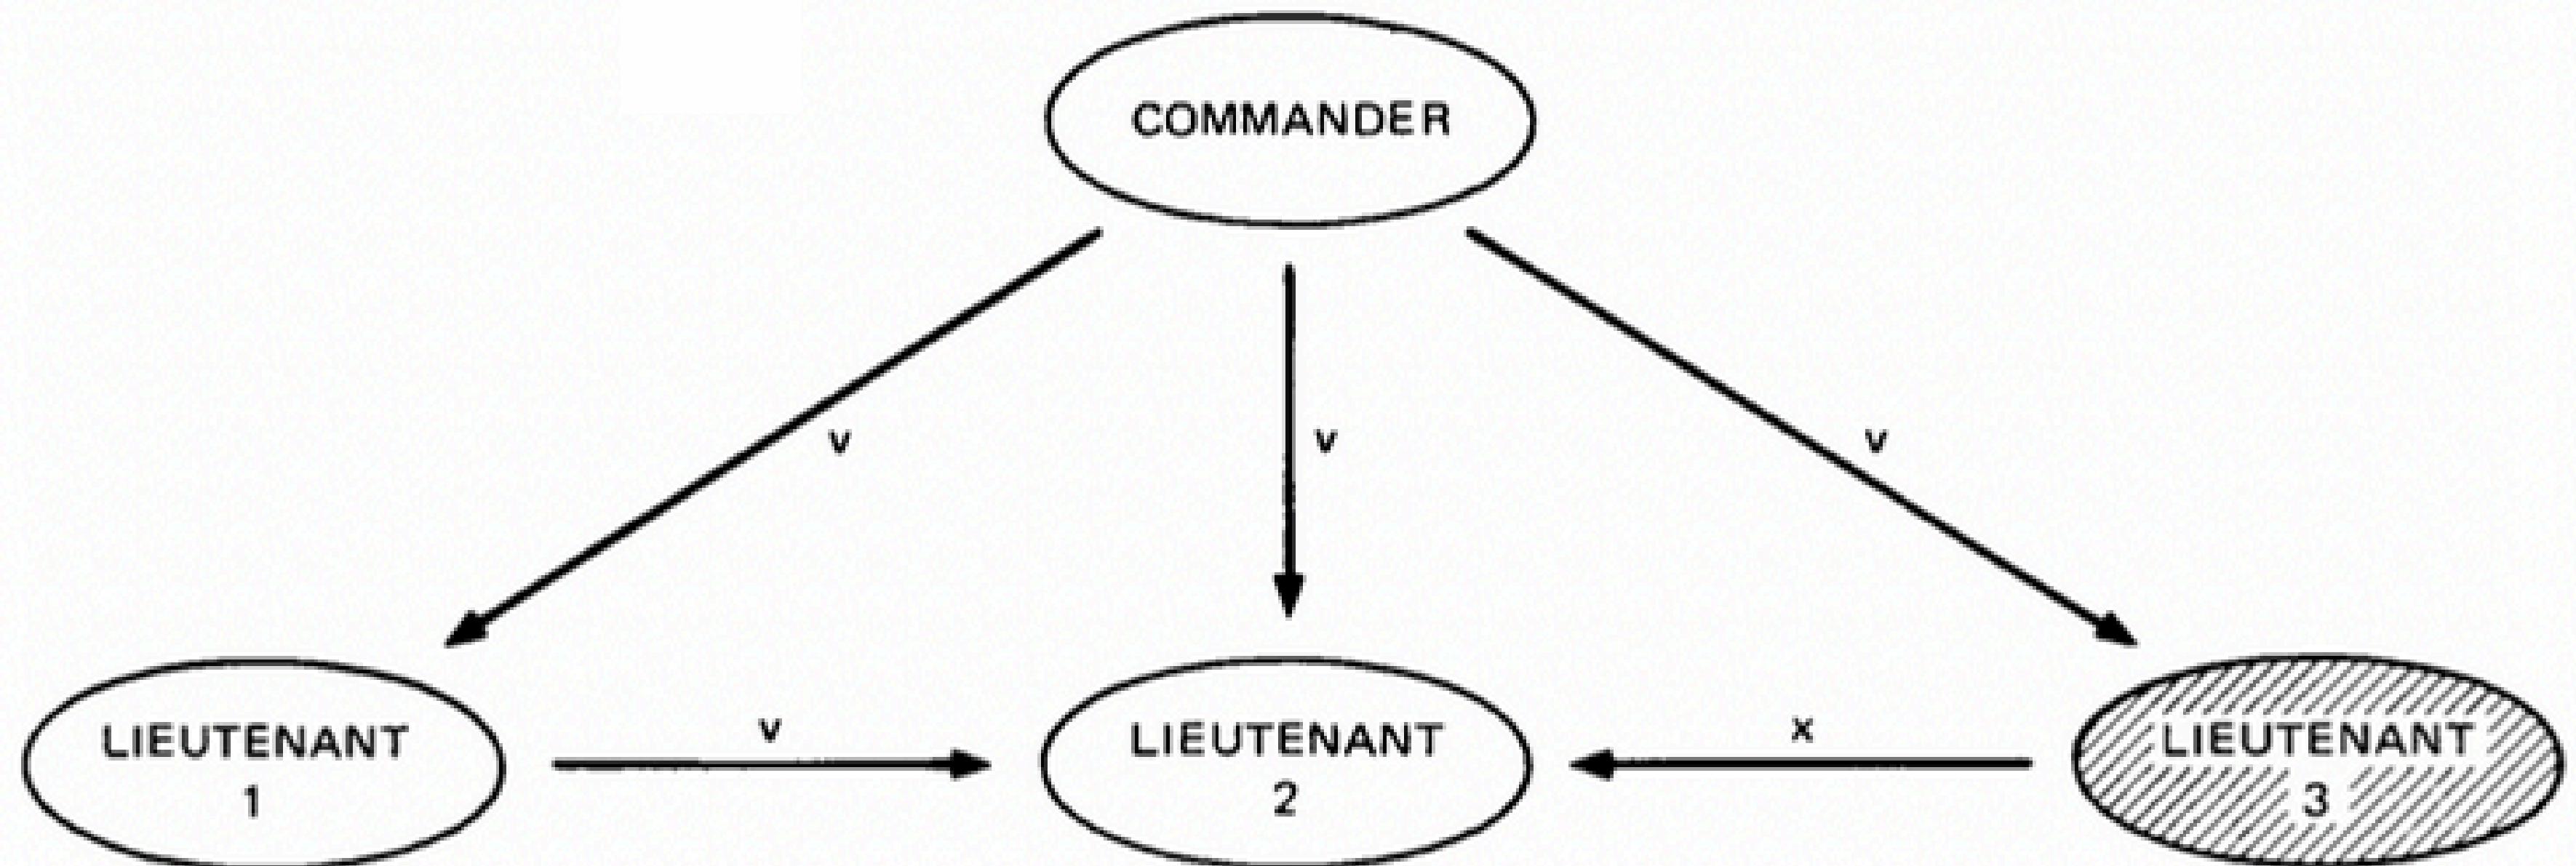
# SOLUTION WITH ORAL MESSAGES

## SCENARIO 1



# SOLUTION WITH ORAL MESSAGES

## SCENARIO 1



# **SOLUTION WITH ORAL MESSAGES**

## **SCENARIO 2**

LIEUTENANT  
1

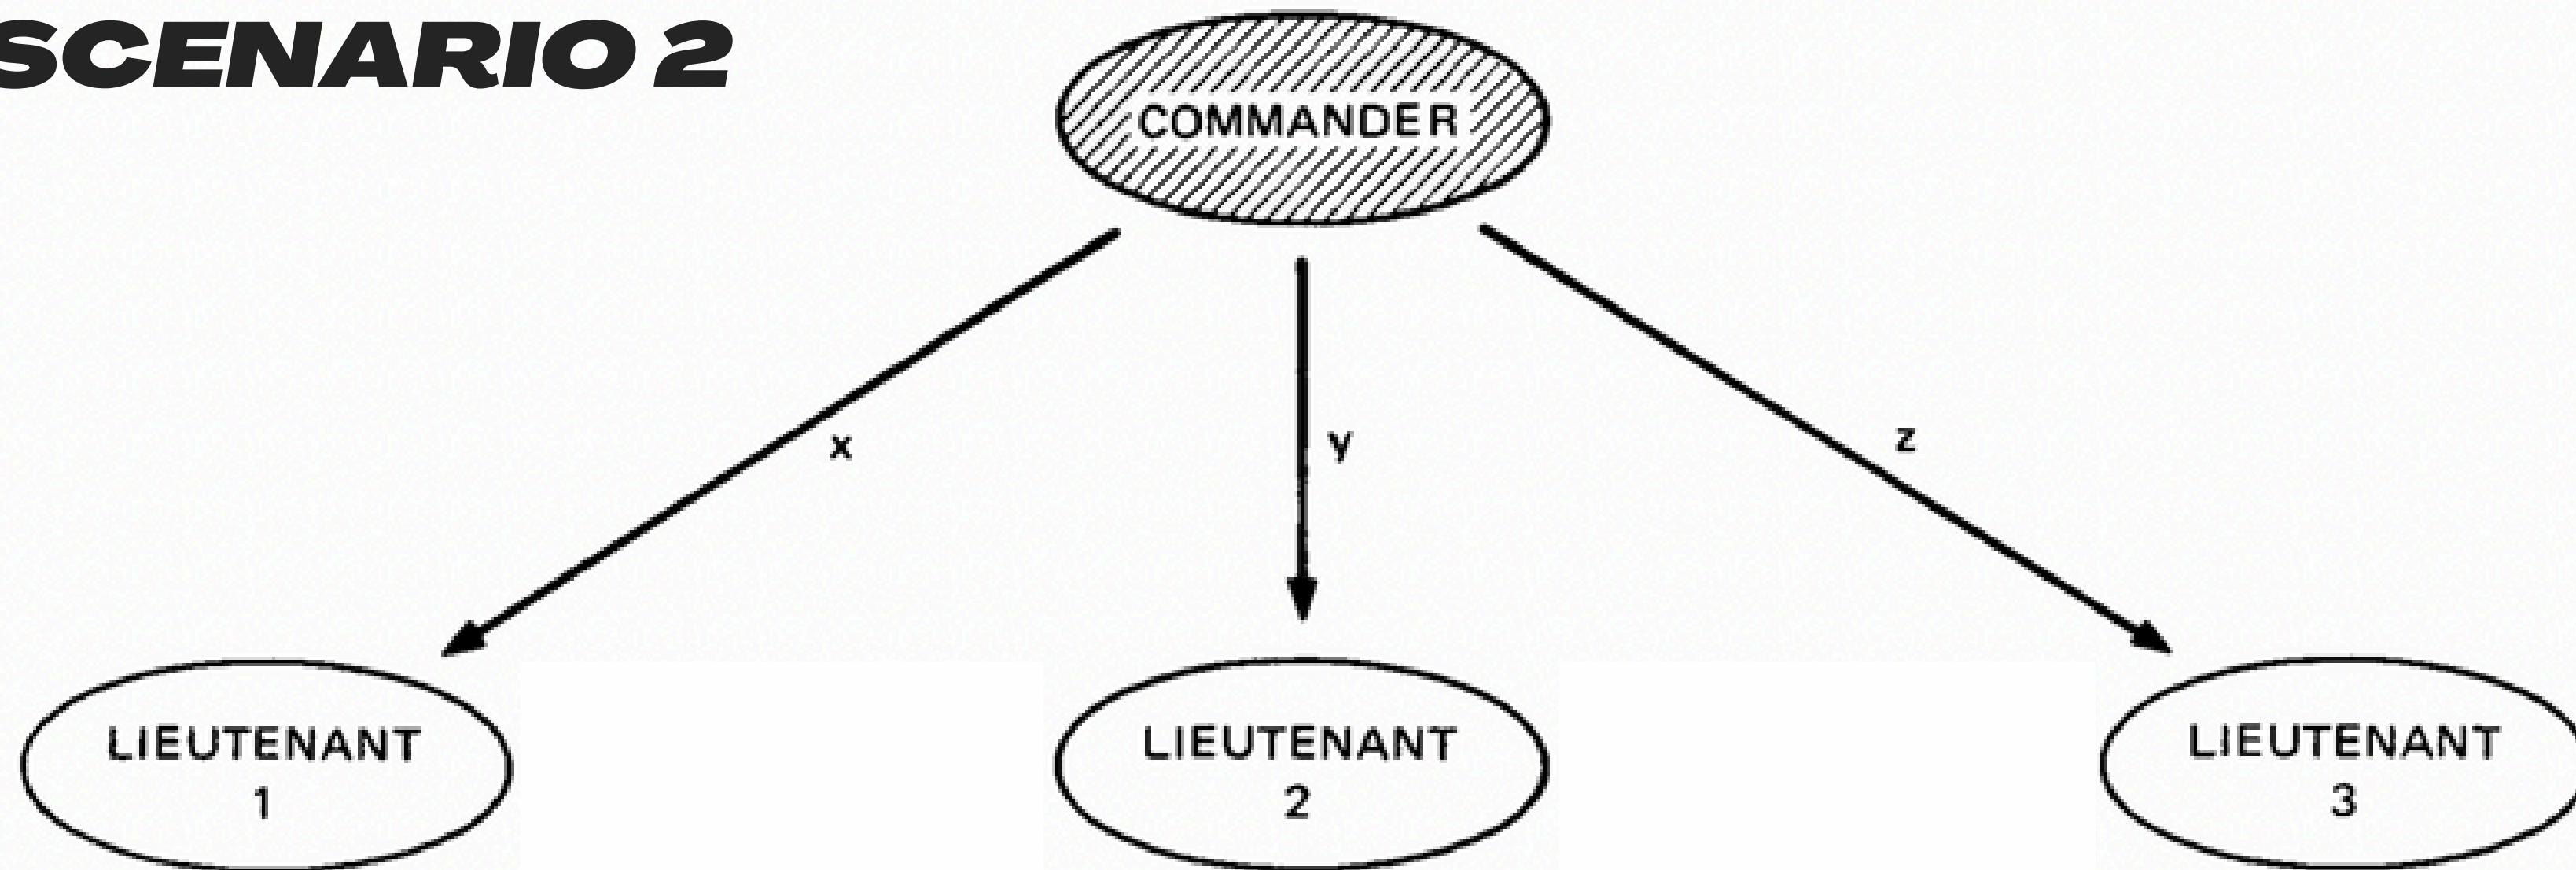
LIEUTENANT  
2

LIEUTENANT  
3

COMMANDER

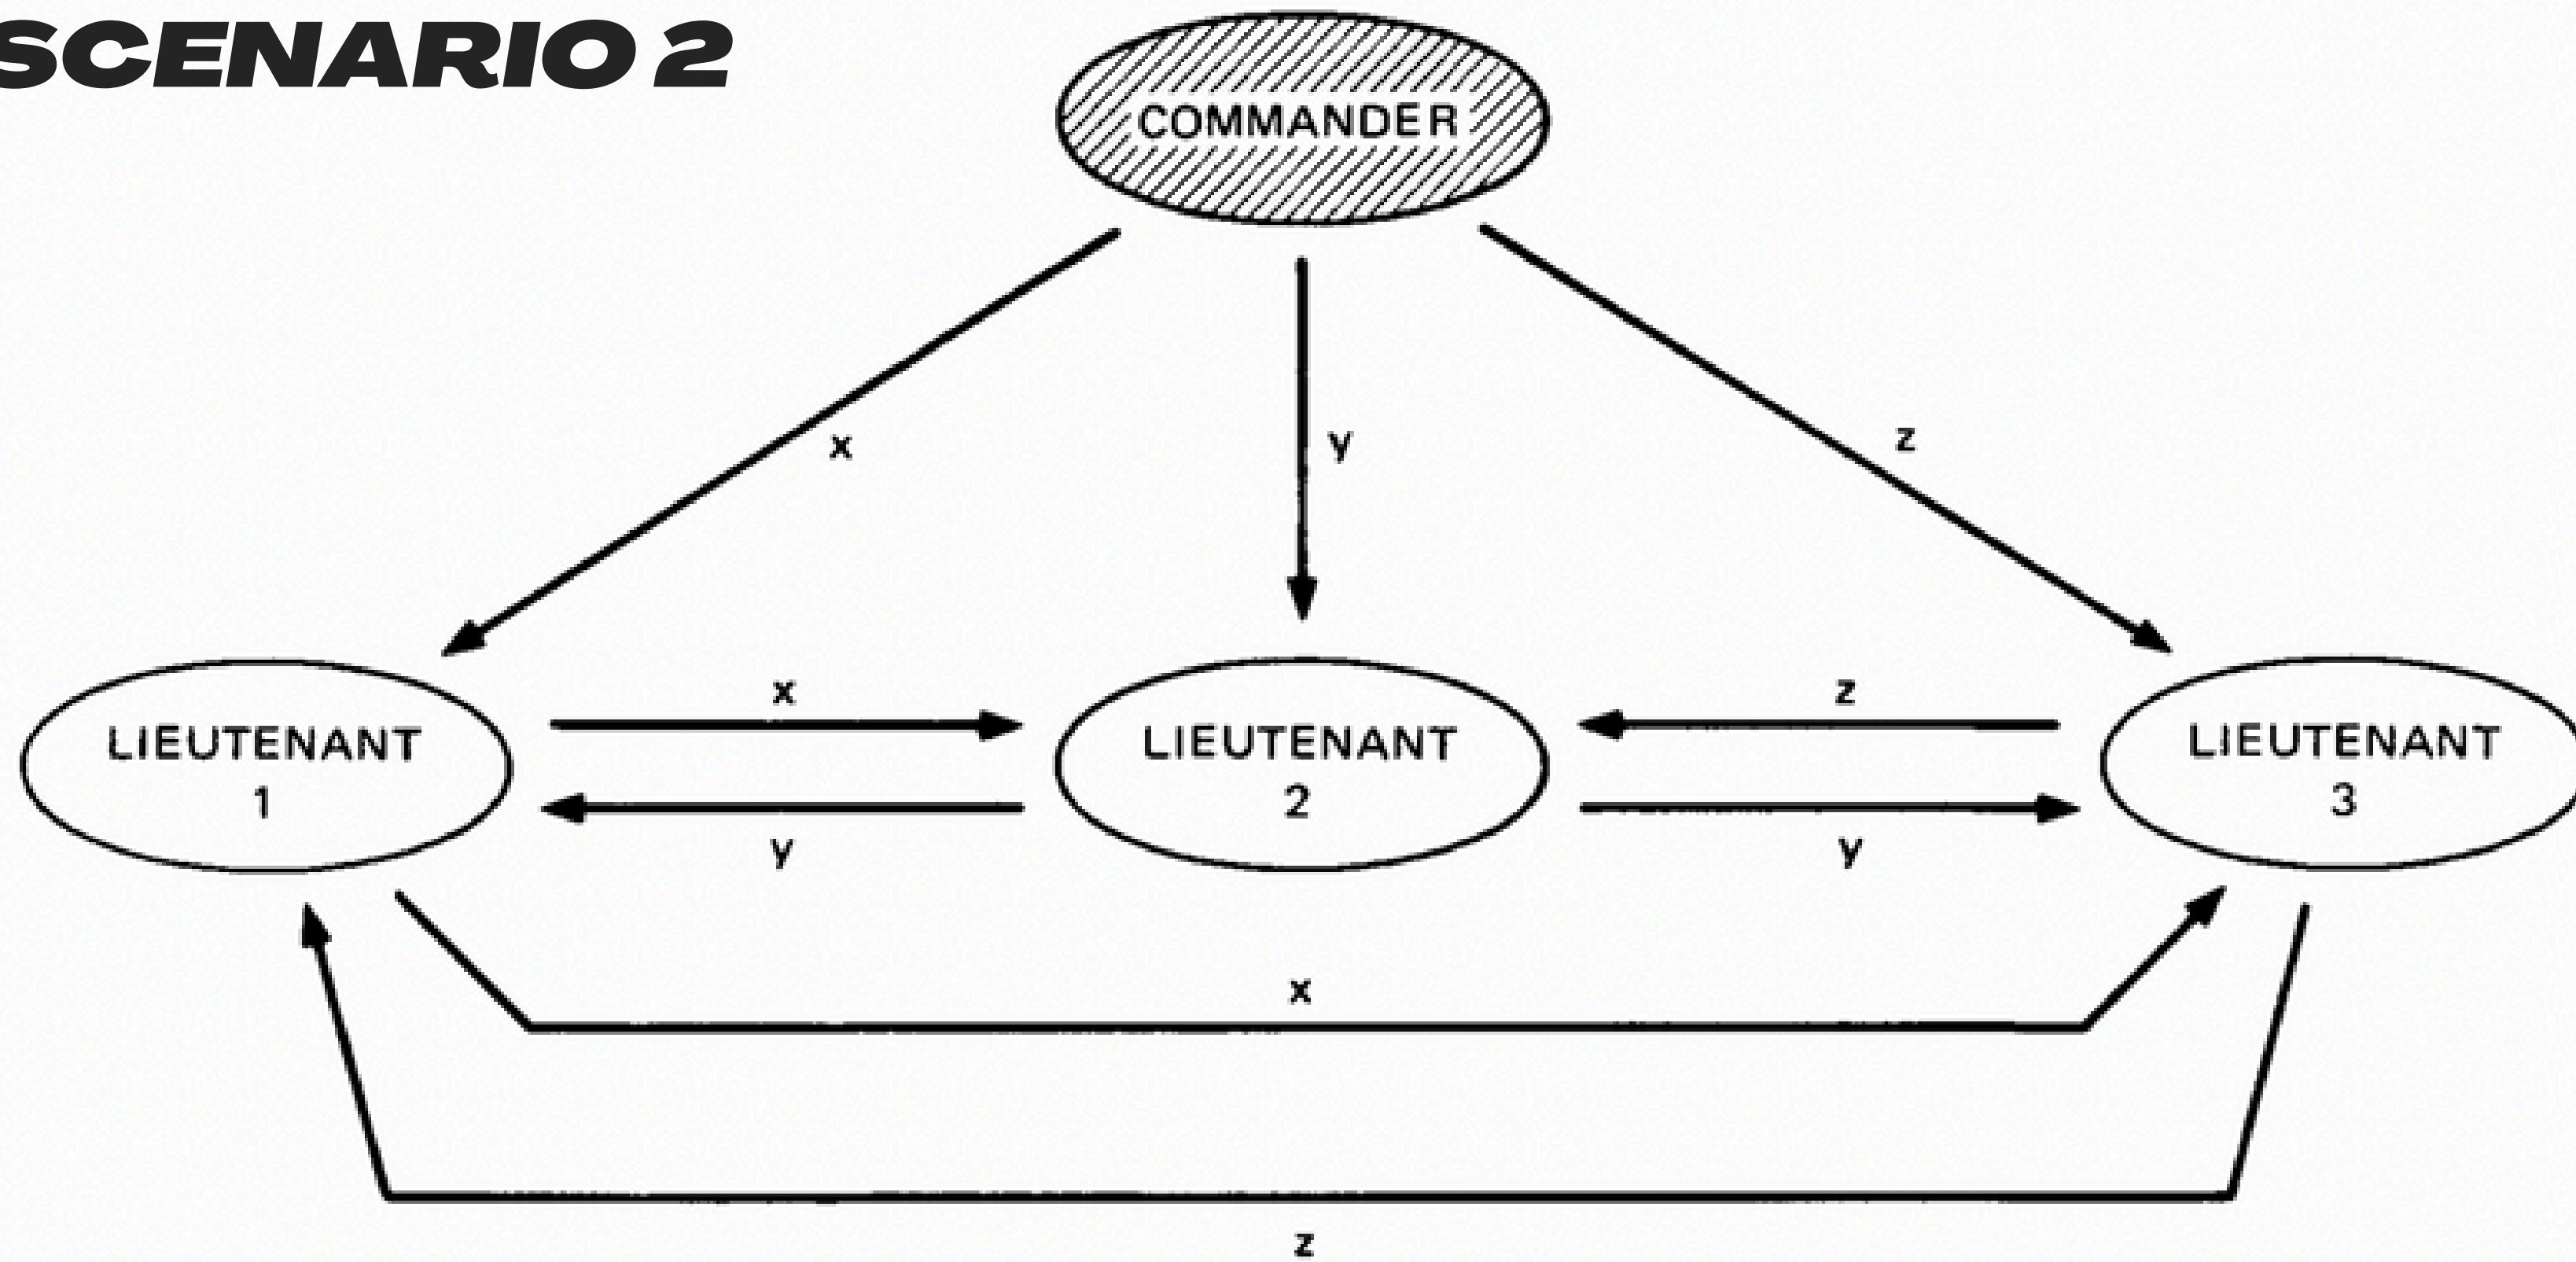
# SOLUTION WITH ORAL MESSAGES

## SCENARIO 2



# SOLUTION WITH ORAL MESSAGES

## SCENARIO 2



# SOLUTION WITH ORAL MESSAGES

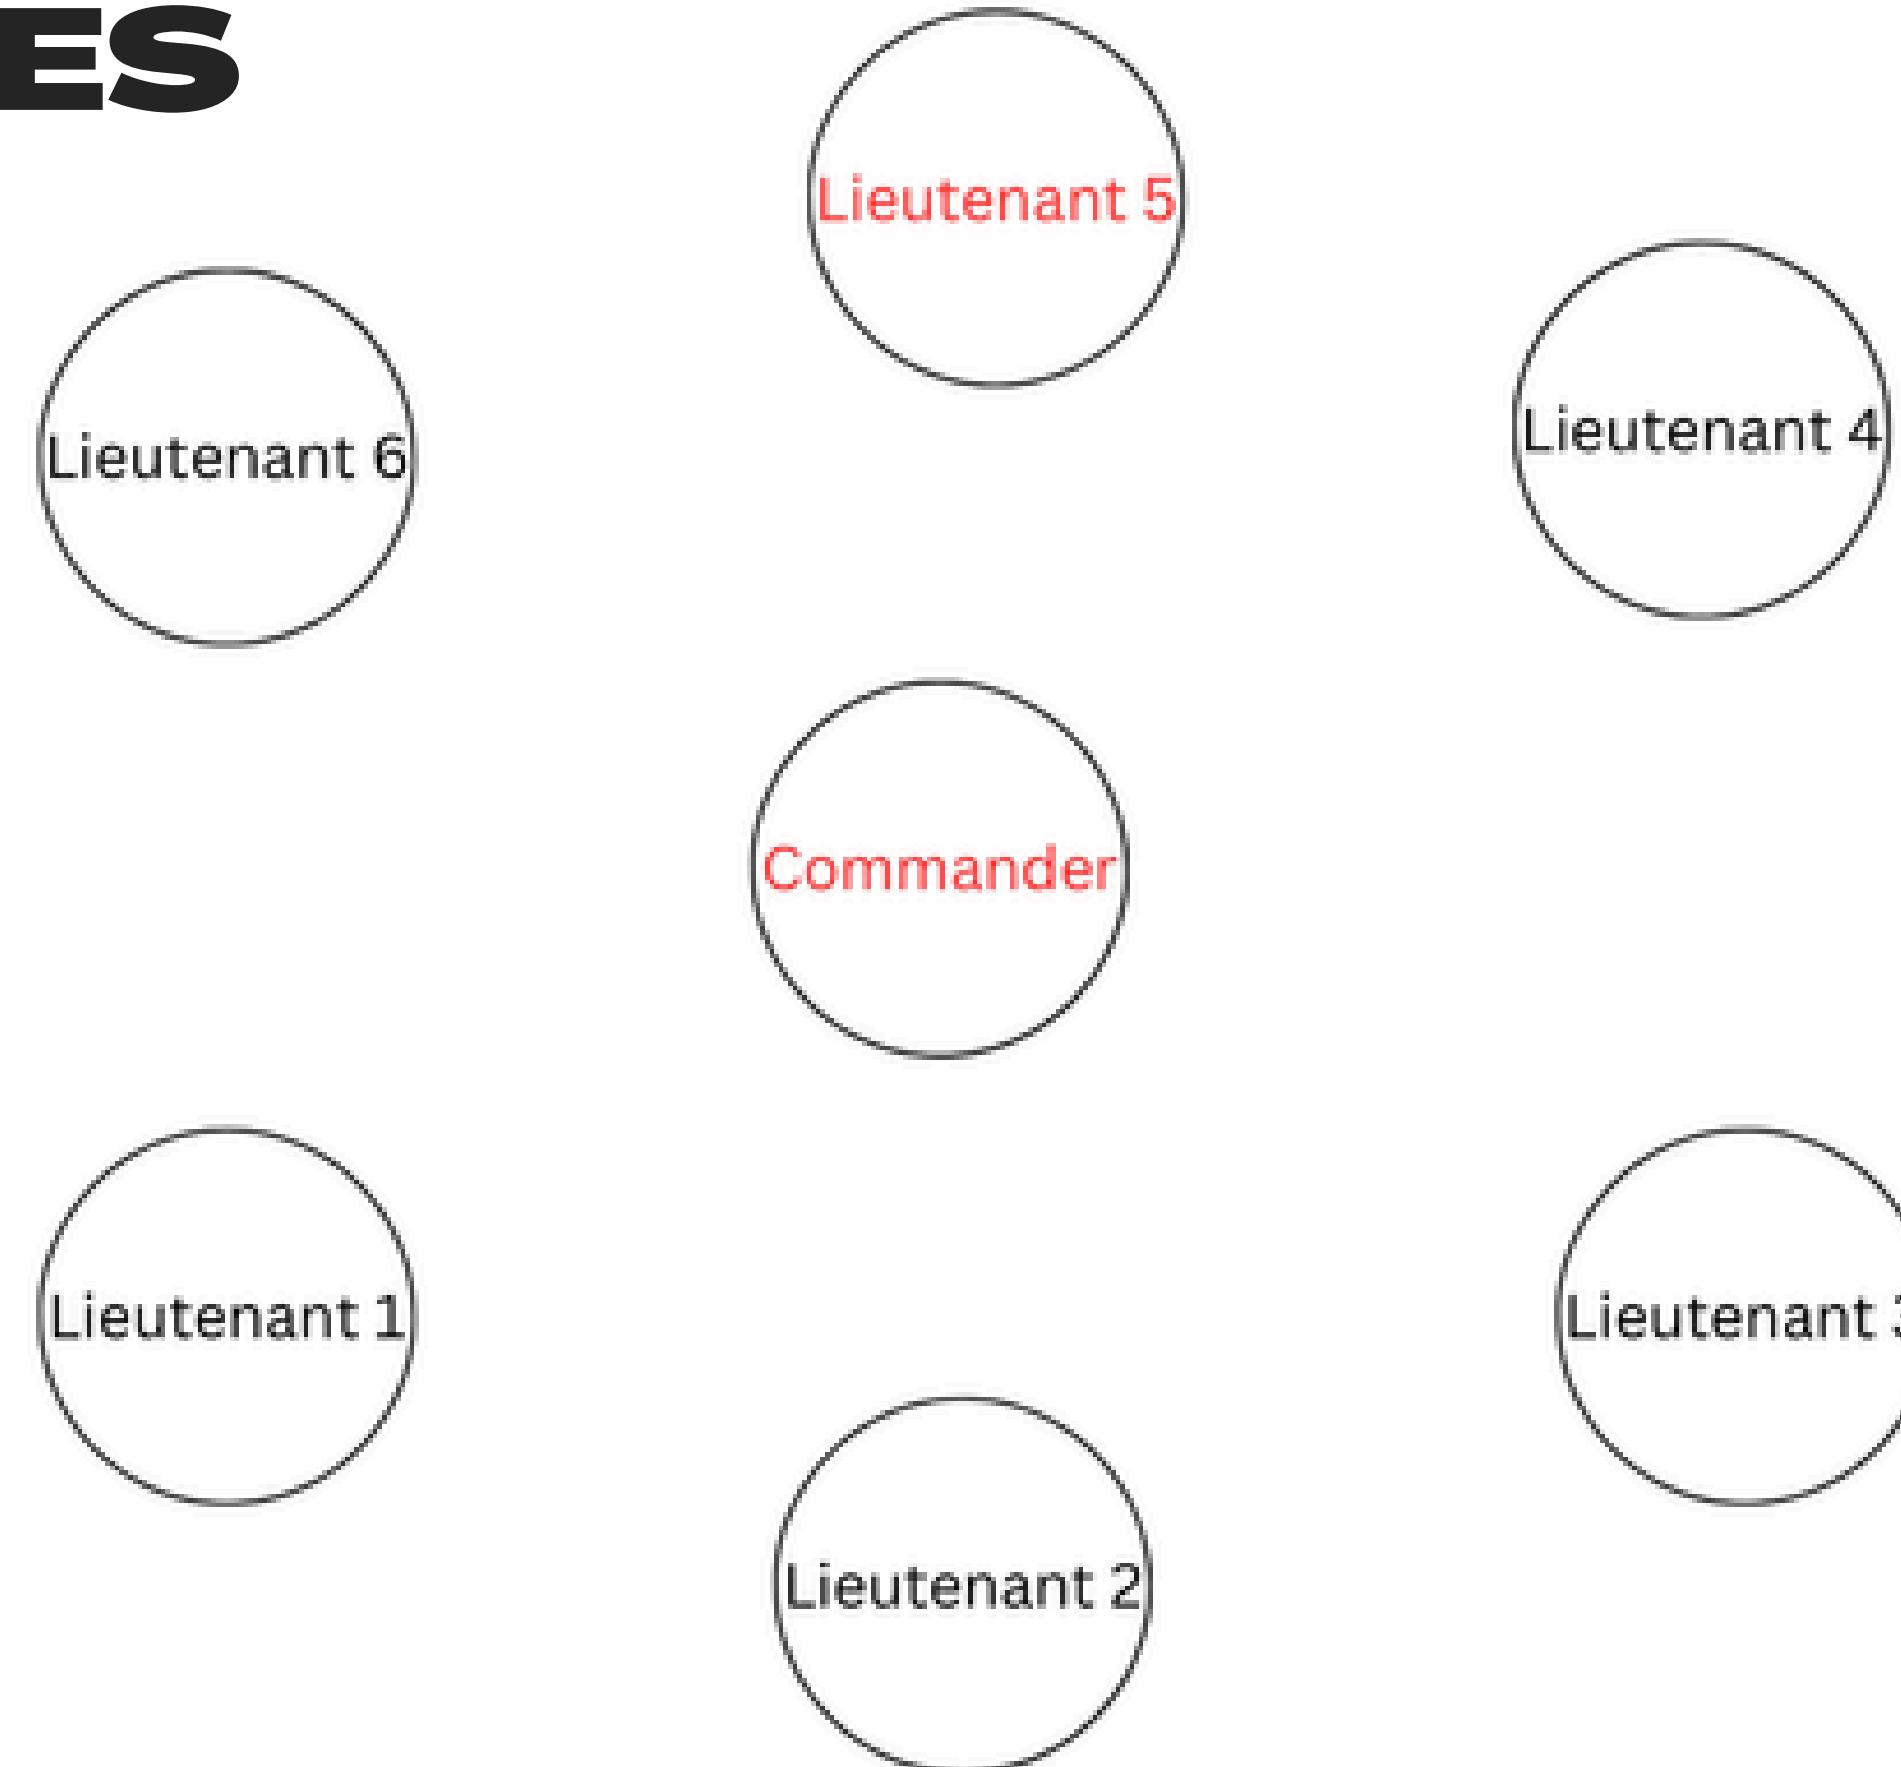
## THEOREM 1

For any  $m$ , Algorithm  $OM(m)$  satisfies conditions IC1 and IC2 if more than  $3m$  generals and at most  $m$  traitors.

Proof by induction:

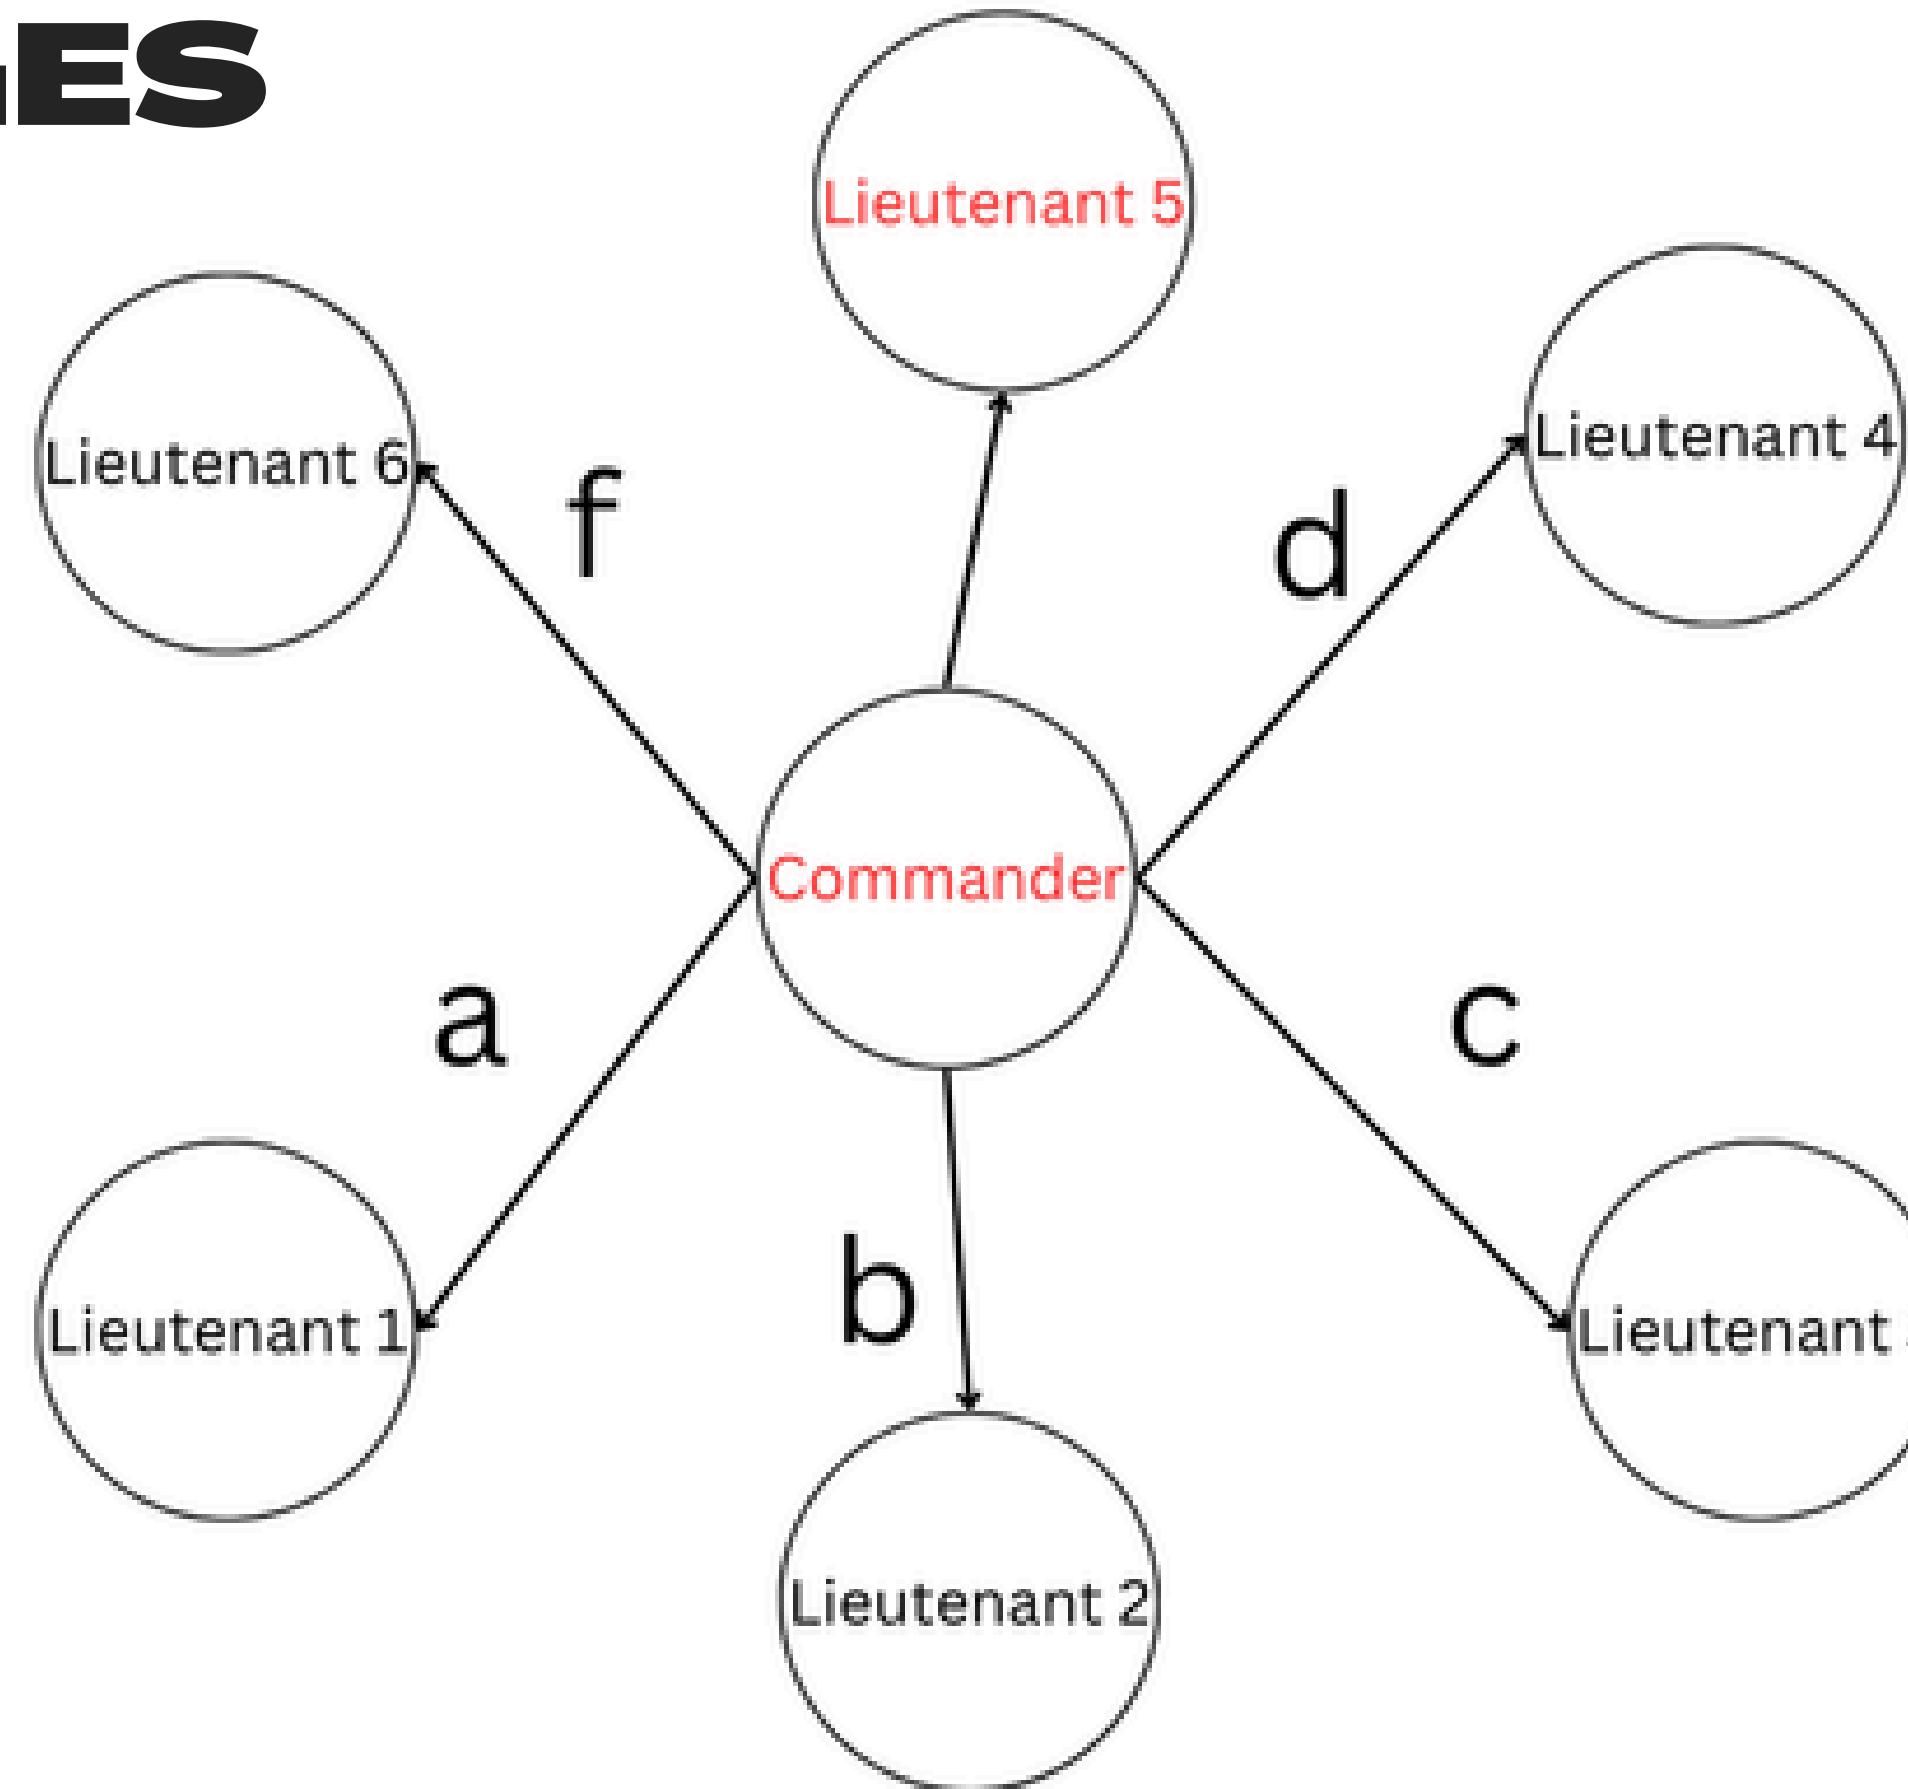
- $OM(0)$  obviously satisfies IC1 and IC2
- $OM(m)$  satisfies IC2 (majority of lieutenants will be loyal every level down until  $m=0$ )
- IC1 obviously follows IC2 if commander is loyal, so just need to prove IC1 if commander isn't loyal
- Using induction, can assume  $OM(m-1)$  true, meaning every loyal lieutenant gets same vector of values, proving IC1

# SOLUTION WITH ORAL MESSAGES



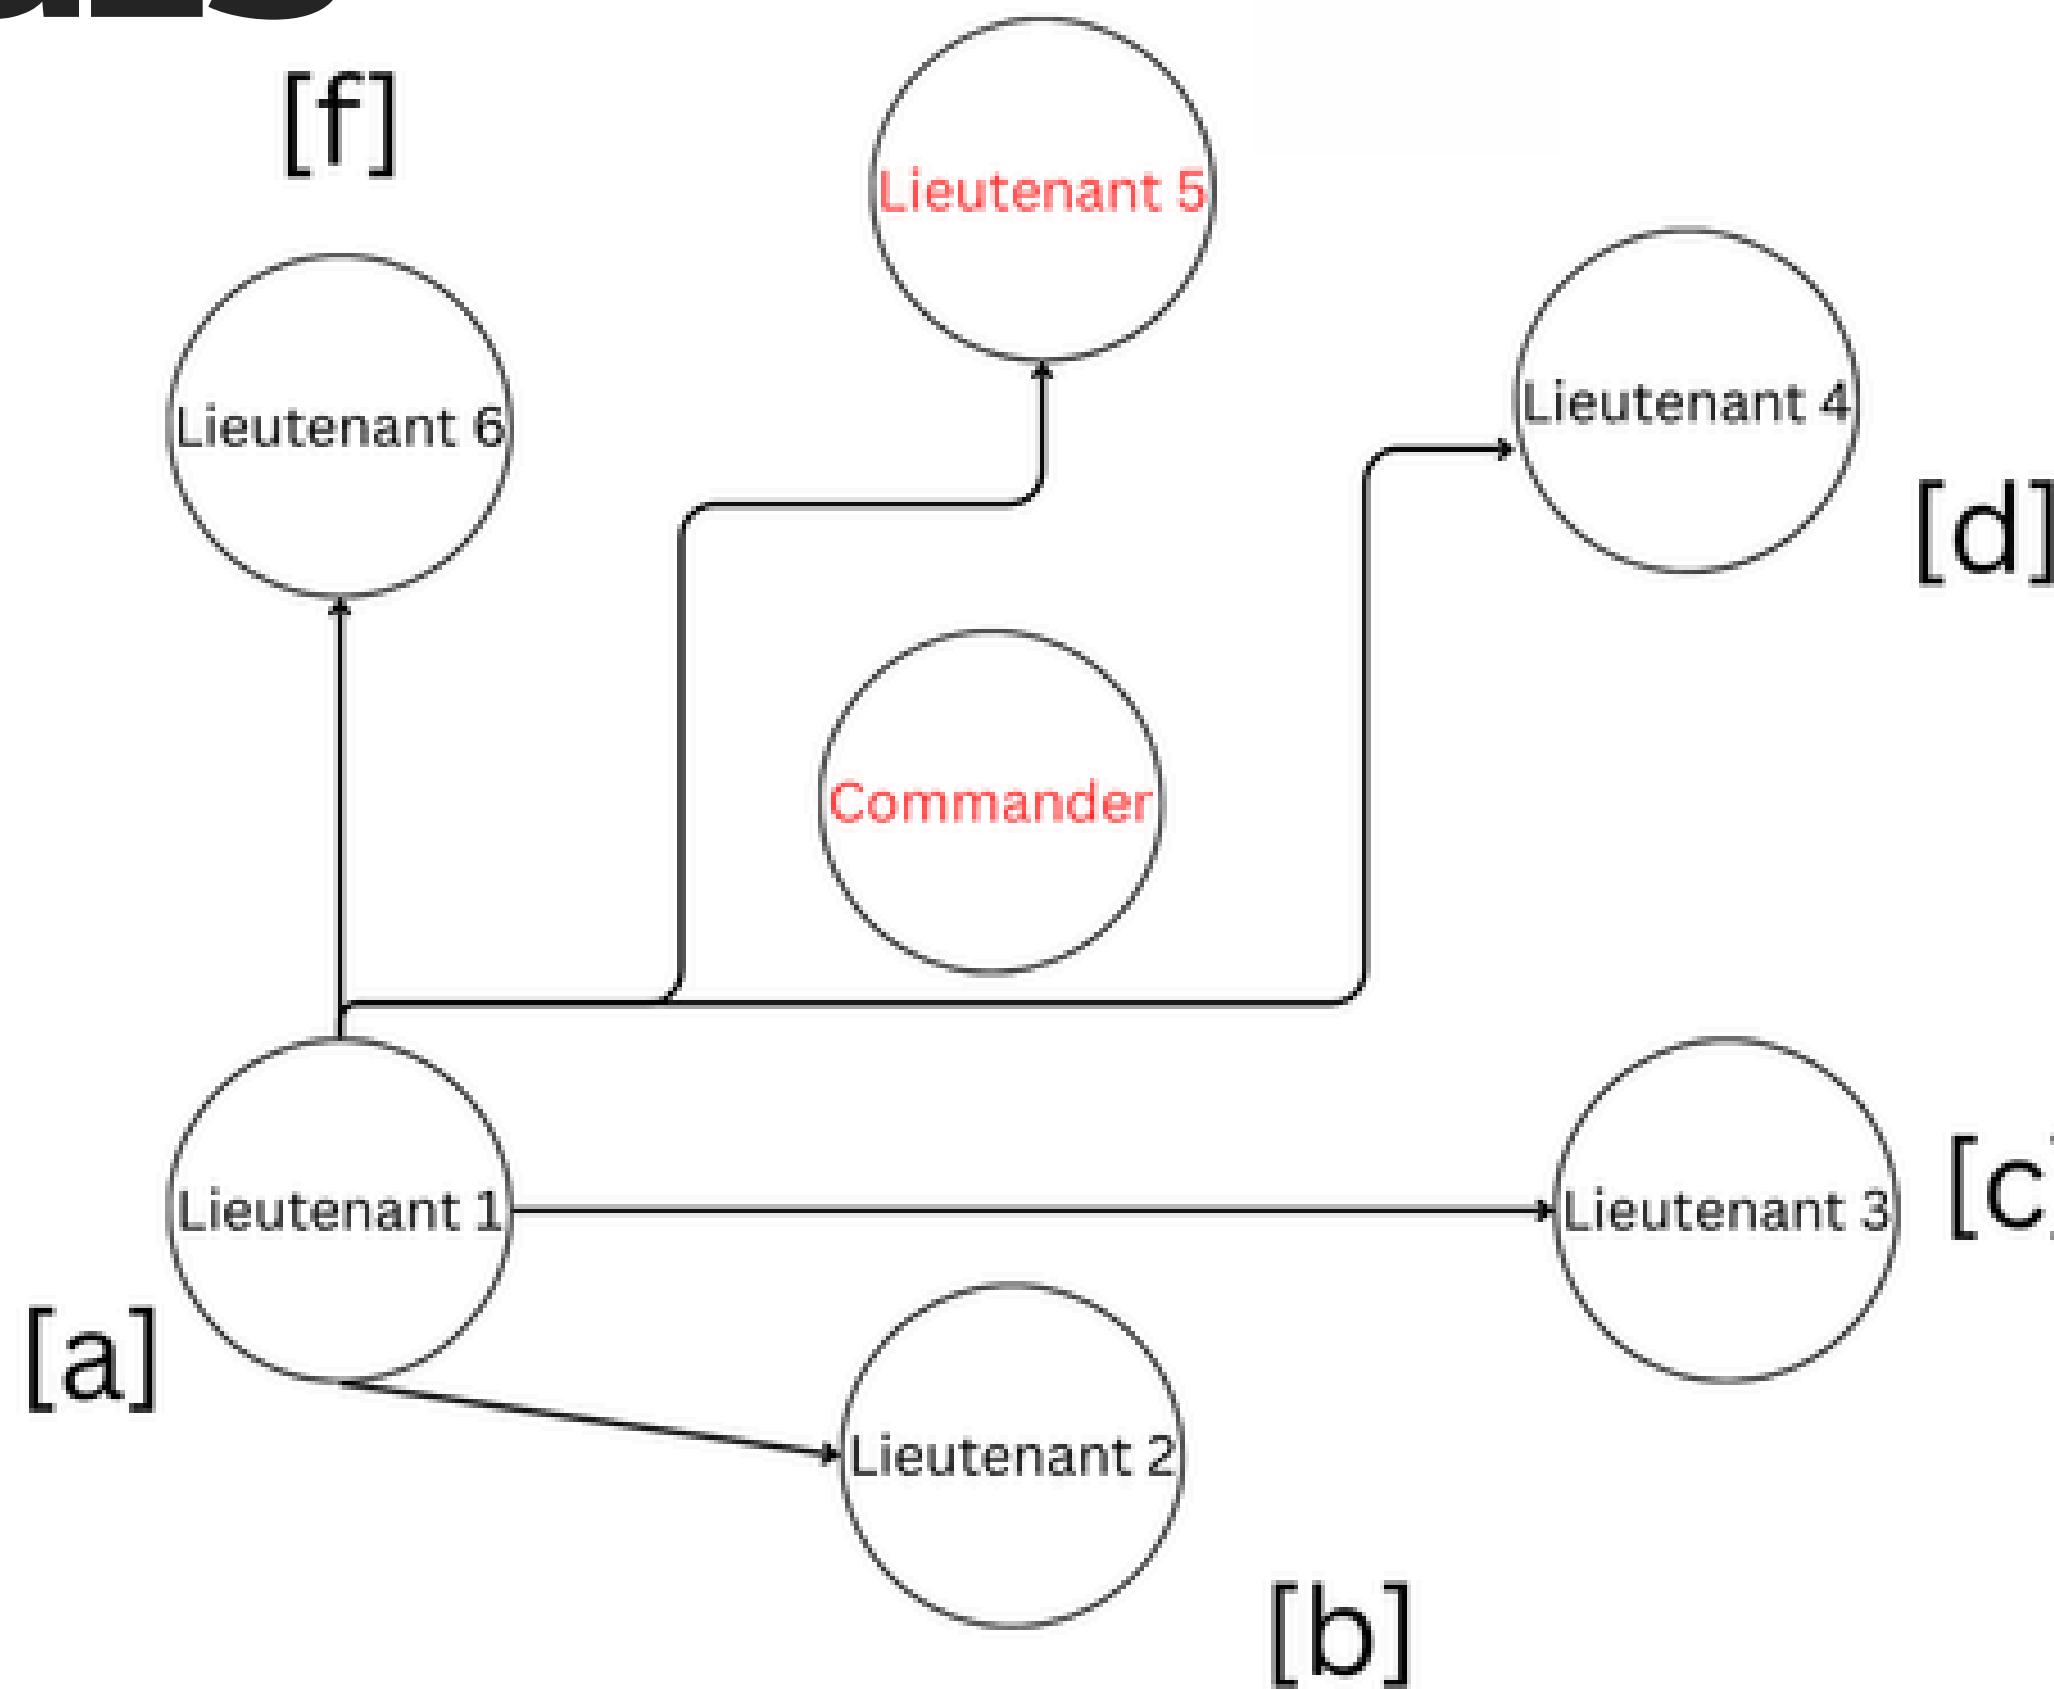
# SOLUTION WITH ORAL MESSAGES

**OM(2)**



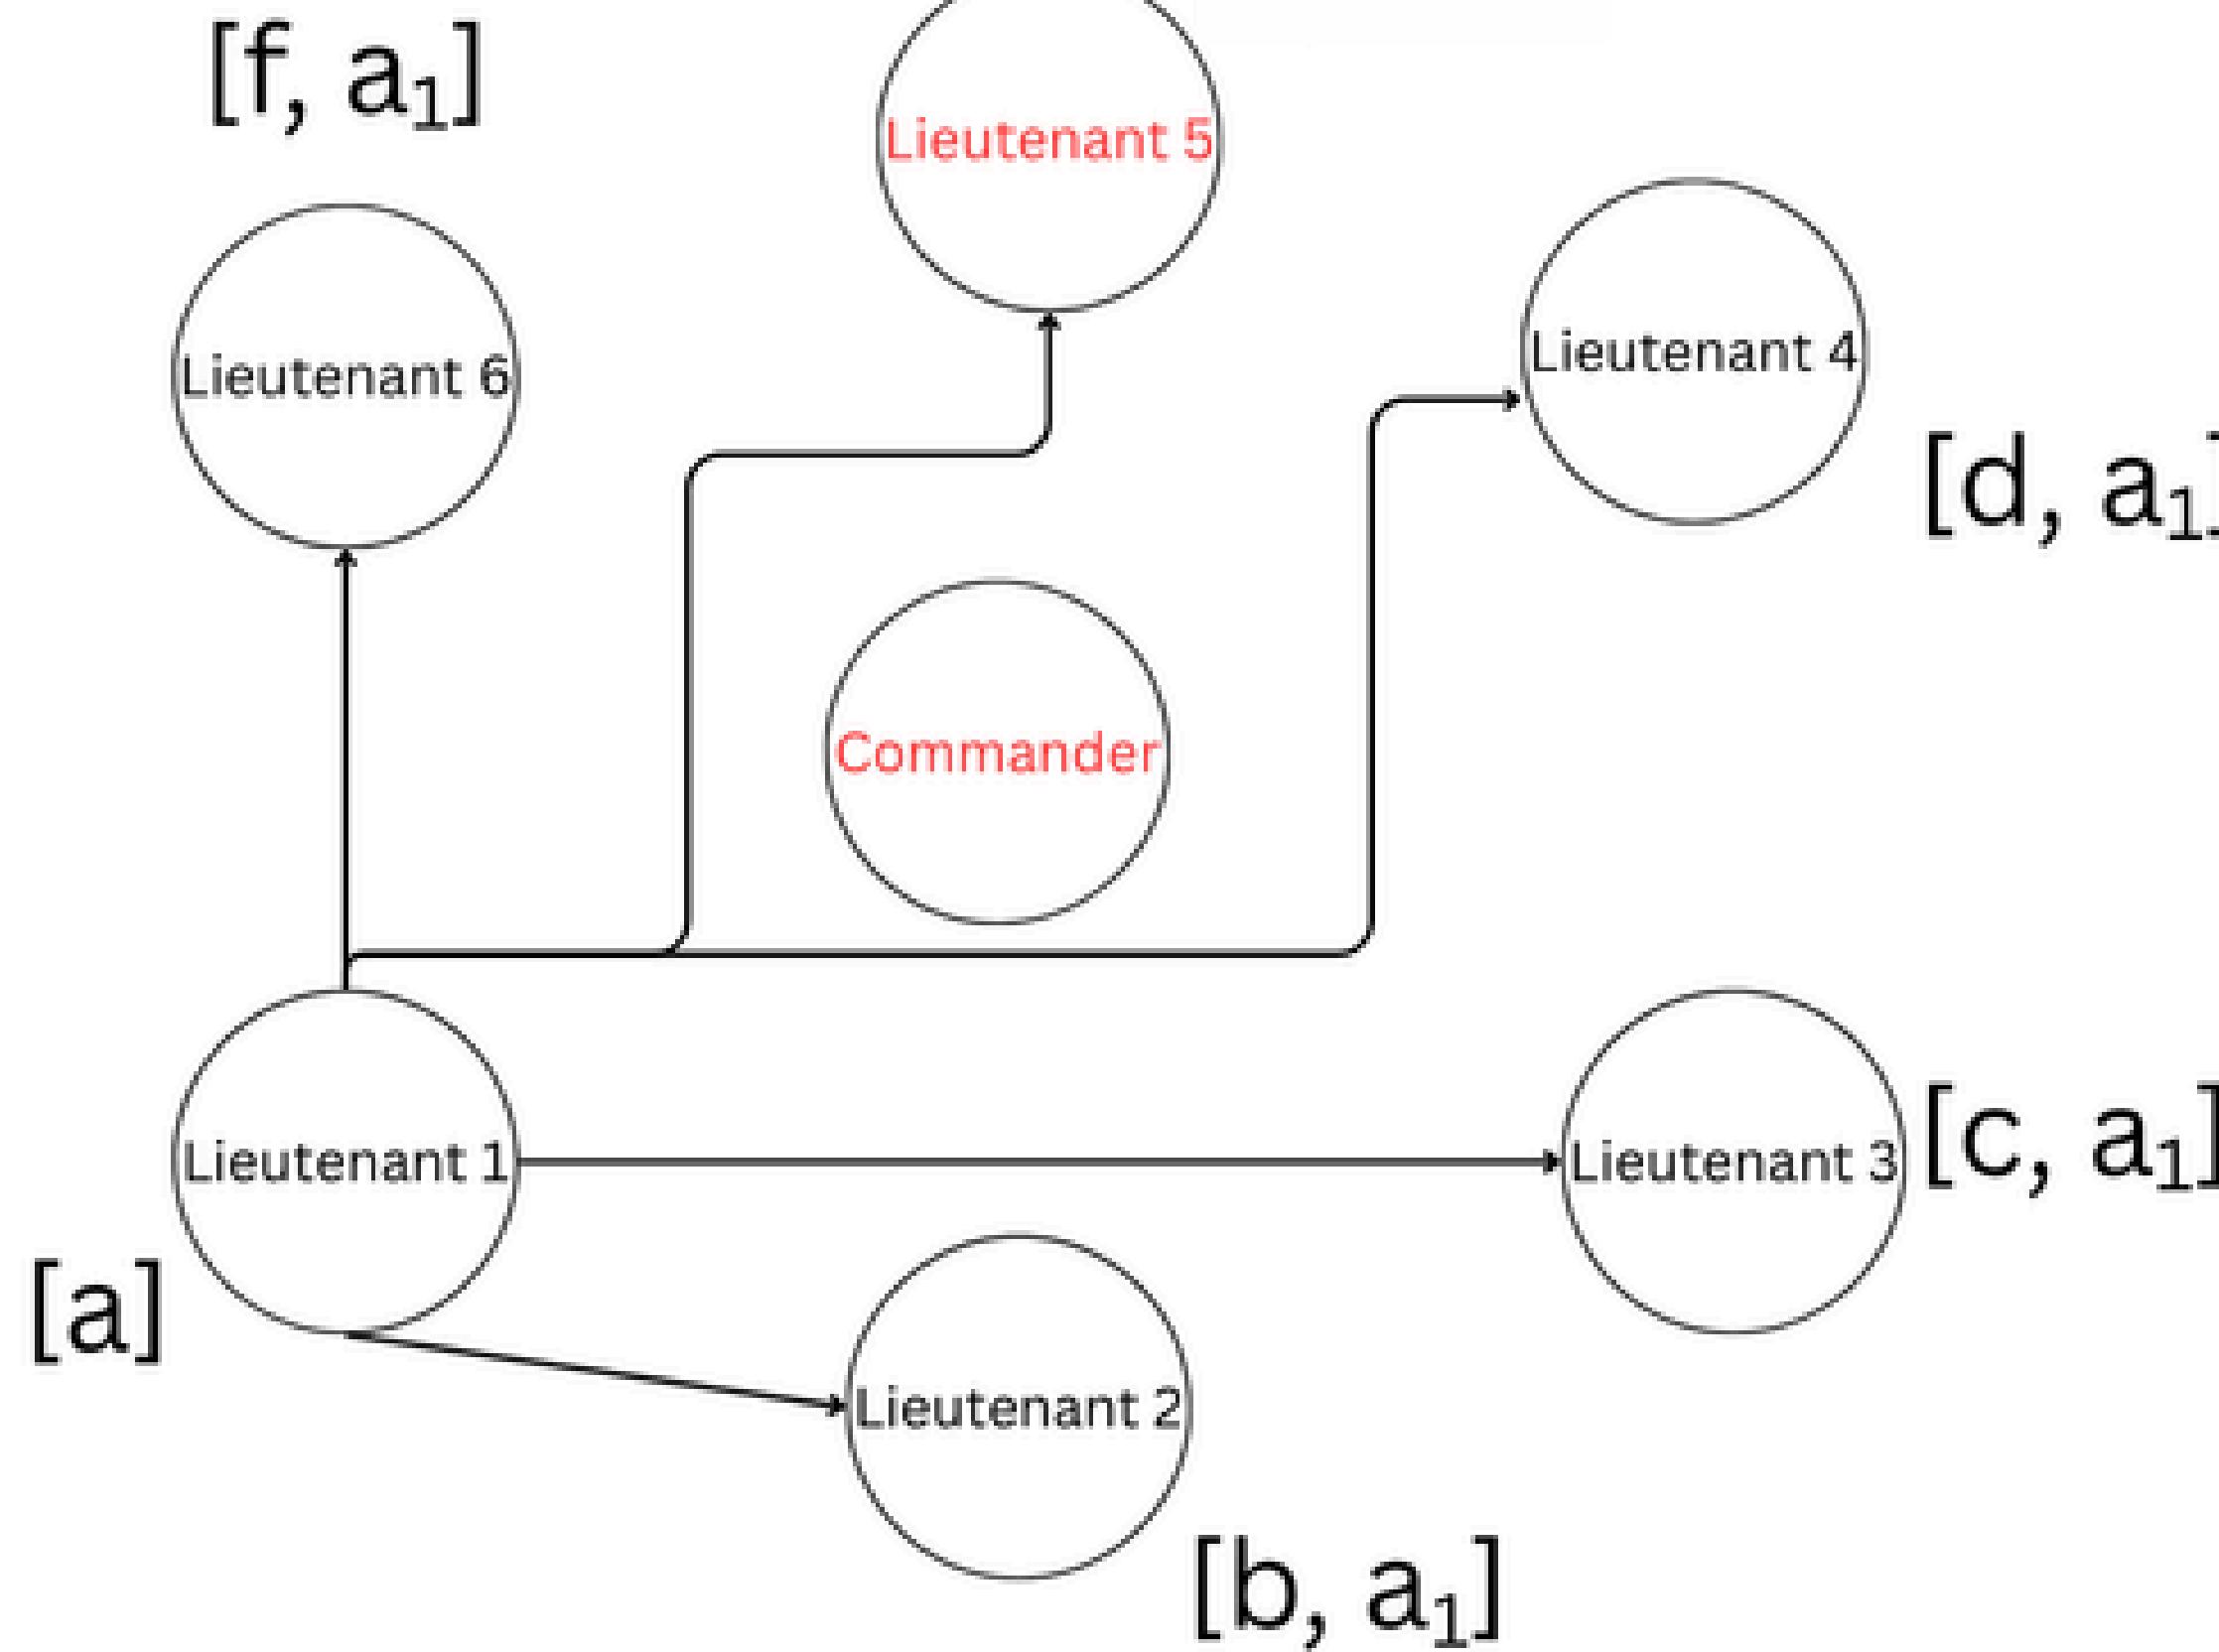
# SOLUTION WITH ORAL MESSAGES

**OM(1)**



# SOLUTION WITH ORAL MESSAGES

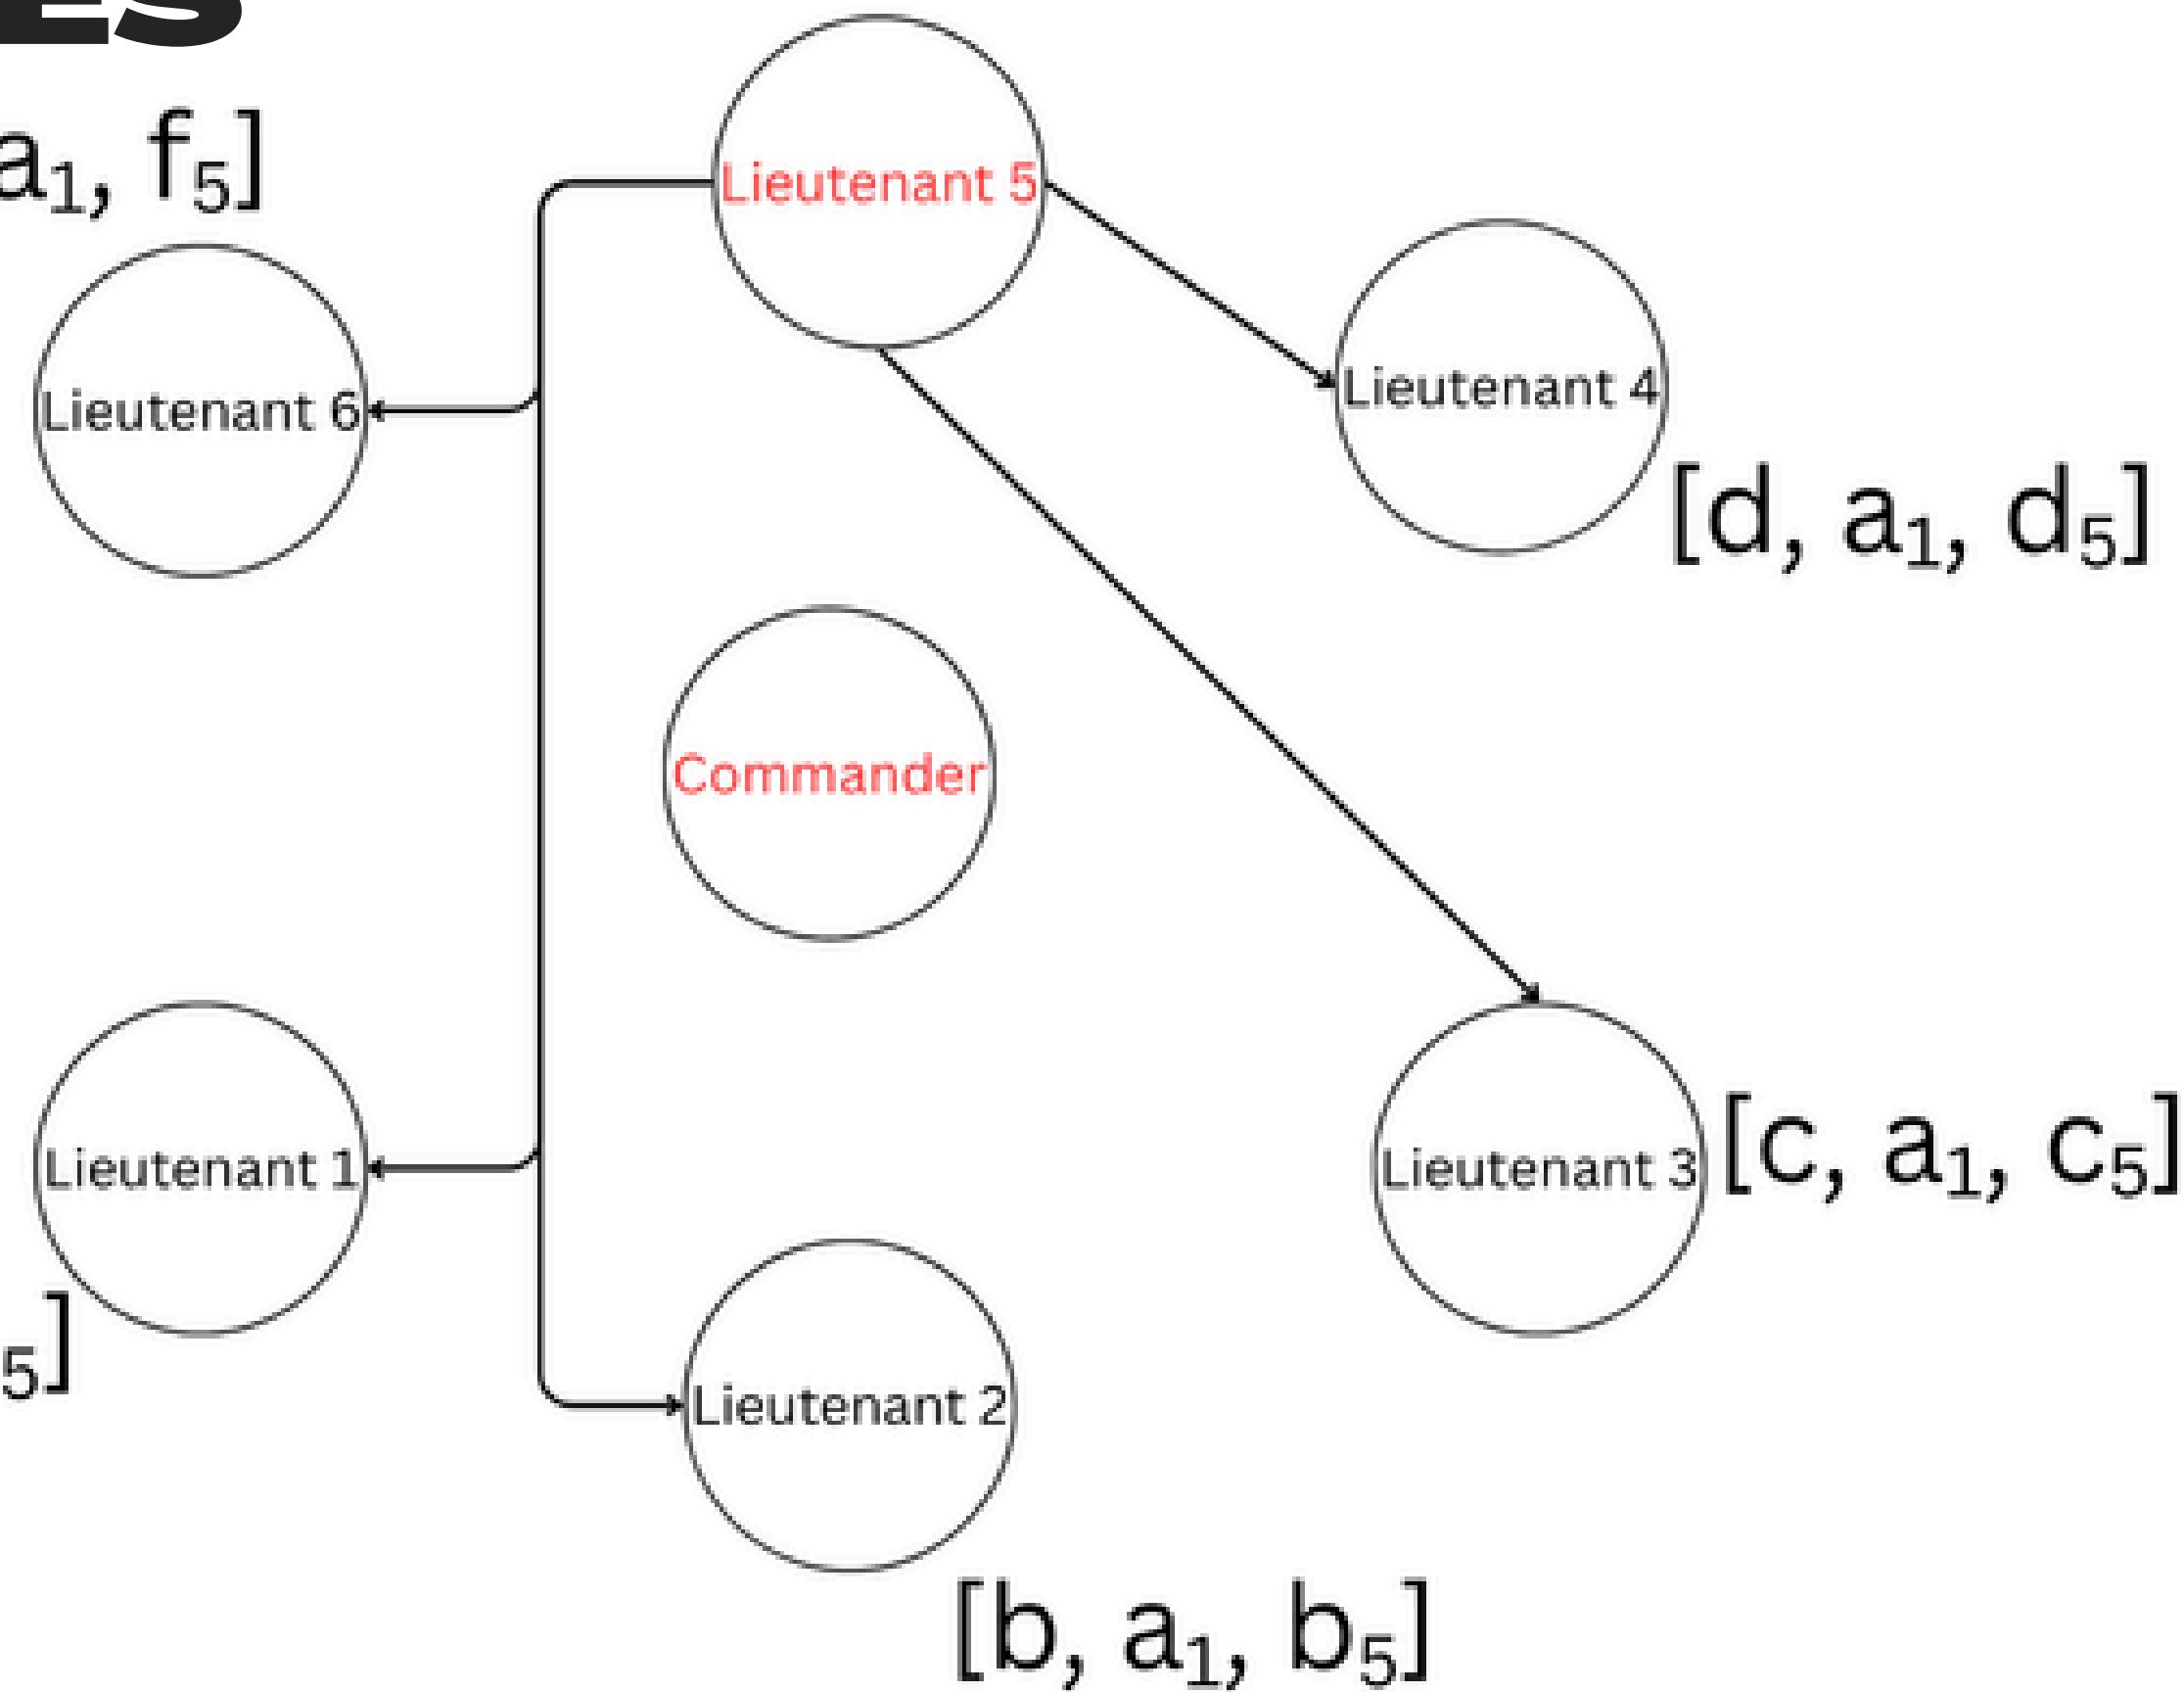
**OM(1)**



# SOLUTION WITH ORAL MESSAGES

**OM(1)**

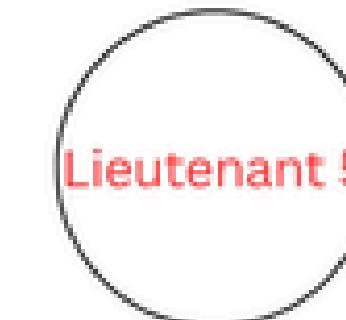
[f, a<sub>1</sub>, f<sub>5</sub>]



# SOLUTION WITH ORAL MESSAGES

**OM(1)**

[f, a<sub>1</sub>, b<sub>2</sub>, c<sub>3</sub>, d<sub>4</sub>, f<sub>5</sub>]



[d, a<sub>1</sub>, b<sub>2</sub>, c<sub>3</sub>, f<sub>6</sub>, d<sub>5</sub>]



[c, b<sub>2</sub>, d<sub>4</sub>, f<sub>6</sub>, a<sub>1</sub>, c<sub>5</sub>]

[a, b<sub>2</sub>, c<sub>3</sub>, d<sub>4</sub>, f<sub>6</sub>, a<sub>5</sub>]

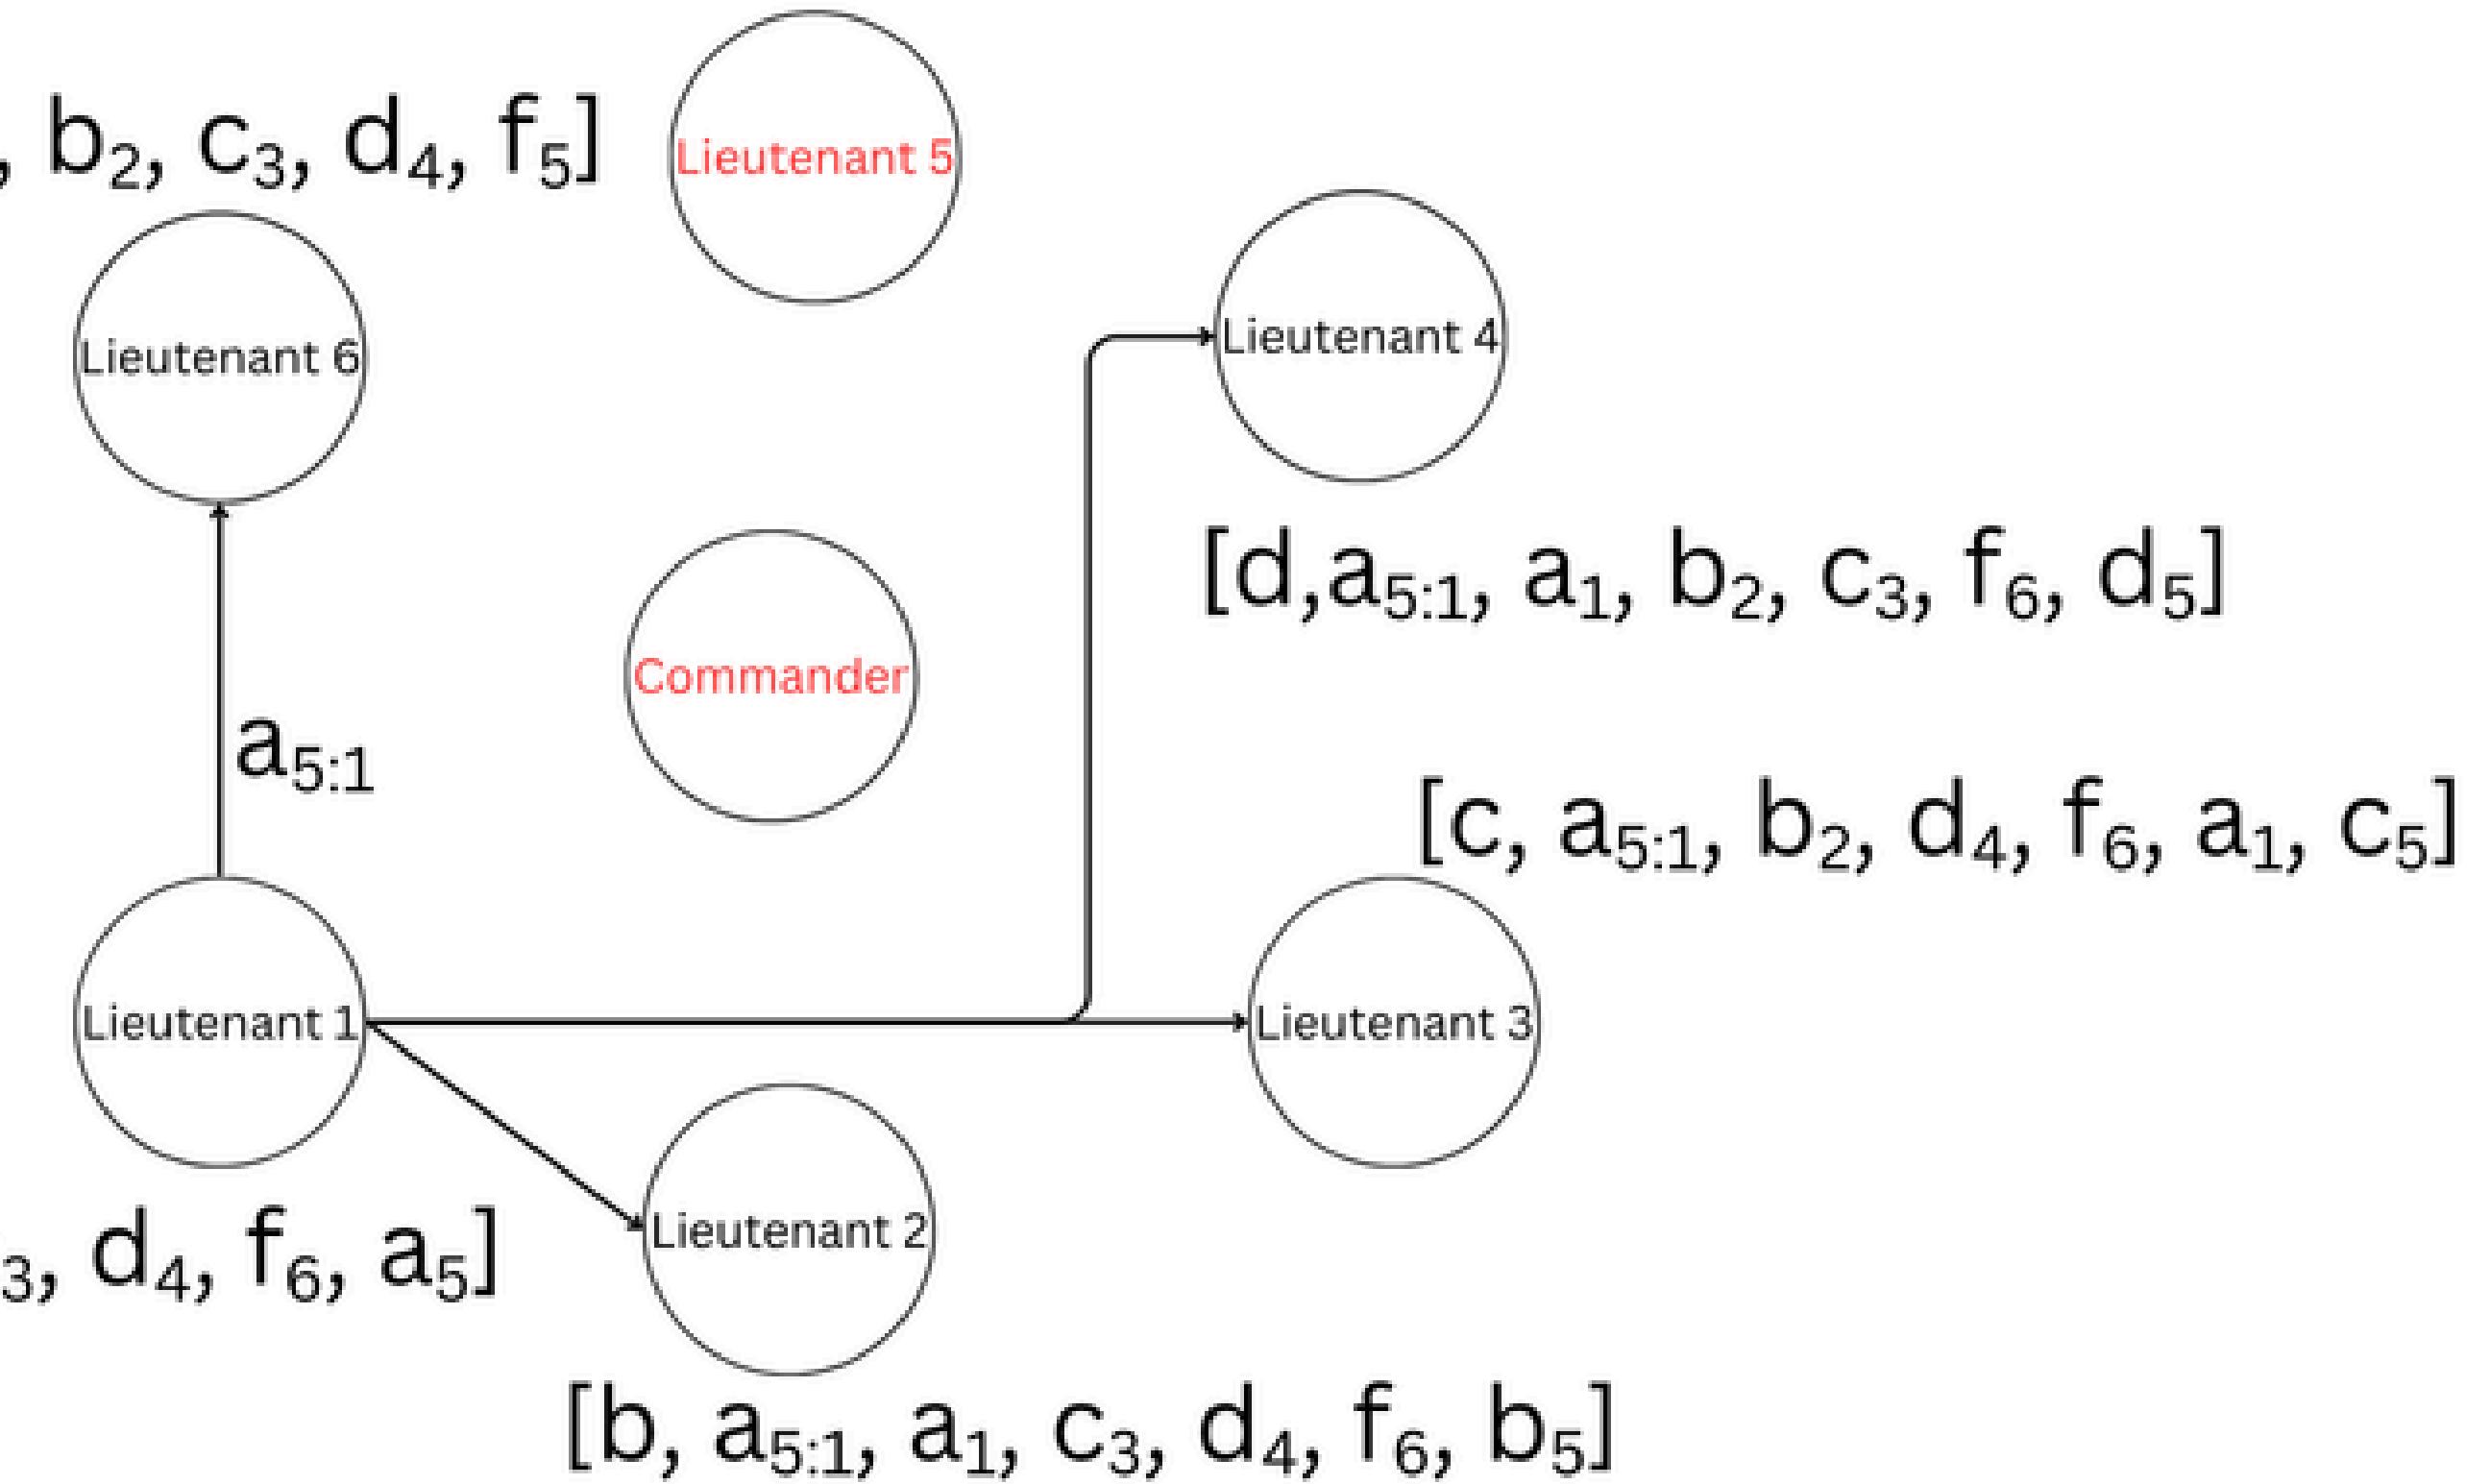


[b, a<sub>1</sub>, c<sub>3</sub>, d<sub>4</sub>, f<sub>6</sub>, b<sub>5</sub>]

# SOLUTION WITH ORAL MESSAGES

**OM(O)**

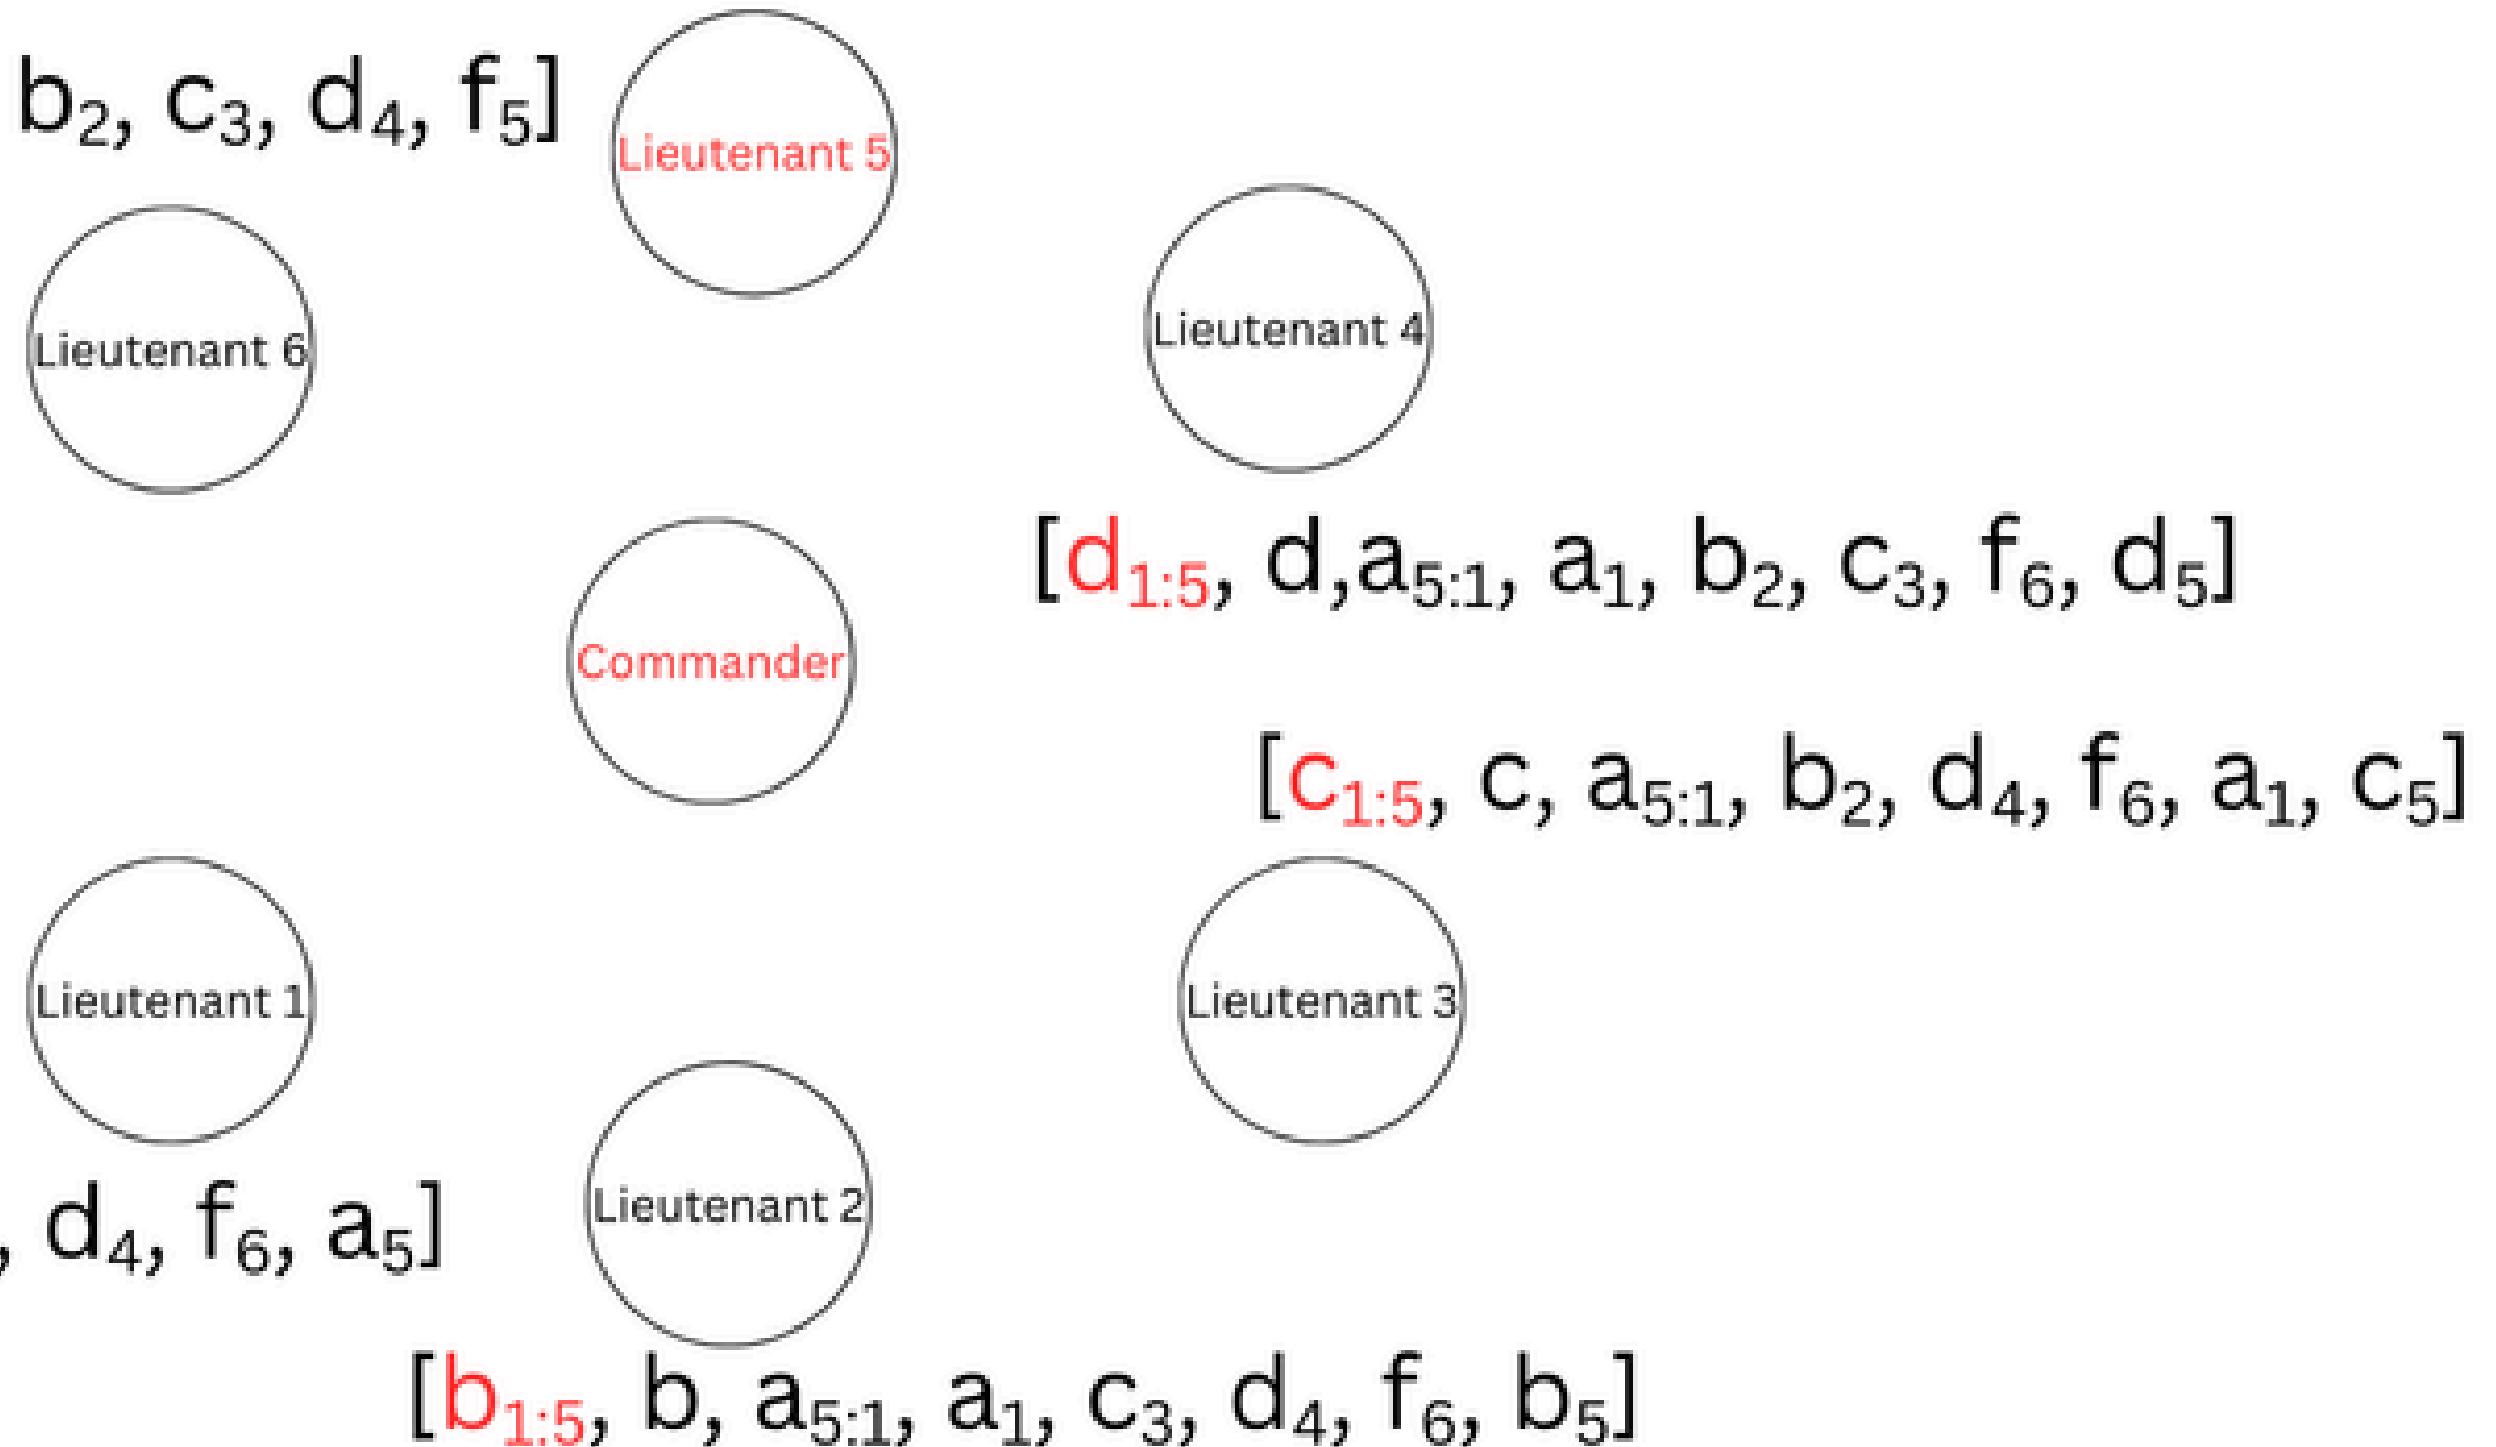
[ $a_{5:1}, f, a_1, b_2, c_3, d_4, f_5$ ]



# SOLUTION WITH ORAL MESSAGES

**OM(O)**

[**f<sub>1:5</sub>**, f, a<sub>1</sub>, a<sub>5:1</sub>, b<sub>2</sub>, c<sub>3</sub>, d<sub>4</sub>, f<sub>5</sub>]



# SOLUTION WITH ORAL MESSAGES

**OM(O)**

Lieutenant 1	Lieutenant 2	Lieutenant 3	Lieutenant 4	Lieutenant 6
5a, 5b, 5c, 5d, 5f	5a, <b>5b</b> , 5c, 5d, 5f	5a, 5b, <b>5c</b> , 5d, 5f	5a, 5b, 5c, <b>5d</b> , 5f	5a, 5b, 5c, 5d, <b>5f</b>

# SOLUTION WITH ORAL MESSAGES

## ***PROBLEM RECAP: OM(M) LIMITATIONS***

- $OM(m)$  works only if  $n \geq 3m + 1$  (needs many generals for few traitors).
- Traitors can send conflicting orders to different lieutenants.
- No authentication → messages can be forged or modified.
- Works only under strong assumptions like complete connectivity (everyone talks to everyone).

# SOLUTION WITH SIGNED MESSAGES

## MAIN IDEA:

*“The traitors’ ability to lie makes the Byzantine Problem difficult”*

To ease the problem, the lying ability of traitors need to be restricted.

### A4: Assumptions of Unforgeable Signatures:

- (a) A loyal general’s signature cannot be forged and message tampering is detectable.
- (b) Anyone can verify the authenticity of a signature.

Traitors’ signatures can be forged by other traitors → allows collusion.

# SOLUTION WITH SIGNED MESSAGES

## KEY CONCEPTS

What Each Lieutenant Maintains

Concept	Meaning
$x : i$ <b>Attack:0:1</b>	“x” signed by General $i$ $0 \rightarrow 1 \rightarrow 2$
$V_i$	Set of properly signed orders received so far by Lieutenant $i$
<b>Choice (<math>V_i</math>)</b>	decision function applied by lieutenant $i$ , used to decide the final action.

# SOLUTION WITH SIGNED MESSAGES

## RULES OF HANDLING MESSAGES:

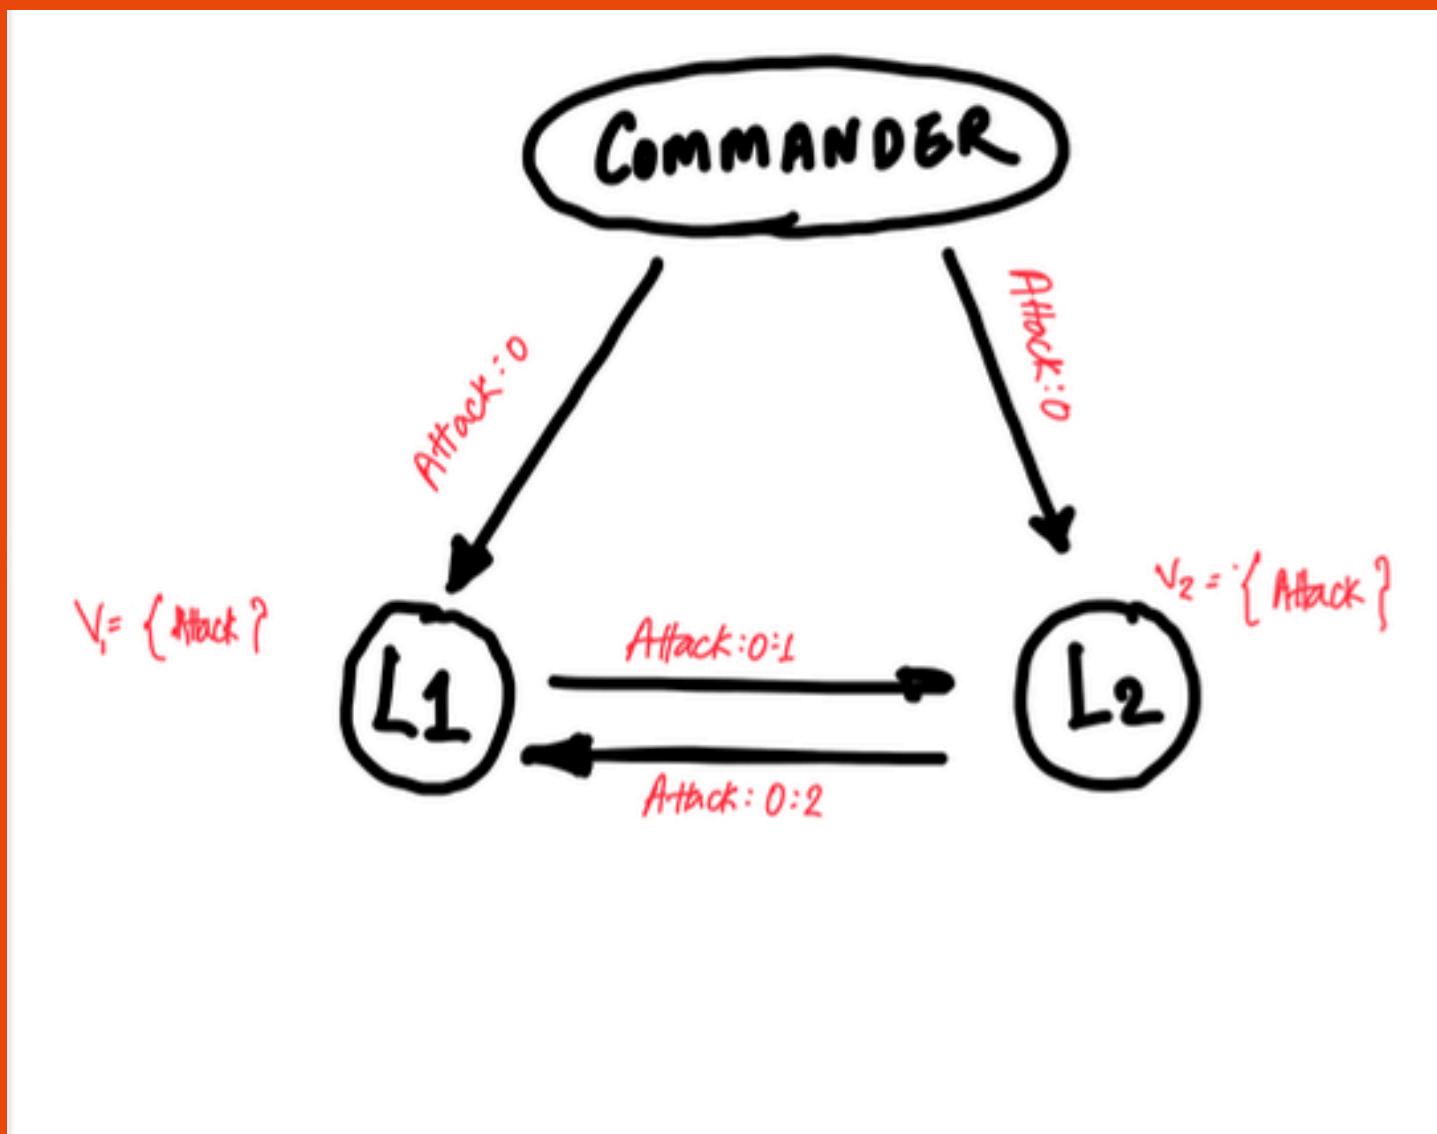
Case	Example Message	What Happens	Explanation
✓ Keep unique, valid messages	Attack:0:1	✓ Keep	Message has valid signatures and hasn't been seen before.
✗ Duplicate message	Two copies of Attack:0:1	✗ Ignore	Already received the same signed message; no need to store twice.
✗ Invalid signature	Attack:1 (forged or tampered)	✗ Discard	Signature check fails → message not trustworthy.
✗ Same order already known	Retreat:0 already in $V_3$ , then Retreat:0:2 arrives	✗ Ignore	Order "Retreat" already exists; no need to store forwarded copies.
✓ Enough signatures → stop forwarding	Attack:0:2 in $SM(1)$	✓ Stop sending further	Message already verified by commander + one lieutenant ( $m+1$ signatures reached).

- **Ignore duplicates or improperly signed messages:**
- **Ignore any message where the order already exists in  $V_i$ :**
- **If an order is signed by  $m$  lieutenants ( $SM(m)$ ), stop forwarding (enough verification)**

# SOLUTION WITH SIGNED MESSAGES

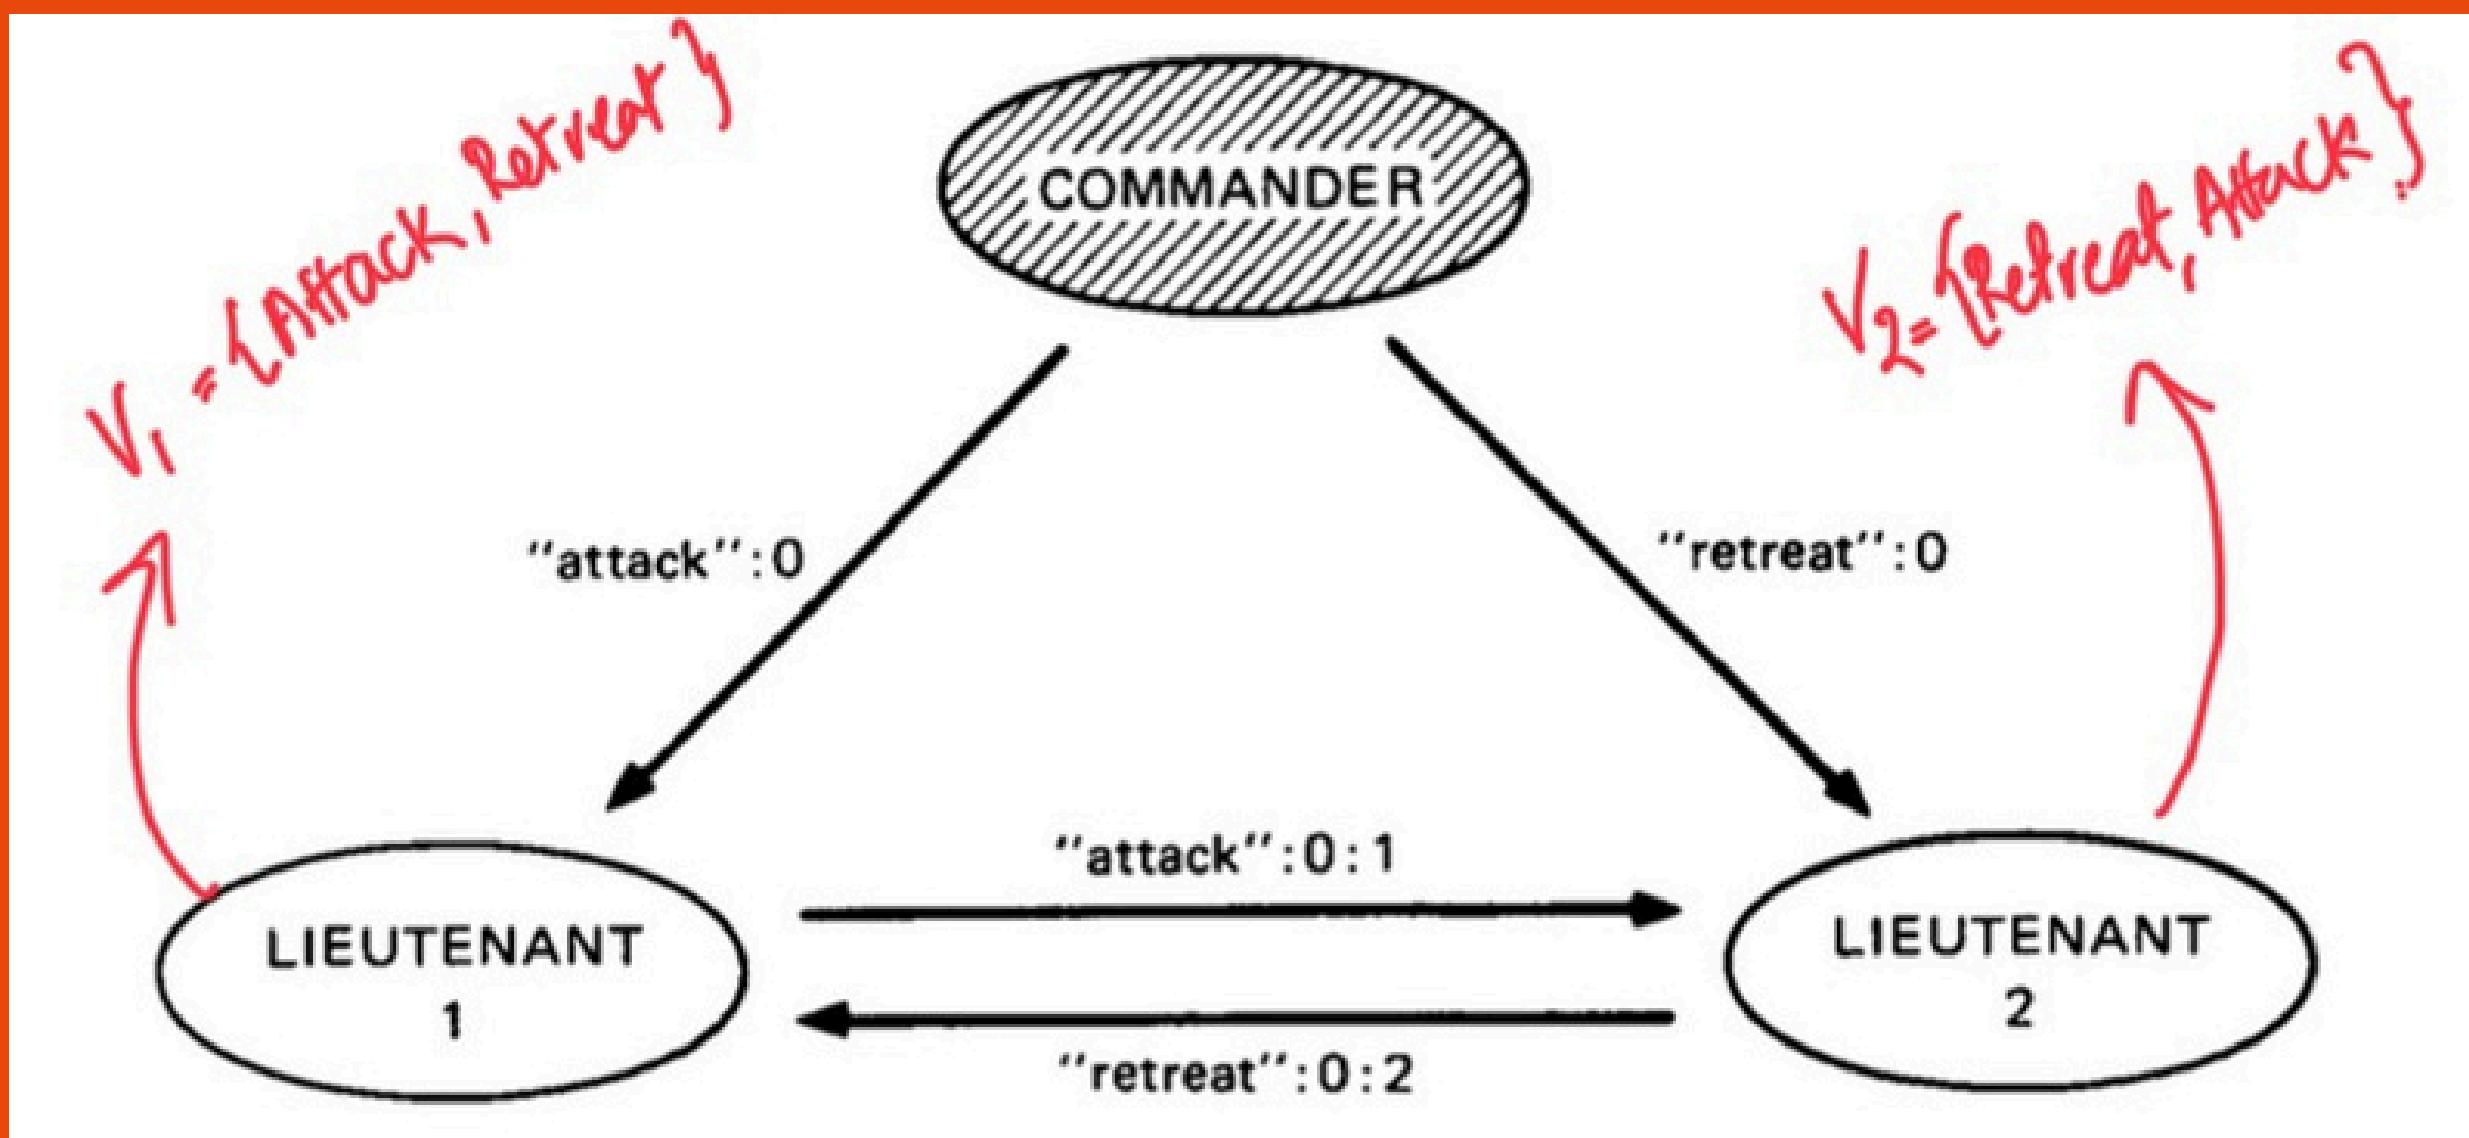
## LIEUTENANTS IN ACTION IN SM(M)

1. Commander signs & sends his order (e.g., “Attack:0”) to all lieutenants.
2. Each lieutenant who receives a new signed order:
  - Adds the order to his set  $V_i$ .
  - Adds his own signature → forwards it to all other lieutenants.
3. When no more messages are incoming, each lieutenant executes **choice( $V_i$ )** to decide final action.



# SOLUTION WITH SIGNED MESSAGES

**EXAMPLE: SM(1) WITH 3 GENERALS**  
(COMMANDER IS A TRAITOR)



# SOLUTION WITH SIGNED MESSAGES

## WHY SM(M) ALWAYS WORKS

Theorem 2:

For any  $m$ , Algorithm SM( $m$ ) solves the Byzantine Generals Problem if  $\leq m$  traitors.

Proof Sketch:

- IC2: (*If commander is loyal*)  $\rightarrow$  all loyal lieutenants receive and obey the same signed order.
- IC1: (*If commander is traitor*)  $\rightarrow$  all loyal lieutenants still receive the same set of signed orders, so they make the same final choice.
- Recursive signature propagation ensures every loyal lieutenant sees the same set of valid orders.

**THANK YOU!**

# RELIABLE SYSTEM

- Redundant processors compute the same output; majority voting selects the correct result.
- CONDITIONS FOR RELIABILITY:
  - All nonfaulty processors use the same input.
  - If the input unit is nonfaulty, all nonfaulty processors must use its value as input.
- Analogy to Byzantine Generals Problem:
  - Commander → Input device
  - Lieutenants → Processors
  - Loyal → Nonfaulty
  - Traitor → Faulty

# RELIABLE SYSTEM

- COMMUNICATION ASSUMPTIONS:
  - A1: Messages delivered correctly
  - A2: Sender can be identified
  - A3: Missing messages detectable (timeouts + synchronized clocks)
  - A4: Messages can be signed to prevent forgery
- Timeouts to work , two things are needed
  - A known maximum delay for message transmission
  - Synchronized clocks between sender and receiver
- If P1's and P2's clocks are not perfectly synchronized:
  - P1 might think it sent the message on time
  - P2 might think the message is late- even when it's not