

# QueCC: Queue-Oriented, Control-Free, Concurrency Architecture



**Thamir Qadah**

*School of Electrical and  
Computer Engineering*



**PURDUE**  
UNIVERSITY



**Mohammad Sadoghi**

*Department of  
Computer Science*

**UCDAVIS**  
UNIVERSITY OF CALIFORNIA



**Exploratory Systems Lab**  
**UCDAVIS**

# Hardware Trends

Large core counts


Large main-memory



HPE Superdome Server  
144 physical cores  
6TB of RAM

# High-Contention Workloads

Challenge ???



High number of  
contented operations





# State-of-the-Art Concurrency Control Protocols

- Optimized for multi-core hardware and main-memory databases
- Non-deterministic



CC	Class	Year
SILO	Optimistic CC	SOSP '13
TICTOC	Timestamp Ordering	SIGMOD '16
FOEDUS-MOCC	Optimistic CC	VLDB '16
ERMIA	MVCC	SIGMOD '16
Cicada	MVCC	SIGMOD '17

# Performance Under High-Contention



Optimize-for-multi-core concurrency control techniques suffer under high-contention due to increasing abort rate

# Performance Under High-Contention



Under high-contention: Non-deterministic aborts dominates

2PL - NoWait

Abort Count: 0

Client Transactions


w <sub>4</sub> (b)	w <sub>3</sub> (b)	w <sub>2</sub> (b)	r <sub>1</sub> (a)
r <sub>4</sub> (d)	r <sub>3</sub> (c)	r <sub>2</sub> (a)	w <sub>1</sub> (b)

each color presents a transaction

Worker Thread #1



Worker Thread #2




2PL - NoWait

Abort Count: 0



Client Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)

Worker  
Thread #1



Worker  
Thread #2




2PL - NoWait

Abort Count: 0



Client Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)

Worker  
Thread #1



Worker  
Thread #2




2PL - NoWait

Abort Count: 0



Client Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)

Worker Thread #1



Worker Thread #2




2PL - NoWait

Abort Count: 0

Client Transactions



w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)

Worker Thread #1



conflict!

Worker Thread #2



2PL - NoWait



Abort Count: 0

Abort transaction (to avoid potential deadlocks)


Client Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)

Worker  
Thread #1



Worker  
Thread #2



2PL - NoWait

Abort Count: 1

Client Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)

r <sub>1</sub> (a)
w <sub>1</sub> (b)


Worker  
Thread #1



Worker  
Thread #2



w <sub>2</sub> (b)
r <sub>2</sub> (a)



2PL - NoWait

Abort Count: 1

Client Transactions

w<sub>4</sub>(b)  
r<sub>4</sub>(d)


r<sub>1</sub>(a)  
w<sub>1</sub>(b)

Worker  
Thread #1

w<sub>3</sub>(b)  
r<sub>3</sub>(c)

Worker  
Thread #2

w<sub>2</sub>(b)  
r<sub>2</sub>(a)



2PL - NoWait

Abort Count: 1

Client Transactions

w<sub>4</sub>(b)  
r<sub>4</sub>(d)


r<sub>1</sub>(a)  
w<sub>1</sub>(b)

Worker  
Thread #1

w<sub>3</sub>(b)  
r<sub>3</sub>(c)

Worker  
Thread #2

w<sub>2</sub>(b)  
r<sub>2</sub>(a)



2PL - NoWait

Abort Count: 1

Client Transactions

w<sub>4</sub>(b)  
r<sub>4</sub>(d)

r<sub>1</sub>(a)  
w<sub>1</sub>(b)


Worker  
Thread #1

w<sub>3</sub>(b)  
r<sub>3</sub>(c)

conflict!

Worker  
Thread #2

w<sub>2</sub>(b)  
r<sub>2</sub>(a)




2PL - NoWait

Abort Count: 1

Abort transaction (to avoid potential deadlocks)

Worker  
Thread #1





Client Transactions

w<sub>4</sub>(b)  
r<sub>4</sub>(d)

r<sub>1</sub>(a)  
w<sub>1</sub>(b)

Worker  
Thread #2



2PL - NoWait

Abort Count: 2

Client Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)

r <sub>1</sub> (a)
w <sub>1</sub> (b)


Worker  
Thread #1



Worker  
Thread #2



w <sub>2</sub> (b)
r <sub>2</sub> (a)



2PL - NoWait

Abort Count: 2

Client Transactions

w<sub>3</sub>(b)  
r<sub>3</sub>(c)


r<sub>1</sub>(a)  
w<sub>1</sub>(b)

Worker  
Thread #1

w<sub>4</sub>(b)  
r<sub>4</sub>(d)

Worker  
Thread #2

w<sub>2</sub>(b)  
r<sub>2</sub>(a)



2PL - NoWait

Abort Count: 2

Client Transactions

w<sub>3</sub>(b)  
r<sub>3</sub>(c)


r<sub>1</sub>(a)  
w<sub>1</sub>(b)

Worker  
Thread #1

w<sub>4</sub>(b)  
r<sub>4</sub>(d)

Worker  
Thread #2

w<sub>2</sub>(b)  
r<sub>2</sub>(a)



2PL - NoWait

Abort Count: 2

Client Transactions

w<sub>3</sub>(b)  
r<sub>3</sub>(c)


r<sub>1</sub>(a)  
w<sub>1</sub>(b)

Worker  
Thread #1

w<sub>4</sub>(b)  
r<sub>4</sub>(d)

Worker  
Thread #2

w<sub>2</sub>(b)  
r<sub>2</sub>(a)



2PL - NoWait

Abort Count: 2

Client Transactions

w<sub>3</sub>(b)  
r<sub>3</sub>(c)

r<sub>1</sub>(a)  
w<sub>1</sub>(b)


Worker  
Thread #1

w<sub>4</sub>(b)  
r<sub>4</sub>(d)

conflict!

Worker  
Thread #2

w<sub>2</sub>(b)  
r<sub>2</sub>(a)




2PL - NoWait


Abort Count: 2

Abort transaction (to avoid potential deadlocks)



Worker  
Thread #1



Client Transactions



Worker  
Thread #2



2PL - NoWait

Abort Count: 3

Client Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)
r <sub>4</sub> (d)	r <sub>3</sub> (c)


r<sub>1</sub>(a)

w<sub>1</sub>(b)

Worker  
Thread #1



Worker  
Thread #2



Committed Transactions

w <sub>2</sub> (b)
r <sub>2</sub> (a)

2PL - NoWait


Abort Count: 5

Client Transactions

Worker  
Thread #1



Worker  
Thread #2




- Eventually transactions commit in some serial order!
- Many aborts due to high contention on record b
- Non-determinism in CC cause these aborts
- Wasted work

Committed Transactions

w <sub>4</sub> (b)	r <sub>1</sub> (a)	w <sub>3</sub> (b)	w <sub>2</sub> (b)
r <sub>4</sub> (d)	w <sub>1</sub> (b)	r <sub>3</sub> (c)	r <sub>2</sub> (a)

# Key Insights

- Many aborts due to high contention
- Non-determinism in CC cause these aborts
- Can we do better?
- Is it possible to eliminate non-deterministic concurrency control from transaction execution?



# Deterministic Transaction Execution

- H-Store [Kallman et al. '08]
- Designed and optimized for horizontal scalability, multi-core hardware and in-memory databases
- Stored procedure transaction model
- Static partitioning of database
- Assigns a single core to each partition
- Execute transaction serially without concurrency control within each partition

H-Store

Abort Count: 0

Client Transactions

w <sub>4</sub> (d)	w <sub>3</sub> (b)	w <sub>2</sub> (c)	r <sub>1</sub> (a)
r <sub>4</sub> (c)	r <sub>3</sub> (a)	r <sub>2</sub> (d)	w <sub>1</sub> (b)

Single-partition  
transactions

Worker  
Thread #1




P1 is assigned to  
Worker Thread #1

Worker  
Thread #2



P2 is assigned to  
Worker Thread #2




H-Store

Abort Count: 0



Client Transactions

w <sub>4</sub> (d)	w <sub>3</sub> (b)
r <sub>4</sub> (c)	r <sub>3</sub> (a)

Worker Thread #1



Worker Thread #2



Committed Transactions

H-Store

Abort Count: 0


Client Transactions

w <sub>4</sub> (d)	w <sub>3</sub> (b)
r <sub>4</sub> (c)	r <sub>3</sub> (a)

Worker Thread #1



Worker Thread #2



Committed Transactions

w <sub>2</sub> (c)	r <sub>1</sub> (a)
r <sub>2</sub> (d)	w <sub>1</sub> (b)

H-Store

Abort Count: 0


Client Transactions

Worker  
Thread #1

w<sub>3</sub>(b)  
r<sub>3</sub>(a)

Worker  
Thread #2

w<sub>4</sub>(d)  
r<sub>4</sub>(c)



Committed Transactions

w <sub>2</sub> (c)	r <sub>1</sub> (a)
r <sub>2</sub> (d)	w <sub>1</sub> (b)

H-Store


Abort Count: 0

Client Transactions

Worker  
Thread #1



Worker  
Thread #2



Committed Transactions

w <sub>4</sub> (d)	w <sub>3</sub> (b)	w <sub>2</sub> (c)	r <sub>1</sub> (a)
r <sub>4</sub> (c)	r <sub>3</sub> (a)	r <sub>2</sub> (d)	w <sub>1</sub> (b)

H-Store

Abort Count: 0

Client Transactions

Worker Thread #1




Worker Thread #2



- ✓ Deterministic Execution
- ✓ No aborts because of CC
- ✓ Minimal coordination among threads


● Performs well only when transactions are single-partitioned



Committed Transactions

w <sub>4</sub> (d)	w <sub>3</sub> (b)	w <sub>2</sub> (c)	r <sub>1</sub> (a)
r <sub>4</sub> (c)	r <sub>3</sub> (a)	r <sub>2</sub> (d)	w <sub>1</sub> (b)

# Effect of Increasing Percentage of Multi-Partition Transactions in the Workload



H-Store is sensitive to the percentage of multi-partition transactions in the workload


# Can We Do Better?

Our motivations are

- Efficiently exploits multi-core and large main-memory systems
- Provide serializable multi-statement transactions for key-value stores
- Scales well under high-contention workloads

Desired Properties

- Concurrent execution over shared data
- Not limited to partitionable workloads
- Without any concurrency controls



*Is it possible to have concurrent execution over shared data without having any concurrency controls?*

# Introducing: QueCC

## Queue-Oriented, Control-Free, Concurrency Architecture

*A two parallel & independent phases of priority-driven planning & execution*

**Phase 1:** Deterministic priority-based planning of transaction operations in parallel

- *Plans take the form of **Prioritized Execution Queues***
- Execution Queues inherits predetermined priority of its planner
- Results in a deterministic plan of execution

**Phase 2:** Priority driven execution of plans in parallel

- Satisfies the ***Execution Priority Invariance***


*“For each record (or a queue), operations that belong to higher priority queues (created by a higher priority planner) must always be executed before executing any lower priority operations.”*

# QueCC Architecture


Priority-based Parallel Planning Phase

Batching Client  
Transactions


# QueCC Architecture





# QueCC Architecture







# QueCC Architecture









## Priority Groups





Low-priority  
QueuesHigh-priority  
Queues




Committed Transactions

QueCC

Abort Count: 0










QueCC

Abort Count: 0




## Priority Groups

### Low-priority Queues




### High-priority Queues






Committed Transactions




## Priority Groups




Execution Priority  
Invariance

Low-priority  
QueuesHigh-priority  
Queues



Committed Transactions



## Priority Groups











Execution Priority  
Invariance





**Low-priority  
Queues****High-priority  
Queues**


Committed Transactions

w<sub>2</sub>(b) r<sub>1</sub>(a)  
r<sub>2</sub>(a) w<sub>1</sub>(b)












- ✓ Deterministic Execution
- ✓ No aborts because of CC
- ✓ Minimal coordination among threads
- ✓ Not sensitive to multi-partition transactions
- ✓ Exploits Intra-transaction parallelism

## Priority Groups

### Low-priority Queues




### High-priority Queues



## Committed Transactions

w <sub>4</sub> (b)	w <sub>3</sub> (b)	w <sub>2</sub> (b)	r <sub>1</sub> (a)
r <sub>4</sub> (d)	r <sub>3</sub> (c)	r <sub>2</sub> (a)	w <sub>1</sub> (b)

# ExpoDB Fabric



# Evaluation Environment

## Hardware

Microsoft Azure instance with 32 core  
CPU: Intel Xeon E5-2698B v3  
*32KB L1 data and instruction caches*  
*256KB L2 cache*  
*40MB L3 cache*  
RAM: 448GB

## Workload


YCSB: 1 table, 10 operations, 50% RMW, Zipfian distribution  
TPCC: 9 tables, Payment and NewOrder, 1 Warehouse

## Software

Operating System: Ubuntu LTS 16.04.3  
Compiler: GCC with -O3 compiler optimizations

# Effect of Varying Contention


- 5 write and 5 read operation per transaction
- 32 worker threads



Workload contention resiliency  
Cache locality under high-contention


# Effect of Varying Worker Threads

- 5 write and 5 read operation per transaction
- Zipfian theta = 0.99




Avoiding thread coordination & eliminating all execution-induced aborts

# Effect of Increasing Percentage of Multi-Partition Transactions in the Workload




# Effect of Increasing Percentage of Multi-Partition Transactions in the Workload



QueCC is not sensitive to multi-partitioning

# TPC-C Results

**1 Warehouse (highly contended workload)**  
**50% Payment + 50% NewOrder transaction mix**




QueCC can achieve up to 3x better performance on high-contention  
TPC-C workloads

# QueCC Conclusions

- ✓ Efficient, parallel and deterministic in-memory transaction processing
- ✓ Eliminates almost all aborts by resolving transaction conflicts *a priori*
- ✓ Works extremely well under high-contention workloads



# What's Next: Q-Store




# What's Next: Q-Store



Plan Local and Remote  
Execution Queues


# What's Next: Q-Store



Plan Local and Remote  
Execution Queues

Deliver Remote  
Execution Queues

# What's Next: Q-Store




Plan Local and Remote  
Execution Queues

Deliver Remote  
Execution Queues


Execute  
Queues

# What's Next: Q-Store



Multi-core  
Single-node

# What's Next: QBFT




# What's Next: QBFT



Queue-oriented  
Byzantine Fault-  
Tolerance



Resilient planning  
followed by resilient  
execution





**Mohammad Sadoghi**  
*(Principal Investigator)*

**THANK  
YOU**



**Jelle Hellings, PostDoc**  
*(Blockchain)*



**Suyash Gupta, PhD**  
*(Blockchain)*



**Thamir Qadah, PhD**  
*(Coordination-free Concurrency)*



**Sajjad Rahnama, PhD**  
*(Blockchain)*




**Nikhil Wadhwa, PhD**  
*(Blockchain)*




**Masoud Hemmatpour, PhD**  
*(RDMA KV-Stores)*




**Domenic Cianfichi, MSc**  
*(Blockchain)*



**Shreenath Iyer, MSc**  
*(Blockchain)*



**Robert He, MSc**  
*(Coordination-free Concurrency)*



**Patrick Liao, BSc**  
*(Blockchain)*