



Hive - A Warehousing Solution Over a Map-Reduce Framework

Instructor: Mohammad Sadoghi
msadoghi@purdue.edu

SQL at Facebook: Emergence of Apache Hive



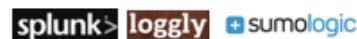
- Developed Hive to address traditional RDBMS limitations.
- 300+ PB of data under management⁽¹⁾.
- 600+ TB of data loaded daily.
- 60,000+ Hive queries per day⁽²⁾.
- More than 1,000 users per day.
- Initial Apache release in April 2009.

Big Data Landscape

Vertical Apps



Log Data Apps



Ad/Media Apps



Business Intelligence



Analytics and Visualization



Data As A Service



Analytics Infrastructure



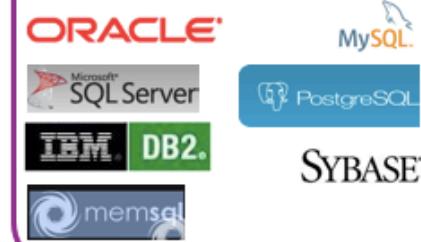
Operational Infrastructure



Infrastructure As A Service



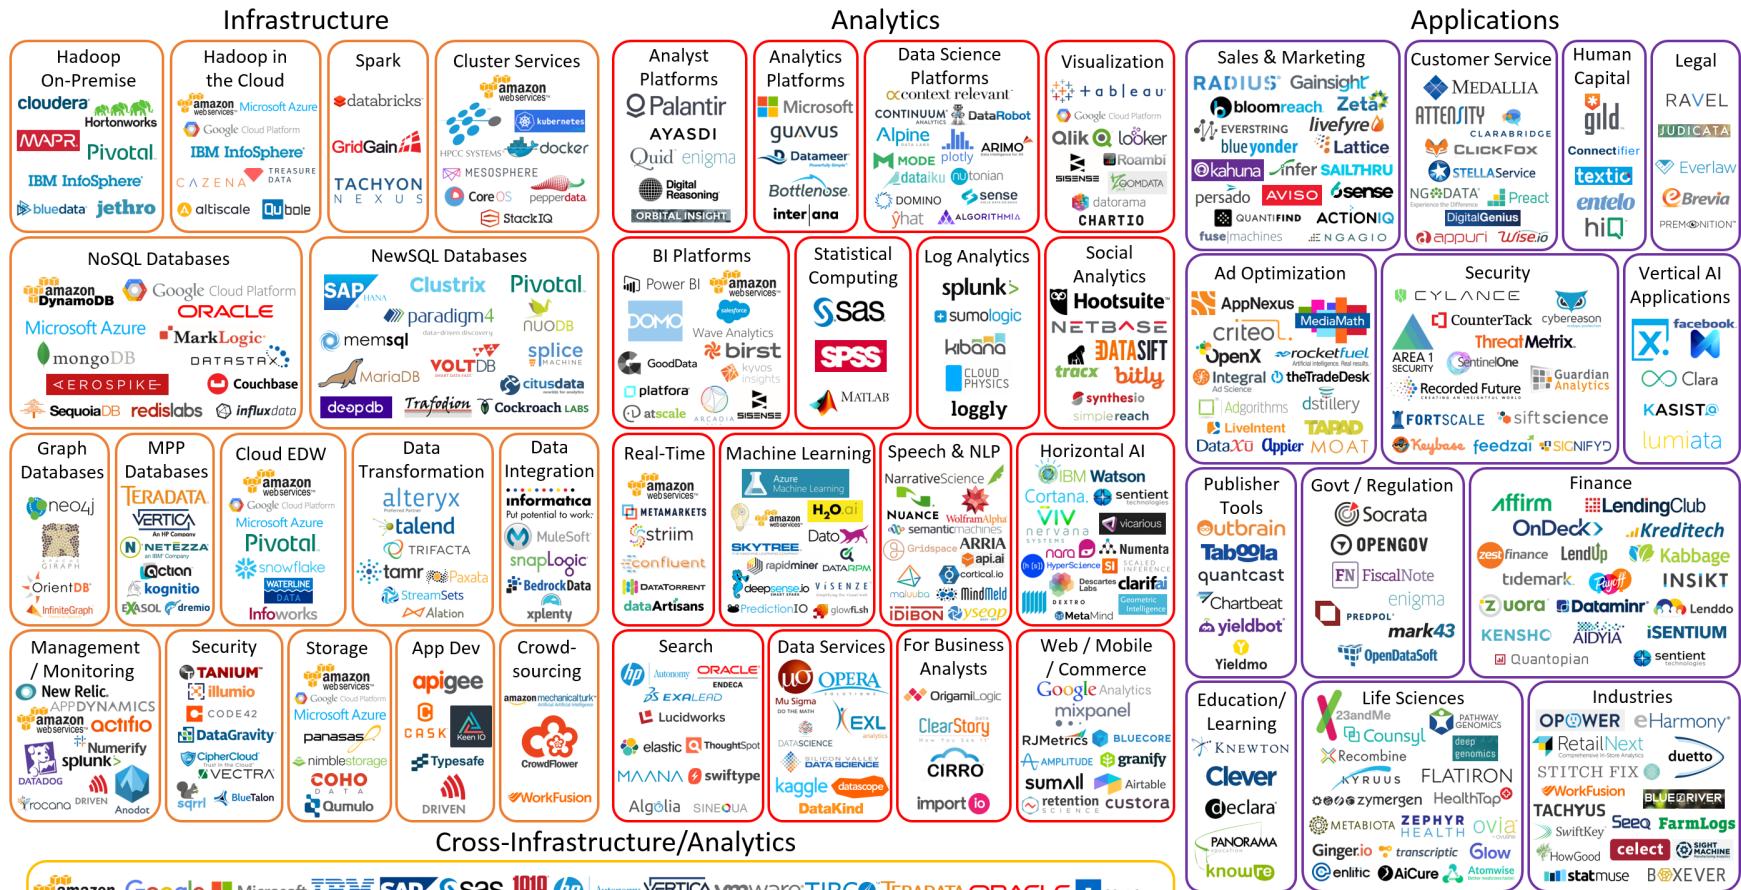
Structured Databases



Technologies



Big Data Landscape 2016 (Version 3.0)



amazon web services, Google, Microsoft, IBM, SAP, SAS, 1010 data, hp, Anapty, VERTICAL, vmware, TIBCO, TERADATA, ORACLE, NetApp



Hive Classic: Strengths and Challenges



Familiar SQL Interface



Economical Processing of Petabytes



Hive Classic tied to MapReduce, leading to latency

Traditional SQL Workloads Needed Higher Performance!

Challenges that Data Analysts faced

- Data Explosion
 - TBs of data generated everyday

Solution – HDFS to store data and Hadoop Map-Reduce framework to parallelize processing of Data

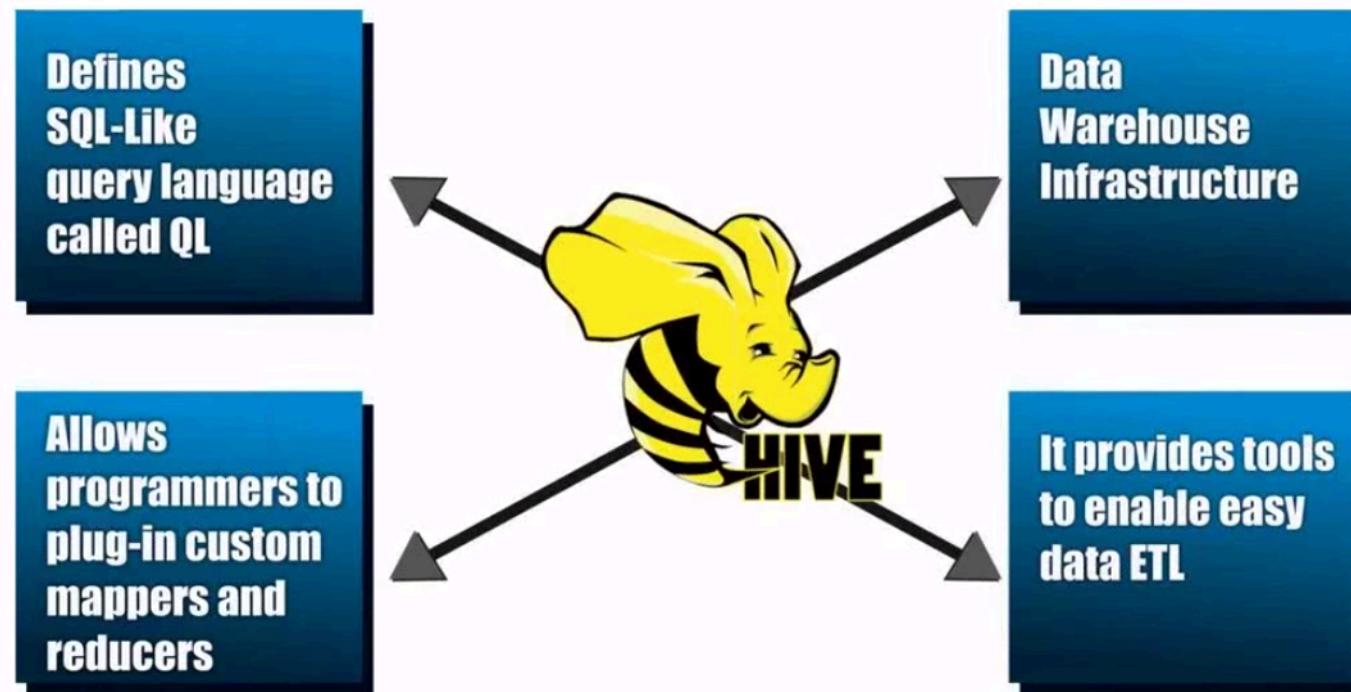
What is the catch?

- Hadoop Map Reduce is Java intensive
- Thinking in Map Reduce paradigm can get tricky

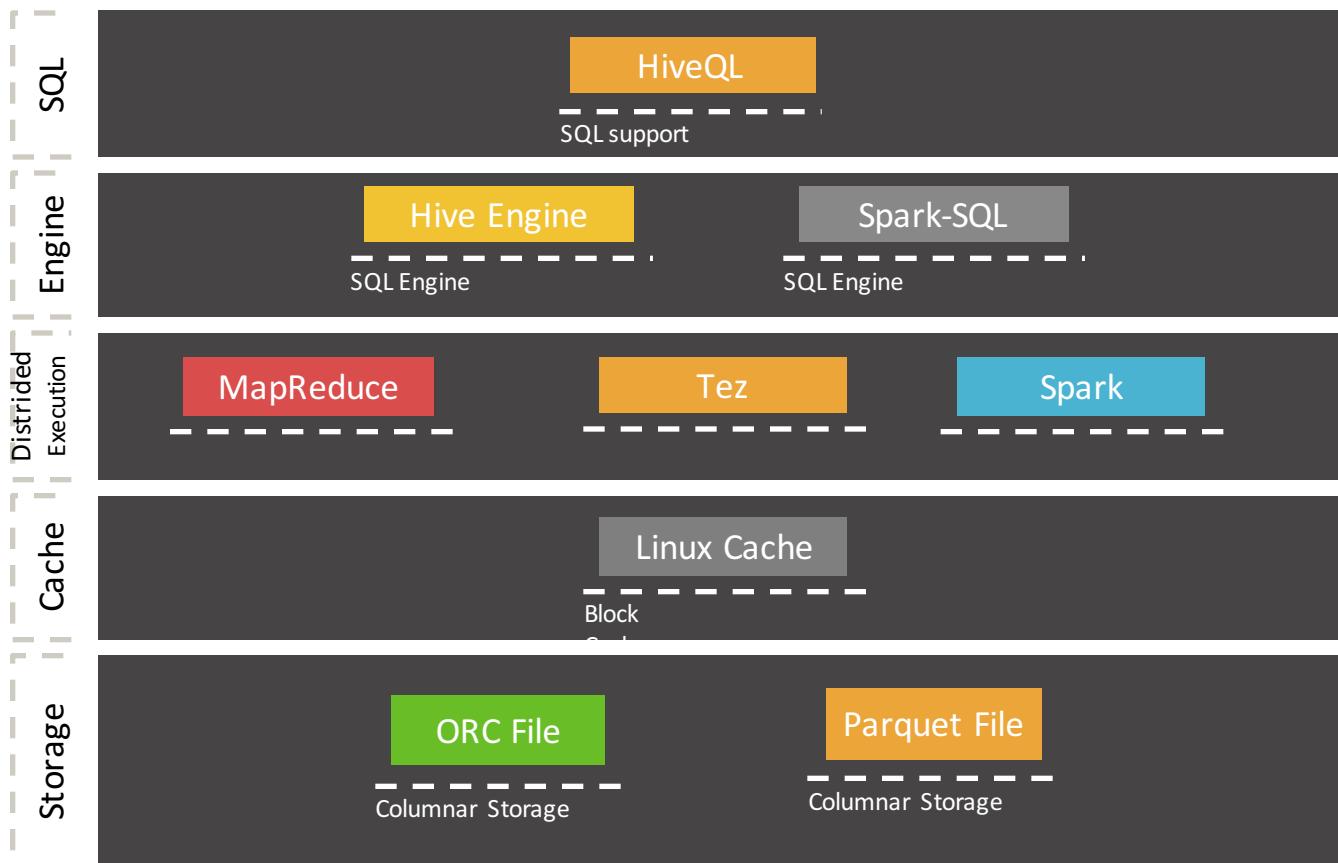
What is HIVE?

- A system for querying and managing structured data built on top of Hadoop
 - Uses Map-Reduce for execution
 - HDFS for storage – but any system that implements Hadoop FS API
- Key Building Principles:
 - Structured data with rich data types (structs, lists and maps)
 - Directly query data from different formats (text/binary) and file formats (Flat/Sequence)
 - SQL as a familiar programming tool and for standard analytics
 - Allow embedded scripts for extensibility and for non standard applications

Hive Key Principles

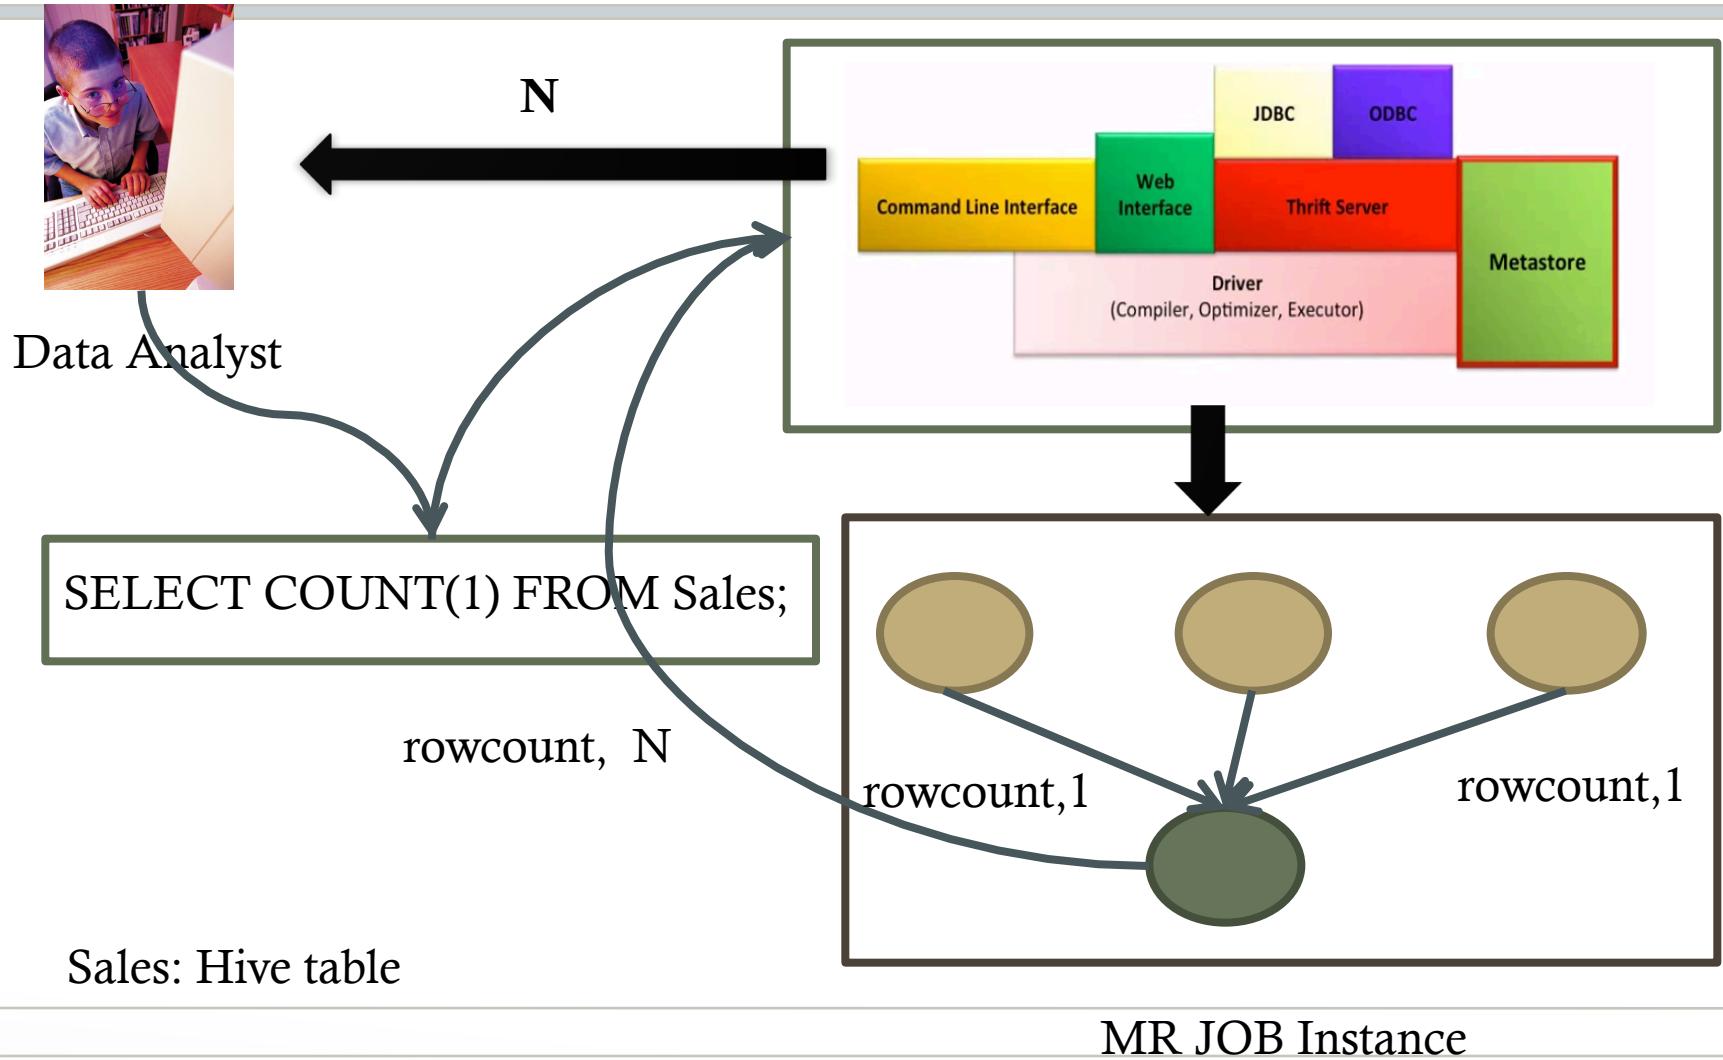


Query processing in Hadoop



HiveQL to MapReduce

Hive Framework

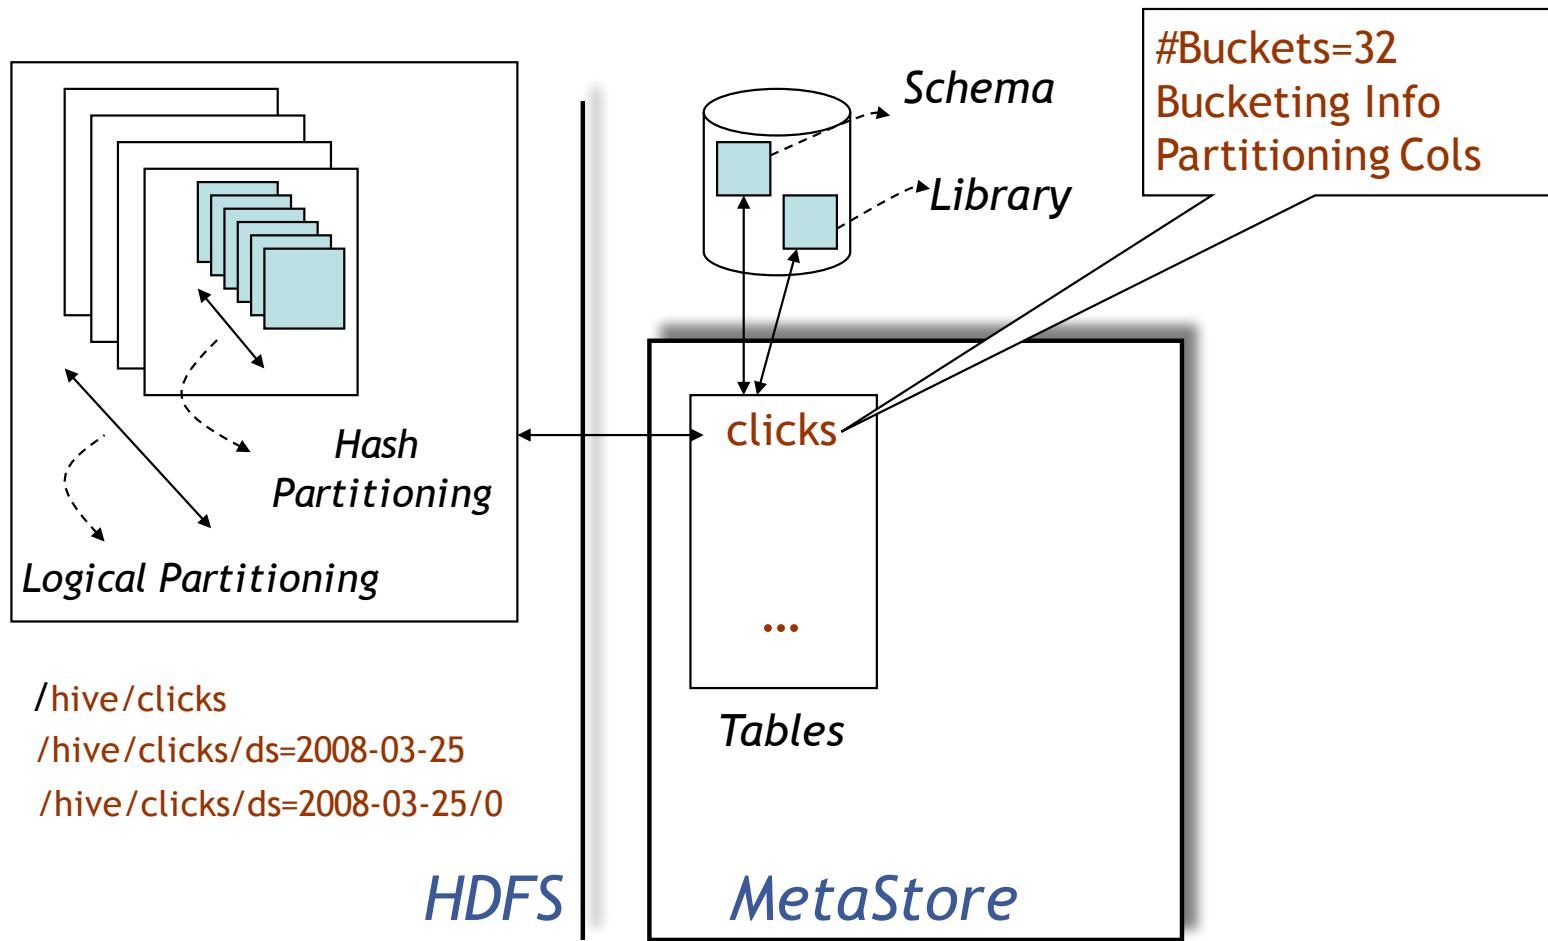


Hive Data Model

Data in Hive organized into :

- Tables
- Partitions
- Buckets

Hive Data Model



Hive Data Model Contd.

- Tables
 - Analogous to relational tables
 - Each table has a corresponding directory in HDFS
 - Data serialized and stored as files within that directory
 - Hive has default serialization built in which supports compression and lazy deserialization
 - Users can specify custom serialization –deserialization schemes (**SerDe's**)

Hive Data Model Contd.

- Partitions
 - Each table can be broken into partitions
 - Partitions determine distribution of data within subdirectories

Example -

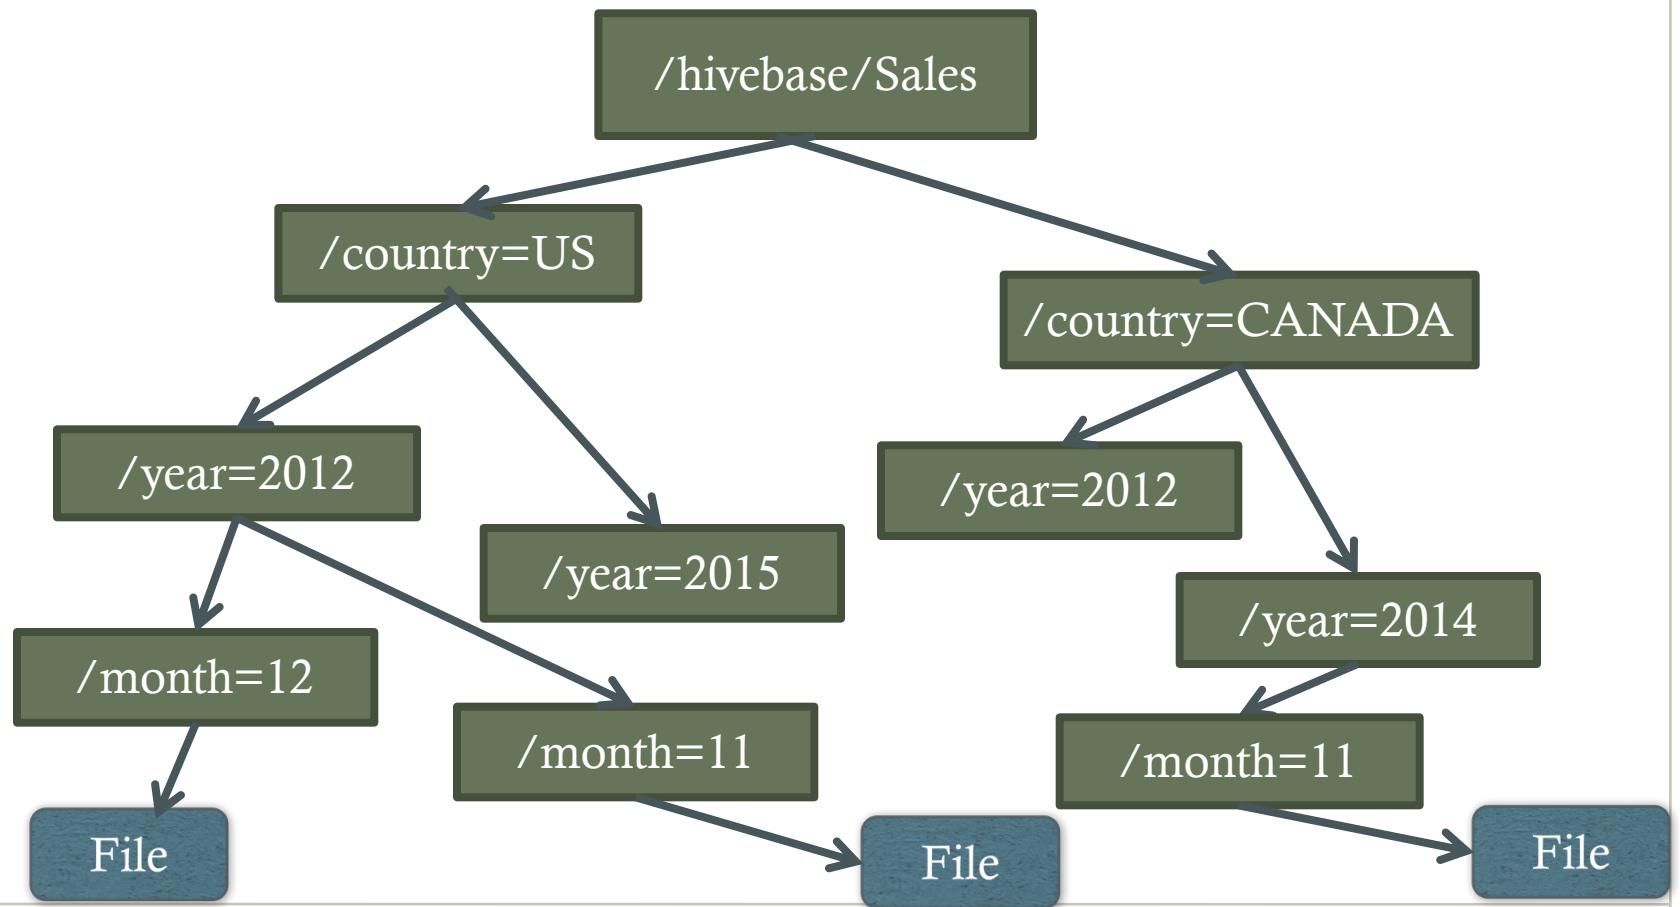
CREATE_TABLE Sales (sale_id INT, amount FLOAT)

PARTITIONED BY (country STRING, year INT, month INT)

So each partition will be split out into different folders like

Sales/country=US/year=2012/month=12

Hierarchy of Hive Partitions



Hive Data Model Contd.

- Buckets
 - Data in each partition divided into buckets
 - Based on a hash function of the column
 - **$H(\text{column}) \bmod \text{NumBuckets} = \text{bucket number}$**
 - Each bucket is stored as a file in partition directory

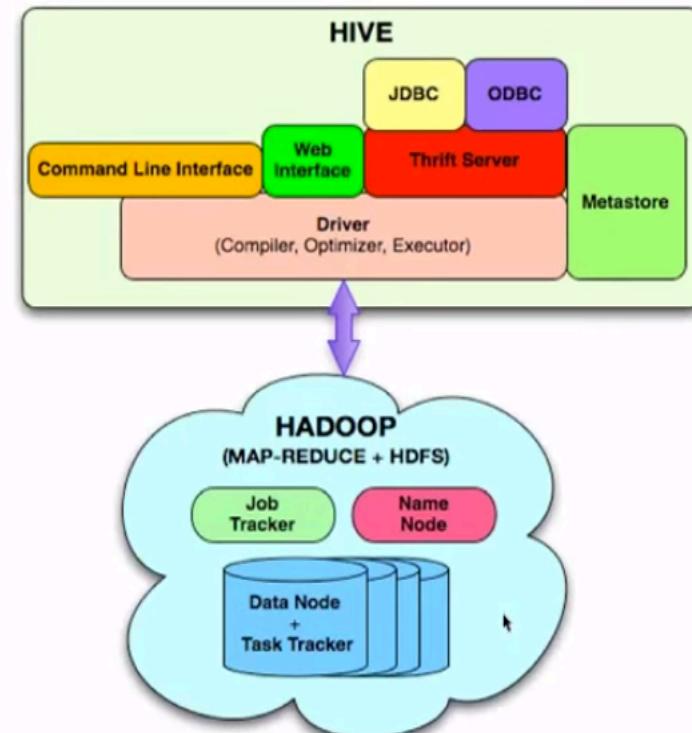
Architecture

External Interfaces- CLI, WebUI, JDBC, ODBC programming interfaces

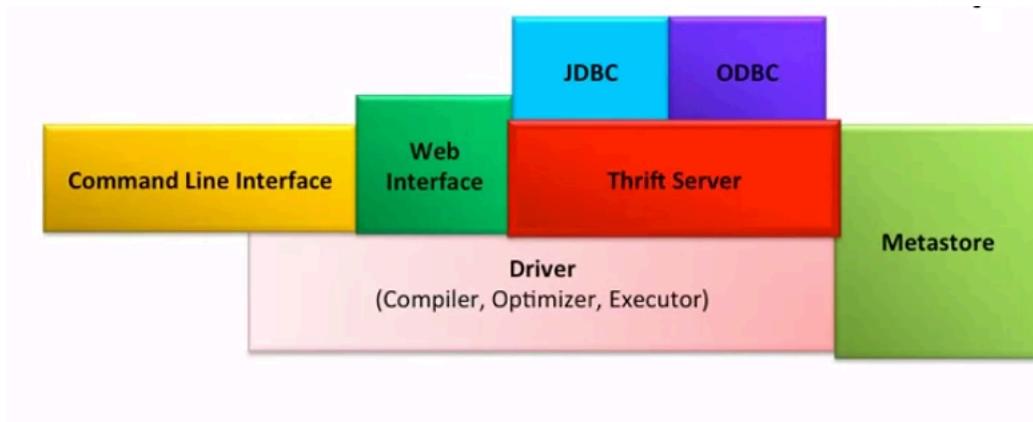
Thrift Server – Cross Language service framework .

Metastore - Meta data about the Hive tables, partitions

Driver - Brain of Hive! Compiler, Optimizer and Execution engine

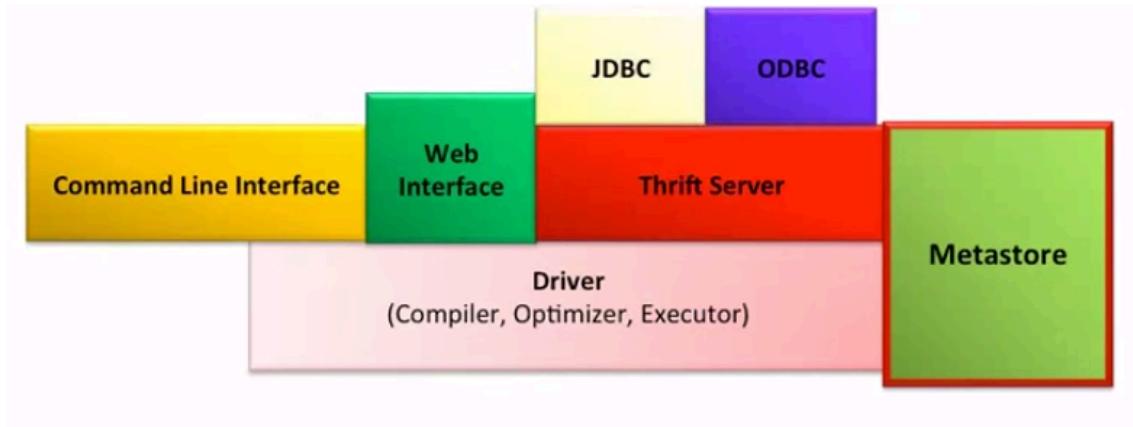


Hive Thrift Server



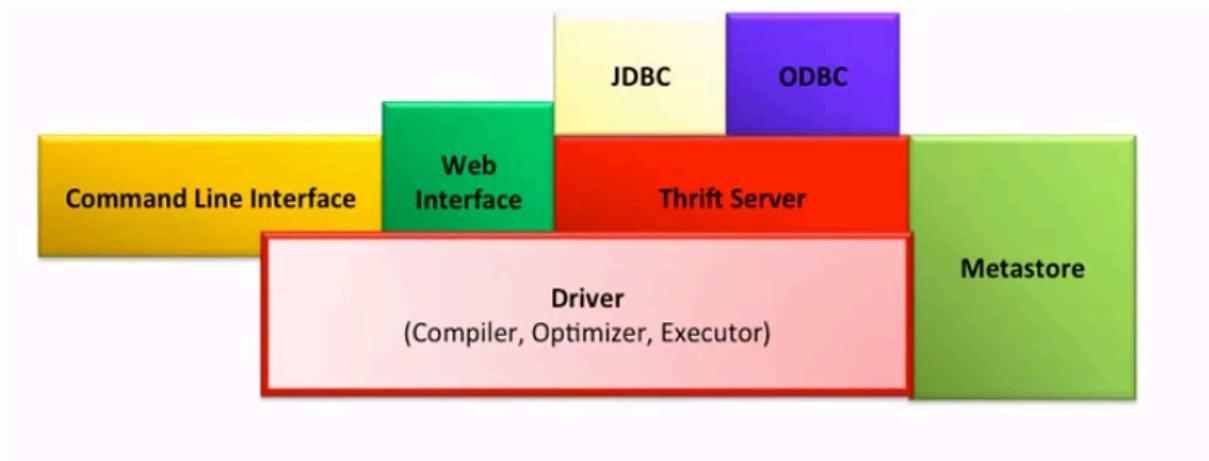
- Framework for cross language services
- Server written in Java
- Support for clients written in different languages
 - JDBC(java), ODBC(c++), php, perl, python scripts

Metastore



- System catalog which contains metadata about the Hive tables
- Stored in RDBMS/local fs. HDFS too slow(not optimized for random access)
- Objects of Metastore
 - Database - Namespace of tables
 - Table - list of columns, types, owner, storage, SerDes
 - Partition – Partition specific column, Serdes and storage

Hive Driver



- **Driver** - Maintains the lifecycle of HiveQL statement
- **Query Compiler** – Compiles HiveQL in a DAG of map reduce tasks
- **Executor** - Executes the tasks plan generated by the compiler in proper dependency order. Interacts with the underlying Hadoop instance

Compiler

- Converts the HiveQL into a plan for execution
- Plans can
 - Metadata operations for DDL statements e.g. CREATE
 - HDFS operations e.g. LOAD
- Semantic Analyzer – checks schema information, type checking, implicit type conversion, column verification
- Optimizer – Finding the best logical plan e.g. Combines multiple joins in a way to reduce the number of map reduce jobs, Prune columns early to minimize data transfer
- Physical plan generator – creates the DAG of map-reduce jobs

Hive Query Language

- Basic SQL
 - From clause subquery
 - ANSI JOIN (equi-join only)
 - Multi-table Insert
 - Multi group-by
 - Sampling
 - Objects traversal
- Extensibility
 - Pluggable Map-reduce scripts using TRANSFORM

HiveQL

DDL :

- CREATE DATABASE
- CREATE TABLE
- ALTER TABLE
- SHOW TABLE
- DESCRIBE

DML:

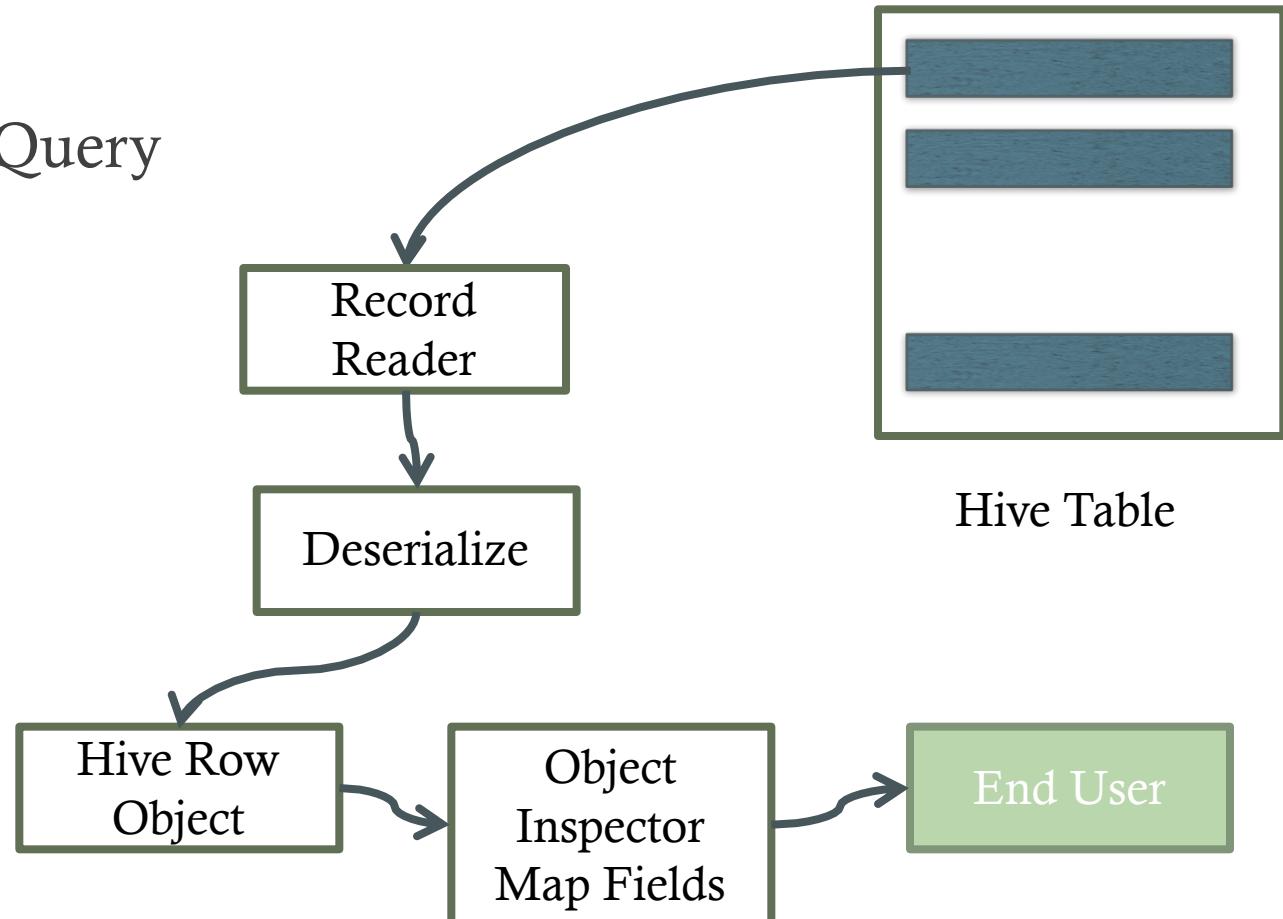
- LOAD TABLE
- INSERT

QUERY:

- SELECT
- GROUP BY
- JOIN
- MULTI TABLE INSERT

Hive SerDe

- SELECT Query
 - Hive built in Serde: Avro, ORC, Regex etc
 - Can use Custom SerDe's (e.g. for unstructured data like audio/video data, semistructured XML data)



Good Things

- Suitable for Data Analysts
- Easy Learning curve
- Completely transparent to underlying Map-Reduce
- Partitions(speed!)
- Flexibility to load data from localFS/HDFS into Hive Tables

Cons and Possible Improvements

- Extending the SQL queries support(Updates, Deletes)
- Parallelize firing independent jobs from the work DAG
- Table Statistics in Metastore
- Explore methods for multi query optimization
- Perform N- way generic joins in a single map reduce job
- Better debug support in shell



Hive v/s Pig



Similarities:

- Both High level Languages which work on top of map reduce framework
- Can coexist since both use the under lying HDFS and map reduce

Differences:

◆ **Language**

- Pig is a procedural ; (A = load 'mydata'; dump A)
- Hive is Declarative (select * from A)

◆ **Work Type**

- Pig more suited for adhoc analysis (on demand analysis of click stream search logs)
- Hive a reporting tool (e.g. weekly BI reporting)



Hive v/s Pig



Differences:

◆ Users

- Pig – Researchers, Programmers (build complex data pipelines, machine learning)
- Hive – Business Analysts

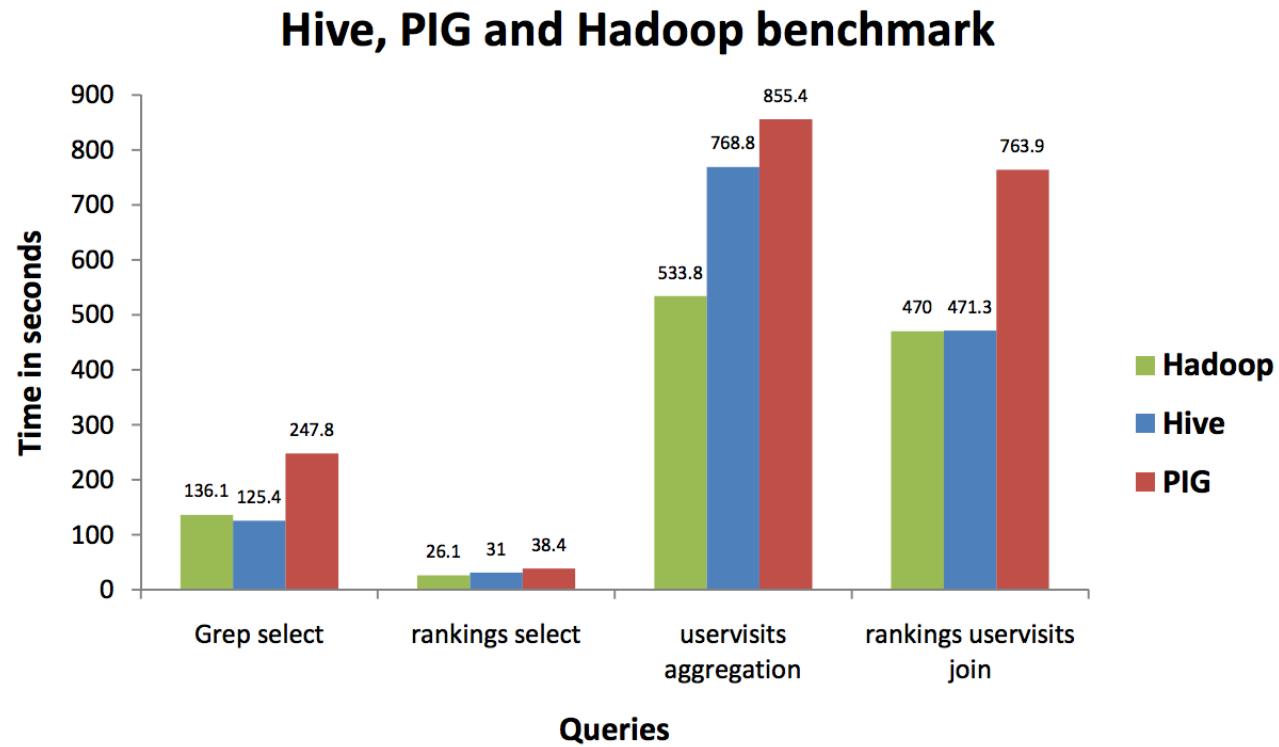
◆ Integration

- Pig - Doesn't have a thrift server(i.e no/limited cross language support)
- Hive - Thrift server

◆ User's need

- Pig – Better dev environments, debuggers expected
- Hive - Better integration with technologies expected(e.g JDBC, ODBC)

Head-to-Head (the bee, the pig, the elephant)



Version: Hadoop – 0.18x, Pig:786346, Hive:786346

REFERENCES

- <https://hive.apache.org/>
- <https://cwiki.apache.org/confluence/display/Hive/Presentations>
- <https://developer.yahoo.com/blogs/hadoop/comparing-pig-latin-sql-constructing-data-processing-pipelines-444.html>
- <http://www.qubole.com/blog/big-data/hive-best-practices/>
- Hortonworks tutorials (youtube)
- Graph :
https://issues.apache.org/jira/secure/attachment/12411185/hive_benchmark_2009-06-18.pdf