

# Fault-Tolerant Distributed Transactions on Blockchain

## *Toward Scalable Blockchain*



Suyash Gupta



Jelle Hellings



Mohammad Sadoghi

**UCDAVIS**  
UNIVERSITY OF CALIFORNIA



## Scalability versus Fully-Replicated Blockchains

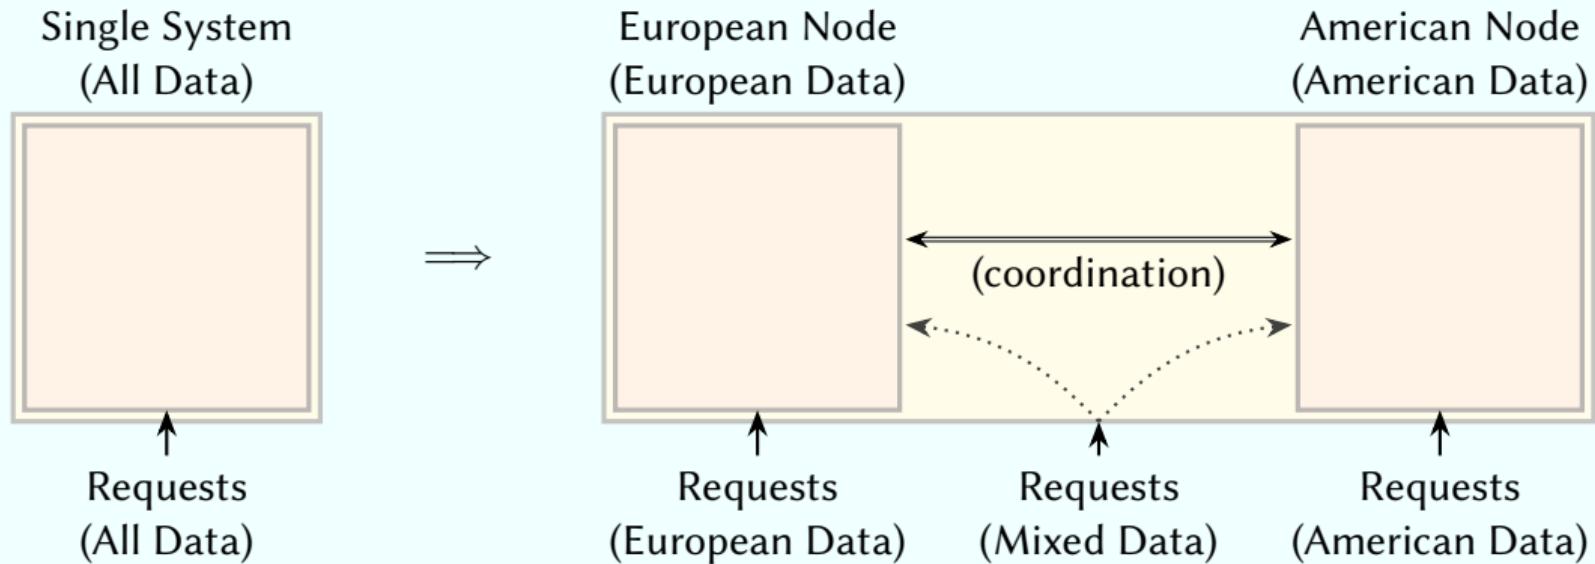
*Scalability: adding resources  $\implies$  adding performance.*

## Scalability versus Fully-Replicated Blockchains

*Scalability: adding resources  $\implies$  adding performance.*

Full replication: adding resources (replicas)  $\implies$  less performance!

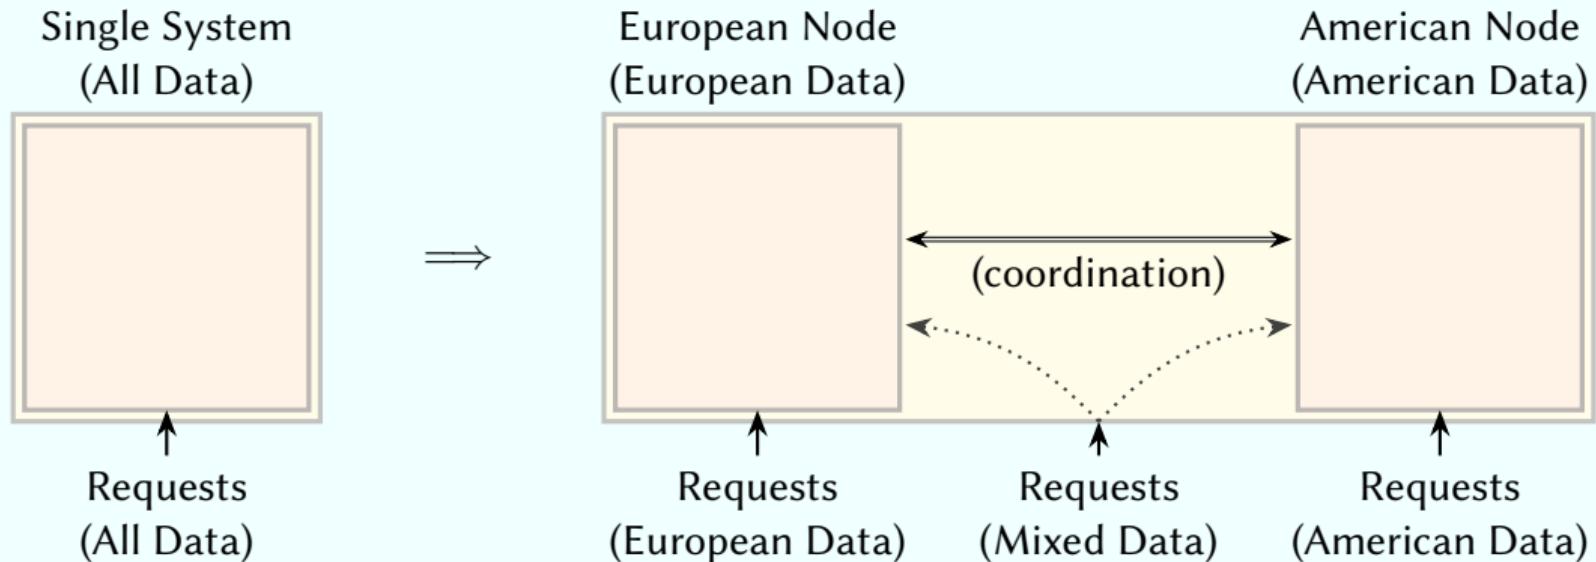
# Distributed Systems: Scalability



Partition the system: More storage and *potentially* more performance.

Potentially *lower latencies* if data ends up closer to users.

# Distributed Systems: Scalability

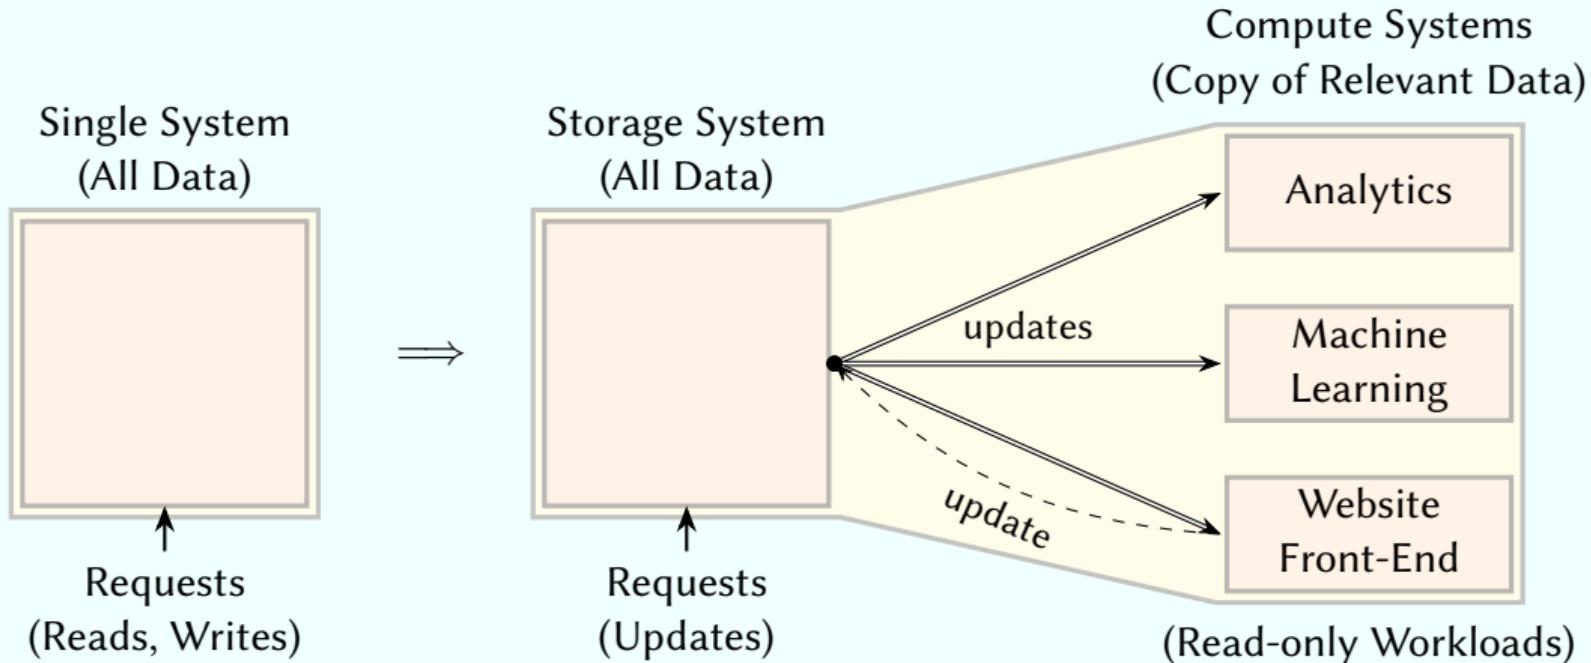


Partition the system: More storage and *potentially* more performance.

Potentially *lower latencies* if data ends up closer to users.

Adding shards  $\implies$  adding throughput (parallel processing), adding storage.

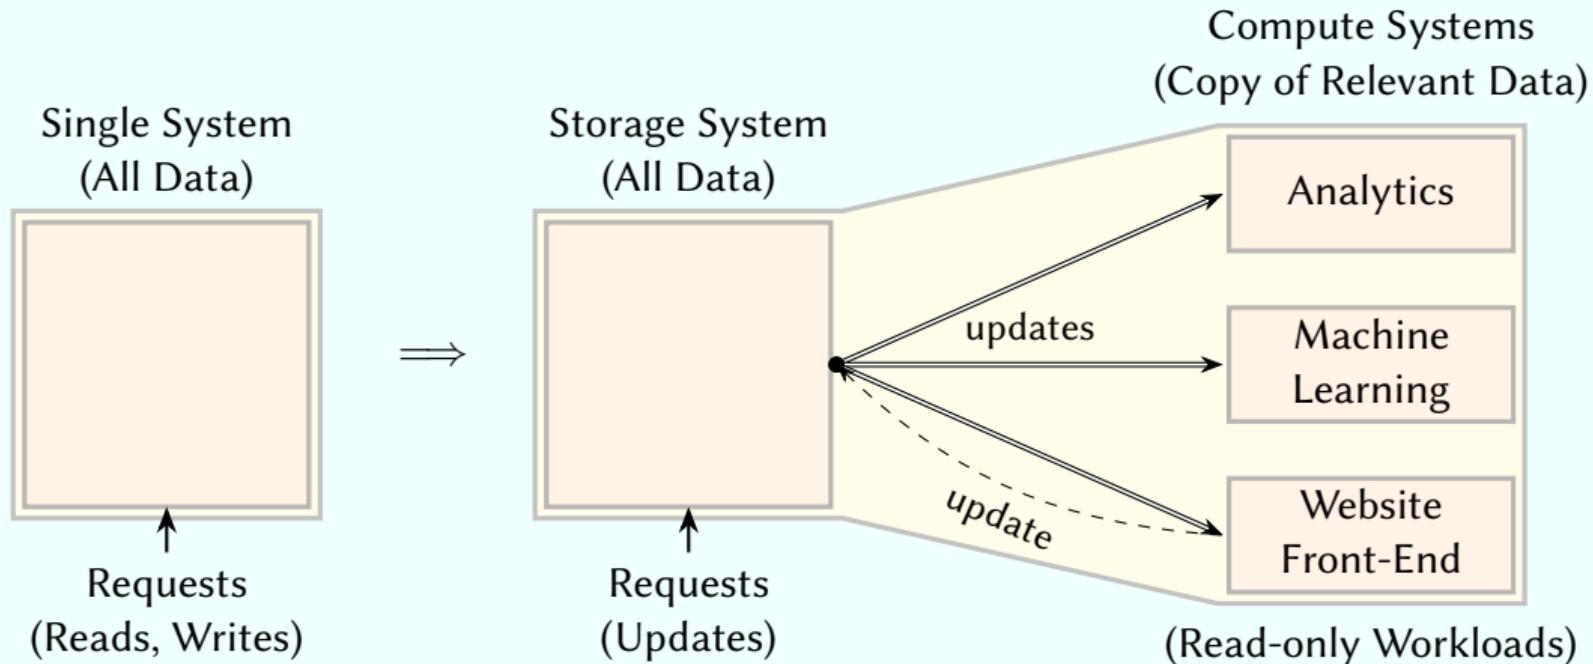
# Distributed Systems: Specialization



Specialize the system: Different nodes have distinct tasks.

Specialized hardware and software *per* task.

# Distributed Systems: Specialization



Specialize the system: Different nodes have distinct tasks.

Specialized hardware and software *per* task.

Specializing roles  $\Rightarrow$  adding throughput (parallel processing, specialized hardware, ...).

# Central Ideas for Improvement

## Reminder

We can make a resilient system that manages data: e.g., fully-replicated blockchains.

- ▶ **Role Specialization:** make the storage system a blockchain.  
Requires: *reliable read-only updates of the blockchain.*  
Permissionless blockchains: light clients!
- ▶ **Sharding:** make each shard an independent blockchain.  
Requires: *reliable communication between blockchains.*  
Permissionless blockchains: relays, atomic swaps!

# Central Ideas for Improvement

## Reminder

We can make a resilient system that manages data: e.g., fully-replicated blockchains.

- ▶ **Role Specialization:** make the storage system a blockchain.  
Requires: *reliable read-only updates of the blockchain.*  
Permissionless blockchains: light clients!
- ▶ **Sharding:** make each shard an independent blockchain.  
Requires: *reliable communication between blockchains.*  
Permissionless blockchains: relays, atomic swaps!

Consensus is of no use here if we want efficiency.

# Reliable Read-Only Updates of Fault-Tolerant Clusters

## Definition

Let  $\mathcal{C}$  be a cluster deciding on a sequence of transactions  $\mathcal{L}$  and  $\mathsf{L}$  be a learner.

The *Byzantine learning problem* is the problem of sending  $\mathcal{L}$  from  $\mathcal{C}$  to  $\mathsf{L}$  such that:

- ▶ the learner  $\mathsf{L}$  will eventually *receive all* decided transactions;
- ▶ the learner  $\mathsf{L}$  will *only receive* decided transactions.

# Reliable Read-Only Updates of Fault-Tolerant Clusters

## Definition

Let  $\mathcal{C}$  be a cluster deciding on a sequence of transactions  $\mathcal{L}$  and  $L$  be a learner.

The *Byzantine learning problem* is the problem of sending  $\mathcal{L}$  from  $\mathcal{C}$  to  $L$  such that:

- ▶ the learner  $L$  will eventually *receive all* decided transactions;
- ▶ the learner  $L$  will *only receive* decided transactions.

## Practical requirements

- ▶ Minimizing overall communication.
- ▶ Load balancing among all replicas in  $\mathcal{C}$ .

# Background: Information Dispersal Algorithms

## Definition

Let  $v$  be a value with storage size  $s = \|v\|$ .

An *information dispersal algorithm* can encode  $v$  in  $n$  pieces  $v'$  such that  $v$  can be *decoded* from every set of  $n - f$  such pieces.

## Theorem (Rabin 1989)

The IDA algorithm is an *optimal* information dispersal algorithm:

- ▶ Each piece  $v'$  has size  $\left\lceil \frac{\|v\|}{n-f} \right\rceil$ .
- ▶ The  $n - f$  pieces necessary for decoding have a total size of  $(n - f) \left\lceil \frac{\|v\|}{(n-f)} \right\rceil \approx \|v\|$ .

# The Delayed-Replication Algorithm

Idea:  $\mathcal{C}$  sends a ledger to learner  $L$

# The Delayed-Replication Algorithm

Idea:  $\mathcal{C}$  sends a ledger to learner  $L$

1. Partition the ledger in sequences  $S$  of  $n$  transactions.

# The Delayed-Replication Algorithm

Idea:  $\mathcal{C}$  sends a ledger to learner  $\mathsf{L}$

1. Partition the ledger in sequences  $S$  of  $\mathbf{n}$  transactions.
2. Replica  $\mathsf{R}_i \in \mathcal{C}$  encodes  $S$  into the  $i$ -th IDA piece  $S_i$ .

# The Delayed-Replication Algorithm

Idea:  $\mathcal{C}$  sends a ledger to learner  $\mathsf{L}$

1. Partition the ledger in sequences  $S$  of  $\mathbf{n}$  transactions.
2. Replica  $\mathsf{r}_i \in \mathcal{C}$  encodes  $S$  into the  $i$ -th IDA piece  $S_i$ .
3. Replica  $\mathsf{r}_i \in \mathcal{C}$  sends  $S_i$  with a checksum  $C_i(S)$  of  $S$  to  $\mathsf{L}$ .

# The Delayed-Replication Algorithm

Idea:  $\mathcal{C}$  sends a ledger to learner  $\mathsf{L}$

1. Partition the ledger in sequences  $S$  of  $\mathbf{n}$  transactions.
2. Replica  $\mathsf{r}_i \in \mathcal{C}$  encodes  $S$  into the  $i$ -th IDA piece  $S_i$ .
3. Replica  $\mathsf{r}_i \in \mathcal{C}$  sends  $S_i$  with a checksum  $C_i(S)$  of  $S$  to  $\mathsf{L}$ .
4.  $\mathsf{L}$  receives at least  $\mathbf{n} - \mathbf{f}$  distinct pieces and decodes  $S$ .

# The Delayed-Replication Algorithm

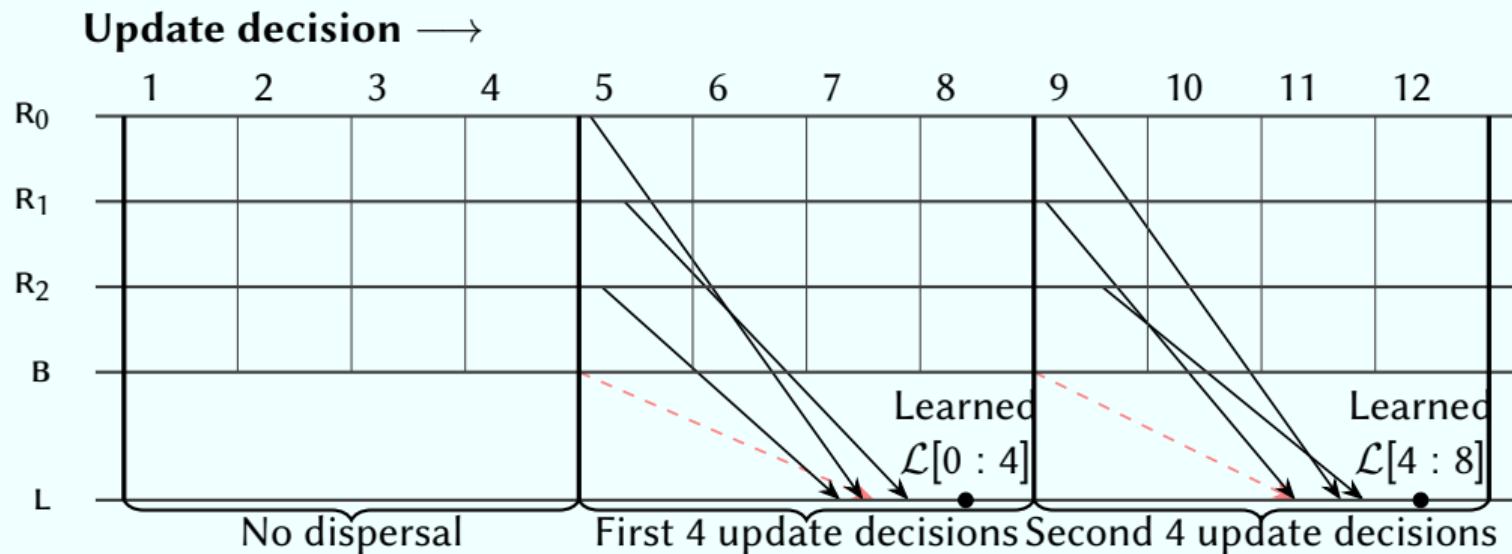
Idea:  $\mathcal{C}$  sends a ledger to learner  $\mathsf{L}$

1. Partition the ledger in sequences  $S$  of  $\mathbf{n}$  transactions.
2. Replica  $\mathsf{R}_i \in \mathcal{C}$  encodes  $S$  into the  $i$ -th IDA piece  $S_i$ .
3. Replica  $\mathsf{R}_i \in \mathcal{C}$  sends  $S_i$  with a checksum  $C_i(S)$  of  $S$  to  $\mathsf{L}$ .
4.  $\mathsf{L}$  receives at least  $\mathbf{n} - \mathbf{f}$  distinct pieces and decodes  $S$ .

Observation ( $\mathbf{n} > 2\mathbf{f}$ )

- ▶ Replica  $\mathsf{R}_i$  sends at most  $B = \left\lceil \frac{\|S\|}{\mathbf{n}-\mathbf{f}} \right\rceil + c \leq \frac{2\|S\|}{\mathbf{n}} + 1 + c = \mathcal{O}\left(\frac{\|S\|}{\mathbf{n}} + c\right)$  bytes.
- ▶ Learner  $\mathsf{L}$  receives at most  $\mathbf{n} \cdot B = \mathcal{O}(\|S\| + c\mathbf{n})$  bytes.

# Communication by the Delayed-Replication Algorithm



## Decoding $S$ Using Simple Checksums ( $n > 2f$ )

## Decoding $S$ Using Simple Checksums ( $n > 2f$ )

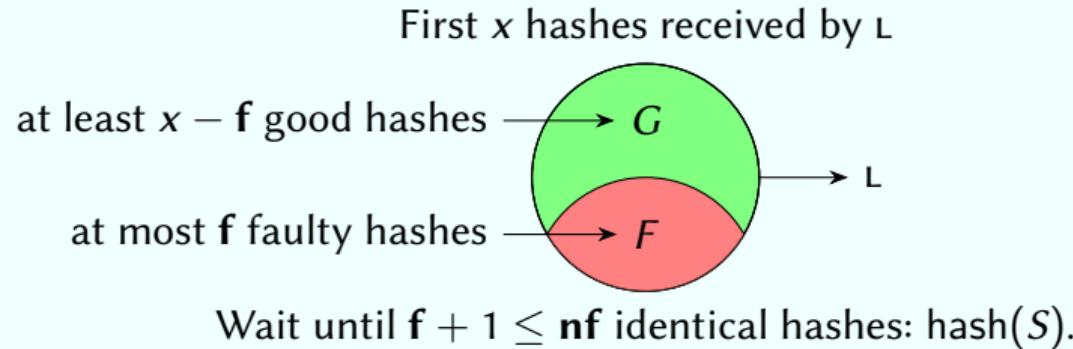
- ▶ Use checksums  $\text{hash}(S)$ .

## Decoding $S$ Using Simple Checksums ( $n > 2f$ )

- ▶ Use checksums  $\text{hash}(S)$ .
- ▶ The  $n - f$  non-faulty replicas will provide correct *pieces*.

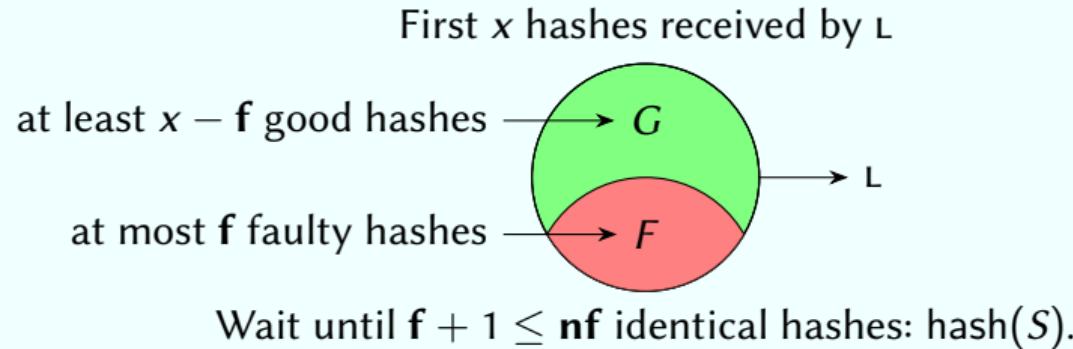
## Decoding $S$ Using Simple Checksums ( $n > 2f$ )

- ▶ Use checksums  $\text{hash}(S)$ .
- ▶ The  $n - f$  non-faulty replicas will provide correct *pieces*.
- ▶ At least  $n - f > f$  messages with correct *checksums*.



## Decoding $S$ Using Simple Checksums ( $n > 2f$ )

- ▶ Use checksums  $\text{hash}(S)$ .
- ▶ The  $n - f$  non-faulty replicas will provide correct *pieces*.
- ▶ At least  $n - f > f$  messages with correct *checksums*.



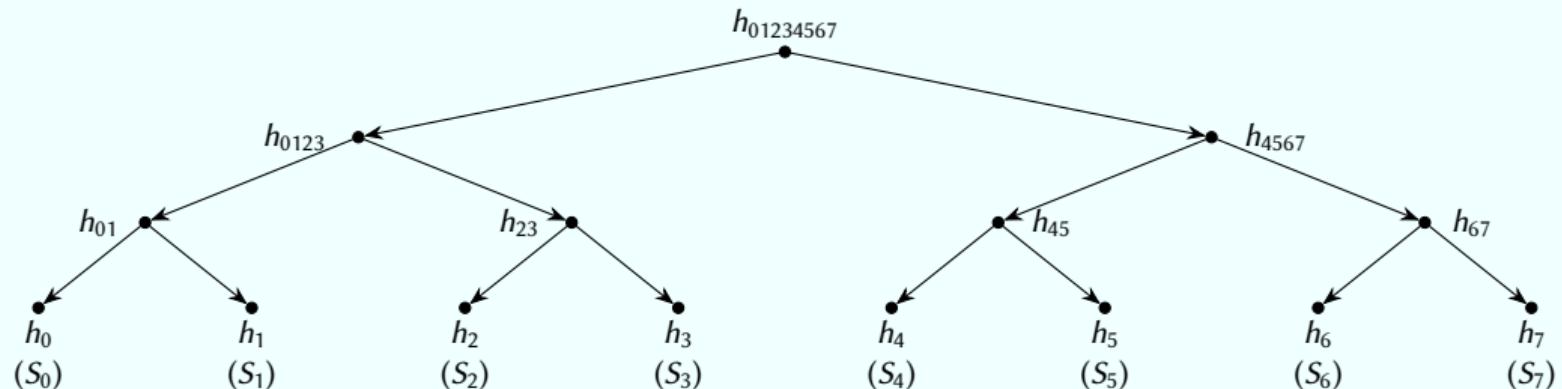
- ▶ Intensive for learners: one can choose  $n - f$  out of  $n$  messages in  $\binom{n}{n-f}$  ways  
*only one such choice is guaranteed to be correct!*

# Decoding $S$ Using Tree Checksums

Use Merkle-trees to construct checksums

Consider 8 replicas and a sequence  $S$ .

We construct the checksum  $C_5(S)$  of  $S$  (used by  $R_5$ ).



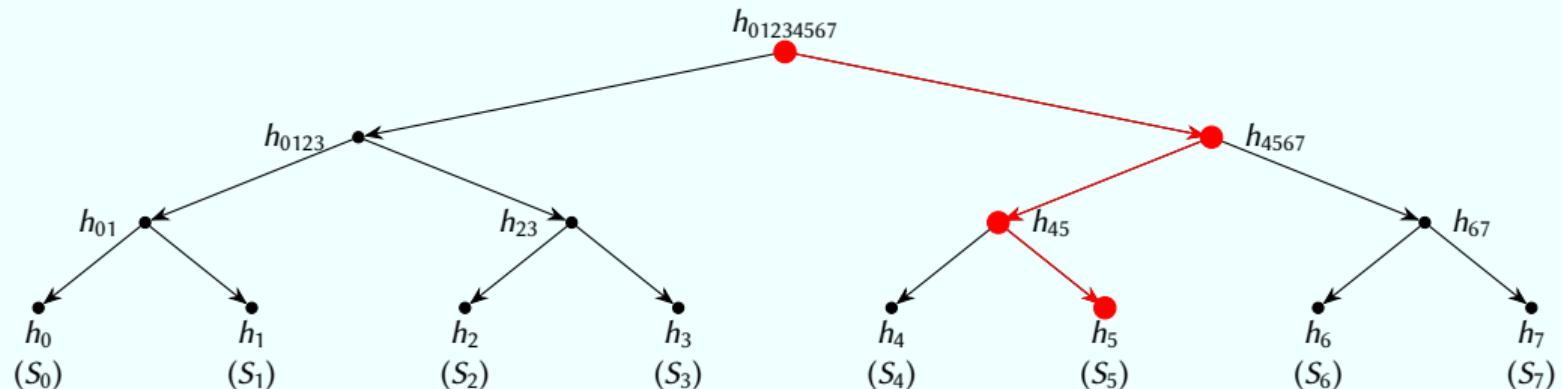
Construct a Merkle tree for pieces  $S_0, \dots, S_7$ .

# Decoding $S$ Using Tree Checksums

Use Merkle-trees to construct checksums

Consider 8 replicas and a sequence  $S$ .

We construct the checksum  $C_5(S)$  of  $S$  (used by  $R_5$ ).



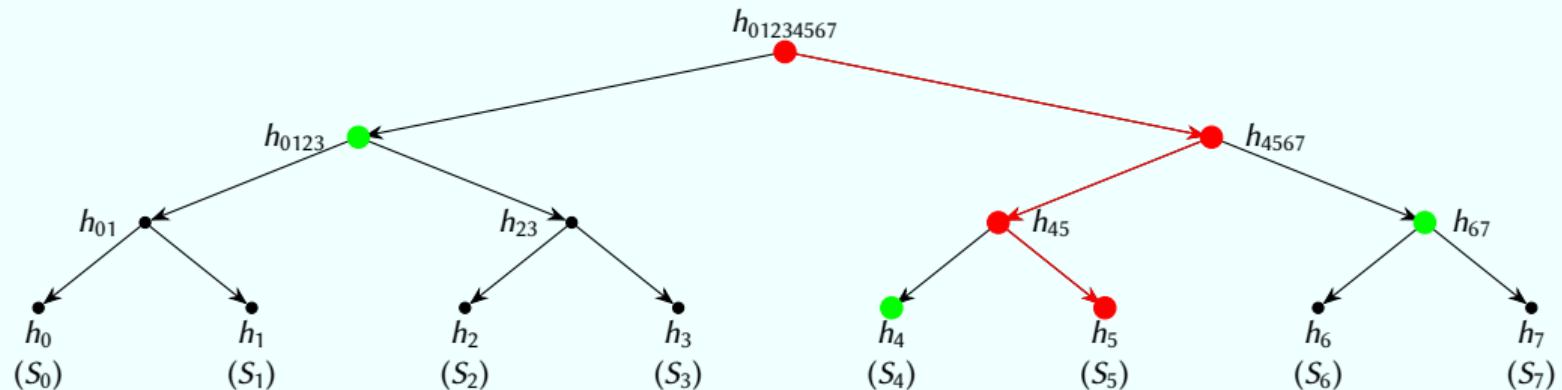
Determine the path from root to  $S_5$ .

# Decoding $S$ Using Tree Checksums

Use Merkle-trees to construct checksums

Consider 8 replicas and a sequence  $S$ .

We construct the checksum  $C_5(S)$  of  $S$  (used by  $R_5$ ).



Select *root* and *neighbors*:  $C_5(S) = [h_4, h_{67}, h_{0123}, h_{01234567}]$ .

## Delayed-Replication: Main Result ( $n > 2f$ )

### Theorem

Consider the learner  $L$ , replica  $R$ , and decided transactions  $\mathcal{T}$ . The delayed-replication algorithm with tree checksums guarantees

1.  $L$  will learn  $\mathcal{T}$ ;
2.  $L$  will receive at most  $|\mathcal{T}|$  messages with a total size of  $\mathcal{O}(|\mathcal{T}| + |\mathcal{T}| \log n)$ ;
3.  $L$  will only need at most  $\frac{|\mathcal{T}|}{n}$  decode steps;
4.  $R$  will send at most  $\frac{|\mathcal{T}|}{n}$  messages to  $L$  of size  $\mathcal{O}\left(\frac{|\mathcal{T}| + |\mathcal{T}| \log n}{n}\right)$ .

## Application: Scalable Storage for Resilient Systems

- ▶ Replicas typically only need the *current data*  $V$  to decide on future updates.
- ▶ Replicas only need the full ledger  $\mathcal{L}$  for *recovery*.
- ▶ We can use *delayed-replication* to reduce the data each replica has to store.

### Theorem

*The storage cost per replica can be reduced from*

$$\mathcal{O}(\|\mathcal{L}\| + \|V\|) \quad \text{to} \quad \mathcal{O}\left(\frac{\|\mathcal{L}\|}{n - f} + \frac{|\mathcal{L}|}{n} \log(n) + \|V\|\right).$$

# Reliable Communication between Fault-Tolerant Clusters

## Definition

Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters, both having non-faulty replicas.

The *cluster-sending problem* is the problem of sending a value  $v$  from  $\mathcal{C}_1$  to  $\mathcal{C}_2$  such that:

1. non-faulty replicas in  $\mathcal{C}_2$  *receive*  $v$ ;
2. non-faulty replicas in  $\mathcal{C}_1$  *confirm* that  $v$  was received by the non-faulty replicas in  $\mathcal{C}_2$ ;
3. replicas in  $\mathcal{C}_2$  only receive  $v$  if all non-faulty replicas in  $\mathcal{C}_1$  *agree* upon sending  $v$ .

# Reliable Communication between Fault-Tolerant Clusters

## Definition

Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters, both having non-faulty replicas.

The *cluster-sending problem* is the problem of sending a value  $v$  from  $\mathcal{C}_1$  to  $\mathcal{C}_2$  such that:

1. non-faulty replicas in  $\mathcal{C}_2$  *receive*  $v$ ;
2. non-faulty replicas in  $\mathcal{C}_1$  *confirm* that  $v$  was received by the non-faulty replicas in  $\mathcal{C}_2$ ;
3. replicas in  $\mathcal{C}_2$  only receive  $v$  if all non-faulty replicas in  $\mathcal{C}_1$  *agree* upon sending  $v$ .

## Informal Definition

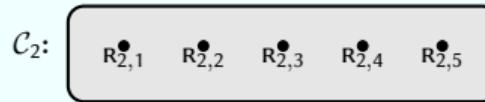
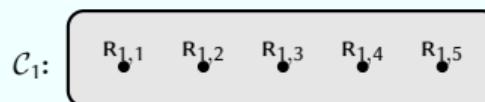
Successfully sending a value  $v$  from a cluster  $\mathcal{C}_1$  to a  $\mathcal{C}_2$  without any faulty replicas being able to *disrupt sending* or send *alternative forged values*.

# Basic Cluster-Sending via Broadcasting

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica in  $\mathcal{C}_1$  has a *certificate*  $\text{cert}(v, \mathcal{C}_1)$  that proves agreement.
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.

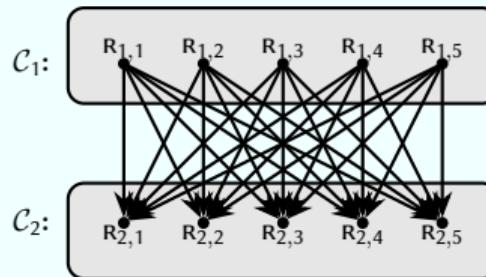


# Basic Cluster-Sending via Broadcasting

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica in  $\mathcal{C}_1$  has a *certificate*  $\text{cert}(v, \mathcal{C}_1)$  that proves agreement.
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



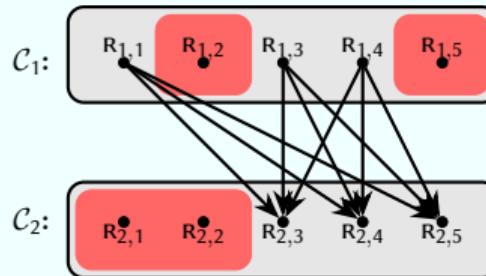
*Broadcast:* every replica in  $\mathcal{C}_1$  sends pairs  $(v, \text{cert}(v, \mathcal{C}_1))$  to every replica in  $\mathcal{C}_2$ .

# Basic Cluster-Sending via Broadcasting

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica in  $\mathcal{C}_1$  has a *certificate*  $\text{cert}(v, \mathcal{C}_1)$  that proves agreement.
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



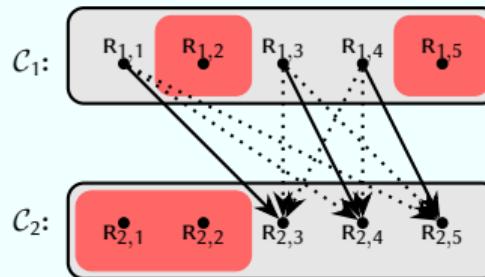
Faulty replicas can *fail* to send (in  $\mathcal{C}_1$ ) or to receive (in  $\mathcal{C}_2$ ).

# Basic Cluster-Sending via Broadcasting

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica in  $\mathcal{C}_1$  has a *certificate*  $\text{cert}(v, \mathcal{C}_1)$  that proves agreement.
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



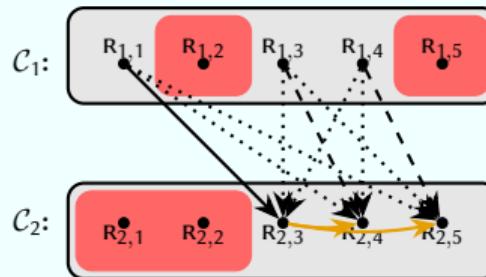
Non-faulty replicas in  $\mathcal{C}_2$  only need at-least one message  $(v, \text{cert}(v, \mathcal{C}_1))$ .

# Basic Cluster-Sending via Broadcasting

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica in  $\mathcal{C}_1$  has a *certificate*  $\text{cert}(v, \mathcal{C}_1)$  that proves agreement.
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



Replicas in  $\mathcal{C}_2$  can redistribute  $(v, \text{cert}(v, \mathcal{C}_1))$ .

# Basic Cluster-Sending via Broadcasting

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica in  $\mathcal{C}_1$  has a *certificate*  $\text{cert}(v, \mathcal{C}_1)$  that proves agreement.
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



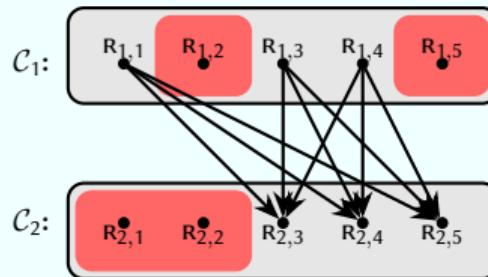
With certificates: a *single* message between non-faulty sender and receiver is sufficient!

# Basic Cluster-Sending via Broadcasting (Without Certificates)

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica  $R \in \mathcal{C}_1$  can only *claim* agreement via a digital signature  $\text{cert}(v, R)$ .
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.

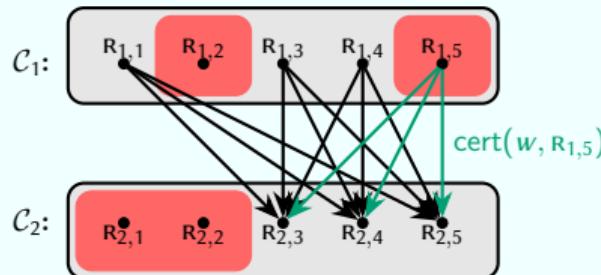


# Basic Cluster-Sending via Broadcasting (Without Certificates)

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica  $r \in \mathcal{C}_1$  can only *claim* agreement via a digital signature  $\text{cert}(v, r)$ .
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



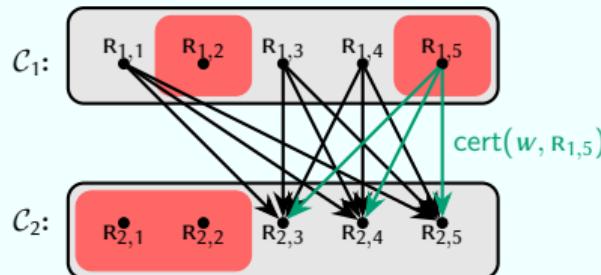
Faulty replicas can *lie* and send  $\text{cert}(w, r)$  without agreement on  $w$ .

# Basic Cluster-Sending via Broadcasting (Without Certificates)

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica  $R \in \mathcal{C}_1$  can only *claim* agreement via a digital signature  $\text{cert}(v, R)$ .
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



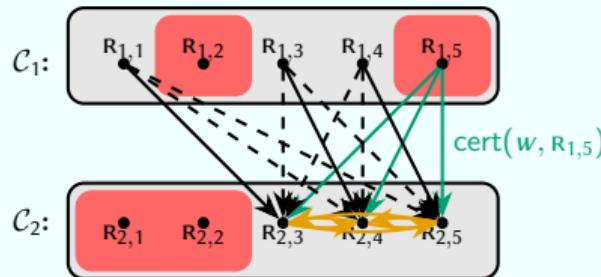
Claims from *three* distinct replicas in  $\mathcal{C}_1$ : at-least one from a non-faulty replica.

# Basic Cluster-Sending via Broadcasting (Without Certificates)

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica  $R \in \mathcal{C}_1$  can only *claim* agreement via a digital signature  $\text{cert}(v, R)$ .
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



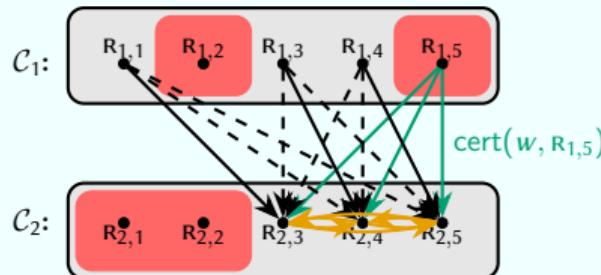
Replicas in  $\mathcal{C}_2$  can redistribute  $(v, \text{cert}(v, R))$ .

# Basic Cluster-Sending via Broadcasting (Without Certificates)

*Goal:* send a value  $v$  from cluster  $\mathcal{C}_1$  to cluster  $\mathcal{C}_2$ .

## Assumptions

- ▶ Every replica  $R \in \mathcal{C}_1$  can only *claim* agreement via a digital signature  $\text{cert}(v, R)$ .
- ▶ Communication is *reliable*.
- ▶ At-most *two* replicas faulty in each cluster.



Without certificates: *at least  $f_{\mathcal{C}_1} + 1$*  distinct received messages by non-faulty senders!

## Efficient Cluster-Sending

Cluster-Sending via broadcasting: straightforward, *not efficient*:

- ▶ With certificates:  $(f_{C_1} + 1)(f_{C_2} + 1) \approx f_{C_1} \times f_{C_2}$  messages.
- ▶ With claims:  $(2f_{C_1} + 1)(f_{C_2} + 1) \approx 2f_{C_1} \times f_{C_2}$  messages.

# Efficient Cluster-Sending

Cluster-Sending via broadcasting: straightforward, *not efficient*:

- ▶ With certificates:  $(f_{C_1} + 1)(f_{C_2} + 1) \approx f_{C_1} \times f_{C_2}$  messages.
- ▶ With claims:  $(2f_{C_1} + 1)(f_{C_2} + 1) \approx 2f_{C_1} \times f_{C_2}$  messages.

## Local communication versus global communication

	Ping round-trip times (ms)						Bandwidth (Mbit/s)					
	OR	IA	Mont.	BE	TW	Syd.	OR	IA	Mont.	BE	TW	Syd.
Oregon	≤ 1	38	65	136	118	161	7998	669	371	194	188	136
Iowa		≤ 1	33	98	153	172		10004	752	243	144	120
Montreal			≤ 1	82	186	202			7977	283	111	102
Belgium				≤ 1	252	270				9728	79	66
Taiwan					≤ 1	137					7998	160
Sydney						≤ 1						7977

*Goal:* Minimize communication *between* clusters.

## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{C_1} = 15$$

$$f_{C_1} = 7$$

$$n_{C_2} = 5$$

$$f_{C_2} = 2$$

Proposition (assuming certificates)

Any correct algorithm needs to send at least 14 messages.



## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{C_1} = 15$$

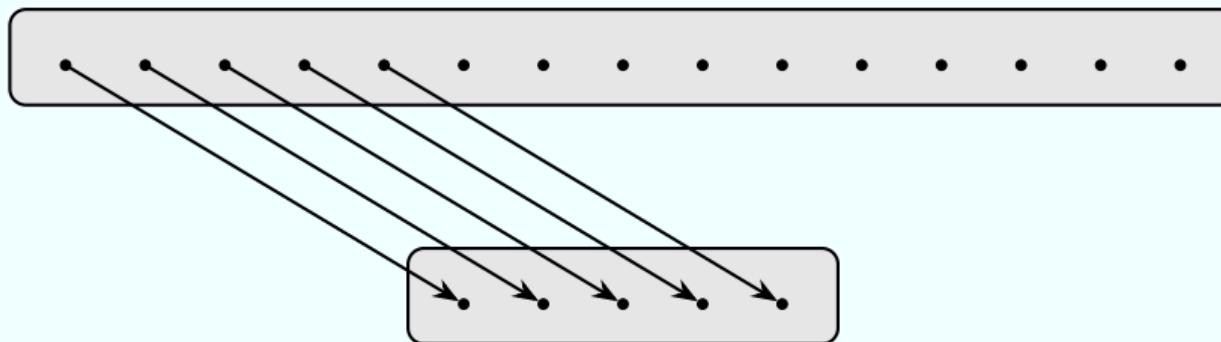
$$f_{C_1} = 7$$

$$n_{C_2} = 5$$

$$f_{C_2} = 2$$

Proposition (assuming certificates)

Any correct algorithm needs to send at least 14 messages.



Minimize impact of faulty replicas: minimum number of messages per participant.

## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{C_1} = 15$$

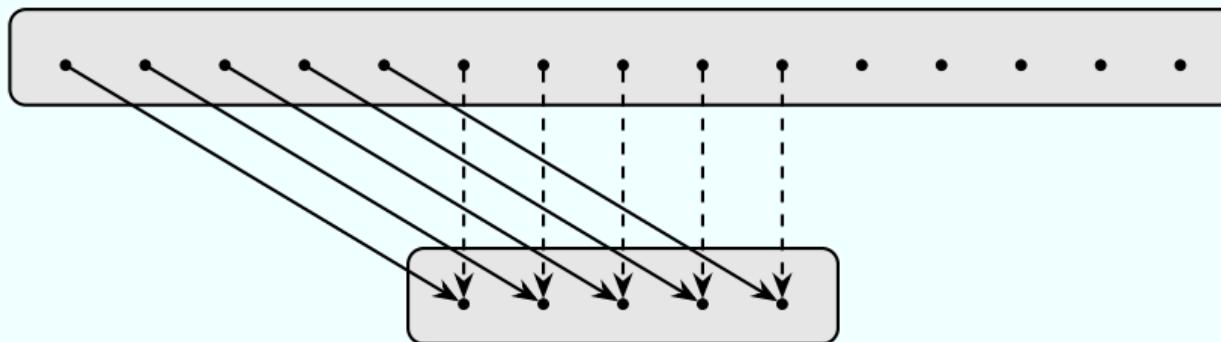
$$f_{C_1} = 7$$

$$n_{C_2} = 5$$

$$f_{C_2} = 2$$

Proposition (assuming certificates)

Any correct algorithm needs to send at least 14 messages.



Minimize impact of faulty replicas: minimum number of messages per participant.

## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{C_1} = 15$$

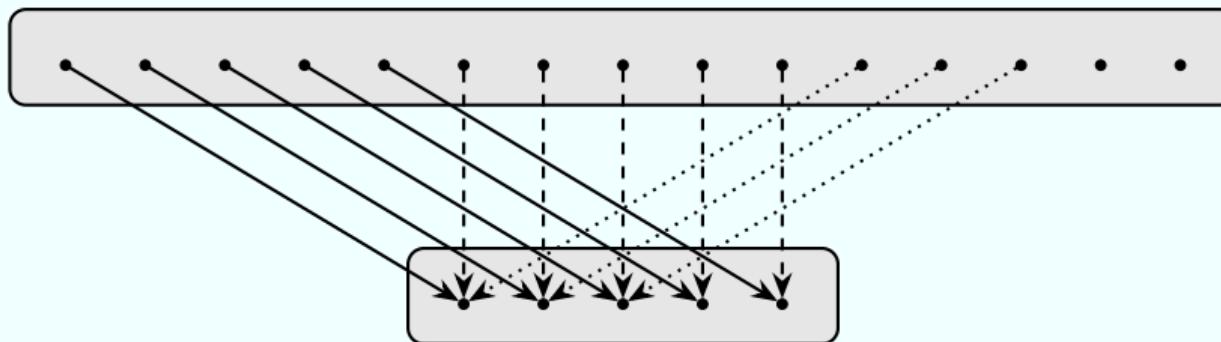
$$f_{C_1} = 7$$

$$n_{C_2} = 5$$

$$f_{C_2} = 2$$

Proposition (assuming certificates)

Any correct algorithm needs to send at least 14 messages.



Minimize impact of faulty replicas: minimum number of messages per participant.

## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{\mathcal{C}_1} = 15$$

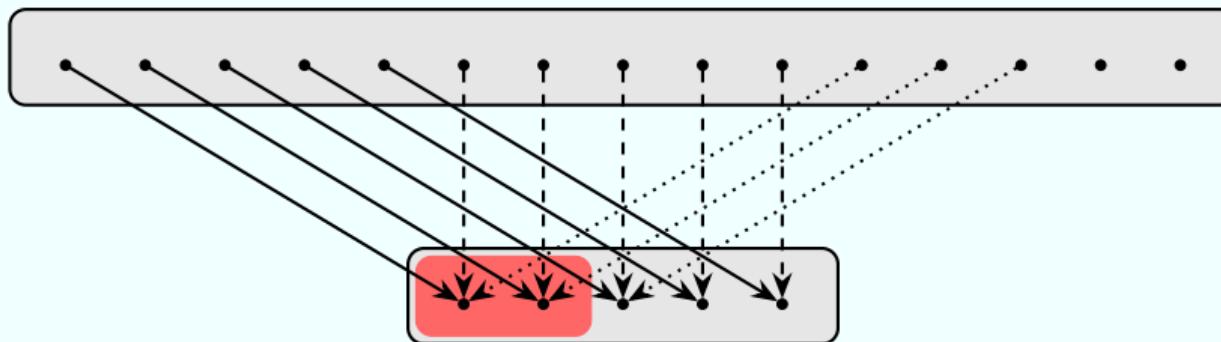
$$f_{\mathcal{C}_1} = 7$$

$$n_{\mathcal{C}_2} = 5$$

$$f_{\mathcal{C}_2} = 2$$

**Proposition (assuming certificates)**

Any correct algorithm needs to send at least *14 messages*.



Any  $f_{\mathcal{C}_2}$  replicas in  $\mathcal{C}_2$  can be faulty: top  $f_{\mathcal{C}_2}$  receivers receive at-least 6 messages.

## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{C_1} = 15$$

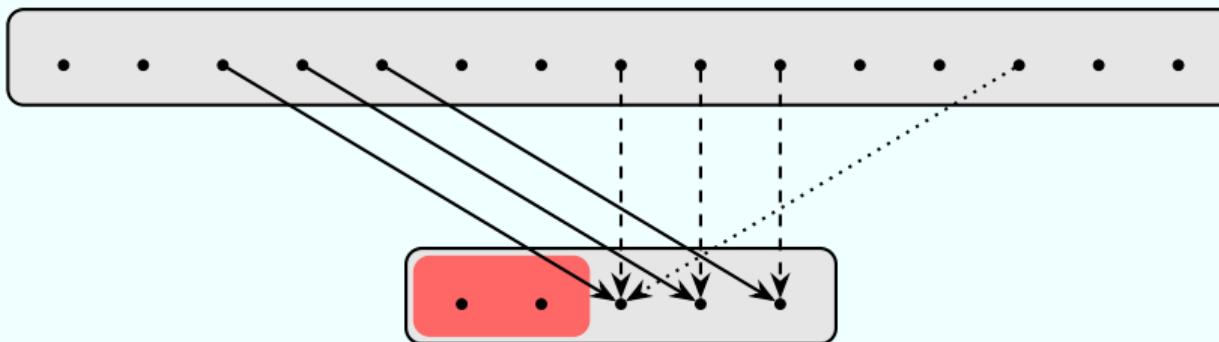
$$n_{C_2} = 5$$

$$f_{C_1} = 7$$

$$f_{C_2} = 2$$

Proposition (assuming certificates)

Any correct algorithm needs to send at least 14 messages.



## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{\mathcal{C}_1} = 15$$

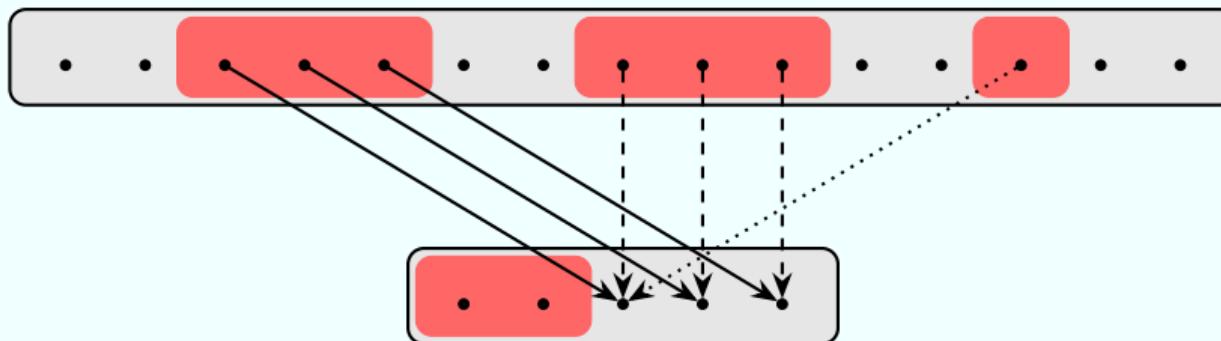
$$f_{\mathcal{C}_1} = 7$$

$$n_{\mathcal{C}_2} = 5$$

$$f_{\mathcal{C}_2} = 2$$

Proposition (assuming certificates)

Any correct algorithm needs to send at least 14 messages.



Only  $f_{\mathcal{C}_1}$  messages remaining, can all be sent by faulty replicas in  $\mathcal{C}_1$ .

## Towards a Lower-Bound for Cluster-Sending (Example)

$$n_{C_1} = 15$$

$$f_{C_1} = 7$$

$$n_{C_2} = 5$$

$$f_{C_2} = 2$$

Proposition (assuming certificates)

Any correct algorithm needs to send at least 14 messages.



# Lower-Bound for Cluster-Sending with Certificates

## Basic Idea

- ▶ One message needs to be exchanged between a non-faulty sender and receiver.
- ▶ Have to deal with size imbalances between clusters.

# Lower-Bound for Cluster-Sending with Certificates

## Basic Idea

- ▶ One message needs to be exchanged between a non-faulty sender and receiver.
- ▶ Have to deal with size imbalances between clusters.

## Theorem

Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters and let  $\{i, j\} = \{1, 2\}$  such that  $\mathbf{n}_{\mathcal{C}_i} \geq \mathbf{n}_{\mathcal{C}_j}$ . Let

$$\begin{aligned} q_i &= (\mathbf{f}_{\mathcal{C}_i} + 1) \operatorname{div} \mathbf{n}_{\mathcal{C}_j}, \\ r_i &= (\mathbf{f}_{\mathcal{C}_i} + 1) \operatorname{mod} \mathbf{n}_{\mathcal{C}_j}, \\ \sigma_i &= q_i \mathbf{n}_{\mathcal{C}_j} + r_i + \mathbf{f}_{\mathcal{C}_j} \operatorname{sgn} r_i. \end{aligned}$$

Any protocol that solves the cluster-sending problem in which  $\mathcal{C}_1$  sends a value  $v$  to  $\mathcal{C}_2$  needs to exchange at least  $\sigma_i$  messages.

# Lower-Bound for Cluster-Sending with Certificates (Example)

## Theorem

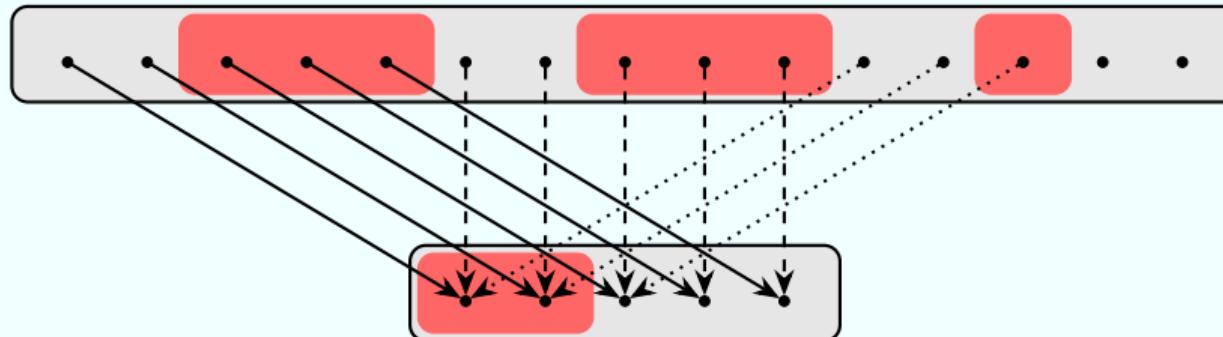
Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters and let

$$q_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ div } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ div } 3 = 2,$$

$$r_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ mod } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ mod } 3 = 1,$$

$$\sigma_1 = q_1 \mathbf{n}_{\mathcal{C}_2} + r_1 + \mathbf{f}_{\mathcal{C}_2} \text{ sgn } r_1 = 2 \cdot 5 + 1 + 3 = 14.$$

Any protocol that solves the cluster-sending problem in which  $\mathcal{C}_1$  sends a value  $v$  to  $\mathcal{C}_2$  needs to exchange at least  $\sigma_1 = 14$  messages.



# Lower-Bound for Cluster-Sending with Certificates (Example)

## Theorem

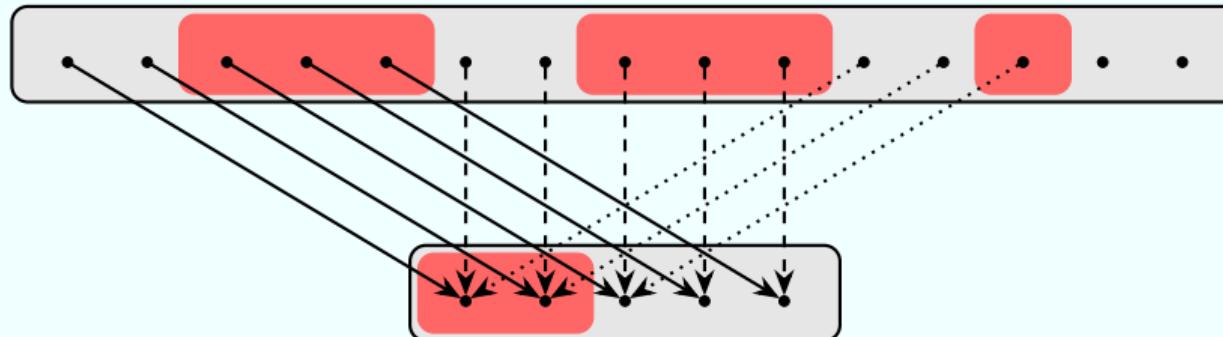
Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters and let

$$q_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ div } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ div } 3 = 2,$$

$$r_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ mod } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ mod } 3 = 1,$$

$$\sigma_1 = \boxed{q_1 \mathbf{n}\mathbf{f}_{\mathcal{C}_2}} + r_1 + \mathbf{f}_{\mathcal{C}_2} \text{ sgn } r_1 = \boxed{2 \cdot 5} + 1 + 3 = 14.$$

Any protocol that solves the cluster-sending problem in which  $\mathcal{C}_1$  sends a value  $v$  to  $\mathcal{C}_2$  needs to exchange at least  $\sigma_1 = 14$  messages.



# Lower-Bound for Cluster-Sending with Certificates (Example)

## Theorem

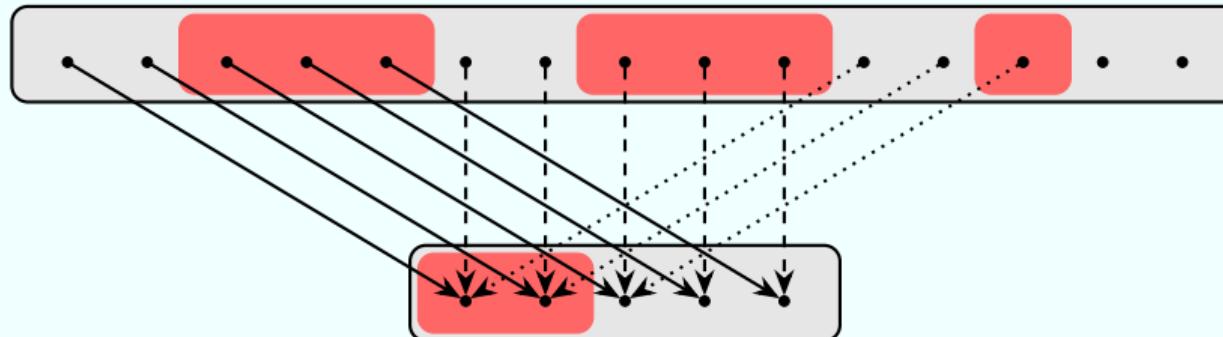
Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters and let

$$q_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ div } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ div } 3 = 2,$$

$$r_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ mod } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ mod } 3 = 1,$$

$$\sigma_1 = q_1 \mathbf{n}_{\mathcal{C}_2} + [r_1 + \mathbf{f}_{\mathcal{C}_2} \text{ sgn } r_1] = 2 \cdot 5 + [1 + 3] = 14.$$

Any protocol that solves the cluster-sending problem in which  $\mathcal{C}_1$  sends a value  $v$  to  $\mathcal{C}_2$  needs to exchange at least  $\sigma_1 = 14$  messages.



# Lower-Bound for Cluster-Sending with Certificates (Example)

## Theorem

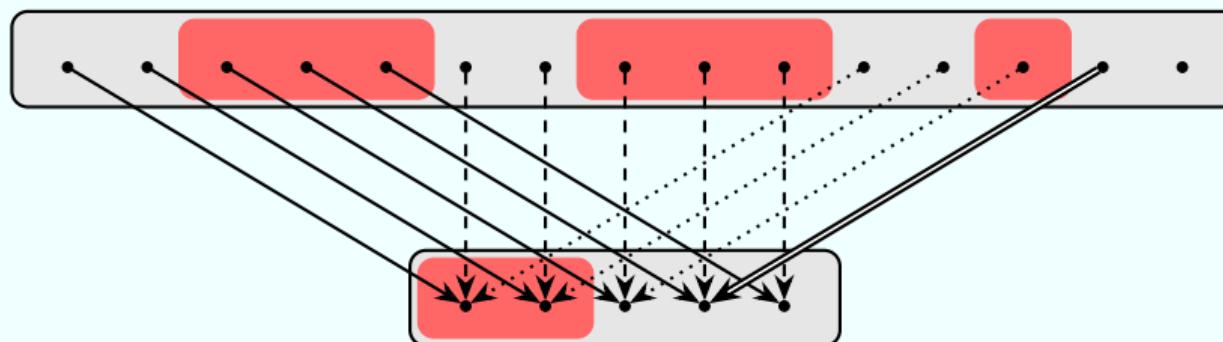
Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters and let

$$q_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ div } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ div } 3 = 2,$$

$$r_1 = (\mathbf{f}_{\mathcal{C}_1} + 1) \text{ mod } \mathbf{n}\mathbf{f}_{\mathcal{C}_2} = 7 \text{ mod } 3 = 1,$$

$$\sigma_1 = q_1 \mathbf{n}_{\mathcal{C}_2} + r_1 + \mathbf{f}_{\mathcal{C}_2} \text{ sgn } r_1 = 2 \cdot 5 + 1 + 3 = 14.$$

Any protocol that solves the cluster-sending problem in which  $\mathcal{C}_1$  sends a value  $v$  to  $\mathcal{C}_2$  needs to exchange at least  $\sigma_1 = 14$  messages.



# Lower-Bound for Cluster-Sending with Claims

## Basic Idea

- ▶  $\mathbf{f}_{\mathcal{C}_1} + 1$  message needs to be sent by distinct non-faulty senders to non-faulty receiver.
- ▶ Have to deal with size imbalances between clusters.

## Theorem

Let  $\mathcal{C}_1, \mathcal{C}_2$  be two clusters and let  $\{i, j\} = \{1, 2\}$  such that  $\mathbf{n}_{\mathcal{C}_i} \geq \mathbf{n}_{\mathcal{C}_j}$ . Let

$$\begin{array}{ll} q_1 = (2\mathbf{f}_{\mathcal{C}_1} + 1) \text{ div } \mathbf{n}_{\mathcal{C}_2}, & q_2 = (\mathbf{f}_{\mathcal{C}_2} + 1) \text{ div } (\mathbf{n}_{\mathcal{C}_1} - \mathbf{f}_{\mathcal{C}_1}) \\ r_1 = (2\mathbf{f}_{\mathcal{C}_1} + 1) \text{ mod } \mathbf{n}_{\mathcal{C}_2}, & r_2 = (\mathbf{f}_{\mathcal{C}_2} + 1) \text{ mod } (\mathbf{n}_{\mathcal{C}_1} - \mathbf{f}_{\mathcal{C}_1}) \\ \tau_1 = q_1 \mathbf{n}_{\mathcal{C}_2} + r_1 + \mathbf{f}_{\mathcal{C}_2} \operatorname{sgn} r_1 & \tau_2 = q_2 \mathbf{n}_{\mathcal{C}_1} + r_2 + 2\mathbf{f}_{\mathcal{C}_1} \operatorname{sgn} r_2. \end{array}$$

Any protocol that solves the cluster-sending problem in which  $\mathcal{C}_1$  sends a value  $v$  to  $\mathcal{C}_2$  needs to exchange at least  $\tau_i$  messages.

# Bijective Sending with Certificates

Assume  $f_{\mathcal{C}_1} + f_{\mathcal{C}_2} + 1 \leq \min(n_{\mathcal{C}_1}, n_{\mathcal{C}_2})$ .

We have  $\sigma_1 = \sigma_2 = f_{\mathcal{C}_1} + f_{\mathcal{C}_2} + 1$ .

## Protocol for the sending cluster $\mathcal{C}_1$ :

- 1: All replicas in  $\mathcal{G}_{\mathcal{C}_1}$  agree on  $v$  and construct  $\text{cert}(v, \mathcal{C}_1)$ .
- 2: Choose replicas  $S_1 \subseteq \mathcal{C}_1$  and  $S_2 \subseteq \mathcal{C}_2$  with  $n_{S_2} = n_{S_1} = f_{\mathcal{C}_1} + f_{\mathcal{C}_2} + 1$ .
- 3: Choose a bijection  $b : S_1 \rightarrow S_2$ .
- 4: **for**  $r_1 \in S_1$  **do**
- 5:    $r_1$  sends  $(v, \text{cert}(v, \mathcal{C}_1))$  to  $b(r_1)$ .

## Protocol for the receiving cluster $\mathcal{C}_2$ :

- 6: **event**  $r_2 \in \mathcal{G}_{\mathcal{C}_2}$  receives  $(w, \text{cert}(w, \mathcal{C}_1))$  from  $r_1 \in \mathcal{C}_1$  **do**
- 7:   Broadcast  $(w, \text{cert}(w, \mathcal{C}_1))$  to all replicas in  $\mathcal{C}_2$ .
- 8: **event**  $r'_2 \in \mathcal{G}_{\mathcal{C}_2}$  receives  $(w, \text{cert}(w, \mathcal{C}_1))$  from  $r_2 \in \mathcal{C}_2$  **do**
- 9:    $r'_2$  considers  $w$  *received*.

## Bijective Sending with Certificates: Example

$$n_{\mathcal{C}_1} = 8$$

$$f_{\mathcal{C}_1} = 3$$

$$n_{\mathcal{C}_2} = 7$$

$$f_{\mathcal{C}_2} = 2$$

$$\sigma_1 = 6.$$

$\mathcal{C}_1:$

$R_{1,1}$     $R_{1,2}$     $R_{1,3}$     $R_{1,4}$     $R_{1,5}$     $R_{1,6}$     $R_{1,7}$     $R_{1,8}$

$\mathcal{C}_2:$

$R_{2,1}$     $R_{2,2}$     $R_{2,3}$     $R_{2,4}$     $R_{2,5}$     $R_{2,6}$     $R_{2,7}$

## Bijective Sending with Certificates: Example

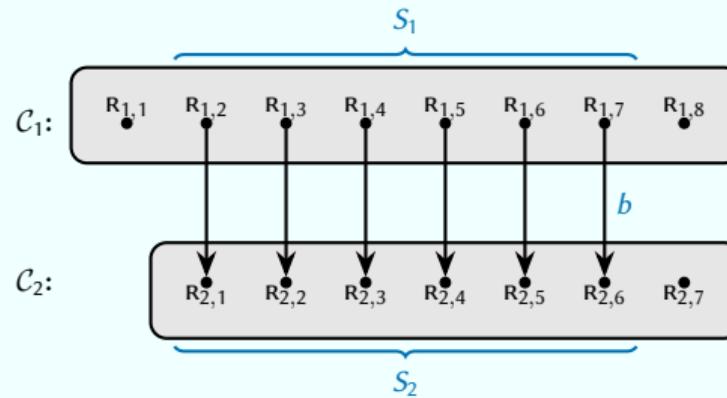
$$n_{\mathcal{C}_1} = 8$$

$$f_{\mathcal{C}_1} = 3$$

$$n_{\mathcal{C}_2} = 7$$

$$f_{\mathcal{C}_2} = 2$$

$$\sigma_1 = 6.$$



## Bijective Sending with Certificates: Example

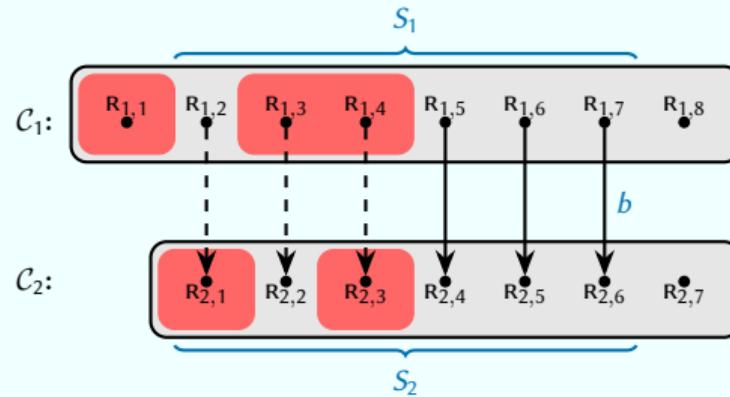
$$n_{\mathcal{C}_1} = 8$$

$$f_{\mathcal{C}_1} = 3$$

$$n_{\mathcal{C}_2} = 7$$

$$f_{\mathcal{C}_2} = 2$$

$$\sigma_1 = 6.$$



## Bijective Sending with Claims

Assume  $2\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1 \leq \min(\mathbf{n}_{\mathcal{C}_1}, \mathbf{n}_{\mathcal{C}_2})$ .

We have  $\tau_1 = \tau_2 = 2\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1$ .

### **Protocol for the sending cluster $\mathcal{C}_1$ :**

- 1: All replicas in  $\mathcal{G}_{\mathcal{C}_1}$  agree on  $v$ .
- 2: Choose replicas  $S_1 \subseteq \mathcal{C}_1$  and  $S_2 \subseteq \mathcal{C}_2$  with  $\mathbf{n}_{S_2} = \mathbf{n}_{S_1} = 2\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1$ .
- 3: Choose bijection  $b : S_1 \rightarrow S_2$ .
- 4: **for**  $r_1 \in S_1$  **do**
- 5:      $r_1$  sends  $(v, \text{cert}(v, r_1))$  to  $b(r_1)$ .

### **Protocol for the receiving cluster $\mathcal{C}_2$ :**

- 6: ....

# Bijective Sending with Claims

Assume  $2f_{\mathcal{C}_1} + f_{\mathcal{C}_2} + 1 \leq \min(n_{\mathcal{C}_1}, n_{\mathcal{C}_2})$ .

We have  $\tau_1 = \tau_2 = 2f_{\mathcal{C}_1} + f_{\mathcal{C}_2} + 1$ .

## Protocol for the sending cluster $\mathcal{C}_1$ :

1: ....

## Protocol for the receiving cluster $\mathcal{C}_2$ :

- 6: **event**  $R_2 \in \mathcal{G}_{\mathcal{C}_2}$  receives  $(w, \text{cert}(w, R'_1))$  from  $R'_1 \in \mathcal{C}_1$  **do**
- 7: Broadcast  $(w, \text{cert}(w, R'_1))$  to all replicas in  $\mathcal{C}_2$ .
- 8: **event**  $R'_2 \in \mathcal{G}_{\mathcal{C}_2}$  receives  $f_{\mathcal{C}_1} + 1$  messages  $(w, \text{cert}(w, R'_1))$ :
  - (i) each message is sent by a replica in  $\mathcal{C}_2$ ;
  - (ii) each message carries the same value  $w$ ; and
  - (iii) each message has a distinct signature  $\text{cert}(w, R'_1)$ ,  $R'_1 \in \mathcal{C}_1$**do**
- 9:  $R'_2$  considers  $w$  *received*.

## Generalizing Bijective Sending

Consider bijective sending from  $\mathcal{C}_1$  to  $\mathcal{C}_2$ ,  $\mathbf{n}_{\mathcal{C}_1} \geq \sigma_1 > \mathbf{n}_{\mathcal{C}_2}$ , with certificates.

- ▶ Bijective sending requires  $\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1$  distinct replicas in both clusters.
- ▶ Restrictive: clusters of roughly the same size.

# Generalizing Bijective Sending

Consider bijective sending from  $\mathcal{C}_1$  to  $\mathcal{C}_2$ ,  $\mathbf{n}_{\mathcal{C}_1} \geq \sigma_1 > \mathbf{n}_{\mathcal{C}_2}$ , with certificates.

- ▶ Bijective sending requires  $\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1$  distinct replicas in both clusters.
- ▶ Restrictive: clusters of roughly the same size.

## Generalize bijective sending



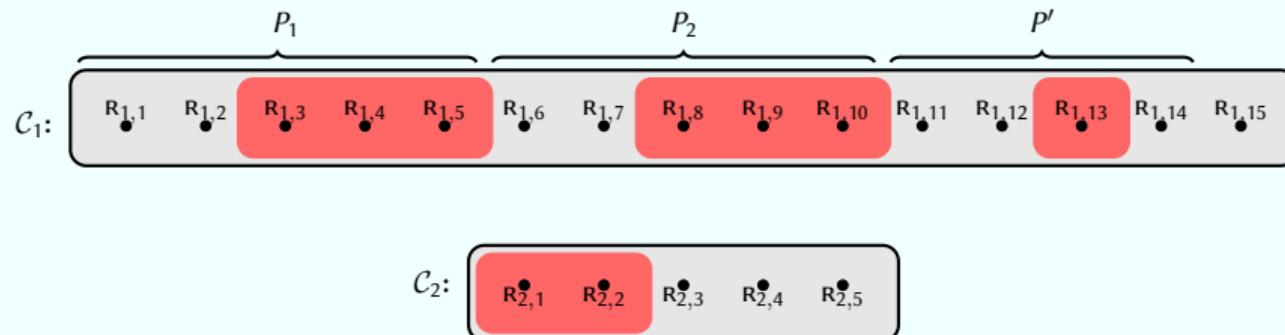
# Generalizing Bijective Sending

Consider bijective sending from  $\mathcal{C}_1$  to  $\mathcal{C}_2$ ,  $\mathbf{n}_{\mathcal{C}_1} \geq \sigma_1 > \mathbf{n}_{\mathcal{C}_2}$ , with certificates.

- ▶ Bijective sending requires  $\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1$  distinct replicas in both clusters.
- ▶ Restrictive: clusters of roughly the same size.

## Generalize bijective sending

- ▶ Partition  $\sigma_1$  replicas of  $\mathcal{C}_1$  into  $\mathbf{n}_{\mathcal{C}_2}$ -sized clusters.



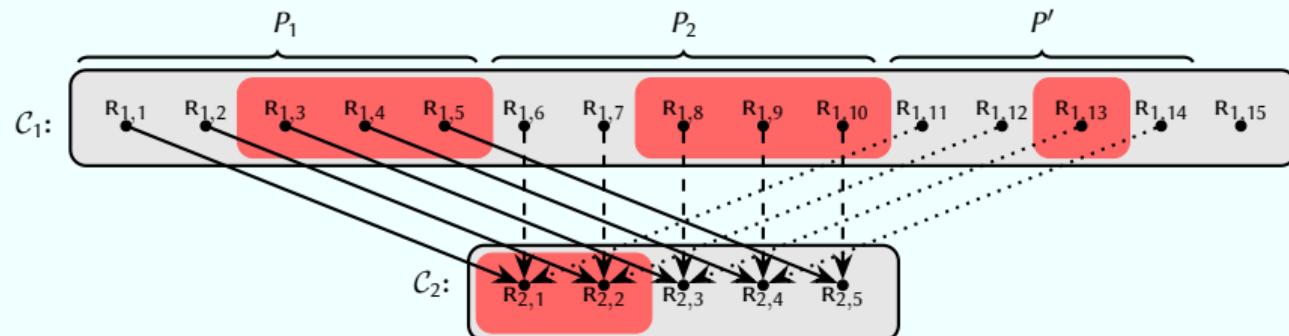
# Generalizing Bijective Sending

Consider bijective sending from  $\mathcal{C}_1$  to  $\mathcal{C}_2$ ,  $\mathbf{n}_{\mathcal{C}_1} \geq \sigma_1 > \mathbf{n}_{\mathcal{C}_2}$ , with certificates.

- ▶ Bijective sending requires  $\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1$  distinct replicas in both clusters.
- ▶ Restrictive: clusters of roughly the same size.

## Generalize bijective sending

- ▶ Partition  $\sigma_1$  replicas of  $\mathcal{C}_1$  into  $\mathbf{n}_{\mathcal{C}_2}$ -sized clusters.
- ▶ Bijective send from each cluster in the partition to  $\mathcal{C}_2$ .



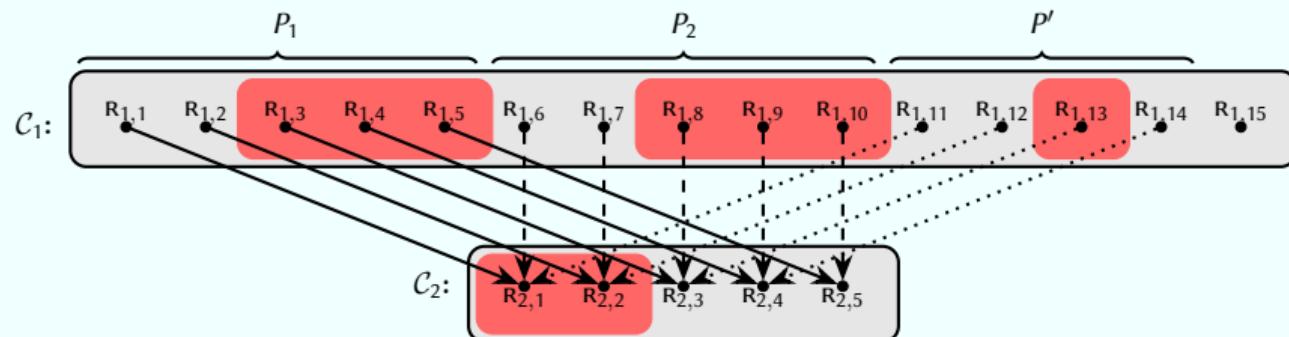
# Generalizing Bijective Sending

Consider bijective sending from  $\mathcal{C}_1$  to  $\mathcal{C}_2$ ,  $\mathbf{n}_{\mathcal{C}_1} \geq \sigma_1 > \mathbf{n}_{\mathcal{C}_2}$ , with certificates.

- ▶ Bijective sending requires  $\mathbf{f}_{\mathcal{C}_1} + \mathbf{f}_{\mathcal{C}_2} + 1$  distinct replicas in both clusters.
- ▶ Restrictive: clusters of roughly the same size.

## Generalize bijective sending

- ▶ Partition  $\sigma_1$  replicas of  $\mathcal{C}_1$  into  $\mathbf{n}_{\mathcal{C}_2}$ -sized clusters.
- ▶ Bijective send from each cluster in the partition to  $\mathcal{C}_2$ .
- ▶  $\mathbf{n}_{\mathcal{C}_1} \geq \sigma_1$  holds always if  $\mathbf{n}_{\mathcal{C}_1} > 3\mathbf{f}_{\mathcal{C}_1}$  and  $\mathbf{n}_{\mathcal{C}_2} > 3\mathbf{f}_{\mathcal{C}_2}$ .



# Partitioned Bijective Sending

## Corollary

Consider the cluster-sending problem in which  $\mathcal{C}_1$  sends a value  $v$  to  $\mathcal{C}_2$ .

1. If  $n_{\mathcal{C}} > 3f_{\mathcal{C}}$  for all clusters  $\mathcal{C}$  and replicas only have crash failures or omit failures, then (partitioned) bijective sending solves cluster-sending with optimal message complexity.
2. If  $n_{\mathcal{C}} > 3f_{\mathcal{C}}$  for all clusters  $\mathcal{C}$  and clusters can produce certificates, then (partitioned) bijective sending solves cluster-sending with optimal message complexity.
3. If  $n_{\mathcal{C}} > 4f_{\mathcal{C}}$  for all clusters  $\mathcal{C}$  and replicas can digitally sign claims, then (partitioned) bijective sending solves cluster-sending with optimal message complexity.

These protocols solve cluster-sending using  $\mathcal{O}(\max(n_{\mathcal{C}_1}, n_{\mathcal{C}_2}))$  messages of size  $\mathcal{O}(\|v\|)$  each.

## Cluster-sending: Can we do Better?

### Pessimistic

**No:** these algorithms are worst-case optimal.

Cannot do better than *linear communication* in the size of the clusters.

# Cluster-sending: Can we do Better?

## Pessimistic

**No:** these algorithms are worst-case optimal.

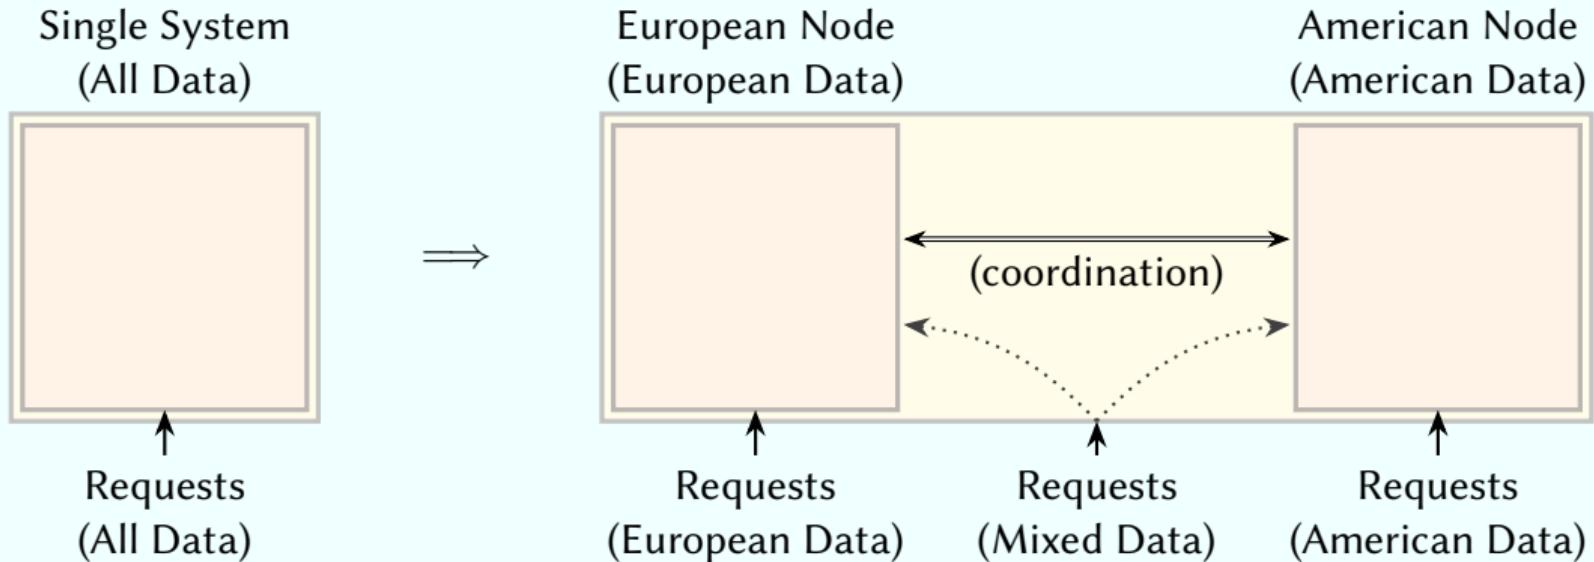
Cannot do better than *linear communication* in the size of the clusters.

## Probabilistic

**Yes:** if we randomly choose sender and receiver, then we often do much better!

Probabilistic approach: expected-case only *constant communication* (four steps).

# Motivation: High-Performance Resilient Systems



Partition the system: More storage and *potentially* more performance.

Potentially *lower latencies* if data ends up closer to users.

Adding shards  $\implies$  adding throughput (parallel processing), adding storage.

# Motivation: High-Performance Resilient Systems

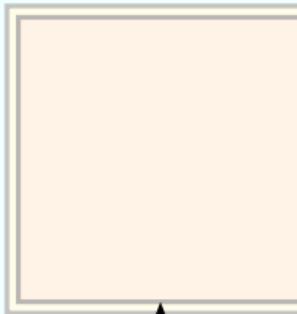
Single System  
(All Data)



Requests  
(All Data)

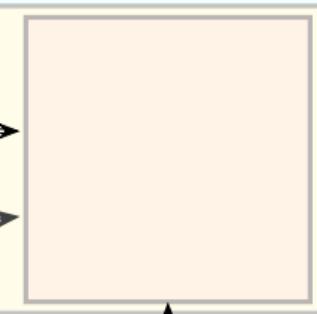


European Node  
(European Data)



Requests  
(European Data)

American Node  
(American Data)



(coordination)

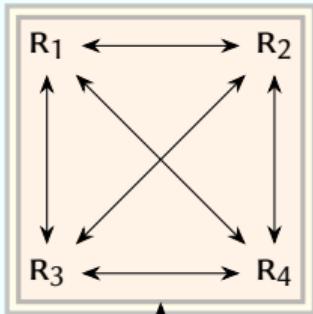
Requests  
(Mixed Data)

Requests  
(American Data)

Resilient system

# Motivation: High-Performance Resilient Systems

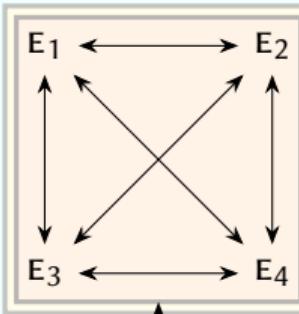
Single System  
(All Data)



Requests  
(All Data)



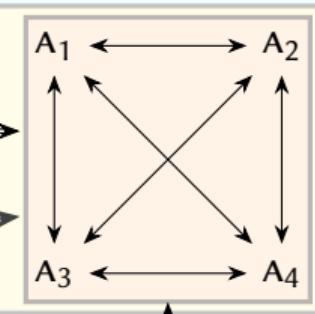
European Node  
(European Data)



Requests  
(European Data)

(coordination)

American Node  
(American Data)



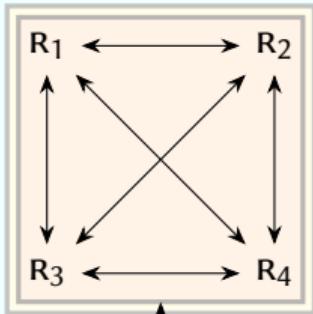
Requests  
(American Data)

## Resilient system

- ▶ Individual shards are consensus-operated *blockchains*.

# Motivation: High-Performance Resilient Systems

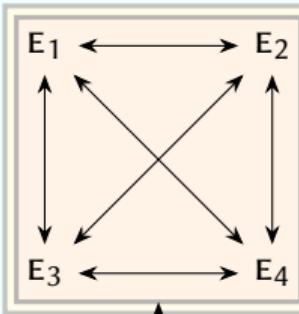
Single System  
(All Data)



Requests  
(All Data)



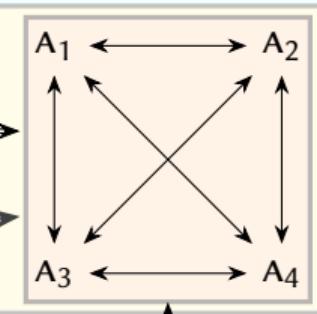
European Node  
(European Data)



Requests  
(European Data)

*cluster-sending*  
(coordination)

American Node  
(American Data)



Requests  
(American Data)

## Resilient system

- ▶ Individual shards are consensus-operated *blockchains*.
- ▶ Communication between shards via *cluster-sending*.

# Transactions

A user interaction with a DBMS: *transaction*.

## Definition

A *transaction* is any one execution of a user program in a DBMS:  
the basic unit of change as seen by the DBMS.

# Transactions

A user interaction with a DBMS: *transaction*.

## Definition

A *transaction* is any one execution of a user program in a DBMS:  
the basic unit of change as seen by the DBMS.

A transaction can be

- ▶ a single query;

# Transactions

A user interaction with a DBMS: *transaction*.

## Definition

A *transaction* is any one execution of a user program in a DBMS:  
the basic unit of change as seen by the DBMS.

A transaction can be

- ▶ a single query;
- ▶ a set of queries;

# Transactions

A user interaction with a DBMS: *transaction*.

## Definition

A *transaction* is any one execution of a user program in a DBMS:  
the basic unit of change as seen by the DBMS.

A transaction can be

- ▶ a single query;
- ▶ a set of queries;
- ▶ a interactive dialog between DBMS and program;
- ▶ ....

# The ACID Properties

Contract between a DBMS and its users.

# The ACID Properties

Contract between a DBMS and its users.

Given a *transaction*  $\tau$ , a DBMS maintains

**Atomicity.** Either all or none of the operations of  $\tau$  are reflected in the database.

**Consistency** Execution of  $\tau$  in *isolation* preserves data consistency.

E.g., integrity constraints—this is *stronger* than CAP-Consistency.

**Isolation**  $\tau$  is “unaware” of other transactions executing concurrently

“As-if” all transactions are executed in a *sequential order*.

**Durability** After  $\tau$  completes successfully, the changes  $\tau$  made persist.

If  $\tau$  fails, then *no* changes persist due to atomicity.

# The ACID Properties

Contract between a DBMS and its users.

Given a *transaction*  $\tau$ , a DBMS maintains

**Atomicity.** Either all or none of the operations of  $\tau$  are reflected in the database.

**Consistency** Execution of  $\tau$  in *isolation* preserves data consistency.

E.g., integrity constraints—this is *stronger* than CAP-Consistency.

**Isolation**  $\tau$  is “unaware” of other transactions executing concurrently  
“As-if” all transactions are executed in a *sequential order*.

**Durability** After  $\tau$  completes successfully, the changes  $\tau$  made persist.

If  $\tau$  fails, then *no* changes persist due to atomicity.

Assumption: individual transactions *make sense* (do not violate consistency).

# The ACID Properties

Contract between a DBMS and its users.

Given a *transaction*  $\tau$ , a DBMS maintains

**Atomicity.** Either all or none of the operations of  $\tau$  are reflected in the database.

**Consistency** Execution of  $\tau$  in *isolation* preserves data consistency.

E.g., integrity constraints—this is *stronger* than CAP-Consistency.

**Isolation**  $\tau$  is “unaware” of other transactions executing concurrently  
“As-if” all transactions are executed in a *sequential order*.

**Durability** After  $\tau$  completes successfully, the changes  $\tau$  made persist.

If  $\tau$  fails, then *no* changes persist due to atomicity.

Assumption: individual transactions *make sense* (do not violate consistency).

Durability is *strong*: crashing or killing the DBMS program, power outage, ....

# The ACID Properties

Contract between a DBMS and its users.

Given a *transaction*  $\tau$ , a DBMS maintains

**Atomicity.** Either all or none of the operations of  $\tau$  are reflected in the database.

**Consistency** Execution of  $\tau$  in *isolation* preserves data consistency.

E.g., integrity constraints—this is *stronger* than CAP-Consistency.

**Isolation**  $\tau$  is “unaware” of other transactions executing concurrently  
“As-if” all transactions are executed in a *sequential order*.

**Durability** After  $\tau$  completes successfully, the changes  $\tau$  made persist.

If  $\tau$  fails, then *no* changes persist due to atomicity.

Assumption: individual transactions *make sense* (do not violate consistency).

Durability is *strong*: crashing or killing the DBMS program, power outage, ....

Typical assumption: *storage* is permanent & reliable.

## Background on Resilience

Consider a transaction  $\tau$  requested by client c in a resilient system.

## Background on Resilience

Consider a transaction  $\tau$  requested by client c in a resilient system.

$\tau$  is processed in *five* steps

1.  $\tau$  needs to be *received* by the system;
2.  $\tau$  must be *replicated* among all replicas in the system;
3. the replicas need to agree on an *execution order* for  $\tau$ ;
4. the replicas each need to *execute*  $\tau$  and *update* their current state accordingly;
5. the client c needs to be *informed* about the result.

## Background on Resilience

Consider a transaction  $\tau$  requested by client c in a resilient system.

$\tau$  is processed in *five* steps

1.  $\tau$  needs to be *received* by the system;
2.  $\tau$  must be *replicated* among all replicas in the system;
3. the replicas need to agree on an *execution order* for  $\tau$ ;
4. the replicas each need to *execute*  $\tau$  and *update* their current state accordingly;
5. the client c needs to be *informed* about the result.

### Non-sharded resilient systems

- ▶ Consensus solves all of the above.
- ▶ In particular *replication order* is *execution order*.
- ▶ Consecutive execution guarantees ACID.

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

Ana	\$0
Bo	\$0
Elisa	\$0

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

Ana	\$0
Bo	\$0
Elisa	\$0

$\xrightarrow{\tau_1}$

Ana	\$500
Bo	\$0
Elisa	\$0

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

Ana	\$0
Bo	\$0
Elisa	\$0

$\xrightarrow{\tau_1}$

Ana	\$500
Bo	\$0
Elisa	\$0

$\xrightarrow{\tau_2}$

Ana	\$500
Bo	\$200
Elisa	\$300

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

Ana	\$0
Bo	\$0
Elisa	\$0

$\xrightarrow{\tau_1}$

Ana	\$500
Bo	\$0
Elisa	\$0

$\xrightarrow{\tau_2}$

Ana	\$500
Bo	\$200
Elisa	\$300

$\xrightarrow{\tau_3}$

Ana	\$470
Bo	\$200
Elisa	\$330

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

Ana	\$0		$\xrightarrow{\tau_1}$	Ana	\$500	
Bo	\$0		$\xrightarrow{\tau_2}$	Bo	\$200	
Elisa	\$0			Elisa	\$300	
			$\xrightarrow{\tau_3}$	Ana	\$470	
				Bo	\$200	
				Elisa	\$330	
			$\xrightarrow{\tau_4}$	Ana	\$470	
				Bo	\$200	
				Elisa	\$260	

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

$\tau_5$  = “move \$500 from *Ana* to *Bo*”.

Ana	\$0		$\xrightarrow{\tau_1}$	Ana	\$500	
Bo	\$0		$\xrightarrow{\tau_2}$	Bo	\$200	
Elisa	\$0			Elisa	\$300	
			$\xrightarrow{\tau_3}$	Ana	\$470	
				Bo	\$200	
			$\xrightarrow{\tau_4}$	Elisa	\$330	
				Ana	\$470	
				Bo	\$200	
				Elisa	\$260	

## Running Example: A Banking System

Setting: Transactions change the balance of one or more accounts

The *current state* is the balance of each account obtained by executing transactions.

$\tau_1$  = “add \$500 to *Ana*”;

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”;

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”;

$\tau_5$  = *aborted* (would invalidate balances).

Ana	\$0		$\xrightarrow{\tau_1}$	Ana	\$500	
Bo	\$0		$\xrightarrow{\tau_2}$	Bo	\$200	
Elisa	\$0			Elisa	\$300	
			$\xrightarrow{\tau_3}$	Ana	\$470	
			$\xrightarrow{\tau_4}$	Bo	\$200	
				Elisa	\$260	

# Toward a Sharded and Resilient System

Consider a transaction  $\tau$  requested by client c in a resilient system.

$\tau$  is processed in *five* steps

1.  $\tau$  needs to be *received* by the system;
2.  $\tau$  must be *replicated* among all replicas in the system;
3. the replicas need to agree on an *execution order* for  $\tau$ ;
4. the replicas each need to *execute*  $\tau$  and *update* their current state accordingly;
5. the client c needs to be *informed* about the result.

# Toward a Sharded and Resilient System

Consider a transaction  $\tau$  requested by client c in a resilient system.

$\tau$  is processed in *five* steps

1.  $\tau$  needs to be *received* by the system;
2.  $\tau$  must be *replicated* among all replicas in the system;
3. the replicas need to agree on an *execution order* for  $\tau$ ;
4. the replicas each need to *execute*  $\tau$  and *update* their current state accordingly;
5. the client c needs to be *informed* about the result.

$\tau$  must be *replicated* among all replicas of all shards affected by  $\tau$ !

# Toward a Sharded and Resilient System

Consider a transaction  $\tau$  requested by client c in a resilient system.

$\tau$  is processed in *five steps*

1.  $\tau$  needs to be *received* by the system;
2.  $\tau$  must be *replicated* among all replicas in the system;
3. the replicas need to agree on an *execution order* for  $\tau$ ;
4. the replicas each need to *execute*  $\tau$  and *update* their current state accordingly;
5. the client c needs to be *informed* about the result.

What is a consistent execution order *across* shards? Does it relate to the *replication order*?

# Toward a Sharded and Resilient System

Consider a transaction  $\tau$  requested by client c in a resilient system.

$\tau$  is processed in *five* steps

1.  $\tau$  needs to be *received* by the system;
2.  $\tau$  must be *replicated* among all replicas in the system;
3. the replicas need to agree on an *execution order* for  $\tau$ ;
4. the replicas each need to *execute*  $\tau$  and *update* their current state accordingly;
5. the client c needs to be *informed* about the result.

Dependencies on data in *other shards*? Writes to data in *other shards*?

# Toward a Sharded and Resilient System

Consider a transaction  $\tau$  requested by client c in a resilient system.

$\tau$  is processed in *five* steps

1.  $\tau$  needs to be *received* by the system;
2.  $\tau$  must be *replicated* among all replicas in the system;
3. the replicas need to agree on an *execution order* for  $\tau$ ;
4. the replicas each need to *execute*  $\tau$  and *update* their current state accordingly;
5. the client c needs to be *informed* about the result.

A single consensus does no longer solve all of the above!

## Sharding Data

Sharded system: Data is distributed over all shards.

### A sharded banking system

Say we have 26 shards:  $\mathcal{C}_a, \mathcal{C}_b, \dots, \mathcal{C}_z$ ,  
such that shard  $\mathcal{C}_\xi$  holds accounts of people whose name starts with  $\xi$ .

## Sharding Data

Sharded system: Data is distributed over all shards.

### A sharded banking system

Say we have 26 shards:  $\mathcal{C}_a, \mathcal{C}_b, \dots, \mathcal{C}_z$ ,  
such that shard  $\mathcal{C}_\xi$  holds accounts of people whose name starts with  $\xi$ .

We write  $\text{shards}(\tau)$  to denote the shards affected by transaction  $\tau$ .

# Sharding Data

Sharded system: Data is distributed over all shards.

## A sharded banking system

Say we have 26 shards:  $\mathcal{C}_a, \mathcal{C}_b, \dots, \mathcal{C}_z$ ,  
such that shard  $\mathcal{C}_\xi$  holds accounts of people whose name starts with  $\xi$ .

We write  $\text{shards}(\tau)$  to denote the shards affected by transaction  $\tau$ .

$\tau_1$  = “add \$500 to *Ana*”,

$\tau_2$  = “add \$200 to *Bo* and \$300 to *Elisa*”,

$\tau_3$  = “move \$30 from *Ana* to *Elisa*”;

$\tau_4$  = “remove \$70 from *Elisa*”,

# Sharding Data

Sharded system: Data is distributed over all shards.

## A sharded banking system

Say we have 26 shards:  $\mathcal{C}_a, \mathcal{C}_b, \dots, \mathcal{C}_z$ ,  
such that shard  $\mathcal{C}_\xi$  holds accounts of people whose name starts with  $\xi$ .

We write  $\text{shards}(\tau)$  to denote the shards affected by transaction  $\tau$ .

$\tau_1 = \text{"add \$500 to } \textcolor{blue}{\textit{Ana}}\text{"}$ ,	$\text{shards}(\tau_1) = \{\mathcal{C}_a\}$ ;
$\tau_2 = \text{"add \$200 to } \textcolor{blue}{\textit{Bo}}\text{ and \$300 to } \textcolor{blue}{\textit{Elisa}}\text{"}$ ,	$\text{shards}(\tau_2) = \{\mathcal{C}_b, \mathcal{C}_e\}$ ;
$\tau_3 = \text{"move \$30 from } \textcolor{blue}{\textit{Ana}}\text{ to } \textcolor{blue}{\textit{Elisa}}\text{"}$ ,	$\text{shards}(\tau_3) = \{\mathcal{C}_a, \mathcal{C}_e\}$ ;
$\tau_4 = \text{"remove \$70 from } \textcolor{blue}{\textit{Elisa}}\text{"}$ ,	$\text{shards}(\tau_4) = \{\mathcal{C}_e\}$ .

# Sharding Data

Sharded system: Data is distributed over all shards.

## A sharded banking system

Say we have 26 shards:  $\mathcal{C}_a, \mathcal{C}_b, \dots, \mathcal{C}_z$ ,  
such that shard  $\mathcal{C}_\xi$  holds accounts of people whose name starts with  $\xi$ .

We write  $\text{shards}(\tau)$  to denote the shards affected by transaction  $\tau$ .

$\tau_1 = \text{"add \$500 to } \textcolor{blue}{\textit{Ana}}\text{"}$ ,	$\text{shards}(\tau_1) = \{\mathcal{C}_a\}$ ;	(single-shard)
$\tau_2 = \text{"add \$200 to } \textcolor{blue}{\textit{Bo}}\text{ and \$300 to } \textcolor{blue}{\textit{Elisa}}\text{"}$ ,	$\text{shards}(\tau_2) = \{\mathcal{C}_b, \mathcal{C}_e\}$ ;	(multi-shard)
$\tau_3 = \text{"move \$30 from } \textcolor{blue}{\textit{Ana}}\text{ to } \textcolor{blue}{\textit{Elisa}}\text{"}$ ;	$\text{shards}(\tau_3) = \{\mathcal{C}_a, \mathcal{C}_e\}$ ;	(multi-shard)
$\tau_4 = \text{"remove \$70 from } \textcolor{blue}{\textit{Elisa}}\text{"}$ ,	$\text{shards}(\tau_4) = \{\mathcal{C}_e\}$ .	(single-shard)

## An Example of Concurrent Execution

Consider a banking example in which

- ▶ Bo wants to transfer \$400 to Ana
  - if* Ana has at-least \$100 and Bo has at-least \$700,
- ▶ Ana wants to transfer \$300 to Elisa
  - if* Ana has at-least \$500,

and no account is allowed to have a negative balance.

## An Example of Concurrent Execution

Consider a banking example in which

- ▶ Bo wants to transfer \$400 to Ana
  - if* Ana has at-least \$100 and Bo has at-least \$700,
- ▶ Ana wants to transfer \$300 to Elisa
  - if* Ana has at-least \$500,

and no account is allowed to have a negative balance.

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

## An Example of Concurrent Execution

$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$

$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$

$A$	\$100
$B$	\$300
$E$	\$0

## An Example of Concurrent Execution

$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$

$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$

$A$	\$100
$B$	\$300
$E$	\$0

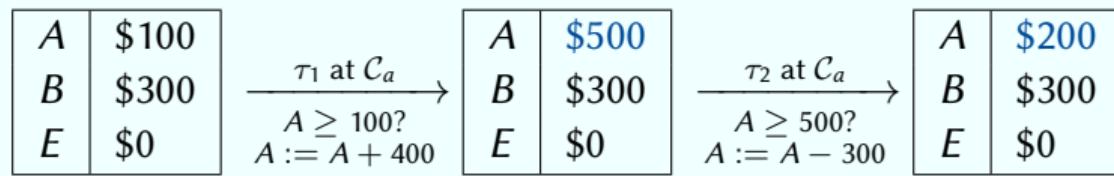
$\xrightarrow{\begin{array}{l} \tau_1 \text{ at } C_a \\ A \geq 100? \\ A := A + 400 \end{array}}$

$A$	\$500
$B$	\$300
$E$	\$0

## An Example of Concurrent Execution

$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$

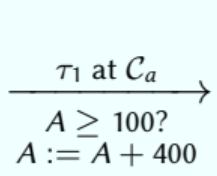
$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$



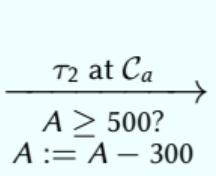
# An Example of Concurrent Execution

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$
$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

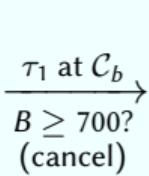
A	\$100
B	\$300
E	\$0



A	\$500
B	\$300
E	\$0



A	\$200
B	\$300
E	\$0

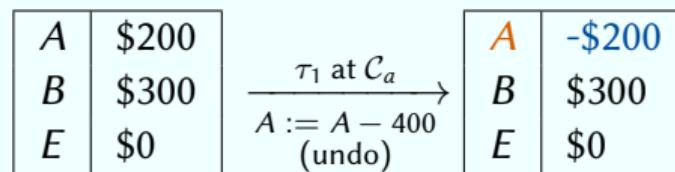
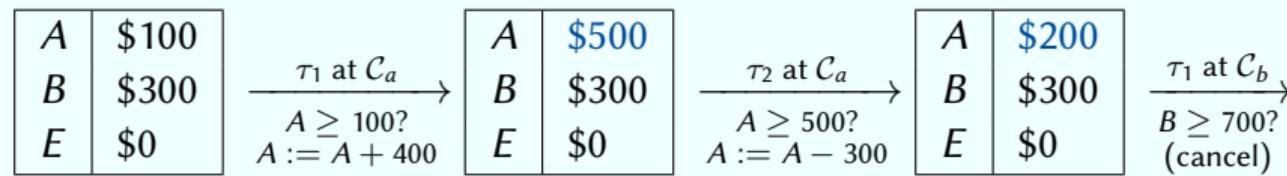


A	\$200
B	\$300
E	\$0

# An Example of Concurrent Execution

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$



# An Example of Concurrent Execution

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

A	\$100
B	\$300
E	\$0

$$\xrightarrow{\begin{array}{l} \tau_1 \text{ at } C_a \\ A \geq 100? \\ A := A + 400 \end{array}}$$

A	\$500
B	\$300
E	\$0

$$\xrightarrow{\begin{array}{l} \tau_2 \text{ at } C_a \\ A \geq 500? \\ A := A - 300 \end{array}}$$

A	\$200
B	\$300
E	\$0

$$\xrightarrow{\begin{array}{l} \tau_1 \text{ at } C_b \\ B \geq 700? \\ (\text{cancel}) \end{array}}$$

A	\$200
B	\$300
E	\$0

$$\xrightarrow{\begin{array}{l} \tau_1 \text{ at } C_a \\ A := A - 400 \\ (\text{undo}) \end{array}}$$

A	-\$200
B	\$300
E	\$0

$$\xrightarrow{\begin{array}{l} \tau_2 \text{ at } C_e \\ E := E + 300 \end{array}}$$

A	-\$200
B	\$300
E	\$300

## An Example of Concurrent Execution–Revisited

Consider a banking example in which

- ▶ Bo wants to transfer \$400 to Ana  
*if* Ana has at-least \$100 and Bo has at-least \$700,
- ▶ Ana wants to transfer \$300 to Elisa  
*if* Ana has at-least \$500,

and no account is allowed to have a negative balance.

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

## An Example of Concurrent Execution–Revisited

Consider a banking example in which

- ▶ Bo wants to transfer \$400 to Ana  
*if* Ana has at-least \$100 and Bo has at-least \$700,
- ▶ Ana wants to transfer \$300 to Elisa  
*if* Ana has at-least \$500,

and no account is allowed to have a negative balance.

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

Transactions  $\tau_1$  and  $\tau_2$  make sense:

their isolated execution will never make balances negative.

## An Example of Concurrent Execution–Revisited

Consider a banking example in which

- ▶ Bo wants to transfer \$400 to Ana  
*if* Ana has at-least \$100 and Bo has at-least \$700,
- ▶ Ana wants to transfer \$300 to Elisa  
*if* Ana has at-least \$500,

and no account is allowed to have a negative balance.

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

Transactions  $\tau_1$  and  $\tau_2$  make sense:

their isolated execution will never make balances negative.

**Guarantee by an ACID-compliant system**

No account will ever have a negative balance.

## Serializability: a High Standard for Isolation

Consider a set of transactions  $S = \{\tau_1, \dots, \tau_n\}$ .

# Serializability: a High Standard for Isolation

Consider a set of transactions  $S = \{\tau_1, \dots, \tau_n\}$ .

## Definition

A *serial schedule* is an execution of  $S$  without *interleaving* of transaction steps.  
Hence, each transaction is executed in sequence, one at a time.

# Serializability: a High Standard for Isolation

Consider a set of transactions  $S = \{\tau_1, \dots, \tau_n\}$ .

## Definition

A *serial schedule* is an execution of  $S$  without *interleaving* of transaction steps.  
Hence, each transaction is executed in sequence, one at a time.

## Definition

A *serializable schedule* is a schedule whose effect on any consistent instance is guaranteed to be identical to that of some serial schedule over the *committed transactions* in  $S$ .

# Serializability: a High Standard for Isolation

Consider a set of transactions  $S = \{\tau_1, \dots, \tau_n\}$ .

## Definition

A *serial schedule* is an execution of  $S$  without *interleaving* of transaction steps. Hence, each transaction is executed in sequence, one at a time.

## Definition

A *serializable schedule* is a schedule whose effect on any consistent instance is guaranteed to be identical to that of some serial schedule over the *committed transactions* in  $S$ .

Serializability assumes *aborted* transactions have no side effects.

# Serializability: a High Standard for Isolation

Consider a set of transactions  $S = \{\tau_1, \dots, \tau_n\}$ .

## Definition

A *serial schedule* is an execution of  $S$  without *interleaving* of transaction steps.  
Hence, each transaction is executed in sequence, one at a time.

## Definition

A *serializable schedule* is a schedule whose effect on any consistent instance is guaranteed to be identical to that of some serial schedule over the *committed transactions* in  $S$ .

Serializability assumes *aborted* transactions have no side effects.  
This is not always the case (example later).

## Simplified Transaction Notation

Consider the transaction  $\tau$ :

$\tau = \text{"if } Ana \text{ has \$500 and } Bo \text{ has \$200, then}$   
 $\text{move \$400 from } Ana \text{ to } Elisa;$   
 $\text{move \$100 from } Bo \text{ to } Elisa".}$

# Simplified Transaction Notation

Consider the transaction  $\tau$ :

$\tau = \text{"if } Ana \text{ has \$500 and } Bo \text{ has \$200, then}$   
 $\text{move \$400 from } Ana \text{ to } Elisa;$   
 $\text{move \$100 from } Bo \text{ to } Elisa".}$

What are the operations of  $\tau$ ?

Depending on *how* the system executes  $\tau$  and the database state:

- ▶ Might read from *Ana*'s account.
- ▶ Might read from *Bo*'s account.
- ▶ Might write to *Ana*'s account.
- ▶ Might write to *Bo*'s account.
- ▶ Might write to *Elisa*'s account.

## Simplified Transaction Notation

Consider the transaction  $\tau$ :

$\tau = \text{"if } Ana \text{ has \$500 and } Bo \text{ has \$200, then}$   
 $\text{move \$400 from } Ana \text{ to } Elisa;$   
 $\text{move \$100 from } Bo \text{ to } Elisa".}$

### Simplifying assumption

Each transaction is a sequence of read and write operations ending in *commit* or *abort*.

# Simplified Transaction Notation

Consider the transaction  $\tau$ :

$\tau = \text{"if } Ana \text{ has \$500 and } Bo \text{ has \$200, then}$   
 $\text{move \$400 from } Ana \text{ to } Elisa;$   
 $\text{move \$100 from } Bo \text{ to } Elisa".}$

## Simplifying assumption

Each transaction is a sequence of read and write operations ending in *commit* or *abort*.

$\text{Read}_\tau(Ana)$ ,  $\text{Read}_\tau(Bo)$ ,  $\text{Write}_\tau(Ana)$ ,  $\text{Write}_\tau(Bo)$ ,  $\text{Read}_\tau(Elisa)$ ,  $\text{Write}_\tau(Elisa)$ ,  $\text{Commit}_\tau$ .

## An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

## An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_1$ , then  $\tau_2$  (insufficient funds)

Instance  
(initial)

A	\$100
B	\$300
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_1$ , then  $\tau_2$  (insufficient funds)

Instance (initial)		<i>Schedule</i>	
A	\$100	Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (A) Abort <sub><math>\tau_1</math></sub>	

		<i>Schedule</i>	
E	\$0	Read <sub><math>\tau_2</math></sub> (A) Abort <sub><math>\tau_2</math></sub>	

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_1$ , then  $\tau_2$  (insufficient funds)

		<i>Schedule</i>	
Instance (initial)			Instance (final)
A	\$100	Read <sub><math>\tau_1</math></sub> (A)	A
B	\$300	Write <sub><math>\tau_1</math></sub> (A)	B
E	\$0	Read <sub><math>\tau_1</math></sub> (B)	E
		Write <sub><math>\tau_1</math></sub> (A)	
		Abort <sub><math>\tau_1</math></sub>	
			Read <sub><math>\tau_2</math></sub> (A)
			Abort <sub><math>\tau_2</math></sub>

## An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_1$ , then  $\tau_2$  (Bob has sufficient funds)

Instance  
(initial)

A	\$100
B	\$800
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_1$ , then  $\tau_2$  (Bob has sufficient funds)

Schedule

Instance  
(initial)

A	\$100
B	\$800
E	\$0

Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B) Commit <sub><math>\tau_1</math></sub>	Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>
---	---

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_1$ , then  $\tau_2$  (Bob has sufficient funds)

Schedule

Instance  
(initial)

A	\$100
B	\$800
E	\$0

Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B) Commit <sub><math>\tau_1</math></sub>	Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>
---	---

Instance  
(final)

A	\$200
B	\$400
E	\$300

## An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_2$ , then  $\tau_1$  (Bob has sufficient funds)

Instance  
(initial)

A	\$100
B	\$800
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_2$ , then  $\tau_1$  (Bob has sufficient funds)

Instance (initial)		<i>Schedule</i>	
A	\$100	Read <sub><math>\tau_2</math></sub> (A) Abort <sub><math>\tau_2</math></sub>	Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B) Commit <sub><math>\tau_1</math></sub>

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_2$ , then  $\tau_1$  (Bob has sufficient funds)

		<i>Schedule</i>	
Instance (initial)			Instance (final)
A	\$100	Read <sub><math>\tau_2</math></sub> (A) Abort <sub><math>\tau_2</math></sub>	A \$500
B	\$800	Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B)	B \$400
E	\$0	Commit <sub><math>\tau_1</math></sub>	E \$0

## An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_2$ , then  $\tau_1$  (Ana has sufficient funds)

Instance  
(initial)

A	\$500
B	\$300
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_2$ , then  $\tau_1$  (Ana has sufficient funds)

*Schedule*

Instance  
(initial)

A	\$500
B	\$300
E	\$0

Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>	
Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (A) Abort <sub><math>\tau_1</math></sub>	

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Serial schedule:  $\tau_2$ , then  $\tau_1$  (Ana has sufficient funds)

*Schedule*

Instance  
(initial)

A	\$500
B	\$300
E	\$0

Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>
Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (A) Abort <sub><math>\tau_1</math></sub>

Instance  
(final)

A	\$200
B	\$300
E	\$300

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

Non-serial schedule—Earlier example

Instance  
(initial)

A	\$100
B	\$300
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Non-serial schedule—Earlier example

		<i>Schedule</i>
Instance (initial)		
$A$	\$100	Read $_{\tau_1}(A)$ Write $_{\tau_1}(A)$
$B$	\$300	Read $_{\tau_2}(A)$ Write $_{\tau_2}(A)$ Read $_{\tau_2}(E)$ Write $_{\tau_2}(E)$ Commit $_{\tau_2}$
$E$	\$0	Read $_{\tau_1}(B)$ Read $_{\tau_1}(A)$ Write $_{\tau_1}(A)$ Abort $_{\tau_1}$

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Non-serial schedule—Earlier example

		<i>Schedule</i>		
Instance (initial)			Instance (final)	
A	\$100	Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A)	Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>	A   -\$200 B   \$300 E   \$300
B	\$300	Read <sub><math>\tau_1</math></sub> (B) Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Abort <sub><math>\tau_1</math></sub>		
E	\$0			

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Non-serial schedule—Another example

Instance  
(initial)

A	\$500
B	\$800
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Non-serial schedule—Another example

		<i>Schedule</i>
Instance (initial)		
A	\$500	Read <sub><math>\tau_1</math></sub> (A)
B	\$800	Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>
E	\$0	Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B) Commit <sub><math>\tau_1</math></sub>

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Non-serial schedule—Another example

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

Read <sub><math>\tau_1</math></sub> (A)	Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>
Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B) Commit <sub><math>\tau_1</math></sub>	

Instance  
(final)

A	\$900
B	\$400
E	\$300

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Non-serial schedule—A third example

Instance  
(initial)

A	\$500
B	\$800
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

## Non-serial schedule—A third example

		<i>Schedule</i>
Instance (initial)		
A	\$500	Read <sub><math>\tau_2</math></sub> (A)
B	\$800	Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B) Commit <sub><math>\tau_1</math></sub>
E	\$0	Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Non-serial schedule—A third example

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B) Commit <sub><math>\tau_1</math></sub>	Read <sub><math>\tau_2</math></sub> (A)     Write <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E) Commit <sub><math>\tau_2</math></sub>
---	---

Instance  
(final)

A	\$200
B	\$400
E	\$300

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100?, A := A + 400, B \geq 700?, B := B - 400;$$

$$\tau_2 = A \geq 500?, A := A - 300, E := E + 300.$$

A serializable schedule (that is non-serial)

Instance  
(initial)

A	\$500
B	\$800
E	\$0

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

A serializable schedule (that is non-serial)

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A)	Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E)
Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B)	Commit <sub><math>\tau_2</math></sub>
Commit <sub><math>\tau_1</math></sub>	

# An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

A serializable schedule (that is non-serial)

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A)	Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A)
Read <sub><math>\tau_2</math></sub> (E) Write <sub><math>\tau_2</math></sub> (E)	Read <sub><math>\tau_1</math></sub> (B) Write <sub><math>\tau_1</math></sub> (B)
Commit <sub><math>\tau_2</math></sub>	Commit <sub><math>\tau_1</math></sub>

Instance  
(final)

A	\$600
B	\$400
E	\$300

## An Example of Schedules

Consider again the transactions

$$\tau_1 = A \geq 100? , A := A + 400, B \geq 700? , B := B - 400;$$

$$\tau_2 = A \geq 500? , A := A - 300, E := E + 300.$$

Key observation: Serial schedules

Individual transactions *make sense* (do not violate consistency):

- ▶ No balance will ever get negative.
- ▶ No money disappears or appears out of thin air.

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.
- ▶ All transactions operate on *shared data* (the database instance).

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.
- ▶ All transactions operate on *shared data* (the database instance).
- ▶ We need to coordinate access to this shared data!

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.
- ▶ All transactions operate on *shared data* (the database instance).
- ▶ We need to coordinate access to this shared data!

In traditional multi-threaded programs:

- ▶ Use *critical sections* in which shared data is accessed.
- ▶ Enforce *critical sections* with locks (e.g., mutex).
- ▶ Ensure proper lock usage to avoid deadlocks, ....

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.
- ▶ All transactions operate on *shared data* (the database instance).
- ▶ We need to coordinate access to this shared data!

In traditional multi-threaded programs:

- ▶ Use *critical sections* in which shared data is accessed.
- ▶ Enforce *critical sections* with locks (e.g., mutex).
- ▶ Ensure proper lock usage to avoid deadlocks, ....

As all data is shared: should the entire transaction be a single critical section?

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.
- ▶ All transactions operate on *shared data* (the database instance).
- ▶ We need to coordinate access to this shared data!

In traditional multi-threaded programs:

- ▶ Use *critical sections* in which shared data is accessed.
- ▶ Enforce *critical sections* with locks (e.g., mutex).
- ▶ Ensure proper lock usage to avoid deadlocks, ....

As all data is shared: should the entire transaction be a single critical section?

What if each transaction *locks the system*, executes, *releases the system*?

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.
- ▶ All transactions operate on *shared data* (the database instance).
- ▶ We need to coordinate access to this shared data!

In traditional multi-threaded programs:

- ▶ Use *critical sections* in which shared data is accessed.
- ▶ Enforce *critical sections* with locks (e.g., mutex).
- ▶ Ensure proper lock usage to avoid deadlocks, ....

As all data is shared: should the entire transaction be a single critical section?

What if each transaction *locks the system*, executes, *releases the system*?

This will enforce a *serial schedule*.

# Guaranteeing Isolation

## Simplified point-of-view

- ▶ A transaction is a *thread* in a multi-threaded program.
- ▶ All transactions operate on *shared data* (the database instance).
- ▶ We need to coordinate access to this shared data!

In traditional multi-threaded programs:

- ▶ Use *critical sections* in which shared data is accessed.
- ▶ Enforce *critical sections* with locks (e.g., mutex).
- ▶ Ensure proper lock usage to avoid deadlocks, ....

As all data is shared: should the entire transaction be a single critical section?

What if each transaction *locks the system*, executes, *releases the system*?

This will enforce a *serial schedule* and eliminate any concurrency.

## Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

## Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

## Using fine-grained locks

A transaction  $\tau$  that wants to access database object  $O$  will:

- ▶ waits until it obtains a lock on  $O$  ( $\text{Lock}_\tau(O)$ ),
- ▶ then perform its operations on  $O$  (e.g.,  $\text{Read}_\tau(O)$  and  $\text{Write}_\tau(O)$ ), and
- ▶ finally release the lock on  $O$  ( $\text{Release}_\tau(O)$ ).

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

Read <sub>$\tau_1$</sub> (A)

Write <sub>$\tau_1$</sub> (A)  
Read <sub>$\tau_1$</sub> (B)

Write <sub>$\tau_1$</sub> (B)  
Commit <sub>$\tau_1$</sub>

Read <sub>$\tau_2$</sub> (A)

Write <sub>$\tau_2$</sub> (A)  
Read <sub>$\tau_2$</sub> (E)  
Write <sub>$\tau_2$</sub> (E)

Commit <sub>$\tau_2$</sub>

Instance  
(final)

A	\$900
B	\$400
E	\$300

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

$\text{Lock}_{\tau_1}(A)$ $\text{Read}_{\tau_1}(A)$	

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

$\text{Lock}_{\tau_1}(A)$ $\text{Read}_{\tau_1}(A)$	$\text{Lock}_{\tau_2}(A)$
--	---------------------------

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance (initial)	
A	\$500
B	\$800
E	\$0

$\text{Lock}_{\tau_1}(A)$ $\text{Read}_{\tau_1}(A)$  $\text{Write}_{\tau_1}(A)$ $\text{Release}_{\tau_1}(A)$	$\text{Lock}_{\tau_2}(A)$
--	---------------------------

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0

$\text{Lock}_{\tau_1}(A)$ $\text{Read}_{\tau_1}(A)$  $\text{Write}_{\tau_1}(A)$ $\text{Release}_{\tau_1}(A)$	$\text{Lock}_{\tau_2}(A)$  $\text{Read}_{\tau_2}(A)$
--	--

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

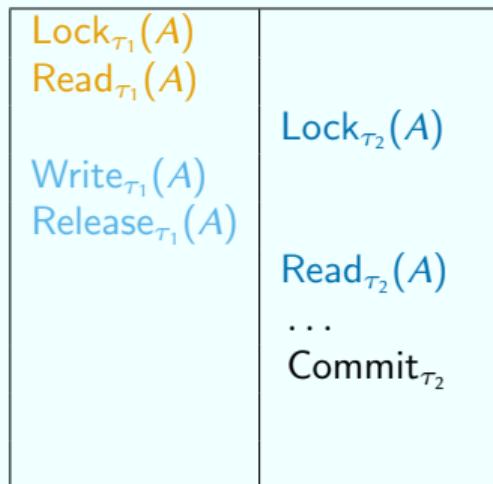
E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance (initial)	
A	\$500
B	\$800
E	\$0



# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

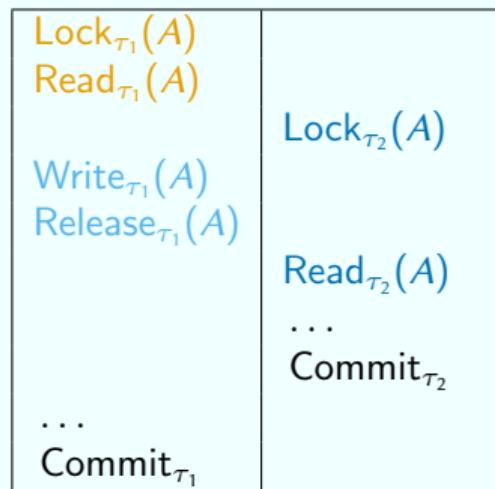
E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance (initial)	
A	\$500
B	\$800
E	\$0



# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

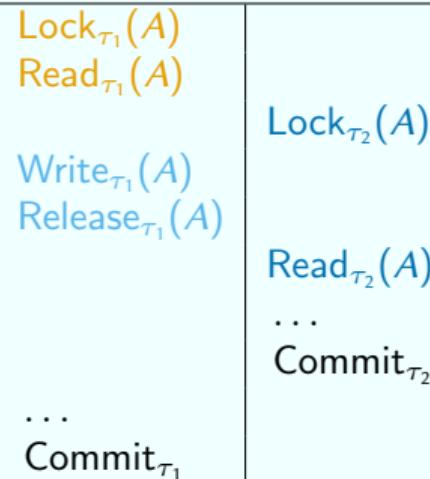
In our examples we abstract from details: *accounts* are database objects.

Lock-based access solves *some* issues ...

*Schedule*

Instance  
(initial)

A	\$500
B	\$800
E	\$0



Instance  
(final)

A	\$600
B	\$400
E	\$300

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

...but not *all* issues ...

Instance  
(initial)

A	\$100
B	\$300
E	\$0

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

...but not *all* issues ...

Instance (initial)		<i>Schedule</i>	
$A$	\$100	$\text{Lock}_{\tau_1}(A)$	
$B$	\$300	$\text{Read}_{\tau_1}(A)$	
$E$	\$0	$\text{Write}_{\tau_1}(A)$	
		$\text{Release}_{\tau_1}(A)$	
			$\text{Lock}_{\tau_2}(A)$
			$\text{Read}_{\tau_2}(A)$
			$\text{Write}_{\tau_2}(A)$
			...
			$\text{Commit}_{\tau_2}$
		...	
			$\text{Abort}_{\tau_1}$

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

...but not *all* issues ...

		<i>Schedule</i>	
			Instance (final)
			A    \$-200
		Lock <sub><math>\tau_1</math></sub> (A) Read <sub><math>\tau_1</math></sub> (A) Write <sub><math>\tau_1</math></sub> (A) Release <sub><math>\tau_1</math></sub> (A)	B    \$300
		Lock <sub><math>\tau_2</math></sub> (A) Read <sub><math>\tau_2</math></sub> (A) Write <sub><math>\tau_2</math></sub> (A) ...	E    \$300
		Commit <sub><math>\tau_2</math></sub>	
		Abort <sub><math>\tau_1</math></sub>	

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

...and introduces *new issues*.

Consider two transactions that both want to access *Ana* and *Bo*:

$$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \dots; \quad \tau_2 = \text{Lock}_{\tau_2}(B), \text{Lock}_{\tau_1}(A), \dots$$

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

...and introduces *new issues*.

Consider two transactions that both want to access *Ana* and *Bo*:

$$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \dots; \quad \tau_2 = \text{Lock}_{\tau_2}(B), \text{Lock}_{\tau_1}(A), \dots$$

Schedule	
$\text{Lock}_{\tau_1}(A)$	$\text{Lock}_{\tau_2}(B)$
$\text{Lock}_{\tau_1}(B)$	$\text{Lock}_{\tau_2}(A)$

# Improving Isolation using Locks

Idea: Use a fine-grained set of locks on *database objects*.

E.g., accounts, tables, rows, ....

In our examples we abstract from details: *accounts* are database objects.

...and introduces *new issues*.

Consider two transactions that both want to access *Ana* and *Bo*:

$$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \dots; \quad \tau_2 = \text{Lock}_{\tau_2}(B), \text{Lock}_{\tau_1}(A), \dots$$

Schedule	
$\text{Lock}_{\tau_1}(A)$	$\text{Lock}_{\tau_2}(B)$
$\text{Lock}_{\tau_1}(B)$	$\text{Lock}_{\tau_2}(A)$

Both transactions will wait forever: a deadlock!

## Achieving Serializability with Locks

Locking itself does not guarantee *serializability*.

Some *locking protocols* (sets of rules on when to use locks) that do guarantee *serializability*.

# Achieving Serializability with Locks

Locking itself does not guarantee *serializability*.

Some *locking protocols* (sets of rules on when to use locks) that do guarantee *serializability*.

## Two-phase locking protocol (2PL)

Execution of transaction  $\tau$  adheres to 2PL if the execution is performed in two phases:

**Growing phase** during which execution can obtain locks, and *not* release them; and

**Shrinking phase** during which execution can release locks, and *not* obtain them,

and any database object  $O$  is only operated on while holding lock  $\text{Lock}_\tau(O)$ .

# Achieving Serializability with Locks

Locking itself does not guarantee *serializability*.

Some *locking protocols* (sets of rules on when to use locks) that do guarantee *serializability*.

## Two-phase locking protocol (2PL)

Execution of transaction  $\tau$  adheres to 2PL if the execution is performed in two phases:

**Growing phase** during which execution can obtain locks, and *not* release them; and

**Shrinking phase** during which execution can release locks, and *not* obtain them,

and any database object  $O$  is only operated on while holding lock  $\text{Lock}_\tau(O)$ .

**Strict** 2PL: locks are only released after completion ( $\text{Commit}_\tau$  or  $\text{Abort}_\tau$ ).

# Achieving Serializability with Locks

Locking itself does not guarantee *serializability*.

Some *locking protocols* (sets of rules on when to use locks) that do guarantee *serializability*.

## Two-phase locking protocol (2PL)

Execution of transaction  $\tau$  adheres to 2PL if the execution is performed in two phases:

**Growing phase** during which execution can obtain locks, and *not* release them; and

**Shrinking phase** during which execution can release locks, and *not* obtain them,

and any database object  $O$  is only operated on while holding lock  $\text{Lock}_\tau(O)$ .

**Strict** 2PL: locks are only released after completion ( $\text{Commit}_\tau$  or  $\text{Abort}_\tau$ ).

Notice—Nothing to deal with *deadlocks*.

## An Example of 2PL

Consider again the transactions

$$\begin{aligned}\tau_1 &= A \geq 100?, A := A + 400, B \geq 700?, B := B - 400; \\ \tau_2 &= A \geq 500?, A := A - 300, E := E + 300.\end{aligned}$$

## An Example of 2PL

Consider again the transactions

$$\begin{aligned}\tau_1 &= A \geq 100?, A := A + 400, B \geq 700?, B := B - 400; \\ \tau_2 &= A \geq 500?, A := A - 300, E := E + 300.\end{aligned}$$

Assumption: Both transactions will succeed (Alice and Bob have sufficient funds)

$$\begin{aligned}\tau_1 &= \text{Lock}_{\tau_1}(A), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(B), \text{Write}_{\tau_1}(B), \\ &\quad \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(A), \text{Release}_{\tau_1}(B); \\ \tau_2 &= \text{Lock}_{\tau_2}(E), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(A), \text{Write}_{\tau_2}(A), \text{Read}_{\tau_2}(E), \text{Write}_{\tau_2}(E), \\ &\quad \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(E).\end{aligned}$$

## An Example of 2PL

Consider again the transactions

$$\begin{aligned}\tau_1 &= A \geq 100?, A := A + 400, B \geq 700?, B := B - 400; \\ \tau_2 &= A \geq 500?, A := A - 300, E := E + 300.\end{aligned}$$

Assumption: Both transactions will succeed (Alice and Bob have sufficient funds)

$$\begin{aligned}\tau_1 &= \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(A), \text{Read}_{\tau_1}(B), \text{Write}_{\tau_1}(B), \\ &\quad \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(A), \text{Release}_{\tau_1}(B); \\ \tau_2 &= \text{Lock}_{\tau_2}(E), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(A), \text{Write}_{\tau_2}(A), \text{Read}_{\tau_2}(E), \text{Write}_{\tau_2}(E), \\ &\quad \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(E).\end{aligned}$$

## An Example of 2PL

Consider again the transactions

$$\begin{aligned}\tau_1 &= A \geq 100?, A := A + 400, B \geq 700?, B := B - 400; \\ \tau_2 &= A \geq 500?, A := A - 300, E := E + 300.\end{aligned}$$

Assumption: Both transactions will succeed (Alice and Bob have sufficient funds)

$$\begin{aligned}\tau_1 &= \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(A), \text{Read}_{\tau_1}(B), \text{Write}_{\tau_1}(B), \\ &\quad \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(B), \text{Release}_{\tau_1}(A); \\ \tau_2 &= \text{Lock}_{\tau_2}(E), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(A), \text{Write}_{\tau_2}(A), \text{Read}_{\tau_2}(E), \text{Write}_{\tau_2}(E), \\ &\quad \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(E).\end{aligned}$$

## An Example of 2PL

Consider again the transactions

$$\begin{aligned}\tau_1 &= A \geq 100?, A := A + 400, B \geq 700?, B := B - 400; \\ \tau_2 &= A \geq 500?, A := A - 300, E := E + 300.\end{aligned}$$

Assumption: Both transactions will succeed (Alice and Bob have sufficient funds)

$$\begin{aligned}\tau_1 &= \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(A), \text{Read}_{\tau_1}(B), \text{Write}_{\tau_1}(B), \\ &\quad \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(B), \text{Release}_{\tau_1}(A); \\ \tau_2 &= \text{Lock}_{\tau_2}(E), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(A), \text{Write}_{\tau_2}(A), \text{Read}_{\tau_2}(E), \text{Write}_{\tau_2}(E), \\ &\quad \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(E).\end{aligned}$$

These are all *strict* 2PL: locks are released after the transactions commit.

## An Example of 2PL

Assumption: Both transactions will succeed (Alice and Bob have sufficient funds)

$$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(A), \text{Read}_{\tau_1}(B), \text{Write}_{\tau_1}(B), \\ \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(B), \text{Release}_{\tau_1}(A);$$
$$\tau_2 = \text{Lock}_{\tau_2}(E), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(A), \text{Write}_{\tau_2}(A), \text{Read}_{\tau_2}(E), \text{Write}_{\tau_2}(E), \\ \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(E).$$

Consider any schedule with any interleaving of operations of  $\tau_1$  and  $\tau_2$

## An Example of 2PL

Assumption: Both transactions will succeed (Alice and Bob have sufficient funds)

$$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(A), \text{Read}_{\tau_1}(B), \text{Write}_{\tau_1}(B), \\ \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(B), \text{Release}_{\tau_1}(A);$$
$$\tau_2 = \text{Lock}_{\tau_2}(E), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(A), \text{Write}_{\tau_2}(A), \text{Read}_{\tau_2}(E), \text{Write}_{\tau_2}(E), \\ \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(E).$$

Consider any schedule with any interleaving of operations of  $\tau_1$  and  $\tau_2$

- If  $\tau_1$  executes  $\text{Lock}_{\tau_1}(A)$  *before*  $\tau_2$  executes  $\text{Lock}_{\tau_2}(A)$ :  
all read and write operations of  $\tau_1$  effectively happen before those of  $\tau_2$ .

## An Example of 2PL

Assumption: Both transactions will succeed (Alice and Bob have sufficient funds)

$$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(A), \text{Read}_{\tau_1}(B), \text{Write}_{\tau_1}(B), \\ \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(B), \text{Release}_{\tau_1}(A);$$
$$\tau_2 = \text{Lock}_{\tau_2}(E), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(A), \text{Write}_{\tau_2}(A), \text{Read}_{\tau_2}(E), \text{Write}_{\tau_2}(E), \\ \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(E).$$

Consider any schedule with any interleaving of operations of  $\tau_1$  and  $\tau_2$

- ▶ If  $\tau_1$  executes  $\text{Lock}_{\tau_1}(A)$  *before*  $\tau_2$  executes  $\text{Lock}_{\tau_2}(A)$ :  
all read and write operations of  $\tau_1$  effectively happen before those of  $\tau_2$ .
- ▶ If  $\tau_2$  executes  $\text{Lock}_{\tau_2}(A)$  *before*  $\tau_1$  executes  $\text{Lock}_{\tau_1}(A)$ :  
all read and write operations of  $\tau_2$  effectively happen before those of  $\tau_1$ .

## Two-Phase Locking and Deadlocks

Consider the transactions

$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(B), \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(A), \text{Release}_{\tau_1}(B);$

$\tau_2 = \text{Lock}_{\tau_2}(B), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(B), \text{Write}_{\tau_2}(A), \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(B).$

These transactions are strict 2PL.

## Two-Phase Locking and Deadlocks

Consider the transactions

$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(B), \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(A), \text{Release}_{\tau_1}(B);$

$\tau_2 = \text{Lock}_{\tau_2}(B), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(B), \text{Write}_{\tau_2}(A), \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(B).$

These transactions are strict 2PL.

Some schedules will cause a deadlock

<i>Schedule</i>	
$\text{Lock}_{\tau_1}(A)$	$\text{Lock}_{\tau_2}(B)$
$\text{Lock}_{\tau_1}(B)$	$\text{Lock}_{\tau_2}(A)$

## Two-Phase Locking and Deadlocks

Consider the transactions

$\tau_1 = \text{Lock}_{\tau_1}(A), \text{Lock}_{\tau_1}(B), \text{Read}_{\tau_1}(A), \text{Write}_{\tau_1}(B), \text{Commit}_{\tau_1}, \text{Release}_{\tau_1}(A), \text{Release}_{\tau_1}(B);$

$\tau_2 = \text{Lock}_{\tau_2}(B), \text{Lock}_{\tau_2}(A), \text{Read}_{\tau_2}(B), \text{Write}_{\tau_2}(A), \text{Commit}_{\tau_2}, \text{Release}_{\tau_2}(A), \text{Release}_{\tau_2}(B).$

These transactions are strict 2PL.

Some schedules will cause a deadlock

Schedule	
$\text{Lock}_{\tau_1}(A)$	$\text{Lock}_{\tau_2}(B)$
$\text{Lock}_{\tau_1}(B)$	$\text{Lock}_{\tau_2}(A)$

Deadlocks are one of the issues arising from *lock contention*.

## Dealing with Deadlocks: Pessimistic Approach

Pessimistic: make sure deadlocks *cannot happen*

## Dealing with Deadlocks: Pessimistic Approach

Pessimistic: make sure deadlocks *cannot happen*

Enforce that all transactions obtain their locks in a unique predetermined order.

E.g., first locks on Ana, then Bo, then Celeste, then Dafni, then Elisa, ....

## Dealing with Deadlocks: Pessimistic Approach

Pessimistic: make sure deadlocks *cannot happen*

Enforce that all transactions obtain their locks in a unique predetermined order.

E.g., first locks on Ana, then Bo, then Celeste, then Dafni, then Elisa, ....

### Example

Consider the transaction

$\tau$  = “if *Bo* has \$500, then move \$200 from *Bo* to *Ana*”.

Any schedule for  $\tau$  needs to start with:

$\text{Lock}_\tau(\text{Ana}), \text{Lock}_\tau(\text{Bo}), \dots,$

we even lock Ana if Bo does *not have funds*.

## Dealing with Deadlocks: Optimistic Approach

Optimistic: Optimize for no lock-contention

If a transaction tries to obtain a lock that is already held: *abort the transaction entirely.*

## Dealing with Deadlocks: Optimistic Approach

Optimistic: Optimize for no lock-contention

If a transaction tries to obtain a lock that is already held: *abort the transaction entirely.*

- ▶ No need for *deadlock detection* or *prevention*.

# Dealing with Deadlocks: Optimistic Approach

## Optimistic: Optimize for no lock-contention

If a transaction tries to obtain a lock that is already held: *abort the transaction entirely.*

- ▶ No need for *deadlock detection* or *prevention*.
- ▶ Very easy to implement.

# Dealing with Deadlocks: Optimistic Approach

## Optimistic: Optimize for no lock-contention

If a transaction tries to obtain a lock that is already held: *abort the transaction entirely.*

- ▶ No need for *deadlock detection* or *prevention*.
- ▶ Very easy to implement.
- ▶ Minimizes the costs for transactions that are able to commit.

# Dealing with Deadlocks: Optimistic Approach

## Optimistic: Optimize for no lock-contention

If a transaction tries to obtain a lock that is already held: *abort the transaction entirely.*

- ▶ No need for *deadlock detection* or *prevention*.
- ▶ Very easy to implement.
- ▶ Minimizes the costs for transactions that are able to commit.
- ▶ Will perform badly when there is a high amount of lock-contention.

## Practice: Read and Write locks

- ▶ Locks need to be *fine-grained* to maximize concurrency.
- ▶ Concurrency issues only arise when a transaction is writing.
- ▶ In most workloads: reads are much more frequent than writes.

## Practice: Read and Write locks

- ▶ Locks need to be *fine-grained* to maximize concurrency.
- ▶ Concurrency issues only arise when a transaction is writing.
- ▶ In most workloads: reads are much more frequent than writes.

Goal: prevent writes concurrent with other activity, but minimize cost for reads.

## Practice: Read and Write locks

- ▶ Locks need to be *fine-grained* to maximize concurrency.
- ▶ Concurrency issues only arise when a transaction is writing.
- ▶ In most workloads: reads are much more frequent than writes.

Goal: prevent writes concurrent with other activity, but minimize cost for reads.

## Introduce separate read and write locks

- ▶ Multiple transactions can hold a lock at the same time *if they all hold read locks.*
- ▶ Only one transaction can hold a lock *if that transaction holds a write lock.*

## Practice: Read and Write locks

- ▶ Locks need to be *fine-grained* to maximize concurrency.
- ▶ Concurrency issues only arise when a transaction is writing.
- ▶ In most workloads: reads are much more frequent than writes.

Goal: prevent writes concurrent with other activity, but minimize cost for reads.

## Introduce separate read and write locks

- ▶ Multiple transactions can hold a lock at the same time *if they all hold read locks.*
- ▶ Only one transaction can hold a lock *if that transaction holds a write lock.*

## Result

- ▶ Many transactions can read at the same time.
- ▶ Read-write, write-read, and write-write conflicts are prevented.

## The Cost of Serializability

- ▶ Serializability provides *strong* isolation guarantees.
- ▶ Providing these guarantees *will* impact concurrency  
(independent of the implementation mechanism, e.g., locks).

## The Cost of Serializability

- ▶ Serializability provides *strong* isolation guarantees.
- ▶ Providing these guarantees *will* impact concurrency  
(independent of the implementation mechanism, e.g., locks).

To improve performance, you can *give up* on serializability!

# Degrees of Isolation in SQL<sup>1</sup>

Level	Dirty Reads	Unrepeatable Read	Phantoms
<b>READ UNCOMMITTED</b>	Possible	Possible	Possible
<b>READ COMMITTED</b>	Not Possible	Possible	Possible
<b>REPEATABLE READ</b>	Not Possible	Not Possible	Possible
<b>SERIALIZABLE</b>	Not Possible	Not Possible	Not Possible

---

<sup>1</sup>There are excellent papers on this topic! E.g., <https://doi.org/10.1145/568271.223785> and [https://doi.org/10.1016/0950-5849\(96\)01109-3](https://doi.org/10.1016/0950-5849(96)01109-3) are recommended.

# Degrees of Isolation in SQL<sup>1</sup>

Level	Dirty Reads	Unrepeatable Read	Phantoms
<b>READ UNCOMMITTED</b>	Possible	Possible	Possible
<b>READ COMMITTED</b>	Not Possible	Possible	Possible
<b>REPEATABLE READ</b>	Not Possible	Not Possible	Possible
<b>SERIALIZABLE</b>	Not Possible	Not Possible	Not Possible

Each *level* can be defined in terms of a locking protocol.

---

<sup>1</sup>There are excellent papers on this topic! E.g., <https://doi.org/10.1145/568271.223785> and [https://doi.org/10.1016/0950-5849\(96\)01109-3](https://doi.org/10.1016/0950-5849(96)01109-3) are recommended.

# Degrees of Isolation in SQL<sup>1</sup>

Level	Dirty Reads	Unrepeatable Read	Phantoms
<b>READ UNCOMMITTED</b>	Possible	Possible	Possible
<b>READ COMMITTED</b>	Not Possible	Possible	Possible
<b>REPEATABLE READ</b>	Not Possible	Not Possible	Possible
<b>SERIALIZABLE</b>	Not Possible	Not Possible	Not Possible

Each *level* can be defined in terms of a locking protocol.

## Locking protocol for **READ UNCOMMITTED**

- ▶ no read locks,
- ▶ *long-duration* write (and predicate) locks before writing data.

---

<sup>1</sup>There are excellent papers on this topic! E.g., <https://doi.org/10.1145/568271.223785> and [https://doi.org/10.1016/0950-5849\(96\)01109-3](https://doi.org/10.1016/0950-5849(96)01109-3) are recommended.

# Degrees of Isolation in SQL<sup>1</sup>

Level	Dirty Reads	Unrepeatable Read	Phantoms
<b>READ UNCOMMITTED</b>	Possible	Possible	Possible
<b>READ COMMITTED</b>	Not Possible	Possible	Possible
<b>REPEATABLE READ</b>	Not Possible	Not Possible	Possible
<b>SERIALIZABLE</b>	Not Possible	Not Possible	Not Possible

Each *level* can be defined in terms of a locking protocol.

## Locking protocol for **READ COMMITTED**

- ▶ *short-duration* read (and predicate) locks before reading data, and
- ▶ *long-duration* write (and predicate) locks before writing data.

---

<sup>1</sup>There are excellent papers on this topic! E.g., <https://doi.org/10.1145/568271.223785> and [https://doi.org/10.1016/0950-5849\(96\)01109-3](https://doi.org/10.1016/0950-5849(96)01109-3) are recommended.

# Degrees of Isolation in SQL<sup>1</sup>

Level	Dirty Reads	Unrepeatable Read	Phantoms
<b>READ UNCOMMITTED</b>	Possible	Possible	Possible
<b>READ COMMITTED</b>	Not Possible	Possible	Possible
<b>REPEATABLE READ</b>	Not Possible	Not Possible	Possible
<b>SERIALIZABLE</b>	Not Possible	Not Possible	Not Possible

Each *level* can be defined in terms of a locking protocol.

## Locking protocol for **REPEATABLE READ**

- ▶ *short-duration* predicate locks and *long-duration* read locks before reading data, and
- ▶ *long-duration* write (and predicate) locks before writing data.

---

<sup>1</sup>There are excellent papers on this topic! E.g., <https://doi.org/10.1145/568271.223785> and [https://doi.org/10.1016/0950-5849\(96\)01109-3](https://doi.org/10.1016/0950-5849(96)01109-3) are recommended.

# Degrees of Isolation in SQL<sup>1</sup>

Level	Dirty Reads	Unrepeatable Read	Phantoms
<b>READ UNCOMMITTED</b>	Possible	Possible	Possible
<b>READ COMMITTED</b>	Not Possible	Possible	Possible
<b>REPEATABLE READ</b>	Not Possible	Not Possible	Possible
<b>SERIALIZABLE</b>	Not Possible	Not Possible	Not Possible

Each *level* can be defined in terms of a locking protocol.

## Locking protocol for **SERIALIZABLE**

- ▶ *long-duration* read (and predicate) locks before reading data, and
- ▶ *long-duration* write (and predicate) locks before writing data.

---

<sup>1</sup>There are excellent papers on this topic! E.g., <https://doi.org/10.1145/568271.223785> and [https://doi.org/10.1016/0950-5849\(96\)01109-3](https://doi.org/10.1016/0950-5849(96)01109-3) are recommended.

# Degrees of Isolation in SQL<sup>1</sup>

Level	Dirty Reads	Unrepeatable Read	Phantoms
<b>READ UNCOMMITTED</b>	Possible	Possible	Possible
<b>READ COMMITTED</b>	Not Possible	Possible	Possible
<b>REPEATABLE READ</b>	Not Possible	Not Possible	Possible
<b>SERIALIZABLE</b>	Not Possible	Not Possible	Not Possible

Each *level* can be defined in terms of a locking protocol.

## Locking protocol for **SERIALIZABLE** (2PL)

- ▶ *long-duration* read (and predicate) locks before reading data, and
- ▶ *long-duration* write (and predicate) locks before writing data.

---

<sup>1</sup>There are excellent papers on this topic! E.g., <https://doi.org/10.1145/568271.223785> and [https://doi.org/10.1016/0950-5849\(96\)01109-3](https://doi.org/10.1016/0950-5849(96)01109-3) are recommended.

## Safe Execution without Locks

Concurrent transaction execution can make sense *without* isolation.  
E.g., one can use application-specific knowledge!

## Safe Execution without Locks

Concurrent transaction execution can make sense *without* isolation.  
E.g., one can use application-specific knowledge!

### A Banking System

Observe: undoing a withdraw increases balance, undoing deposits decreases balance!

# Safe Execution without Locks

Concurrent transaction execution can make sense *without* isolation.  
E.g., one can use application-specific knowledge!

## A Banking System

Observe: undoing a withdraw increases balance, undoing deposits decreases balance!

Consider executions in which all steps can:

- ▶ always withdraw money;
- ▶ only deposit money after either *commit* or *abort* is decided.

# Safe Execution without Locks

Concurrent transaction execution can make sense *without* isolation.  
E.g., one can use application-specific knowledge!

## A Banking System

Observe: undoing a withdraw increases balance, undoing deposits decreases balance!

Consider executions in which all steps can:

- ▶ always withdraw money;
- ▶ only deposit money after either *commit* or *abort* is decided.

These executions guarantee that no account will have a negative balance!

# Ingredients of Sharding in a Resilient Environment

Multi-shard transaction execution of  $\tau$  requires

Replication of  $\tau$  among shards.

E.g., a two-phase commit step.

Concurrency control to guarantee consistent execution of  $\tau$ .

E.g., using *locks* to prevent concurrent access to accounts.

To One needs *computations* within a shard and *communication* between shards.

# Ingredients of Sharding in a Resilient Environment

Multi-shard transaction execution of  $\tau$  requires

Replication of  $\tau$  among shards.

E.g., a two-phase commit step.

Concurrency control to guarantee consistent execution of  $\tau$ .

E.g., using *locks* to prevent concurrent access to accounts.

To One needs *computations* within a shard and *communication* between shards.

Fault-tolerant shards

Each shard is a cluster of replicas that can be faulty.

Consensus for each *computation* within shards.

Cluster-sending for any *communication* between shards.

Consensus is costly: Minimize its use.

# The Orchestrate-Execute Model for Multi-Shard Transactions

Consider a multi-shard transaction  $\tau$ :

- ▶ Processing is broken down into three types of *shard-steps*: vote, commit, and abort.
- ▶ Each shard-step is performed via *one* consensus step.
- ▶ Transfer control between steps using *cluster-sending*.

**Execution method** determines the local operations of a shard-step:  
*locks, checking conditions, updating state, ....*

**Orchestration method** determines how *control is transferred* between shard-steps:  
perform *votes*, collect *votes*, decide *commit* or *abort*  $\tau$ .

## Example of the Orchestrate-Execute Model

Shard accounts by first letter of name

$\tau = \text{"if } Ana \text{ has \$500 and } Bo \text{ has \$200, then}$   
 $\text{move \$400 from } Ana \text{ to } Bo."}$

## Example of the Orchestrate-Execute Model

Shard accounts by first letter of name

$\tau = \text{"if } \textit{Ana} \text{ has \$500 and } \textit{Bo} \text{ has \$200, then}$   
 $\text{move \$400 from } \textit{Ana} \text{ to } \textit{Bo}."$

$\sigma_1 = \text{"Lock}_\tau(\textit{Ana}); \text{ if } \textit{Ana} \text{ has \$500, then forward } \sigma_2 \text{ to } \mathcal{C}_b \text{ (commit vote)}$   
 $\text{else Release}_\tau(\textit{Ana}) \text{ (abort vote).}"$

**vote-step**

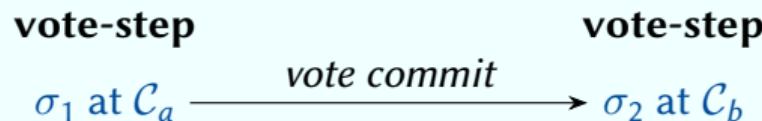
$\sigma_1$  at  $\mathcal{C}_a$

## Example of the Orchestrate-Execute Model

Shard accounts by first letter of name

$\tau$  = “if *Ana* has \$500 and *Bo* has \$200, then  
move \$400 from *Ana* to *Bo*.”

$\sigma_2$  = “Lock $_{\tau}$ (*Bo*); if *Bo* has \$200, then add \$400 to *Bo*; Release $_{\tau}$ (*Bo*); and  
forward  $\sigma_3$  to  $\mathcal{C}_a$  (commit)  
else Release $_{\tau}$ (*Bo*) and forward  $\sigma_4$  to  $\mathcal{C}_a$  (abort).”



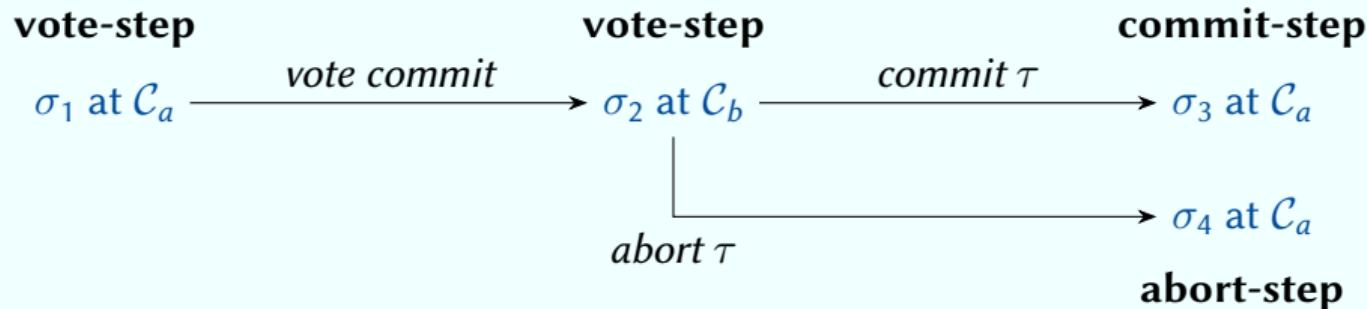
## Example of the Orchestrate-Execute Model

Shard accounts by first letter of name

$\tau = \text{"if } Ana \text{ has \$500 and } Bo \text{ has \$200, then}$   
 $\text{move \$400 from } Ana \text{ to } Bo."$

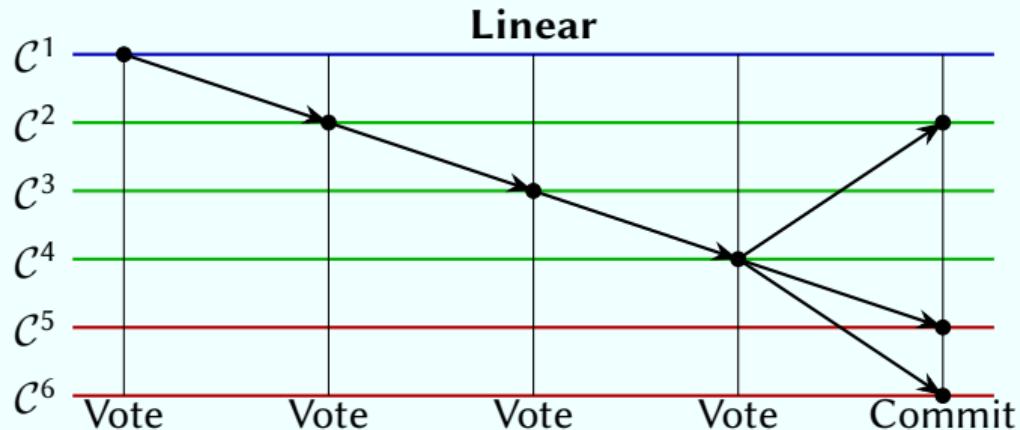
$\sigma_3 = \text{"remove \$400 from } Ana \text{ and Release}_\tau(Ana)."$

$\sigma_4 = \text{"Release}_\tau(Ana).$ "



# Resilient Orchestration Methods

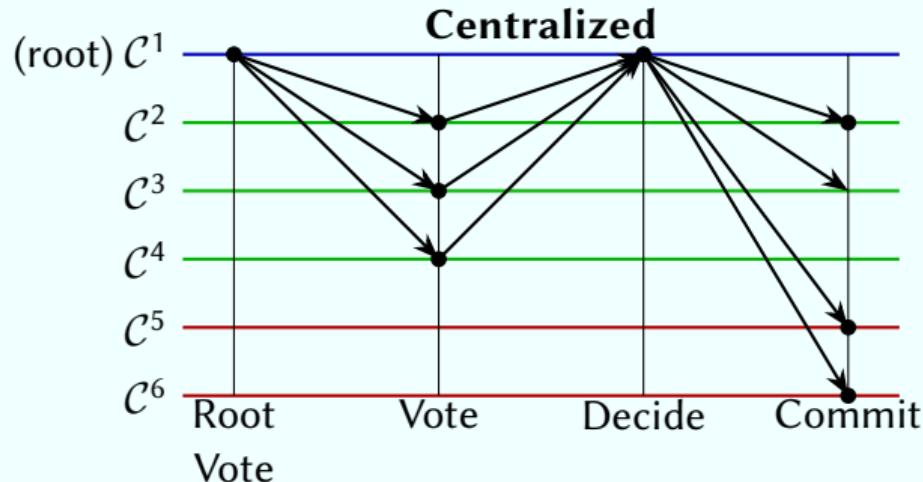
Orchestration  $\approx$  two-phase commit, except that *shards never fail*.



Vote-steps in *sequence*, decide *centralized*, commit or abort in *parallel*.

# Resilient Orchestration Methods

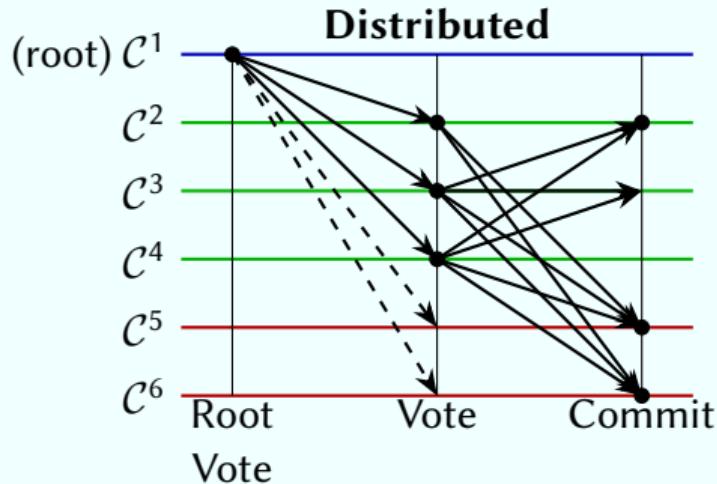
Orchestration  $\approx$  two-phase commit, except that *shards never fail*.



Vote-steps in *parallel*, decide *centralized*, commit or abort in *parallel*.

# Resilient Orchestration Methods

Orchestration  $\approx$  two-phase commit, except that *shards never fail*.



Vote-steps in *parallel*, decide *decentralized*, commit or abort in *parallel*.

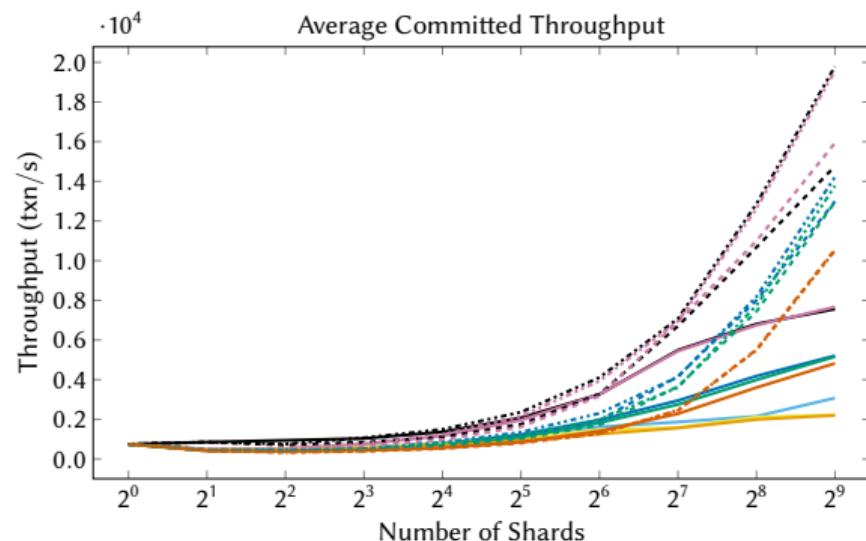
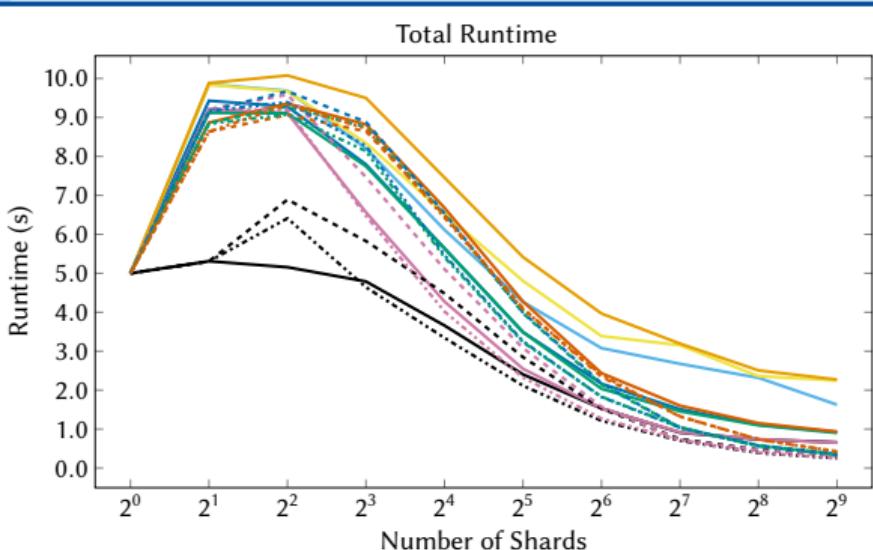
# Resilient Execution Methods

Execution updates state and performs *concurrency control*.

- ▶ Write uncommitted execution for *free*:  
Due to consensus, shard-steps are performed in sequence on that shard.
- ▶ Higher isolation levels via *two-phase locking*:
  - ▶ read uncommitted execution: only *write locks*;
  - ▶ read committed execution: *read locks* during steps;
  - ▶ serializable execution: *read and write locks*.
- ▶ Blocking locks (with linear orchestration) versus non-blocking locks.

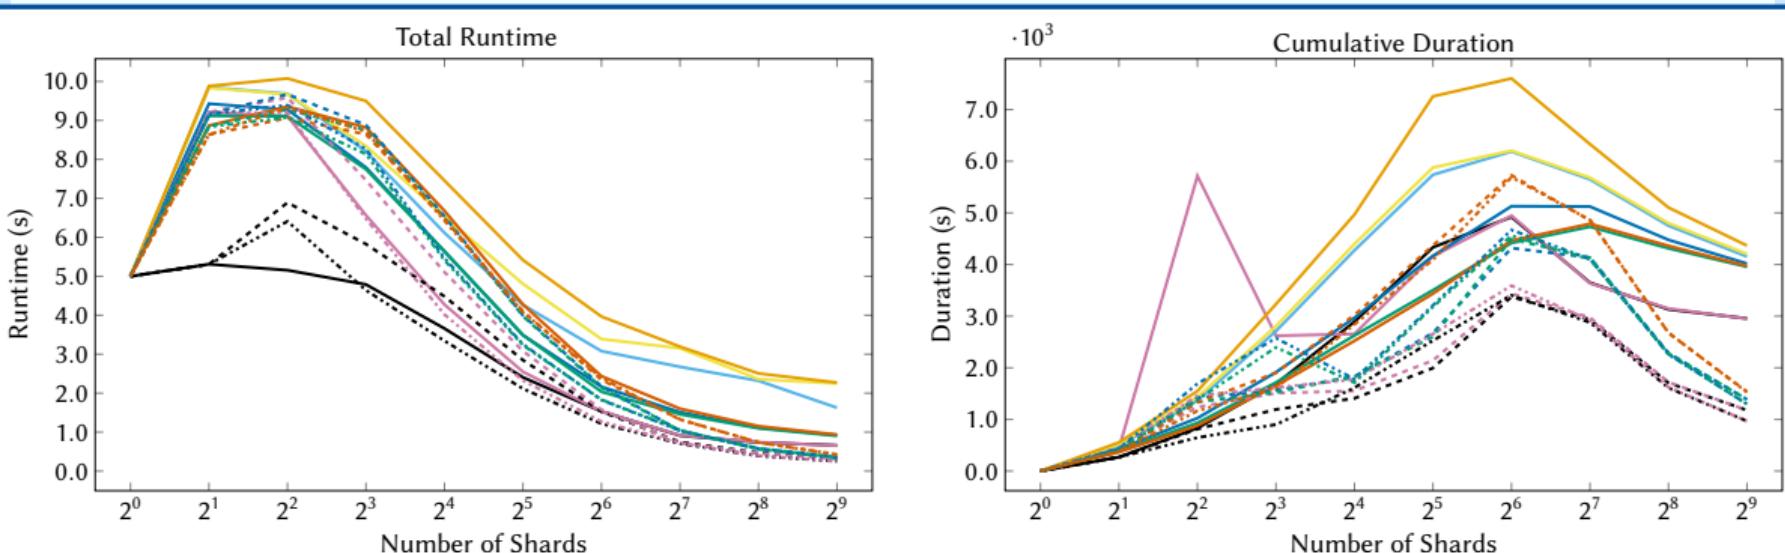
# Evaluation

Isolation-Free execution (write uncommitted)				Lock-based execution			
	unsafe	safe	blocking	Read Uncommitted	non-blocking	blocking	non-blocking
Linear	— LIFu	— LIFs	— LRUb	— LRUnb	— LRCb	— LRCnb	— LSb
Centralized	- - - CIFu	- - - CIFs	- - - CRUnb	- - - CRCnb	- - - DRCnb	- - - CSnb	- - - DSnb
Distributed	--- DIFu	--- DIFs	--- DRUnb				



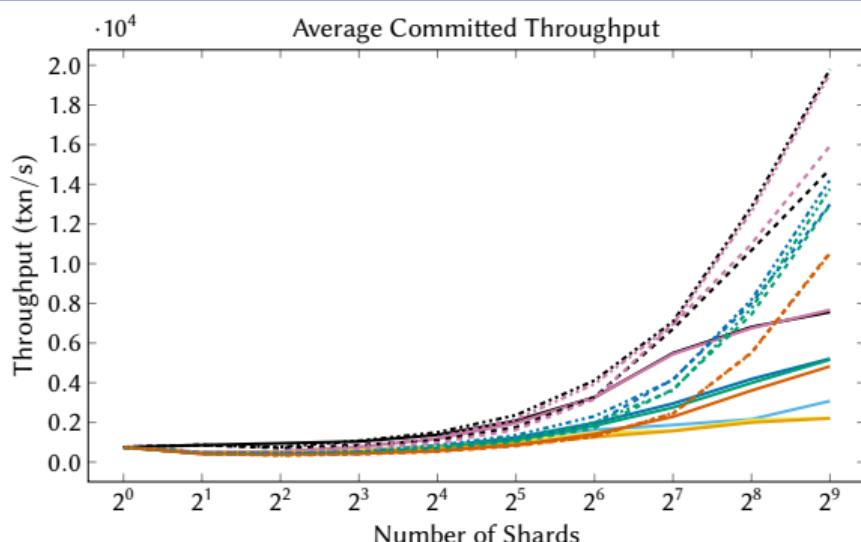
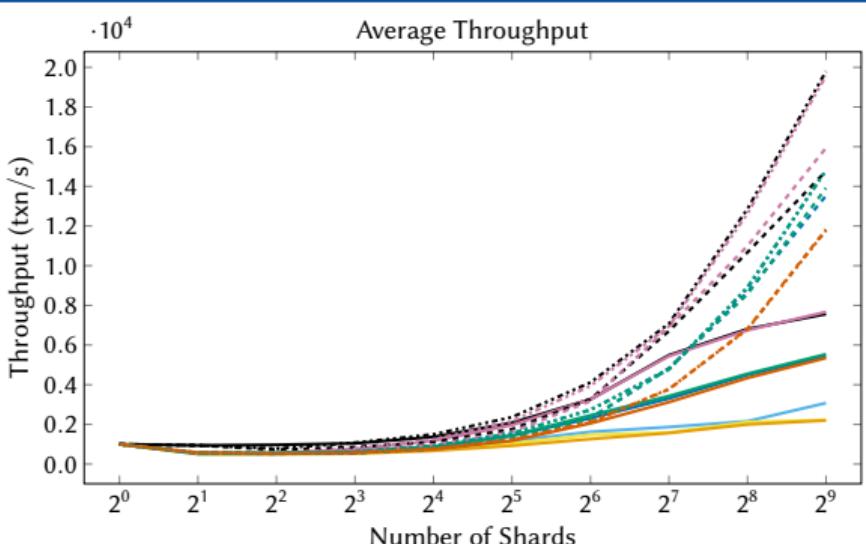
# Evaluation

Isolation-Free execution (write uncommitted)		Lock-based execution							
		Read Uncommitted		Read Committed		Serializable			
		<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>		
Linear	LIFu	LIFs	LRUb	LRUnb	LRCb	LCnb	LSb	LSnb	
Centralized	CIFu	CIFs	CRUnb	CRCnb	DRCnb	CSnb	DSnb		
Distributed	DIFu	DIFs	DRUnb						



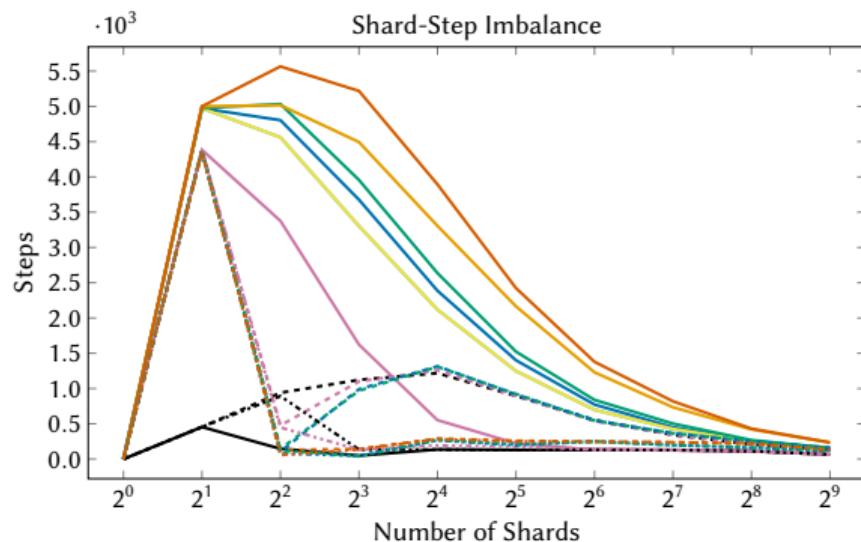
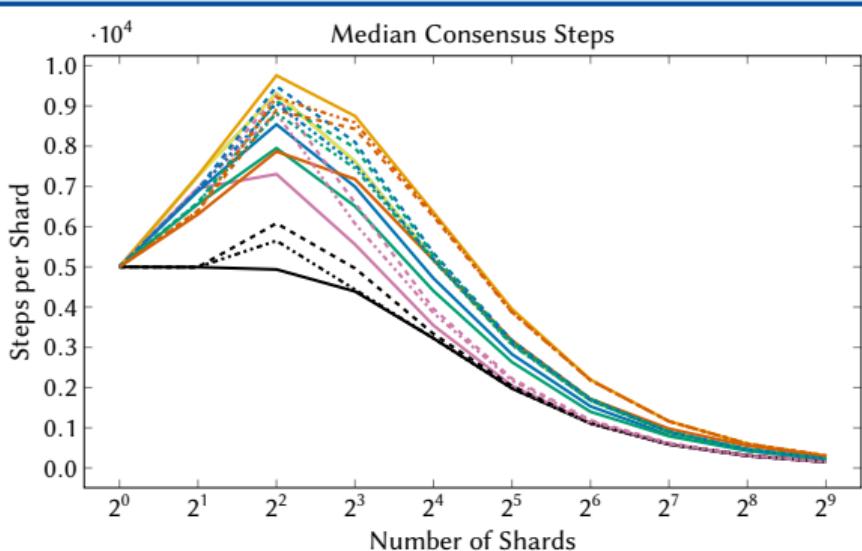
# Evaluation

Isolation-Free execution (write uncommitted)		Lock-based execution					
		Read Uncommitted		Read Committed		Serializable	
		<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>
Linear	LIFu	LIFs	LRUb	LRUnb	LCRcb	LCRnb	LSb
Centralized	CIFu	CIFs	CRUnb	CRCnb	CRCnb	CSnb	CSnb
Distributed	DIFu	DIFs	DRUnb	DRCnb	DRCnb	DSnb	DSnb



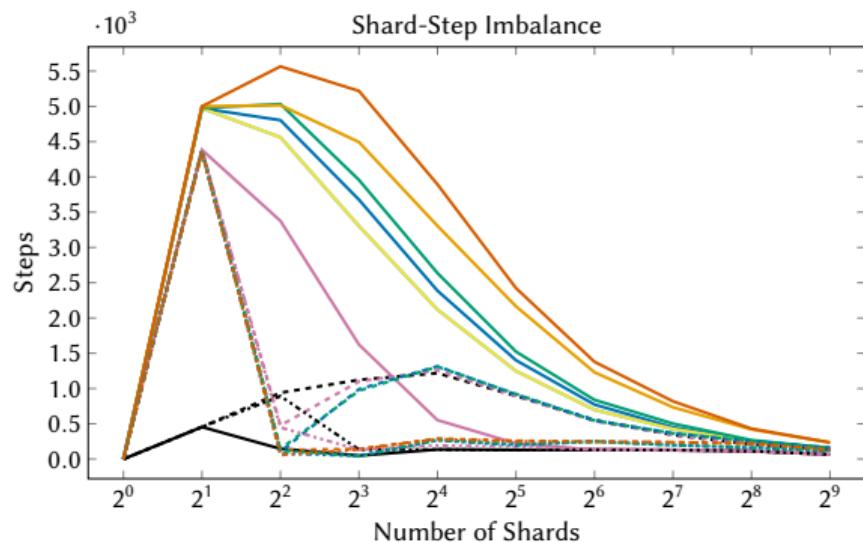
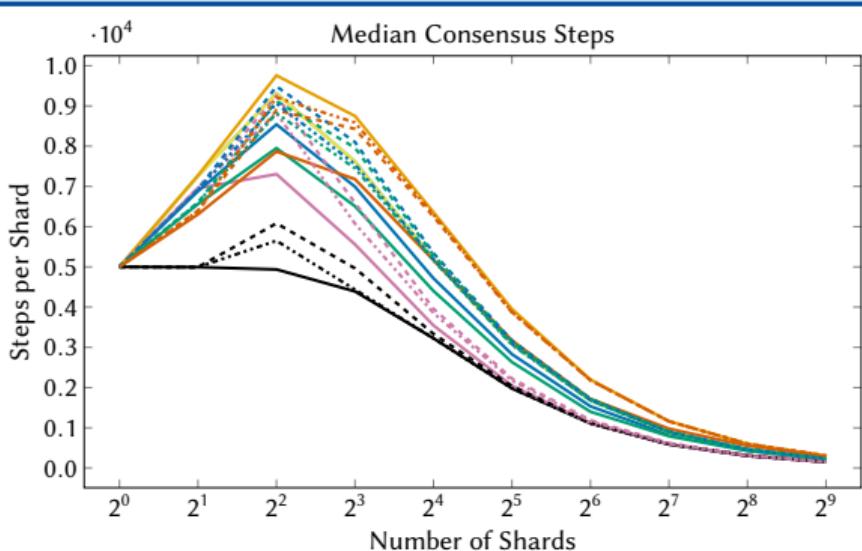
# Evaluation

Isolation-Free execution (write uncommitted)		Lock-based execution					
		Read Uncommitted		Read Committed		Serializable	
		<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>
Linear	LIFu	LIFs	LRUb	LRUnb	LCRcb	LCRnb	LSb
Centralized	CIFu	CIFs	CRUnb	CRCnb	CDRnb	CSnb	DSnb
Distributed	DIFu	DIFs	DRUnb				



# Evaluation

Isolation-Free execution (write uncommitted)		Lock-based execution					
		Read Uncommitted		Read Committed		Serializable	
		<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>
Linear	LIFu	LIFs	LRUb	LRUnb	LCRcb	LCRnb	LSb
Centralized	CIFu	CIFs	CRUnb	CRCnb	CRDnb	CSnb	CDnb
Distributed	DIFu	DIFs	DRUnb	DRCnb	DRDnb	DSnb	DDnb



# Evaluation

Isolation-Free execution (write uncommitted)		Lock-based execution					
		Read Uncommitted		Read Committed		Serializable	
		<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>
Linear	LIFu	LIFs	LRUb	LRUnb	LCRcb	LCRnb	LSb
Centralized	CIFu	CIFs	CRUnb		CRCnb		CSnb
Distributed	DIFu	DIFs	DRUnb		DRCnb		DSnb

