

# An In-Depth Look of BFT Consensus in Blockchain: Challenges and Opportunities



Suyash Gupta



Jelle Hellings



Sajjad Rahnama




Mohammad Sadoghi

MokaBlox LLC  
Exploratory Systems Lab  
University of California Davis




# What is Blockchain?

- A linked list of blocks.
- Each block contains hash of the previous block.
- A block contains information about some client transactions.



# Why Blockchain?




- By User:Pedant, User:Wapcaplet, User:Antonu, User:Vanderlindenma, User:.js. - Composition of File:Barnstar of Diligence Hires.png + File:Voting hand.svg., CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=45960536>


- <https://blog.devolutions.net/2017/10/whats-the-difference-between-2fa-and-mfa>

# Components of a Blockchain System

- Replicas → Store all the data.
- Client → Sends transactions to process.
- Consensus Protocol → Helps ordering transactions.
- Cryptographic Constructs → Authenticate replicas and clients.
- Ledger → Records transactions.



# Consensus



# Types of Blockchain Systems

- **Permissionless** → **Open Access**
  - Anyone can participate.
  - Identities of the replicas are unknown.
  - Applications include crypto-currency and money exchange.
- **Permissioned** → **Restricted Access**
  - Only a select group of replicas, although untrusted can participate.
  - Identities of the replica are known a priori.
  - Applications include health-care and energy trading.


# BITCOIN

- First Crypto-currency → a monetary application.
- Uses Nakamoto consensus → Proof-of-Work beneath the skin.
- Supports permissionless access.
- Requires solving hard cryptographic puzzles.
- Any replica that wants to create a new block proves that it did solve the puzzle.
- Difficulty of the puzzle helps prevent malicious attacks.



# ETHEREUM

- Another Crypto-currency → a token used in variety of applications.
- Uses Proof-of-Work but plans to start using Proof-of-Stake.
- Supports permissionless access.
- Allows programmers to design their transactions or “*smart contracts*”.
- Hard dependency on Ethereum Virtual machine (EVM).
- Envisions design of Permissioned applications.



# Terrorists Turn to Bitcoin for Funding, and They're Learning Fast



Login

Search Q

Disrupt Berlin 2019

THE VERGE

TECH ▾ REVIEWS ▾ SCIENCE ▾ CREATORS ▾ ENTERTAINMENT ▾ VIDEO FEATU

REPORT TECH CYBERSECURITY

## Why the Ethereum Classic hack is a bad omen for the blockchain


The 51 percent attack is real, and it's easier than ever

By Russell Brandom | Jan 9, 2019, 8:47am EST



BTC	ETH	XRP	BCH	LTC
\$7,424 -1.98%	\$152 +0.03%	\$0.22 -0.75%	\$216 -1.00%	\$47.98 +1.24%

News ▾ Features ▾ Price Analysis ▾ Market Tools ▾ Cryptopedia ▾ Industry ▾



By William Suberg

NOV 16, 2019

## Bitcoin Cash Hard Fork Sees Miners 'Waste' Money on 14 Invalid Blocks



Bitcoin 24h  
\$7,355.40 -2.58%

Ethereum 24h  
\$151.13 -0.62%

XRP 24h  
\$0.223687 -


Story from Tech →

## Bitcoin Cash Miners Undo Attacker's Transactions With '51% Attack'

May 24, 2019 at 21:17 UTC • Updated May 25, 2019 at 10:39 UTC



# Permissioned Blockchain Systems

- Require identities of the participating replicas to be known a priori.
- Replicas still untrusted → Consensus through traditional BFT protocols.
- Computationally in-expensive.
- More reliance on *communication* primitives.
- Prevent chain forks.
- Suitable for needs of an industry → JP Morgan, IBM, Oracle
- Advent path for *Blockchain Databases*.





# Transactions, Agreement and Consensus




# The Omniscient Transaction

- A transformation from a *consistent* state to another consistent state.
- A *contract* between two or more parties.
- A collection of *Read* or *Write* operations.
- Types of transactions: nested, compensating, multi-operation etc.

# ACID Properties

- **Atomicity:** A transaction either completes fully or none of its changes take place.
- **Consistency:** The transaction must obey legal protocols
- **Isolation:** The intermediate state of a transaction is invisible to other transactions
- Durability: Once a transaction is committed, it cannot be abrogated



# Consistency vs Availability

- An ongoing struggle that causes *performance tradeoffs*.
- Availability → Database needs to be always available for use.
  - Solution? Replication
  - Issues? Faults, Failures and Attacks.
- Consistency → Database needs to be correct.
  - Solution? All replicas should have same state.
  - Issues? Expensive.



# A Deep Dive into BFT Consensus




**UCDAVIS**  
UNIVERSITY OF CALIFORNIA




# PBFT: Practical Byzantine Fault Tolerance

- First practical Byzantine Fault Tolerant Protocol.
- Tolerates up to  $f$  failure out of  $3f+1$  replicas
- Three phases of which two require quadratic communication complexity.
- Safety is always guaranteed and Liveness is guaranteed in periods of partial synchrony.
- View-Change protocol for replacing malicious primary

# PBFT Failure-Free Flow



# PBFT Primary Failure (View Change)



# Requirements of Existing BFT Protocols

- 1) Require three phases of communication, of which two necessitate quadratic communication (PBFT).
- 2) Expect no failures or dependence on clients (Zyzzyva).
- 3) Incur high client latencies due to many phases of communication (PBFT, HotStuff).
- 4) Require threshold signatures, which are computationally expensive (HotStuff).
- 5) Require more than  $3f+1$  replicas (Q/U, HQ).
- 6) Need trusted components (AHL, Attested Append-only memory).


# Proof-of-Execution (PoE): Reaching Consensus through Fault-Tolerant Speculation

- *Speculative Execution* to reduce the client latency.
- *Out-of-Order message processing* for transactions.
- *Three Linear Phases*.
- *No Dependence* on Clients or requirement of expensive cryptographic primitives.
- *No Requirement* of a *Twin-Path* protocol.


# PoE vs Other Protocols

Protocol	Phases	Messages	Resilience	Requirements
ZYZZYVA	1	$\mathcal{O}(n)$	0	reliable clients and unsafe
PoE (our paper)	3	$\mathcal{O}(3n)$	f	sign. agnostic
PBFT	3	$\mathcal{O}(n + 2n^2)$	f	
HOTSTUFF	4	$\mathcal{O}(n + 3n^2)$	f	
HOTSTUFF-TS	8	$\mathcal{O}(8n)$	f	threshold sign.
SBFT	5	$\mathcal{O}(5n)$	0	threshold sign. and twin path


# PoE Failure-Free Flow



# PoE View Change Protocol



# PoE Scalability under Single Failure





# Scaling Blockchain Databases through Parallel Resilient Consensus Paradigm

- Why should BFT protocols rely on just *one* primary replica?
- Malicious primary can *throttle* the system throughput.
- Malicious primary requires *replacement* → fall in throughput.

# Multiple Byzantine Fault-Tolerance (MultiBFT) Paradigm

- Designate multiple replicas as Primaries!
- Run multiple parallel consensuses on each replica.







MultiBFT with 2 parallel instances on each replica

# Malicious Primaries Collusion

- Multiple malicious primaries can prevent liveness!
- Solution → Optimistic Recovery through State Exchange.




# MultiBFT Scalability



# Global Scale Resilient Blockchain Fabric\*


- Traditional BFT protocols do not scale to geographically large distances.
- Blockchain requires decentralization → replicas can be far apart → expensive communication!
- The underlying BFT consensus protocol should be topology-aware.


# Vision Geo-Scale Byzantine Fault-Tolerance



# GeoBFT Protocol

GeoBFT is a topology-aware protocol, which groups replicas into clusters. Each cluster runs the PBFT consensus protocol, in parallel and independently.






# GeoBFT Takeaways

- To ensure common ordering → linear communication among the clusters is required.
- Primary replica at each cluster sends a secure certificate to  $f+1$  replicas of every other cluster.
- Certificates guarantee common order for execution.
- If primary sends invalid certificates → will be detected as malicious.

# GeoBFT Scalability



# ResilientDB: High Throughput Yielding, Scalable Permissioned Blockchain Fabric

Visit at: <https://resilientdb.com/>




# Why Should You Choose ResilientDB?

- 1) Bitcoin and Ethereum offer low throughputs of *10 txns/s*.
- 2) Existing Permissioned Blockchain Databases still have low throughputs (*20K txns/s*).
- 3) Prior works blame BFT consensus as *expensive*.
- 4) System Design is mostly *overlooked*.
- 5) ResilientDB adopts *well-researched* database and system practices.


# Dissecting Existing Permissioned Blockchains

- 1) Single-threaded Monolithic Design
- 2) Successive Phases of Consensus
- 3) Integrated Ordering and Execution
- 4) Strict Ordering
- 5) Off-Chain Memory Management
- 6) Expensive Cryptographic Practices


# Can a well-crafted system based on a classical BFT protocol outperform a modern protocol?




# ResilientDB Architecture



# ResilientDB Multi-Threaded Deep Pipeline




# Insight 1: Multi-Threaded pipeline Gains




Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

# Insight 1: Multi-Threaded pipeline Gains




Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

# Insight 1: Multi-Threaded pipeline Gains





Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

# Insight 1: Multi-Threaded pipeline Gains





Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

# Insight 2: Optimal Batching Gains





More transactions batched together → increase in throughput  
→ reduced phases of consensus.

# Insight 3: Memory Storage Gains



In-memory blockchain storage → reduces access cost.

# Insight 4: Number of Clients



Too many clients → increases average latency.

# ResilientDB: Hands On

Visit at: <https://github.com/resilientdb/resilientdb>



# How to Run ResilientDB?

- Go to <https://github.com/resilientdb/resilientdb> and Fork it!
- Install Docker-CE and Docker-Compose (Links on git)
- Use the Script "*resilientDB-docker*" as following:

```
./resilientDB-docker --clients=1 --replicas=4
```

```
./resilientDB-docker -d [default 4 replicas and 1 client]
```

- Result will be printed on STDOUT and stored in *res.out* file.

# How to Run ResilientDB?

resilientdb / resilientdb

Code Issues 1 Pull requests 0 Actions Projects 0 Wiki Security Insights

Watch 5 Unstar 11 Fork 13

ResilientDB: A scalable permissioned blockchain fabric

46 commits 1 branch 0 packages 2 releases 4 contributors MIT

Branch: master New pull request Create new file Upload files Find file Clone or download

gupta-suyash readme updated Latest commit f2302e6 3 days ago

benchmarks	Initial Commit	16 days ago
blockchain	ledger archiecture defined	4 days ago
client	Initial Commit	16 days ago
deps	Initial Commit	16 days ago
scripts	added -e to handle multiple clients in docker-ifconfig	13 days ago
statistics	Initial Commit	16 days ago
system	ledger archiecture defined	4 days ago
transport	Initial Commit	16 days ago
.gitignore	Initial Commit	16 days ago
CHANGELOG.md	changelog added	3 days ago
CODE_OF_CONDUCT.md	Create CODE_OF_CONDUCT.md	15 days ago
LICENSE.md	Initial Commit	16 days ago
Makefile	Initial Commit	16 days ago
README.md	readme updated	3 days ago
config.cpp	Initial Commit	16 days ago
config.h	ledger archiecture defined	4 days ago
resilientDB-docker	Initial Commit	16 days ago

# Docker CE

## What is Docker?

*an open-source project that automates the deployment of software applications inside **containers** by providing an additional layer of abstraction and automation of OS-level virtualization on Linux.*

- Run a distributed program on one machine
- Simulate with lightweight virtual machines

# Docker CE

## What is Docker?

*an open-source project that automates the deployment of software applications inside **containers** by providing an additional layer of abstraction and automation of OS-level virtualization on Linux.*

- Run a distributed program on one machine
- Simulate with lightweight virtual machines


# Resilient DB

```
./resilientDB-docker -d
```

- Remove old Containers
- Create new Containers
- Create IP address settings
- Install dependencies
- Compile Code
- Run binary files
- Gather the results

```
sajjad@sajjad-xps:~/WS/expo/resilientdb|master⚡
> ./resilientDB-docker -d
Number of Replicas: 4
Number of Clients: 1
Stopping previous containers...
Stopping s3 ... done
Stopping s1 ... done
Stopping s4 ... done
Stopping c1 ... done
Stopping s2 ... done
Removing s3 ... done
Removing s1 ... done
Removing s4 ... done
Removing c1 ... done
Removing s2 ... done
Removing network resilientdb_default
Successfully stopped
Creating docker compose file ...
Docker compose file created --> docker-compose.yml  Send problem report to the developers?
Starting the containers...
Creating network "resilientdb_default" with the default driver
Creating s4 ... done
Creating c1 ... done
Creating s1 ... done
Creating s2 ... done
Creating s3 ... done
ifconfig file exists... Deleting File
Deleted
Server sequence --> IP
c1 --> 172.21.0.3
s1 --> 172.21.0.4
s2 --> 172.21.0.6
s3 --> 172.21.0.2
s4 --> 172.21.0.5
Put Client IP at the bottom
ifconfig.txt Created!

Checking Dependencies...
Installing dependencies..
/home/sajjad/WS/expo/resilientdb
Dependencies has been installed
```



# Resilient DB

- Throughput
  - Transaction per second
- Average Latency
  - The from client request to client reply
- Working Thread idleness
  - The time that thread is waiting
- WT0: Consensus Messages
- WT1 and WT2: Batch Threads
- WT3: checkpointing Thread
- WT4: Execute Theread

```
Throughputs:  
0: 38525  
1: 38530  
2: 38558  
3: 38551  
4: 38564  
Latencies:  
latency 4: 0.505870  
  
idle times:  
Idleness of node: 0  
Worker THD 0: 116.227  
Worker THD 1: 62.0772  
Worker THD 2: 62.2130  
Worker THD 3: 105.098  
Worker THD 4: 74.9193  
Idleness of node: 1  
Worker THD 0: 39.3157  
Worker THD 1: 0.00000  
Worker THD 2: 0.00000  
Worker THD 3: 104.700  
Worker THD 4: 74.8603  
Idleness of node: 2  
Worker THD 0: 35.0847  
Worker THD 1: 0.00000  
Worker THD 2: 0.00000  
Worker THD 3: 102.415  
Worker THD 4: 78.1078  
Idleness of node: 3  
Worker THD 0: 38.4452  
Worker THD 1: 0.00000  
Worker THD 2: 0.00000  
Worker THD 3: 107.512  
Worker THD 4: 77.6965  
Memory:  
0: 172 MB  
1: 156 MB  
2: 155 MB  
3: 156 MB  
4: 812 MB  
  
avg thp: 4: 38541  
avg lt : 1: .505  
Code Ran successfully ---> res.out
```

# PBFT: Practical Byzantine Fault Tolerance

## Client Request

- Client/client\_main.cpp
- System/client\_thread.cpp
- ClientQueryBatch Class
- Process ClientBatch in primary

```
C++ client_main.cpp ×
client > C++ client_main.cpp > ...
31 int main(int argc, char *argv[])
32 {
33     printf("Running client...\n\n");
34     // 0. initialize global data structure
35     parser(argc, argv);
36     assert(g_node_id >= g_node_cnt);
37     uint64_t seed = get_sys_clock();
38     srand(seed);
39     printf("Random seed: %ld\n", seed);
40
41     int64_t starttime;
42     int64_t endtime;
43     starttime = get_server_clock();
44     // per-partition malloc
45     printf("Initializing stats... ");
46     fflush(stdout);
47     stats.init(g_total_client_thread_cnt);
48     printf("Done\n");
49     printf("Initializing transport manager... ");
50     fflush(stdout);
51     tport_man.init();
52     printf("Done\n");
53     printf("Initializing client manager... ");
54     Workload *m_wl = new YCSBWorkload;
55     m_wl->Workload::init();

C++ client_thread.cpp ×
system > C++ client_thread.cpp > ...
79
80     RC ClientThread::run()
81     {
82
83         tsetup();
84         printf("Running ClientThread %ld\n", _thd_id);
85
86         while (true)
87         {
88             keyMTX.lock();
89             if (keyAvail)
90             {
91                 keyMTX.unlock();
92                 break;
93             }
94             keyMTX.unlock();
95         }
96
97         BaseQuery *m_query;
98         uint64_t iters = 0;
99         uint32_t num_txns_sent = 0;
100        int txns_sent[g_node_cnt];
101        for (uint32_t i = 0; i < g_node_cnt; ++i)
102            txns_sent[i] = 0;
103
104        run_starttime = get_sys_clock();
```

# PBFT: Practical Byzantine Fault Tolerance

## Process Messages

- Transport/message.cpp
- System/worker\_thread.cpp
- System/worker\_thread\_pbft.cpp
- Worker Thread: Run function
- Worker Thread: Process function

```
C++ worker_thread.cpp ×
system > C++ worker_thread.cpp > WorkerThread::run()

626 /**
627  * Starting point for each worker thread.
628 *
629 * Each worker-thread created in the main() starts here. Each worker-thread is alive
630 * till the time simulation is not done, and continuously perform a set of tasks.
631 * These tasks involve, dequeuing a message from its queue and then processing it
632 * through call to the relevant function.
633 */
634 RC WorkerThread::run()
{
    tsetup();
    printf("Running WorkerThread %d\n", _thd_id);

    uint64_t agCount = 0, readystarttime, idlestarttime = 0;

    // Setting batch (only relevant for batching threads).
    next_set = 0;

    while (!simulation->is_done())
    {
        txn_man = NULL;
        heartbeat();
        progress_stats();

#if VIEW_CHANGES
        // Thread 0 continuously monitors the timer for each batch.
        if (get_thd_id() == 0)
        {
            check_for_timeout();
        }

        if (g_node_id != get_current_view(get_thd_id()))
        {
            check_switch_view();
        }
#endif
        // Dequeue a message from its work_queue.
        Message *msg = work_queue.dequeue(get_thd_id());
    }
}
```

```
C++ worker_thread.cpp ×
system > C++ worker_thread.cpp > WorkerThread::process(Message *)
87 void WorkerThread::process(Message *msg)
88 {
89     RC rc __attribute__((unused));
90
91     switch (msg->get_rtype())
92     {
93         case KEYEX:
94             rc = process_key_exchange(msg);
95             break;
96         case CL_BATCH:
97             rc = process_client_batch(msg);
98             break;
99         case BATCH_REQ:
100            rc = process_batch(msg);
101            break;
102        case PBFT_CHKPT_MSG:
103            rc = process_pbft_chkpt_msg(msg);
104            break;
105        case EXECUTE_MSG:
106            rc = process_execute_msg(msg);
107            break;
108 #if VIEW_CHANGES
109        case VIEW_CHANGE:
110            rc = process_view_change_msg(msg);
111            break;
112        case NEW_VIEW:
113            rc = process_new_view_msg(msg);
114            break;
115 #endif
116        case PBFT_PREP_MSG:
117            rc = process_pbft_prep_msg(msg);
118            break;
119        case PBFT_COMMIT_MSG:
120            rc = process_pbft_commit_msg(msg);
121            break;
122        default:
123            printf("Msg: %d\n", msg->get_rtype());
124            fflush(stdout);
125            assert(false);
126            break;
127     }
128 }
```

# PBFT: Practical Byzantine Fault Tolerance

## Process Client Message

- System/worker\_thread\_pbft.cpp
- process\_client\_batch Function
- Create and Send Batch Request
  - create\_and\_send\_batchreq Function
  - Create Transactions
  - Create Digest
- BatchRequest Class
  - Pre-Prepare Message

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > ...
18 /**
19  * Processes an incoming client batch and sends a Pre-prepare message to all replicas.
20  *
21  * This function assumes that a client sends a batch of transactions and
22  * for each transaction in the batch, a separate transaction manager is created.
23  * Next, this batch is forwarded to all the replicas as a BatchRequests Message
24  * which corresponds to the Pre-Prepare stage in the PBFT protocol.
25  *
26  * @param msg Batch of Transactions of type ClientQueryBatch from the client.
27  * @return RC
28 */
29 RC WorkerThread::process_client_batch(Message *msg)
30 {
31     //printf("ClientQueryBatch: %ld, THD: %ld :: CL: %ld :: RQ: %ld\n",msg->txns->size(), msg->txns->begin()->thd_id, msg->txns->begin()->cl_id, msg->txns->begin()->rq_id);
32     //fflush(stdout);
33
34     ClientQueryBatch *clbtch = (ClientQueryBatch *)msg;
35
36     // Authenticate the client signature.
37     validate_msg(clbtch);
38
39 #if VIEW_CHANGES
40     // If message forwarded to the non-primary.
41     if (g_node_id != get_current_view(get_thd_id()))
42     {
43         client_query_check(clbtch);
44         return RCOK;
45     }
46
47     // Partial failure of Primary 0.
48     fail_primary(msg, 9);
49 #endif
50
51     // Initialize all transaction managers and uint64_t Message::txnid.
52     create_and_send_batchreq(clbtch, clbtch->txnid);
53
54     return RCOK;
55 }
```

```
C++ worker_thread.cpp ×
system > C++ worker_thread.cpp > WorkerThread::create_and_send_batchreq(ClientQueryBatch *, uint64_t)
1123 /**
1124  * This function is used by the primary replicas to create and set
1125  * transaction managers for each transaction part of the ClientQueryBatch message
1126  * by the client. Further, to ensure integrity a hash of the complete batch is
1127  * generated, which is also used in future communication.
1128 */
1129 * @param msg Batch of transactions as a ClientQueryBatch message.
1130 * @param tid Identifier for the first transaction of the batch.
1131 */
1132 void WorkerThread::create_and_send_batchreq(ClientQueryBatch *msg, uint64_t tid)
1133 {
1134     // Creating a new BatchRequests Message.
1135     Message *bmsg = Message::create_message(BATCH_REQ);
1136     BatchRequests *breq = (BatchRequests *)bmsg;
1137     breq->init(get_thd_id());
1138
1139     // Starting index for this batch of transactions.
1140     next_set = tid;
1141
1142     // String of transactions in a batch to generate hash.
1143     string batchStr;
1144
1145     // Allocate transaction manager for all the requests in batch.
1146     for (uint64_t i = 0; i < get_batch_size(); i++)
1147     {
1148         uint64_t txn_id = get_next_txnid() + i;
1149
1150         //cout << "Txn: " << txn_id << " :: Thd: " << get_thd_id() << "\n";
1151         //fflush(stdout);
1152         txn_man = get_transaction_manager(txn_id, 0);
1153
1154         // Unset this txn man so that no other thread can concurrently use.
1155         while (true)
1156         {
1157             bool ready = txn_man->unset_ready();
1158             if (!ready)
1159             {
1160                 continue;
1161             }
1162             else
1163             {
1164                 break;
1165             }
1166         }
1167     }
1168 }
```

# PBFT: Practical Byzantine Fault Tolerance

## Process Batch Request (Prepare)

- System/worker\_thread\_pbft.cpp
- process\_batch Function
- Create and Send Prepare Message
  - Create Transactions
  - Save Digest
- PBFTPrepare Class
  - Prepare Message

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > WorkerThread::process_batch(Message *)
57  /**
58   * Process incoming BatchRequests message from the Primary.
59   *
60   * This function is used by the non-primary or backup replicas to process an incoming
61   * BatchRequests message sent by the primary replica. This processing would require
62   * sending messages of type PBFTPrepMessage, which correspond to the Prepare phase of
63   * the PBFT protocol. Due to network delays, it is possible that a replica may have
64   * received some messages of type PBFTPrepMessage and PBFTCommitMessage, prior to
65   * receiving this BatchRequests message.
66   *
67   * @param msg Batch of Transactions of type BatchRequests from the primary.
68   * @return RC
69   */
70 RC WorkerThread::process_batch(Message *msg)
71 {
72     uint64_t cntime = get_sys_clock();
73
74     BatchRequests *breq = (BatchRequests *)msg;
75
76     //printf("BatchRequests: TID:%ld : VIEW: %ld : THD: %ld\n", breq->txn_id, breq->view, get_
77     //fflush(stdout);
78
79     // Assert that only a non-primary replica has received this message.
80     assert(g_node_id != get_current_view(get_thd_id()));
81
82     // Check if the message is valid.
83     validate_msg(breq);
84 }
```

# PBFT: Practical Byzantine Fault Tolerance

## Process Prepare and Commit Messages(Prepare)

- System/worker\_thread\_pbft.cpp
- process\_pbft\_prepare Function
  - Count Prepare Messages
  - Create and Send commit Message
  - PBFTCommit Message
- process\_pbft\_commit Function
  - Count commit messages
  - Create and Send execute Message
  - ExecuteMessage Class

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > ...
200
186 /**
187 * Processes incoming Prepare message.
188 *
189 * This function processes incoming messages of type PBFTPrepMessage. If
190 * received 2f identical Prepare messages from distinct replicas, then it c
191 * and sends a PBFTCommitMessage to all the other replicas.
192 *
193 * @param msg Prepare message of type PBFTPrepMessage from a replica.
194 * @return RC
195 */
196 RC WorkerThread::process_pbft_prep_msg(Message *msg)
197 {
198     //cout << "PBFTPrepMessage: TID: " << msg->txnid << " FROM: " << msg->
199     //fflush(stdout);
200
201     // Start the counter for prepare phase.
202     if (txn_man->prep_rsp_cnt == 2 * g_min_invalid_nodes)
203     {
204         txn_man->txn_stats.time_start_prepare = get_sys_clock();
205     }
206
207     // Check if the incoming message is valid.
208     PBFTPrepMessage *pmsg = (PBFTPrepMessage *)msg;
209     validate_msg(pmsg);
210
211     // Check if sufficient number of Prepare messages have arrived.
212     if (prepared(pmsg))
213     {
214         // Send Commit messages.
215         txn_man->send_pbft_commit_msgs();
216
217         // End the prepare counter.
218         INC_STATS(get_thd_id(), time_prepare, get_sys_clock() - txn_man->tx
219     }
220
221     return RCOK;
222 }
```

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > ⚡ WorkerThread::process_pbft_commit_msg(Message *)
200
275 /**
276 * Processes incoming Commit message.
277 *
278 * This function processes incoming messages of type PBFTCommitMessage
279 * received 2f+1 identical Commit messages from distinct replicas, then :
280 * execute-thread to execute all the transactions in this batch.
281 *
282 * @param msg Commit message of type PBFTCommitMessage from a replica.
283 * @return RC
284 */
285 RC WorkerThread::process_pbft_commit_msg(Message *msg)
286 {
287     //cout << "PBFTCommitMessage: TID: " << msg->txnid << " FROM: " << msg->
288     //fflush(stdout);
289
290     if (txn_man->commit_rsp_cnt == 2 * g_min_invalid_nodes + 1)
291     {
292         txn_man->txn_stats.time_start_commit = get_sys_clock();
293     }
294
295     // Check if message is valid.
296     PBFTCommitMessage *pcmsg = (PBFTCommitMessage *)msg;
297     validate_msg(pcmsg);
298
299     txn_man->add_commit_msg(pcmsg);
300
301     // Check if sufficient number of Commit messages have arrived.
302     if (committed_local(pcmsg))
303     {
304 #if TIMER_ON
305         // End the timer for this client batch.
306         server_timer->endTimer(txn_man->hash);
307    #endif
308
309     // Add this message to execute thread's queue.
310     send_execute_msg();
311
312     INC_STATS(get_thd_id(), time_commit, get_sys_clock() - txn_man->tx
313 }
```

# PBFT: Practical Byzantine Fault Tolerance

## Process Execute Message

- System/worker\_thread.cpp
- Internal Message
- process\_execute Function
- Execute the Transactions in batch in order
- Create and send Client Response
- ClientResponse Class

```
C++ worker_thread.cpp X
system > C++ worker_thread.cpp > WorkerThread::process_execute_msg(Message *)
795
796 /**
797 * Execute transactions and send client response.
798 *
799 * This function is only accessed by the execute-thread, which executes the transactions
800 * in a batch, in order. Note that the execute-thread has several queues, and at any
801 * point of time, the execute-thread is aware of which is the next transaction to
802 * execute. Hence, it only loops on one specific queue.
803 *
804 * @param msg Execute message that notifies execution of a batch.
805 * @ret RC
806 */
807 RC WorkerThread::process_execute_msg(Message *msg)
808 {
809     //cout << "EXECUTE " << msg->txnid << " :: " << get_thd_id() <<"\n";
810     //fflush(stdout);
811
812     uint64_t ctime = get_sys_clock();
813
814     // This message uses txnid of index calling process_execute.
815     Message *rsp = Message::create_message(CL_RSP);
816     ClientResponseMessage *crsp = (ClientResponseMessage *)rsp;
817     crsp->init();
818
819     ExecuteMessage *emsg = (ExecuteMessage *)msg;
820
821     // Execute transactions in a shot
822     uint64_t i;
823     for (i = emsg->index; i < emsg->end_index - 4; i++)
824     {
825         //cout << "i: " << i << " :: next index: " << g_next_index << "\n";
826         //fflush(stdout);
827
828         TxnManager *tman = get_transaction_manager(i, 0);
829
830         inc_next_index();
831
832         // Execute the transaction
833         tman->run_txni();
834
835         // Commit the results.
836         tman->commit();
837
838         crsp->copy_from_txn(tman);
```

# PBFT: Practical Byzantine Fault Tolerance

## Work Queue

- Lock Free queues
- All the messages are being stored in these queues
- System/work\_queue.cpp
- Multiple queues for different Threads
- Dequeue and Enqueue Interfaces
- Enqueue in IOThread
- Dequeue in Worker Thread

```
C++ work_queue.cpp ×  
system > C++ work_queue.cpp > ...  
44     void QWorkQueue::enqueue(uint64_t thd_id, Message *msg, bool busy)  
45     {  
46         uint64_t starttime = get_sys_clock();  
47         assert(msg);  
48         DEBUG_M("QWorkQueue::enqueue work_queue_entry alloc\n");  
49         work_queue_entry *entry = (work_queue_entry *)mem_allocator.align_alloc(sizeof(work_queue_ent  
50             entry->msg = msg;  
51             entry->rtype = msg->rtype;  
52             entry->txn_id = msg->txn_id;  
53             entry->batch_id = msg->batch_id;  
54             entry->starttime = get_sys_clock();  
55             assert(ISSERVER || ISREPLICA);  
56             DEBUG("Work Enqueue (%ld,%ld) %d\n", entry->txn_id, entry->batch_id, entry->rtype);  
57  
58             if (msg->rtype == CL_QRY || msg->rtype == CL_BATCH)  
59             {  
60                 if (g_node_id == get_current_view(thd_id))  
61                 {  
62                     //cout << "Placing \n";  
63                     while (!new_txn_queue->push(entry) && !simulation->is_done())  
64                     {  
65                         //cout << "Pushing \n";  
66                     }  
67                 else  
68                 {  
69                     assert(entry->rtype < 100);  
70                     while (!work_queue[0]->push(entry) && !simulation->is_done())  
71                     {  
72                         //cout << "Pushing \n";  
73                     }  
74                 }  
75             }  
76         }  
77     }  
78 }
```

# PBFT: Practical Byzantine Fault Tolerance

## IO Thread and Transport Layer

- Multiple Input Threads
- Multiple Output Threads
- System/io\_thread.cpp
- Transport Layer: TCP Sockets
- Nano Message Library
- Transport/transport.cpp

```
C++ io_thread.cpp ×  
system > C++ io_thread.cpp > ...  
299     RC InputThread::server_recv_loop()  
300     {  
301  
302         myrand rdm;  
303         rdm.init(get_thd_id());  
304         RC rc = RCOK;  
305         assert(rc == RCOK);  
306         uint64_t starttime = 0;  
307         uint64_t idle_starttime = 0;  
308         std::vector<Message *> *msgs;  
309         while (!simulation->is_done())  
310         {  
311             heartbeat();  
312  
313             #if VIEW_CHANGES  
314                 if (g_node_id != get_current_view(get_thd_id()))  
315                 {  
316                     uint64_t tid = get_thd_id() - 1;  
317                     uint32_t nchange = get_newView(tid);  
318  
319                     if (nchange)  
320                     {  
321                         set_current_view(get_thd_id(), get_current_view(get_thd_id()) + 1);  
322                         set_newView(tid, false);  
323                     }  
324                 }  
325             #endif  
326  
327             msgs = tport_man.recv_msg(get_thd_id());  
328         }
```

# Configuration Parameters to Play

- NODE\_CNT  
Total number of replicas, minimum 4, that is,  $f=1$ .
- THREAD\_CNT  
Total number of threads at primary (at least 5)
- CLIENT\_NODE\_CNT  
Total number of clients (at least 1).
- MAX\_TXN\_IN\_FLIGHT  
Multiple of Batch Size
- DONE\_TIMER  
Amount of time to run the system.
- BATCH\_THREADS  
Number of threads at primary to batch client transactions.
- BATCH\_SIZE  
Number of transactions in a batch (at least 10)
- TXN\_PER\_CHKPT  
Frequency at which garbage collection is done.
- USE\_CRYPTO  
To switch on and off cryptographic signing of messages.
- CRYPTO\_METHOD\_ED25519  
To use ED25519 based digital signatures.
- CRYPTO\_METHOD\_CMAC\_AES  
To use CMAC + AES combination for authentication

# Thank You